

# Certified Mergeable Replicated Data Types

**“KC” Sivaramakrishnan**


joint work with

Vimala Soundarapandian, Adharsh Kamath and Kartik Nagar

IIT  
MADRAS  
SADAM









VISA®



ebay™




NETFLIX


HSBC ◀▶


YouTube






≠






≠



- Weak Consistency & Isolation

- Serializability
- Linearizability

# Even *simple* data structures attract enormous *complexity* when made *distributed*



Lindsey Kuper  
@lindsey

"Oh, you wanted to \*increment a counter\*?! Good luck with that!" -- the distributed systems literature

12:25 AM · Mar 10, 2015 · Twitter Web Client

---

375 Retweets 18 Quote Tweets 614 Likes

# Sequential Counter

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
end
```


# Sequential Counter

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
end
```


- Written in idiomatic style
- Composable

```
type counter_list = Counter.t list
```


# Replicated Counter




# Replicated Counter




# Replicated Counter




# Replicated Counter




# Replicated Counter




# Replicated Counter




# Replicated Counter



# Replicated Counter




# Replicated Counter



- **Idea:** Apply the local operations at all replicas

# Replicated Counter




- **Idea:** Apply the local operations at all replicas

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```


```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```



```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end


```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end


```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end


```



```


module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```



*Diverges*

```
module Counter : sig
  type t
  val read : t -> int
  val add   : t -> int ->
  val sub   : t -> int ->
  val mult  : t -> int ->
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```



## *Addition and multiplication do not commute*

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```

- **Idea:** Capture the effect of multiplication through the *commutative* addition operation


```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```

- **Idea:** Capture the effect of multiplication through the *commutative* addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```




- **Idea:** Capture the effect of multiplication through the *commutative* addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```




- **Idea:** Capture the effect of multiplication through the *commutative* addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```




- **Idea:** Capture the effect of multiplication through the *commutative* addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```




- **Idea:** Capture the effect of multiplication through the *commutative* addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```




- **Idea:** Capture the effect of multiplication through the *commutative* addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```



- **Idea:** Capture the effect of multiplication through the *commutative* addition operation
- *CRDTs*

# Convergent Replicated Data Types (CRDT)

# Convergent Replicated Data Types (CRDT)

- CRDT is guaranteed to ensure *strong eventual consistency (SEC)*
  - ★ G-counters, PN-counters, OR-Sets, Graphs, Ropes, docs, sheets
  - ★ Simple interface for the clients of CRDTs

# Convergent Replicated Data Types (CRDT)

- CRDT is guaranteed to ensure *strong eventual consistency (SEC)*
  - ★ G-counters, PN-counters, OR-Sets, Graphs, Ropes, docs, sheets
  - ★ Simple interface for the clients of CRDTs
- Need to reengineer every datatype to ensure SEC (commutativity)
  - ★ Do not mirror sequential counter parts => implementation & proof burden
  - ★ Do not compose!
    - ◆ `counter set` is not a composition of `counter` and `set` CRDTs

Can we *program & reason about* replicated data types as an extension of their sequential counterparts?

Can we *program & reason about* replicated data types  
as an extension of their sequential counterparts?

**MRDT**

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

```




```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end


```



```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end


```



```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end


```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end


```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

```




$$22 = 7 + (8-1) + (21 - 7)$$

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

```




$$22 = 7 + (8-1) + (21 - 7)$$

- 3-way merge function makes the counter suitable for distribution

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

```



$$22 = 7 + (8-1) + (21 - 7)$$

- 3-way merge function makes the counter suitable for distribution
- Does not appeal to individual operations => independently extend data-type

# Systems → PL

# Systems → PL


- CRDTs need to take care of systems level concerns such as *message loss, duplication and reordering*

# Systems → PL

- CRDTs need to take care of systems level concerns such as *message loss, duplication and reordering*
- 3-way merge is oblivious to these
  - ◆ By leaving those concerns to MRDT middleware


# Systems → PL

- CRDTs need to take care of systems level concerns such as *message loss, duplication and reordering*
- 3-way merge is oblivious to these
  - ◆ By leaving those concerns to MRDT middleware




# Systems → PL

- CRDTs need to take care of systems level concerns such as *message loss, duplication and reordering*
- 3-way merge is oblivious to these
  - ◆ By leaving those concerns to MRDT middleware



# Systems → PL

- CRDTs need to take care of systems level concerns such as *message loss, duplication and reordering*
- 3-way merge is oblivious to these
  - ◆ By leaving those concerns to MRDT middleware



$$22 = 21 + (21-21) + (22 - 21)$$

Does the 3-way merge idea generalise?

Does the 3-way merge idea generalise?

*Sort of*

# Observed-Removed Set

# Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

# Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

Kaki et al. “Mergeable Replicated Data Types”,  
OOPSLA 2019

# Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```


{1}

Kaki et al. “Mergeable Replicated Data Types”,  
OOPSLA 2019

# Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```




Kaki et al. “Mergeable Replicated Data Types”,  
OOPSLA 2019

# Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

Kaki et al. “Mergeable Replicated Data Types”,  
OOPSLA 2019




# Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

Kaki et al. “Mergeable Replicated Data Types”,  
OOPSLA 2019

$$\begin{aligned} & \{ \} \cup (\{1\} - \{1\}) \cup (\{ \} - \{1\}) \\ &= \{ \} \cup \{ \} \cup \{ \} \\ &= \{ \} \text{ (expected } \{1\}) \end{aligned}$$




# Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

Kaki et al. “Mergeable Replicated Data Types”,  
OOPSLA 2019

$$\begin{aligned} & \{ \} \cup (\{1\} - \{1\}) \cup (\{ \} - \{1\}) \\ &= \{ \} \cup \{ \} \cup \{ \} \\ &= \{ \} \text{ (expected } \{1\}) \end{aligned}$$




- Convergence is not sufficient; *Intent* is not preserved




# Concretising Intent

- Intent is a woolly term
  - ★ *How can we formalise the intent of operations on a data structure?*



# Concretising Intent

- Intent is a woolly term
  - ★ *How can we formalise the intent of operations on a data structure?*
- We need
  - ★ A *formal language* to specify the *intent* of an RDT
  - ★ *Mechanization* to bridge the air gap between specification and implementation due to distributed system complexity



# Peepul — Certified MRDTs



# Peepul — Certified MRDTs

- An F\* library implementing and proving MRDTs
  - ★ <https://github.com/prismlab/peepul>



# Peepul — Certified MRDTs

- An F\* library implementing and proving MRDTs
  - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
  - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014



# Peepul — Certified MRDTs

- An F\* library implementing and proving MRDTs
  - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
  - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014
- *Replication-aware simulation* to connect *specification* with *implementation*



# Peepul — Certified MRDTs

- An F\* library implementing and proving MRDTs
  - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
  - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014
- *Replication-aware simulation* to connect *specification* with *implementation*
- Composition of MRDTs and their proofs!



# Peepul — Certified MRDTs

- An F\* library implementing and proving MRDTs
  - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
  - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014
- *Replication-aware simulation* to connect *specification* with *implementation*
- Composition of MRDTs and their proofs!
- Extracted RDTs are compatible with Irmin — a Git-like distributed database



# Fixing OR-Set

- Discriminate duplicate additions by associating a unique id


# Fixing OR-Set

- Discriminate duplicate additions by associating a unique id

{ (a,1) }


# Fixing OR-Set

- Discriminate duplicate additions by associating a unique id



# Fixing OR-Set


- Discriminate duplicate additions by associating a unique id



# Fixing OR-Set

- Discriminate duplicate additions by associating a unique id

$$\begin{aligned} & \{ \} \cup ( \{ (a,1); (a,2) \} - \{ (a,1) \} ) \cup ( \{ \} - \{ (a,1) \} ) \\ &= \{ \} \cup \{ (a,2) \} \cup \{ \} \\ &= \{ (a,2) \} \end{aligned}$$




# Fixing OR-Set

- Discriminate duplicate additions by associating a unique id

$$\begin{aligned} & \{ \} \cup ( \{ (a,1); (a,2) \} - \{ (a,1) \} ) \cup ( \{ \} - \{ (a,1) \} ) \\ &= \{ \} \cup \{ (a,2) \} \cup \{ \} \\ &= \{ (a,2) \} \end{aligned}$$

- MRDT implementation

$$D_\tau = (\Sigma, \sigma_0, do, merge)$$



# Fixing OR-Set


- Discriminate duplicate additions by associating a unique id

$$\begin{aligned}
 & \{ \} \cup ( \{ (a,1); (a,2) \} - \{ (a,1) \} ) \cup ( \{ \} - \{ (a,1) \} ) \\
 &= \{ \} \cup \{ (a,2) \} \cup \{ \} \\
 &= \{ (a,2) \}
 \end{aligned}$$

- MRDT implementation

$$D_\tau = (\Sigma, \sigma_0, do, merge)$$

- 1:  $\Sigma = \mathcal{P}(\mathbb{N} \times \mathbb{N})$
- 2:  $\sigma_0 = \{ \}$
- 3:  $do(rd, \sigma, t) = (\sigma, \{a \mid (a, t) \in \sigma\})$
- 4:  $do(add(a), \sigma, t) = (\sigma \cup \{(a, t)\}, \perp)$
- 5:  $do(remove(a), \sigma, t) = (\{e \in \sigma \mid fst(e) \neq a\}, \perp)$
- 6:  $merge(\sigma_{lca}, \sigma_a, \sigma_b) =$   
 $(\sigma_{lca} \cap \sigma_a \cap \sigma_b) \cup (\sigma_a - \sigma_{lca}) \cup (\sigma_b - \sigma_{lca})$



# Fixing OR-Set


- Discriminate duplicate additions by associating a unique id

$$\begin{aligned}
 & \{ \} \cup ( \{ (a,1); (a,2) \} - \{ (a,1) \} ) \cup ( \{ \} - \{ (a,1) \} ) \\
 &= \{ \} \cup \{ (a,2) \} \cup \{ \} \\
 &= \{ (a,2) \}
 \end{aligned}$$

- MRDT implementation

$$D_\tau = (\Sigma, \sigma_0, do, merge)$$

- $\Sigma = \mathcal{P}(\mathbb{N} \times \mathbb{N})$
- $\sigma_0 = \{ \}$
- $do(rd, \sigma, t) = (\sigma, \{a \mid (a, t) \in \sigma\})$
- $do(add(a), \sigma, t) = (\sigma \cup \{(a, t)\}, \perp)$
- $do(remove(a), \sigma, t) = (\{e \in \sigma \mid fst(e) \neq a\}, \perp)$
- $merge(\sigma_{lca}, \sigma_a, \sigma_b) =$   
 $(\sigma_{lca} \cap \sigma_a \cap \sigma_b) \cup (\sigma_a - \sigma_{lca}) \cup (\sigma_b - \sigma_{lca})$





# Specifying OR-Set

Abstract state  $I = \langle E, oper, rval, time, vis \rangle$



# Specifying OR-Set

Abstract state  $I = \langle E, oper, rval, time, vis \rangle$



# Specifying OR-Set


Abstract state  $I = \langle E, oper, rval, time, vis \rangle$



$$\begin{aligned}
 \mathcal{F}_{orset}(\text{rd}, \langle E, oper, rval, time, vis \rangle) &= \{a \mid \exists e \in E. \text{oper}(e) \\
 &= \text{add}(a) \wedge \neg(\exists f \in E. \text{oper}(f) = \text{remove}(a) \wedge e \xrightarrow{vis} f)\}
 \end{aligned}$$

# Specifying OR-Set

Abstract state  $I = \langle E, oper, rval, time, vis \rangle$



$$\begin{aligned}
 \mathcal{F}_{orset}(\text{rd}, \langle E, oper, rval, time, vis \rangle) &= \{a \mid \exists e \in E. \text{oper}(e) \\
 &= \text{add}(a) \wedge \neg(\exists f \in E. \text{oper}(f) = \text{remove}(a) \wedge e \xrightarrow{vis} f)\}
 \end{aligned}$$


# Simulation Relation

- Connects the abstract execution with the concrete state
- For the OR-set,

$$\begin{aligned}\mathcal{R}_{sim}(I, \sigma) \iff (\forall (a, t) \in \sigma \iff \\ (\exists e \in I.E \wedge I.\text{oper}(e) = \text{add}(a) \wedge I.\text{time}(e) = t \wedge \\ \neg(\exists f \in I.E \wedge I.\text{oper}(f) = \text{remove}(a) \wedge e \xrightarrow{\text{vis}} f)))\end{aligned}$$


# Verifying Operations

I. Show that the simulation holds for operations




# Verifying Operations

## I. Show that the simulation holds for operations



# Verifying Operations


1. Show that the simulation holds for operations




2. Show that the simulation holds for merge

# Verifying Operations

## 1. Show that the simulation holds for operations




## 2. Show that the simulation holds for merge




# Verifying Operations

## 1. Show that the simulation holds for operations



## 2. Show that the simulation holds for merge



# Verifying Operations

3. Show that the specification and the implementation agree on the return values of operations

$$\Phi_{spec}(\mathcal{R}_{sim}) \quad \begin{aligned} & \forall I, \sigma, e, op, a, t. \mathcal{R}_{sim}(I, \sigma) \wedge do^{\#}(I, e, op, a, t) = I' \\ & \wedge \mathcal{D}_{\tau}.do(op, \sigma, t) = (\sigma', a) \wedge \Psi_{ts}(I) \implies a = \mathcal{F}_{\tau}(o, I) \end{aligned}$$

# Verifying Operations

3. Show that the specification and the implementation agree on the return values of operations

$$\Phi_{spec}(\mathcal{R}_{sim}) \quad \begin{aligned} & \forall I, \sigma, e, op, a, t. \mathcal{R}_{sim}(I, \sigma) \wedge do^{\#}(I, e, op, a, t) = I' \\ & \wedge \mathcal{D}_{\tau}.do(op, \sigma, t) = (\sigma', a) \wedge \Psi_{ts}(I) \implies a = \mathcal{F}_{\tau}(o, I) \end{aligned}$$

4. Convergence

$$\Phi_{con}(\mathcal{R}_{sim}) \quad \forall I, \sigma_a, \sigma_b. \mathcal{R}_{sim}(I, \sigma_a) \wedge \mathcal{R}_{sim}(I, \sigma_b) \implies \sigma_a \sim \sigma_b$$

# Verifying Operations

3. Show that the specification and the implementation agree on the return values of operations

$$\Phi_{spec}(\mathcal{R}_{sim}) \quad \begin{aligned} & \forall I, \sigma, e, op, a, t. \mathcal{R}_{sim}(I, \sigma) \wedge do^{\#}(I, e, op, a, t) = I \\ & \wedge \mathcal{D}_{\tau}.do(op, \sigma, t) = (\sigma', a) \wedge \Psi_{ts}(I) \implies a = \mathcal{F}_{\tau}(o, I) \end{aligned}$$

4. Convergence

$$\Phi_{con}(\mathcal{R}_{sim}) \quad \forall I, \sigma_a, \sigma_b. \mathcal{R}_{sim}(I, \sigma_a) \wedge \mathcal{R}_{sim}(I, \sigma_b) \implies \sigma_a \sim \sigma_b$$

- ◆ Permits the different replicas to converge to states that are *observationally equal* but not *structurally equal*
- ❖ Example: differently balanced BSTs

# Verifying Operations


3. Show that the specification and the implementation agree on the return values of operations

$$\Phi_{spec}(\mathcal{R}_{sim}) \quad \begin{aligned} & \forall I, \sigma, e, op, a, t. \mathcal{R}_{sim}(I, \sigma) \wedge do^{\#}(I, e, op, a, t) = I' \\ & \wedge \mathcal{D}_{\tau}.do(op, \sigma, t) = (\sigma', a) \wedge \Psi_{ts}(I) \implies a = \mathcal{F}_{\tau}(o, I) \end{aligned}$$

## 4. Convergence


$$\Phi_{con}(\mathcal{R}_{sim}) \quad \forall I, \sigma_a, \sigma_b. \mathcal{R}_{sim}(I, \sigma_a) \wedge \mathcal{R}_{sim}(I, \sigma_b) \implies \sigma_a \sim \sigma_b$$

- ◆ Permits the different replicas to converge to states that are *observationally equal* but not *structurally equal*
- ❖ Example: differently balanced BSTs



# Space-efficient OR-Set


- Recall that the OR-set has duplicates



- How can we remove them?

# Space-efficient OR-Set


- Recall that the OR-set has duplicates



- How can we remove them?
- Idea**
  - ★ On addition, replace existing element's timestamp with the new timestamp
  - ★ On merge, pick the larger timestamp

# Space-efficient OR-Set

- Recall that the OR-set has duplicates




- How can we remove them?
- Idea**


Correctness argument is tricky

- ★ On addition, replace existing element's timestamp with the new timestamp
- ★ On merge, pick the larger timestamp

# Space-efficient OR-Set




# Space-efficient OR-Set



$$\begin{aligned}
 \mathcal{R}_{sim}((E, oper, rval, time, vis), \sigma) \iff \\
 (\forall (a, t) \in \sigma \implies (\exists e \in E. oper(e) = add(a) \wedge time(e) = t \\
 \wedge \neg(\exists r \in E. oper(r) = remove(a) \wedge e \xrightarrow{vis} r)) \wedge \\
 (\forall e \in E. oper(e) = add(a) \wedge \neg(\exists r \in E. oper(r) = remove(a) \\
 \wedge e \xrightarrow{vis} r) \implies t \geq time(e))) \wedge \\
 (\forall e \in E. \forall a \in \mathbb{N}. oper(e) = add(a) \\
 \wedge \neg(\exists r \in E. oper(r) = remove(a) \wedge e \xrightarrow{vis} r) \implies (a, \_) \in \sigma)
 \end{aligned}$$

# Space-efficient OR-Set



$$\begin{aligned}
 \mathcal{R}_{sim}((E, oper, rval, time, vis), \sigma) \iff \\
 (\forall (a, t) \in \sigma \implies (\exists e \in E. oper(e) = add(a) \wedge time(e) = t \\
 \wedge \neg(\exists r \in E. oper(r) = remove(a) \wedge e \xrightarrow{vis} r)) \wedge \\
 (\forall e \in E. oper(e) = add(a) \wedge \neg(\exists r \in E. oper(r) = remove(a) \\
 \wedge e \xrightarrow{vis} r) \implies t \geq time(e))) \wedge \\
 (\forall e \in E. \forall a \in \mathbb{N}. oper(e) = add(a) \\
 \wedge \neg(\exists r \in E. oper(r) = remove(a) \wedge e \xrightarrow{vis} r) \implies (a, \_) \in \sigma)
 \end{aligned}$$

*Simulation relation is more intricate as one would expect*

# Verification effort

MRDTs verified	#Lines code	#Lines proof	#Lemmas	Verif. time (s)
Increment-only counter	6	43	2	3.494
PN counter	8	43	2	23.211
Enable-wins flag	20	58 81 89	3 6 7	1074 171 104
LWW register	5	44	1	4.21
G-set	10	23 28 33	0 1 2	4.71 2.462 1.993
G-map	48	26	0	26.089
Mergeable log	39	95	2	36.562
OR-set (§2.1.1)	30	36 41 46	0 1 2	43.85 21.656 8.829
OR-set-space (§2.1.2)	59	108	7	1716
OR-set-spacetime	97	266	7	1854
Queue	32	1123	75	4753

# Composing CRDTs is HARD!



**Martin Kleppmann**  
@martinkl

...

Today in “distributed systems are hard”: I wrote down a simple CRDT algorithm that I thought was “obviously correct” for a course I’m teaching. Only 10 lines or so long. Found a fatal bug only after spending hours trying to prove the algorithm correct. 😭

4:18 AM · Nov 13, 2020 · Tweetbot for iOS

---

41 Retweets 4 Quote Tweets 541 Likes

---



**Martin Kleppmann** @martinkl · Nov 13, 2020

...

The interesting thing about this bug is that it comes about only from the interaction of two features. A LWW map by itself is fine. A set in which you can insert and delete elements (but not update them) is fine. The problem arises only when delete and update interact.



# Composing IRC-style chat

- Build IRC-style group chat
  - ★ Send and read messages in channels
  - ★ For simplicity, channels and messages cannot be deleted
- Represent application state as a **grow-only map** with string (channel name) keys and **mergeable-log** as values
- **Goal:**
  - ★ *map* and *log* proved correct separately
  - ★ Use the proof of underlying RDTs to prove chat application correctness

# Generic Map MRDT

- Specification

# Generic Map MRDT

- Specification

$$\mathcal{F}_{\alpha\text{-}map}(get(k, o_\alpha), I) = \\ \text{let } I_\alpha = \text{project}(k, I) \text{ in } \mathcal{F}_\alpha(o_\alpha, I_\alpha)$$

where

$$\text{project } k \ I_{\alpha\text{-}map} = I_\alpha$$

# Generic Map MRDT

- Specification

$$\mathcal{F}_{\alpha\text{-map}}(\text{get}(k, o_\alpha), I) = \\ \text{let } I_\alpha = \text{project}(k, I) \text{ in } \mathcal{F}_\alpha(o_\alpha, I_\alpha)$$

where

$$\text{project } k \ I_{\alpha\text{-map}} = I_\alpha$$

- Project *filters* the abstract state of the map on the key **k** and returns an abstract state of the underlying data type
  - ★ Provided by the user once for a generic MRDT

# Generic Map MRDT


- Specification

$$\mathcal{F}_{\alpha\text{-map}}(\text{get}(k, o_\alpha), I) = \\ \text{let } I_\alpha = \text{project}(k, I) \text{ in } \mathcal{F}_\alpha(o_\alpha, I_\alpha)$$

where

$$\text{project } k \ I_{\alpha\text{-map}} = I_\alpha$$

- Project *filters* the abstract state of the map on the key *k* and returns an abstract state of the underlying data type
  - ★ Provided by the user once for a generic MRDT



# Generic Map MRDT

## Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$  where

- 1:  $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$
- 2:  $\sigma_0 = \{\}$
- 3:  $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_{0_\alpha}, & \text{otherwise} \end{cases}$
- 4:  $do(\text{set}(k, o_\alpha), \sigma, t) =$   
    let  $(v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$  in  $(\sigma[k \mapsto v], r)$
- 5:  $do(\text{get}(k, o_\alpha), \sigma, t) =$   
    let  $(_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$  in  $(\sigma, r)$
- 6:  $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$   
     $\{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$   
     $v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

## Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

- 1:  $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{oper}(e) = \text{set}(k, \_)) \wedge$
- 2:  $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

# Generic Map MRDT

## Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$  where

1:  $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$

2:  $\sigma_0 = \{\}$

3:  $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_{0_\alpha}, & \text{otherwise} \end{cases}$

4:  $do(\text{set}(k, o_\alpha), \sigma, t) =$   
 $\quad \text{let } (v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma[k \mapsto v], r)$

5:  $do(\text{get}(k, o_\alpha), \sigma, t) =$   
 $\quad \text{let } (\_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma, r)$

6:  $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$   
 $\quad \{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$   
 $\quad \quad v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

*Get applies given operation on the value at key k and returns the value*

## Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

1:  $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{ oper}(e) = \text{set}(k, \_)) \wedge$

2:  $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

# Generic Map MRDT

## Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$  where

1:  $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$

2:  $\sigma_0 = \{\}$

3:  $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_0, & \text{otherwise} \end{cases}$

4:  $do(\text{set}(k, o_\alpha), \sigma, t) =$   
 $\text{let } (v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma[k \mapsto v], r)$

5:  $do(\text{get}(k, o_\alpha), \sigma, t) =$   
 $\text{let } (\_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma, r)$

6:  $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$   
 $\{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$   
 $v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

*Set is Get + update the map with the new state*

*Get applies given operation on the value at key k and returns the value*

## Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

1:  $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{ oper}(e) = \text{set}(k, \_)) \wedge$

2:  $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

# Generic Map MRDT

## Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$  where

1:  $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$

2:  $\sigma_0 = \{\}$

3:  $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_0, & \text{otherwise} \end{cases}$

4:  $do(\text{set}(k, o_\alpha), \sigma, t) =$   
 $\text{let } (v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma[k \mapsto v], r)$

5:  $do(\text{get}(k, o_\alpha), \sigma, t) =$   
 $\text{let } (\_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma, r)$

6:  $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$   
 $\{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$   
 $v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

*Set is Get + update the map with the new state*

*Get applies given operation on the value at key k and returns the value*

*Merge uses the merge of the underlying value type!*

## Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

1:  $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{oper}(e) = \text{set}(k, \_)) \wedge$

2:  $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

# Generic Map MRDT

## Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$  where

1:  $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$

2:  $\sigma_0 = \{\}$

3:  $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_0, & \text{otherwise} \end{cases}$

4:  $do(\text{set}(k, o_\alpha), \sigma, t) =$   
 $\text{let } (v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma[k \mapsto v], r)$

5:  $do(\text{get}(k, o_\alpha), \sigma, t) =$   
 $\text{let } (\_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma, r)$

6:  $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$   
 $\{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$   
 $v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

Set is Get + update the map with the new state

Get applies given operation on the value at key k and returns the value

Merge uses the merge of the underlying value type!

## Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

1:  $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{oper}(e) = \text{set}(k, \_)) \wedge$

2:  $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

Simulation relation appeals to the value type's simulation relation!

# Composing IRC-style chat


- Program state is constructed by instantiating *generic map* with *mergeable log*
  - ★ The proof of correctness of the chat application directly follows from the composition!

# Mergeable Queues

- Replicated queue with *at-least-once* dequeue semantics
  - ★ First verified queue RDT!


# Mergeable Queues

- Replicated queue with *at-least-once* dequeue semantics
  - ★ First verified queue RDT!



# Mergeable Queues


- Replicated queue with *at-least-once* dequeue semantics
  - ★ First verified queue RDT!



- Our aim is to have  $O(1)$  enqueue and dequeue and  $O(n)$  merge

# Mergeable Queues

- Implementation
  - ★ Uses *two-list functional queue* implementation
    - ◆ amortised  $O(1)$  enqueue and dequeue operations
  - ★ Merge uses *longest common contiguous subsequence* algorithm —  $O(n)$




# Mergeable Queues

- Implementation

- ★ Uses *two-list functional queue* implementation
  - ◆ amortised  $O(1)$  enqueue and dequeue operations
- ★ Merge uses *longest common contiguous subsequence* algorithm —  $O(n)$

- Specification

1. Any element *popped* in either A or B does not remain in M
2. Any element *pushed* into either A or B appears in M
3. An element that remains *untouched* in LCA, A, B remains in M
4. *Order* of pairs of elements in LCA, A, B must be preserved in M, if those elements are present in M.




# Mergeable Queues

- Implementation

- ★ Uses *two-list functional queue* implementation
  - ◆ amortised  $O(1)$  enqueue and dequeue operations
- ★ Merge uses *longest common contiguous subsequence* algorithm —  $O(n)$

- Specification

1. Any element *popped* in either A or B does not remain in M
2. Any element *pushed* into either A or B appears in M
3. An element that remains *untouched* in LCA, A, B remains in M
4. *Order* of pairs of elements in LCA, A, B must be preserved in M, if those elements are present in M.



*Implementation far removed from the specification!*

# Verification effort

MRDTs verified	#Lines code	#Lines proof	#Lemmas	Verif. time (s)
Increment-only counter	6	43	2	3.494
PN counter	8	43	2	23.211
Enable-wins flag	20	58 81 89	3 6 7	1074 171 104
LWW register	5	44	1	4.21
G-set	10	23 28 33	0 1 2	4.71 2.462 1.993
G-map	48	26	0	26.089
Mergeable log	39	95	2	36.562
OR-set (§2.1.1)	30	36 41 46	0 1 2	43.85 21.656 8.829
OR-set-space (§2.1.2)	59	108	7	1716
OR-set-spacetime	97	266	7	1854
Queue	32	1123	75	4753

# Summary

# Summary

- Programming and proving with RDTs is complicated due to *concurrency* and the lack of suitable *programming abstractions*

# Summary

- Programming and proving with RDTs is complicated due to *concurrency* and the lack of suitable *programming abstractions*
- MRDTs simplify RDTs by implementing them as extensions of sequential data types
  - ★ Reasoning about correctness is still hard

# Summary

- Programming and proving with RDTs is complicated due to *concurrency* and the lack of suitable *programming abstractions*
- MRDTs simplify RDTs by implementing them as extensions of sequential data types
  - ★ Reasoning about correctness is still hard
- Peepul is an F\* library for certified MRDTs
  - ★ Replication-aware simulation for proving complex MRDTs
  - ★ Complex MRDTs can be constructed and proved using simpler MRDTs

# Summary

- Programming and proving with RDTs is complicated due to *concurrency* and the lack of suitable *programming abstractions*
- MRDTs simplify RDTs by implementing them as extensions of sequential data types
  - ★ Reasoning about correctness is still hard
- Peepul is an F\* library for certified MRDTs
  - ★ Replication-aware simulation for proving complex MRDTs
  - ★ Complex MRDTs can be constructed and proved using simpler MRDTs
- F\* allows us to strike a balance between automated and interactive proofs
  - ★ Extract to OCaml and run on Irmin!

# Backup Slides

# Queue Performance

