

Certified Mergeable Replicated Data Types

“KC” Sivaramakrishnan

joint work with

Vimala Soundarapandian, Adharsh Kamath and Kartik Nagar

IIT
MADRAS
SADAM



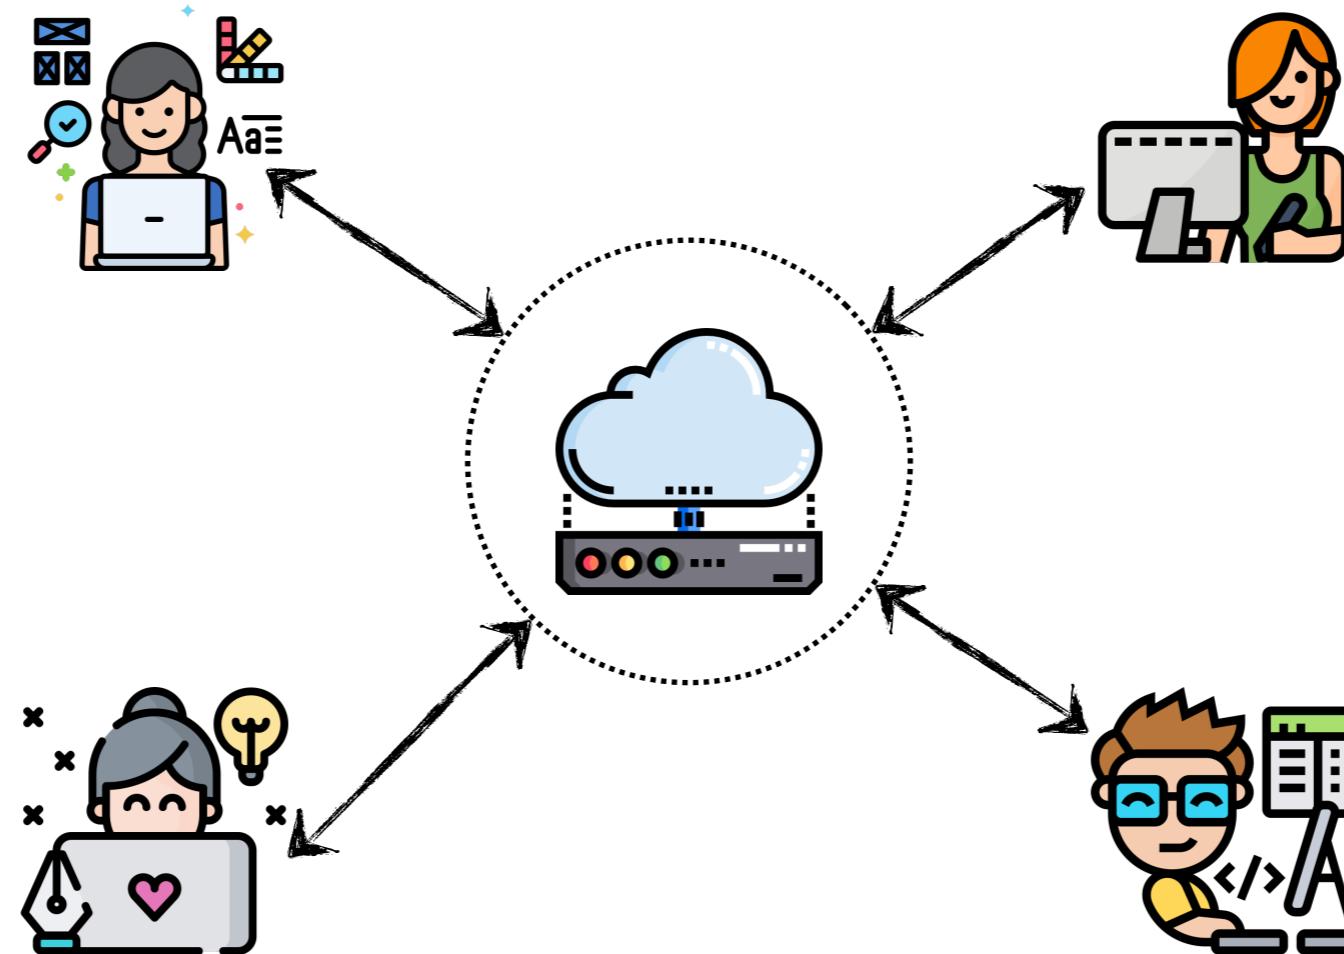
Collaborative Apps



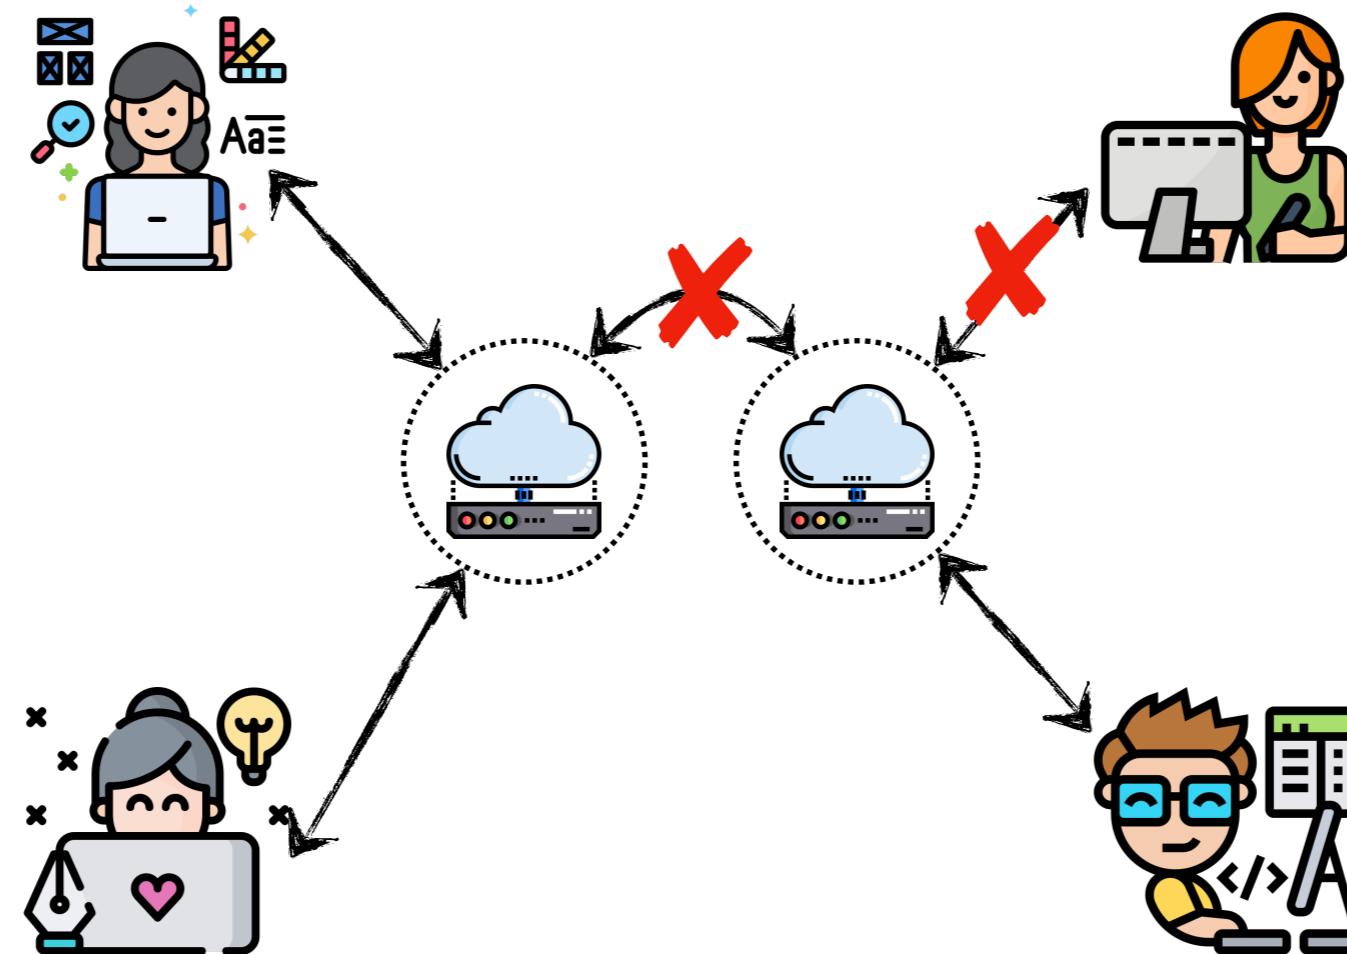
Google Docs



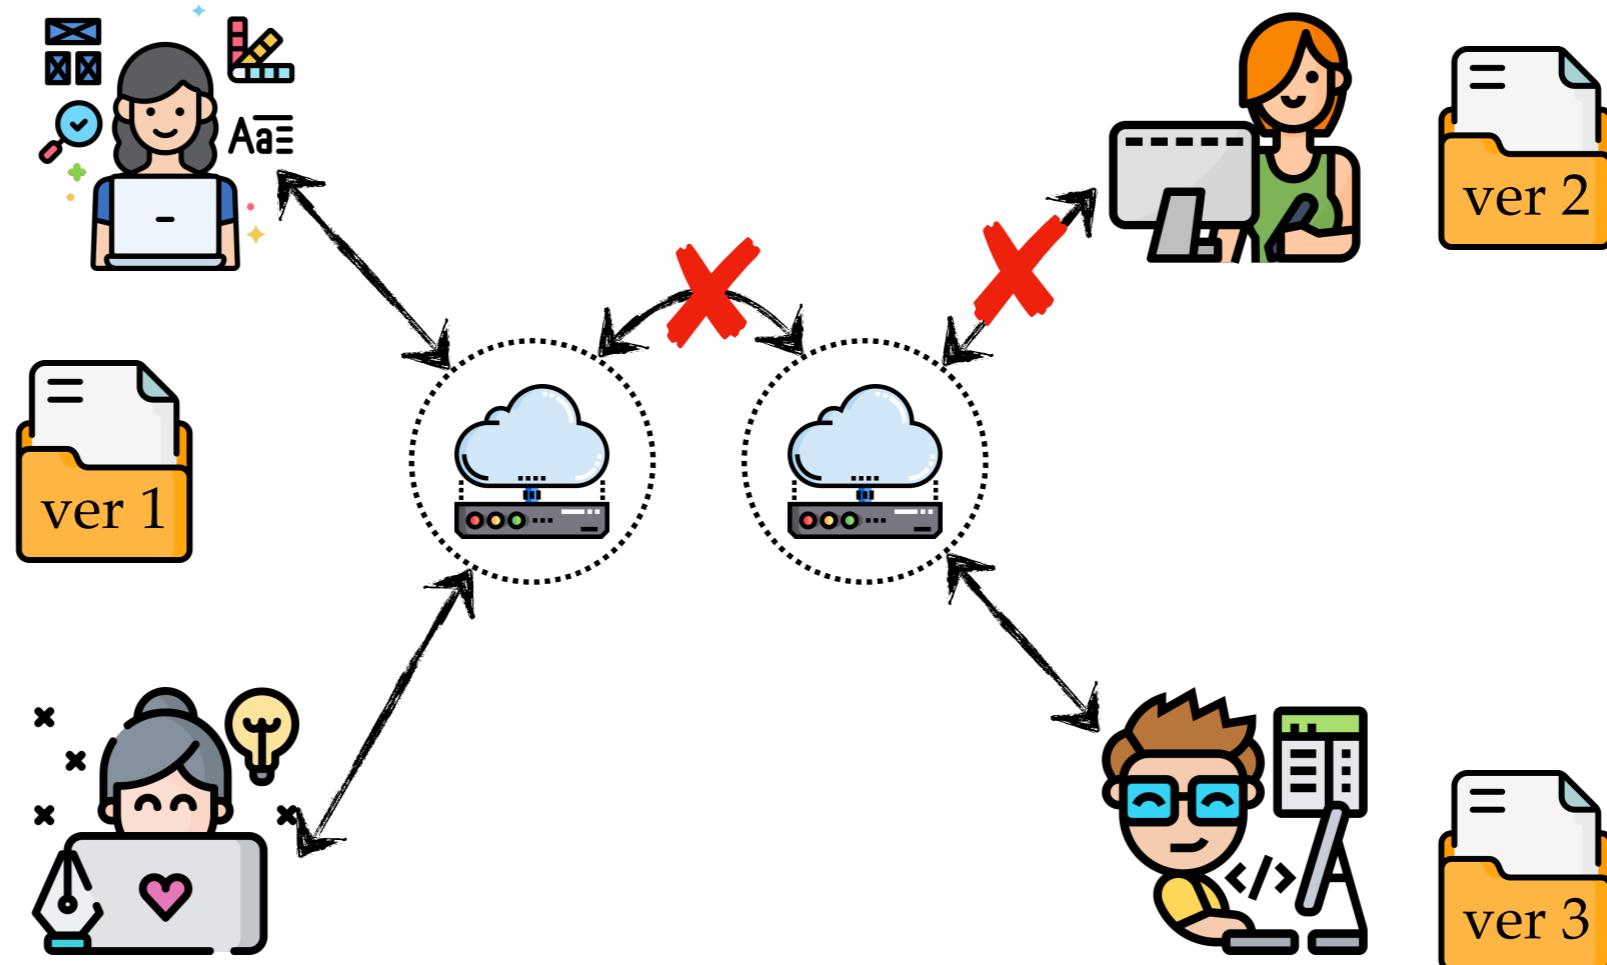
Collaborative Apps



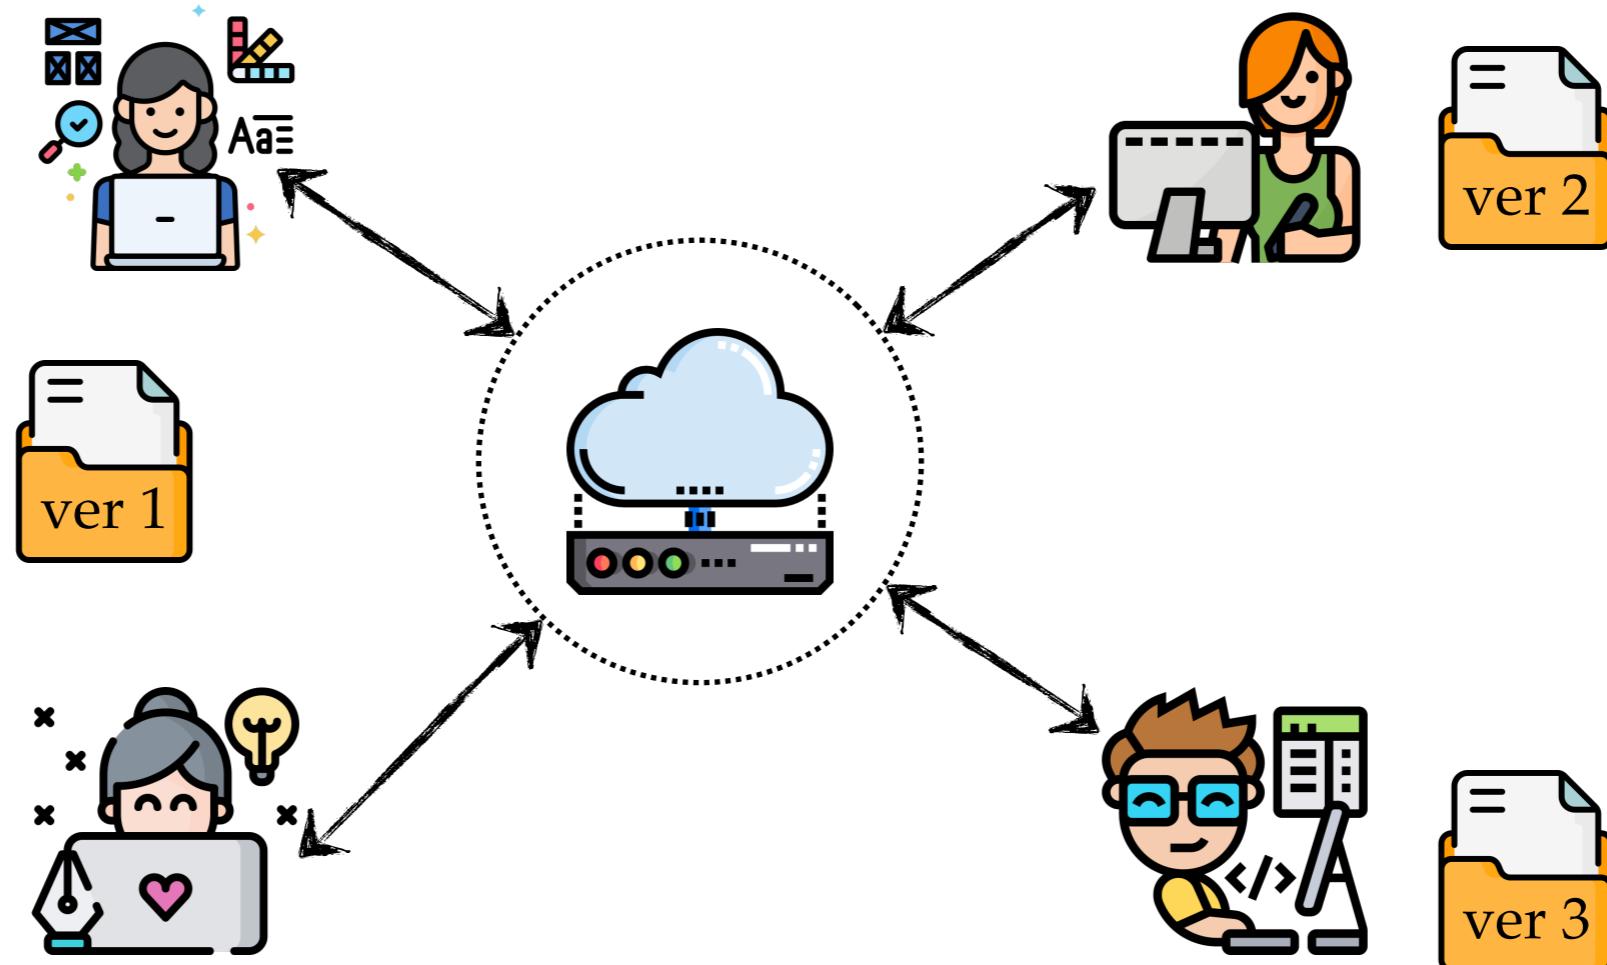
Network Partitions



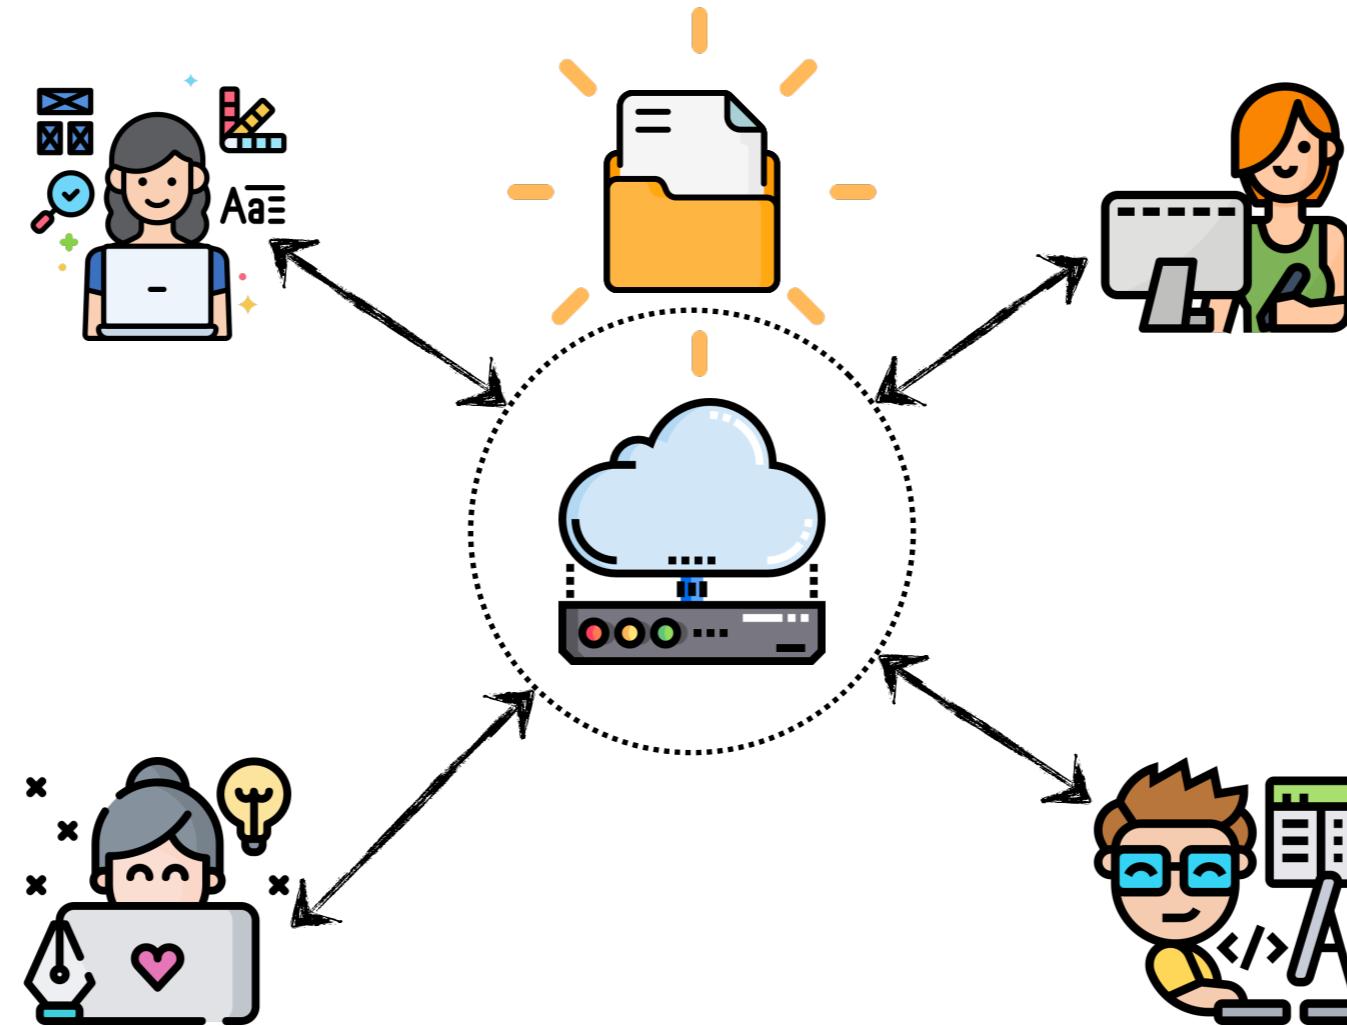
Local-first software



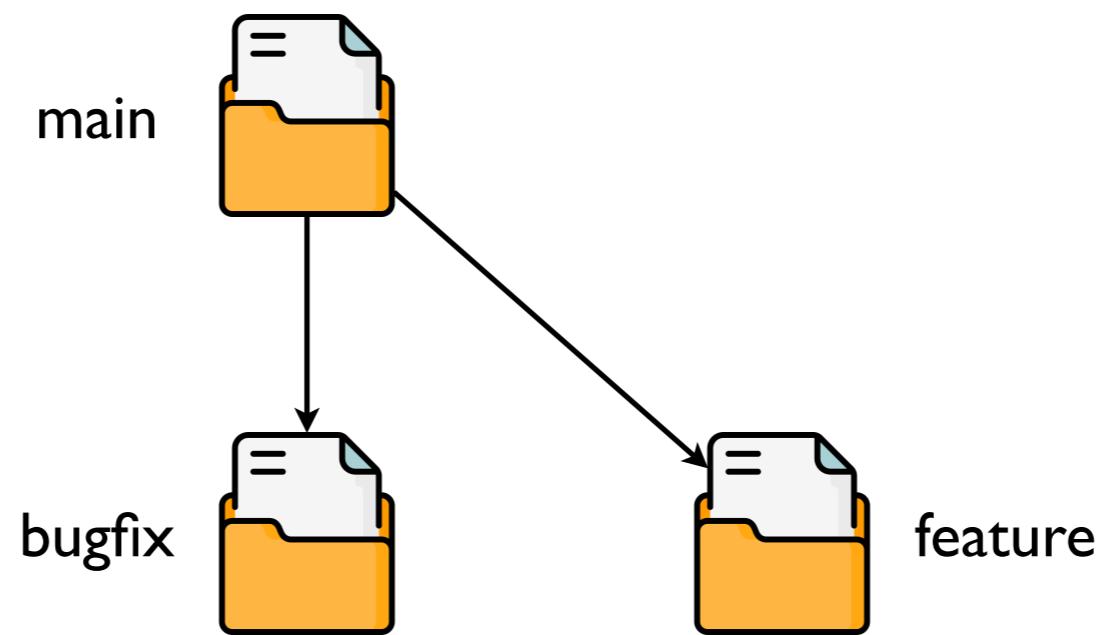
Local-first software



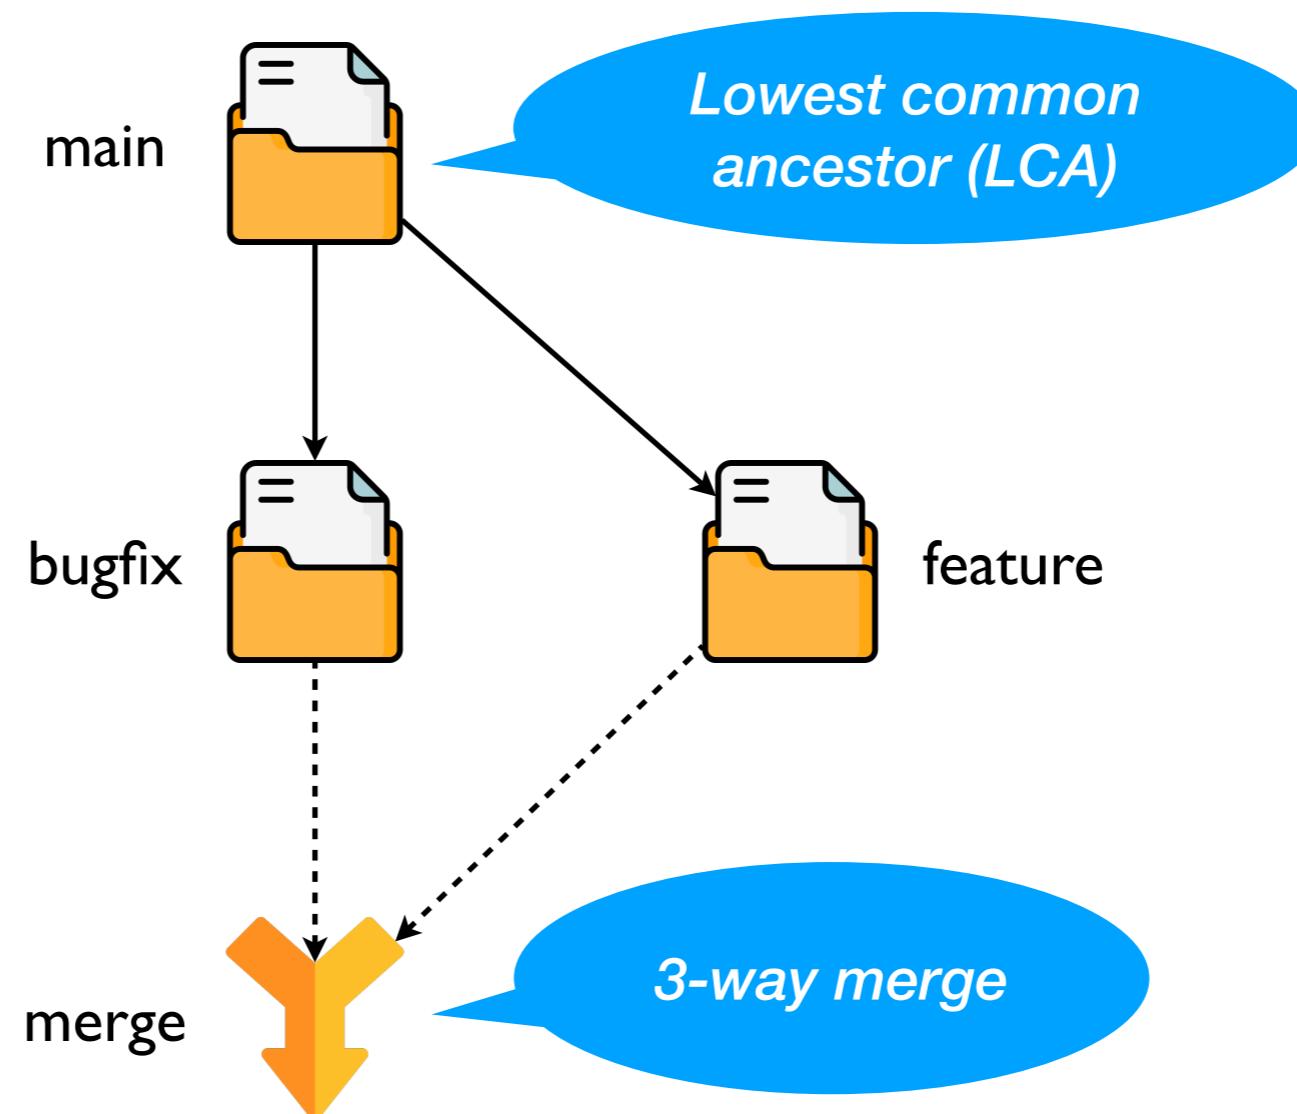
Local-first software



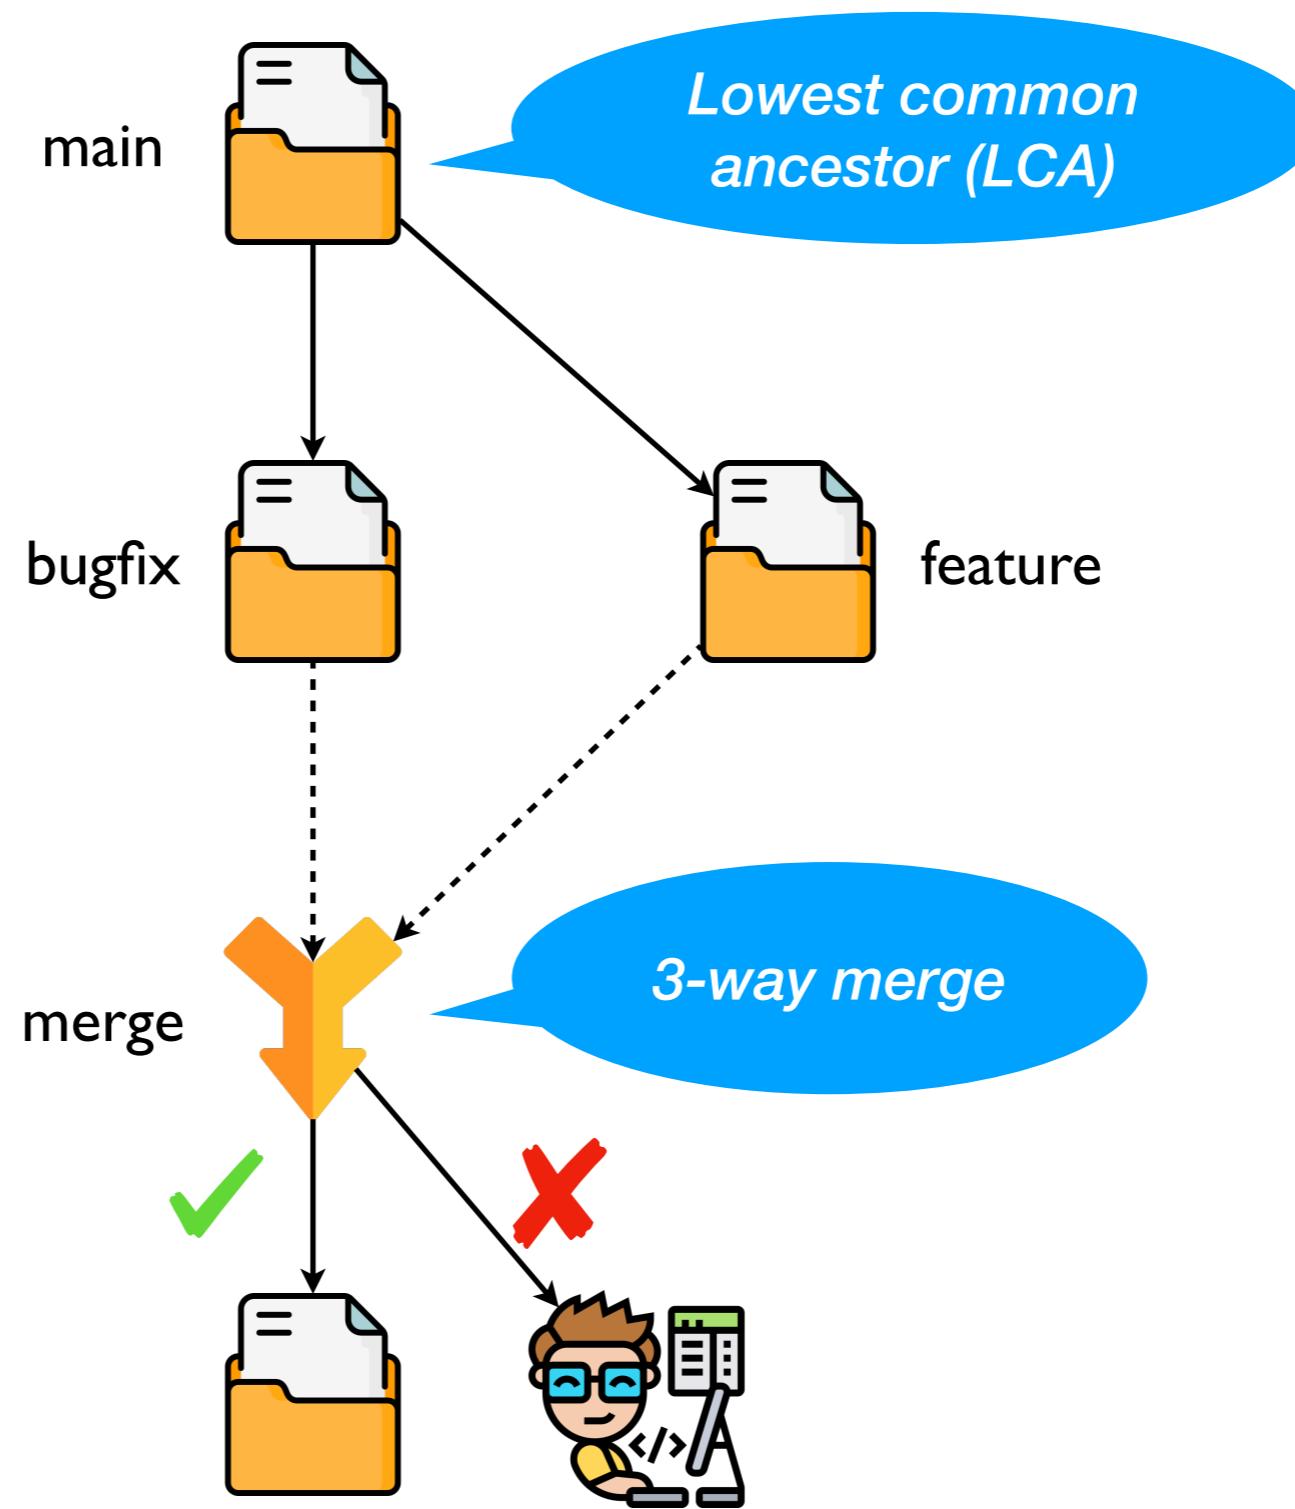
Distributed Version Control Systems



Distributed Version Control Systems



Distributed Version Control Systems



Mergeable Replicated Data Types

- MRDTs — DVCS for *data types* rather than just *text files*
- Sequential data types + 3-way merge = replicated data type!

Mergeable Replicated Data Types

- MRDTs — DVCS for *data types* rather than just *text files*
- Sequential data types + 3-way merge = replicated data type!

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```

Mergeable Replicated Data Types

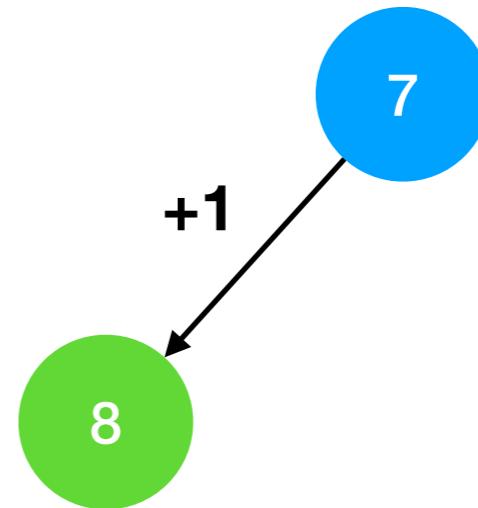
- MRDTs — DVCS for *data types* rather than just *text files*
- Sequential data types + 3-way merge = replicated data type!

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```

Mergeable Replicated Data Types

- MRDTs — DVCS for *data types* rather than just *text files*
- Sequential data types + 3-way merge = replicated data type!

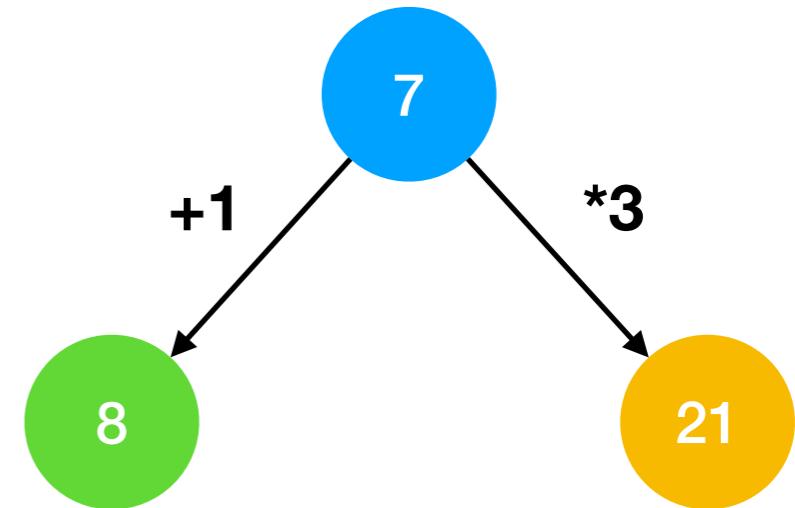
```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```



Mergeable Replicated Data Types

- MRDTs — DVCS for *data types* rather than just *text files*
- Sequential data types + 3-way merge = replicated data type!

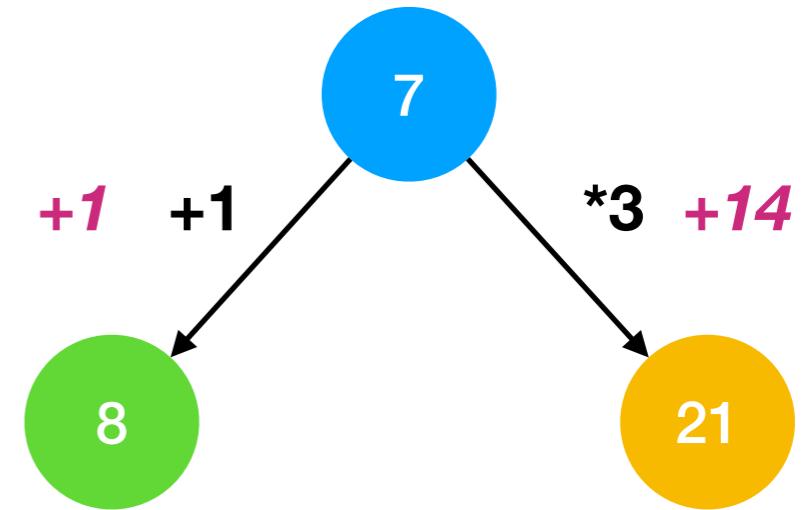
```
module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```



Mergeable Replicated Data Types

- MRDTs — DVCS for *data types* rather than just *text files*
- Sequential data types + 3-way merge = replicated data type!

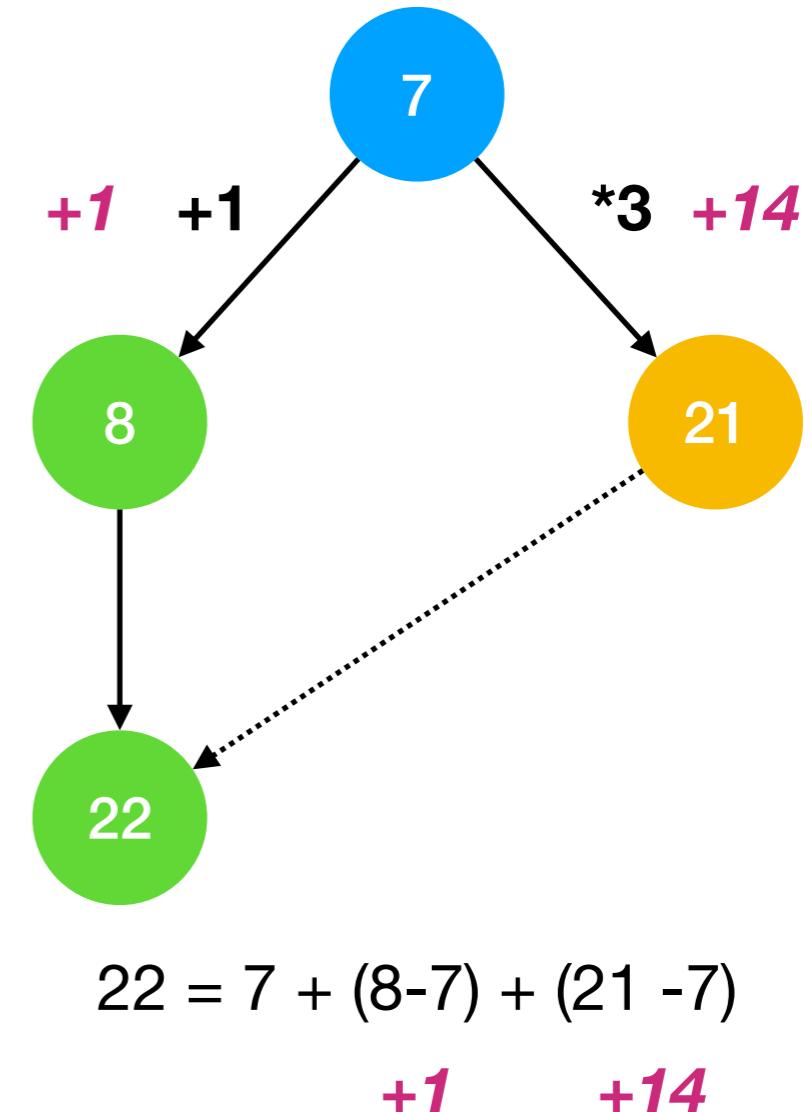
```
module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```



Mergeable Replicated Data Types

- MRDTs — DVCS for *data types* rather than just *text files*
- Sequential data types + 3-way merge = replicated data type!

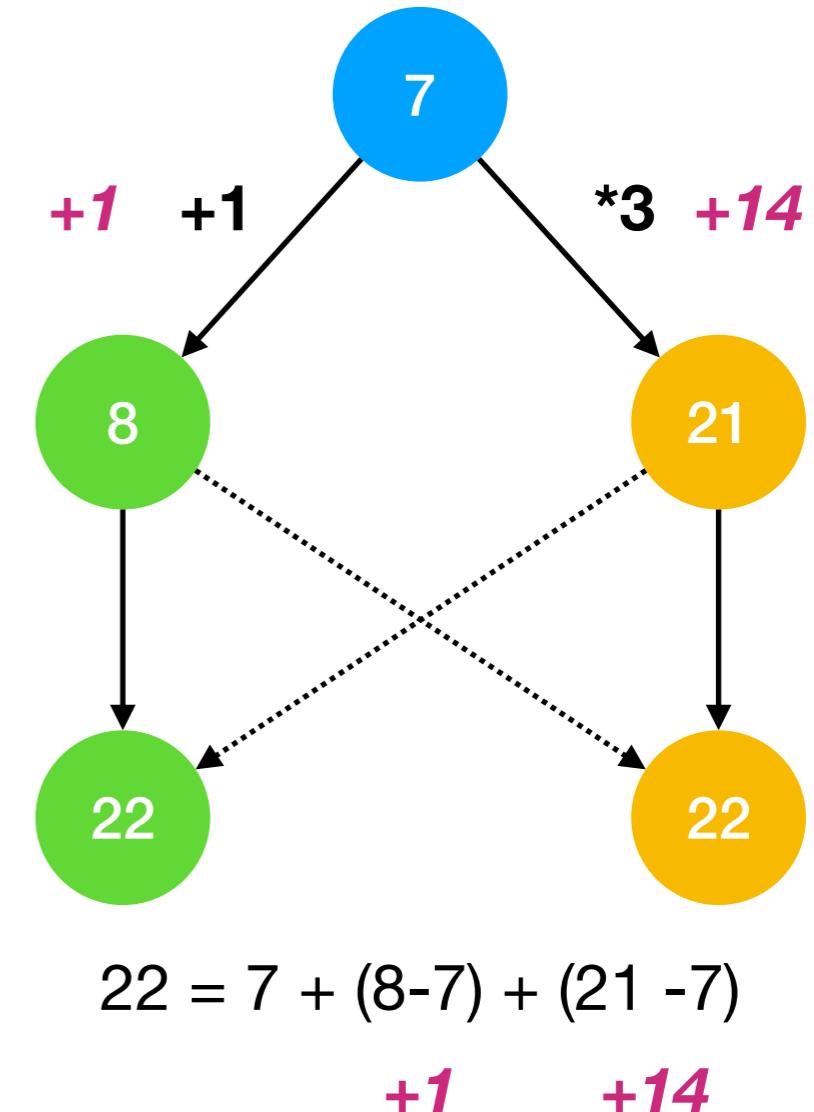
```
module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```



Mergeable Replicated Data Types

- MRDTs — DVCS for *data types* rather than just *text files*
- Sequential data types + 3-way merge = replicated data type!

```
module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```



Does the 3-way merge idea generalise?

Does the 3-way merge idea generalise?

Sort of...

Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

Kaki et al. “Mergeable Replicated Data Types”,
OOPSLA 2019

Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

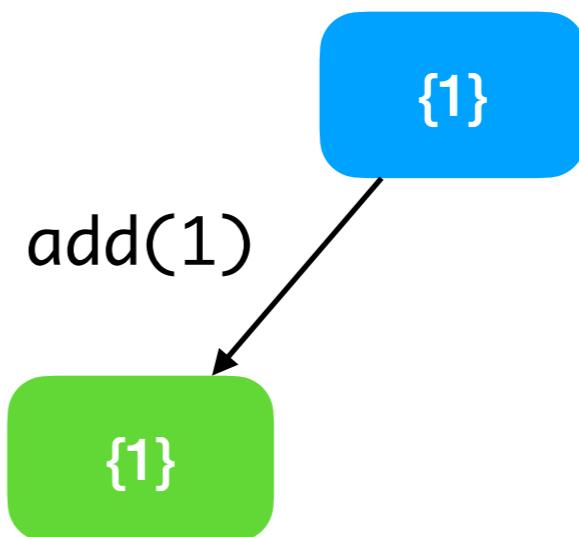
{1}

Kaki et al. “Mergeable Replicated Data Types”,
OOPSLA 2019

Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```



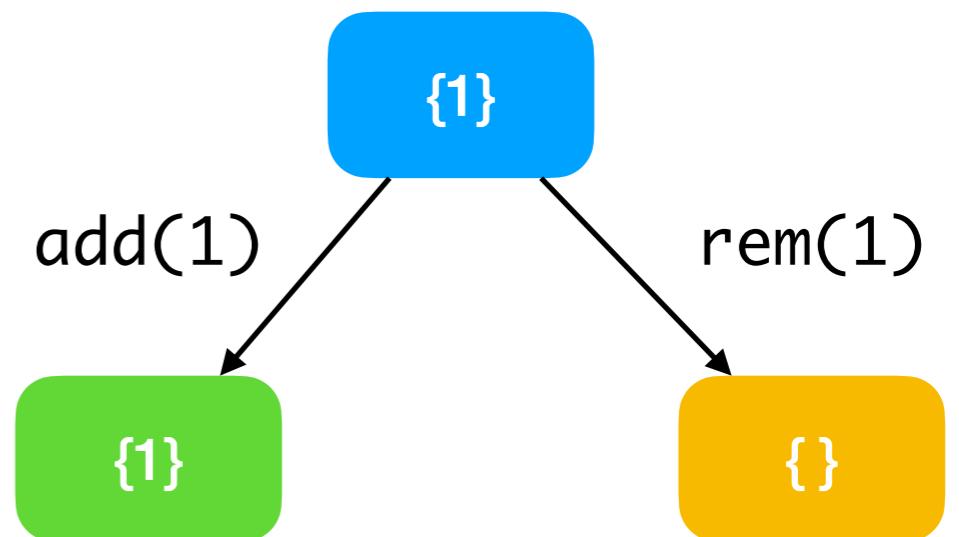
Kaki et al. “Mergeable Replicated Data Types”,
OOPSLA 2019

Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

Kaki et al. “Mergeable Replicated Data Types”,
OOPSLA 2019



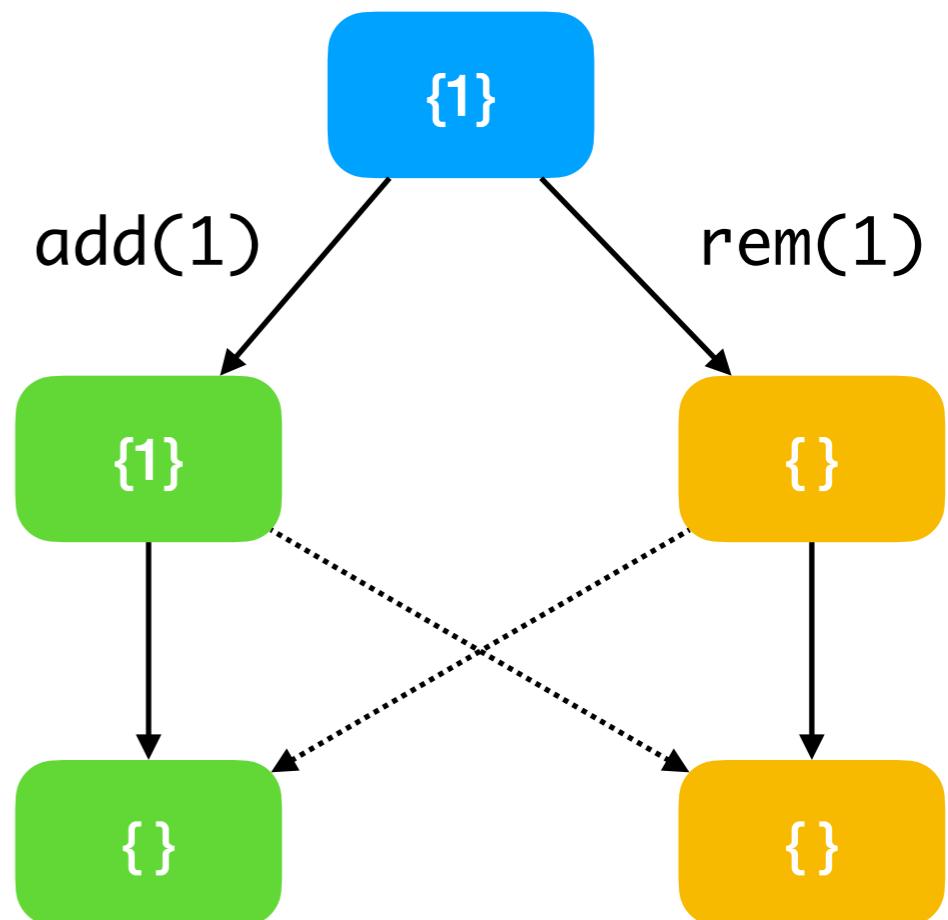
Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

Kaki et al. “Mergeable Replicated Data Types”,
OOPSLA 2019

$$\begin{aligned} & \{ \} \cup (\{1\} - \{1\}) \cup (\{ \} - \{1\}) \\ &= \{ \} \cup \{ \} \cup \{ \} \\ &= \{ \} \text{ (expected } \{1\}) \end{aligned}$$



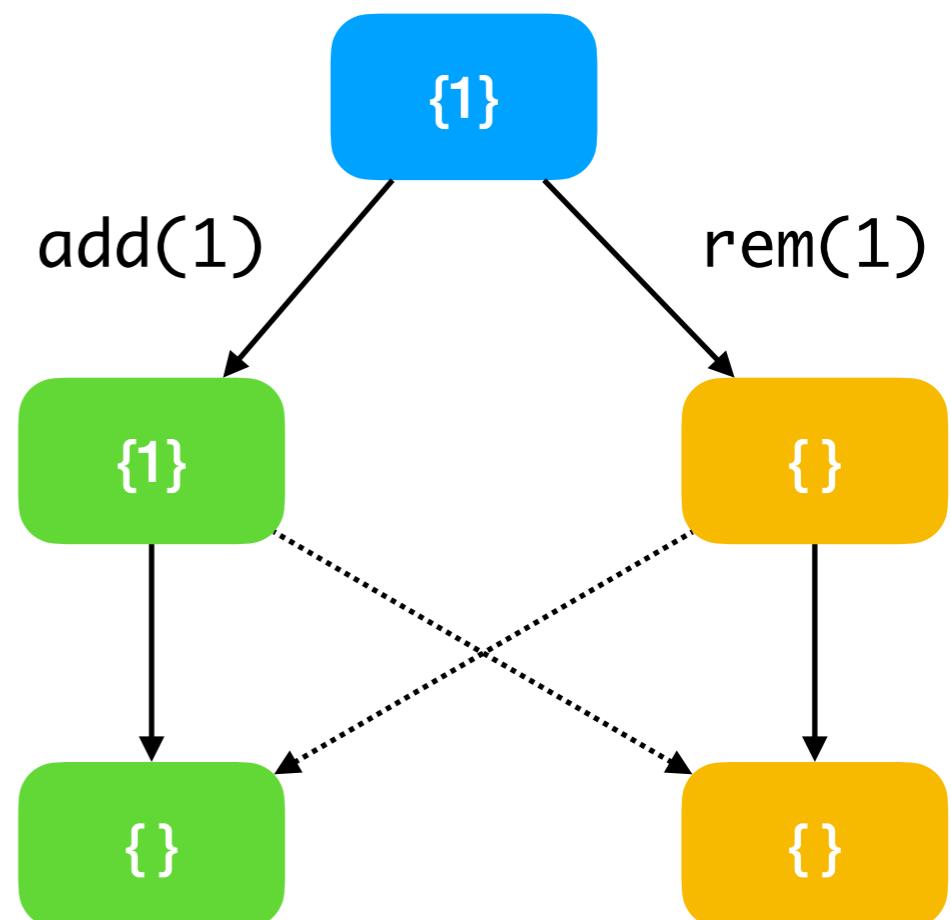
Observed-Removed Set

- OR-set — *add-wins* when there is a concurrent add and remove of the same element

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* unmodified elements *)
  ∪ (v1 - lca) (* added in v1 *)
  ∪ (v2 - lca) (* added in v2 *)
```

Kaki et al. “Mergeable Replicated Data Types”,
OOPSLA 2019

$$\begin{aligned} & \{ \} \cup (\{1\} - \{1\}) \cup (\{ \} - \{1\}) \\ &= \{ \} \cup \{ \} \cup \{ \} \\ &= \{ \} \text{ (expected } \{1\}) \end{aligned}$$



- Convergence is not sufficient; *Intent* is not preserved



Concretising Intent

- A *formal specification language* to capture the *intent* of the MRDT
 - ◆ Must be rich enough to capture eventual consistency

Concretising Intent

- A *formal specification language* to capture the *intent* of the MRDT
 - ◆ Must be rich enough to capture eventual consistency
- Even *simple* data types attract enormous *complexity* when made *distributed*

Concretising Intent

- A *formal specification language* to capture the *intent* of the MRDT
 - ◆ Must be rich enough to capture eventual consistency
- Even *simple* data types attract enormous *complexity* when made *distributed*

Lindsey Kuper
@lindsey

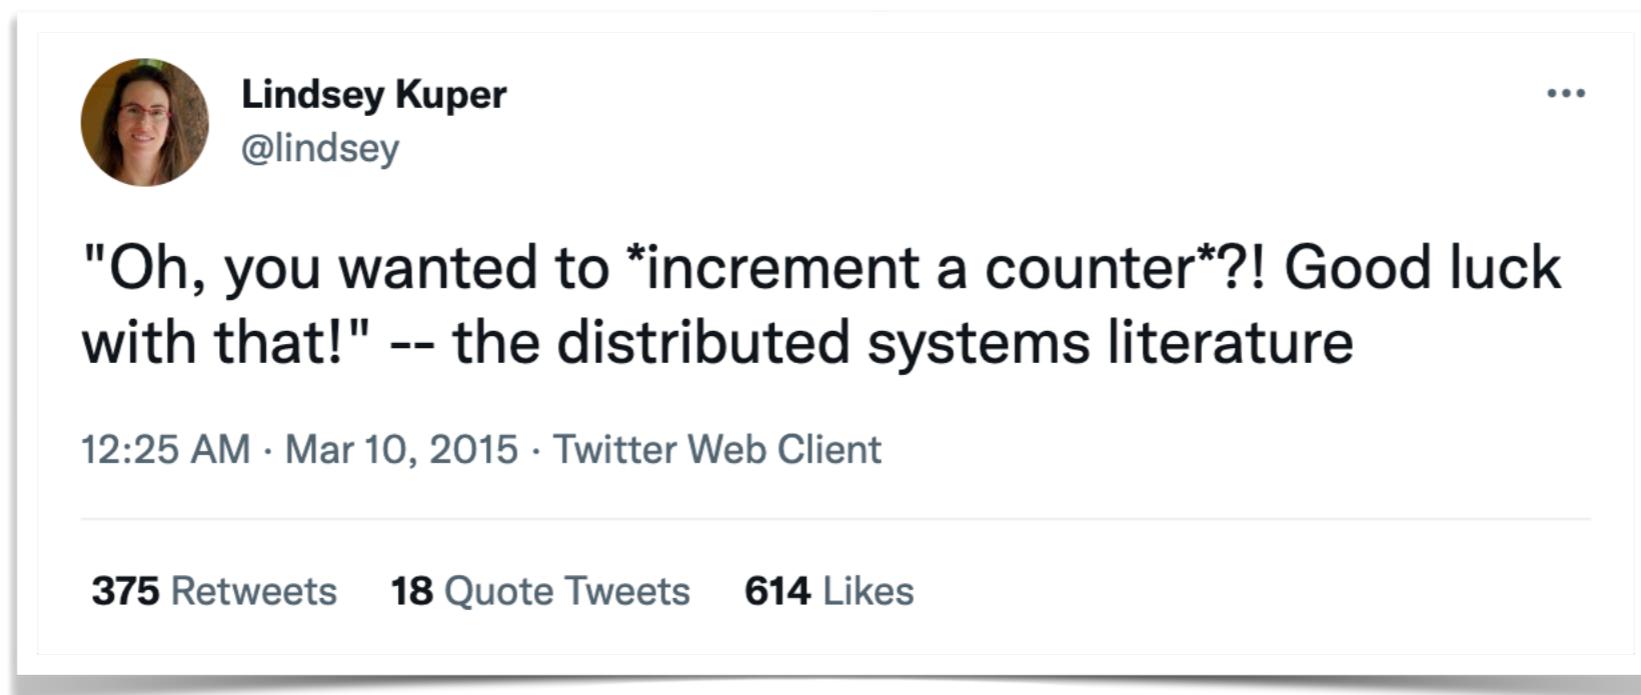
"Oh, you wanted to *increment a counter*?! Good luck with that!" -- the distributed systems literature

12:25 AM · Mar 10, 2015 · Twitter Web Client

375 Retweets 18 Quote Tweets 614 Likes

Concretising Intent

- A *formal specification language* to capture the *intent* of the MRDT
 - ◆ Must be rich enough to capture eventual consistency
- Even *simple* data types attract enormous *complexity* when made *distributed*



- *Mechanization* to bridge the gap between spec and impl

Peepul — Certified MRDTs

- An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>



Peepul — Certified MRDTs

- An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
 - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014



Peepul — Certified MRDTs

- An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
 - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014
- *Replication-aware simulation* to connect *specification* with *implementation*



Peepul — Certified MRDTs

- An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
 - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014
- *Replication-aware simulation* to connect *specification* with *implementation*
- *Space- and time-efficient* implementations
 - ★ 1st certified implementation of a $O(1)$ replicated queue with $O(n)$ merge.



Peepul — Certified MRDTs

- An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
 - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014
- *Replication-aware simulation* to connect *specification* with *implementation*
- *Space- and time-efficient* implementations
 - ★ 1st certified implementation of a $O(1)$ replicated queue with $O(n)$ merge.
- *Composition* of MRDTs and their proofs!



Peepul — Certified MRDTs

- An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>
- Specification language is event-based
 - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014
- *Replication-aware simulation* to connect *specification* with *implementation*
- *Space- and time-efficient* implementations
 - ★ 1st certified implementation of a $O(1)$ replicated queue with $O(n)$ merge.
- *Composition* of MRDTs and their proofs!
- Extracted RDTs are compatible with *Irmin* — a Git-like distributed database



Fixing OR-Set

- Discriminate duplicate additions by associating a unique id

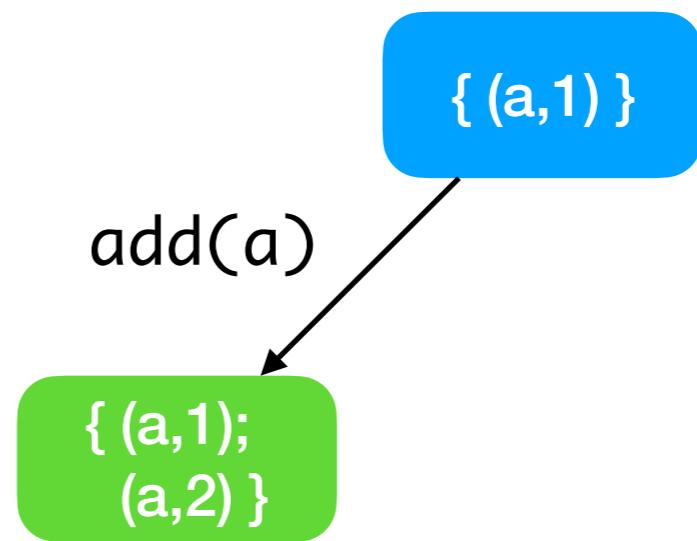
Fixing OR-Set

- Discriminate duplicate additions by associating a unique id

{ (a,1) }

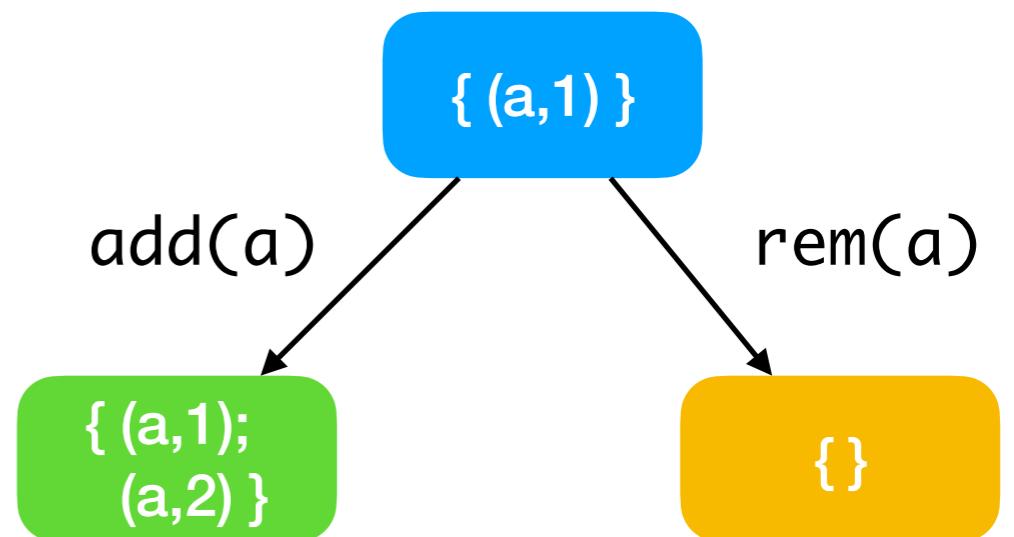
Fixing OR-Set

- Discriminate duplicate additions by associating a unique id



Fixing OR-Set

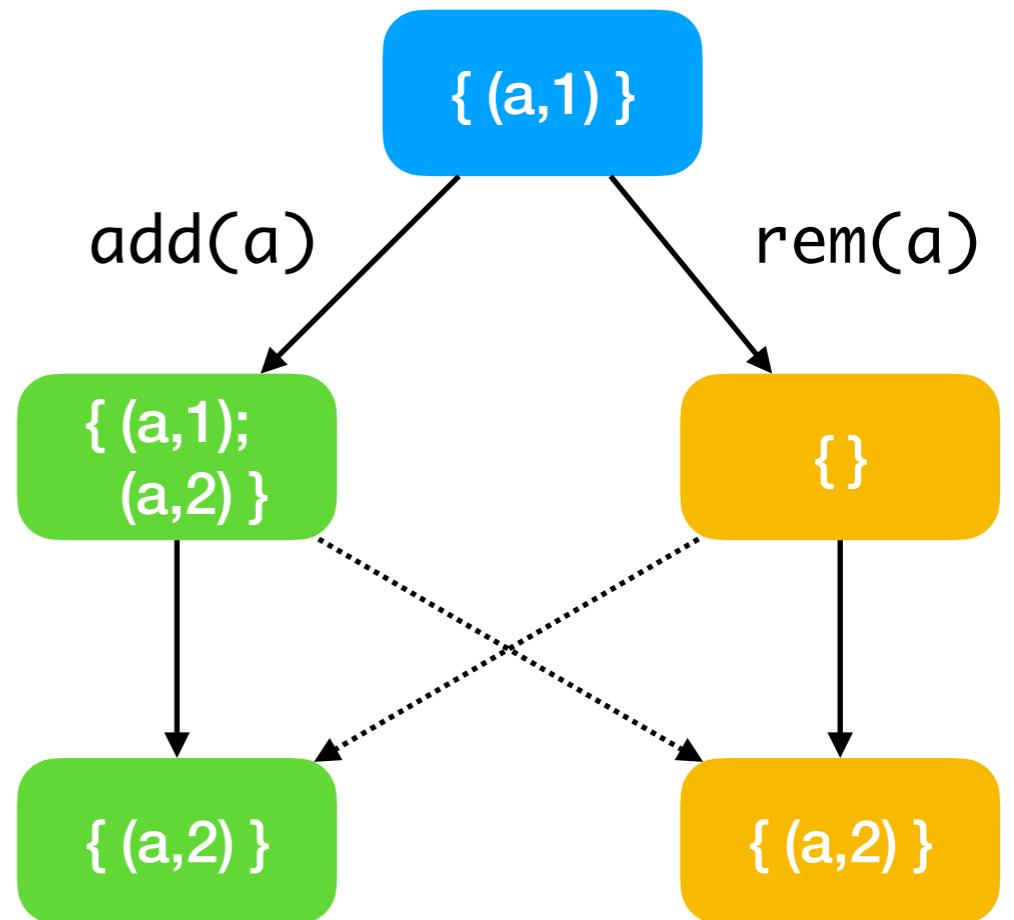
- Discriminate duplicate additions by associating a unique id



Fixing OR-Set

- Discriminate duplicate additions by associating a unique id

$$\begin{aligned} & \{ \} \\ & \cup (\{ (a,1); (a,2) \} - \{ (a,1) \}) \\ & \cup (\{ \} - \{ (a,1) \}) \\ & = \{ \} \cup \{ (a,2) \} \cup \{ \} \\ & = \{ (a,2) \} \end{aligned}$$



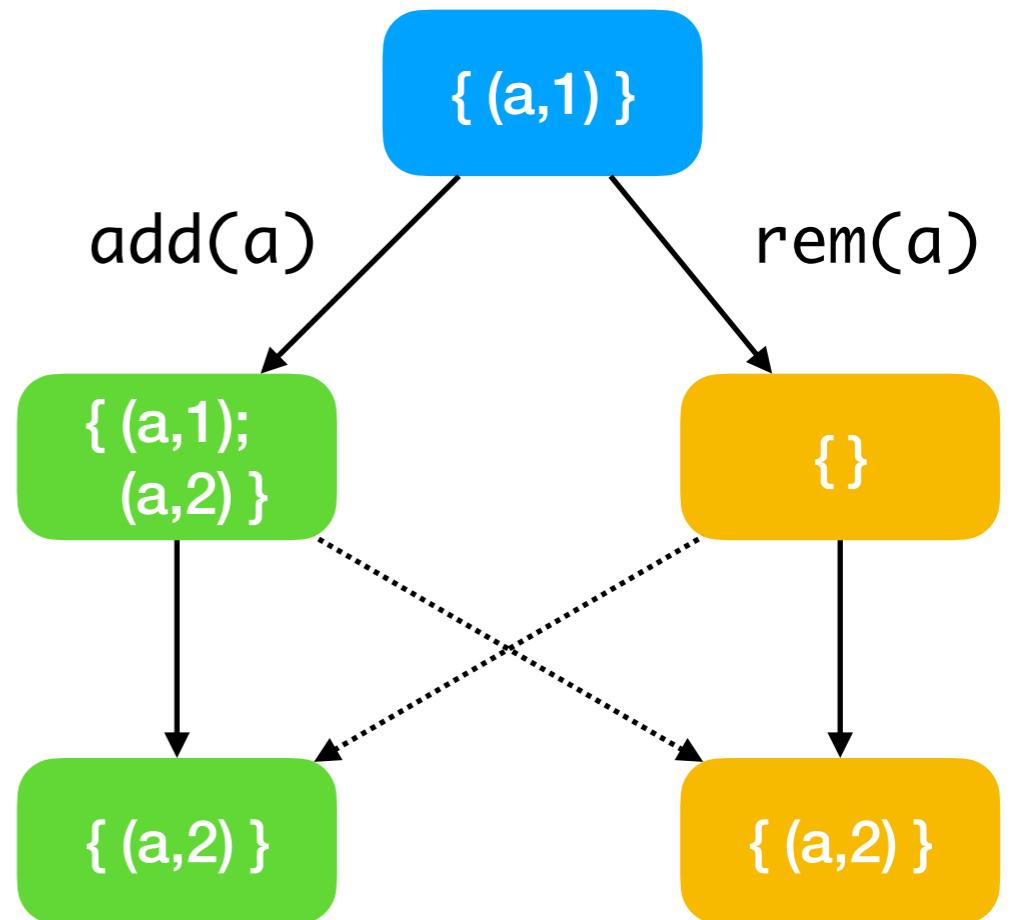
Fixing OR-Set

- Discriminate duplicate additions by associating a unique id

$$\begin{aligned} & \{ \} \\ & \cup (\{ (a,1); (a,2) \} - \{ (a,1) \}) \\ & \cup (\{ \} - \{ (a,1) \}) \\ & = \{ \} \cup \{ (a,2) \} \cup \{ \} \\ & = \{ (a,2) \} \end{aligned}$$

- MRDT implementation

$$D_\tau = (\Sigma, \sigma_0, do, merge)$$



Fixing OR-Set

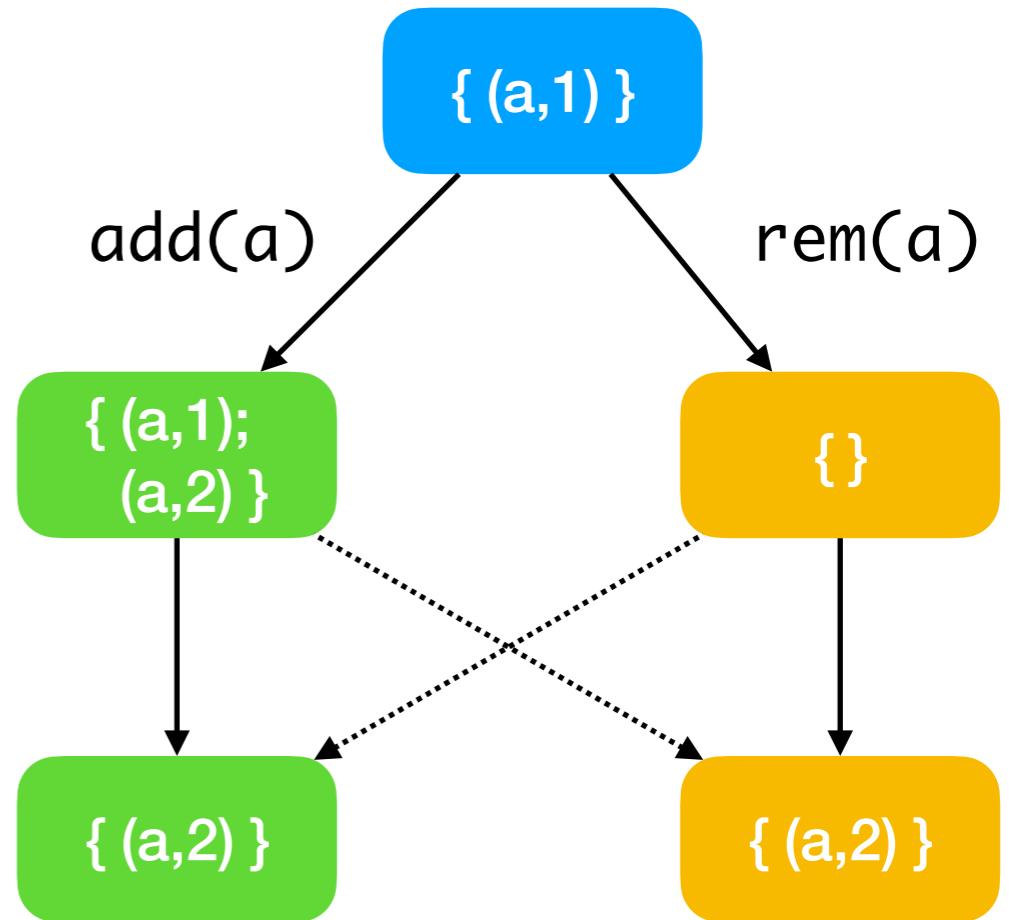
- Discriminate duplicate additions by associating a unique id

$$\begin{aligned}
 & \{ \} \\
 & \cup (\{ (a,1); (a,2) \} - \{ (a,1) \}) \\
 & \cup (\{ \} - \{ (a,1) \}) \\
 & = \{ \} \cup \{ (a,2) \} \cup \{ \} \\
 & = \{ (a,2) \}
 \end{aligned}$$

- MRDT implementation

$$D_\tau = (\Sigma, \sigma_0, do, merge)$$

- 1: $\Sigma = \mathcal{P}(\mathbb{N} \times \mathbb{N})$
- 2: $\sigma_0 = \{ \}$
- 3: $do(rd, \sigma, t) = (\sigma, \{a \mid (a, t) \in \sigma\})$
- 4: $do(add(a), \sigma, t) = (\sigma \cup \{(a, t)\}, \perp)$
- 5: $do(remove(a), \sigma, t) = (\{e \in \sigma \mid fst(e) \neq a\}, \perp)$
- 6: $merge(\sigma_{lca}, \sigma_a, \sigma_b) =$
 $(\sigma_{lca} \cap \sigma_a \cap \sigma_b) \cup (\sigma_a - \sigma_{lca}) \cup (\sigma_b - \sigma_{lca})$



Fixing OR-Set

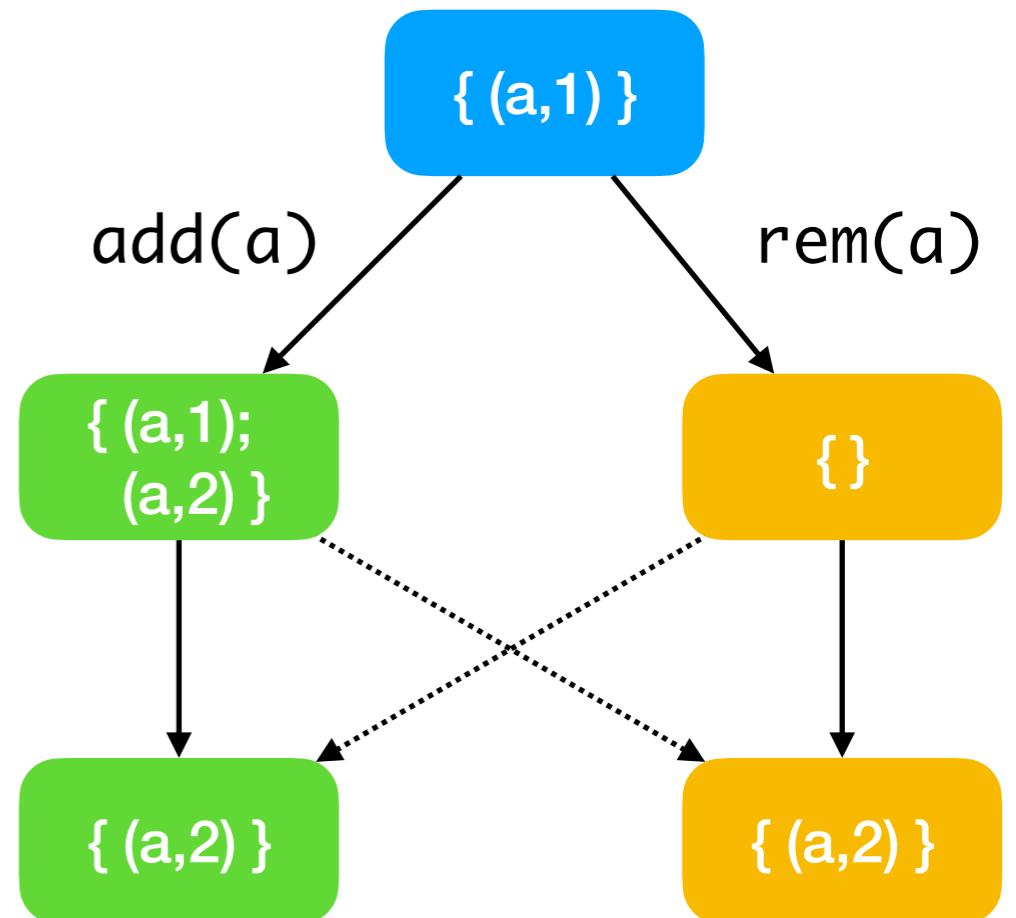
- Discriminate duplicate additions by associating a unique id

$$\begin{aligned}
 & \{ \} \\
 & \cup (\{ (a,1); (a,2) \} - \{ (a,1) \}) \\
 & \cup (\{ \} - \{ (a,1) \}) \\
 & = \{ \} \cup \{ (a,2) \} \cup \{ \} \\
 & = \{ (a,2) \}
 \end{aligned}$$

- MRDT implementation

$$D_\tau = (\Sigma, \sigma_0, do, merge)$$

- 1: $\Sigma = \mathcal{P}(\mathbb{N} \times \mathbb{N})$
 - 2: $\sigma_0 = \{ \}$
 - 3: $do(rd, \sigma, t) = (\sigma, \{a \mid (a, t) \in \sigma\})$
 - 4: $do(add(a), \sigma, t) = (\sigma \cup \{(a, t)\}, \perp)$
 - 5: $do(remove(a), \sigma, t) = (\{e \in \sigma \mid fst(e) \neq a\}, \perp)$
 - 6: $merge(\sigma_{lca}, \sigma_a, \sigma_b) =$
 $(\sigma_{lca} \cap \sigma_a \cap \sigma_b) \cup (\sigma_a - \sigma_{lca}) \cup (\sigma_b - \sigma_{lca})$
- Unique Lamport Timestamps*

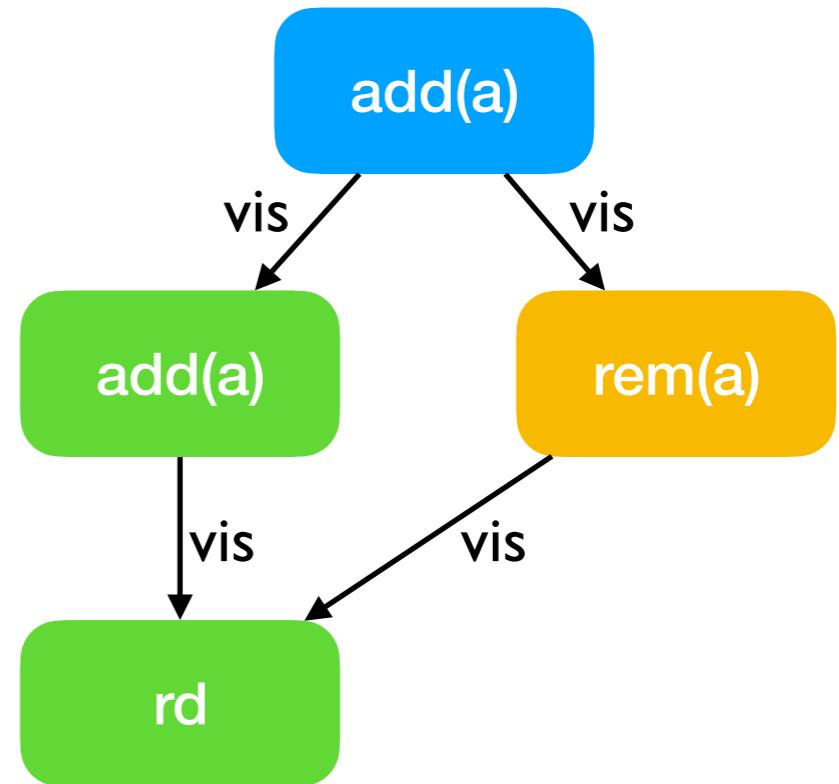
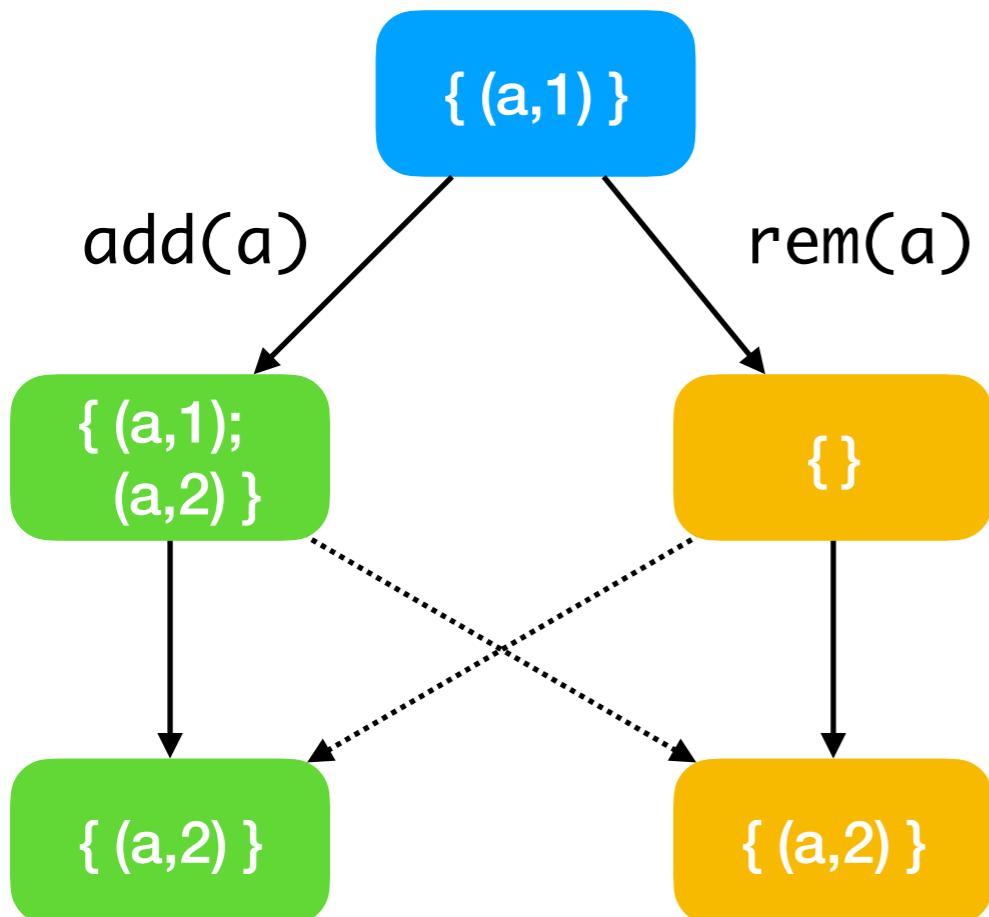


Specifying OR-Set

Abstract state $I = \langle E, oper, rval, time, vis \rangle$

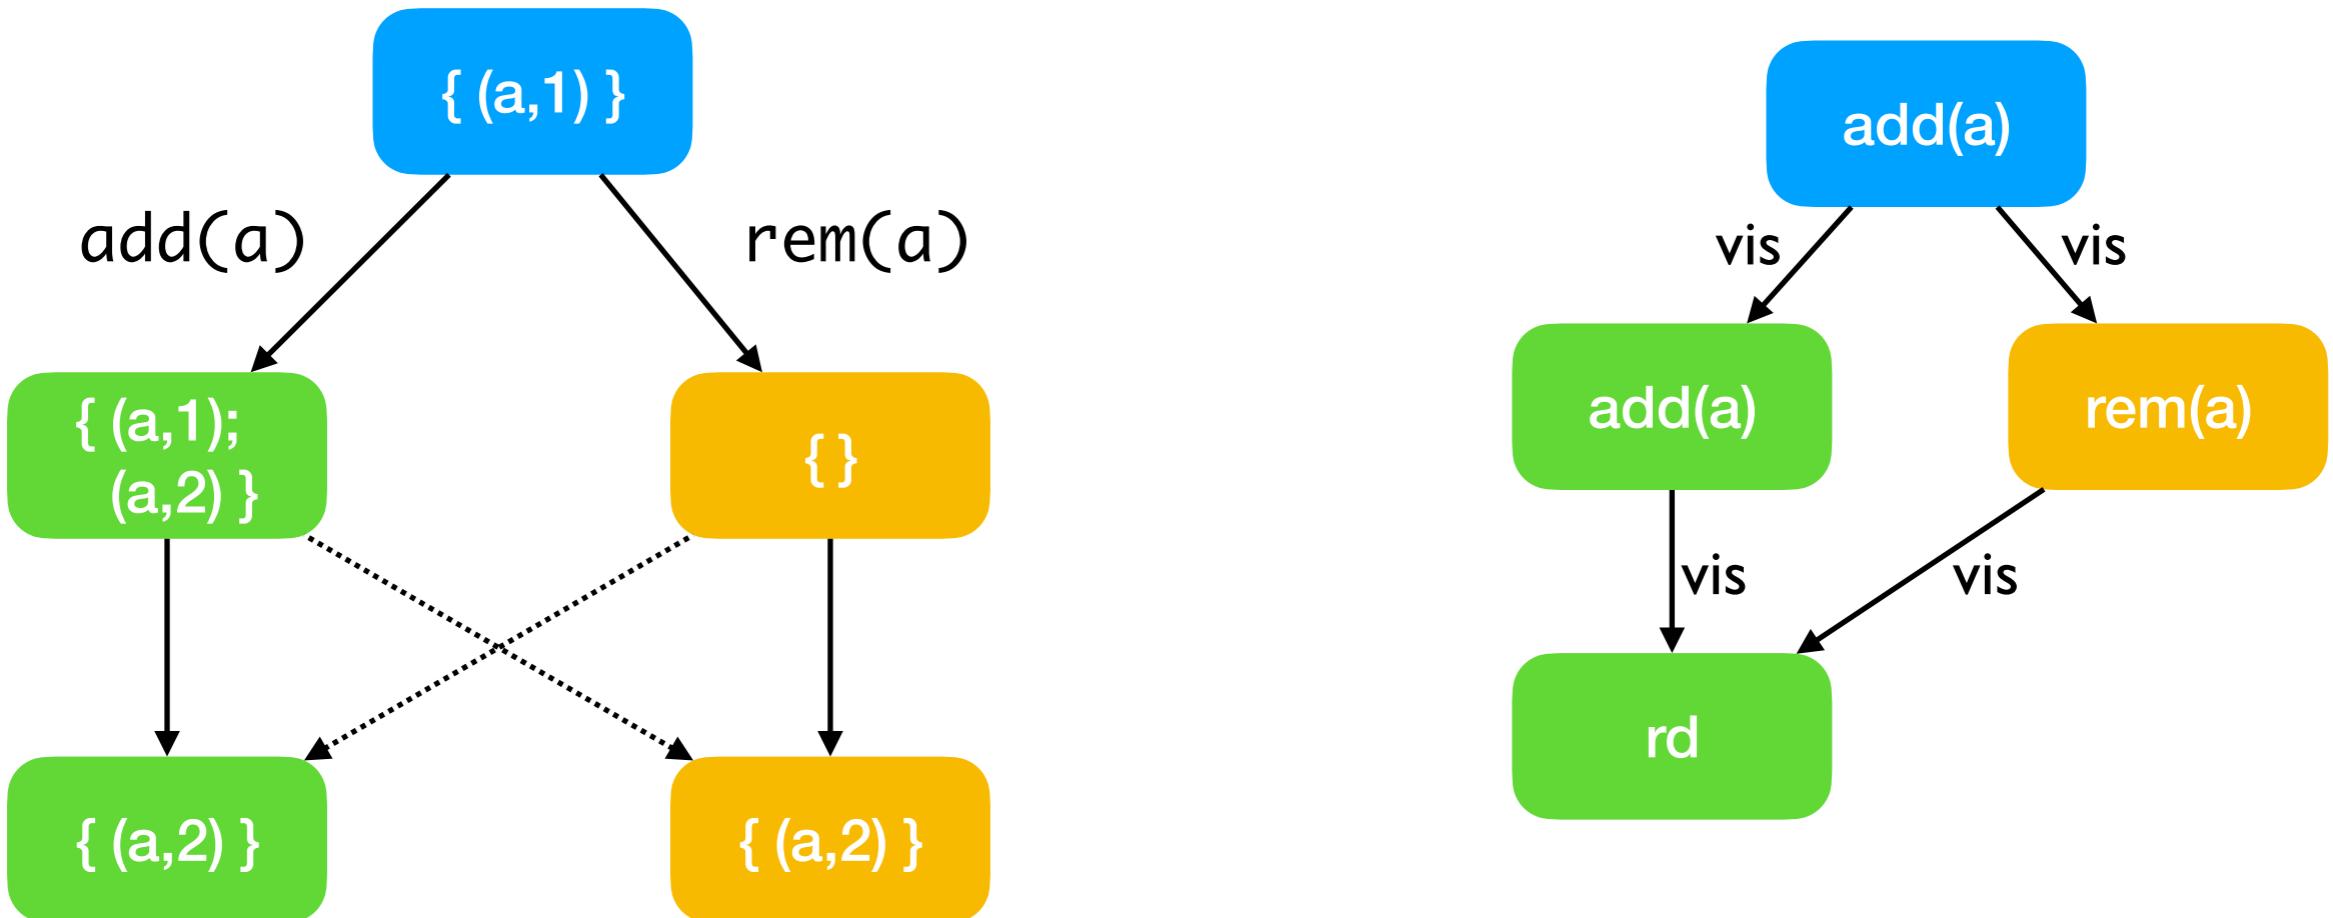
Specifying OR-Set

Abstract state $I = \langle E, oper, rval, time, vis \rangle$



Specifying OR-Set

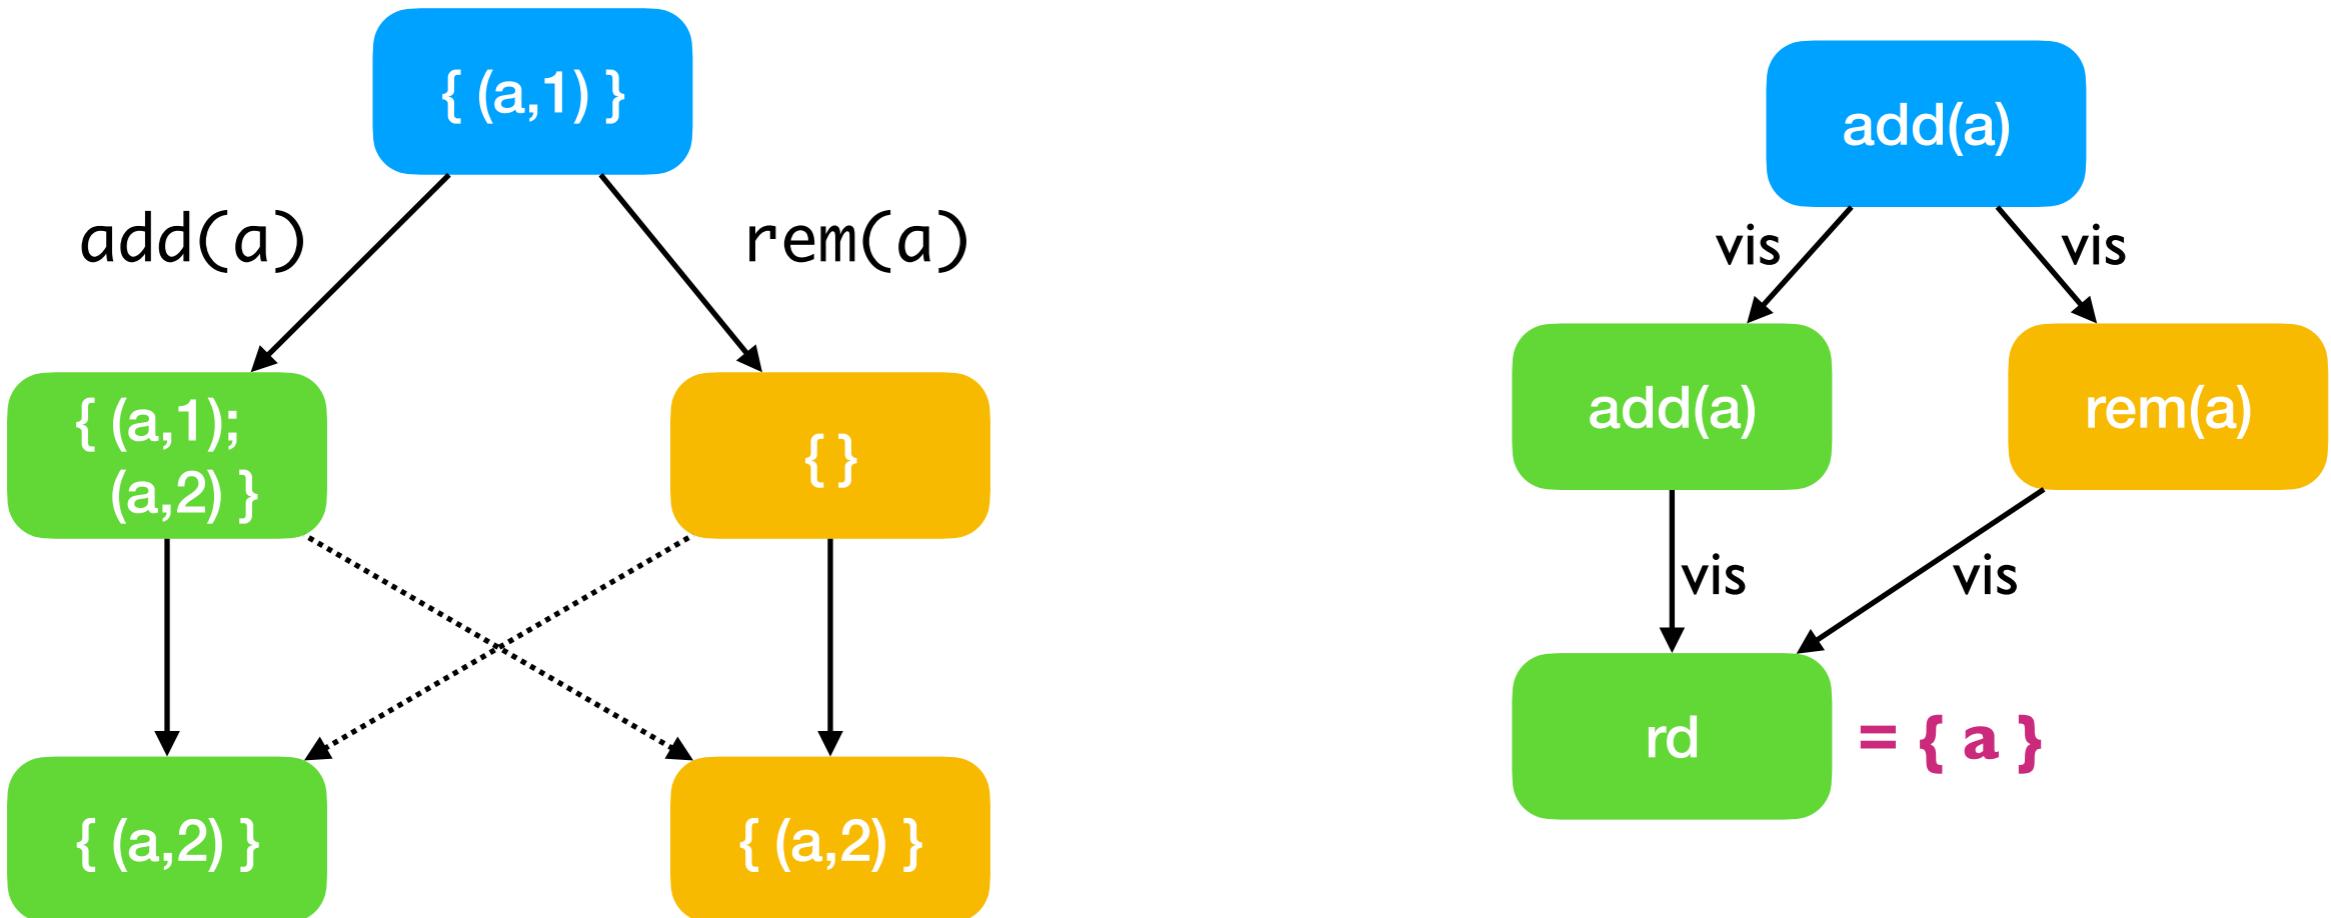
Abstract state $I = \langle E, oper, rval, time, vis \rangle$



$$\begin{aligned}
 \mathcal{F}_{orset}(\text{rd}, \langle E, oper, rval, time, vis \rangle) &= \{a \mid \exists e \in E. oper(e) \\
 &= \text{add}(a) \wedge \neg(\exists f \in E. oper(f) = \text{remove}(a) \wedge e \xrightarrow{vis} f)\}
 \end{aligned}$$

Specifying OR-Set

Abstract state $I = \langle E, oper, rval, time, vis \rangle$



$$\begin{aligned}
 \mathcal{F}_{orset}(\text{rd}, \langle E, oper, rval, time, vis \rangle) &= \{a \mid \exists e \in E. \text{oper}(e) \\
 &= \text{add}(a) \wedge \neg(\exists f \in E. \text{oper}(f) = \text{remove}(a) \wedge e \xrightarrow{vis} f)\}
 \end{aligned}$$

Simulation Relation

Simulation Relation

- Connects the *abstract state* with the *concrete state*

Simulation Relation

- Connects the *abstract state* with the *concrete state*
- For the OR-set,

Simulation Relation

- Connects the *abstract state* with the *concrete state*
- For the OR-set,

$$\begin{aligned}\mathcal{R}_{sim}(I, \sigma) \iff (\forall(a, t) \in \sigma \iff \\ (\exists e \in I.E \wedge I.\text{oper}(e) = \text{add}(a) \wedge I.\text{time}(e) = t \wedge \\ \neg(\exists f \in I.E \wedge I.\text{oper}(f) = \text{remove}(a) \wedge e \xrightarrow{\text{vis}} f)))\end{aligned}$$

Simulation Relation

- Connects the *abstract state* with the *concrete state*
- For the OR-set,

$$\begin{aligned}\mathcal{R}_{sim}(I, \sigma) \iff (\forall(a, t) \in \sigma \iff \\ (\exists e \in I.E \wedge I.\text{oper}(e) = \text{add}(a) \wedge I.\text{time}(e) = t \wedge \\ \neg(\exists f \in I.E \wedge I.\text{oper}(f) = \text{remove}(a) \wedge e \xrightarrow{\text{vis}} f)))\end{aligned}$$

- The main verification effort is to show that the relation above is indeed a *simulation relation*
 - ★ Shown separately for *operations* and *merge function*
 - ★ Proof by induction on the execution trace

Verification effort

MRDTs verified	#Lines code	#Lines proof	#Lemmas	Verif. time (s)
Increment-only counter	6	43	2	3.494
PN counter	8	43	2	23.211
Enable-wins flag	20	58 81 89	3 6 7	1074 171 104
LWW register	5	44	1	4.21
G-set	10	23 28 33	0 1 2	4.71 2.462 1.993
G-map	48	26	0	26.089
Mergeable log	39	95	2	36.562
OR-set (§2.1.1)	30	36 41 46	0 1 2	43.85 21.656 8.829
OR-set-space (§2.1.2)	59	108	7	1716
OR-set-spacetime	97	266	7	1854
Queue	32	1123	75	4753

Verification effort

MRDTs verified	#Lines code	#Lines proof	#Lemmas	Verif. time (s)
Increment-only counter	6	43	2	3.494
PN counter	8	43	2	23.211
Enable-wins flag	20	58 81 89	3 6 7	1074 171 104
LWW register	5	44	1	4.21
G-set	10	23 28 33	0 1 2	4.71 2.462 1.993
G-map	48	26	0	26.089
Mergeable log	39	95	2	36.562
OR-set (§2.1.1)	30	36 41 46	0 1 2	43.85 21.656 8.829
OR-set-space (§2.1.2)	59	108	7	1716
OR-set-spacetime	97	266	7	1854
Queue	32	1123	75	4753

Composing RDTs is HARD!



Martin Kleppmann
@martinkl

...

Today in “distributed systems are hard”: I wrote down a simple CRDT algorithm that I thought was “obviously correct” for a course I’m teaching. Only 10 lines or so long. Found a fatal bug only after spending hours trying to prove the algorithm correct. 😭

4:18 AM · Nov 13, 2020 · Tweetbot for iOS

41 Retweets 4 Quote Tweets 541 Likes



Martin Kleppmann @martinkl · Nov 13, 2020

...

The interesting thing about this bug is that it comes about only from the interaction of two features. A LWW map by itself is fine. A set in which you can insert and delete elements (but not update them) is fine. The problem arises only when delete and update interact.



1

16



Composing IRC-style chat

- Build IRC-style group chat
 - ★ Send and read messages in channels

Composing IRC-style chat

- Build IRC-style group chat
 - ★ Send and read messages in channels
- Represent application state as a map MRDT
 - ★ String (*channel name*) keys → mergeable-log MRDT values

Composing IRC-style chat

- Build IRC-style group chat
 - ★ Send and read messages in channels
- Represent application state as a **map** MRDT
 - ★ String (*channel name*) keys → **mergeable-log** MRDT values
- **Goal:**
 - ★ **map** and **log** proved correct separately
 - ★ Use the proof of underlying RDTs to prove chat application correctness

Generic Map MRDT

Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$ where

- 1: $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$
- 2: $\sigma_0 = \{\}$
- 3: $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_{0_\alpha}, & \text{otherwise} \end{cases}$
- 4: $do(\text{set}(k, o_\alpha), \sigma, t) =$
 let $(v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$ in $(\sigma[k \mapsto v], r)$
- 5: $do(\text{get}(k, o_\alpha), \sigma, t) =$
 let $(_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$ in (σ, r)
- 6: $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$
 $\{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$
 $v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

- 1: $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{oper}(e) = \text{set}(k, _)) \wedge$
- 2: $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

Generic Map MRDT

Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$ where

- 1: $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$ The values in the MRDT map are MRDTs
- 2: $\sigma_0 = \{\}$
- 3: $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_{0_\alpha}, & \text{otherwise} \end{cases}$
- 4: $do(\text{set}(k, o_\alpha), \sigma, t) =$
let $(v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$ in $(\sigma[k \mapsto v], r)$
- 5: $do(\text{get}(k, o_\alpha), \sigma, t) =$
let $(_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$ in (σ, r)
- 6: $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$
 $\{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$
 $v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

- 1: $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{oper}(e) = \text{set}(k, _)) \wedge$
- 2: $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

Generic Map MRDT

Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$ where

1: $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$

2: $\sigma_0 = \{\}$

3: $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_0, & \text{otherwise} \end{cases}$

4: $do(\text{set}(k, o_\alpha), \sigma, t) =$
 $\quad \text{let } (v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma[k \mapsto v], r)$

5: $do(\text{get}(k, o_\alpha), \sigma, t) =$
 $\quad \text{let } (_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t) \text{ in } (\sigma, r)$

6: $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$
 $\quad \{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$
 $\quad \quad v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

The values in the MRDT map are MRDTs

Merge uses the merge of the underlying value type!

Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

1: $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{ oper}(e) = \text{set}(k, _)) \wedge$

2: $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

Generic Map MRDT

Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$ where

1: $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$

2: $\sigma_0 = \{\}$

3: $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_0, & \text{otherwise} \end{cases}$

4: $do(\text{set}(k, o_\alpha), \sigma, t) =$
let $(v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$ in $(\sigma[k \mapsto v], r)$

5: $do(\text{get}(k, o_\alpha), \sigma, t) =$
let $(_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$ in (σ, r)

6: $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$
 $\{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$
 $v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

The values in the MRDT map are MRDTs

Merge uses the merge of the underlying value type!

Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

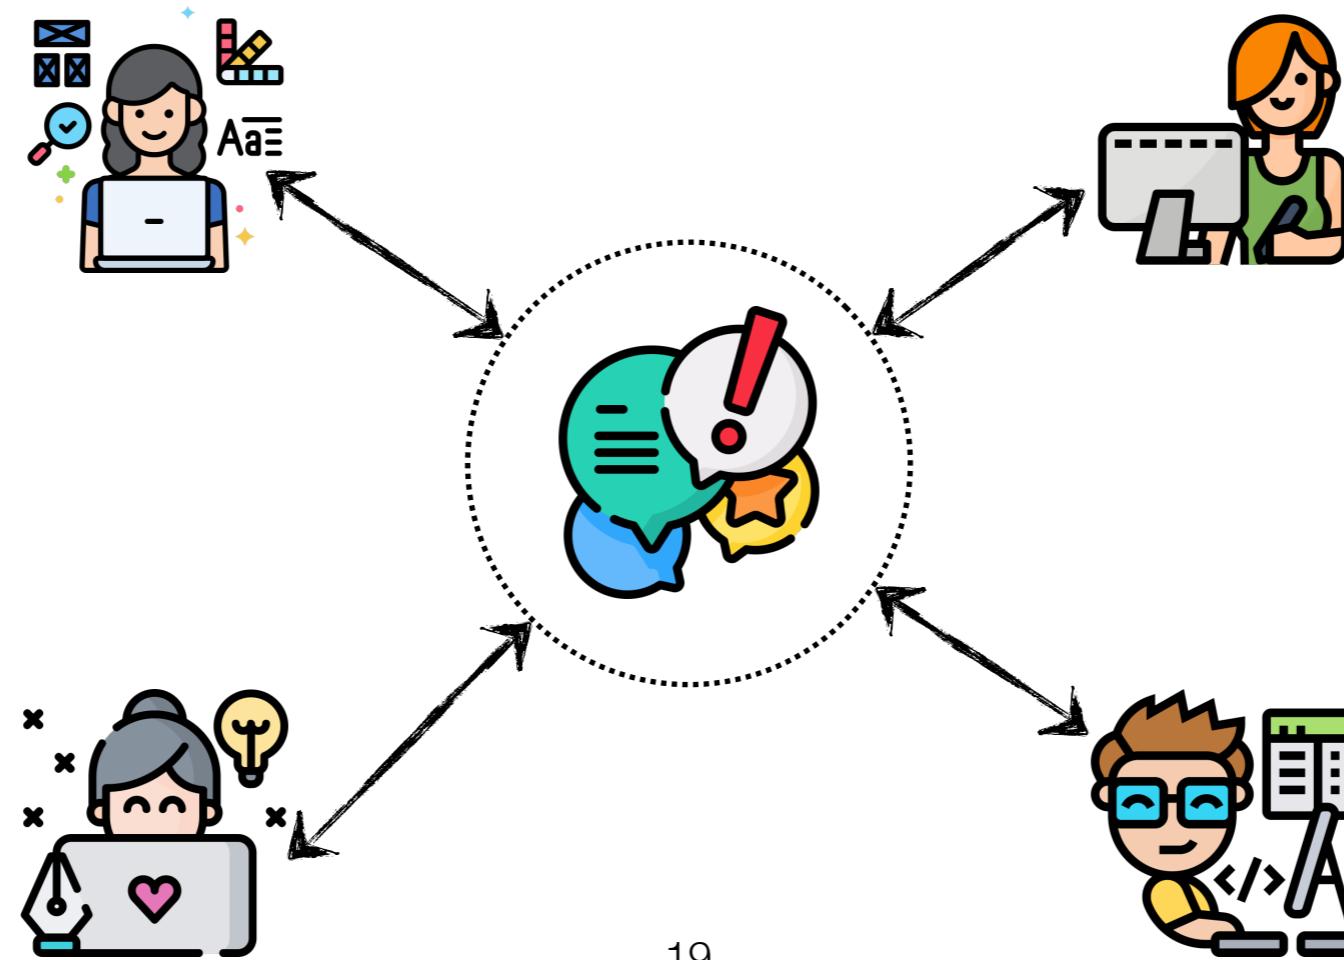
1: $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{oper}(e) = \text{set}(k, _)) \wedge$

2: $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$

Simulation relation appeals to the value type's simulation relation!

Composing IRC-style chat

- IRC app state is constructed by instantiating *generic map* with *mergeable log*
- The proof of correctness of the chat application directly follows from the composition.
 - ★ See paper for details!



Summary

- **Peepul**
 - ◆ An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>



Summary

- **Peepul**
 - ◆ An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>
- *Space- and time-efficient* implementations
 - ★ Certified implementation of a $O(1)$ replicated queue with $O(n)$ merge.



Summary

- **Peepul**
 - ◆ An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>
- *Space- and time-efficient* implementations
 - ★ Certified implementation of a $O(1)$ replicated queue with $O(n)$ merge.
- *Composition* of MRDTs and their proofs!



Summary

- **Peepul**
 - ◆ An F* library implementing and proving MRDTs
 - ★ <https://github.com/prismlab/peepul>
- *Space- and time-efficient* implementations
 - ★ Certified implementation of a $O(1)$ replicated queue with $O(n)$ merge.
- *Composition* of MRDTs and their proofs!
- See paper for
 - ◆ Formal description of the system + soundness proof
 - ◆ Case study on replicated queues
 - ◆ Performance results

