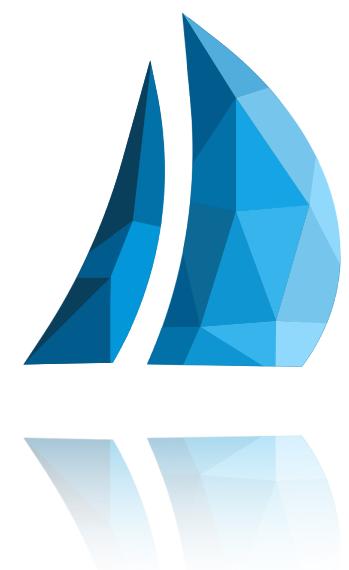


# Evolving the OCaml programming language

KC Sivaramakrishnan  
[kcsrk.info](http://kcsrk.info)

Krea University  
4th November 2025



Tarides  
SADAM

# Who am I – KC Sivaramakrishnan

- CS Prof at IIT Madras
  - Programming languages, formal verification and systems
- A core maintainer of the *OCaml* programming language
- CTO at Tarides
  - Building functional systems using *OCaml*
  - Maintainers of the OCaml compiler and platform tools

- Turbo C++ IDE
- Learnt to program C here
- Believed the C language was "perfect & final"
- ...like mountains and oceans
- Grew up and realised neither was!
- This talk is about the evolution of programming languages
- Specifically, OCaml



## Language



- Functional-first but multi-paradigm (imperative, OO)
- Static-type system with Hindley-Milner type inference
- Advanced features – powerful module system, GADTs, Polymorphic variants
- Multicore support and *effect handlers*



## Platform

- Fast, native code – x86, ARM, RISC-V, etc.
- JavaScript and WebAssembly (using *WasmGC*) compilation
- Platform tools – editor (LSP), build system (dune), package manager (opam), docs generator (odoc), etc.

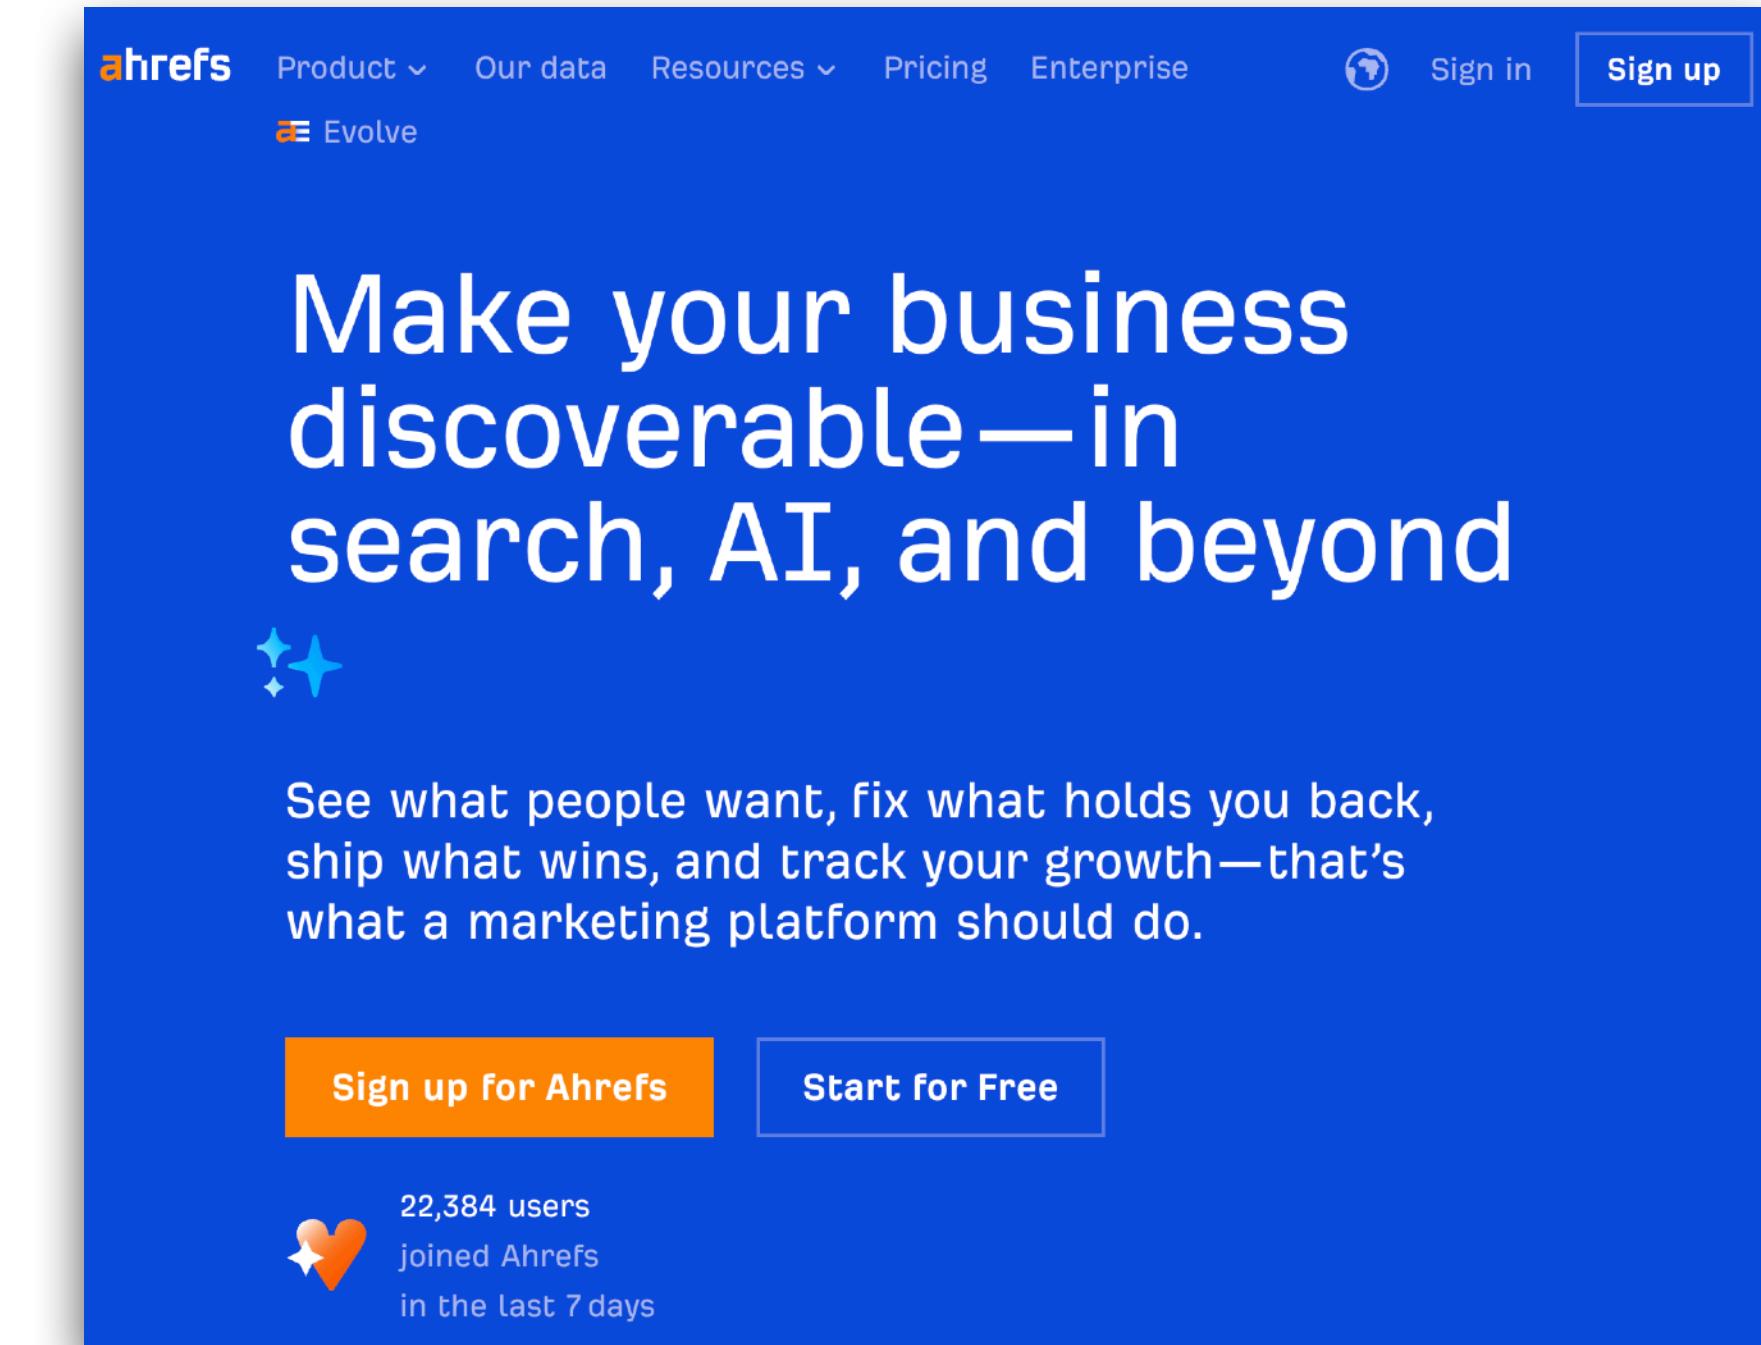
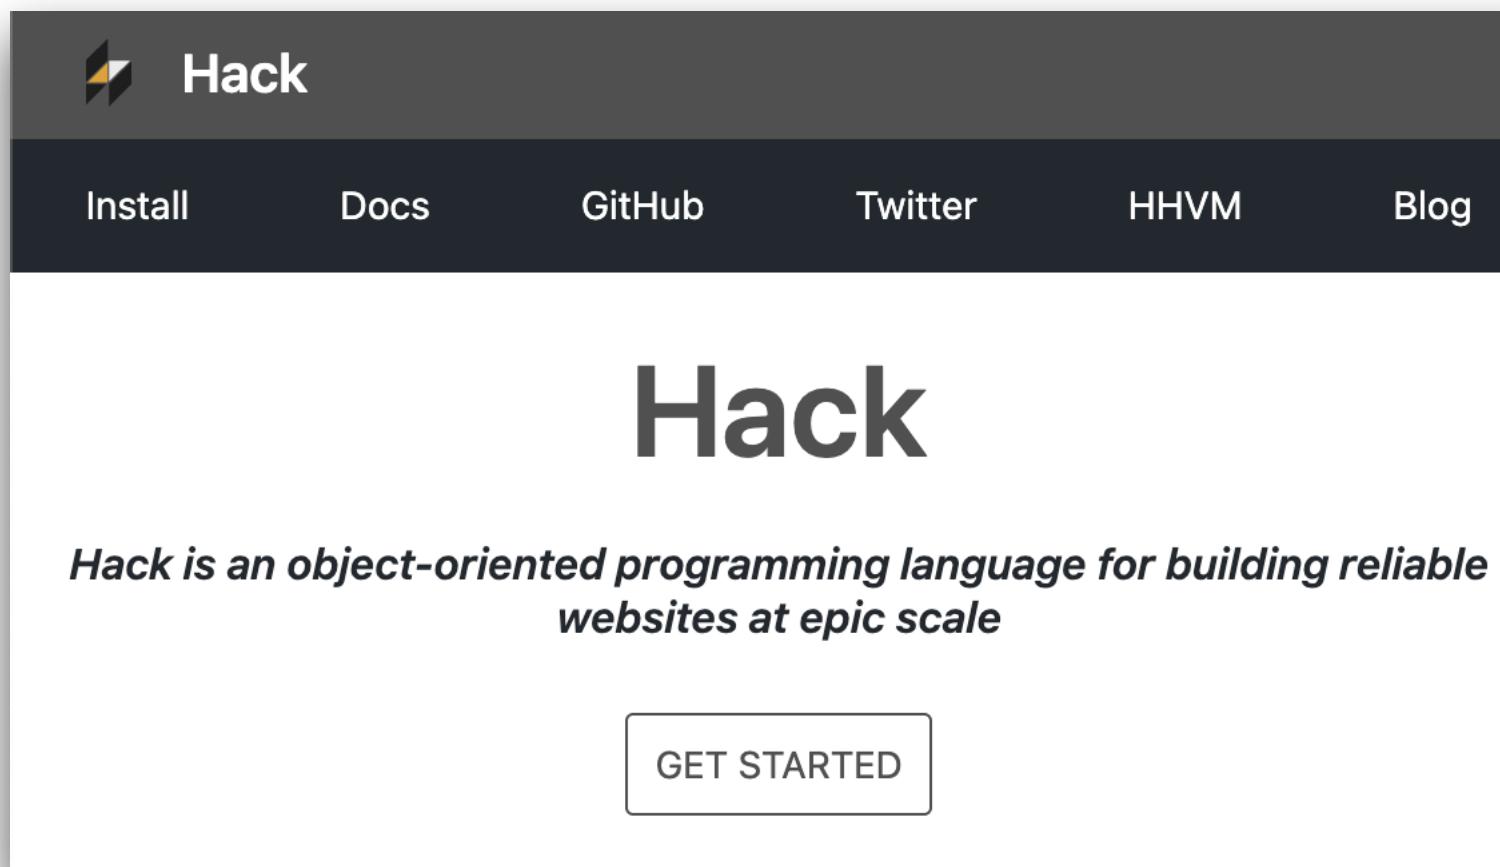


## Ecosystem

- Opam repository – small but mature package ecosystem
- Notable Industrial users – Jane Street, Meta, Microsoft, Ahrefs, Citrix, Tezos, Bloomberg, Docker

# High dynamic range

*From scripts to scalable systems, research prototypes to production infrastructure*



***Compilers***

***Web Frontend***

# High dynamic range

*From scripts to scalable systems, research prototypes to production infrastructure*

## Functional Networking for Millions of Docker Desktops (Experience Report)

ANIL MADHAVAPEDDY, University of Cambridge, United Kingdom

DAVID J. SCOTT, Docker, Inc., United Kingdom

PATRICK FERRIS, University of Cambridge, United Kingdom

RYAN T. GIBB, University of Cambridge, United Kingdom

THOMAS GAZAGNAIRE, Tarides, France



Docker is a developer tool used by millions of developers to build, share and run software stacks. The Docker Desktop clients for Mac and Windows have long used a novel combination of virtualisation and OCaml unikernels to seamlessly run Linux containers on these non-Linux hosts. We reflect on a decade of shipping this functional OCaml code into production across hundreds of millions of developer desktops, and discuss the lessons learnt from our experiences in integrating OCaml deeply into the container architecture that now drives much of the global cloud. We conclude by observing just how good a fit for systems programming that the unikernel approach has been, particularly when combined with the OCaml module and type system.

CCS Concepts: • Software and its engineering → *Software system structures; Functional languages;* • Computer systems organization → *Cloud computing.*

## OCaml in Space 🚀



***Virtualisation and Networking***

# High dynamic range

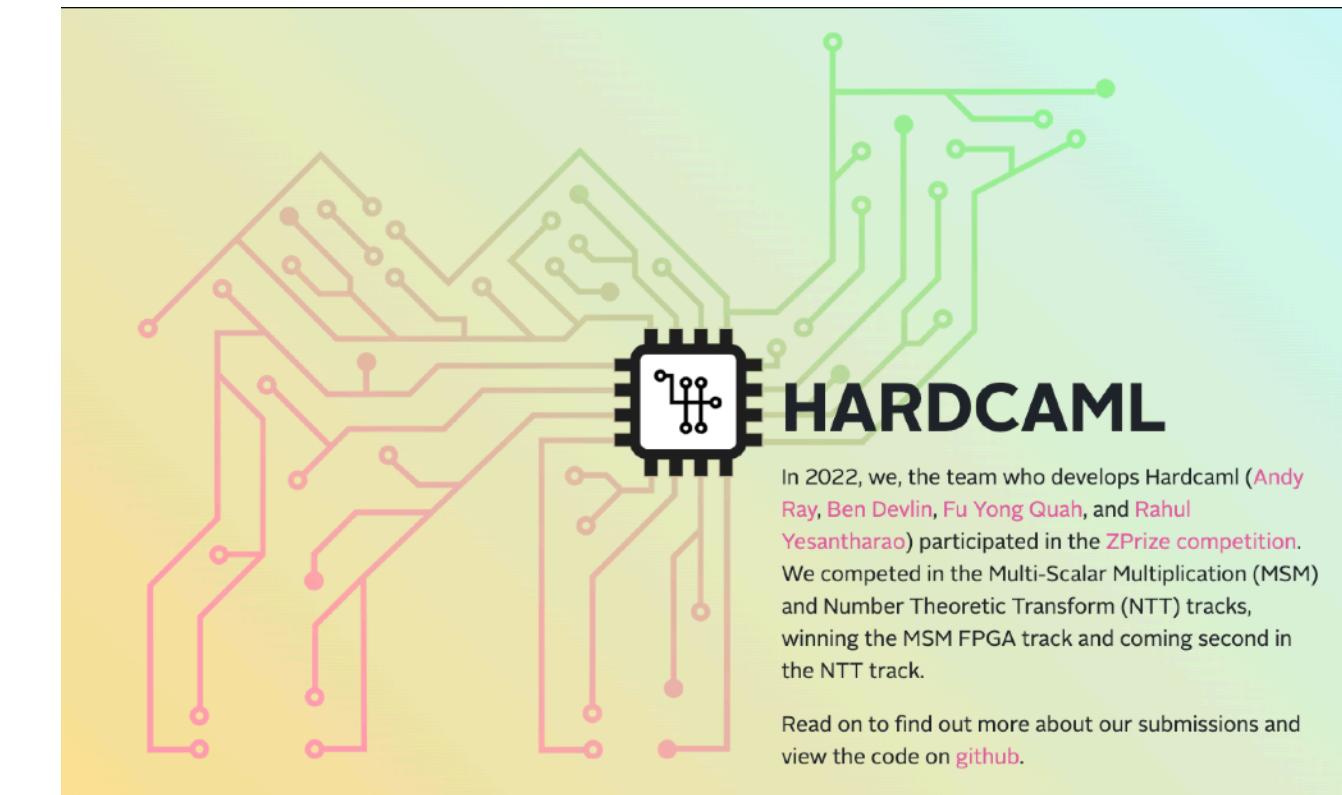
*From scripts to scalable systems, research prototypes to production infrastructure*



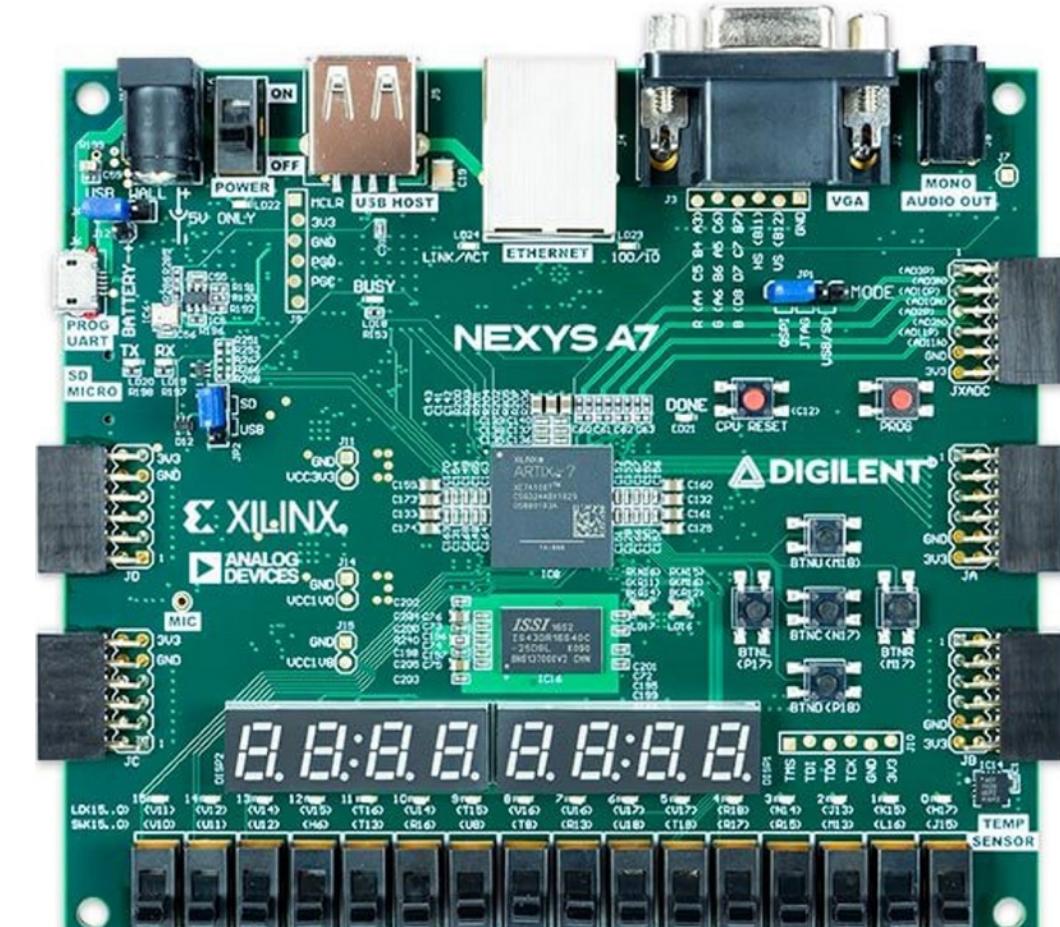
Bloomberg

*Finance*

60+M lines of OCaml code!



**Hardware  
design**





29 years old!



Steady evolution  
over **50+** years

## 1973 – Robin Milner’s “ML” for LCF

*Type system, type inference*



## 1985 – Guy Cousineau & co’s CAML

*Categorical abstract machine (CAM) as IR*



## 1996 – OCaml 1.0

*Object system, low-latency GC, fast native backend, module system*



## 2012 – OCaml 4.0

*Generalized Algebraic Data types (GADTs)*



## 2022 – OCaml 5.0

*Multicore parallelism, effect handlers*



2025

**How to *thrive* not just  
*survive* after ~30 years?**

# *Simplicity* and *stability*

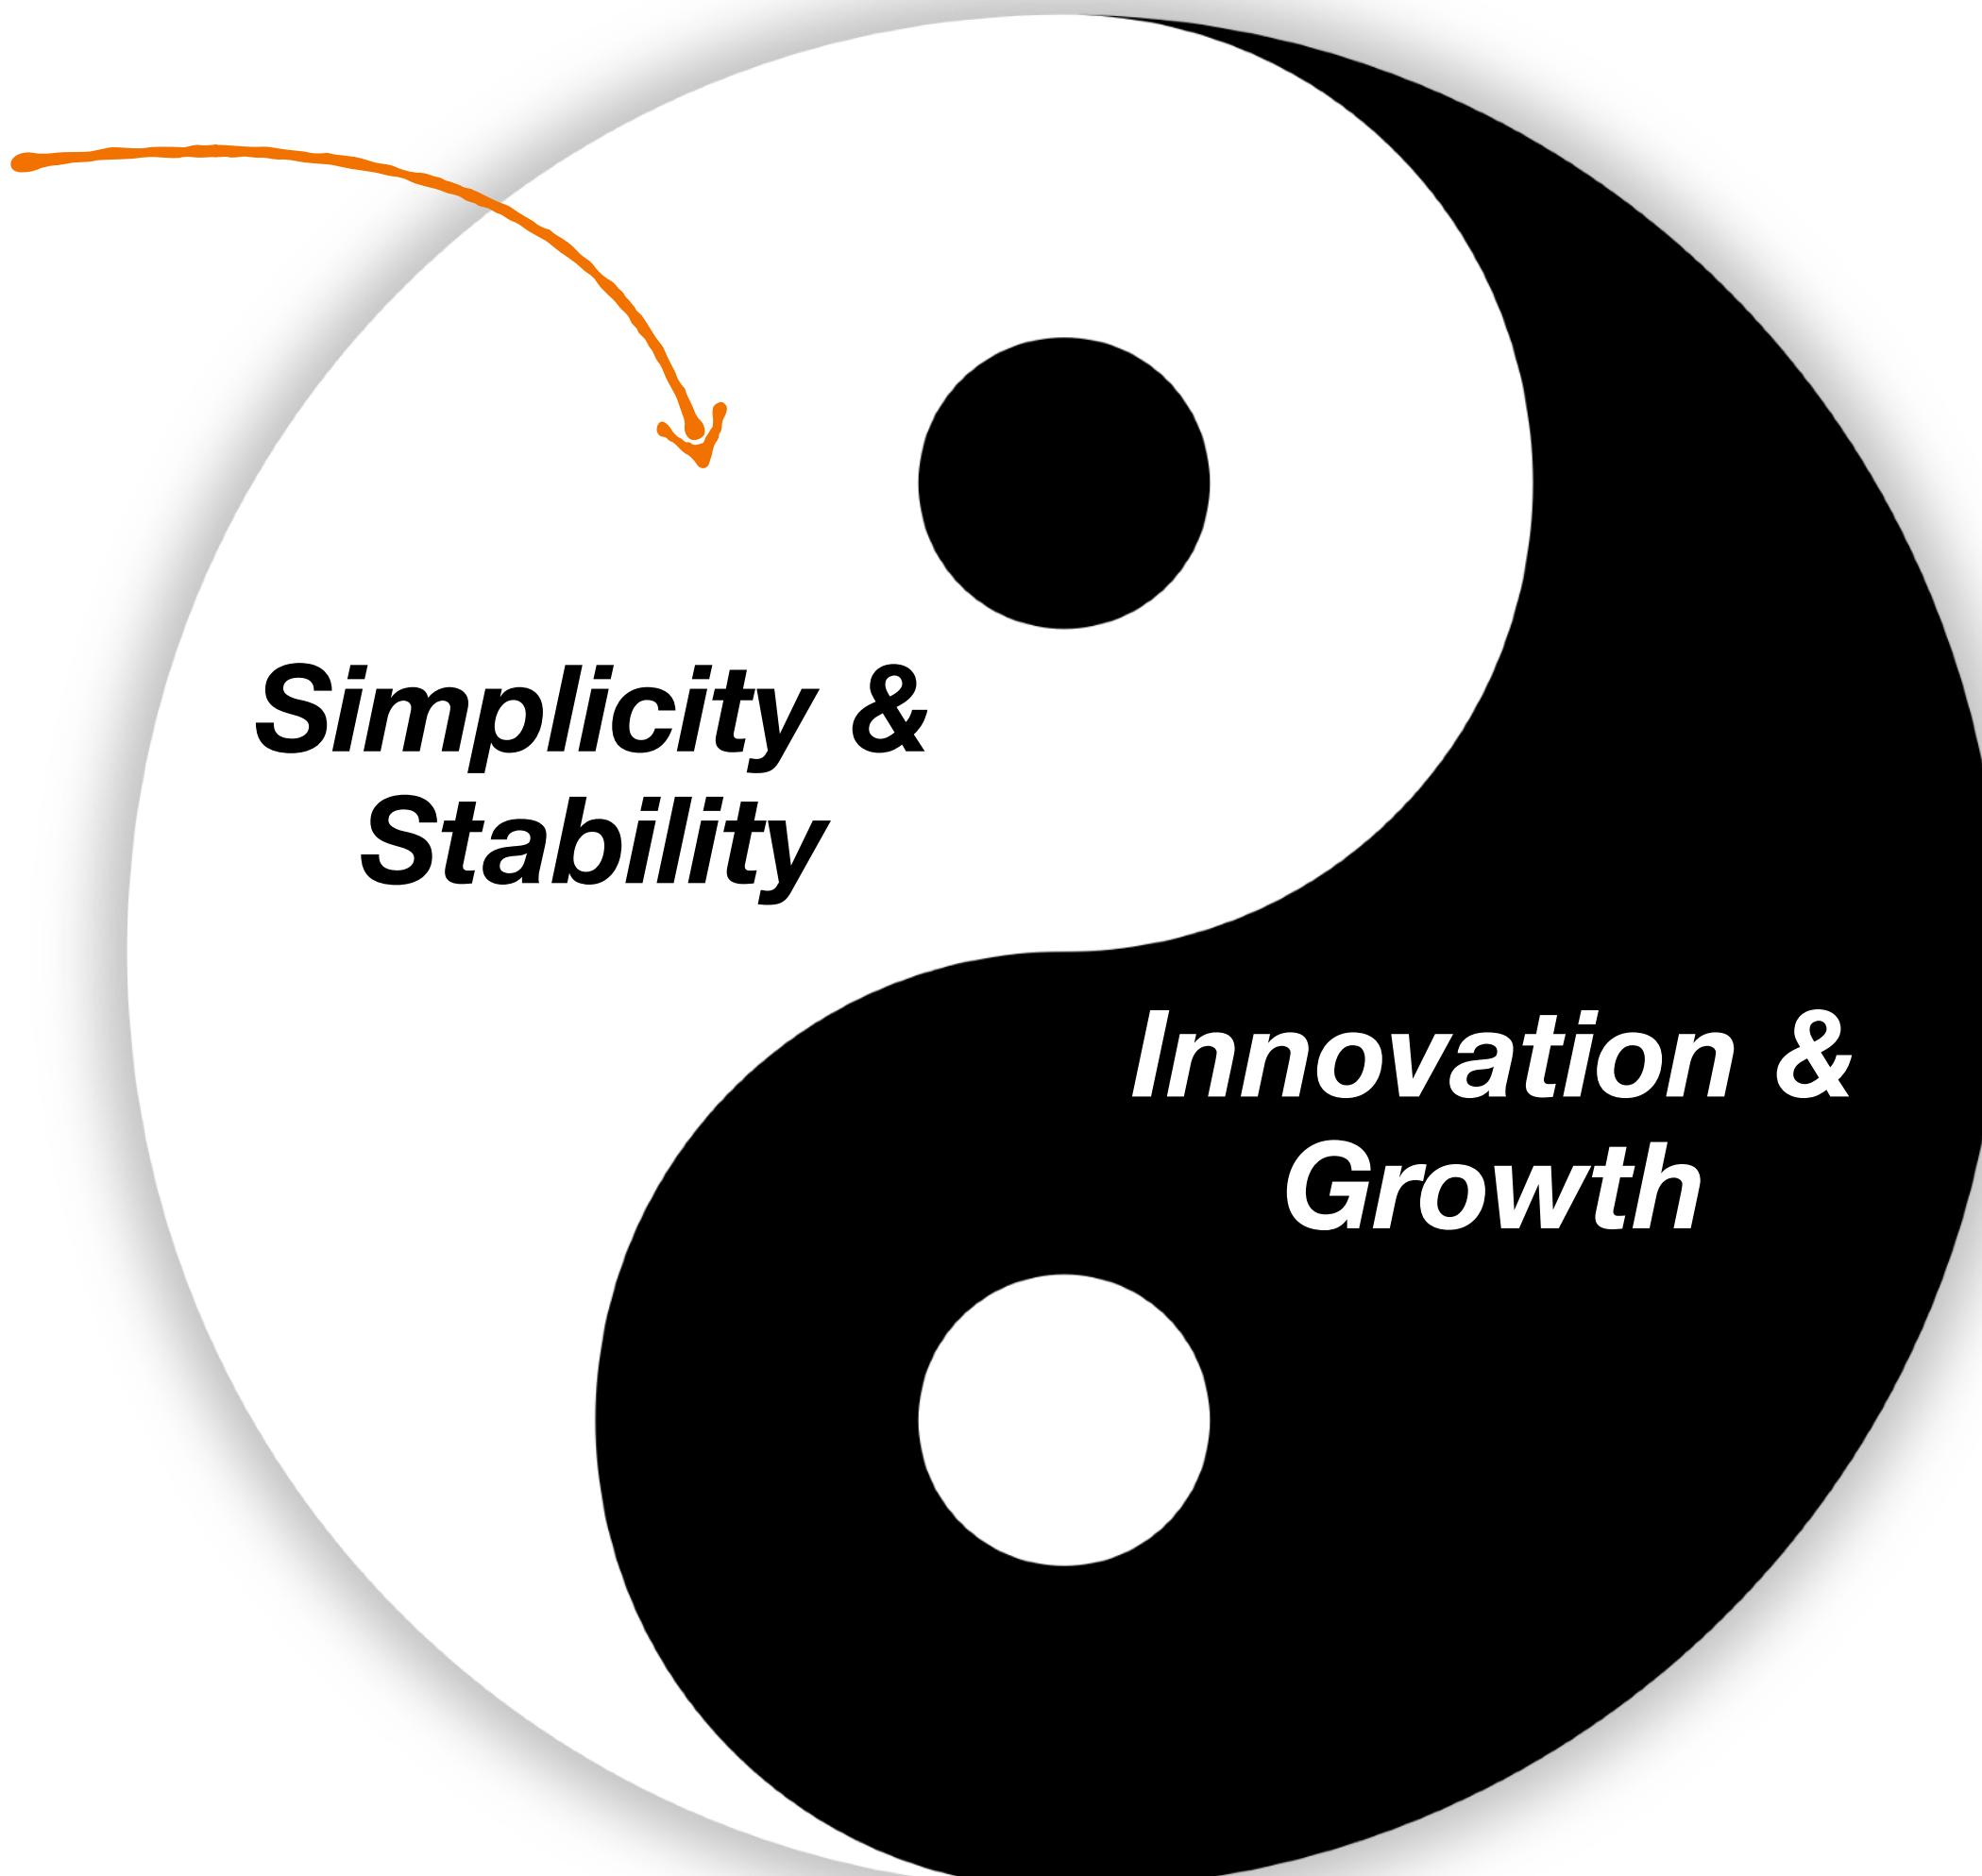
Xavier Leroy, 2023 SIGPLAN programming languages software award! 

What made that possible? Not just fancy types and nice modules – even though systems programmers value type safety and modularity highly – but also basic properties of OCaml:

- a language with a simple cost model, where it's easy to track how much time and how much space is used;
- a compiler that produces efficient code that looks like the source code, with only predictable optimizations;
- a low-latency garbage collector, usable for soft real-time applications.

- If you take OCaml from 20 years ago, the code will likely **continue to work!**
- No recent releases for some popular packages
  - They are **good enough**, and continue to be so.
  - Nothing to be done to keep it working!

OCaml



# OCaml Maintainers

Abigael	Richard Eisenberg	Nicolás Ojeda Bär
Alain Frisch	Jacques-Henri Jourdan	Florian Angeletti
Armaël Guéneau	KC Sivaramakrishnan	Olivier Nicole
Anil Madhavapeddy	Frédéric Bour	Sadiq Jaffer
Pierre Chambart	Leo White	Sébastien Hinderer
Damien Doligez	Vincent Lviron	Stephen Dolan
David Allsopp	Luc Maranget	Thomas Refis
Jacques Garrigue	Mark Shinwell	Xavier Leroy
Gabriel Scherer	Nick Barnes	Jeremy Yallop

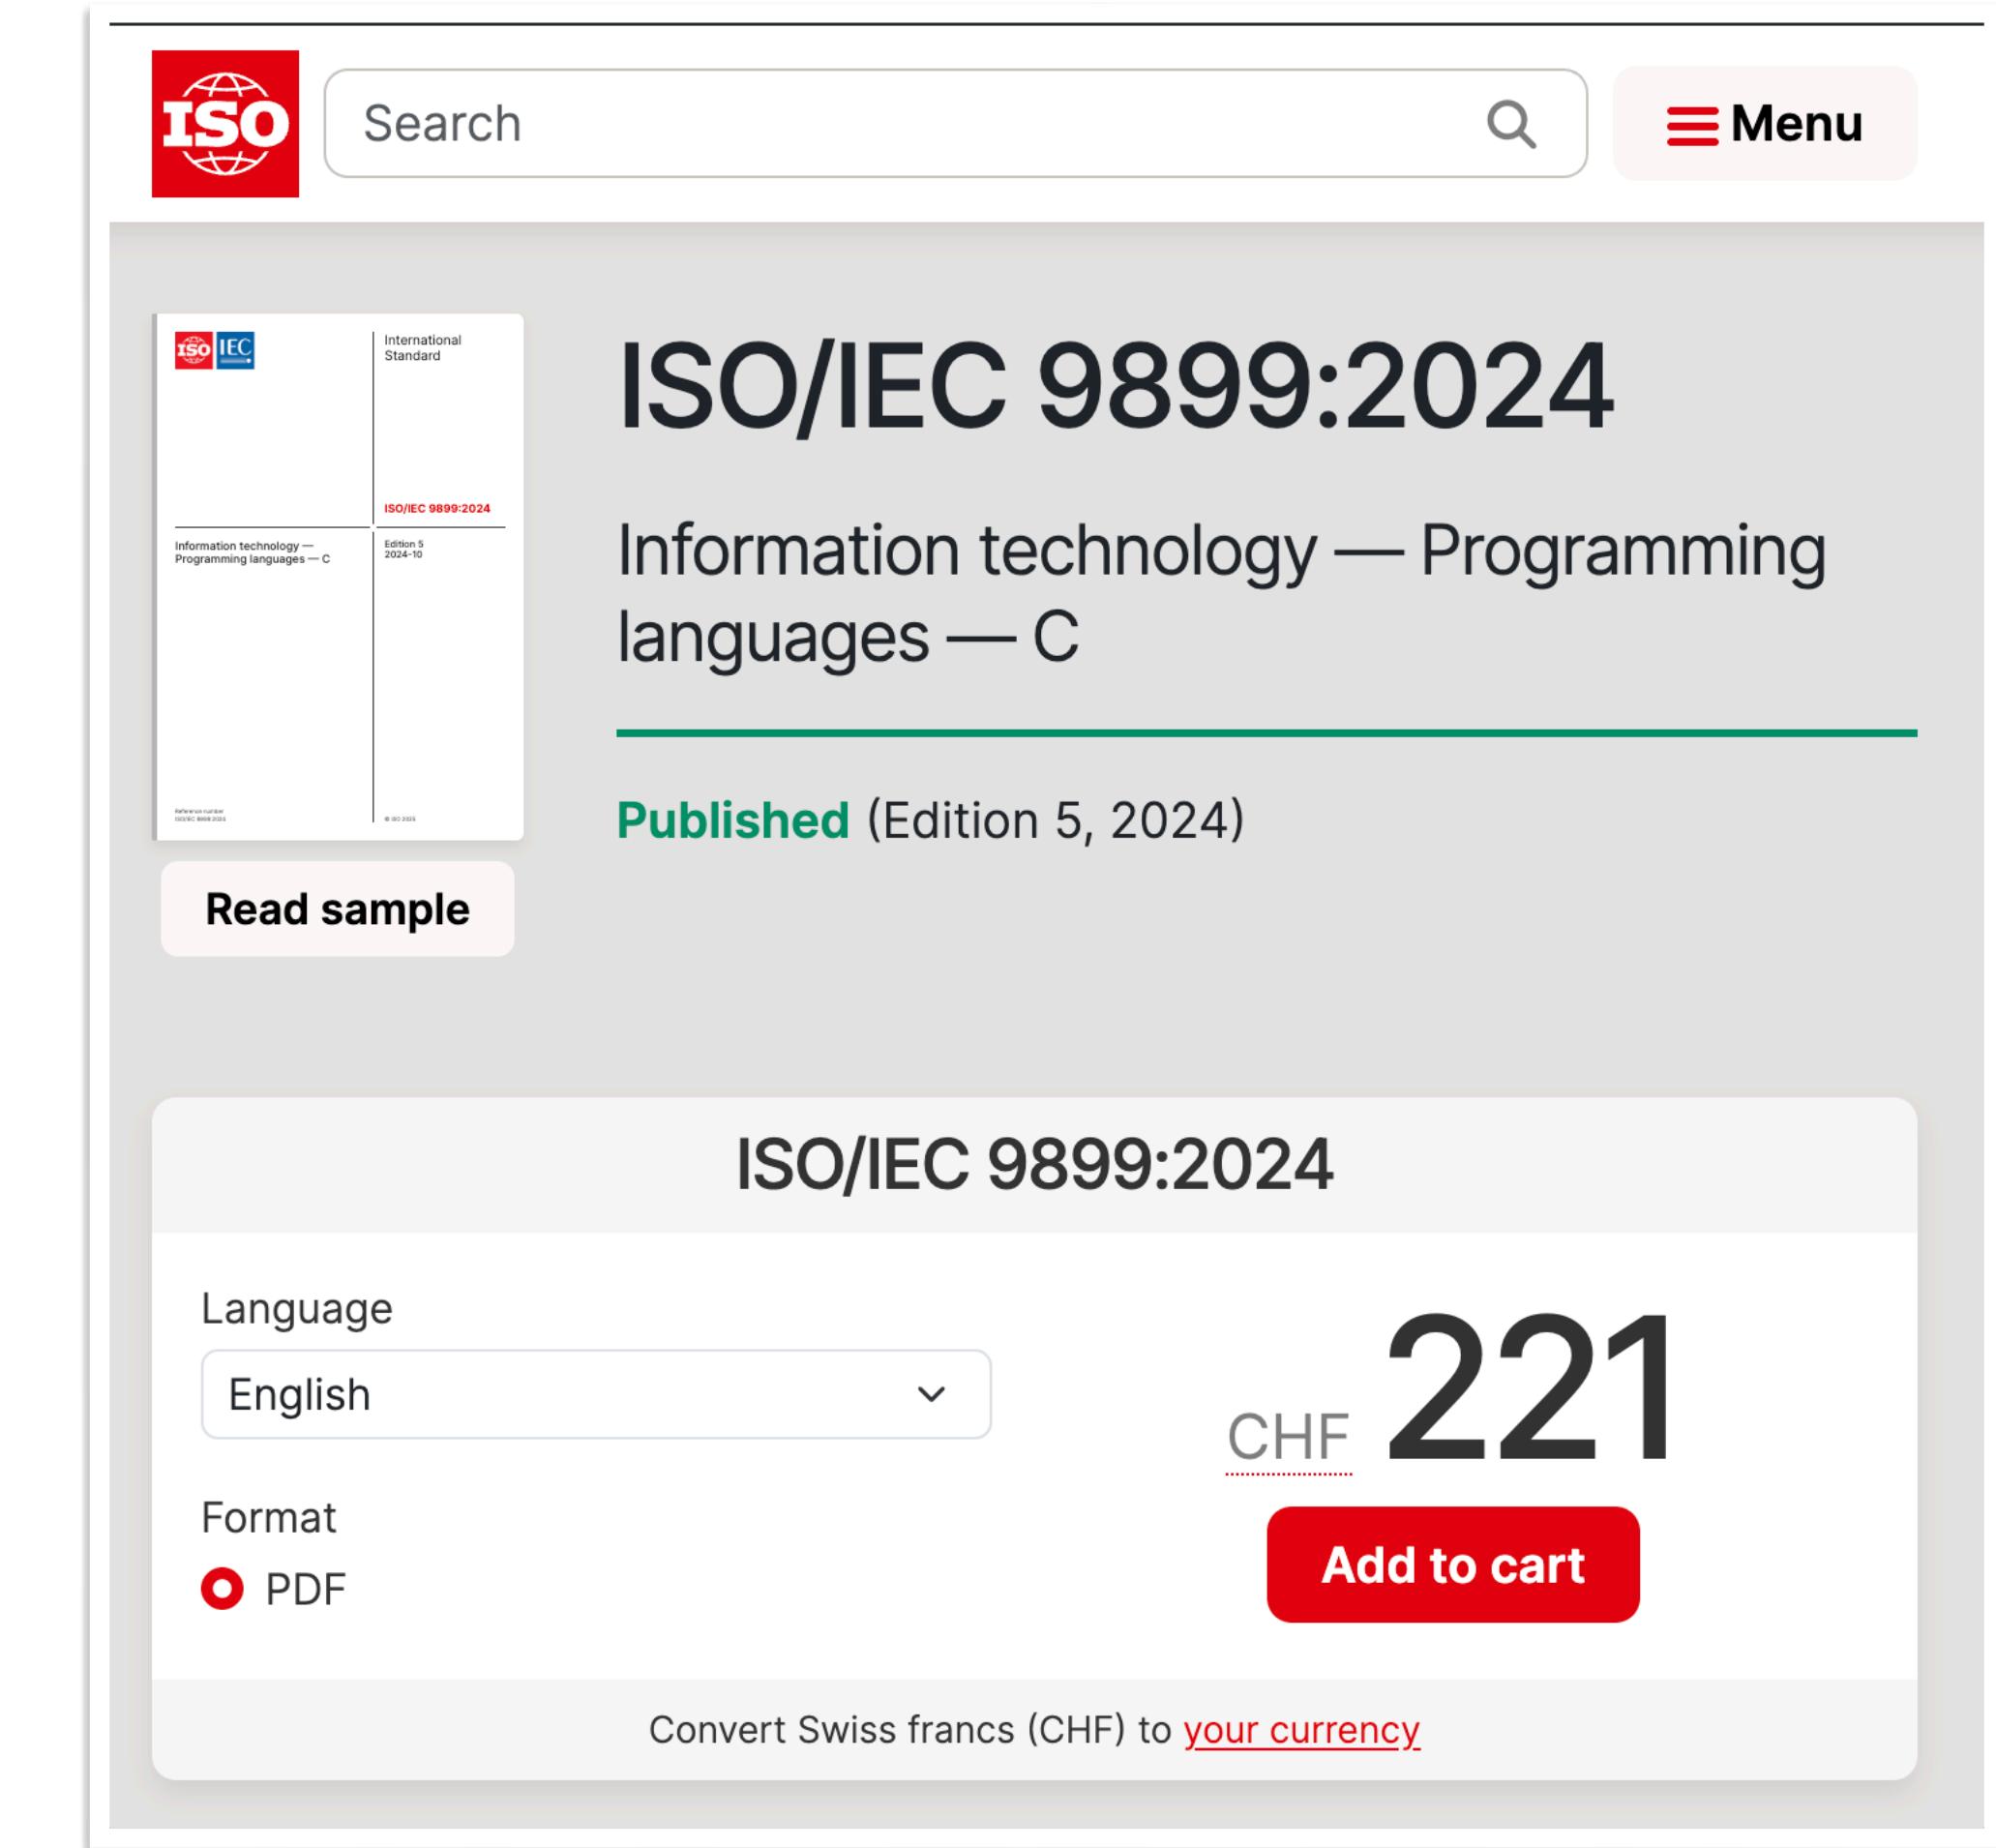
- 27 maintainers from France, UK, Japan, India and USA, across industry and academia.
- Custodians of the compiler
  - *Not the ones deciding how the language should evolve!*

# Who decides how OCaml evolves?



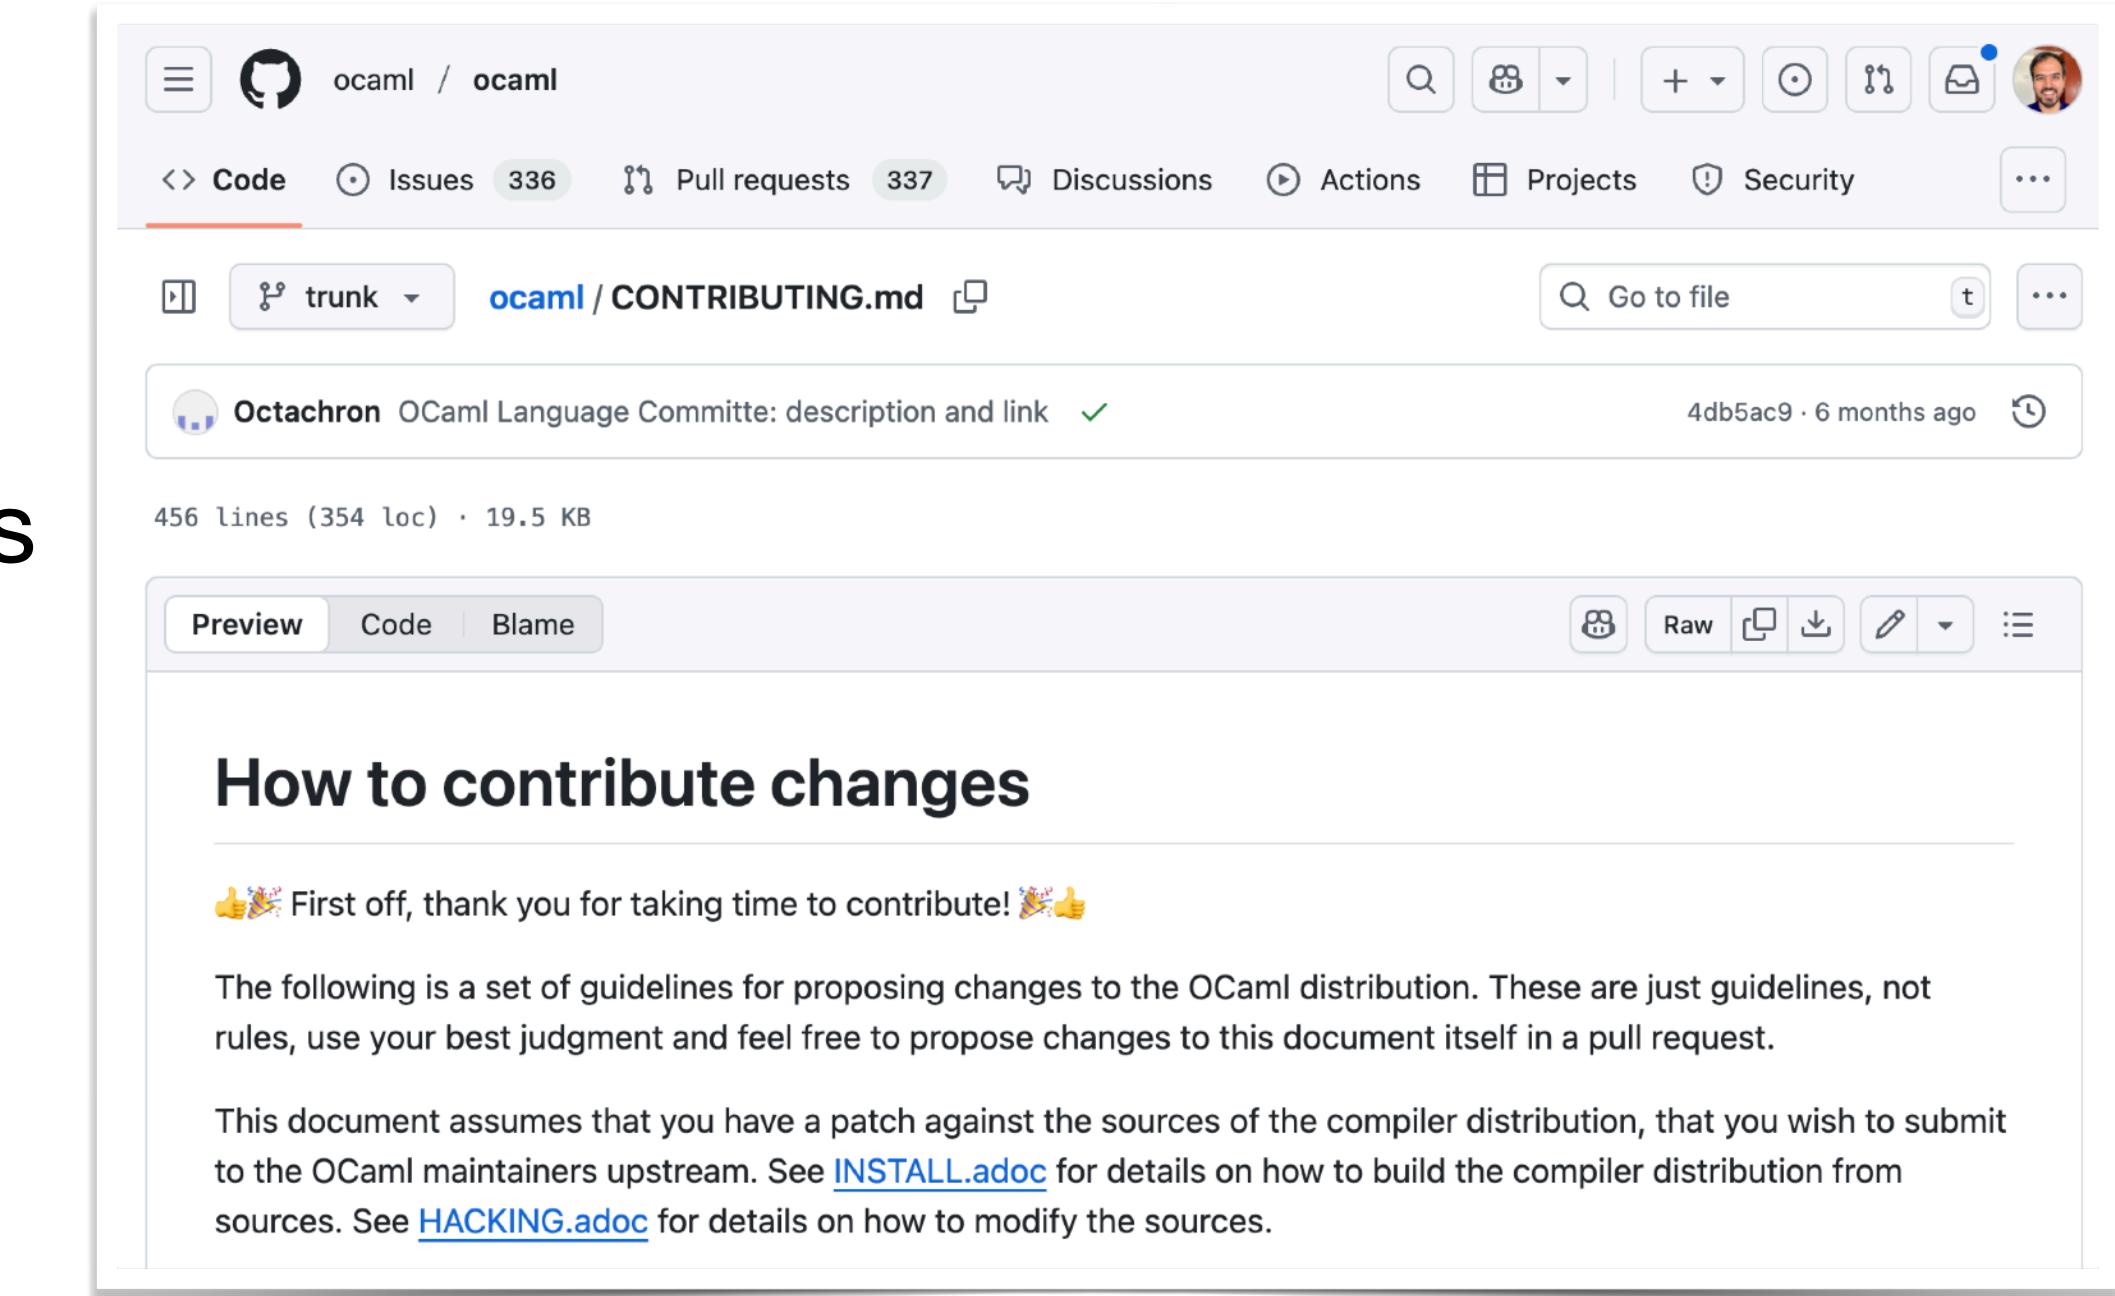
# Who decides how OCaml evolves?

- Evolution
  - **User-driven:** OCaml, Python
  - **Committee-driven:** ISO/IEC evolving C and C++
  - **Vendor-driven consensus:** WebAssembly
- **Language and compiler** aren't distinct
  - OCaml compiler implementation **IS** the language.
    - Unlike C, Wasm, JavaScript
  - *Bar is lower to change the language*



# Evolving OCaml

- Open process
  - OCaml compiler is maintained on GitHub
  - All discussions are public in the PRs, Issues and RFCs on GitHub
- Multi-speed model
  - **Small fixes/features** → Make an issue (“feature request”), open a PR, discuss and get that merged
    - Every PR needs a maintainer's approval before merging
  - **Large features** → Bespoke based on the features
    - May need publishing papers, extensive performance evaluation, formalised/mechanised soundness results, etc.
- ***Often, presumably small feature requests take a life of their own!***



# A small(?) change – Dynamic Arrays

The screenshot shows a vertical stack of five GitHub pull request cards:

- Added dynamic arrays #9122**: Opened on Nov 15, 2019, Closed on Nov 15 2019. Status: Closed. Mathilde411 wants to merge 1 commit into `ocaml:trunk` from `Mathilde411:trunk`. Conversation: 12, Commits: 1, Checks: 0, Files changed: 7. +198 -5.
- add Dynarray to the stdlib. #11563**: Opened on Sep 25, 2022, Closed on Jan 18, 2023. Status: Closed. c-cube wants to merge 29 commits into `ocaml:trunk` from `c-cube:dyn-array`. Conversation: 92, Commits: 29, Checks: 0, Files changed: 10. +792 -1.
- Dynarrays, boxed #11882**: Opened on Jan 11, 2023, Merged on Oct 21, 2023. Status: Merged. gasche merged 51 commits into `ocaml:trunk` from `gasche:dyn-array-boxed` on Oct 21, 2023. Conversation: 92, Commits: 29, Checks: 0, Files changed: 10. +792 -1.
- Dynarrays, unboxed (with local dummies) #12885**: Opened on Jan 5, 2024, Merged on May 2, 2024. Status: Merged. gasche merged 9 commits into `ocaml:trunk` from `gasche:dynarray-unboxed-dummy` on May 2, 2024. Conversation: 53, Commits: 9, Checks: 0, Files changed: 6. +455 -235.
- gasche commented on Jan 5, 2024 · edited**: A comment from gasche on the last pull request.

**Opened: Nov 15, 2019, Closed: Nov 15 2019**

Implementation rather naive, room for improvements

**Opened: Sep 25, 2022, Closed: Jan 18, 2023**

Clean API, **but** multicore safety, performance

**Opened: Jan 11, 2023, Merged: Oct 21, 2023**

Clean API **and** simple implementation

**Opened: Jan 5, 2024, Merged: May 2, 2024**

Clean API and **optimised** implementation

# Dynamic Arrays

**Added dynamic arrays #9122**

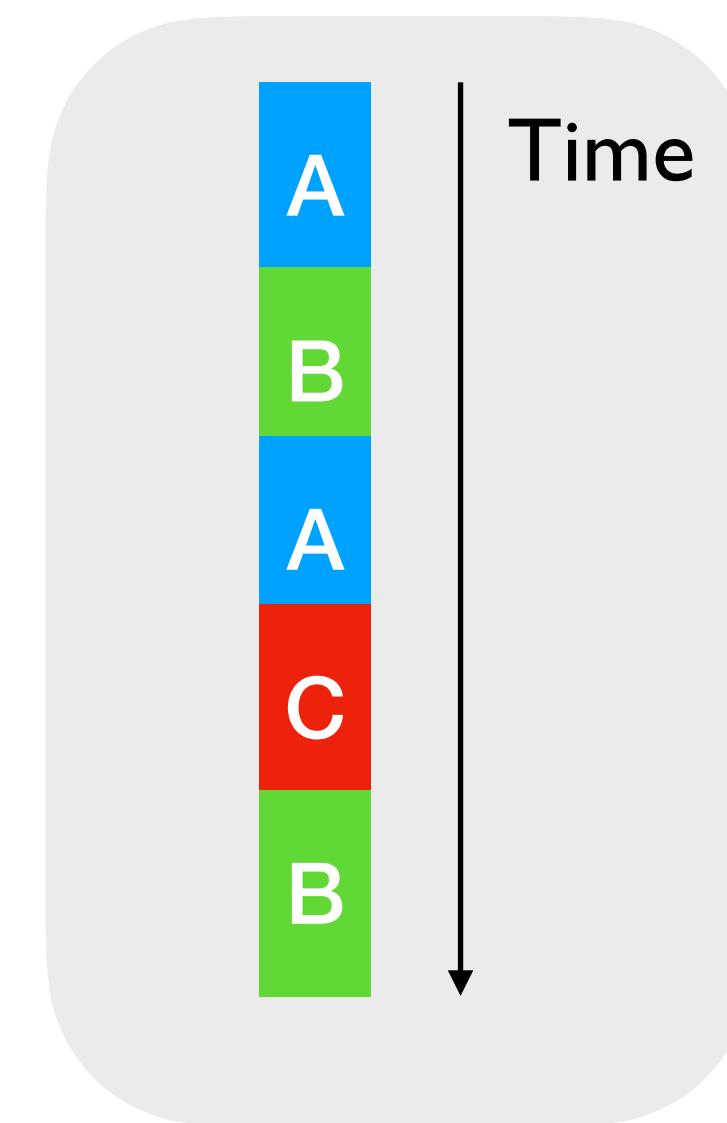
**Dynarrays, unboxed (with local dummies) #12885**

- **Summary**
  - Proposed – Nov 2019, Merged – (PR#1) Jan 2024; (PR#2) May 2024
  - Initially – 198 loc, finally – ~2500 loc
  - 500+ comments in the various PRs
- **Worth it?**
  - **Yes!** Should work for the next couple of decades.
  - Harder to undo changes after the release.

# A large change – Multicore OCaml

- Native support for concurrency and parallelism to OCaml

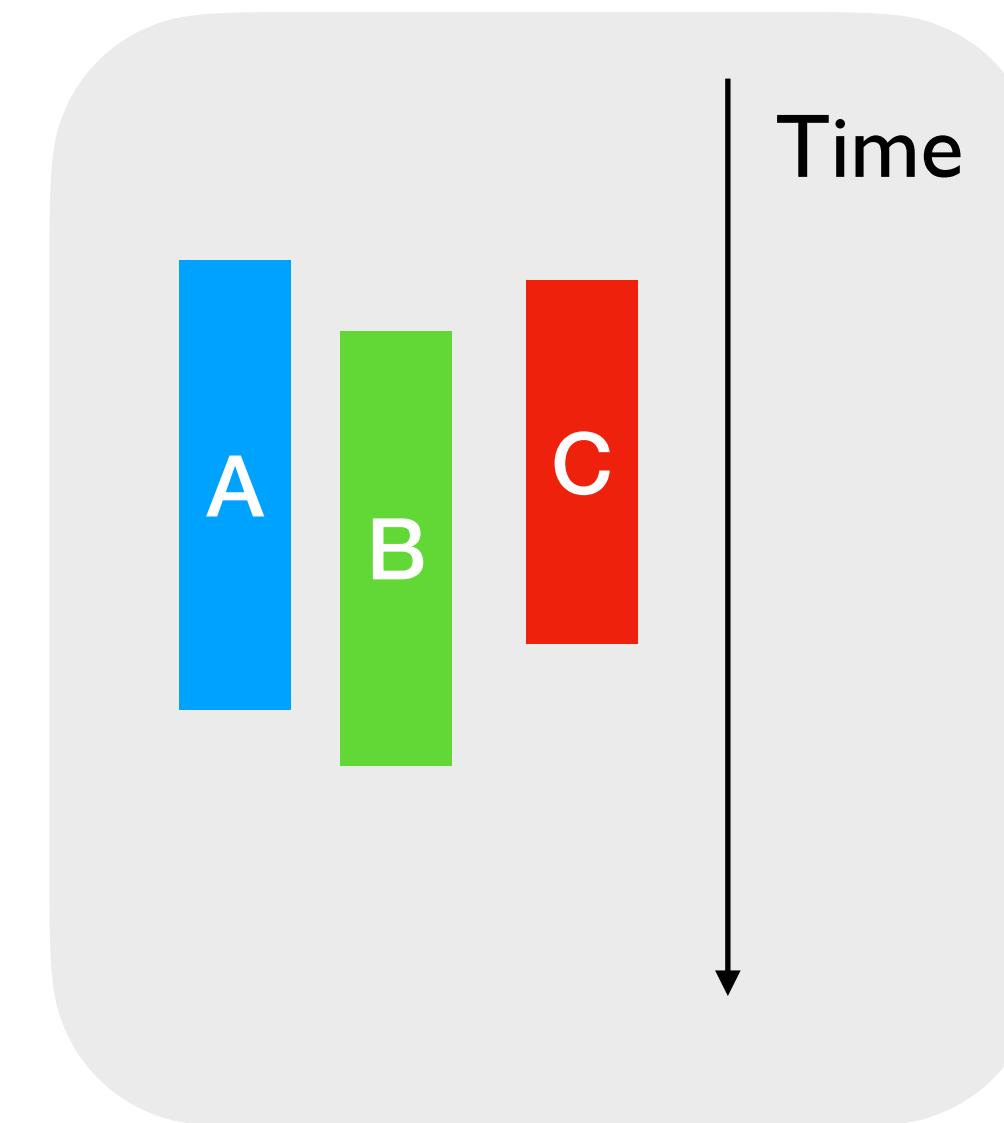
*Concurrency*



*Interleaved execution*

*Effect Handlers*

*Parallelism*



*Simultaneous execution*

*Domains*

# Challenges

- A new multicore garbage collector and multicore runtime system
  - *Replacing a car engine with a new one!*
- Make the language itself thread-safe
  - OCaml is a safe language! (Unlike C/C++, Go)
- Maintain feature and performance backwards compatibility!
  - Most OCaml programs will continue to remain single-threaded

Build credibility by ***publishing key results*** and ***rigorous evaluation***

# Starting out

## Multicore OCaml

Stephen Dolan

Leo White

Anil Madhavapeddy

*Currently, threading is supported in OCaml only by means of a global lock, allowing at most one thread to run OCaml code at any time. We present ongoing work to design and implement an OCaml runtime capable of shared-memory parallelism.*

### 1 Introduction

Adding shared-memory parallelism to an existing lan-

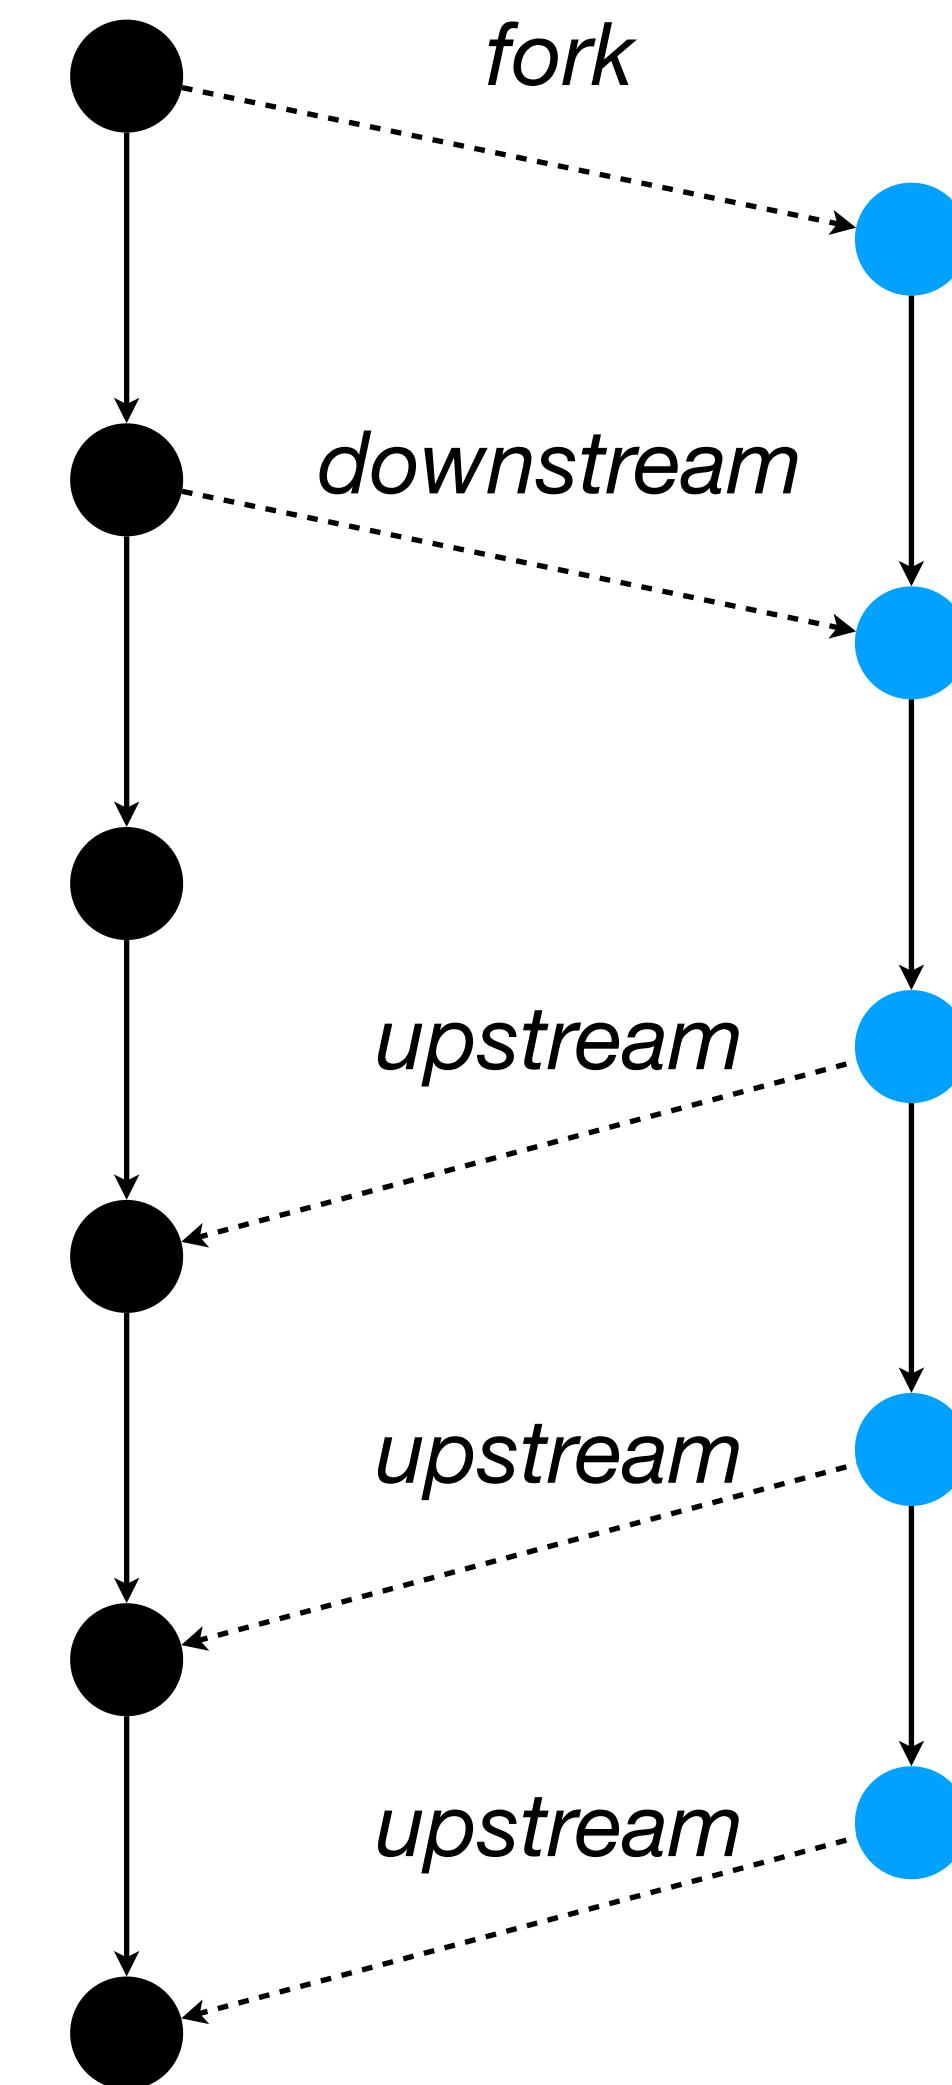
all objects reachable from it to be promoted to the shared heap en masse. Unfortunately this eagerly promotes many objects that were never really shared: just because an object is pointed to by a shared object does not mean another thread is actually going to attempt to access it.

Our design is similar but lazier, along the lines of the multicore Haskell work [2], where objects are promoted to the shared heap whenever another thread

**OCaml Workshop 2014**

## Upstream OCaml

## Multicore OCaml



# Building confidence – Papers

Multicore GC and  
runtime system

## Retrofitting Parallelism onto OCaml

KC

STE

LEO

SAD

TO

AN

SU

AT

AN

OCa

OCa

memory parallel pro

f

1

## Bounding Data Races in Space and Time

(Extended version, with appendices)

## Retrofitting Effect Handlers onto OCaml

Stephen Dolan

OCaml Labs  
Cambridge, UK

stephen.dolan@cl.cam.ac.uk

Leo White

Jane Street  
London, UK  
leo@lpw25.net

Sadiq Jaffer

Opsian and OCaml Labs  
Cambridge, UK  
sadiq@toao.com

Anil Madhavapeddy

University of Cambridge and OCaml Labs  
Cambridge, UK  
avsm2@cl.cam.ac.uk

Relaxed Memory  
Model

Concurrency  
story

### Abstract

Effect handlers have been gathering momentum as a mechanism for modular programming with user-defined effects.

### 1 Introduction

Effect handlers [45] provide a modular foundation for user-defined effects. The key idea is to separate the definition of

**Peer-reviewed ideas build confidence**

# Diving deeper – Concurrency

## Retrofitting Effect Handlers onto OCaml

KC Sivaramakrishnan  
IIT Madras  
Chennai, India  
kcsrk@cse.iitm.ac.in

Stephen Dolan  
OCaml Labs  
Cambridge, UK  
stephen.dolan@cl.cam.ac.uk

Leo White  
Jane Street  
London, UK  
leo@lpw25.net

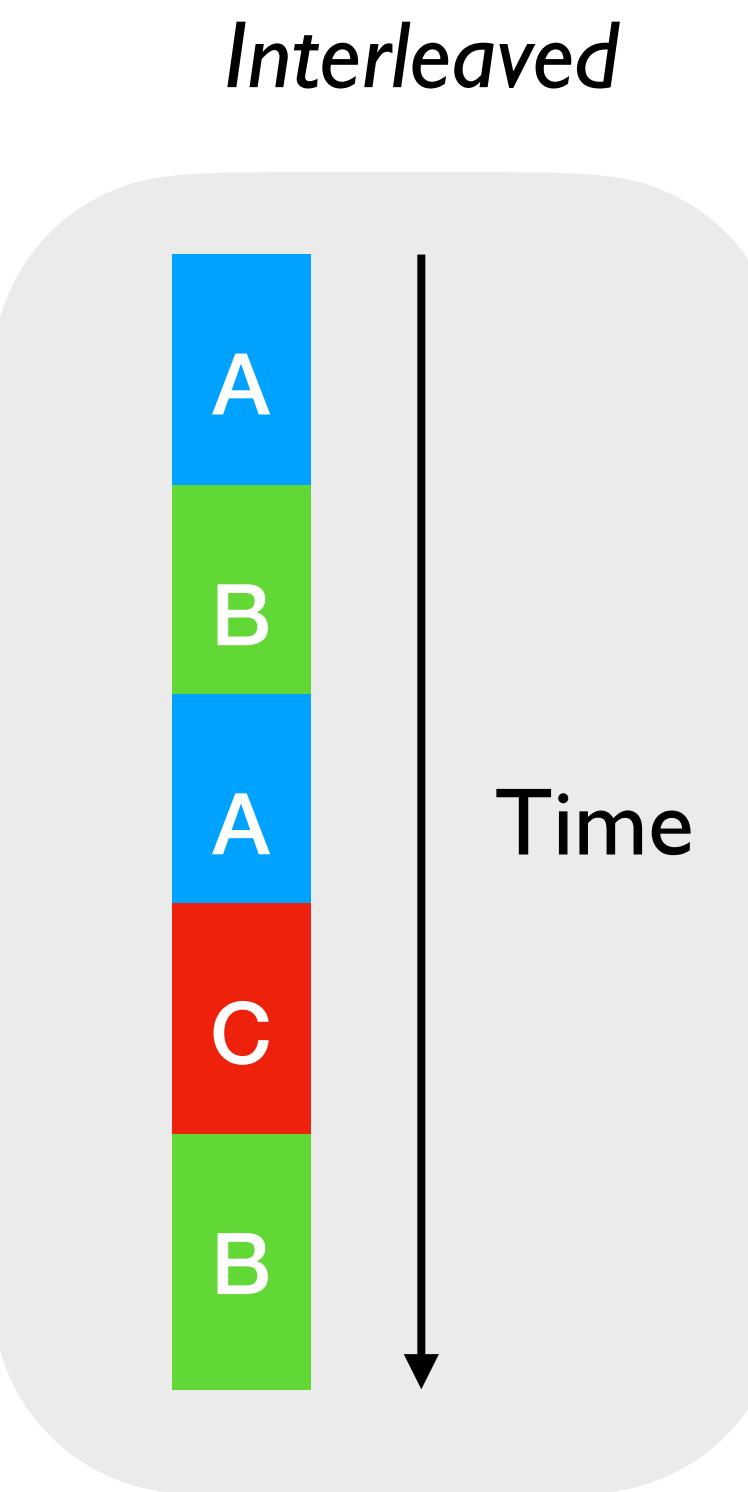
Tom Kelly  
OCaml Labs  
Cambridge, UK  
tom.kelly@cantab.net

Sadiq Jaffer  
Opsian and OCaml Labs  
Cambridge, UK  
sadiq@toao.com

Anil Madhavapeddy  
University of Cambridge and OCaml Labs  
Cambridge, UK  
avsm2@cl.cam.ac.uk

**Abstract**  
Effect handlers have been gathering momentum as a mechanism for modular programming with user-defined effects.

1 Introduction  
Effect handlers [45] provide a modular foundation for user-defined effects. The key idea is to separate the definition of



# Concurrent Programming

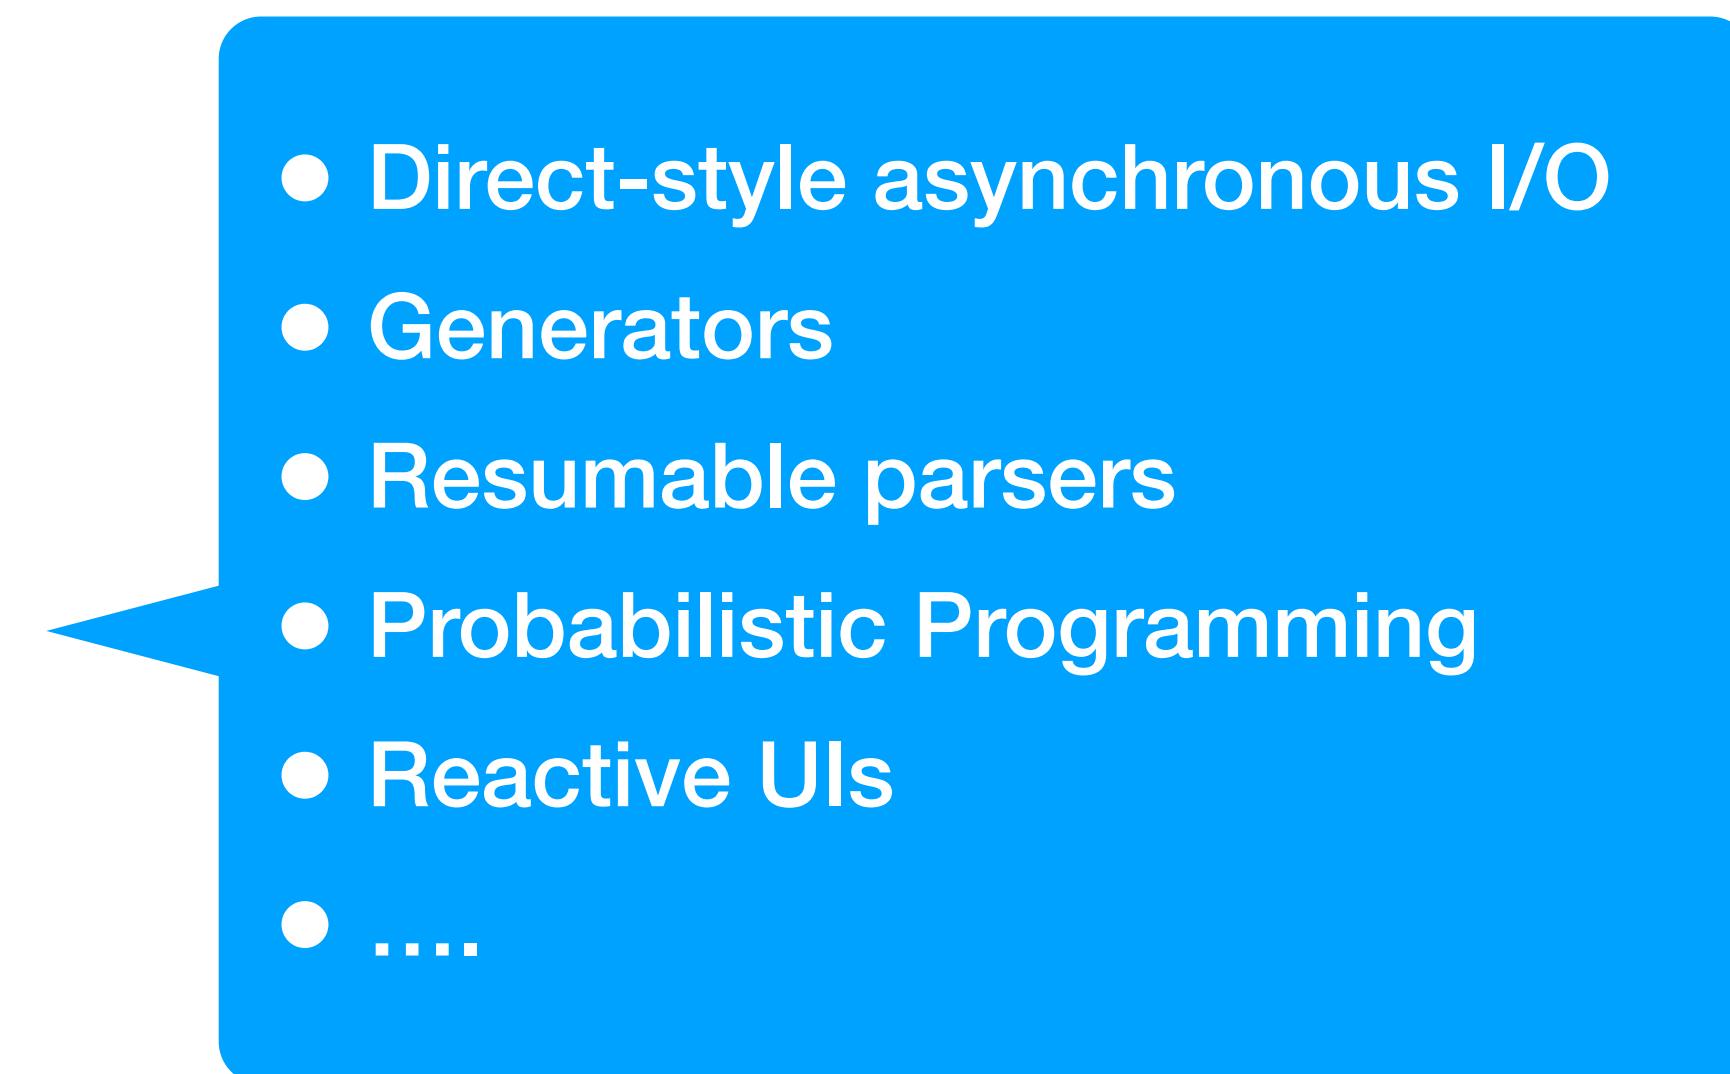
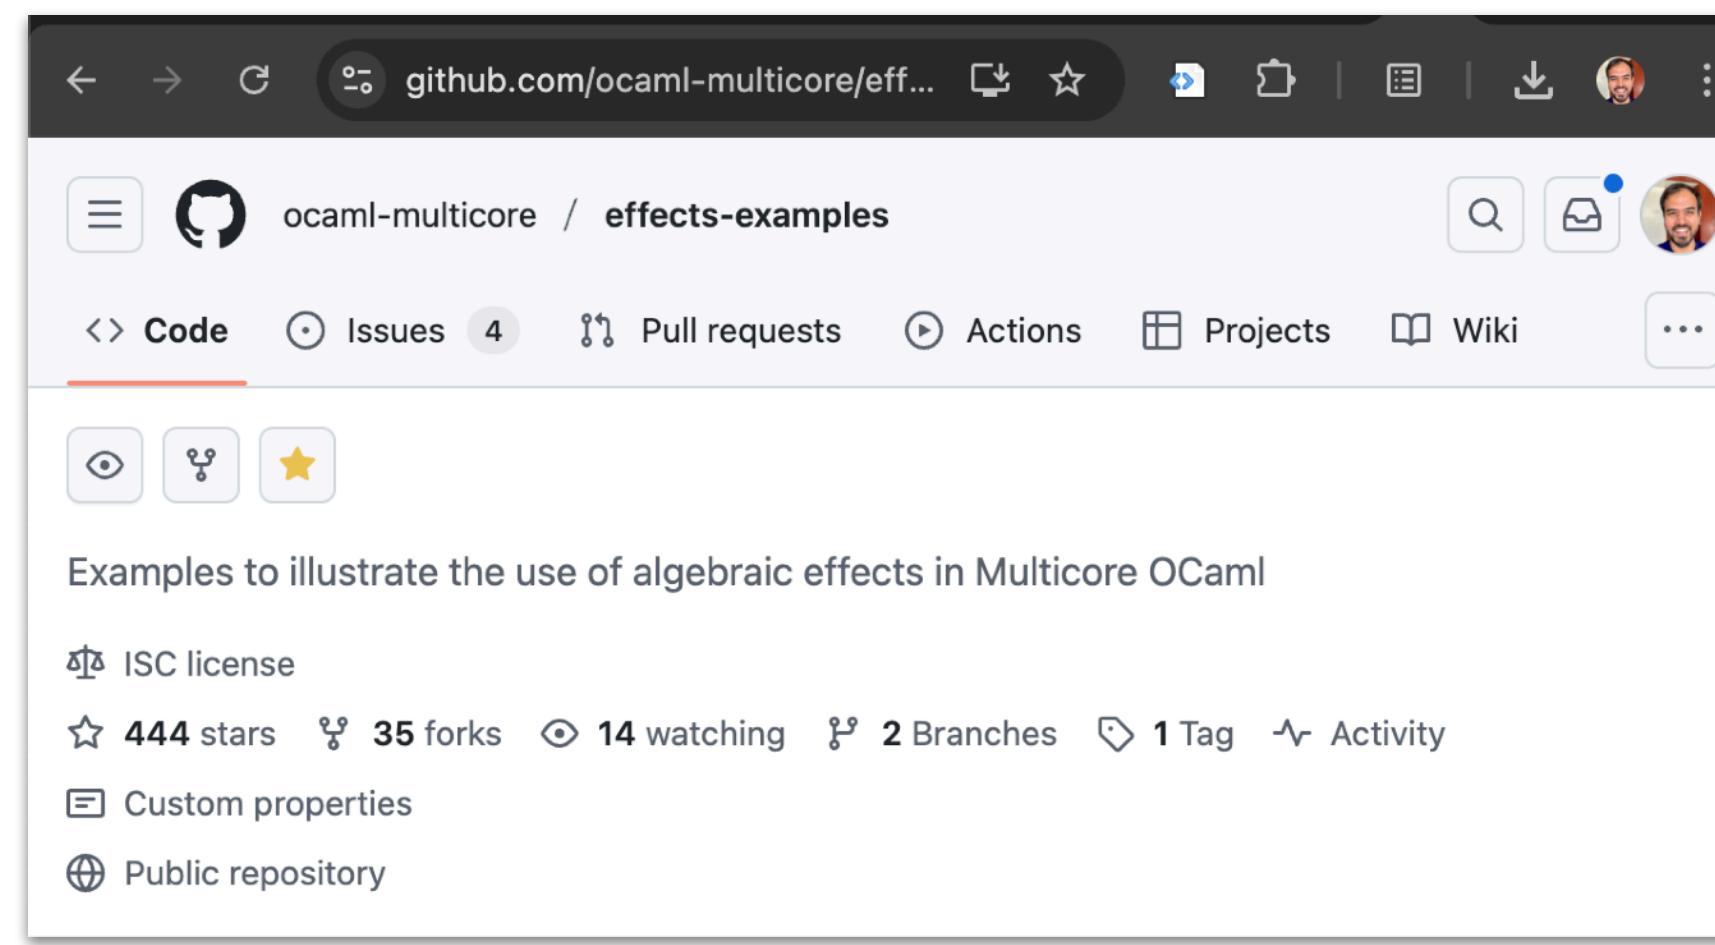
- Computations may be *suspended* and *resumed* later
- Many languages provide concurrent programming mechanisms as *primitives*
  - ◆ **async/await** – JavaScript, Python, Rust, C# 5.0, F#, Swift, ...
  - ◆ **generators** – Python, Javascript, ...
  - ◆ **coroutines** – C++, Kotlin, Lua, ...
  - ◆ **futures & promises** – JavaScript, Swift, ...
  - ◆ **Lightweight threads/processes** – Haskell, Go, Erlang
- *Often include many different primitives in the same language!*
  - ◆ JavaScript has async/await, generators, promises, and callbacks

Don't want a **zoo** of primitives but  
want **expressivity**

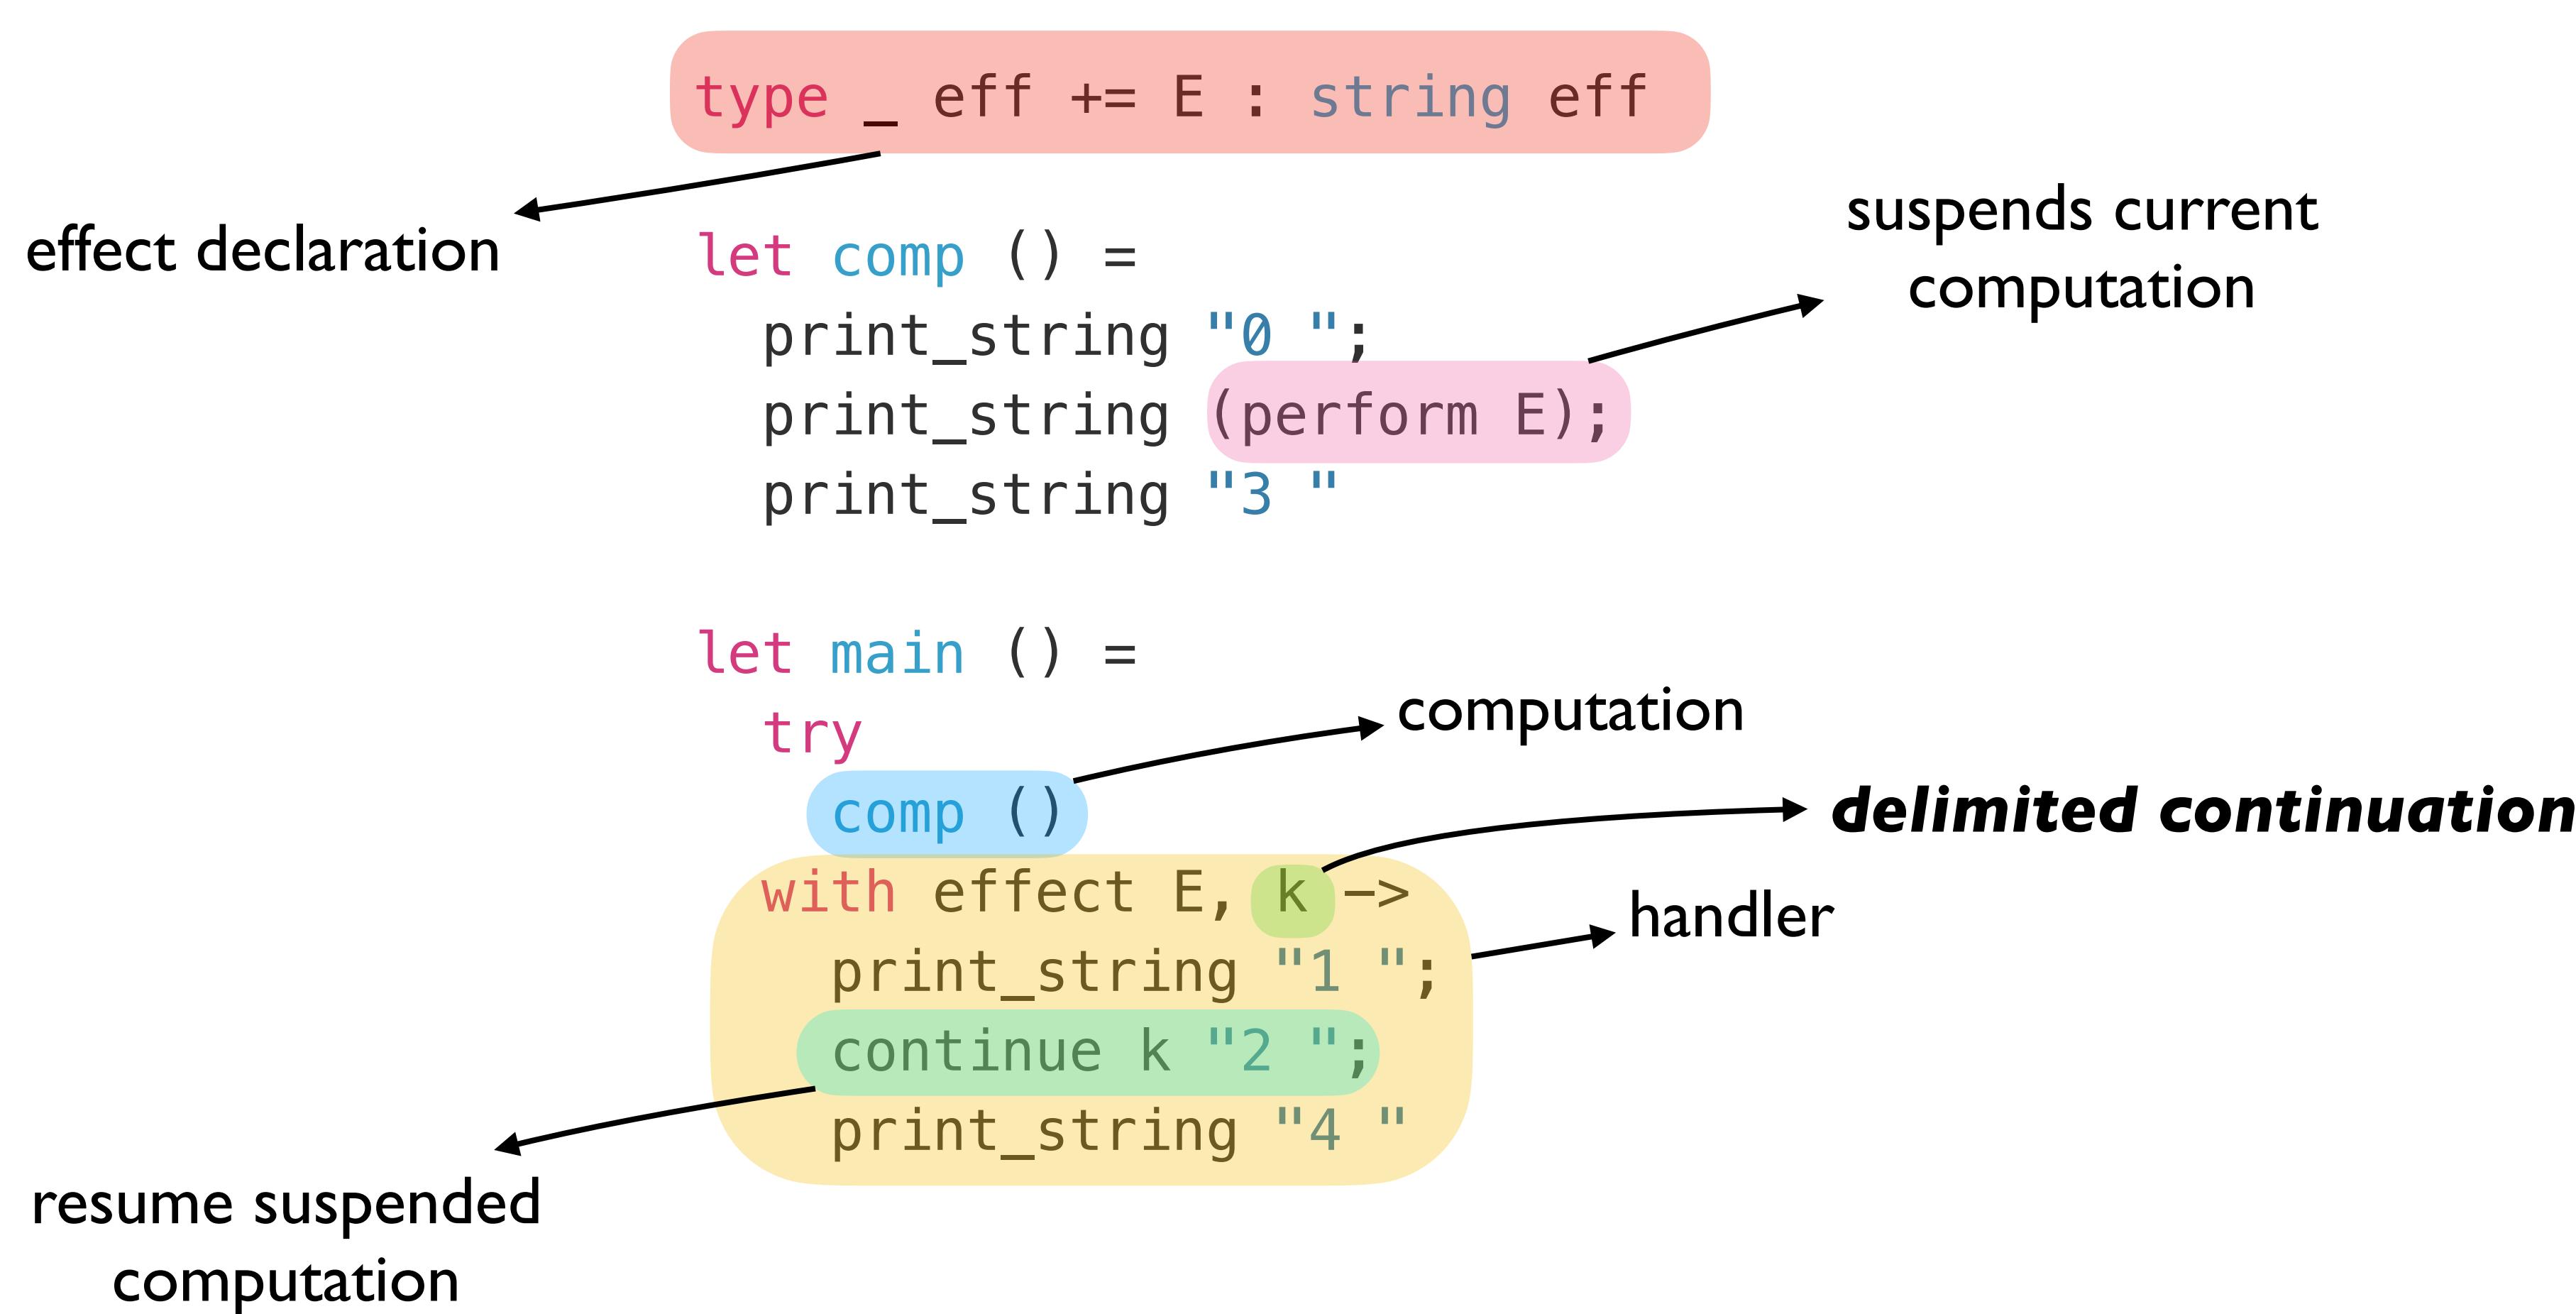
What's the **smallest** primitive that  
expresses **many** concurrency patterns?

# Effect handlers

- A mechanism for programming with *user-defined effects*
- *Modular* and *composable* basis of non-local control-flow mechanisms
  - ◆ Exceptions, generators, lightweight threads, promises, asynchronous IO, coroutines as *libraries*
- Effect handlers  $\sim=$  *first-class, restartable exceptions*
  - ◆ Structured programming with *delimited continuations*



# Effect handlers



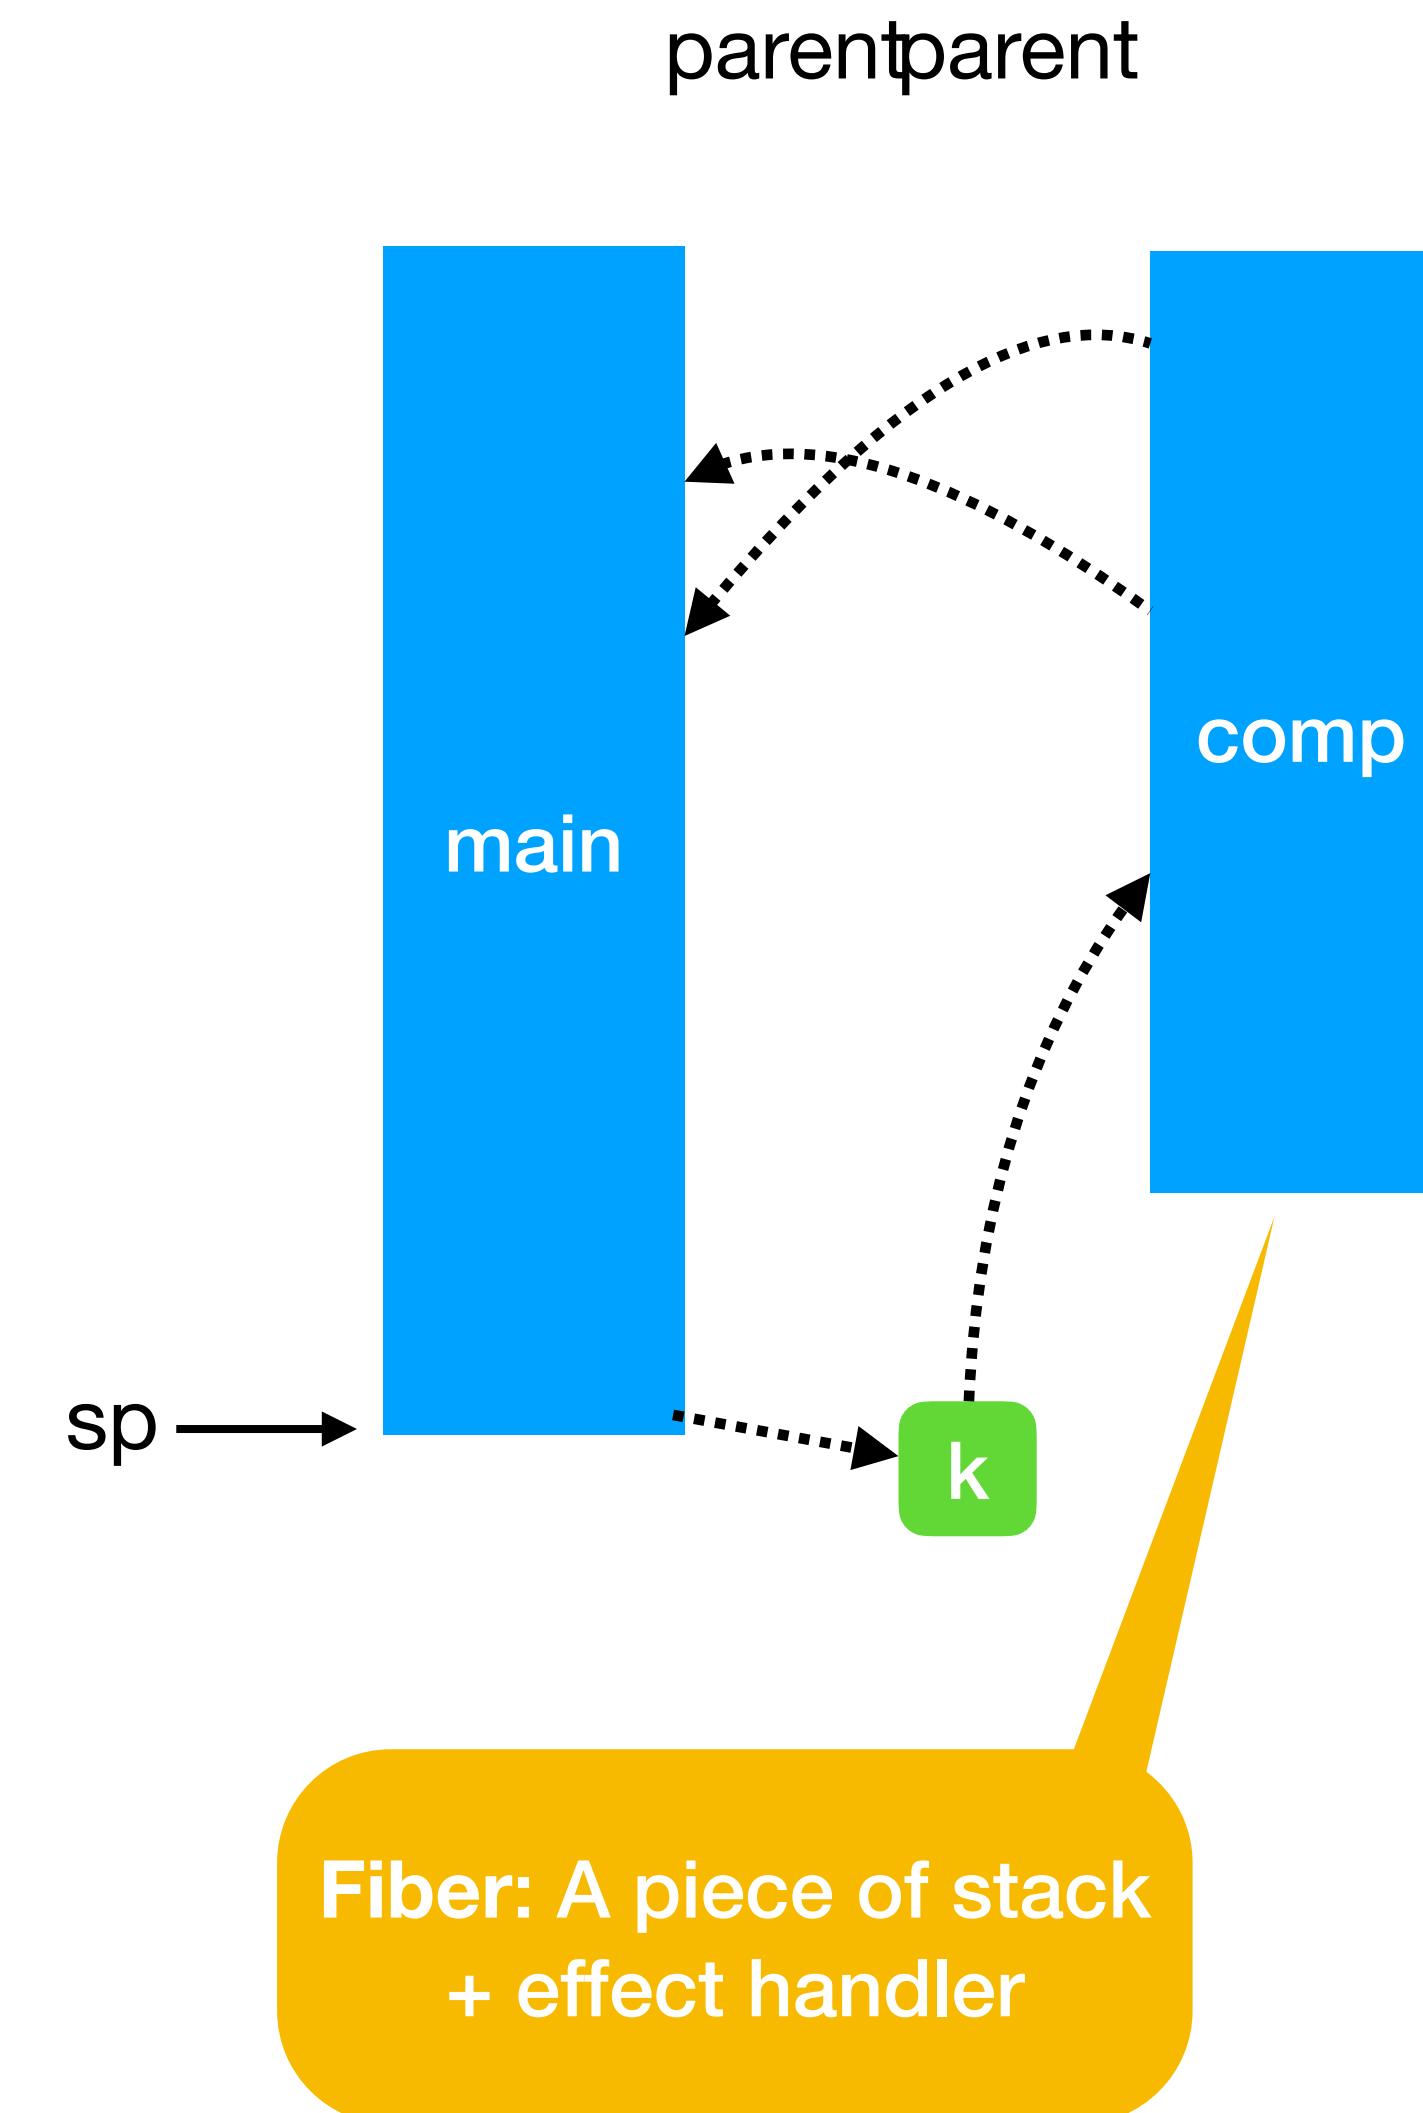
# Stepping through the example

```
type 'a eff += E : string eff

let comp () =
    print_string "0 ";
    print_string (perform E);
    print_string "3 "

let main () =
    try
        comp ()
    with effect E, k ->
        print_string "1 ";
        continue k "2 ";
        print_string "4 "

pc →
0 1 2 3 4
```



# Lightweight threading

```
type _ eff += Fork : (unit -> unit) -> unit eff
             | Yield : unit eff

let run main =
  ... (* assume queue of continuations *)
let run_next () =
  match dequeue () with
  | Some k -> continue k ()
  | None -> ()
in
let rec spawn f =
  match f () with
  | () -> run_next () (* value case *)
  | effect Yield, k -> enqueue k; run_next ()
  | effect (Fork f), k -> enqueue k; spawn f
in
spawn main

let fork f = perform (Fork f)
let yield () = perform Yield
```

Effect Handler {

# Lightweight threading

```
let main () =
  fork (fun _ ->
    print_endline "1.a";
    yield ();
    print_endline "1.b");
  fork (fun _ ->
    print_endline "2.a";
    yield ();
    print_endline "2.b")
;;
run main
```

```
1.a
2.a
1.b
2.b
```

# Lightweight threading

```
let main () =
  fork (fun _ ->
    print_endline "1.a";
    yield ();
    print_endline "1.b");
  fork (fun _ ->
    print_endline "2.a";
    yield ();
    print_endline "2.b")
;;
run main
```

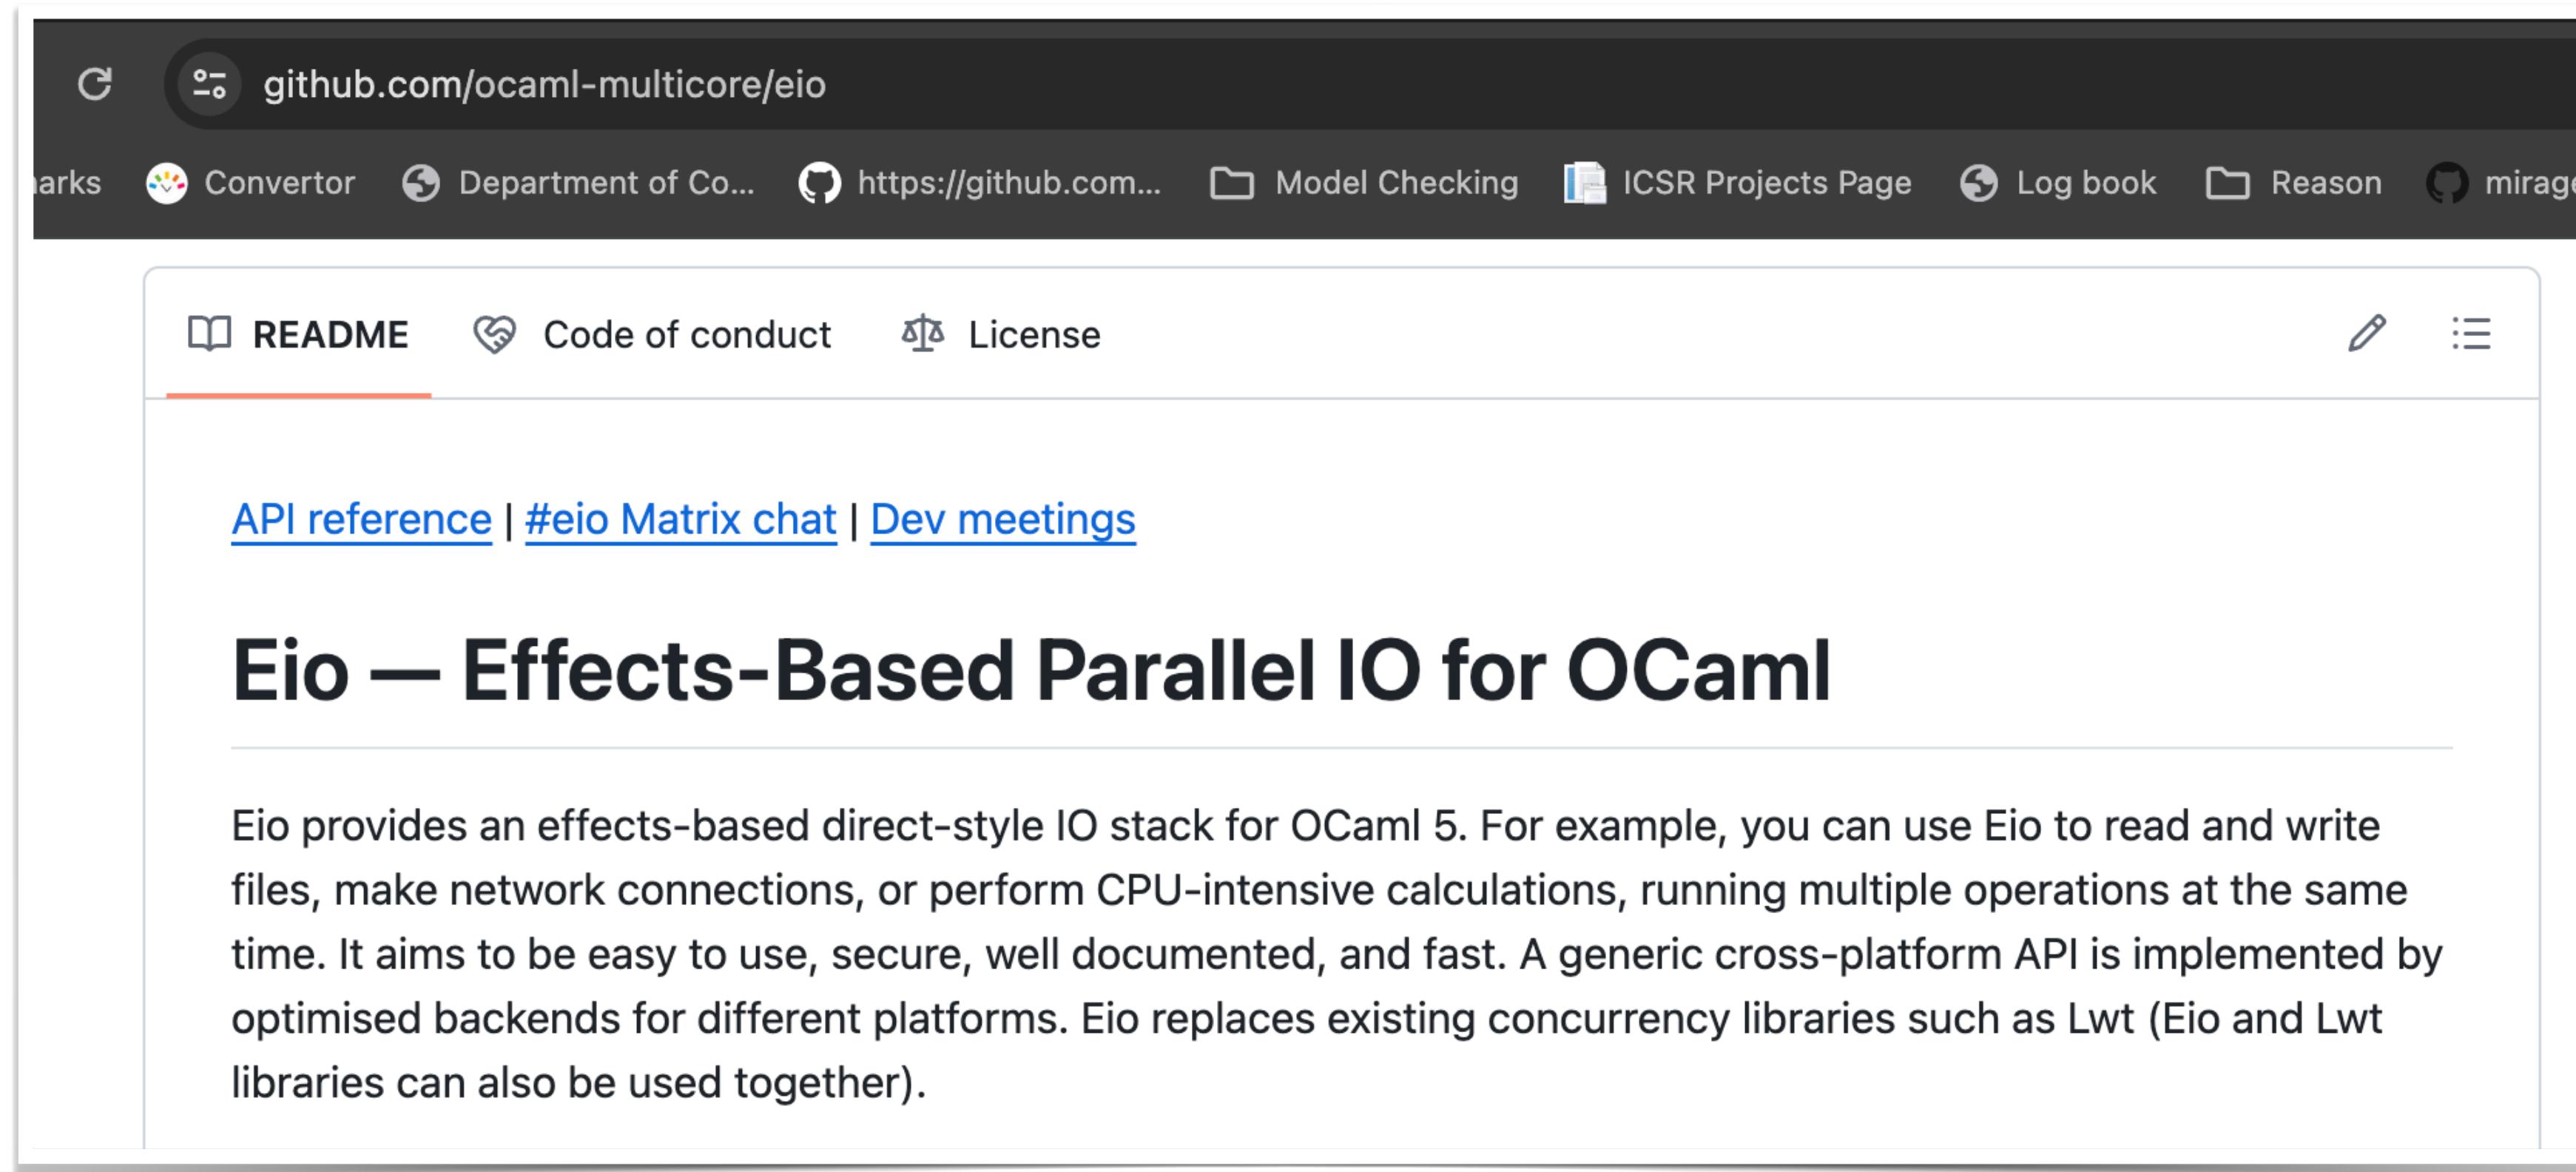
Ability to specialise  
scheduler  
unlike GHC Haskell / Go

1.a  
2.a  
1.b  
2.b

User-code need not be  
aware of effects

# Industrial-strength concurrency

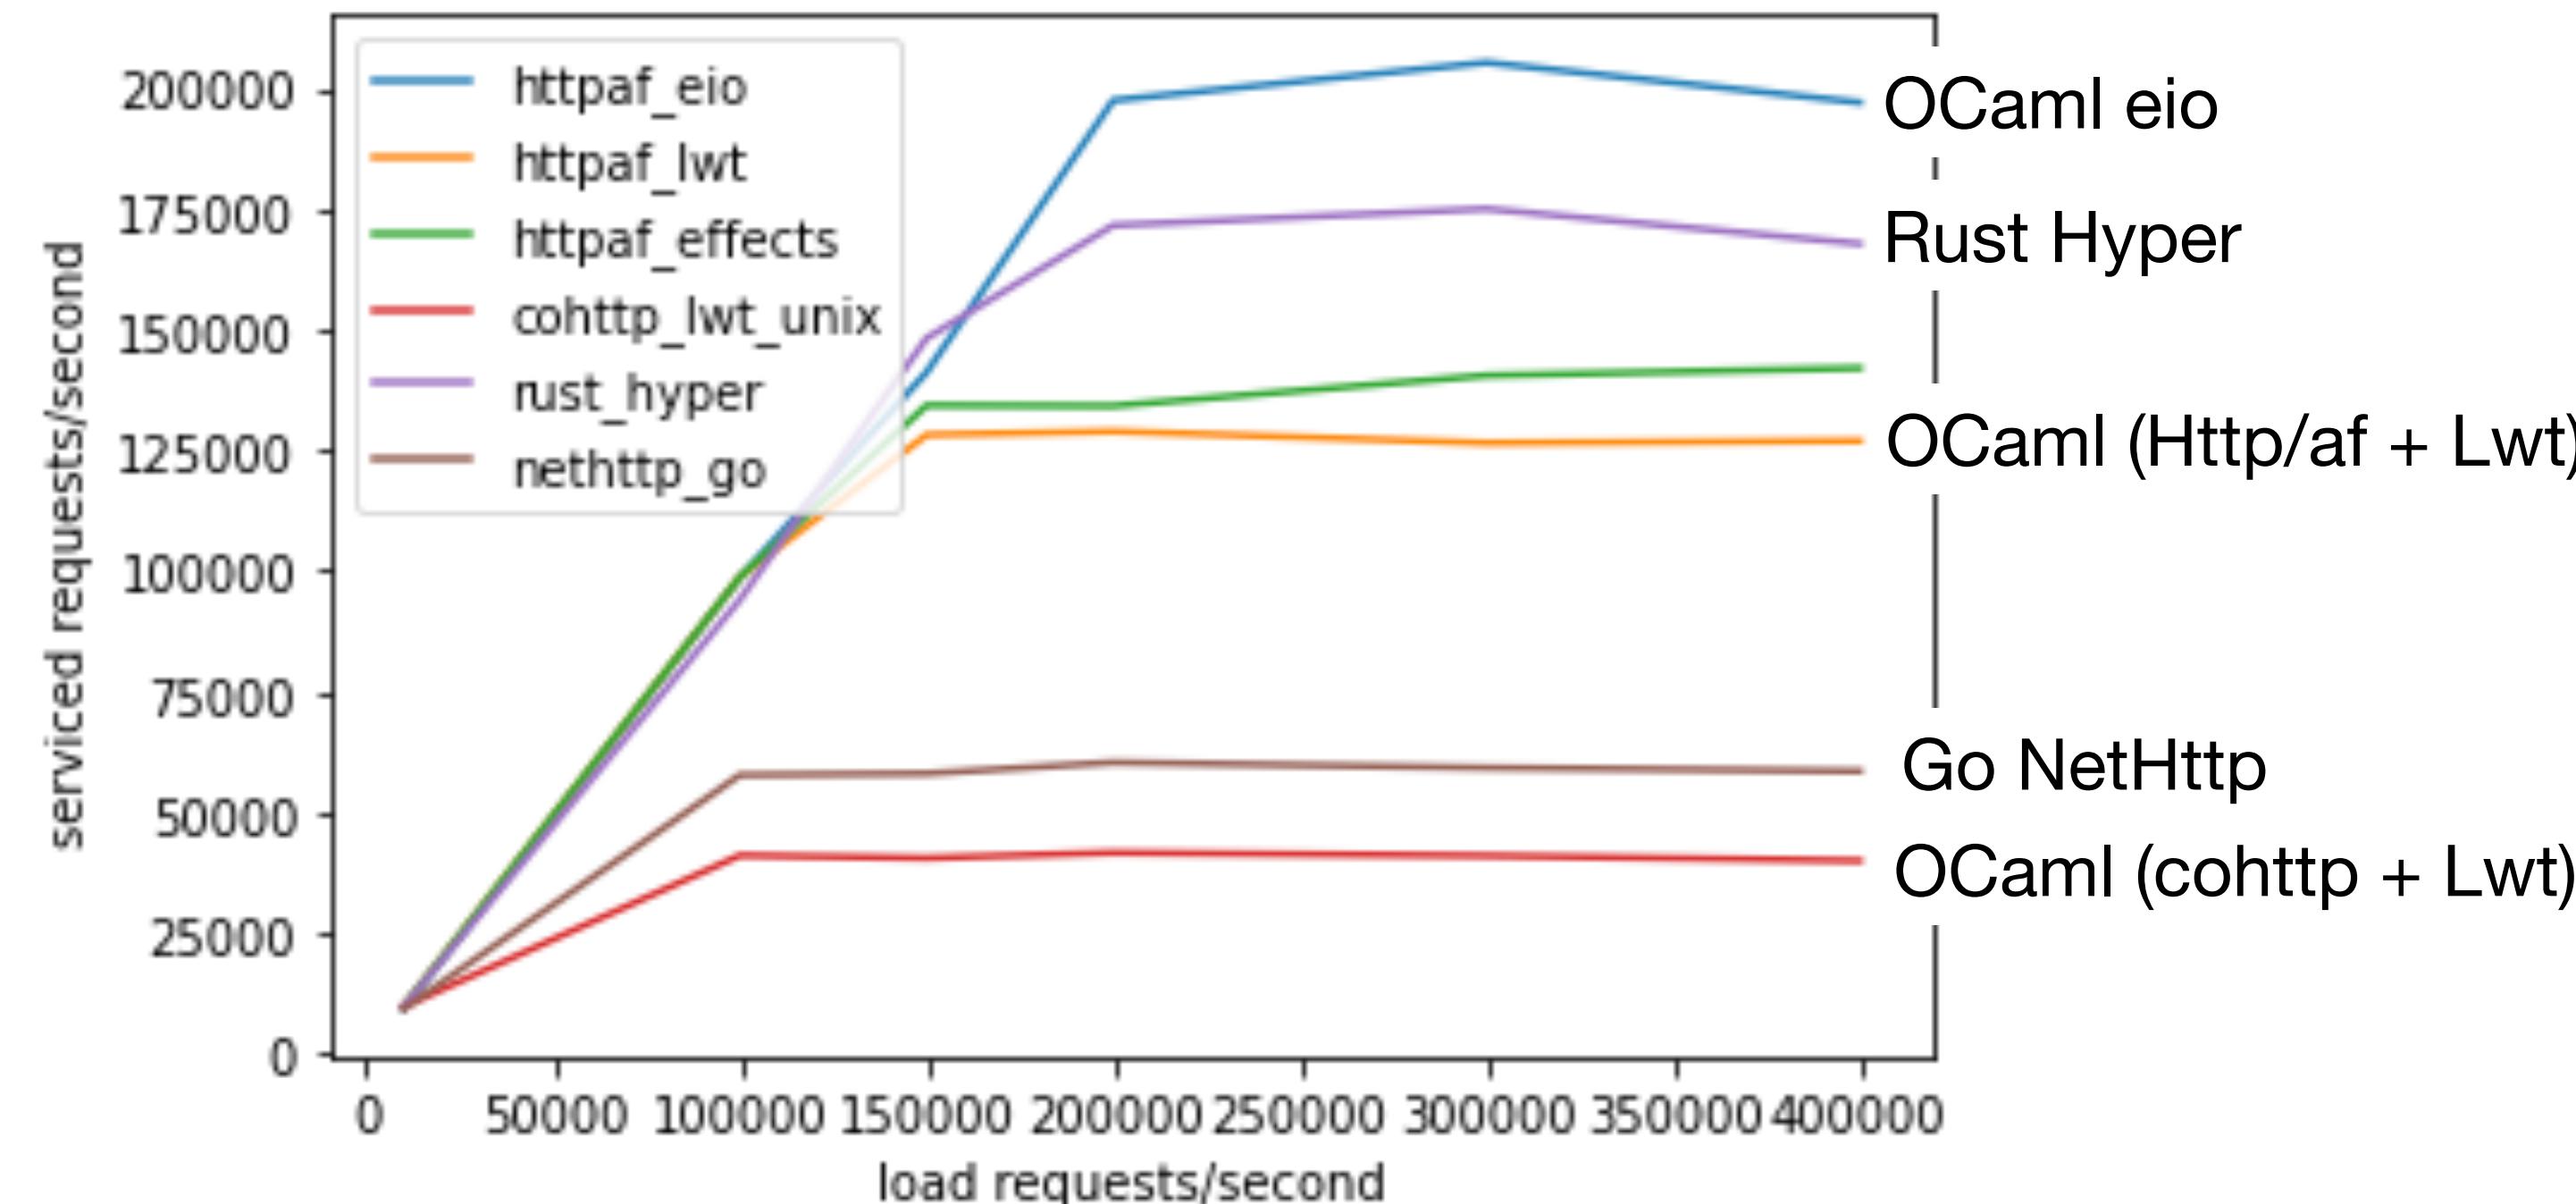
- **eio**: effects-based direct-style I/O
  - ◆ Multiple backends – epoll, select, *io\_uring* (*new async io in Linux kernel*)



<https://github.com/ocaml-multicore/eio>

# Industrial-strength concurrency

- **eio:** effects-based direct-style I/O
  - ◆ Multiple backends – epoll, select, *io\_uring* (*new async io in Linux kernel*)



100 open connections, 60 seconds w/ *io\_uring*

<https://github.com/ocaml-multicore/eio>

# Further reading

## Control structures in programming languages: from goto to algebraic effects

Xavier Leroy

This book is a journey through the design space and history of programming languages from the perspective of control structures: the language mechanisms that enable programs to control their execution flows. Starting with the “goto” jumps of early programming languages and the emergence of structured programming in the 1960s, the book explores advanced control structures for imperative languages such as generators and coroutines, then develops alternate views of control in functional languages, first as continuations and their control operators, then as algebraic effects and effect handlers. Blending history, code examples, and theory, the book offers an original, comparative perspective on programming languages, as well as an extensive introduction to algebraic effects and other contemporary research topics in P.L.

### Publication history

To be published by Cambridge University Press.

### Book preview

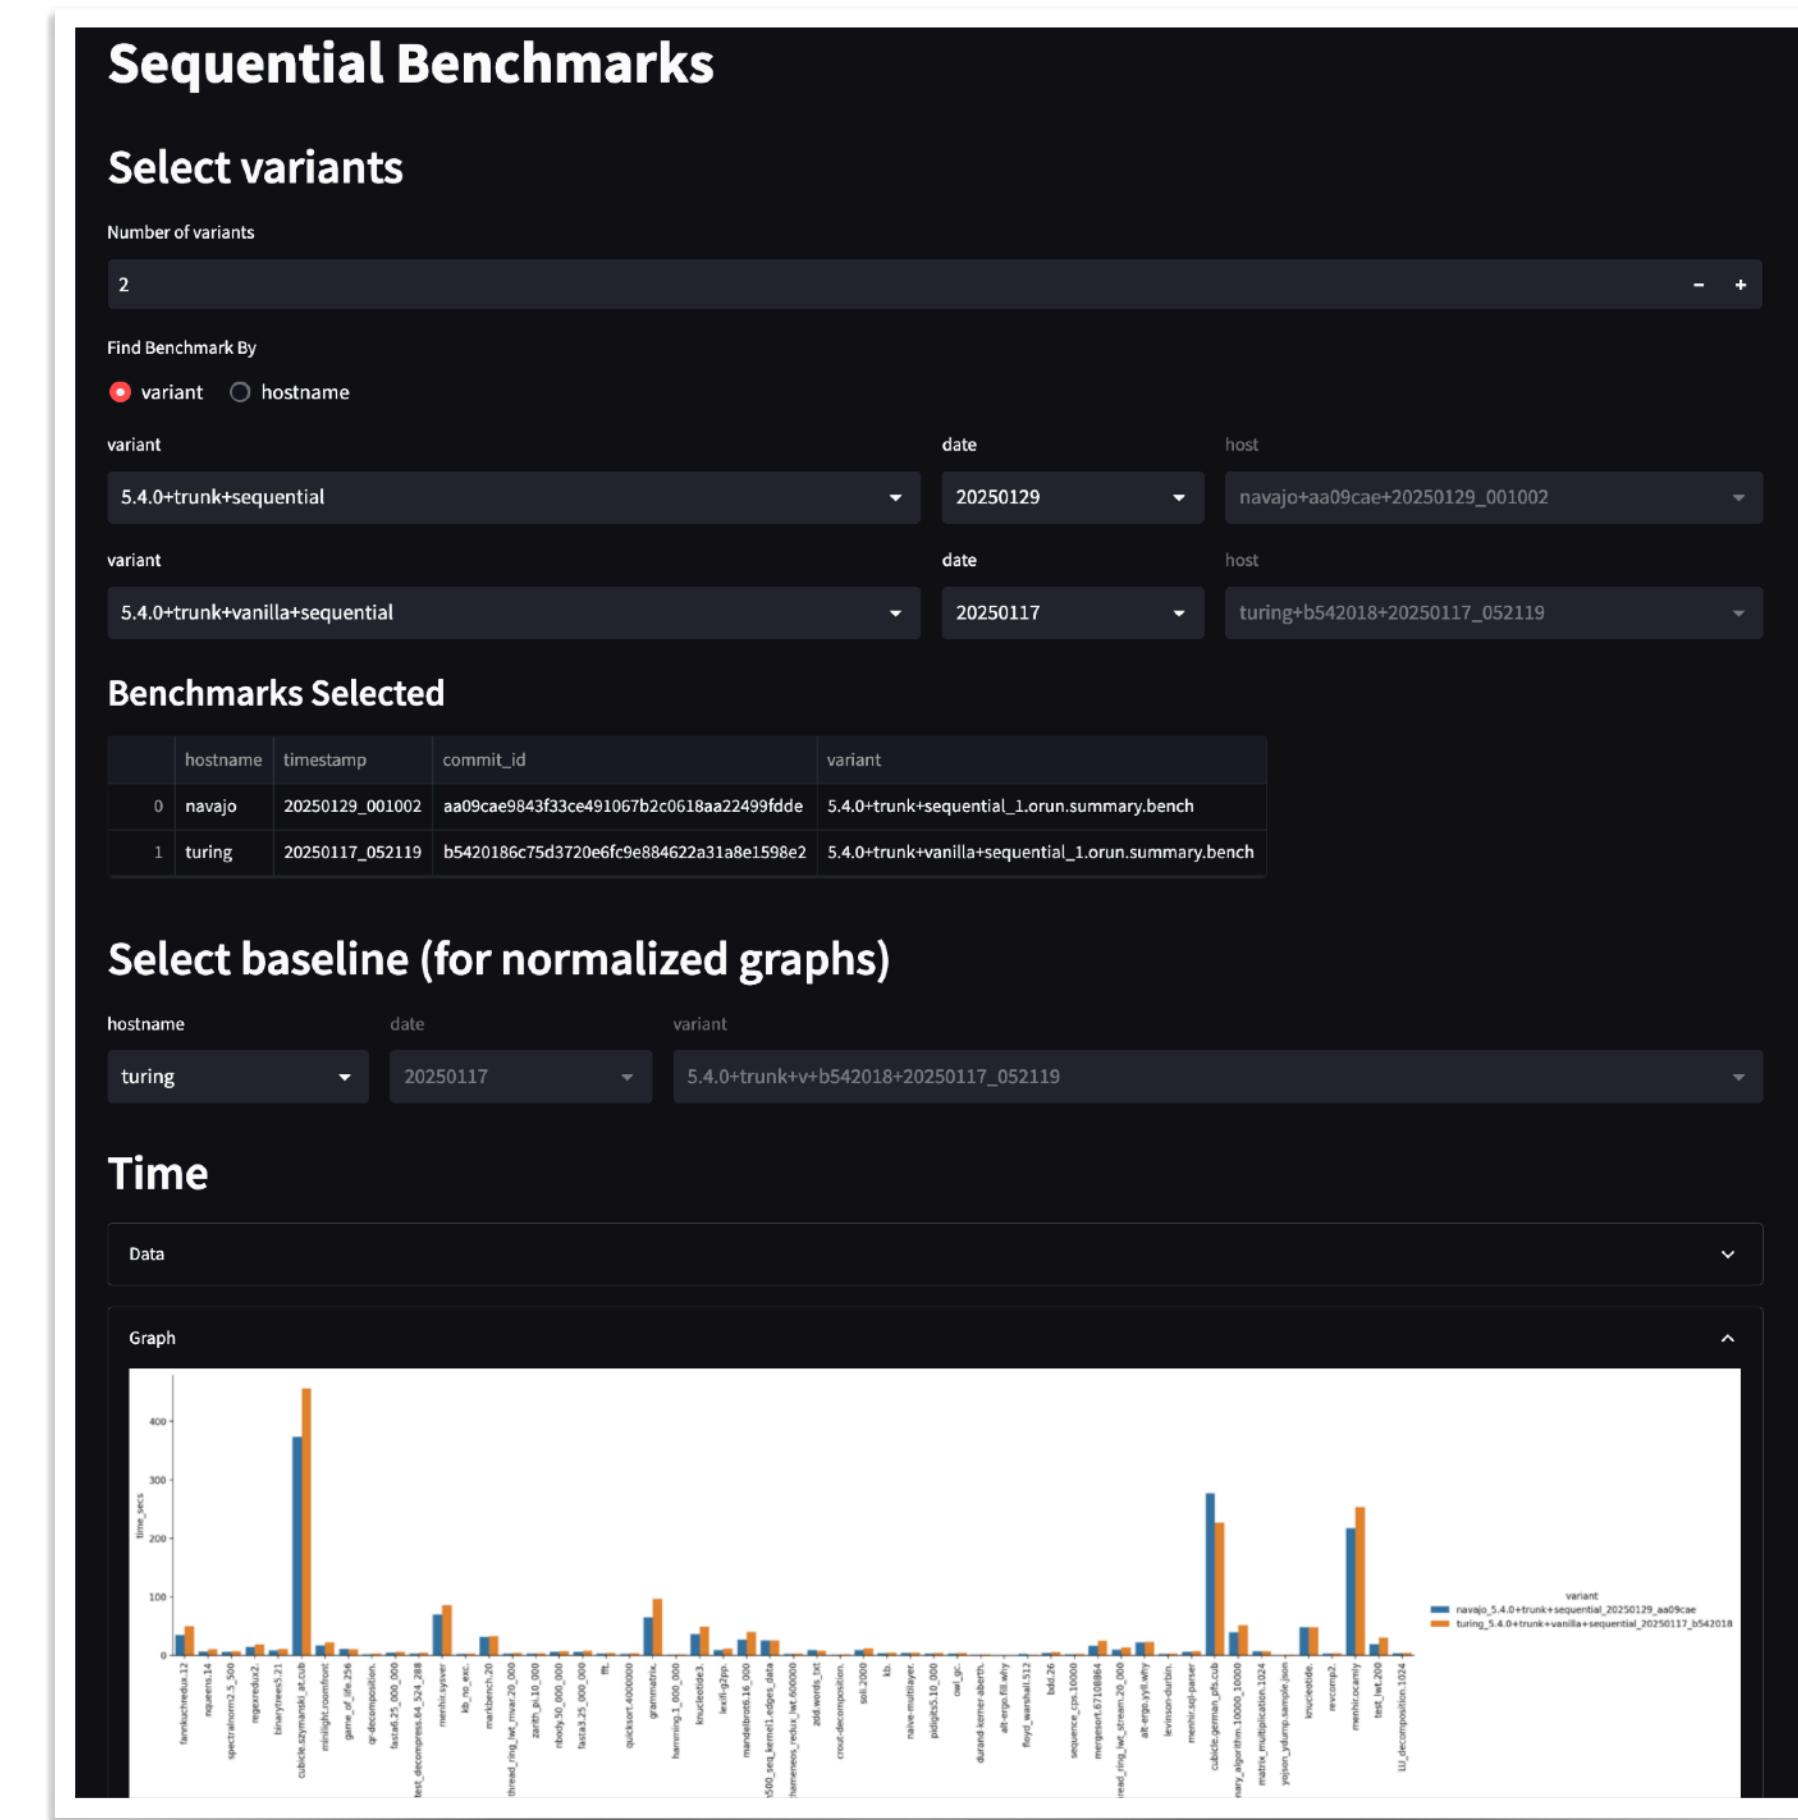
This is an HTML preview of the book, generated with [Hevea](#). License: [CC-BY-NC-ND 4.0](#).

- [Table of contents](#)
- [Introduction](#)

<https://xavierleroy.org/control-structures/>

# Building confidence – Benchmarking

- **Rigorous, continuous** benchmarking on *real-world programs*
  - [sandmark.tarides.com](http://sandmark.tarides.com) – Benchmark suite, Infra and runners



# Building confidence – CI for package universe

- *Can the new compiler build the existing universe?*
  - Build the OPAM universe of packages against *upstream* and *multicore* compilers

	4.14	5.0+alpha-repo	number of revdeps
0install.2.18	✓	✗	1
BetterErrors.0.0.1	✓	✗	7
TCSLib.0.3	✓	✗	1
absolute.0.1	✓	✗	0
acgtk.1.5.3	✓	✗	0
advi.2.0.0	✓	✗	0
aez.0.3	✓	✗	0
ahrocksdb.0.2.2	✗	✗	0
aio.0.0.3	✓	✗	0
alt-ergo-free.2.2.0	✓	✗	7
amqp-client-async.2.2.2	✓	✗	0
amqp-client-lwt.2.2.2	✓	✗	0
ancient.0.9.1	✓	✗	0
apron.v0.9.13	✓	✗	17

*You can contribute to the compiler development without hacking on the compiler*

# Release and Long Tail

- **Opened – Dec 2021, Merged – Jan 2022**
  - ....*A few months of iteration to fix design issues and bugs....*

Multicore OCaml #10831

Merged xavierleroy merged 4,103 commits into [ocaml:trunk](#) from [ocaml-multicore:multicore-pr](#) on Jan 10, 2022

Conversation 393 Commits 250 Checks 0 Files changed 300+ +22,955 -14,062

kayceesrk commented on Dec 21, 2021 · edited Member ...

This PR adds support for shared-memory parallelism through domains and direct-style concurrency through effect handlers (without syntactic support). It intends to have backwards compatibility in terms of language features, C API, and also the performance of single-threaded code.

**For users**

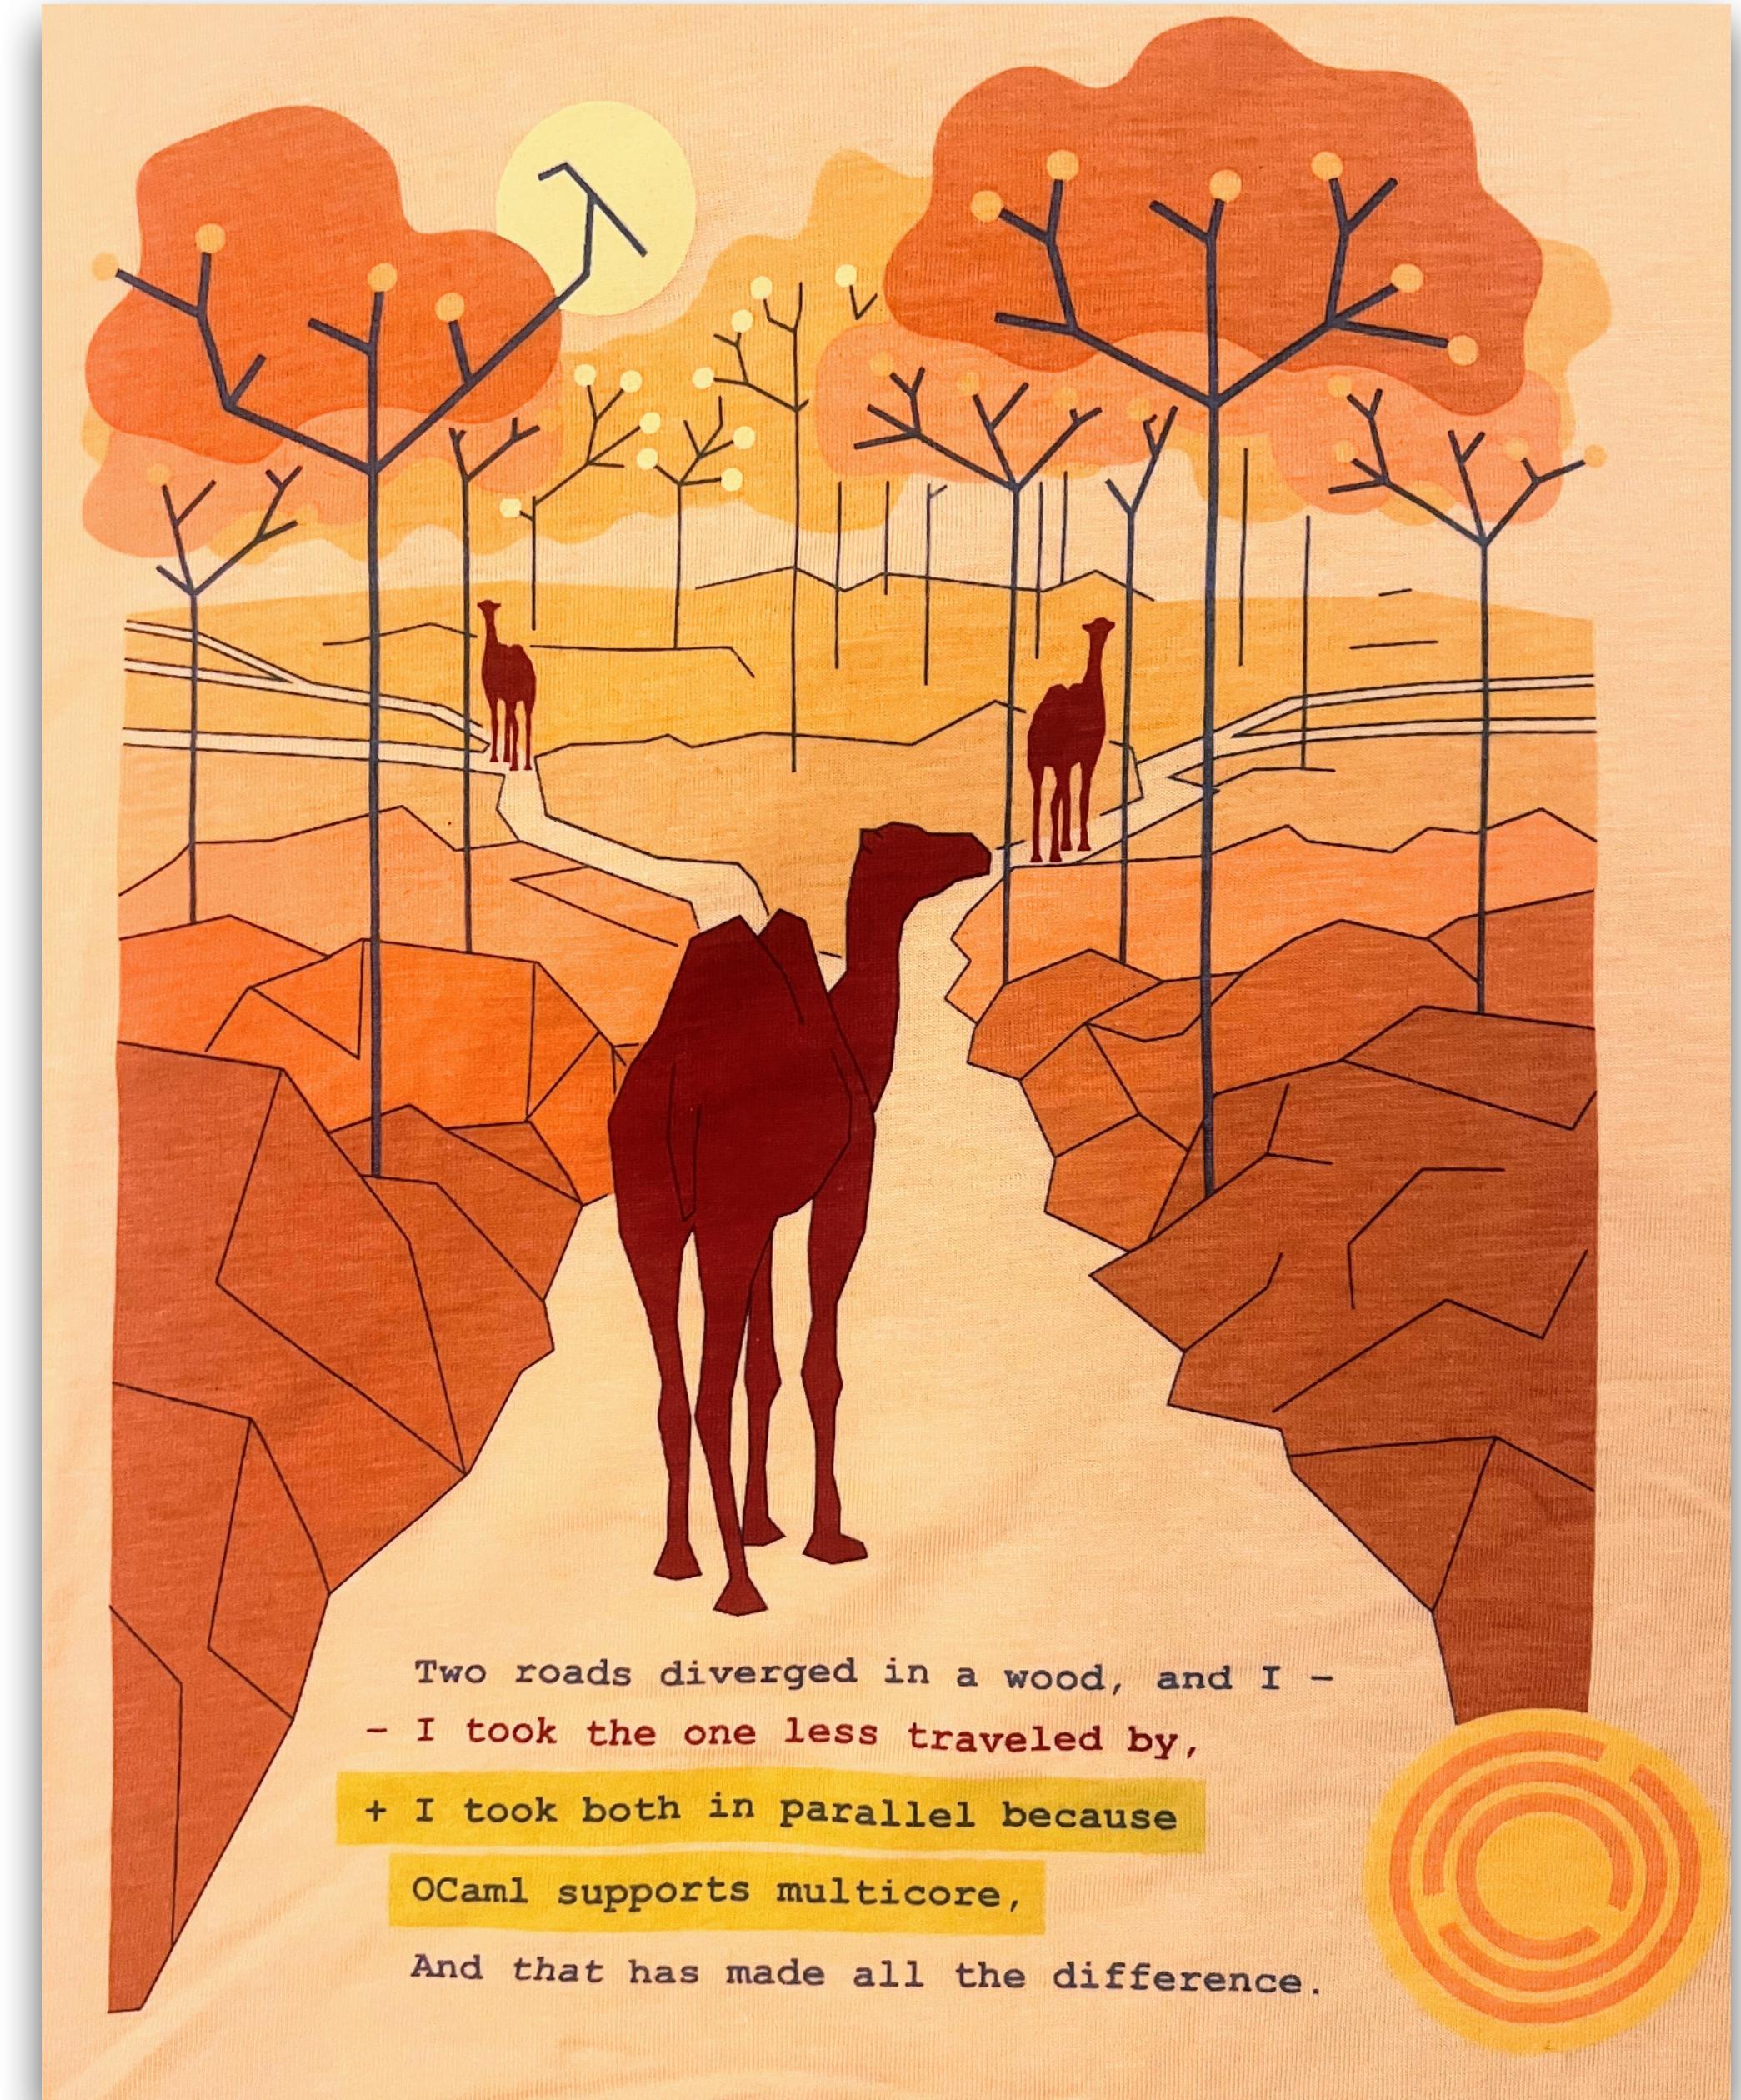
If you want to learn more about Multicore OCaml, please have a look at the [multicore](#)

Reviewers

- abbysmal
- gasche
- sadiqj
- avsm
- xavierleroy
- damiendoligez
- dra27

# Release and Long Tail

- **Opened** – Dec 2021, **Merged** – Jan 2022
  - ....A few months of iteration to fix design issues and bugs....
- **Released** – Dec 16 2022, as OCaml 5.0
- **Long tail** of adding missing features, bug fixes and performance improvements
  - 5.1 – Sep 2023
  - 5.2 – May 2024
  - 5.3 – Jan 2025
  - 5.4 – Sep 2025



# What's next for OCaml?

- **OxCaml** – Bridging the performance and safety gap between OCaml and Rust
  - *Data-race-free parallelism* through *modes*
  - Better control over object layout, allocations and GC
- Draws lessons from Multicore OCaml execution
  - Several award-winning papers at POPL, ICFP, OOPSLA
  - CI for the external universe – <https://oxcaml.check.ci.dev/>
- But different in other ways...
  - In production at Jane Street
  - Valuable user-feedback-oriented design



# OxCaml tutorial

The screenshot shows a GitHub README page with the following content:

**A Guided Tour Through Oxidized Caml**

This will be held at ICFP/SPLASH 2025, with inputs variously from Gavin Gray, Anil Madhavapeddy, KC Sivaramkrishnan, Richard Eisenberg, Chris Casinghino, Will Crichton, Shriram Krishnamurthi, Patrick Ferris.

URL: <https://conf.researchr.org/track/icfp-splash-2025/icfp-splash-2025-tutorials>

We have two 90 minute sessions that repeat, first at 1400-1530 and then 1600-1730.

## Getting started

Use the pre-built Docker Hub image.

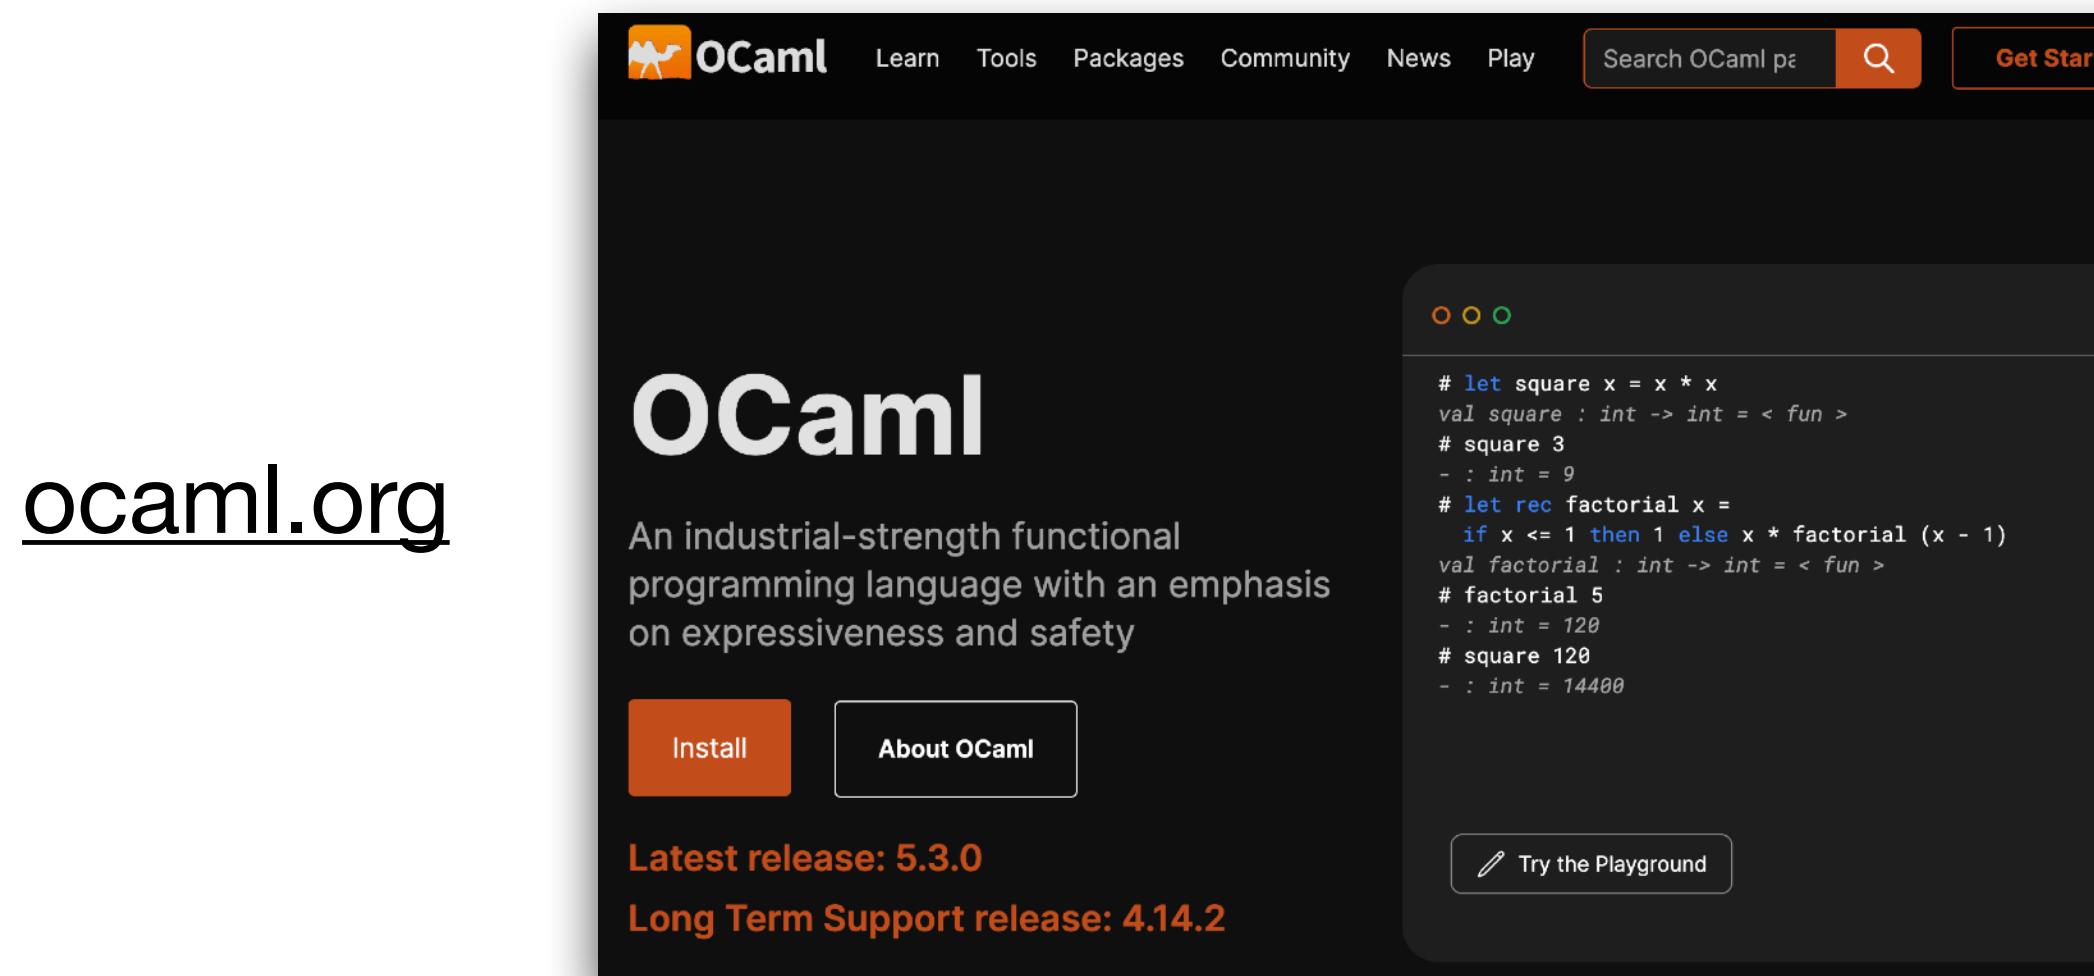
1. Open this folder in VS Code
2. Click "Reopen in Container" when prompted
  - OR: Cmd/Ctrl + Shift + P → "Dev Containers: Reopen in Container"

The docker image has three switches installed

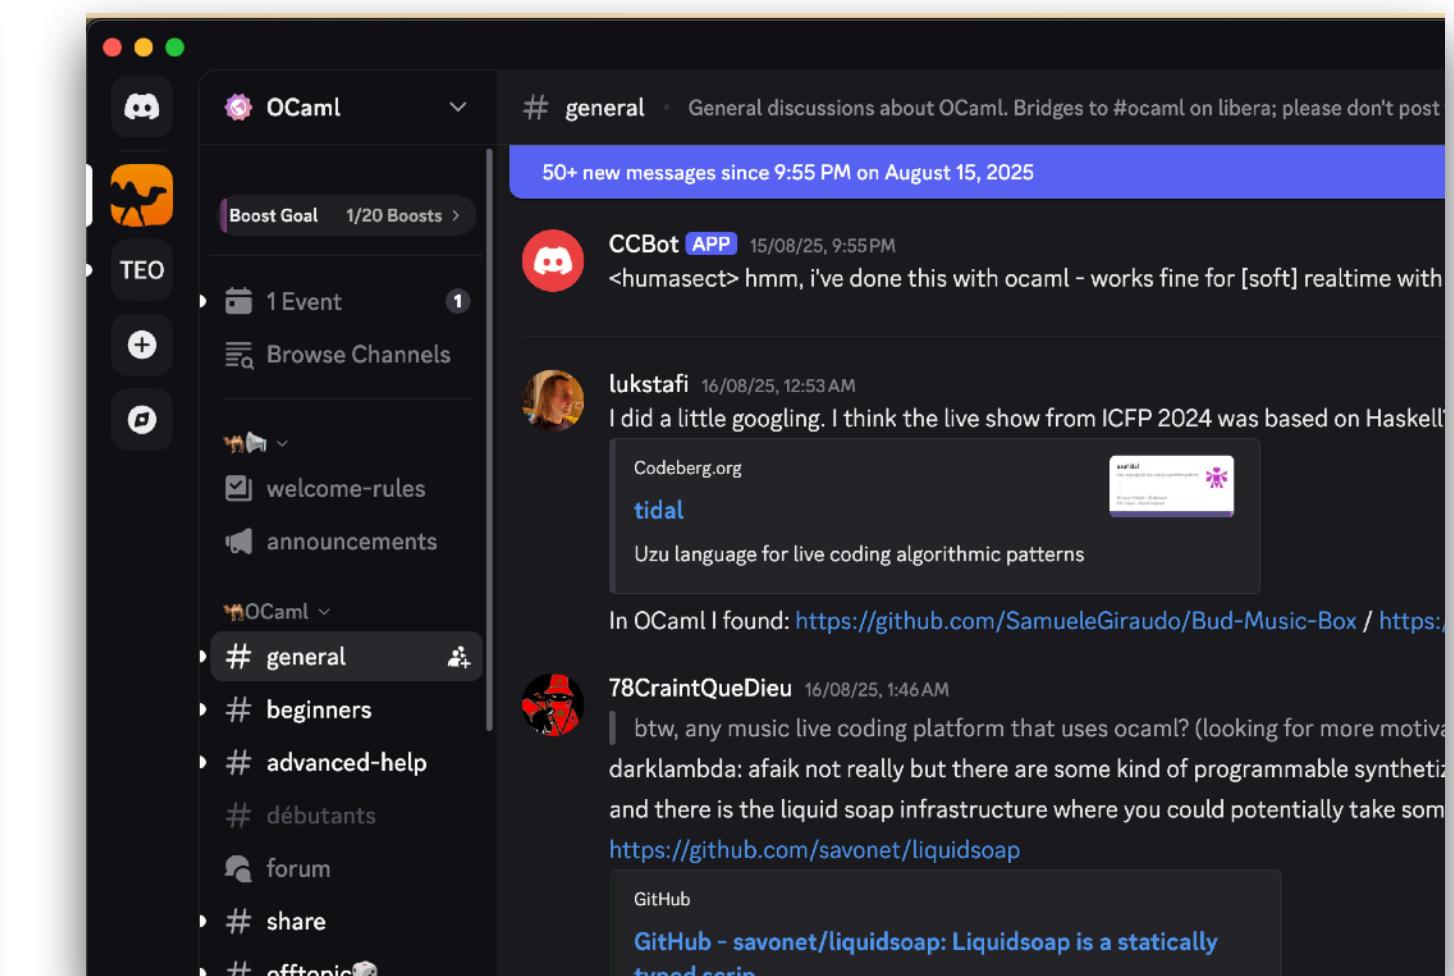
```
$ opam switch
# switch      compiler                               description
→ 5.2.0+ox   ocaml-variants.5.2.0+ox              5.2.0+ox
             ocaml-base-compiler.5.3.0                5.3
             ocaml-option-tsan.1,ocaml-variants.5.3.0+options 5.3.0+tsan
```

<https://github.com/oxcaml/tutorial-icfp25>

# Get Involved!



The OCaml website homepage features a dark background with the OCaml logo at the top. Below it, the word "OCaml" is prominently displayed in large white letters. A subdomain "ocaml.org" is shown below. A central text block states: "An industrial-strength functional programming language with an emphasis on expressiveness and safety". It includes links for "Install" and "About OCaml". Below this, release information is provided: "Latest release: 5.3.0" and "Long Term Support release: 4.14.2". On the right side, there's a code editor window showing OCaml code for calculating the factorial of 128.



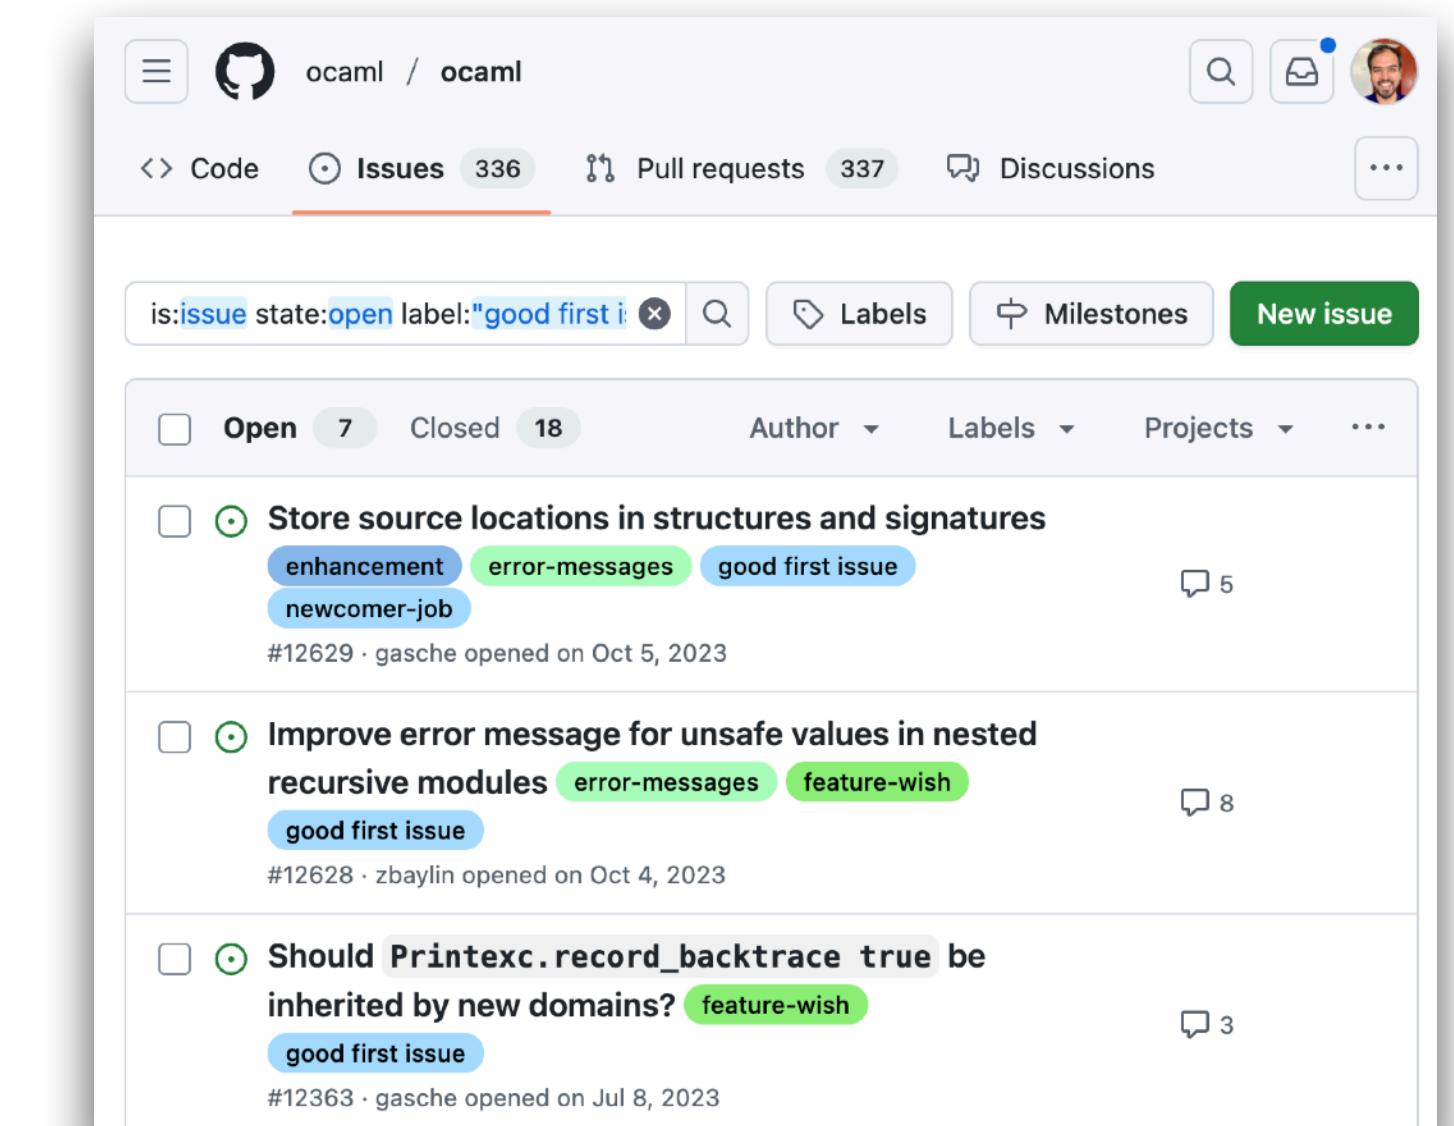
A screenshot of the OCaml Discord server interface. The left sidebar shows various channels like "# general", "# beginners", and "# advanced-help". The main #general channel has a message from CCBot (@CCBot) about a live show. Other messages include one from lukstafi (@lukstafi) about ICFP 2024, and one from 78CraintQueDieu (@78CraintQueDieu) about a music live coding platform.

OCaml  
Discord



The OCaml Outreachy Internships page is part of the ocaml.org website. It features a dark header with the OCaml logo and the text "Community > Outreachy Internships". Below this, a large image shows a person working on a laptop. The text "OCaml Outreachy Internships" is prominently displayed. It explains that Outreachy offers internship projects for people subject to systemic bias and impacted by underrepresentation in the technical industry. It mentions that Outreachy internship projects include programming, research, documentation, data science, and more. A button at the bottom says "Learn more at Outreachy".

ocaml.org/outreachy



A screenshot of the GitHub repository "ocaml/ocaml" showing the "Issues" tab. The search bar contains the query "is:issue state:open label:'good first issue'". Three issues are listed: 1) "Store source locations in structures and signatures" (opened by gasche on Oct 5, 2023), 2) "Improve error message for unsafe values in nested recursive modules" (opened by zbaylin on Oct 4, 2023), and 3) "Should Printexc.record\_backtrace true be inherited by new domains?" (opened by gasche on Jul 8, 2023).

github.com/ocaml