

# Irmin: Immutable, Fast, Distributed Storage

KC Sivaramakrishnan ([kcsrk.info](http://kcsrk.info))

Winter School on Decentralised Trust and Blockchains 2025  
19th December 2025



IIT  
MADRAS  
SAGDA M





- A distributed database built on the same principles as **Git**
  - Irmin is an **OCaml** library for building *mergeable, branchable distributed data stores.*



#### Built-In Snapshotting

Backup and restore your data at any point in time.



#### Storage Agnostic

You can use Irmin on top of your own storage layer.



#### Custom Datatypes

Automatic (de)serialisation for custom data types.



#### Highly Portable

Runs anywhere from Linux to web browsers and Xen unikernels.



#### Git Compatibility

Bidirectional compatibility with the Git on-disk format. Irmin state can be inspected and modified using the Git command-line tool.



#### Dynamic Behaviour

Allows users to define custom merge functions and create event-driven workflows using a notification mechanism.



## Language

- Functional-first but multi-paradigm (imperative, OO)
- Static-type system with Hindley-Milner type inference
- Advanced features – powerful module system, GADTs, Polymorphic variants
- Multicore support and ***effect handlers***

## Platform

- Fast, native code— x86, ARM, RISC-V, etc.
- JavaScript and WebAssembly (using ***WasmGC***) compilation
- Platform tools — editor (LSP), build system (dune), package manager (opam), docs generator (odoc), etc.

## Ecosystem

- Opam repository – small but mature package ecosystem
- Notable Industrial users – Jane Street, Meta, Microsoft, Ahrefs, Citrix, ***Tezos***, Bloomberg, Docker




- A blockchain founded in 2018, \$232M in ICO
- Innovations
  - The first proof-of-stake blockchain
  - On-chain governance
    - Protocol upgrades are built into the system
    - Soft upgrades without hard forks
- **Michelson** is the smart contract language
- **Protocol safety and formal verification** through strong typing, OCaml, and Rocq semantics
- **Irmin** is the distributed database and storage layer
  - Peers gossip blocks and state info, receive blocks, update state




- A distributed version control system
  - .... also a great model for a local-first, asynchronous, distributed system
- Do some work locally and then ...
  - **Branch** to create temporary copies efficiently
  - **Pull** remote branches to get remote updates
  - **Push** to remotes to get your local changes to upstream
- How does this work efficiently?

# Git store data model




- Branches / tags
- Mutable
- Stores the files under version control
- Immutable, append-only & content addressed
  - hash → object

# Block store – Persistence and Merkle DAG




Example from Pro Git book: <https://git-scm.com/book/en/v2/Git-Internals-Git-Objects>


# Merge commits



# Excellent distributed model



Replica 1




Replica 2

- Distributed communication
  - Pull changes from remote — block store has no conflicts!
  - Merge is local uses merge commit
- Strong integrity guarantees
  - Content-addressed storage (hashes), Merkle DAG of commits
  - History is immutable once referenced
  - This ensures that you can detect tampering

# Commit history forms a DAG

- Captures the causal history of the distribution
- ***Lowest common ancestor (LCA)*** commit always exists
  - Assuming there was an origin commit
- LCA represents the point in the history where the two commits/branches diverged
- What is the LCA of
  - E & B?
  - E & C?




# Merkle Proofs


- The Tezos network builds trust between its nodes by using two components
  - a **tamper-proof database** (**Irmin**) that can generate cryptographic hashes, which uniquely and compactly represent the state of its contents; and
  - a consensus algorithm to share these cryptographic hashes across the network of (potentially adversarial) nodes.
- **Merkle proofs** let Tezos prove facts about the blockchain state without sharing the entire state.
  - Essential for scalability, decentralisation, and trust minimisation.

# Merkle Proofs

Merkle Tree




Merkle Proof




- Merkle proof is a compact representation of the Merkle tree
  - 100 ops, Merkle Proof = 46kB, Merkle Tree = 3.4 GB

# Local-first software

# Collaborative Applications




# Collaborative Applications




# Network Partitions

- Centralised Apps provide limited support for offline editing




*Enabling offline sync for one account prevents other accounts from working offline*

# Local-first software



# Local-first software




*How do we build such applications?*

Embed the notion of **replication** into the  
**data types**

# Mergeable Replicated Data Types (MRDTs)

- MRDTs = Sequential data types + 3-way merge function à la Git




*Thanks to  
Irmin!*

```
let merge lca v1 v2 =  
  lca + (v1 - lca) + (v2 - lca)
```

# MRDT Set

- A replicated set with ***add***, ***remove*** and ***read*** operations




```
let merge lca v1 v2 =  
  (lca ∩ v1 ∩ v2) (* unchanged or removed elements *)  
  ∪ (v1 \ lca) (* added elements on left *)  
  ∪ (v2 \ lca) (* added elements on right *)
```

- Do you foresee any issues with the implementation?
  - *Conflicts between concurrent add and remove of same elements!*

# MRDT Set


- By default, “remove wins”



```
let merge lca v1 v2 =  
  (lca ∩ v1 ∩ v2) (* unchanged or removed elements *)  
  ∪ (v1 \ lca) (* added elements on left *)  
  ∪ (v2 \ lca) (* added elements on right *)
```

# MRDT Set

- How to do “add wins”?



- Add associates a **fresh id** with the element
- Remove removes all matching elements with **any id**
- Read returns the set removing ids
- Merge remains unchanged

```
let merge lca v1 v2 =  
  (lca ∩ v1 ∩ v2) (* unchanged or removed elements *)  
  ∪ (v1 \ lca) (* added elements on left *)  
  ∪ (v2 \ lca) (* added elements on right *)
```

Why is this  
correct?

How do we  
automatically verify it?

# Algebraic properties are insufficient

$$\mu(a, b) = \mu(b, c)$$

$$\mu(a, a) = a$$

$$\mu(\mu(a, b), c) = \mu(a, \mu(b, c))$$

Commutativity


Idempotence

Associativity


```
let merge v1 v2 = max v1 v2
```

Satisfies  
algebraic  
properties

*Intent is not  
captured*



# *Is there a more natural spec?*



$\sigma_5 = \sigma_6 = \text{linearization}(\{op_1, op_2, op_3, op_4\}) \sigma_0$

# Replication-aware Linearizability

RESEARCH-ARTICLE

## Replication-aware linearizability

Authors:  Chao Wang,  Constantin Enea,  Suha Orhun Mutluergil,  Gustavo Petri | [Authors Info & Clai](#)

PLDI 2019: Proceedings of the 40th ACM SIGPLAN Conference on Programming Language Design and Implementation  
<https://doi.org/10.1145/3314221.3314617>

- Replica states should be a *linearisation* of observed *update* operations
  - Linearisation total order *lo* compatible with partially-ordered visibility relation *vis*
  - No real-time ordering requirement unlike traditional linearizability
- Payoff
  - If a replicated object is RA-linearizable, reason about it using sequential semantics

# Using RA-linearizability for verification


add(a);	add(a);
rem(a);	
X = read();	Y = read();

$$a \in X \implies a \in Y$$

- Since Add-wins set is RA-linearizable, you can use *totally ordered trace* and the *sequential spec* to reason about correctness

add(a);	rem(a);	add(a);	X = read();	Y = read()
$\{a\}$	$\{\}$	$\{a\}$	$X = \{a\}$	$Y = \{a\}$

# Using RA-linearizability for verification



- Let's try to make the statement false
  - Make  $a \in X$  true and  $a \in Y$  false

# Replication-aware Linearizability

RESEARCH-ARTICLE

## Replication-aware linearizability

**Authors:**  Chao Wang,  Constantin Enea,  Suha Orhun Mutluergil,  Gustavo Petri | [Authors Info & Clai](#)

[PLDI 2019: Proceedings of the 40th ACM SIGPLAN Conference on Programming Language Design and Implementation](#)  
<https://doi.org/10.1145/3314221.3314617>

- Presented a proof methodology to show that an RDT is linearisable
- Not automated or mechanised

# Neem – Automatic verification of RDTs


- What's in the box?
  - Definition of RA-linearizability for MRDTs
  - A novel induction scheme for MRDTs and state-based CRDTs to **automatically** verify RA-linearizability
  - Implemented in F\*

The screenshot shows a research article page from a digital library. At the top, it says "RESEARCH-ARTICLE | OPEN ACCESS |" followed by four small icons. Below that is the title "Automatically Verifying Replication-Aware Linearizability". Under the title, it lists "Authors: Vimala Soundarapandian, Kartik Nagar, Aseem Rastogi, KC Sivaramakrishnan" and links to "Authors Info & Claims". It also mentions the publication source: "Proceedings of the ACM on Programming Languages, Volume 9, Issue OOPSLA1 • Article No.: 111, Pages 871 - 897" and the DOI: "https://doi.org/10.1145/3720452". The date "Published: 09 April 2025" is shown with a link to "Publication History". There is a "Check for updates" button. In the bottom right corner, there are icons for a bell, a folder, a double quote, a PDF file, and an eReader.

The screenshot shows a GitHub repository page for "neem". The top bar shows the URL "github.com/prismlab/neem". The main content area has a "README" tab selected, showing the title "Neem" and a brief description: "Neem is a framework for automated verification of mergeable replicated data types (MRDTs) and state-based convergent replicated data types (CRDTs). See <https://dl.acm.org/doi/10.1145/3720452>". Below this is a "Development Environment" section with instructions: "Easiest way to get started is to use the devcontainer." and a code block: "\$ git clone https://github.com/prismlab/neem\n\$ cd neem\n\$ code . # Start VSCode". To the right, there are sections for "Packages", "Contributors" (with three entries: vimcy7, kayceesrk, aseemr), "Languages" (F\* 97.2%, Shell 2.4%, Dockerfile 0.4%), and "Suggested workflows".

# Verified MRDTs

MRDT	rc Policy	#LOC	Verification Time (s)
Increment-only counter [12]	none	6	0.72
PN counter [23]	none	10	1.64
Enable-wins flag*	disable $\xrightarrow{\text{rc}}$ enable	30	29.80
Disable-wins flag*	enable $\xrightarrow{\text{rc}}$ disable	30	37.91
Grows-only set [12]	none	6	0.45
Grows-only map [23]	none	11	4.65
OR-set [23]	$\text{rem}_a \xrightarrow{\text{rc}} \text{add}_a$	20	4.53
OR-set (efficient)*	$\text{rem}_a \xrightarrow{\text{rc}} \text{add}_a$	34	660.00
Remove-wins set*	$\text{add}_a \xrightarrow{\text{rc}} \text{rem}_a$	22	9.60
Set-wins map*	$\text{del}_k \xrightarrow{\text{rc}} \text{set}_k$	20	5.06
Replicated Growable Array [1]	none	13	1.51
Optional register*	$\text{unset} \xrightarrow{\text{rc}} \text{set}$	35	200.00
Multi-valued Register*	none	7	0.65
JSON-style MRDT*	Fig. 13	26	148.84



*Neem also supports verification of RA-linearizability of state-based CRDTs*

<https://github.com/prismlab/neem>

# Conclusion



- **Irmin** is an excellent foundation for efficient and principled distributed applications
  - Rich data model using MRDTs
  - Ability to verify the correctness of the application layer automatically through Neem
  - Can be compiled to various storage-efficient backends transparently
- Questions
  - Can we use replication-aware linearizability to **reconcile the DAG view to linear blockchain?**
    - Hedera Hashgraph and other DAGs are interesting alternatives to blockchain
  - Could Irmin act as a **state substrate** for L2s, rollups, or light clients?
  - Can we reason about **Byzantine behaviour** when merges are application-defined?
    - Should RA-linearizability be a prerequisite for user-defined merges?