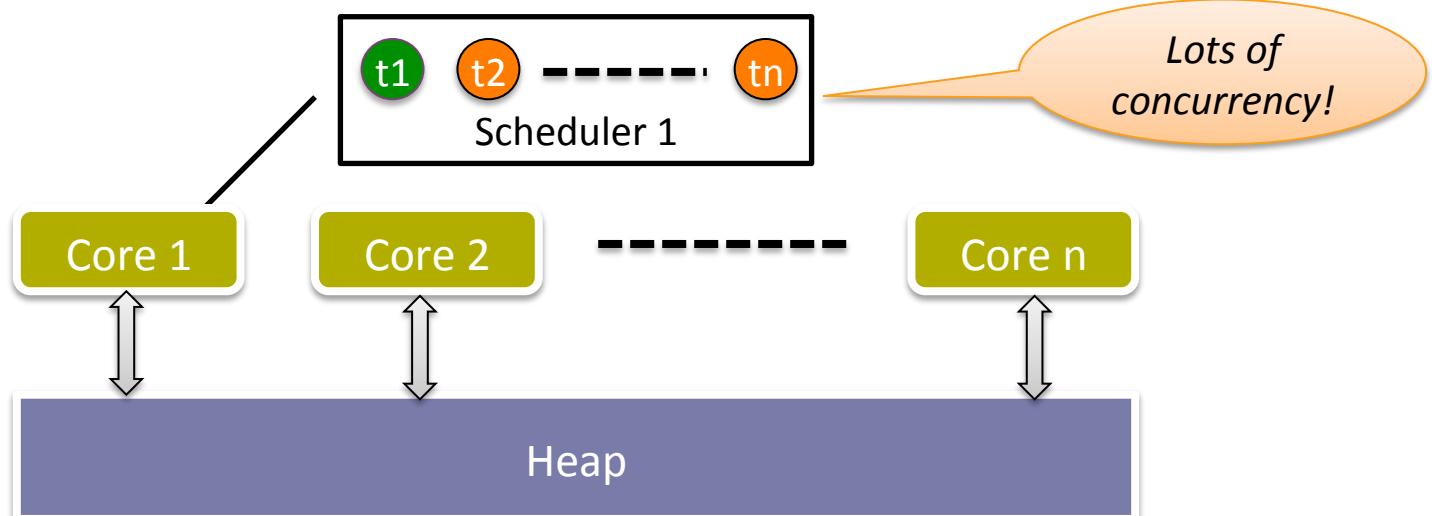
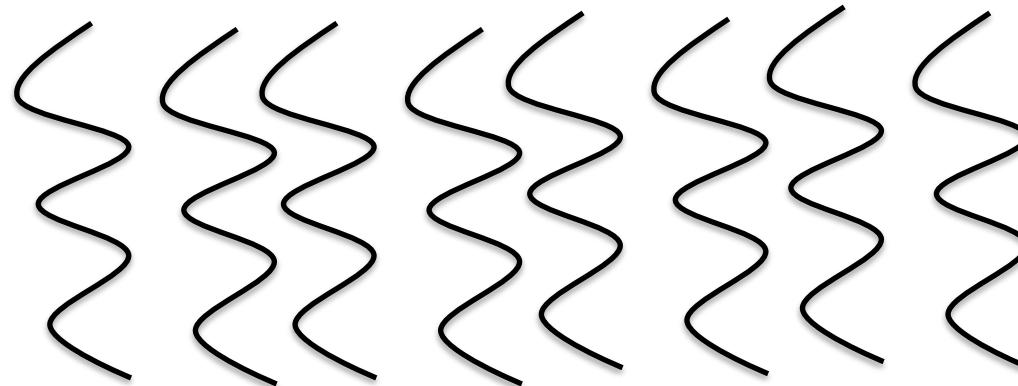


# **Eliminating Read Barriers through Procrastination and Cleanliness**

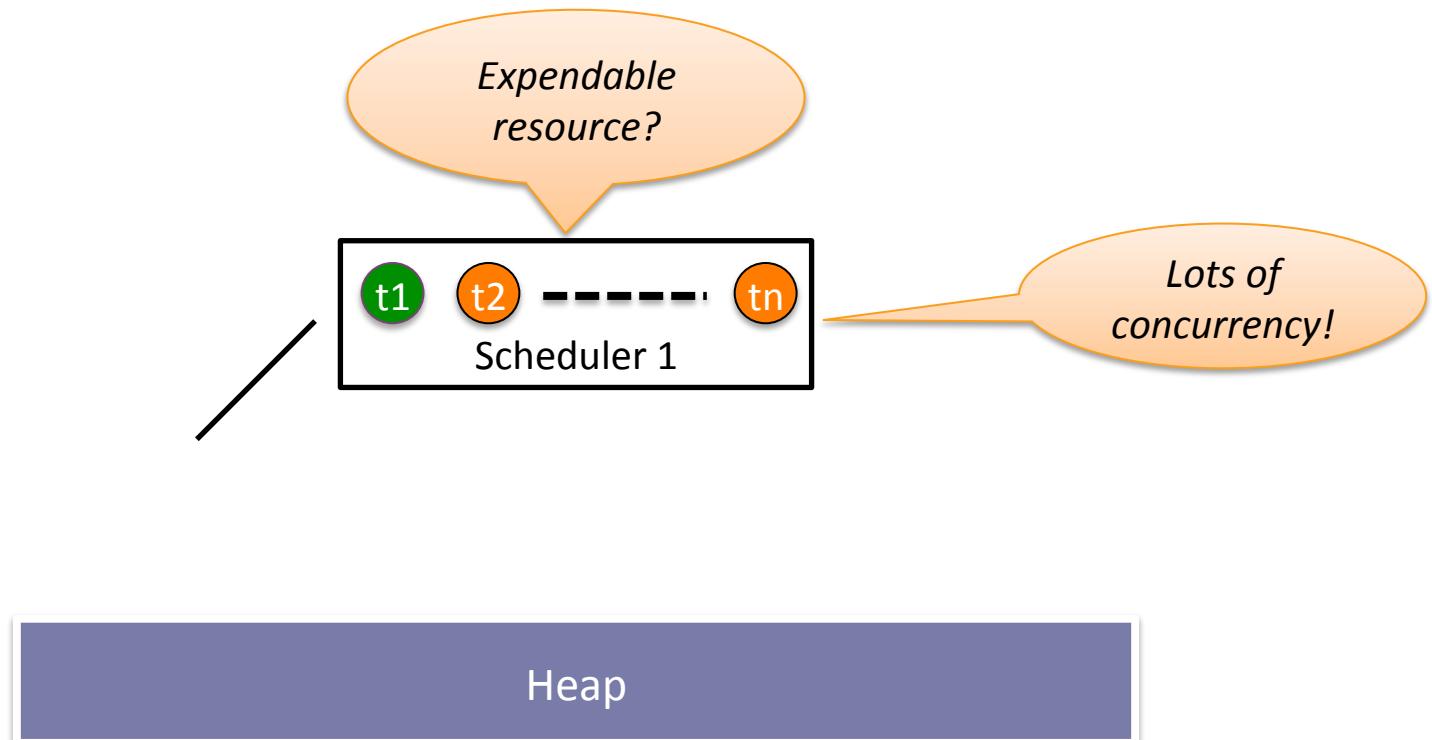
**KC Sivaramakrishnan**  
**Lukasz Ziarek**  
**Suresh Jagannathan**

# Big Picture

Lightweight user-level threads

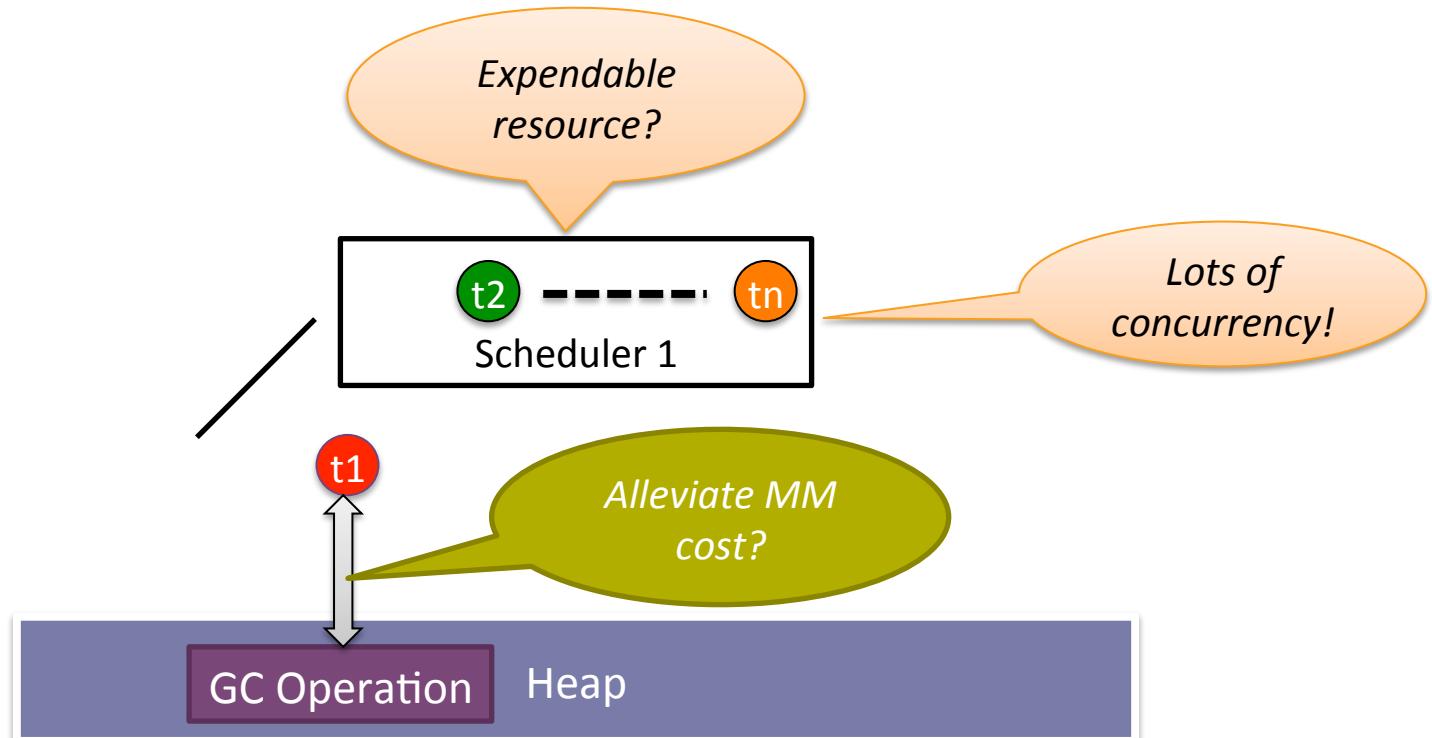


# Big Picture



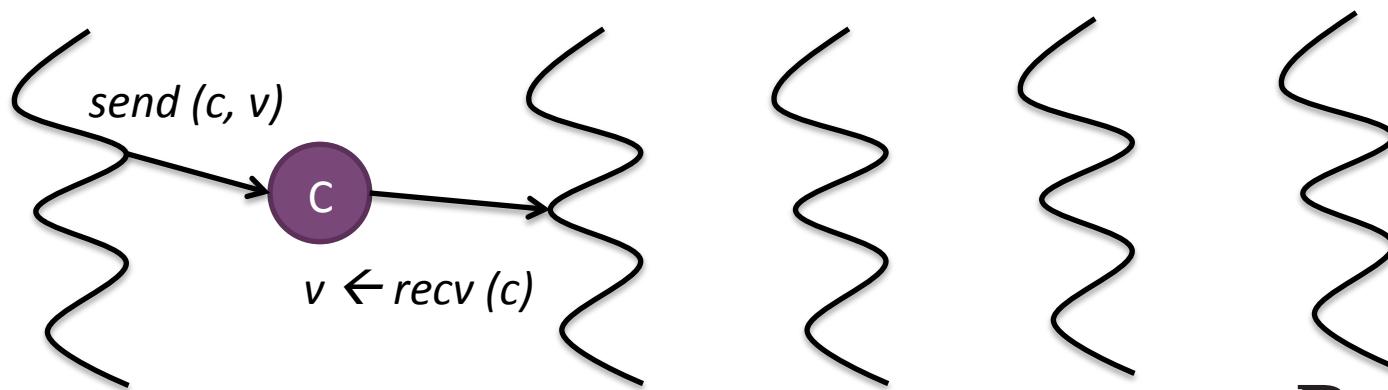
# Big Picture

Exploit program concurrency  
to  
eliminate read barriers from thread-local collectors



# MultiMLton

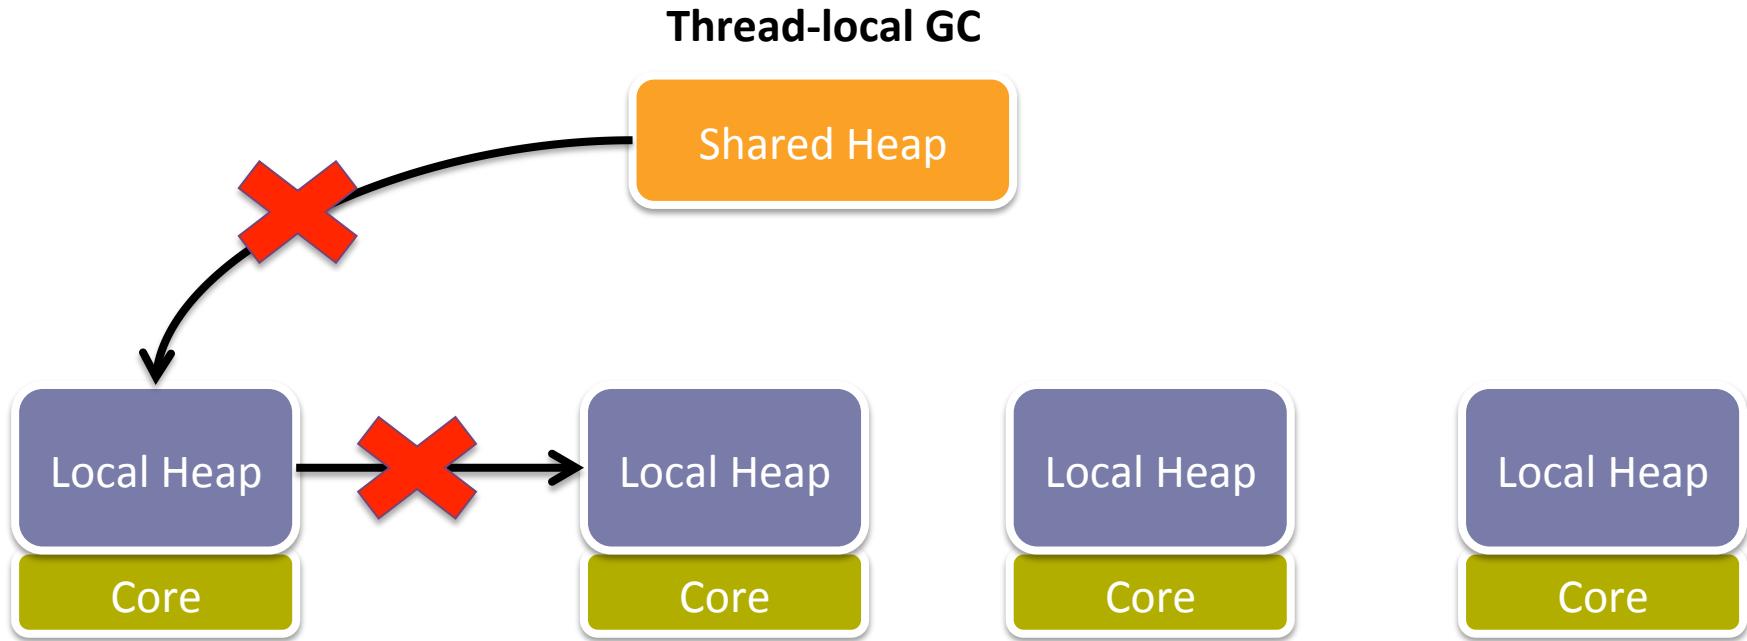
- Goals
  - Safety, Scalability, ready for future manycore processors
- Parallel extension of MLton – a whole-program, optimizing SML compiler
- Parallel extension of Concurrent ML
  - *Lots of Concurrency!*
  - Interact by sending messages over first-class channels



# MultiMLton GC: Considerations

- Standard ML – functional PL with side-effects
  - Most objects are small and ephemeral
    - Independent generational GC
  - # Mutations << # Reads
    - Keep cost of reads to be low
- Minimize NUMA effects
- Run on non-cache coherent HW

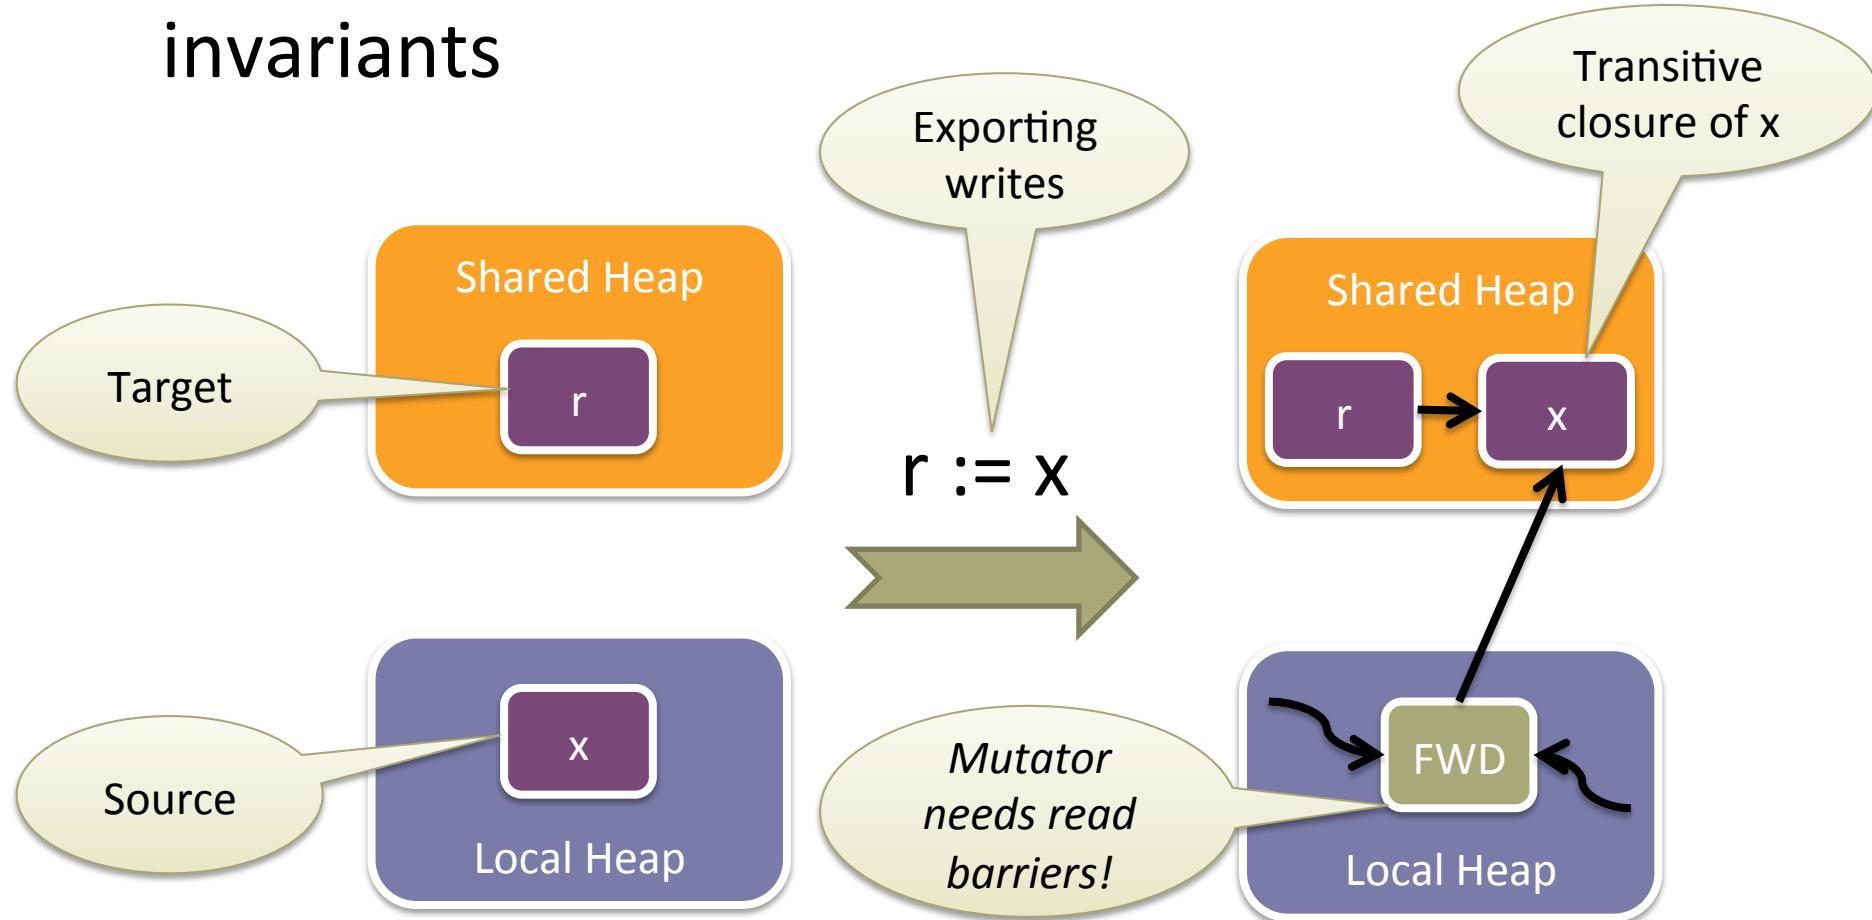
# MultiMLton GC: Design



- NUMA Awareness
- Circumvent cache-coherence issues

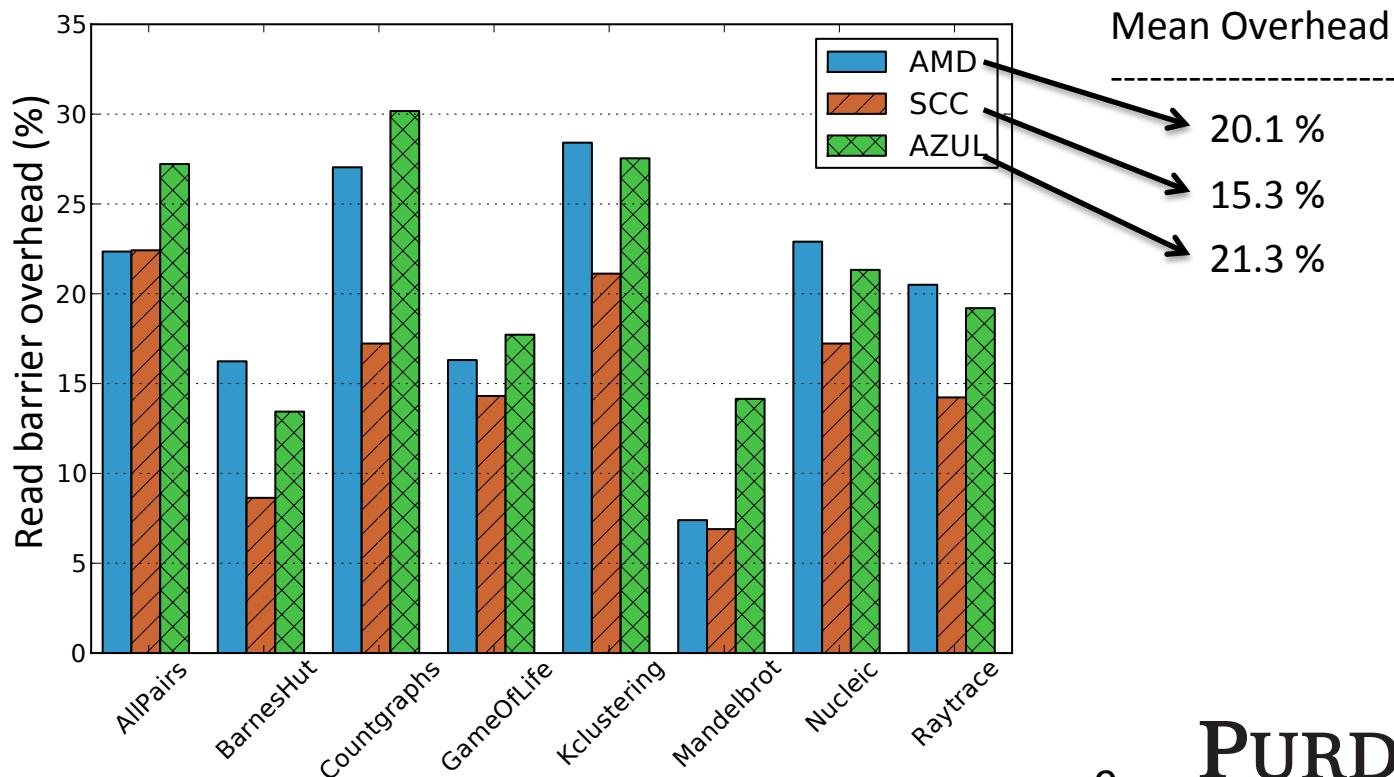
# Invariant Preservation

- Read and write barriers for preserving invariants



# Challenge

- Object reads are pervasive
  - RB overhead  $\propto$  cost (RB) \* frequency (RB)
- Read barrier optimization
  - Stacks and Registers never point to forwarded objects



# Mutator and Forwarded Objects

# Encountered  
forwarded objects

---

< 0.00001

# RB invocations

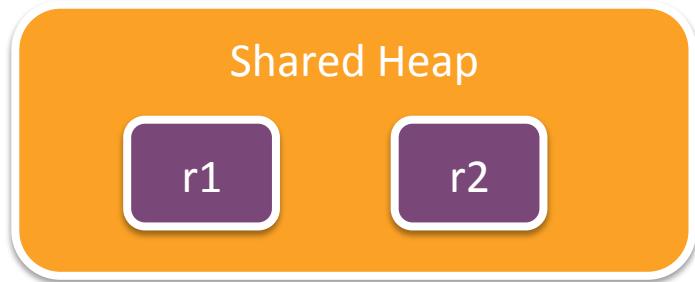
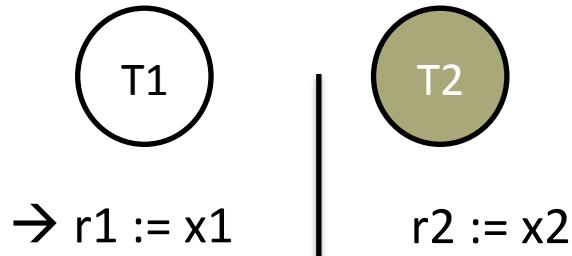


*Eliminate read barriers altogether*

# RB Elimination

- Visibility Invariant
  - Mutator does not encounter forwarded objects
- Observation
  - No forwarded objects created  $\Rightarrow$  visibility invariant  $\Rightarrow$  No read barriers
- Exploit concurrency  $\rightarrow$  ***Procrastination!***

# Procrastination

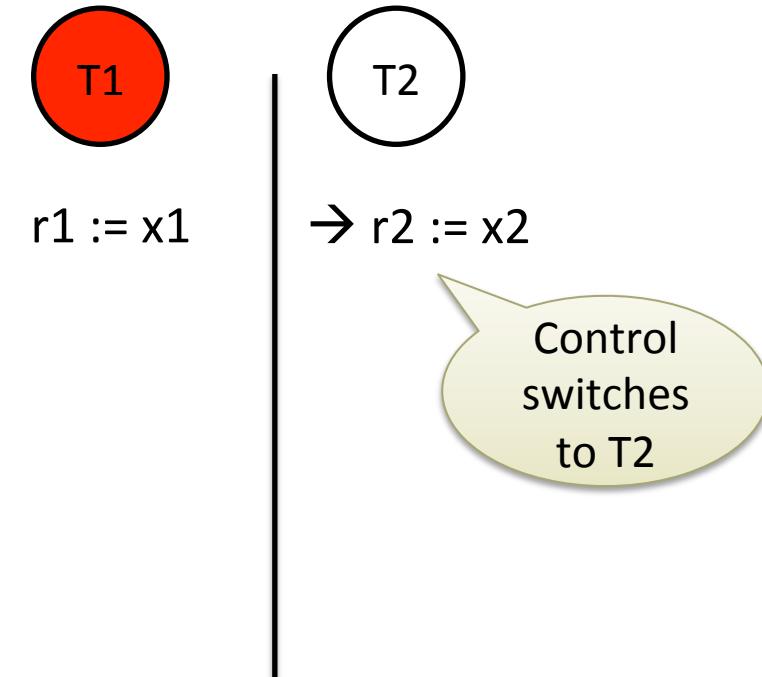


T → T is *running*

T → T is *suspended*

T → T is *blocked*

# Procrastination

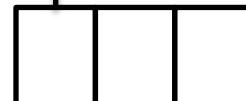


T → T is *running*

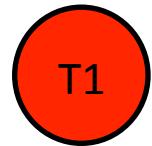
T → T is *suspended*

T → T is *blocked*

Delayed write list →



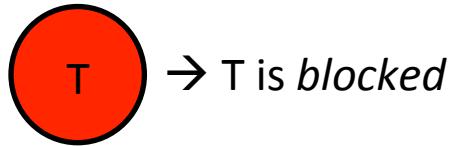
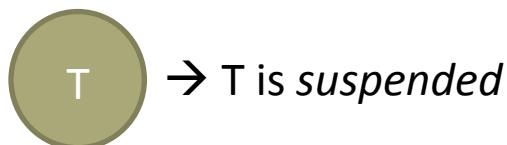
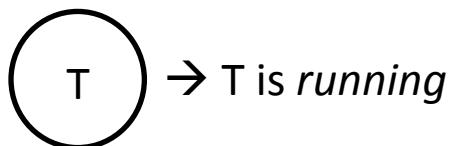
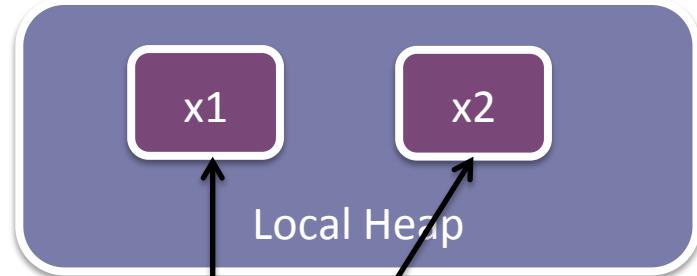
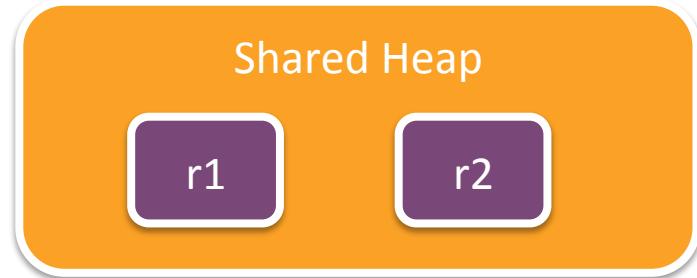
# Procrastination



r1 := x1



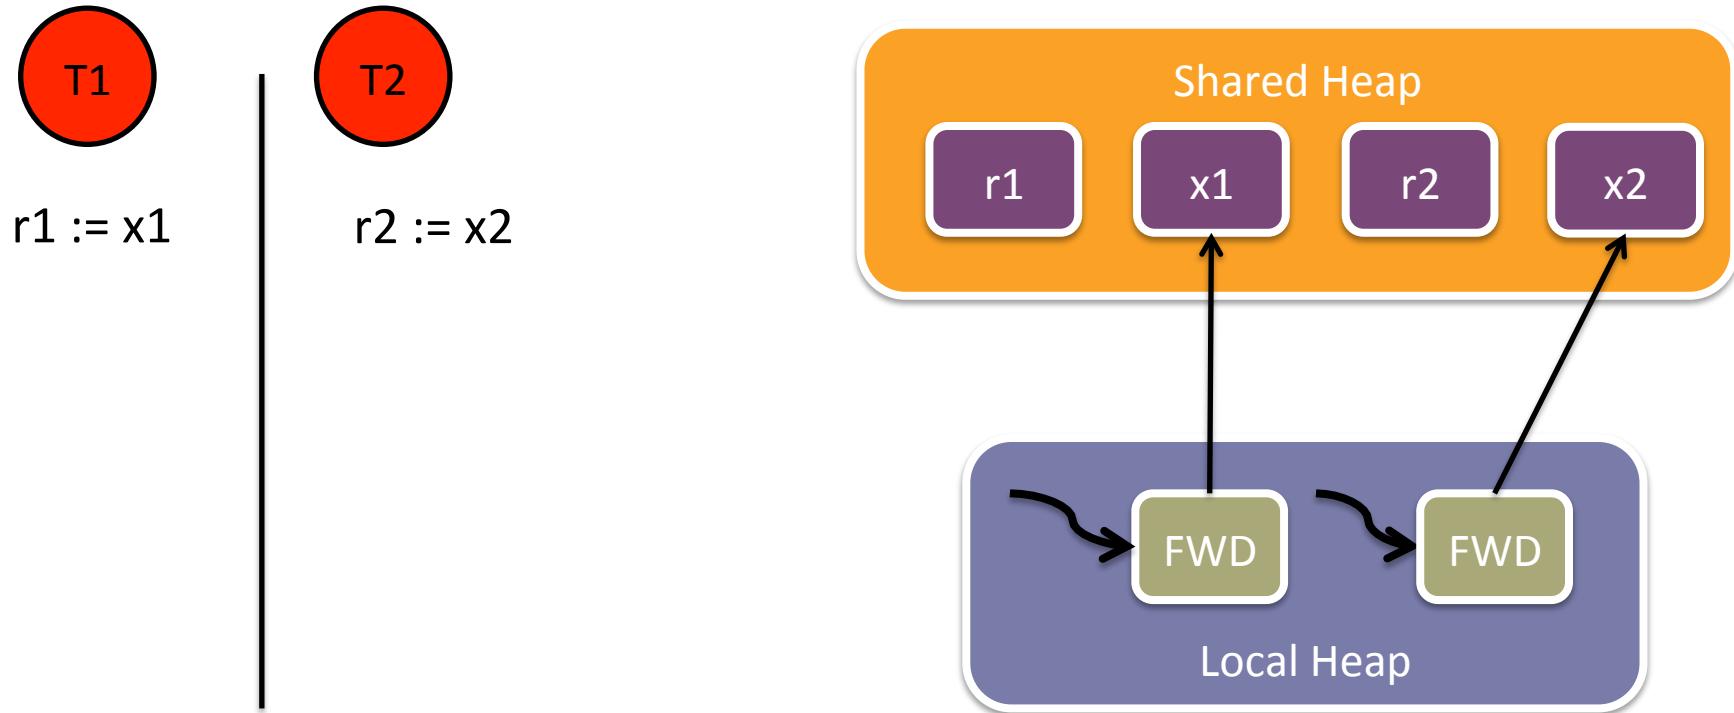
r2 := x2



Delayed write list →



# Procrastination



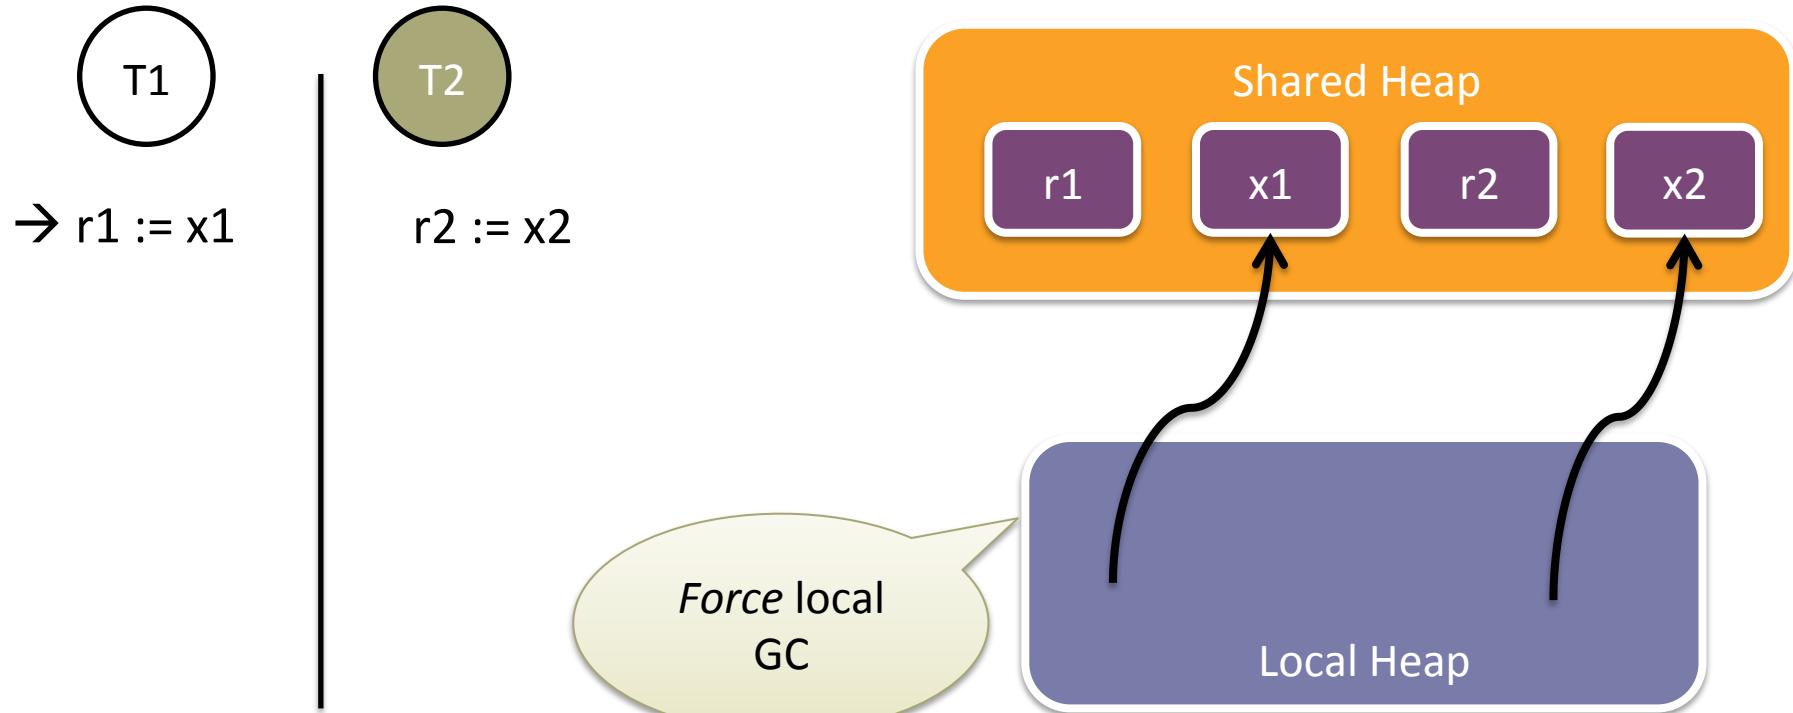
T → T is *running*

T → T is *suspended*

T → T is *blocked*

Delayed write list →  A horizontal row of three empty boxes separated by vertical lines, representing a list of delayed writes.

# Procrastination



T → T is *running*

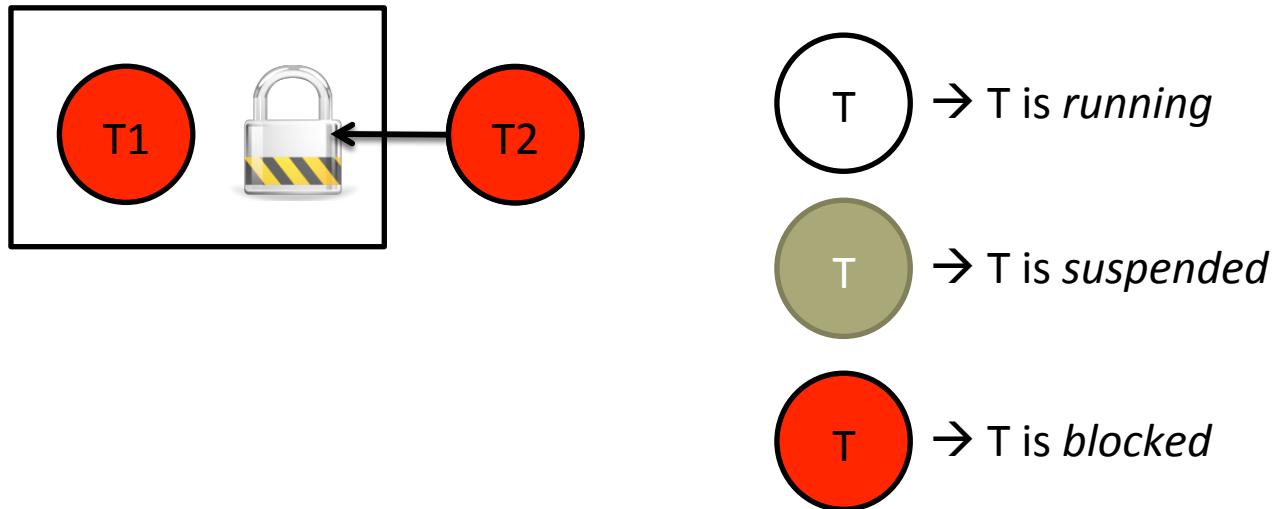
T → T is *suspended*

T → T is *blocked*

Delayed write list →  (three empty boxes)

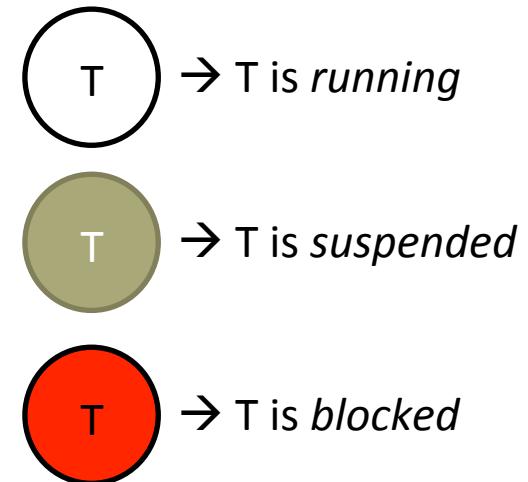
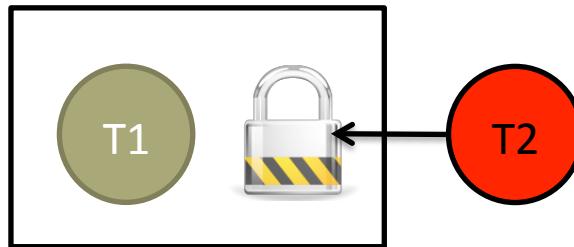
# Correctness

- Does Procrastination introduce deadlocks?
  - Threads can be procrastinated while holding a lock!



# Correctness

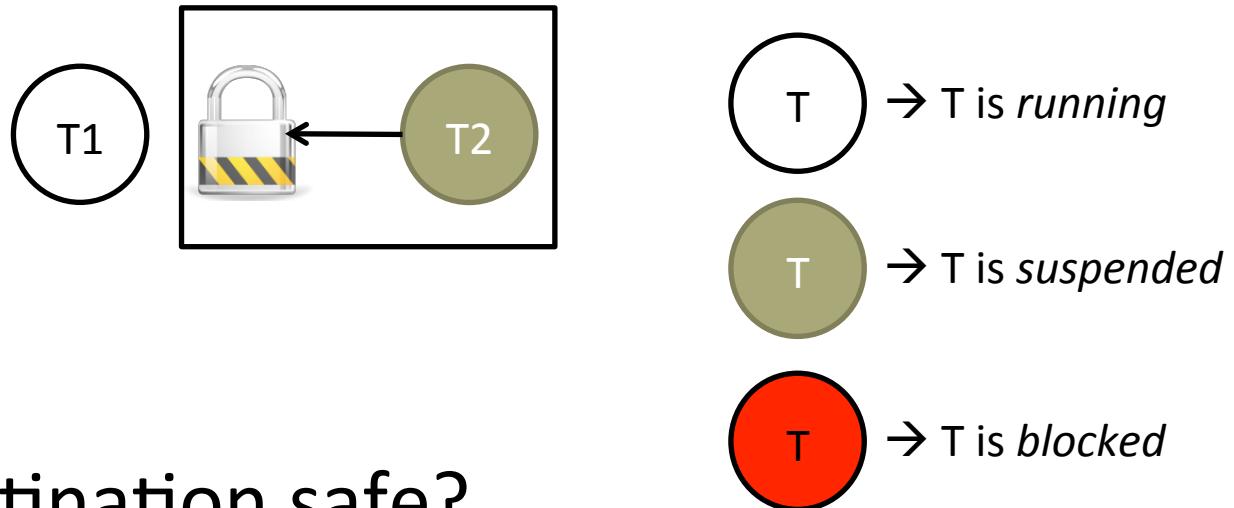
- Does Procrastination introduce deadlocks?
  - Threads can be procrastinated while holding a lock!



- Is Procrastination safe?
  - Yes. Forcing a local GC unblocks the threads.
  - No deadlocks or livelocks!

# Correctness

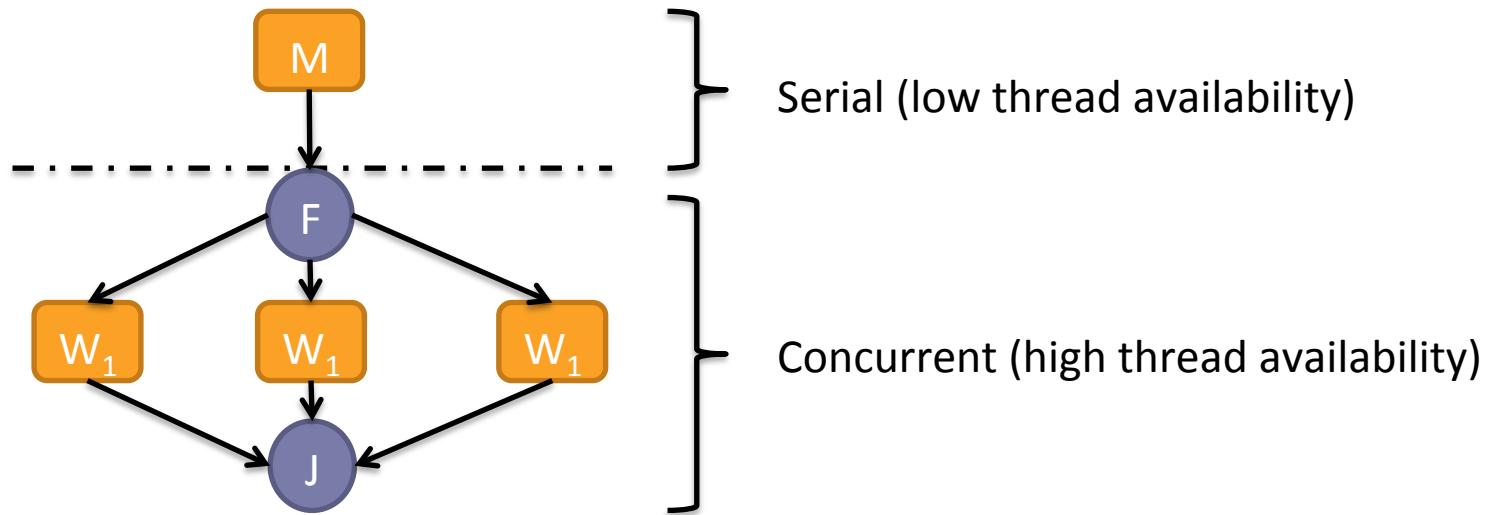
- Does Procrastination introduce deadlocks?
  - Threads can be procrastinated while holding a lock!



- Is Procrastination safe?
  - Yes. Forcing a local GC unblocks the threads.
  - No deadlocks or livelocks!

# Is Procrastination alone enough?

- Efficacy (Procrastination)  $\propto$  # Available runnable threads



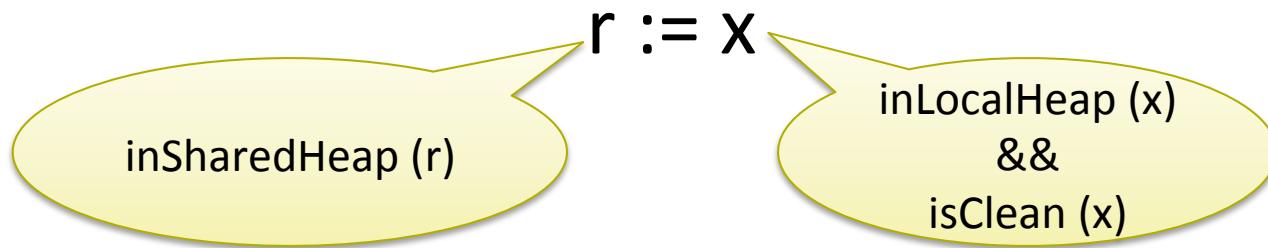
- With Procrastination, half of local major GCs were *forced*



Eager exporting writes while preserving visibility invariant

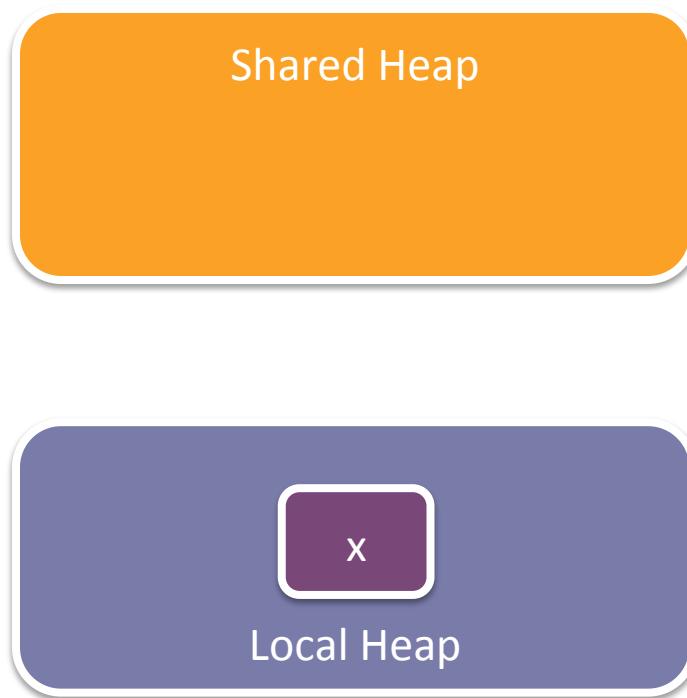
# Cleanliness

- A clean object closure can be lifted to the shared heap without breaking the visibility invariant



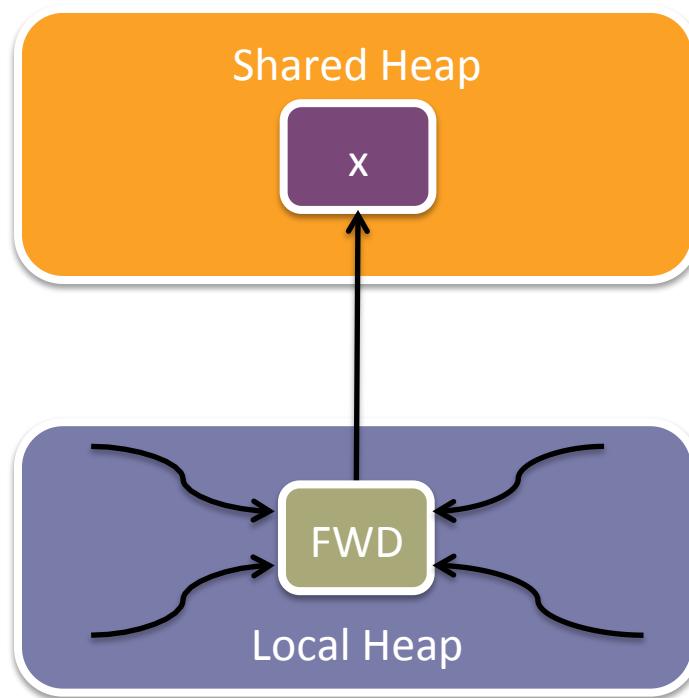
Eager write (no Procrastination)

# Cleanliness: Intuition



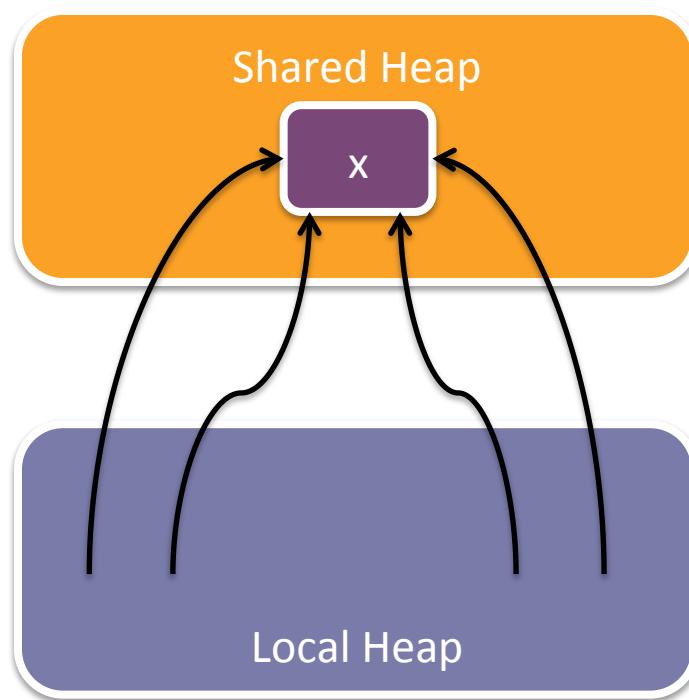
lift ( $x$ ) to shared  
heap

# Cleanliness: Intuition



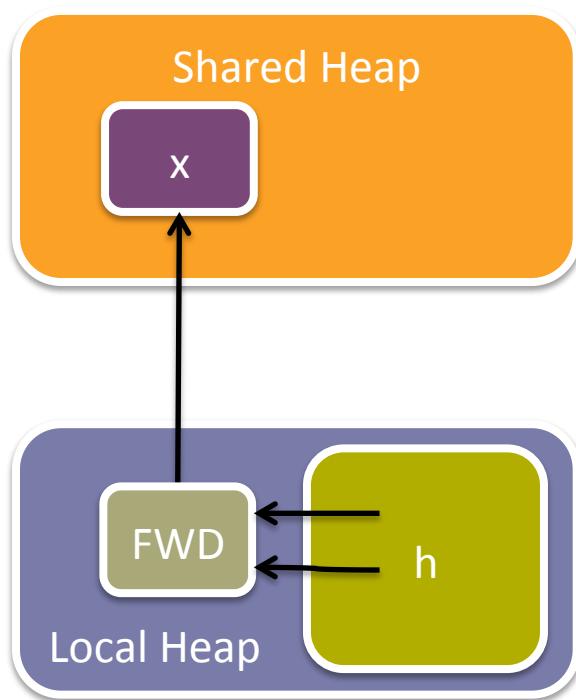
find all references  
to FWD

# Cleanliness: Intuition



Need to scan the entire local heap

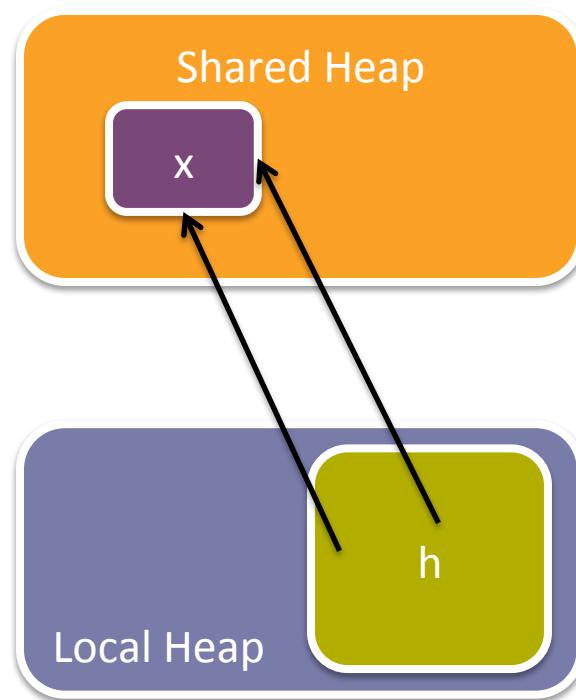
# Cleanliness: Simpler question



Do all references originate from  
heap region *h*?

`sizeof (h) << sizeof (local heap)`

# Cleanliness: Simpler question



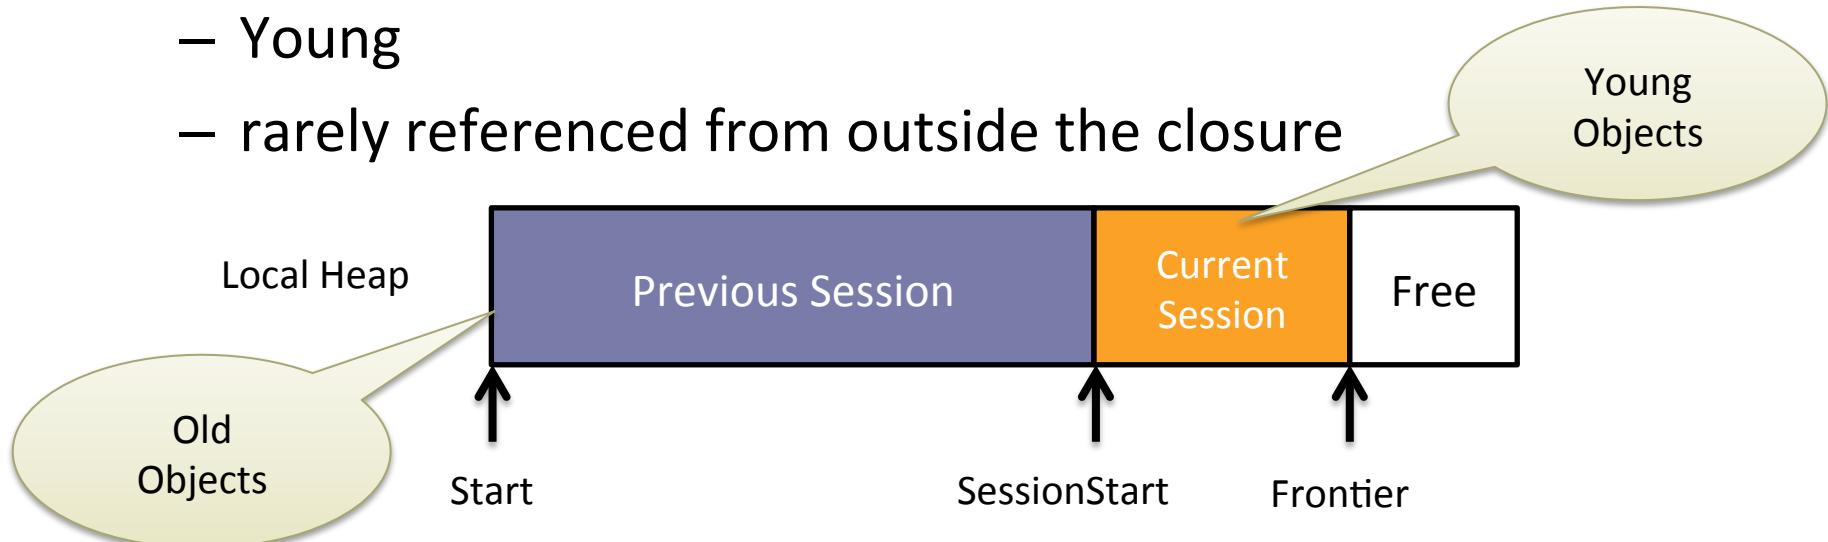
Only scan the  
heap region *h*.

Heap  
session!

`sizeof (h) << sizeof (local heap)`

# Heap Sessions

- Source of an exporting write is often
  - Young
  - rarely referenced from outside the closure



- Current session closed & new session opened
  - After an exporting write, a user-level context switch, a local GC

# Heap Sessions

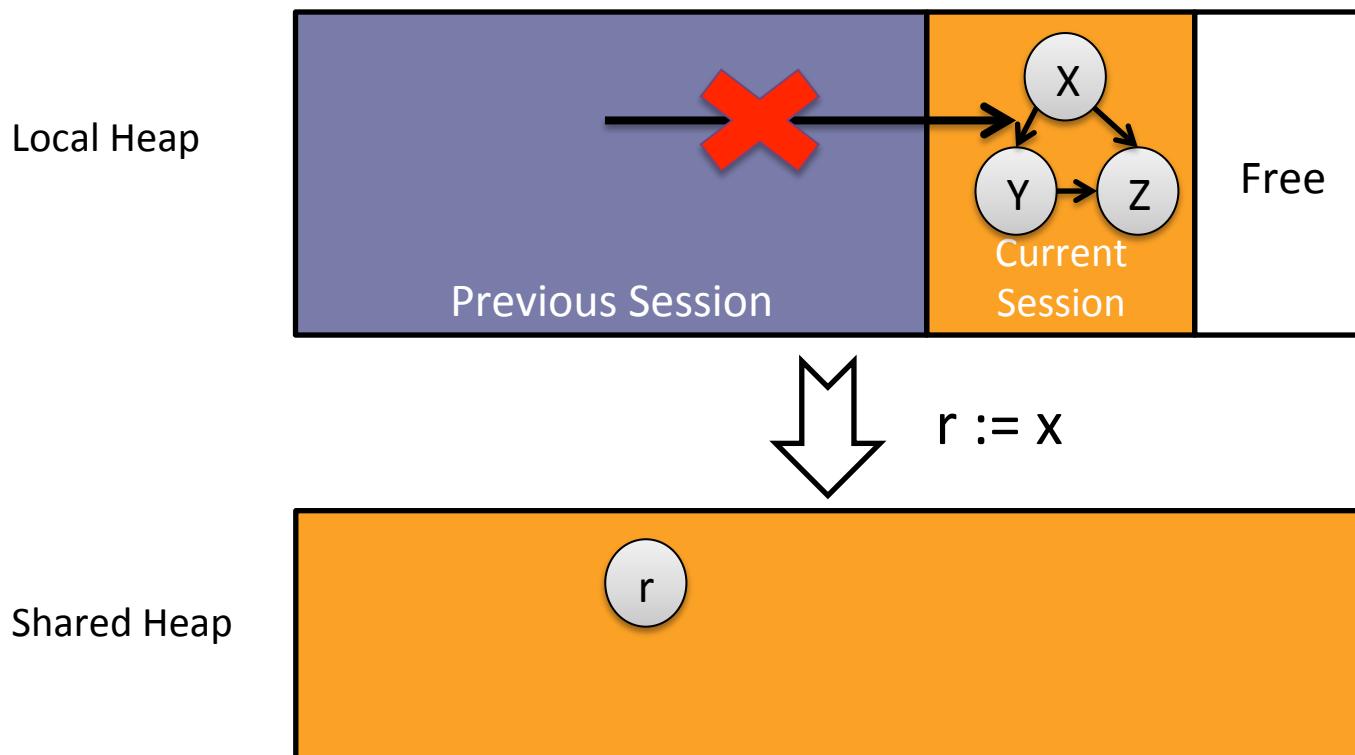
- Source of an exporting write is often
  - Young
  - rarely referenced from outside the closure



- Current session closed & new session opened
  - After an exporting write, a user-level context switch, a local GC
  - SessionStart is moved to Frontier
- Average current session size < 4KB

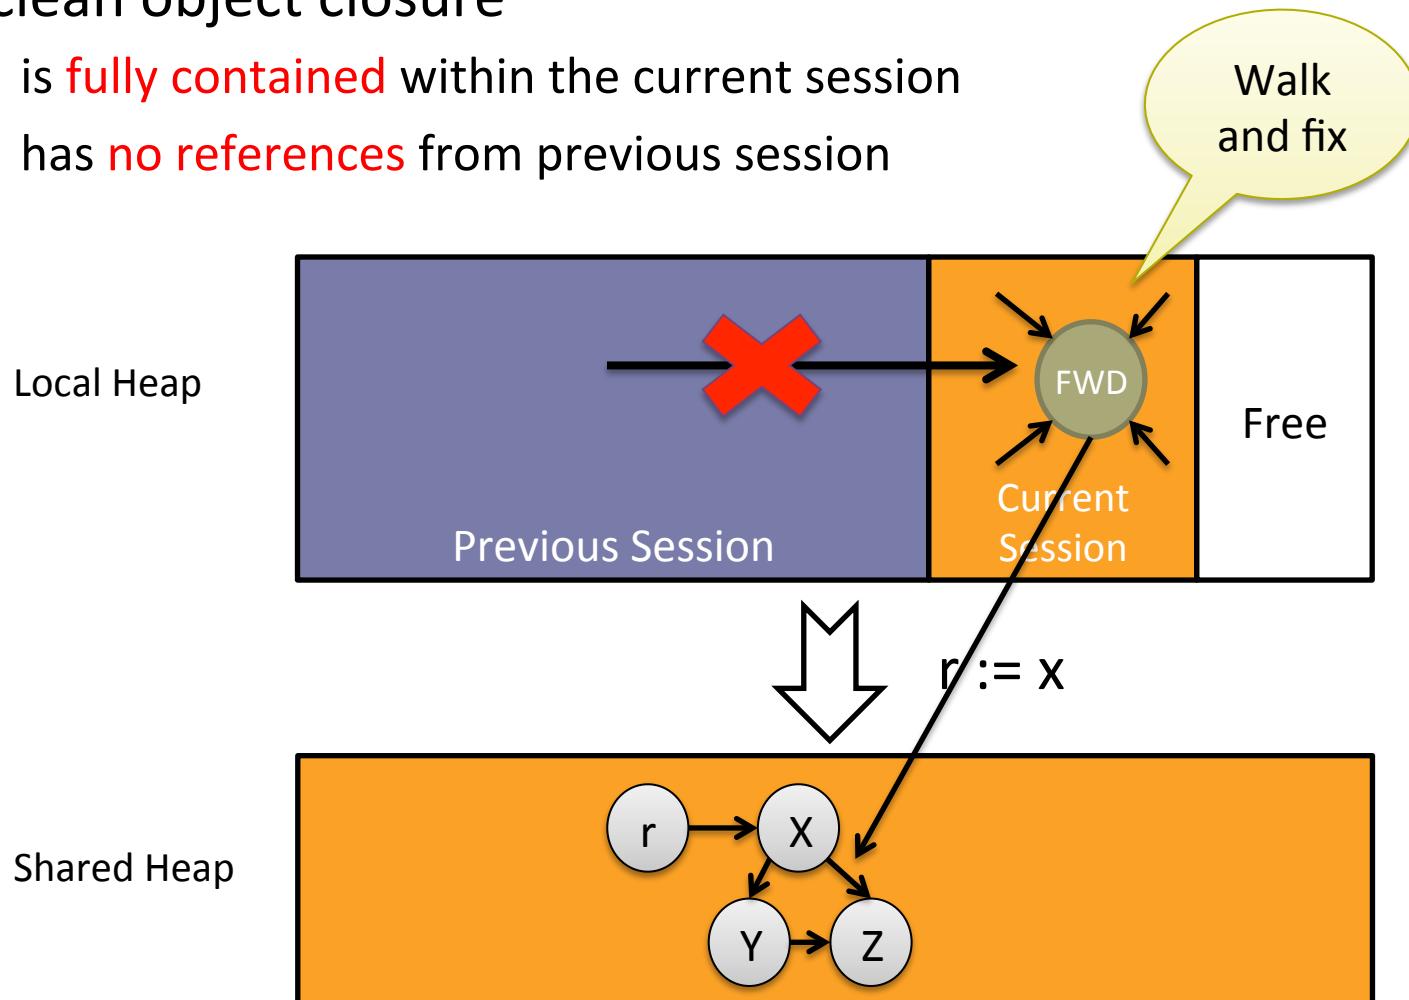
# Cleanliness: Eager exporting writes

- A clean object closure
  - is **fully contained** within the current session
  - has **no references** from previous session



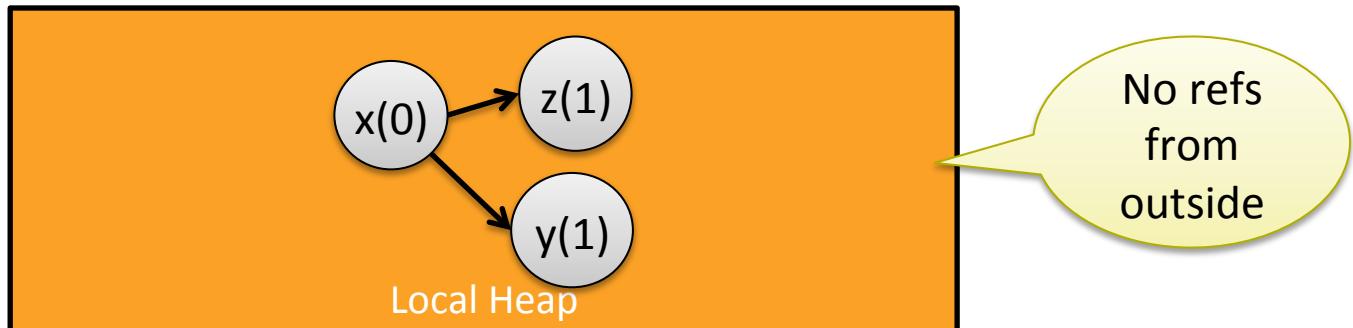
# Cleanliness: Eager exporting writes

- A clean object closure
  - is **fully contained** within the current session
  - has **no references** from previous session



# Avoid tracing current session?

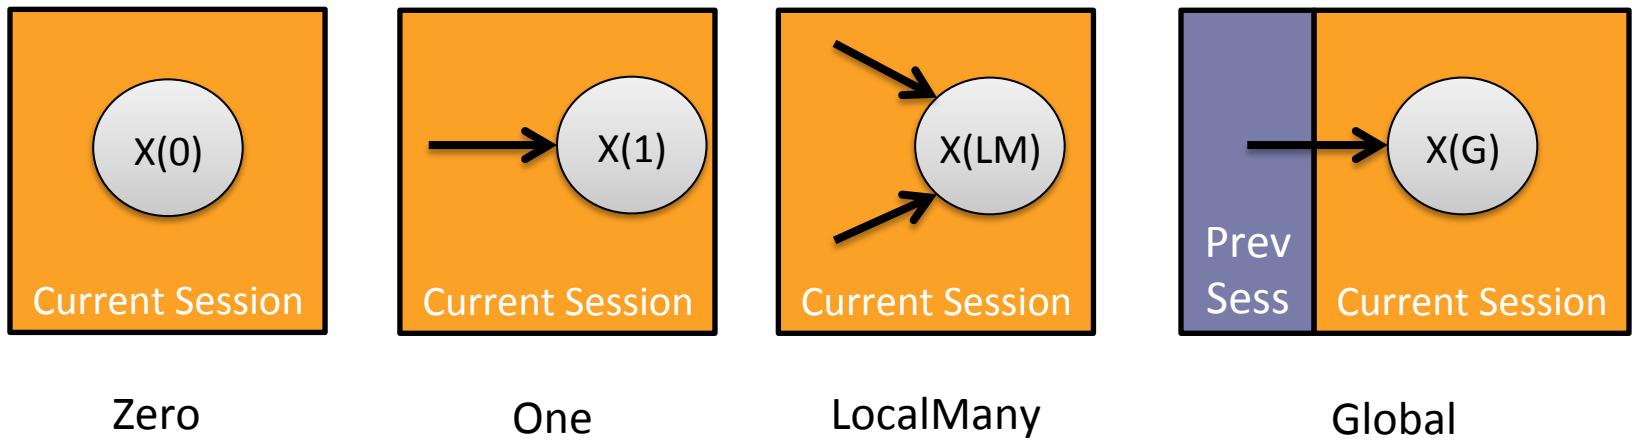
- Many SML objects are tree-structured (List, Tree, etc.,.)
  - Specialize for no pointers from outside the object closure
- $\forall x' \in \text{transitive object closure } (x),$   
 $\text{ref\_count}(x) = 0 \ \&\& \ \text{ref\_count}(x') = 1$



- `ref_count` does not consider pointers from stack or registers
- Eager exporting write
  - No current session tracing needed!

# Reference Count

- Purpose
  - Track pointers from previous session to current session
  - Identify tree-structured object

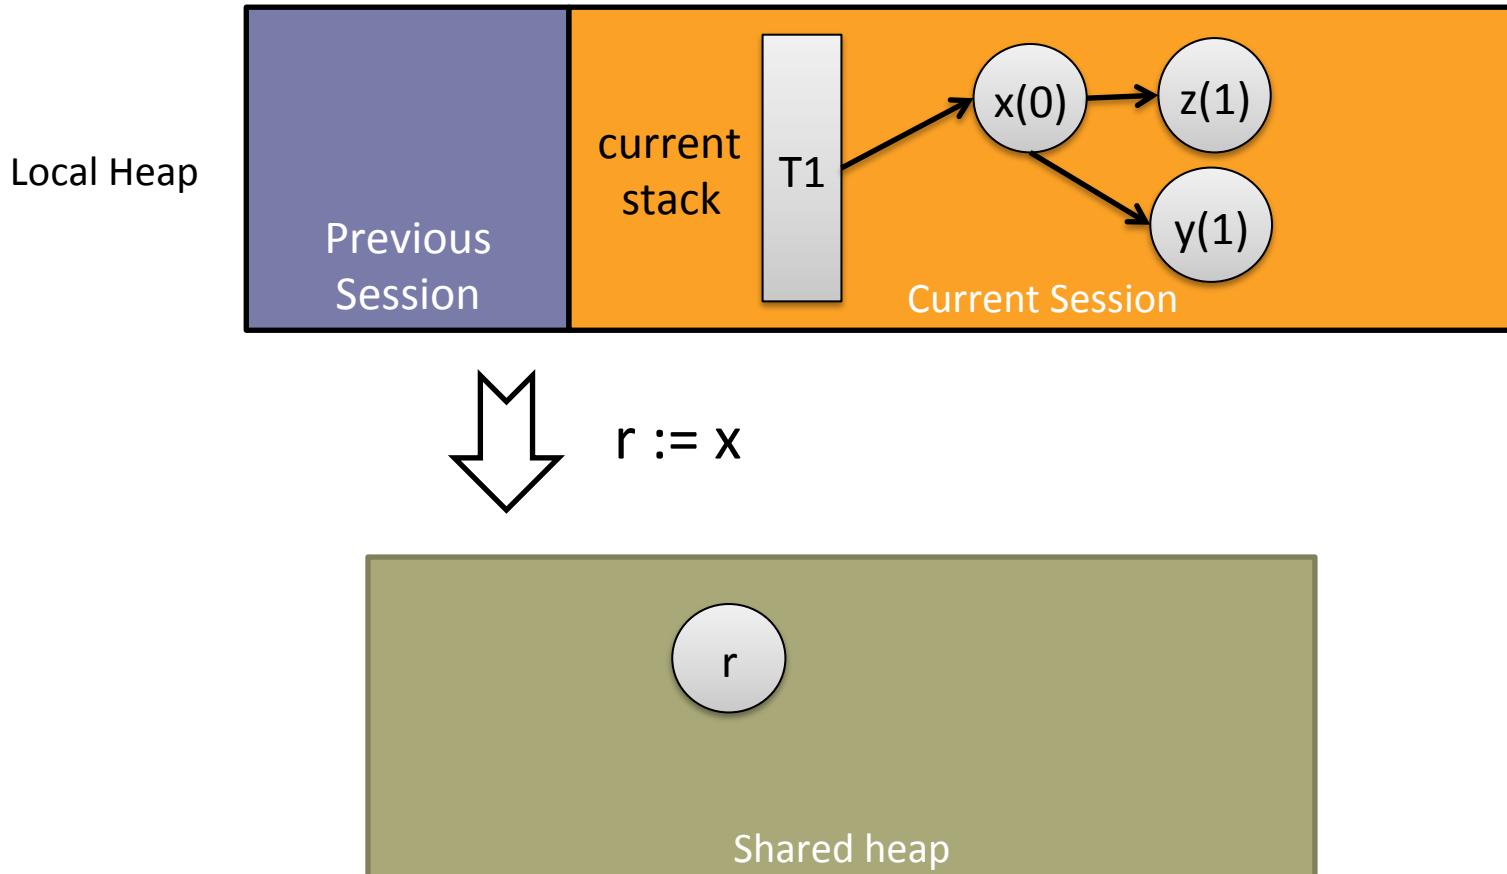


- Does not track pointers from stack and registers
  - Reference count only triggered during object initialization and mutation

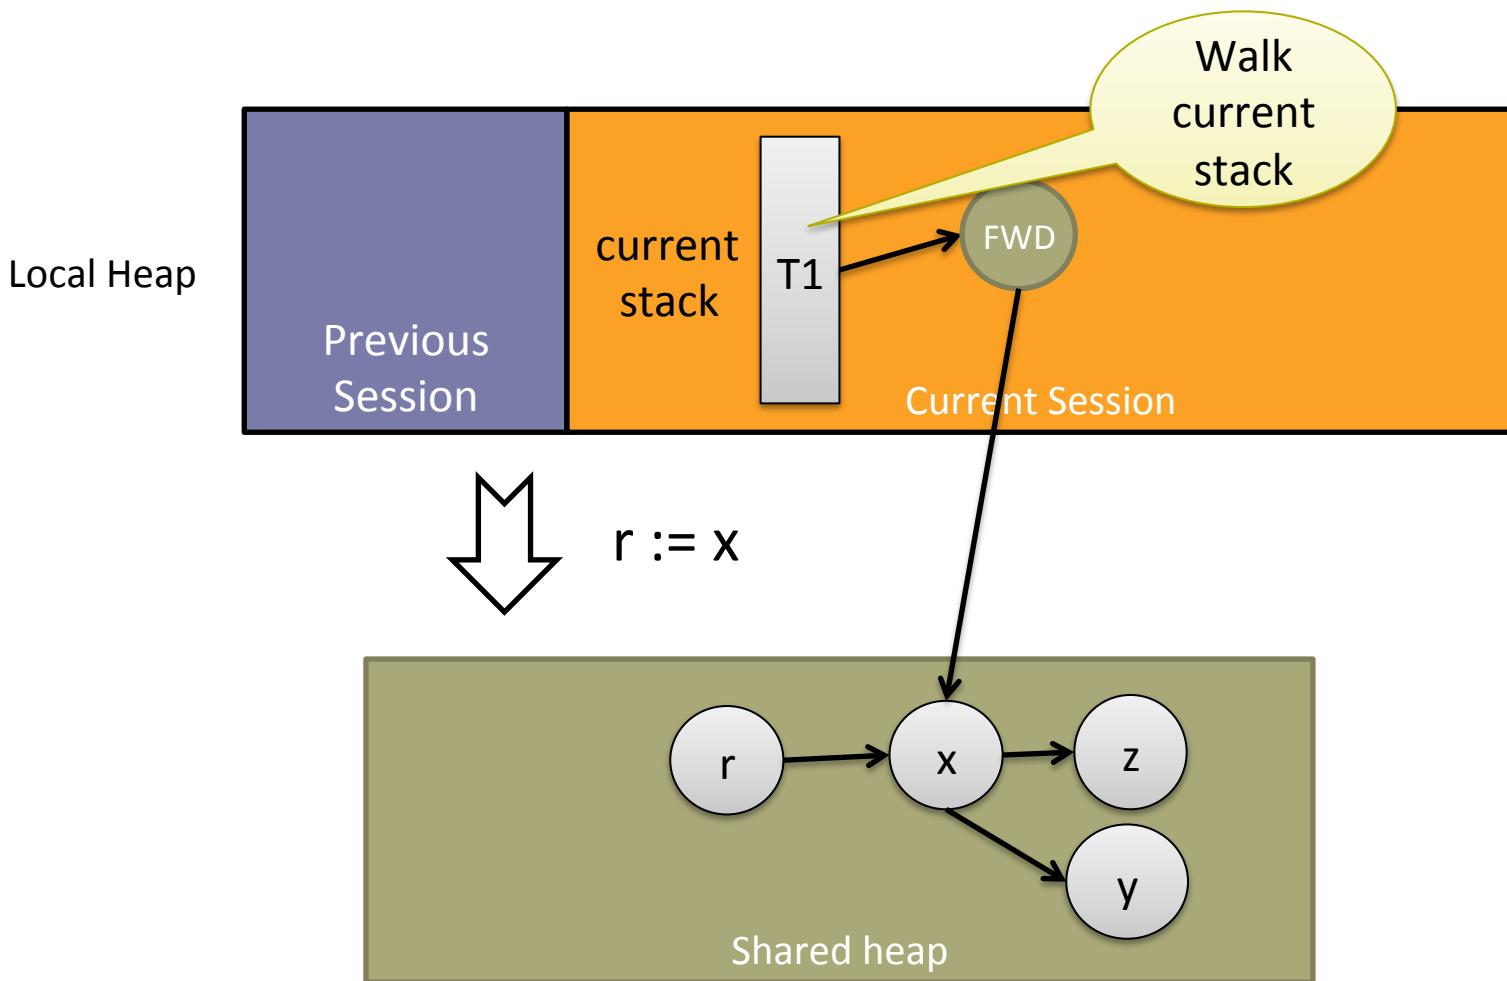
# Bringing it all together

- $\forall x' \in \text{transitive object closure } (x),$   
if **max (ref\_count (x'))**
  - One & ref\_count (x) = 0  $\Rightarrow$  tree-structured (Clean)  
 $\Rightarrow$  *Session tracing not needed*
  - LocalMany  $\Rightarrow$  Clean  $\Rightarrow$  *Trace current session*
  - Global  $\Rightarrow$  1+ pointer from previous session  $\Rightarrow$   
*Procrastinate*

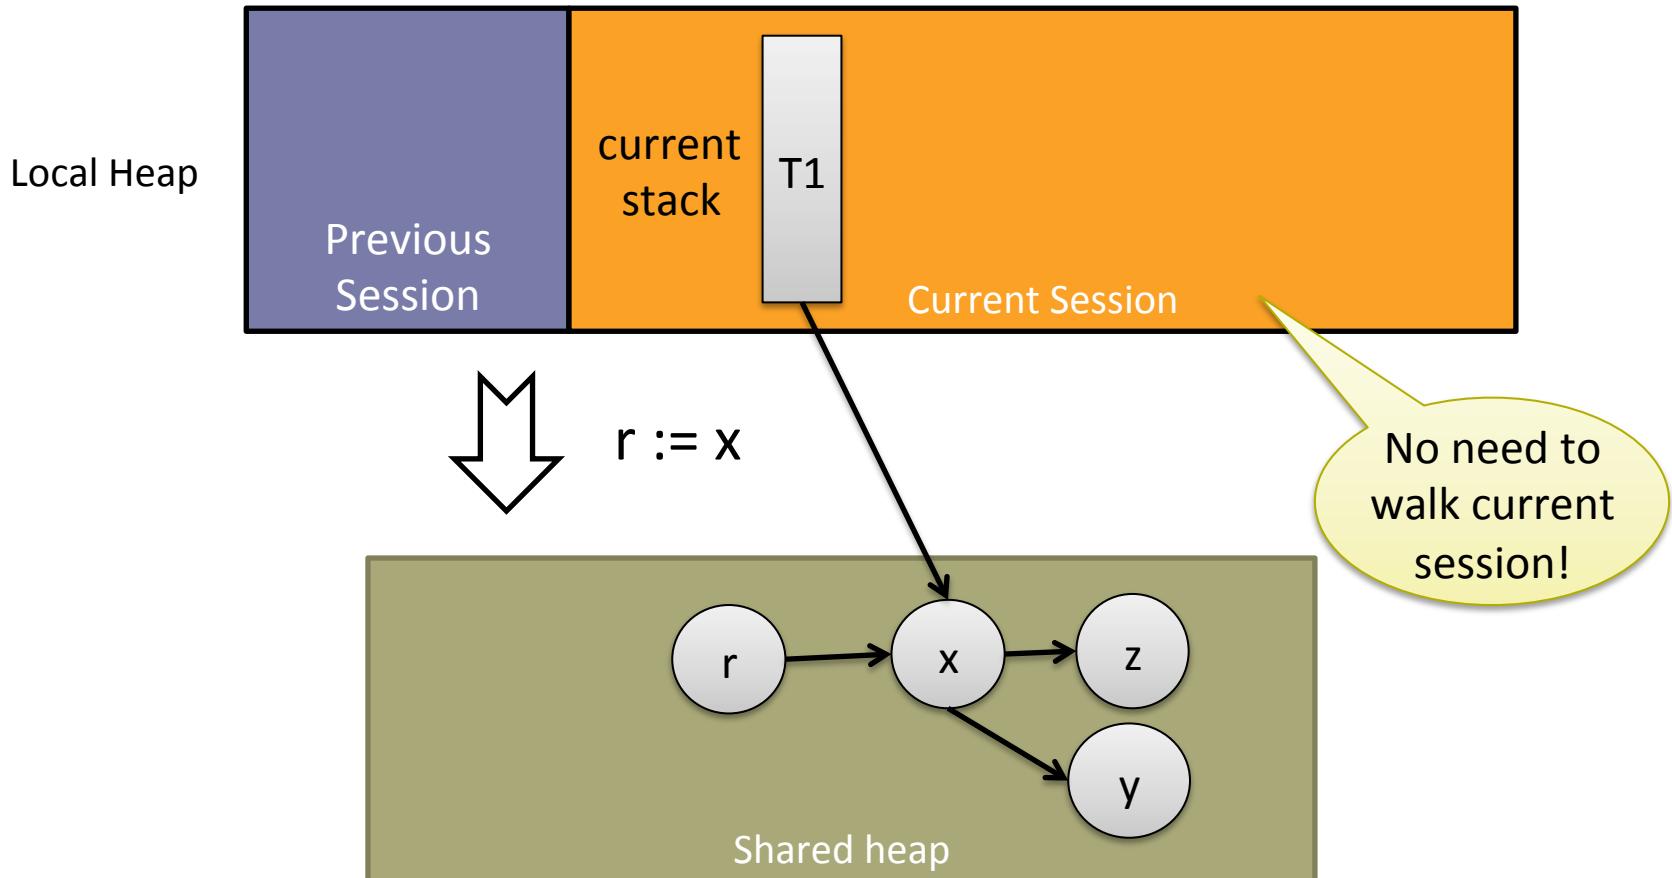
# Example 1: Tree-structured Object



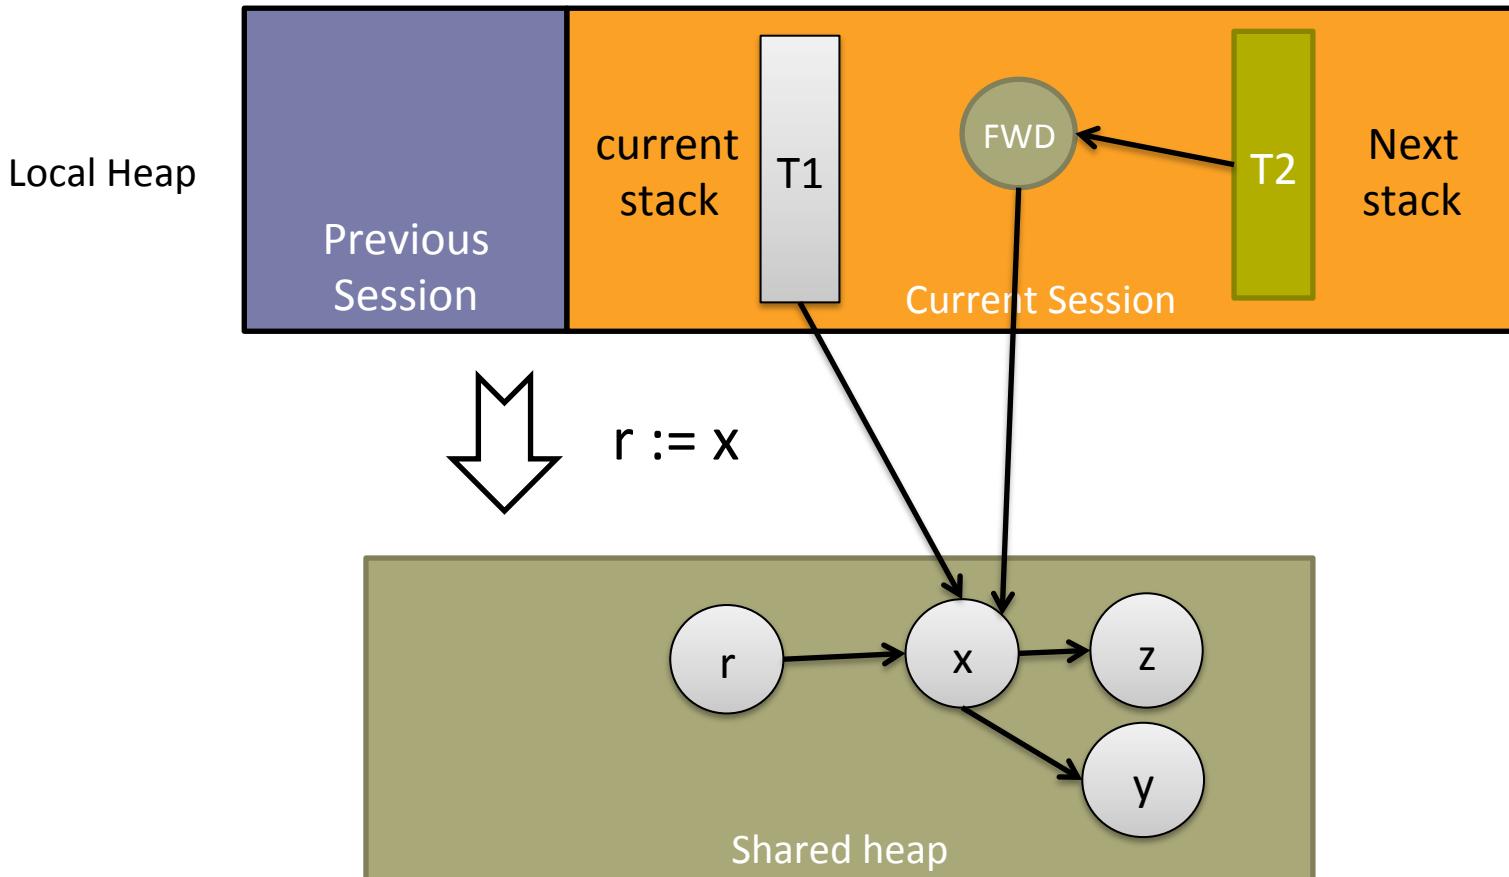
# Example 1: Tree-structured Object



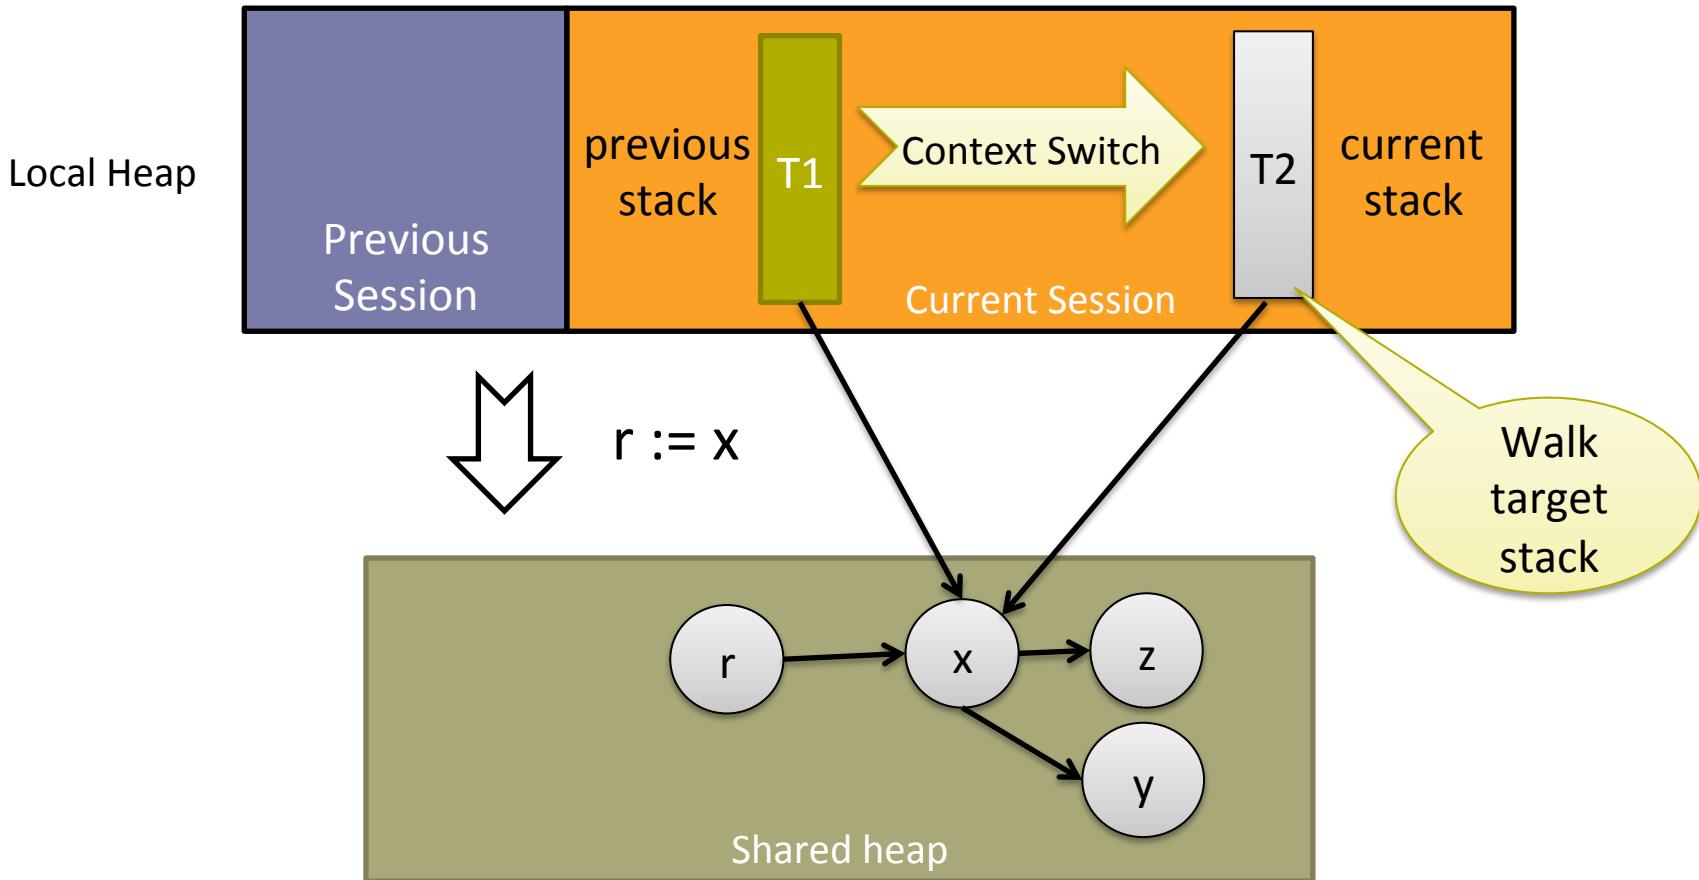
# Example 1: Tree-structured Object



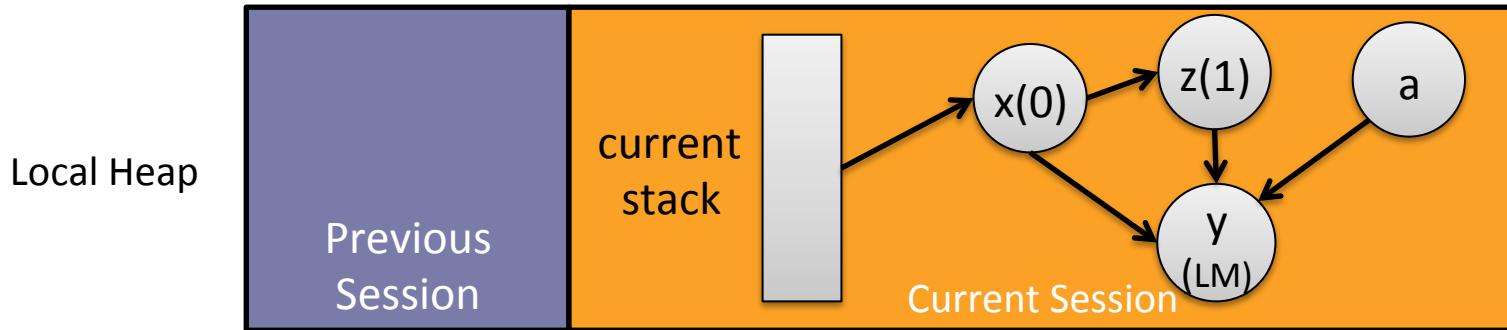
# Example 1: Tree-structured Object



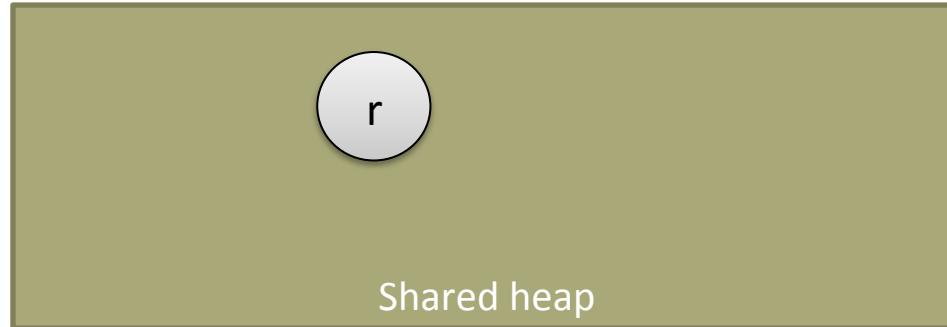
# Example 1: Tree-structured Object



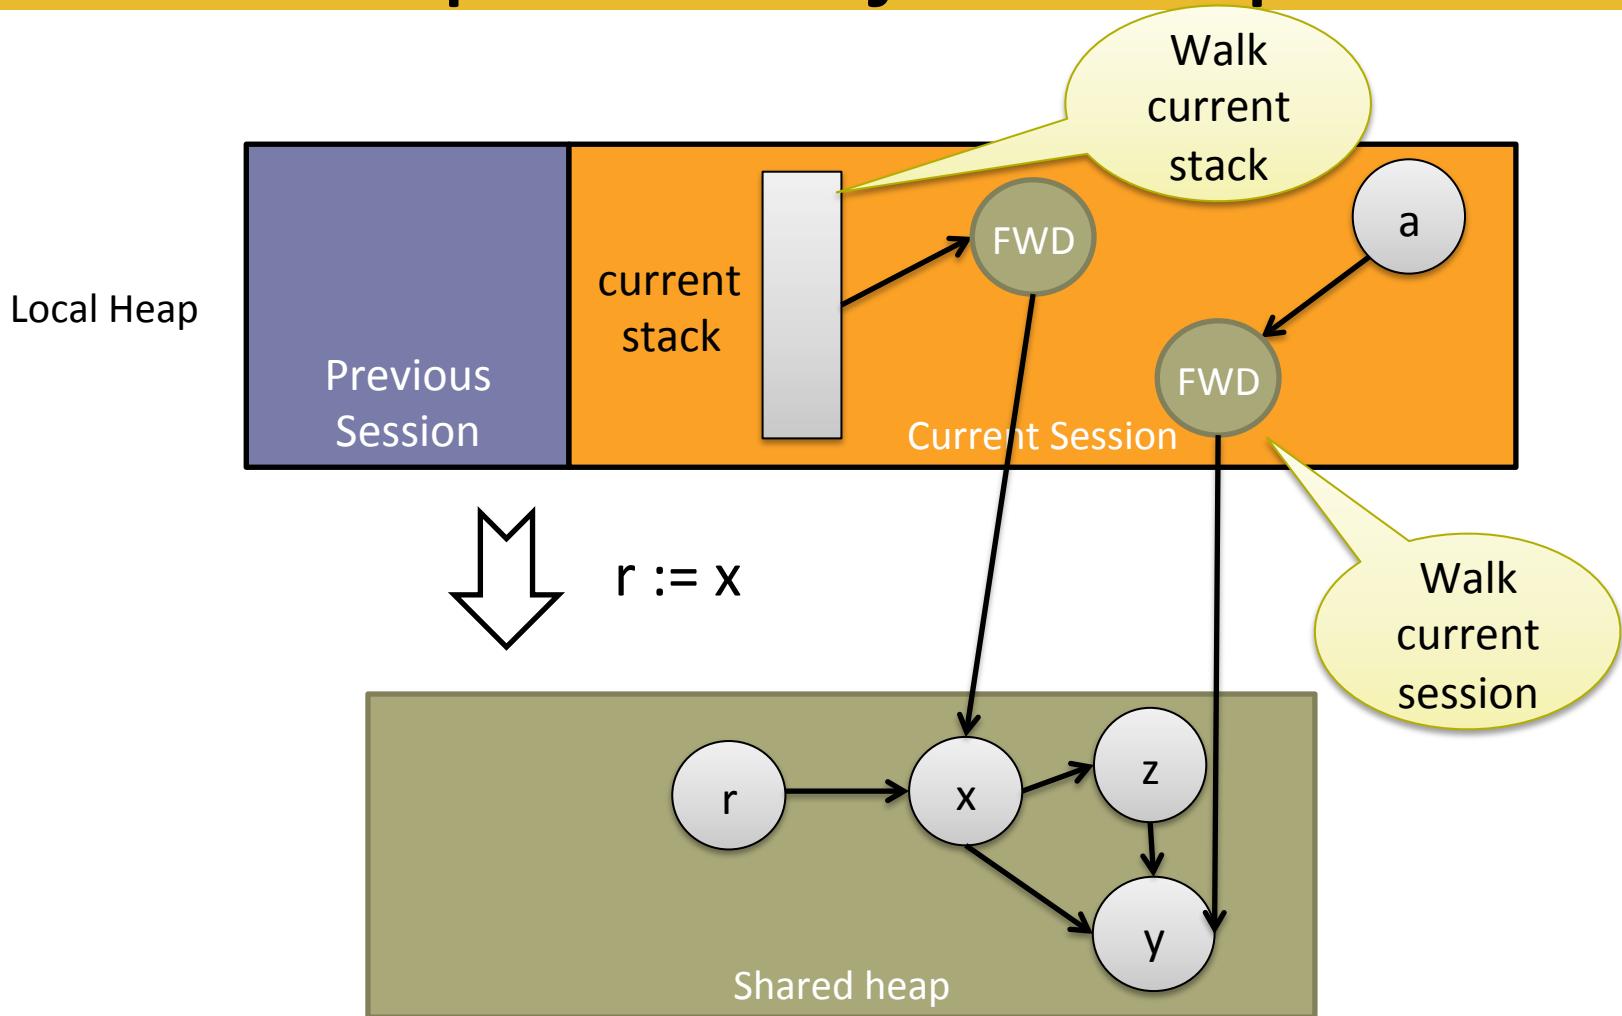
# Example 2: Object Graph



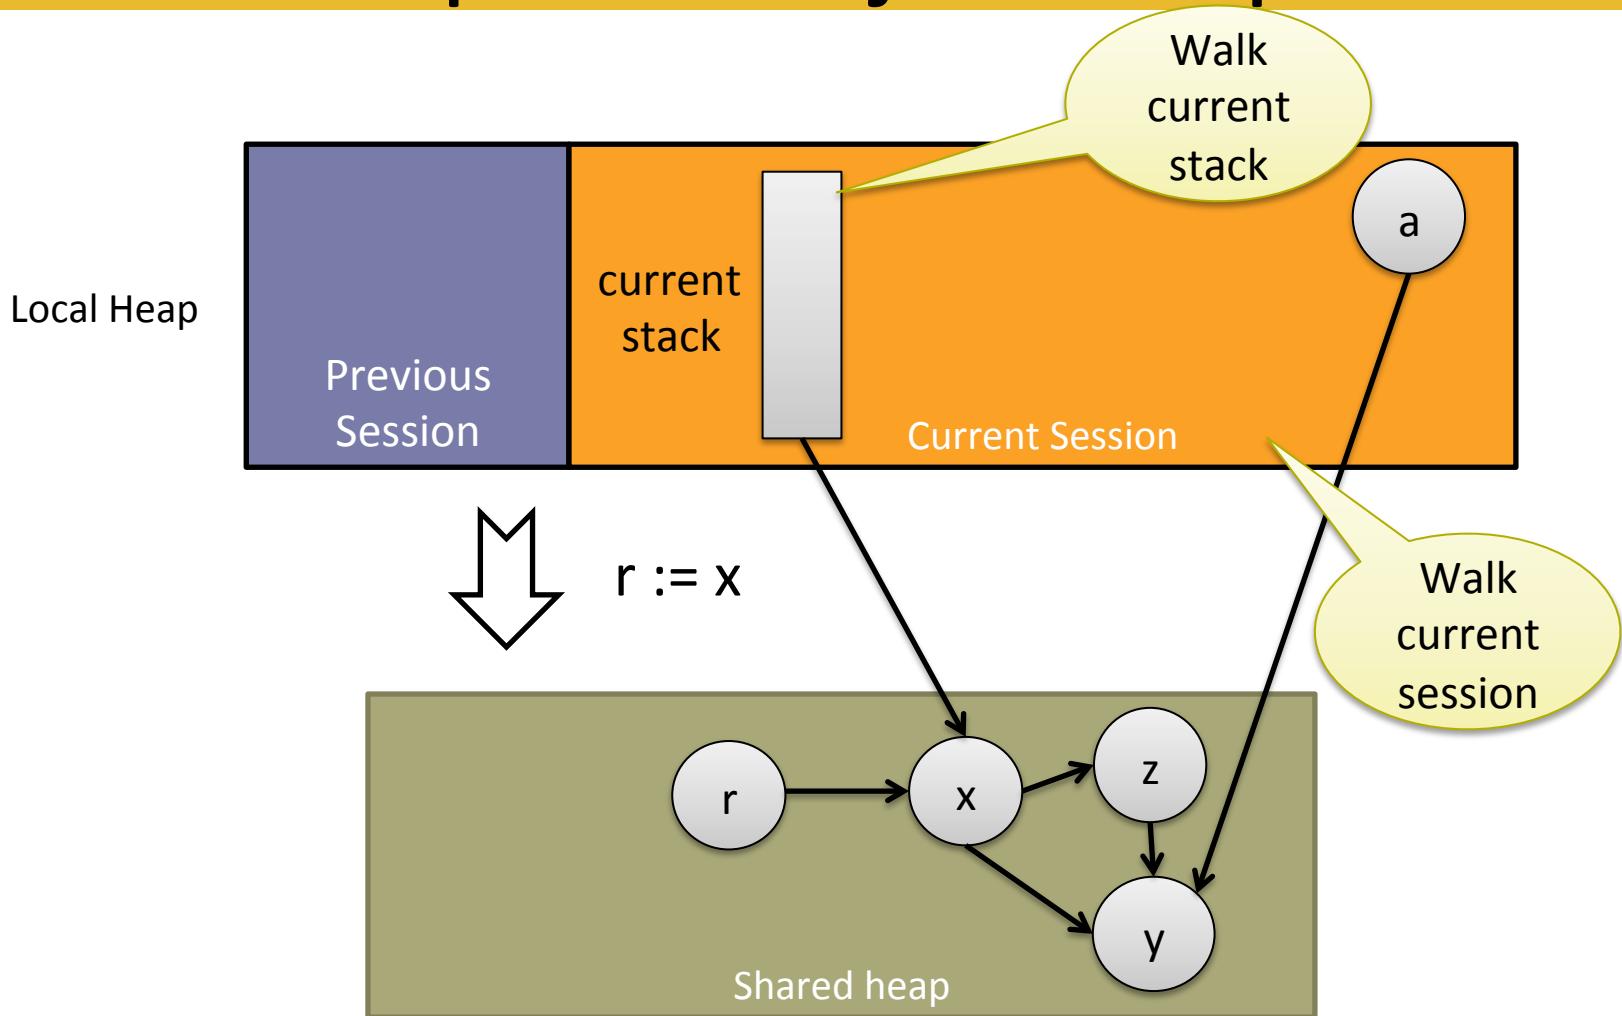
$r := x$



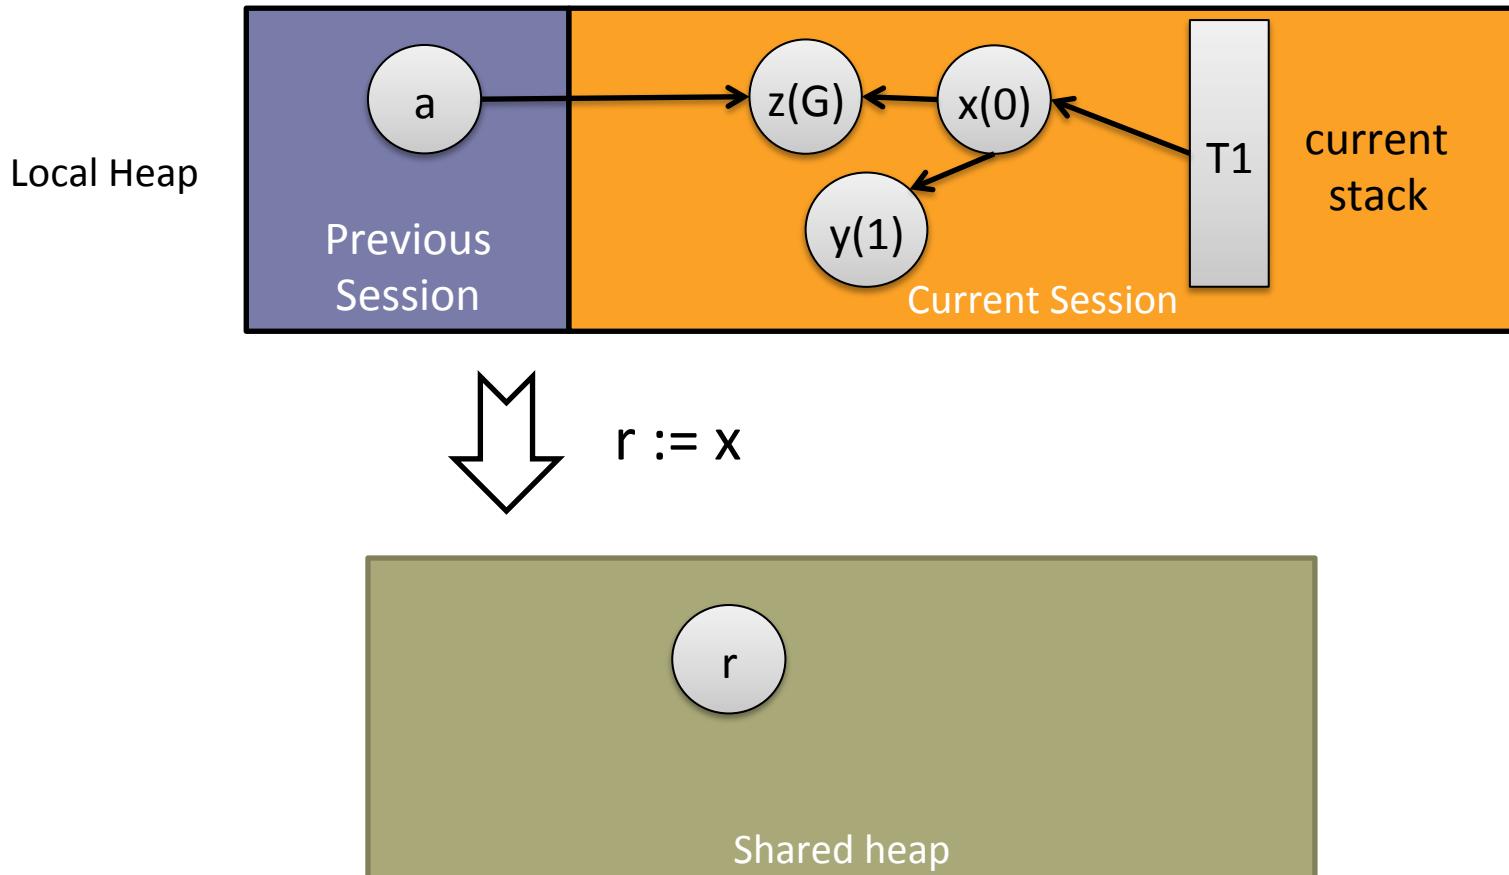
# Example 2: Object Graph



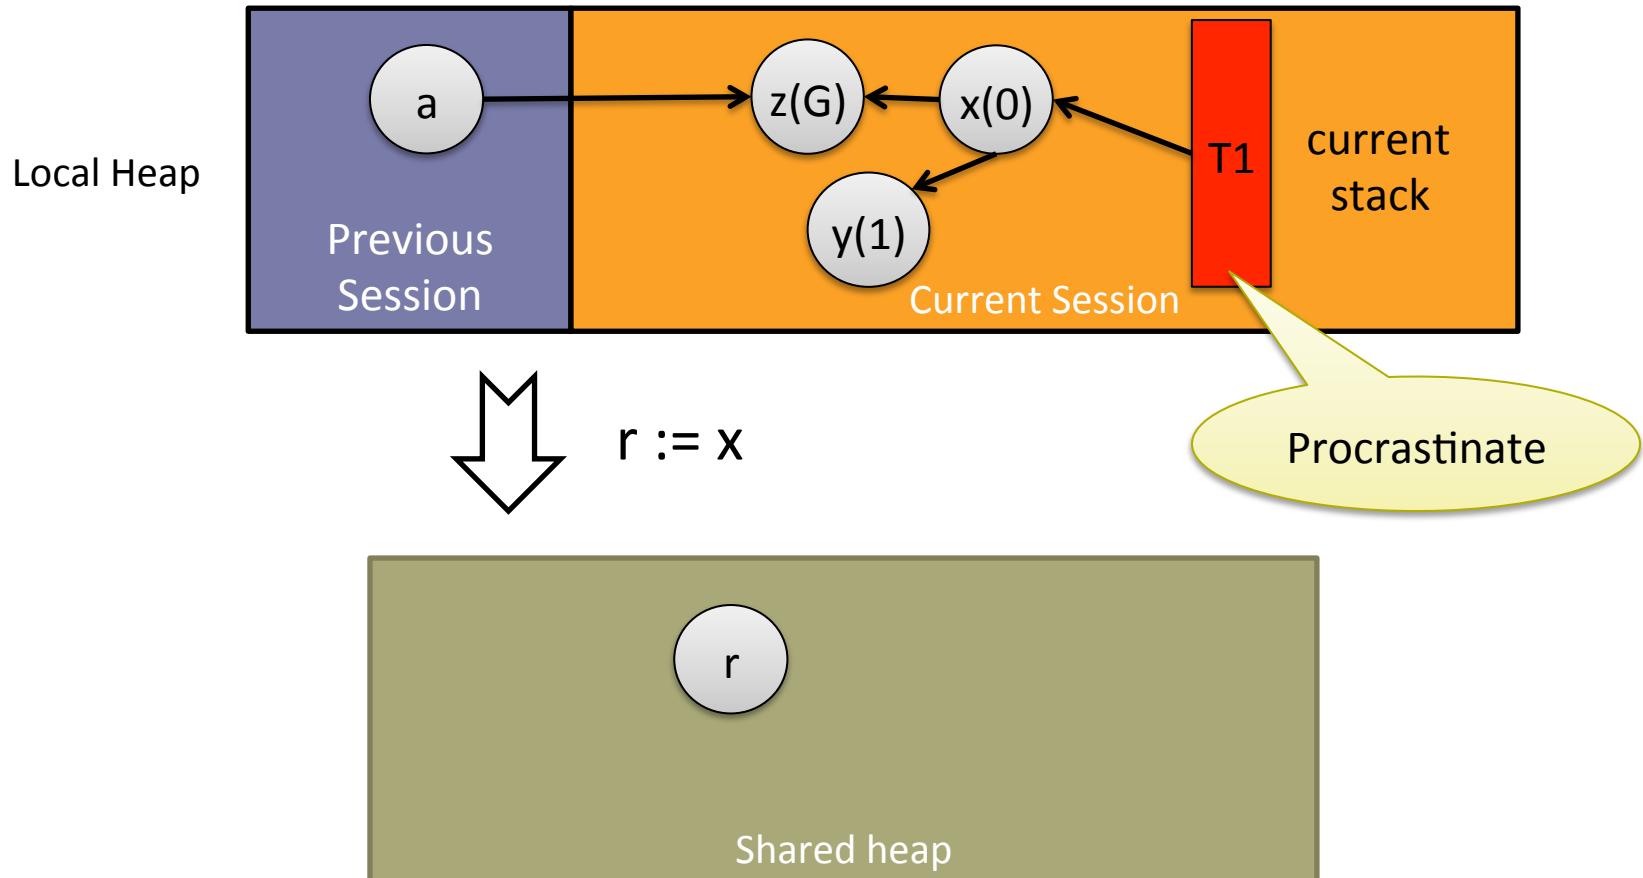
# Example 2: Object Graph



# Example 3: Global Reference



# Example 3: Global Reference



# Immutable Objects

- Specialize exporting writes
- If immutable object in previous session
  - Copy to shared heap
    - Immutable objects in SML do not have *identity*
  - Original object unmodified
- Avoid space leaks
  - Treat large immutable objects as mutable

# Cleanliness: Summary

- Cleanliness allows eager exporting writes while preserving visibility invariant
- With Procrastination + Cleanliness, <1% of local GCs were *forced*

# Evaluation

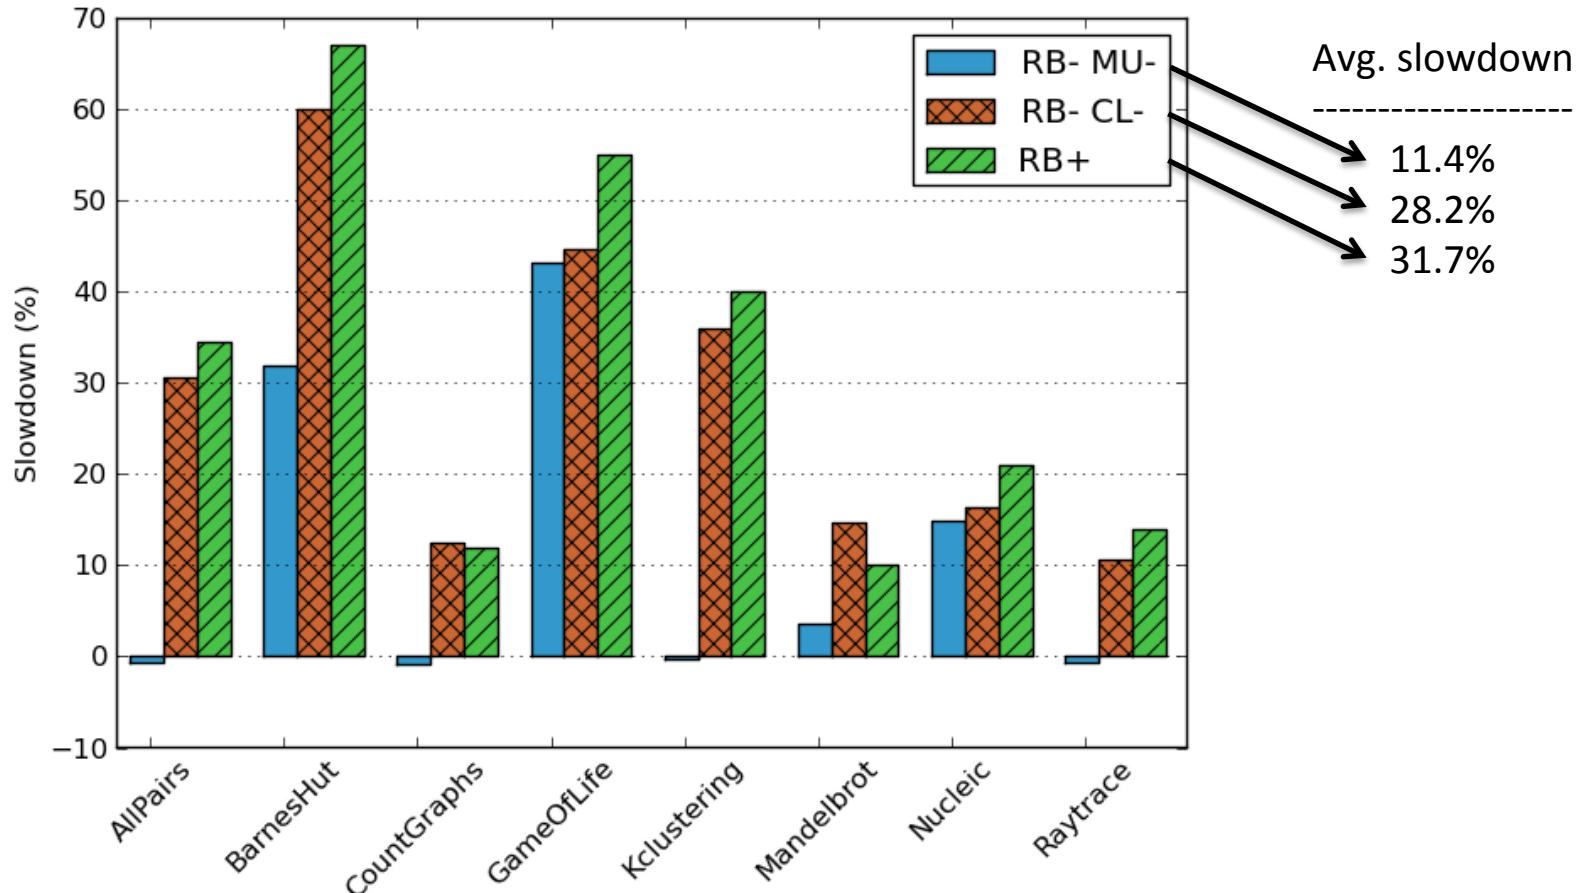
- Variants
  - **RB-** : TLC with Procrastination and Cleanliness
  - **RB+** : TLC with read barriers
- Sansom's dual-mode GC
  - Cheney's 2-space copying collection  $\leftrightarrow$  Jonker's sliding mark-compacting
  - Generational, 2 generations, No aging
- **Target Architectures:**
  - 16-core AMD Opteron server (NUMA)
  - 48-core Intel SCC (non-cache coherent)
  - 864-core Azul Vega3

# Results

- **Speedup:** At 3X min heap size, RB- faster than RB+
  - AMD (16-cores) **32%** (**2X** faster than STW collector)
  - SCC (48-cores) **20%**
  - AZUL (864-cores) **30%**
- **Concurrency**
  - During exporting write, **8** runnable user-level threads/core!

# Cleanliness Impact

- **RB- MU-** : RB- GC ignoring mutability for Cleanliness
- **RB- CL-** : RB- GC ignoring Cleanliness (*Only Procrastination*)



# Conclusion

- Eliminate the need for read barriers by preserving the **visibility invariant**
  - **Procrastination:** **Exploit concurrency** for delaying exporting writes
  - **Cleanliness:** **Exploit generational property** for eagerly perform exporting writes
- Additional niceties
  - Completely dynamic → Portable
  - Does not impose any restriction on the GC strategy

# Questions?



<http://multimlton.cs.purdue.edu>

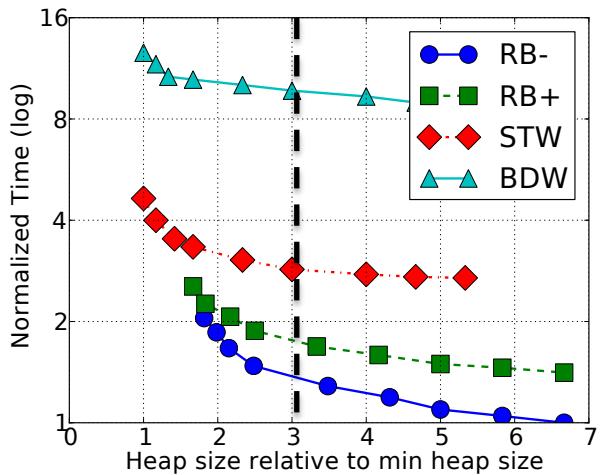
# Results

- On AMD, 16 Cores, 3X minimum heap size
- **Mutator time:**
  - STW GC spends the least amount of time in the mutator
    - No read/write barriers
  - Compared to STW GC, the mutator time of
    - RB- 18% more, RB+ 39% more
- **GC time:**
  - RB- spends the least amount time doing GC
  - RB- within 5% of RB+

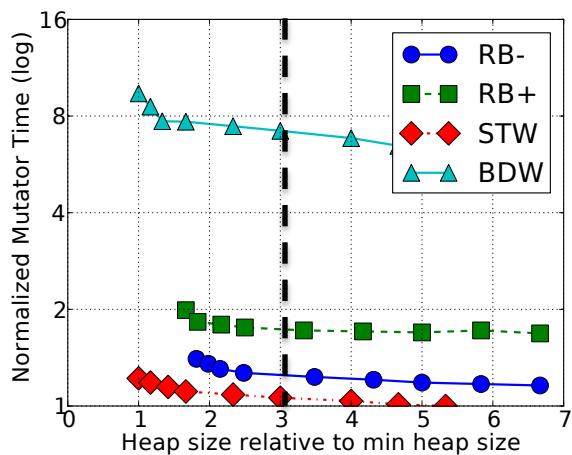
# Performance on AMD (16-cores)

At 3X min  
heap size:

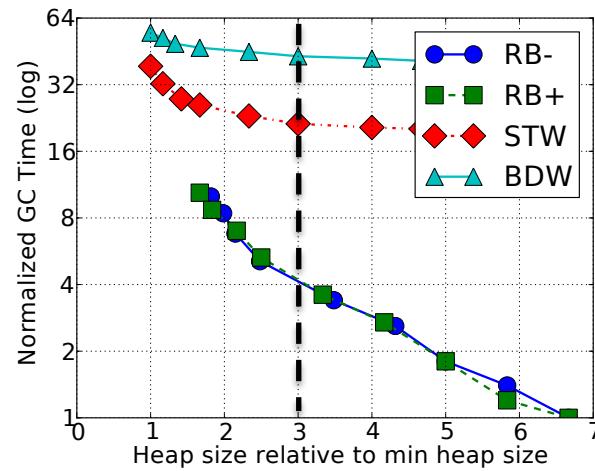
RB+ 32%  
STW 106%  
BDW 584%



(a) Total time

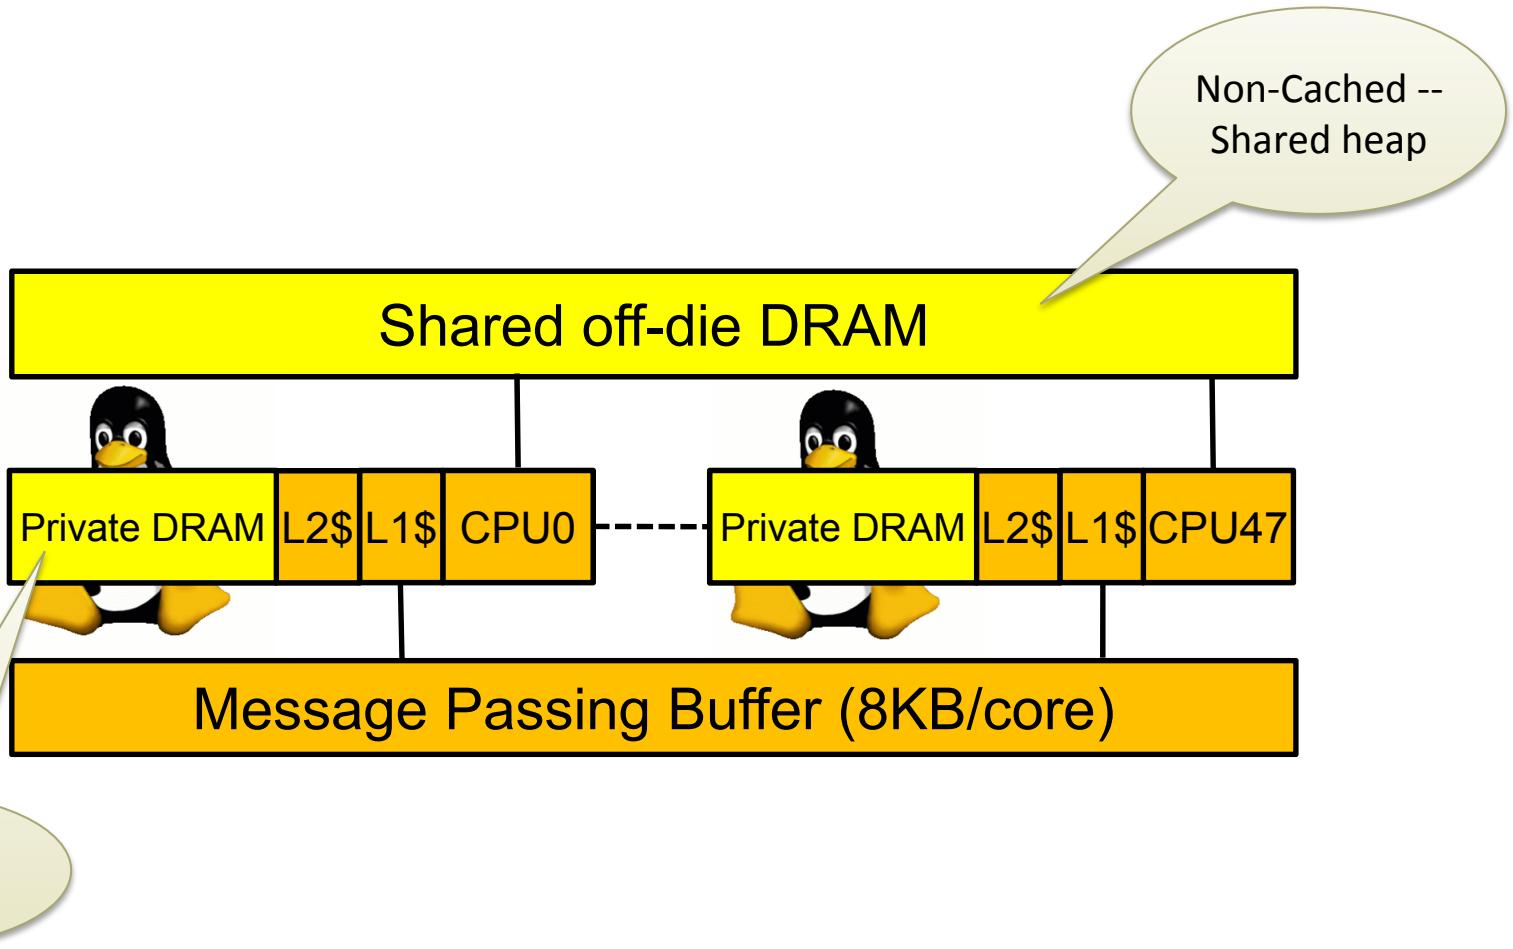


(b) Mutator time

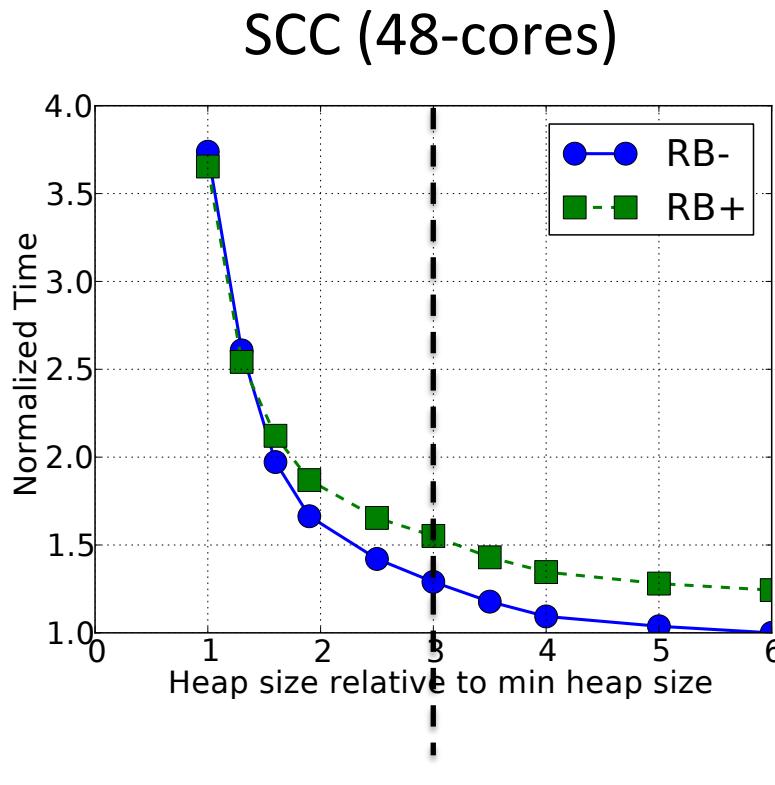


(c) GC time

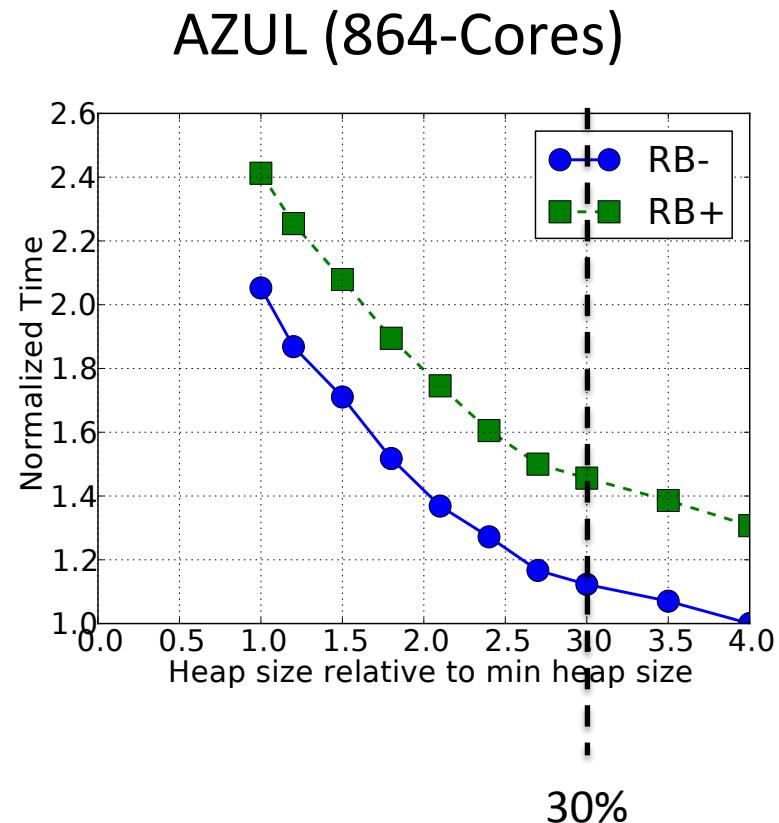
# MultiMLton - SCC implementation



# Total time: SCC and AZUL



*Non-cache coherent*



*Scalable, cache-coherent*

# Cleanliness Impact (1)

Benchmark	AllPairs	BarnesHut	CountGraphs	GameOfLife	Kclustering	Mandelbrot	Nucleic	Raytrace
RB-	1831	46532	154	38621	25812	132	156	3523
RB- MU-	1831	4092312	192	735543	50323	209	433092	3743
RB- CL-	124232	67156821	50178	5867423	27023911	25491	912349	61198

Number of Preemptions on exporting writes

Benchmark	AllPairs	BarnesHut	CountGraphs	GameOfLife	Kclustering	Mandelbrot	Nucleic	Raytrace
RB-	0.08	0.17	0	3.54	0	1.43	0	1.72
RB- MU-	0.08	19.2	0.03	9.47	0.02	2.86	9.37	1.72
RB- CL-	38.55	100	0.18	99.75	21.64	86.22	19.3	24.86

Forced GCs as a % of total number of local major GCs

# Benchmark Characteristics

Benchmark	Allocation Rate (MB/s)			Bytes Allocated (GB)				# Threads		
	AMD	SCC	AZUL	AMD	SCC	AZUL	% Sh	AMD	SCC	AZUL
AllPairs	817	53	1505	16	16	54	11	256	512	32768
Barneshut	772	70	1382	20	20	876	2	512	1024	32768
Countgraphs	2594	144	4475	24	24	1176	1	128	256	16384
GameOfLife	2445	127	4266	21	21	953	13	256	1024	8192
Kclustering	3643	108	8927	32	32	1265	3	256	1024	8192
Mandelbrot	349	43	669	2	2	32	8	128	512	8192
Nucleic	1430	87	4761	13	14	609	1	64	384	16384
Raytrace	809	54	2133	11	12	663	4	128	256	2048

# Session Impact

Benchmark	AllPairs	Barneshut	Countgraphs	GameOfLife	Kclustering	Mandelbrot	Nucleic	Raytrace
% LM clean	5.3	13.4	8.6	23.2	17.6	4.5	13.3	8.2
Avg. session size (bytes)	2908	1580	3612	1344	2318	8723	1264	1123

Figure 17: Impact of heap session: % LM clean represents the fraction of instances when a clean object closure has at least one object with LOCAL\_MANY references.

# Read Barrier

## Conditional (Baker Style)

From

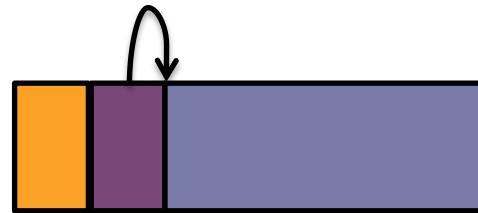


To

-----

## Unconditional (Brooks style)

From



To

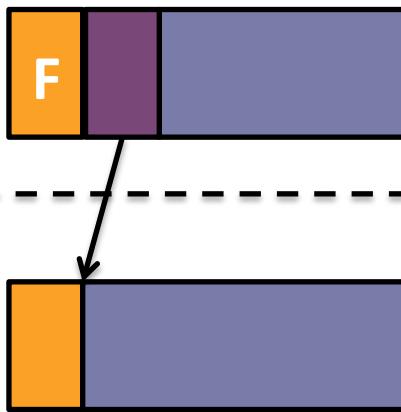
-----

# Read Barrier



## Conditional (Baker Style)

From



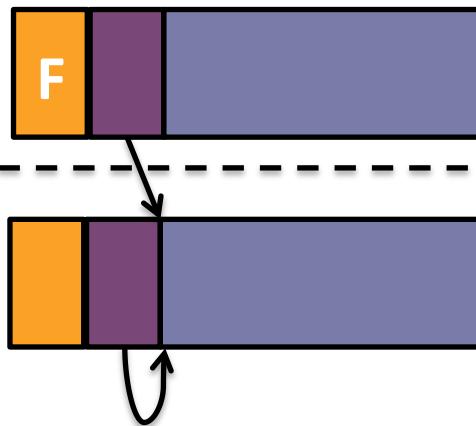
```
pointer readBarrier (pointer *p) {  
    if ((*Header*) (p - HD_OFF) == F)  
        return * (pointer*) p;  
    return p;  
}
```



Has Conditional Check

## Unconditional (Brooks style)

From



```
pointer readBarrier (pointer *p) {  
    return * (pointer*) (p - IND_OFF);  
}
```



Needs extra header word

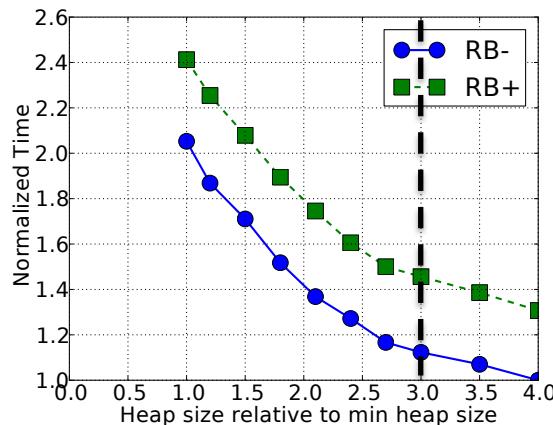
# Read Barrier Optimizations

- Stacks and registers never point to forwarding pointers
- “Eager” read barriers (D.Bacon et al. POPL’93)
- Scan stack after exporting write
- Exporting write is a GC safe-point
- Reduces RB overhead by ~5%

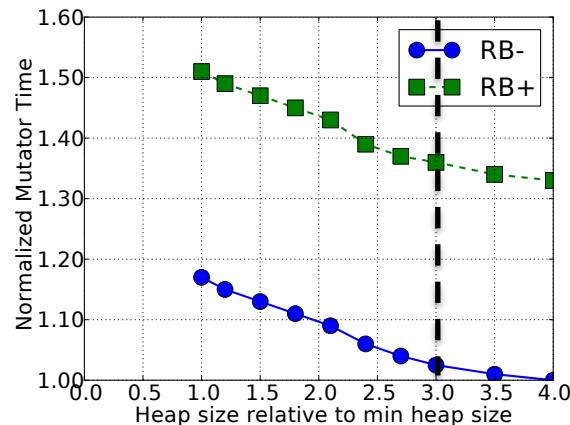
# Performance on AZUL

At 3X min  
heap size:

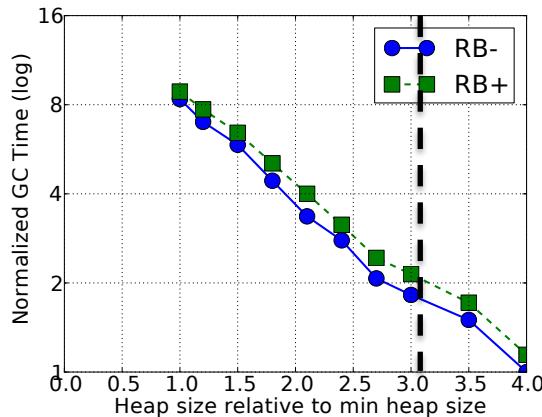
RB+ 30%



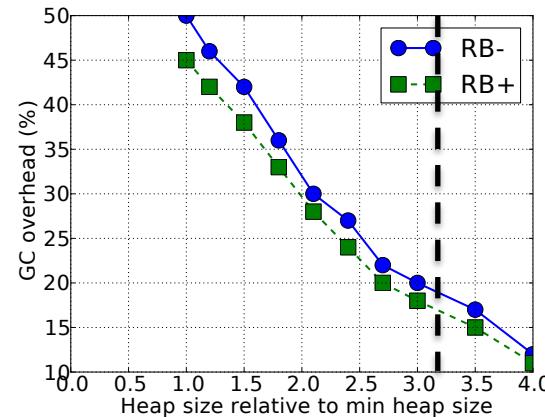
(a) Total time



(b) Mutator time



(c) Garbage collection time

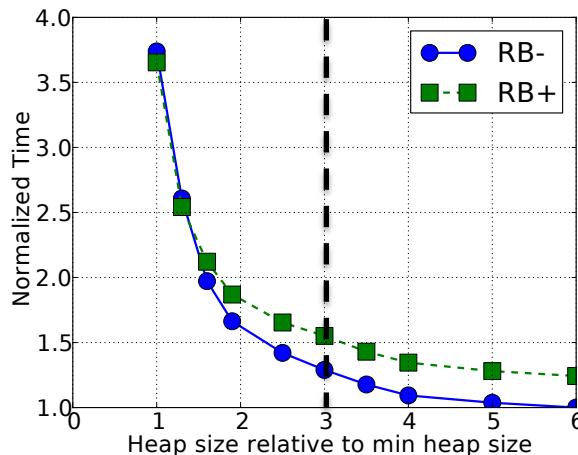


(d) Garbage collection overhead

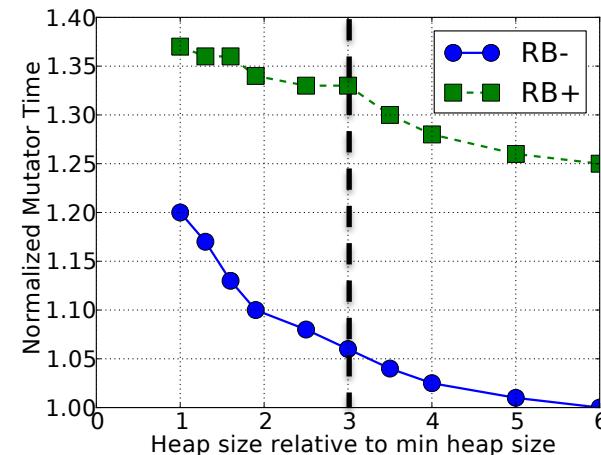
# Performance on SCC

At 3X min  
heap size:

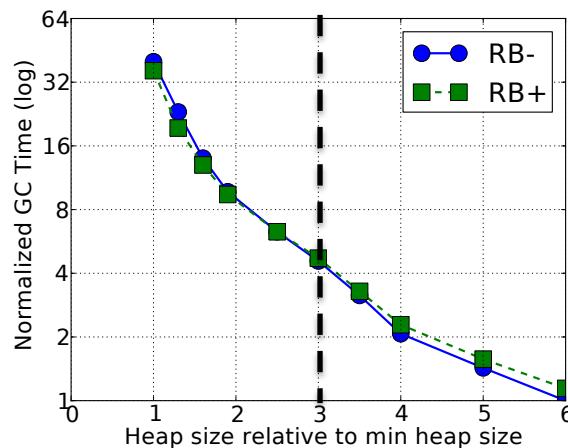
RB+ 20%



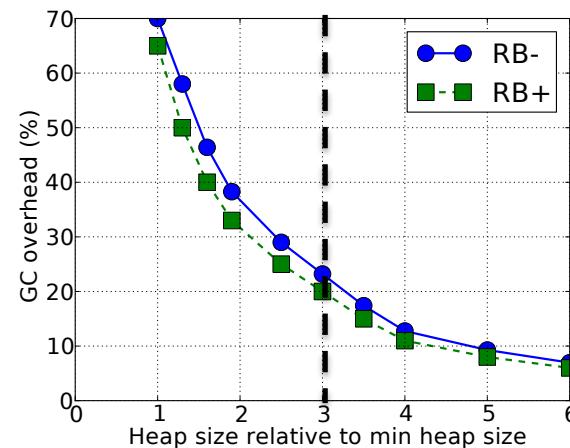
(a) Total time



(b) Mutator time

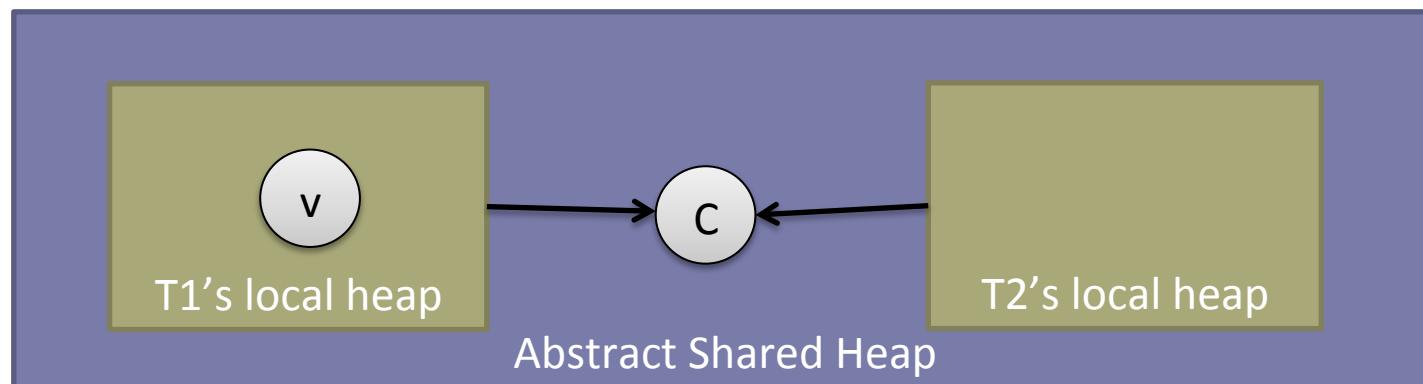
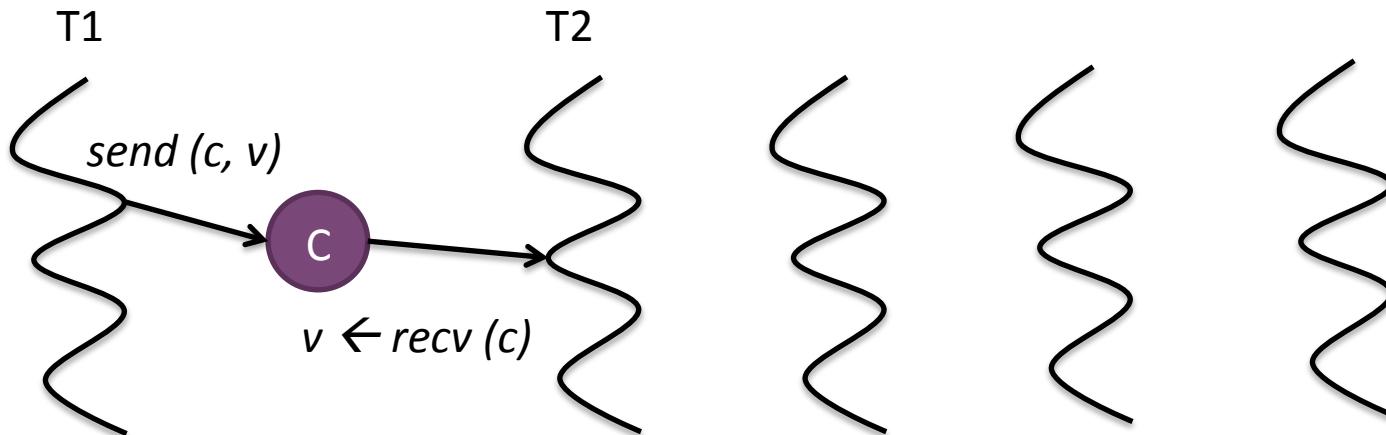


(c) Garbage collection time



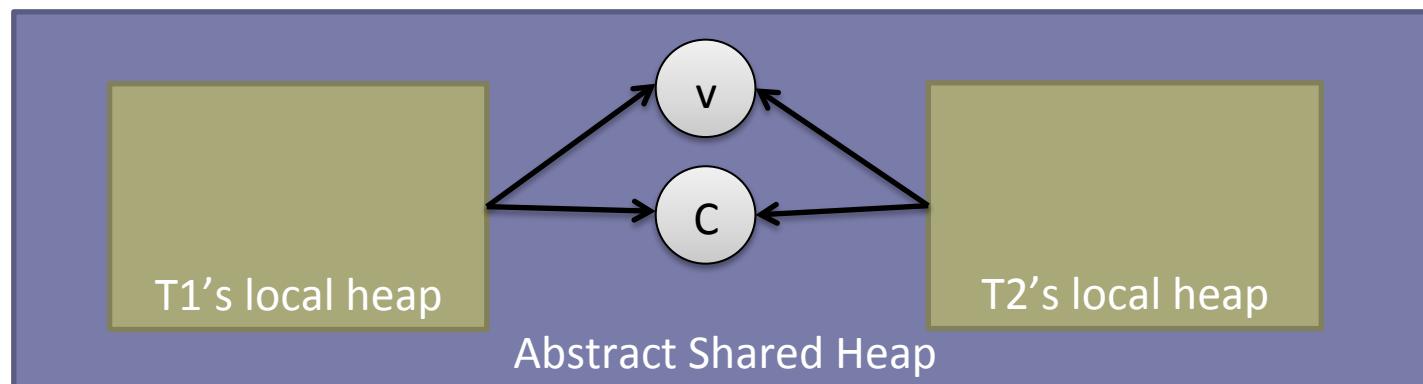
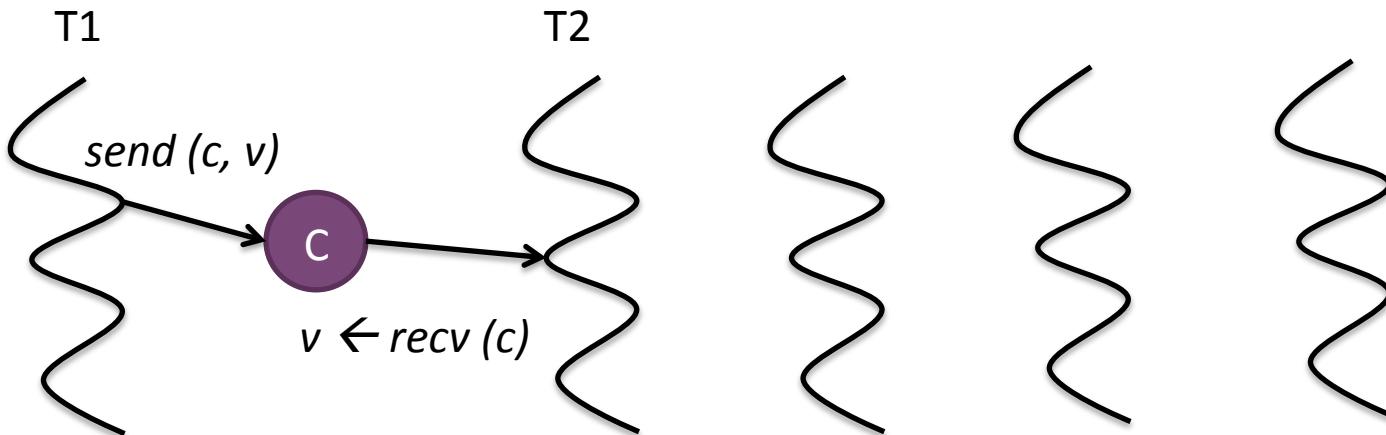
(d) Garbage collection overhead

# Under the hood



Before Communication

# Under the hood



After Communication