


Programming and *proving* distributed systems with *persistent data structures*


“KC” Sivaramakrishnan




Collaborative Apps




Google Docs



Network Partitions




Network Partitions — Google docs




Enabling offline sync for one account prevents other accounts from working offline

Local-first software



Local-first software



How do we build such applications?


Make data types aware of replication

CRDTs

- Conflict-free Replicated Data Types (CRDTs)
 - ◆ Multiple *replicas* of the data types
- Supports *local* operations
- Share updates *asynchronously* and ensure *convergence*
 - ◆ *Strong eventual consistency*

Grow-only Set

- Supports *add* and *lookup* operations




A set with only *add* and *lookup* is *monotonic*

Include remove operation


- Let's include *remove* operation

- ◆ No longer monotonic




Two-phase set

- Represent the set with a pair of sets to track *additions* and *removals* — $A * R$
 - ◆ Lookup is performed in A / R
 - ◆ Merge is pair-wise union of A and R sets



Two-phase set — Observations



- Monotonic — Simulate remove with adds
- Remove-wins semantics
- *Reengineer the set implementation*
- Tombstones — *elements removed by adding to R set!*
- *Removed elements gone forever*

Challenges with CRDTs

- Monotonicity forces *reengineering of data structures from scratch*
 - ◆ Challenges proving correctness of *even sequential operations*
- Not space and time-efficient
 - ◆ *Tombstones* affect time- and space-efficiency
- Little attention has been paid to *composition* of RDTs
 - ◆ Parametric polymorphism for RDTs
 - ✿ Like to compose 'a `set` with a `counter` to get `counter set`
 - ◆ *How to compose proofs of correctness?*


Can we do better?

Sequential data types

+

Git

Distributed Version Control Systems




Mergeable Replicated Data Types

- MRDTs — DVCS for *data types* rather than just *text files*
 - ◆ *Branches* are *replica states*
- Sequential data type + 3-way merge = replicated data type!

Counter MRDT


```
module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```



Set MRDT


```
let merge ~lca ~v1 ~v2 =  
  (lca ∩ v1 ∩ v2) (* common elements *)  
  ∪ (v1 - lca)    (* added in v1 *)  
  ∪ (v2 - lca)    (* added in v2 *)
```

*isomorphic to counter
merge if you squint*




Set MRDT

```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* common elements *)
  ∪ (v1 - lca)    (* added in v1 *)
  ∪ (v2 - lca)    (* added in v2 *)
```



Set MRDT — Add after Remove




MRDTs and Causal History

- How did we get away with no tombstones in MRDT set?
- Tombstones in CRDTs record *history*
- *Git records the causal history in MRDTs!*
 - ◆ Presented via LCA in 3-way merge
- How does Git keep track of causal history efficiently?

Persistent Data Structures

Git store




- Branches / tags
- Mutable
- Stores the files under version control
- Immutable, append-only & content addressed
 - hash → object

Block store and persistence

Added original
`/text.txt` at
`/bak/text.txt`

Added `/new.txt`
Modified `/test.txt`


Added `/test.txt`



Irmin store



- A Git-like distributed database
- MRDTs are executed on top of Irmin




Persistent algebraic data types:
Stacks, Queues, Ropes, Balanced binary
trees, etc.

Kaki et al, “Mergeable Replicated Data Types”,
OOPSLA 2019

Commit DAG

- Commit nodes form a DAG
 - ◆ Captures causal history \Rightarrow *Happens-before* relation
- MRDTs provide *causal* consistency
 - ◆ Strongest consistency level without coordination
- LCA discovered by traversing the commit DAG



Forgetting History

- Git remembers the entire history
 - ◆ Useful if *provenance* is necessary
- If not, keeping entire history is wasteful
 - ◆ Nodes will *run out of storage* quicker
- But, history needed for LCA in a 3-way merge
- *How much history should we keep?*
- *Any commits older than the latest commit “K” known by all replicas can be GCed*

Dubey et al, “Banyan: Coordination-free Distributed Transactions over Mergeable Types”, APLAS 2020

Dubey, “Banyan: Coordination-free Distributed Transactions over Mergeable Types”, MS Thesis, IIT Madras




MRDTs = Sequential data types + 3-way merge

Does this make proving MRDTs correct easier?

Is our set an add-wins set?

- *add-wins* when there is a concurrent add and remove of the same element
- Also known as *Observed-Removed set (OR-set)*


```
let merge ~lca ~v1 ~v2 =
  (lca ∩ v1 ∩ v2) (* common elements *)
  ∪ (v1 - lca)    (* added in v1 *)
  ∪ (v2 - lca)    (* added in v2 *)
  {
    {
      {} (* common elements *)
      ∪ {} (* added in v1 *)
      ∪ {} (* added in v2 *)
    }
    = {} ∪ {} ∪ {}
    = {} (expected {1})
```



- *Our set is not add-wins set!*
- Convergence is not sufficient; *Intent* is not preserved

Concretising Intent

- A *formal specification language* to capture the *intent* of the MRDT
 - ◆ Must be rich enough to capture distributed execution
- Even *simple* data types attract enormous *complexity* when made *distributed*



- *Mechanization* to bridge the gap between spec and impl

Peepul — Certified MRDTs


- F^* library implementing and proving MRDTs
 - ★ F^* — proof-oriented, solver-aided PL
- Specification language is event-based
 - ★ Burckhardt et al. “Replicated Data Types: Specification, Verification and Optimality”, POPL 2014
- *Replication-aware simulation* to connect *specification* with *implementation*
- *Space- and time-efficient* implementations
 - ★ 1st certified implementation of a $O(1)$ replicated queue with $O(n)$ merge.
- *Composition* of MRDTs and their proofs!
- Extracted OCaml RDTs are compatible with *Irmin*



Fixing Add-wins Set

- Discriminate duplicate additions by associating a unique id


```
{ } u      (* common *)  
{ (a,2) } u (* added left *)  
{ }         (* added right *)  
= { (a,2) }
```



MRDT Implementation


$$D_\tau = (\Sigma, \sigma_0, do, merge)$$

- 1: $\Sigma = \mathcal{P}(\mathbb{N} \times \mathbb{N})$
 - 2: $\sigma_0 = \{\}$
 - 3: $do(rd, \sigma, t) = (\sigma, \{a \mid (a, t') \in \sigma\})$
 - 4: $do(add(a), \sigma, t) = (\sigma \cup \{(a, t)\}, \perp)$
 - 5: $do(remove(a), \sigma, t) = (\{e \in \sigma \mid fst(e) \neq a\}, \perp)$
 - 6: $merge(\sigma_{lca}, \sigma_a, \sigma_b) =$
 $(\sigma_{lca} \cap \sigma_a \cap \sigma_b) \cup (\sigma_a - \sigma_{lca}) \cup (\sigma_b - \sigma_{lca})$
- Unique Lamport Timestamp*



Specifying Add-wins Set

Abstract state $I = \langle E, oper, rval, time, vis \rangle$



$$\mathcal{F}_{orset}(rd, \langle E, oper, rval, time, vis \rangle) = \{ a \mid \exists e \in E. oper(e)$$

$$= add(a) \wedge \neg(\exists f \in E. oper(f) = remove(a) \wedge e \xrightarrow{vis} f) \}$$


Simulation Relation

- Connects the *abstract state* with the *concrete state*
- For the add-wins set,


$$\begin{aligned}\mathcal{R}_{sim}(I, \sigma) \iff (\forall(a, t) \in \sigma \iff \\ (\exists e \in I.E \wedge I.\text{oper}(e) = \text{add}(a) \wedge I.\text{time}(e) = t \wedge \\ \neg(\exists f \in I.E \wedge I.\text{oper}(f) = \text{remove}(a) \wedge e \xrightarrow{vis} f)))\end{aligned}$$

- The main verification effort is to show that the relation above is indeed a *simulation relation*
 - ★ Shown separately for *operations* and *merge function*
 - ★ Proof by induction on the execution trace


Verifying operations




Verifying Merge function



Verifying Merge function



Verifying Merge function



Verification effort

MRDTs verified	#Lines code	#Lines proof	#Lemmas	Verif. time (s)
Increment-only counter	6	43	2	3.494
PN counter	8	43	2	23.211
Enable-wins flag	20	58 81 89	3 6 7	1074 171 104
LWW register	5	44	1	4.21
G-set	10	23 28 33	0 1 2	4.71 2.462 1.993
G-map	48	26	0	26.089
Mergeable log	39	95	2	36.562
OR-set (§2.1.1)	30	36 41 46	0 1 2	43.85 21.656 8.829
OR-set-space (§2.1.2)	59	108	7	1716
OR-set-spacetime	97	266	7	1854
Queue	32	1123	75	4753

Composing RDTs is HARD!



Martin Kleppmann
@martinkl

...

Today in “distributed systems are hard”: I wrote down a simple CRDT algorithm that I thought was “obviously correct” for a course I’m teaching. Only 10 lines or so long. Found a fatal bug only after spending hours trying to prove the algorithm correct. 😭

4:18 AM · Nov 13, 2020 · Tweetbot for iOS

41 Retweets 4 Quote Tweets 541 Likes



Martin Kleppmann @martinkl · Nov 13, 2020

...

The interesting thing about this bug is that it comes about only from the interaction of two features. A LWW map by itself is fine. A set in which you can insert and delete elements (but not update them) is fine. The problem arises only when delete and update interact.



1

16



Composing IRC-style chat

- Build IRC-style group chat
 - ★ Send and read messages in channels
- Represent application state as a **map** **MRDT**
 - ★ String (*channel name*) keys → **mergeable log** **MRDT** values
 - ★ Mergeable log — message + timestamp; merge ordered by timestamp
- **Goal:**
 - ★ **map** and **log** proved correct separately
 - ★ Use the proof of underlying RDTs to prove chat application correctness

Generic Map MRDT

Implementation

$\mathcal{D}_{\alpha\text{-map}} = (\Sigma, \sigma_0, do, merge_{\alpha\text{-map}})$ where

1: $\Sigma_{\alpha\text{-map}} = \mathcal{P}(\text{string} \times \Sigma_\alpha)$

2: $\sigma_0 = \{\}$

3: $\delta(\sigma, k) = \begin{cases} \sigma(k), & \text{if } k \in \text{dom}(\sigma) \\ \sigma_0, & \text{otherwise} \end{cases}$

4: $do(\text{set}(k, o_\alpha), \sigma, t) =$
let $(v, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$ in $(\sigma[k \mapsto v], r)$

5: $do(\text{get}(k, o_\alpha), \sigma, t) =$
let $(_, r) = do_\alpha(o_\alpha, \delta(\sigma, k), t)$ in (σ, r)

6: $merge_{\alpha\text{-map}}(\sigma_{lca}, \sigma_a, \sigma_b) =$
 $\{(k, v) \mid (k \in \text{dom}(\sigma_{lca}) \cup \text{dom}(\sigma_a) \cup \text{dom}(\sigma_b)) \wedge$
 $v = merge_\alpha(\delta(\sigma_{lca}, k), \delta(\sigma_a, k), \delta(\sigma_b, k))\}$

The values in the MRDT map are MRDTs

Merge uses the merge of the underlying value type!

Simulation Relation

$\mathcal{R}_{sim\text{-}\alpha\text{-map}}(I, \sigma) \iff \forall k.$

1: $(k \in \text{dom}(\sigma) \iff \exists e \in I.E. \text{oper}(e) = \text{set}(k, _)) \wedge$

2: $\mathcal{R}_{sim\text{-}\alpha}(\text{project}(k, I), \delta(\sigma, k))$


Simulation relation appeals to the value type's simulation relation!

Composing IRC-style chat

- IRC app state is constructed by instantiating *generic map* with *mergeable log*
- The proof of correctness of the chat application directly follows from the composition.

★ See paper for details!

Soundarapandian et al, “Certified Mergeable Replicated Data Types”, PLDI 2022




Verification is still too hard!

MRDTs verified	#Lines code	#Lines proof	#Lemmas	Verif. time (s)
Increment-only counter	6	43	2	3.494
PN counter	8	43	2	23.211
Enable-wins flag	20	58 81 89	3 6 7	1074 171 104
LWW register	5	44	1	4.21
G-set	10	23 28 33	0 1 2	4.71 2.462 1.993
G-map	48	26	0	26.089
Mergeable log	39	95	2	36.562
OR-set (§2.1.1)	30	36 41 46	0 1 2	43.85 21.656 8.829
OR-set-space (§2.1.2)	59	108	7	1716
OR-set-spacetime	97	266	7	1854
Queue	32	1123	75	4753

Queue MRDT — Implementation

- Two-list queue + merge function
- “At least once” semantics for dequeue
- Each element e enqueueued is (e, t) where t is the unique Lamport timestamp



Queue MRDT — Specification

$$\begin{aligned}\text{match}_I(e_1, e_2) \Leftrightarrow & I.\text{oper}(e_1) = \text{enqueue}(a) \\ & \wedge I.\text{oper}(e_2) = \text{dequeue} \wedge a = I.\text{rval}(e_2)\end{aligned}$$


- $\text{AddRem}(I) : \forall e \in I.E. I.\text{oper}(e) = \text{dequeue} \wedge I.\text{rval}(e) \neq \text{EMPTY} \implies \exists e' \in I.E. \text{match}_I(e', e)$
 - $\text{Empty}(I) : \forall e_1, e_2, e_3 \in I.E. I.\text{oper}(e_1) = \text{dequeue} \wedge I.\text{rval}(e_1) = \text{EMPTY} \wedge I.\text{oper}(e_2) = \text{enqueue}(a) \wedge e_2 \xrightarrow{I.\text{vis}} e_1 \implies \exists e_3 \in I.E. \text{match}_I(e_2, e_3) \wedge e_3 \xrightarrow{I.\text{vis}} e_1$
 - $\text{FIFO}_1(I) : \forall e_1, e_2, e_3 \in I.E. I.\text{oper}(e_1) = \text{enqueue}(a) \wedge \text{match}_I(e_2, e_3) \wedge e_1 \xrightarrow{I.\text{vis}} e_2 \implies \exists e_4 \in I.E. \text{match}_I(e_1, e_4)$
 - $\text{FIFO}_2(I) : \forall e_1, e_2, e_3, e_4 \in I.E. \neg(\text{match}_I(e_1, e_4) \wedge \text{match}_I(e_2, e_3) \wedge e_1 \xrightarrow{I.\text{vis}} e_2 \wedge e_3 \xrightarrow{I.\text{vis}} e_4)$
-
- Extremely *hard* to write specs over event-based structures
 - Simulation relations are *harder*

Better Specification

- *Sequential data type + constraints* as the specification for MRDT
 - ◆ Constraints — ordering, commutativity, duplication, ...
 - ◆ MRDT behaviour = constrained linearisation + Sequential DT

Op1	Op2	Order
add(a)	rem(a)	Op2, Op1
rem(a)	add(a)	Op1, Op2
add(_)	add(_)	Any
rem(_)	rem(_)	Any
add(a)	rem(b)	Any
rem(a)	add(b)	Any

Add-wins set ordering constraint



Summary

- MRDT simplify the construction of RDTs
 - ♦ Sequential data types + 3-way merge functions
- *Persistent data structures* to efficiently record causal history
- 3-way merge function is a pathway to verifying MRDTs

