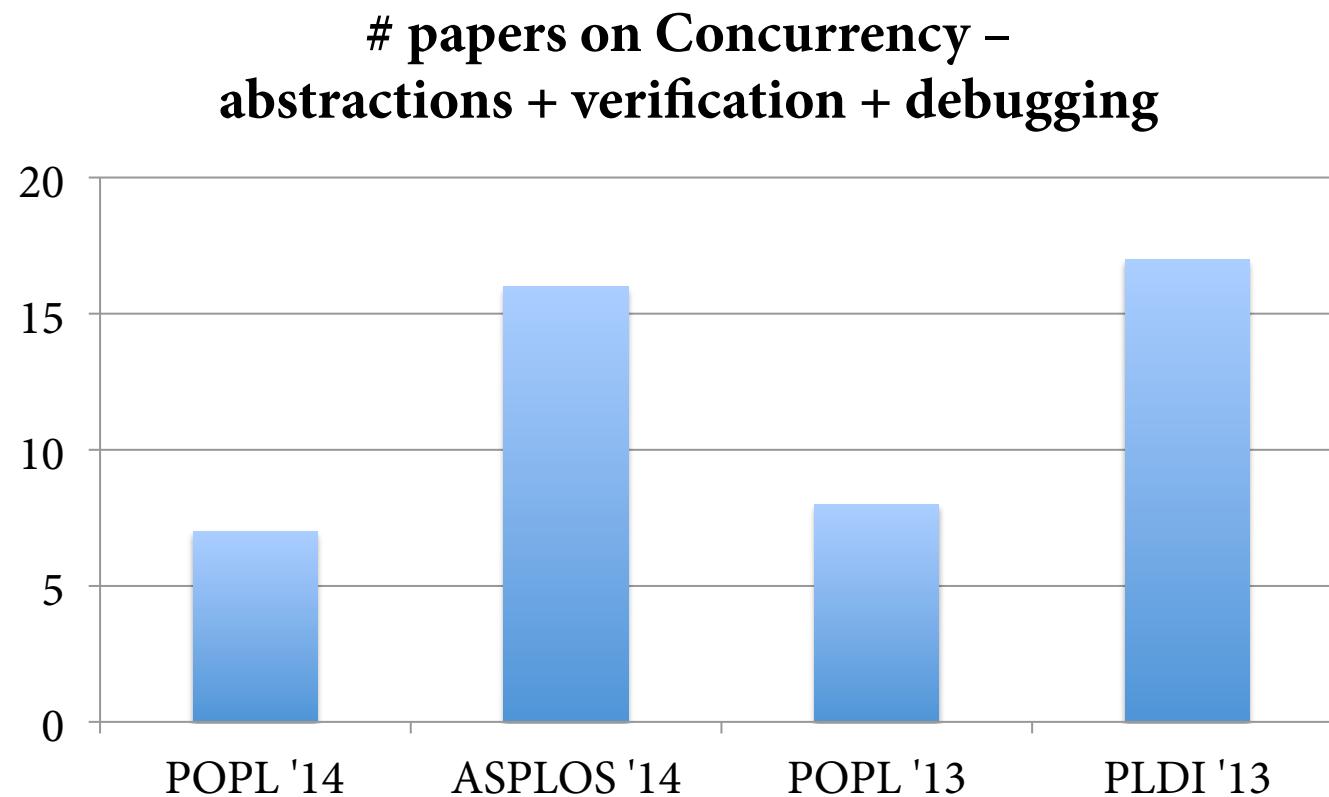


Functional Abstractions for Practical and Scalable Concurrent Programming

KC Sivaramakrishnan
Purdue University

Concurrent programming is (still) hard!



Concurrency bugs can be disastrous

Concurrency bugs can be disastrous

2003 Northeast blackout^[1]



- Data race disables alarm system
- Effect
 - 256 power stations go offline
 - ~7 hr major blackout
 - 11 fatalities, \$6 billion

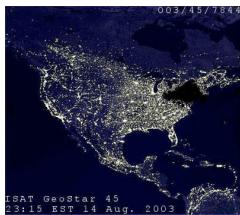
[1] <http://www.scientificamerican.com/article/2003-blackout-five-years-later/>

[2] <http://sunnyday.mit.edu/papers/therac.pdf>

[3] <http://www.cnbc.com/id/100587334>

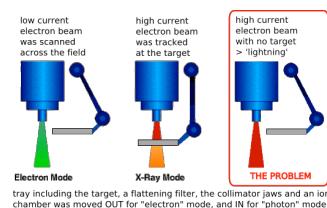
Concurrency bugs can be disastrous

2003 Northeast blackout^[1]



- Data race disables alarm system
- Effect
 - 256 power stations go offline
 - ~7 hr major blackout
 - 11 fatalities, \$6 billion

Therac-25 incidents^[2]



- Data race b/w UI and controller
- 6 fatalities

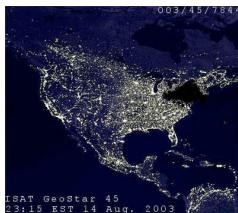
[1] <http://www.scientificamerican.com/article/2003-blackout-five-years-later/>

[2] <http://sunnyday.mit.edu/papers/therac.pdf>

[3] <http://www.cnbc.com/id/100587334>

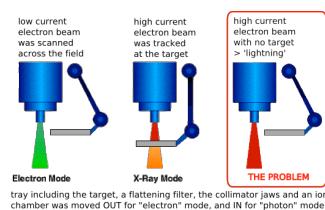
Concurrency bugs can be disastrous

2003 Northeast blackout^[1]



- Data race disables alarm system
- Effect
 - 256 power stations go offline
 - ~7 hr major blackout
 - 11 fatalities, \$6 billion

Therac-25 incidents^[2]



- Data race b/w UI and controller
- 6 fatalities

Facebook IPO @ NASDAQ^[3]



- Race between validation and new orders arrival
- NASDAQ compensation = \$62 million

[1] <http://www.scientificamerican.com/article/2003-blackout-five-years-later/>

[2] <http://sunnyday.mit.edu/papers/therac.pdf>

[3] <http://www.cnbc.com/id/100587334>



Safe and scalable
concurrent program



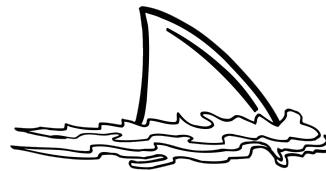
Safe and scalable
concurrent program

Fences, Locks,
Condition variables, etc.

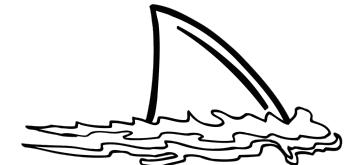




Safe and scalable
concurrent program

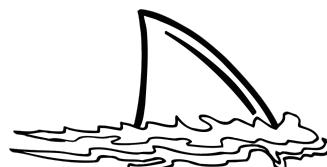


Data race



Deadlock

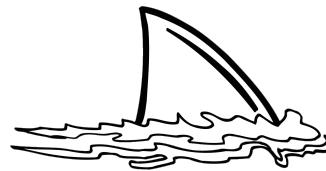
Fences, Locks,
Condition variables, etc.



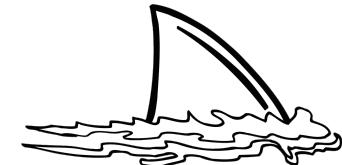
Atomicity
violations



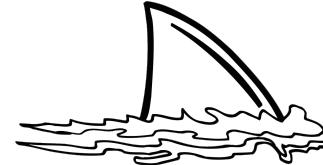
Safe and scalable concurrent program



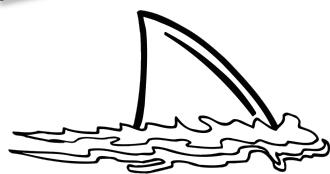
Data race



Deadlock



Weak memory semantics



Eventual consistency

Fences, Locks,
Condition variables, etc.



No cache coherence



Atomicity violations



Stick to
Status Quo



Safe and scalable
concurrent program



Stick to
Status Quo

Testing



Safe and scalable
concurrent program



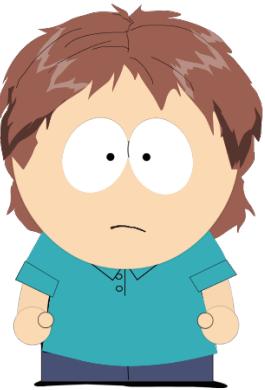
Stick to
Status Quo

Testing

Verification



Safe and scalable
concurrent program



Stick to
Status Quo

Testing

Verification

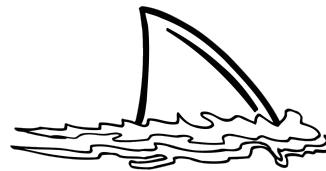
PL Support



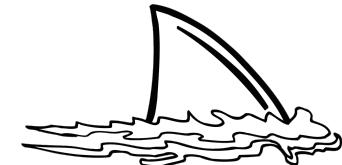
Safe and scalable
concurrent program



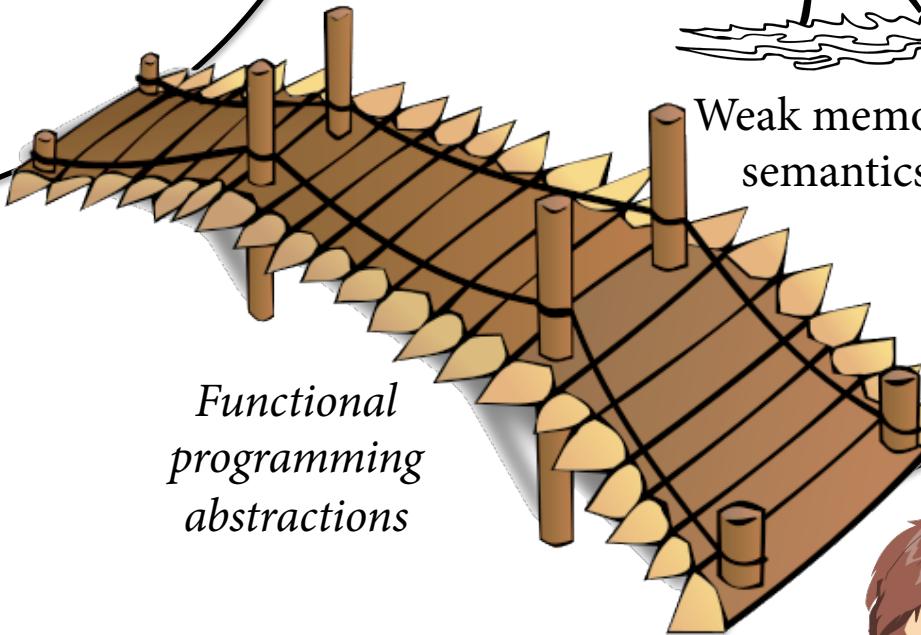
Safe and scalable concurrent program



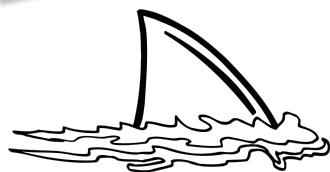
Data race



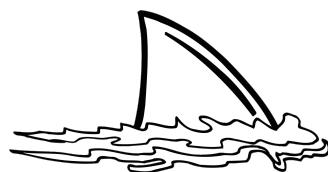
Deadlock



*Functional
programming
abstractions*



Eventual consistency



Non-cache coherence



Atomicity violations

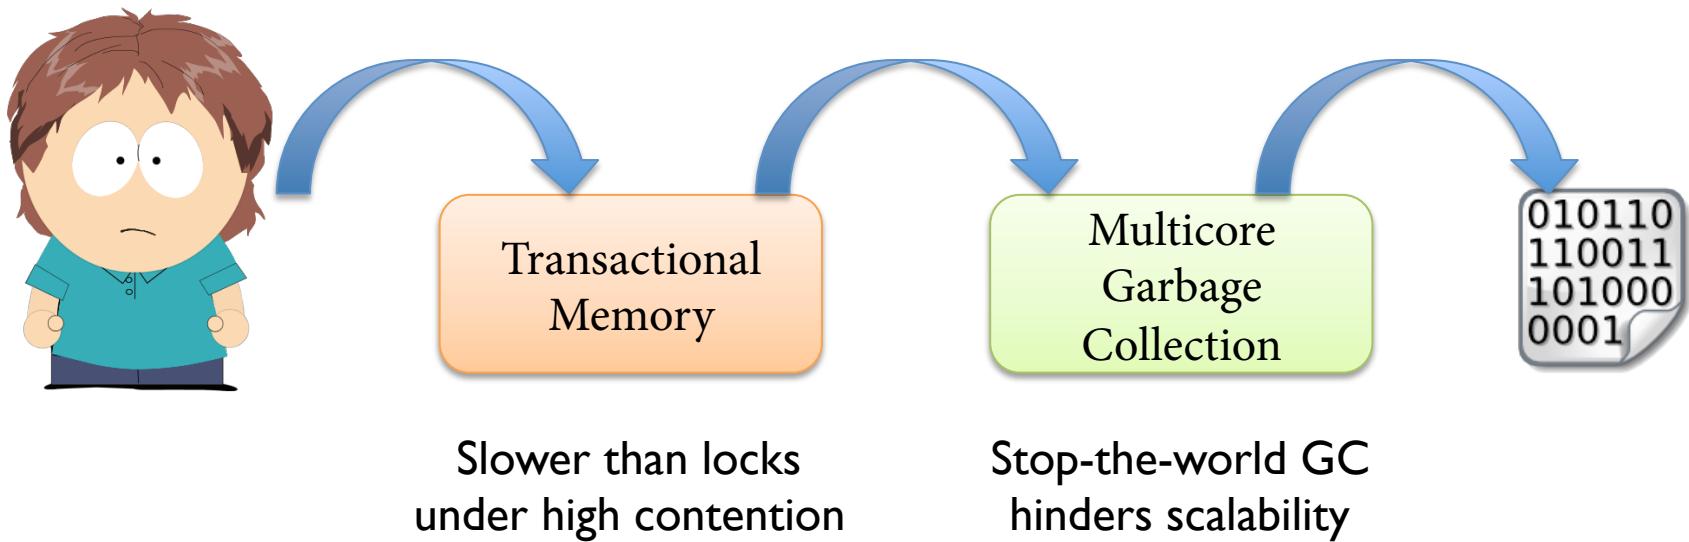


Programming abstractions *simplify* Concurrent Programming

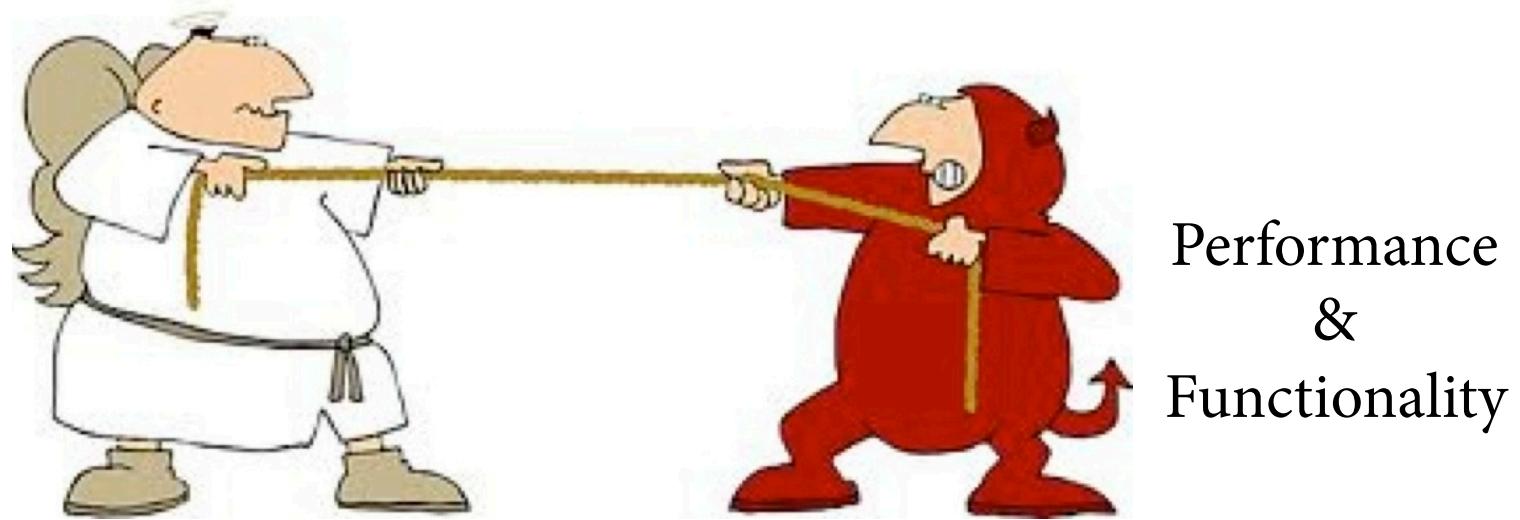
Transactional
Memory

Multicore
Garbage
Collection

But, programming abstractions introduce a *level of indirection*

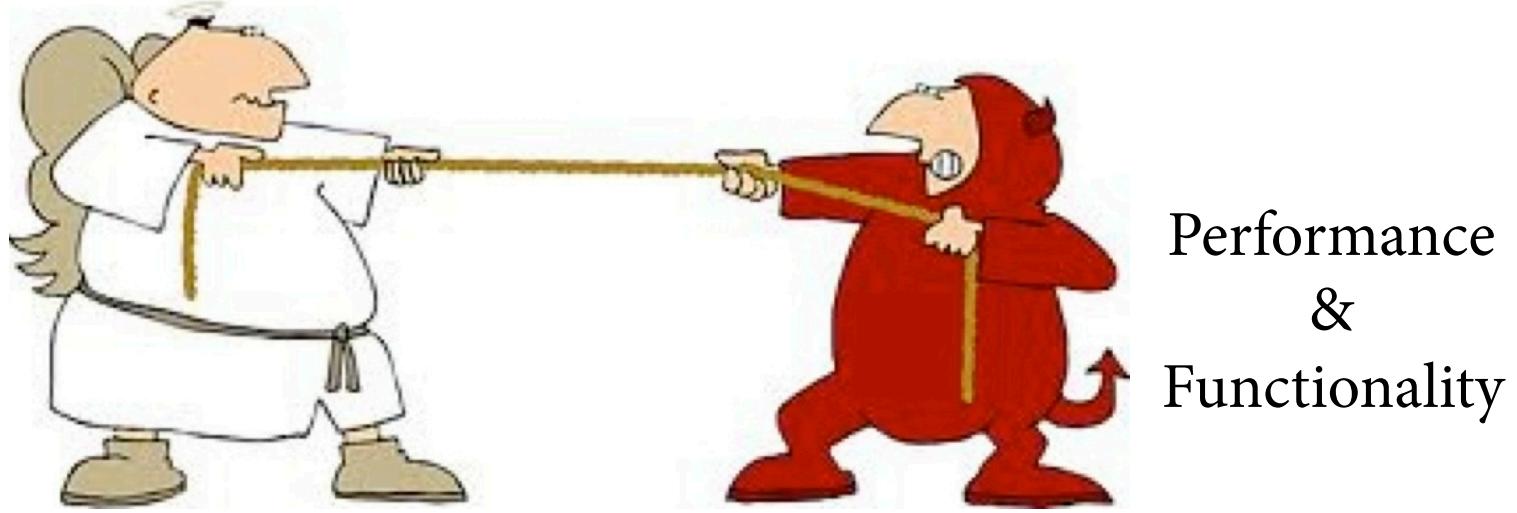


Safety
&
Simplicity



Performance
&
Functionality

Safety
&
Simplicity



Always desirable to marry the two whenever possible!

MultiMLton

Intel SCC
48-core Non-cache-coherent



Azul Vega 3
864-core CC-UMA



Compute clouds



MultiMLton

Intel SCC
48-core Non-cache-coherent



Azul Vega 3
864-core CC-UMA



Compute clouds



Asynchronous CML [PLDI '11]
Memoizing Communication [ICFP '09]
Rx-CML – optimistic CML [PADL '14]

Parasitic threads [DAMP '10, JFP '14]
Thread-local GC [ISMM '12, MARC '12, JFP '14]

Language Design

Runtime Systems

MultiMLton

Intel SCC
48-core Non-cache-coherent



Azul Vega 3
864-core CC-UMA



Compute clouds



Asynchronous CML
Memoizing Communication
Rx-CML – optimistic CML

[PLDI '11]
[ICFP '09]
[PADL '14]

Parasitic threads [DAMP '10, JFP '14]
Thread-local GC [ISMM '12, MARC '12, JFP '14]

Language Design Runtime Systems

Scheduler activations
for Haskell

Schedulers for Haskell threads as Haskell libraries
[In submission to OOPSLA '14]

Sting (Java)

Session type based protocol optimization
[Coordination '10, SCP '13]

MultiMLton

Language Design

Rx-CML – optimistic CML [PADL '14]

Scheduler activations
for Haskell

Schedulers for Haskell threads as Haskell libraries
[In submission to OOPSLA '14]

Rx-CML : A Prescription for Safely Relaxing Synchrony

Abstraction



Domain

Synchronous
communication



Asynchronous
distributed system

Transactional
memory



Shared memory
concurrency

Garbage
collection



Memory
management

Abstraction

Synchronous communication

Transactional memory

Garbage collection



Domain

Asynchronous distributed system

Shared memory concurrency

Memory management



Synchronous communication =
atomic { data transfer +
synchronization }

Abstraction

Synchronous communication

Transactional memory

Garbage collection



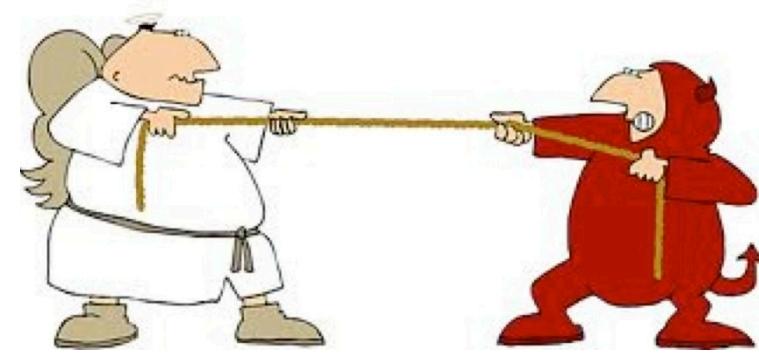
Domain

Asynchronous distributed system

Shared memory concurrency

Memory management

Synchronous communication =
atomic { data transfer +
synchronization }



synchrony

latency

Can we **discharge** synchronous communications
asynchronously and ensure **observable equivalence**?

Can we **discharge** synchronous communications
asynchronously and ensure **observable equivalence**?

Formalize:

$$\llbracket \text{send } (c, v) \rrbracket k \equiv \llbracket \text{asend } (c, v) \rrbracket k$$

Implement:

Distributed Concurrent ML in MultiMLton +
Speculative execution framework

Concurrent ML

val spawn : (unit -> unit) -> thread_id

Thread
creation

val channel : unit -> 'a chan

val send : 'a chan * 'a -> unit

val recv : 'a chan -> 'a

Synchronous
message passing

val sendEvt : 'a chan * 'a -> unit event

val recvEvt : 'a chan -> 'a event

val sync : 'a event -> 'a

val never : 'a event

val alwaysEvt : 'a -> 'a event

val wrap : 'a event -> ('a -> 'b) -> 'b event

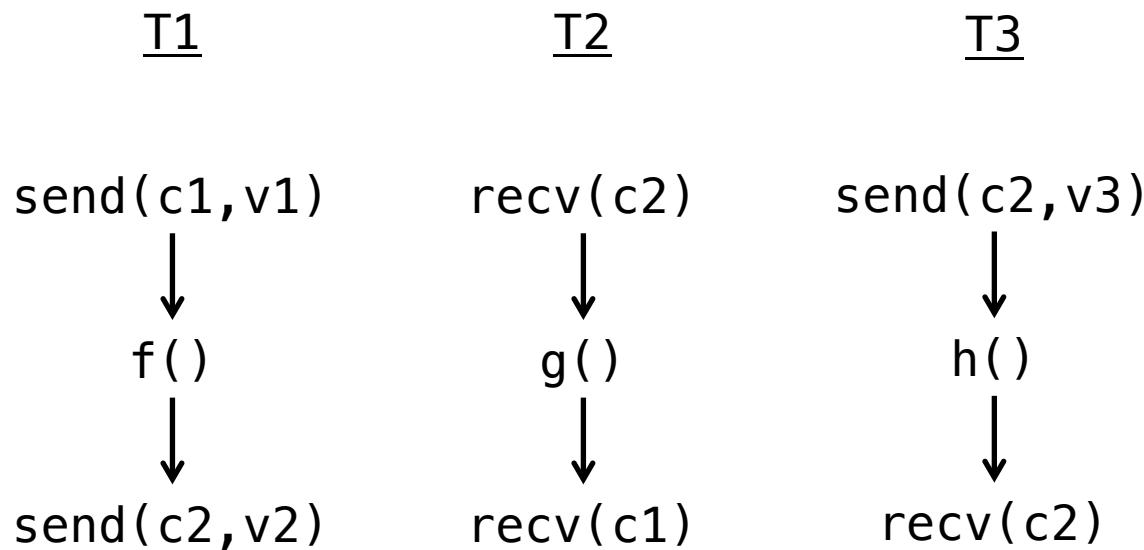
val guard : (unit -> 'a event) -> 'a event

val choose : 'a event list -> 'a event

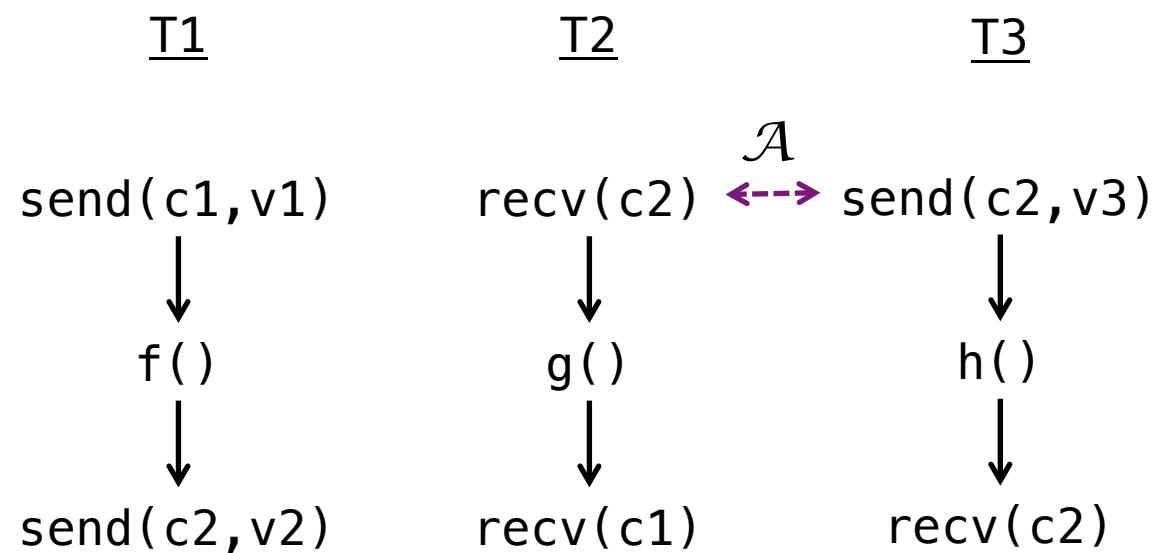
First-class events

...

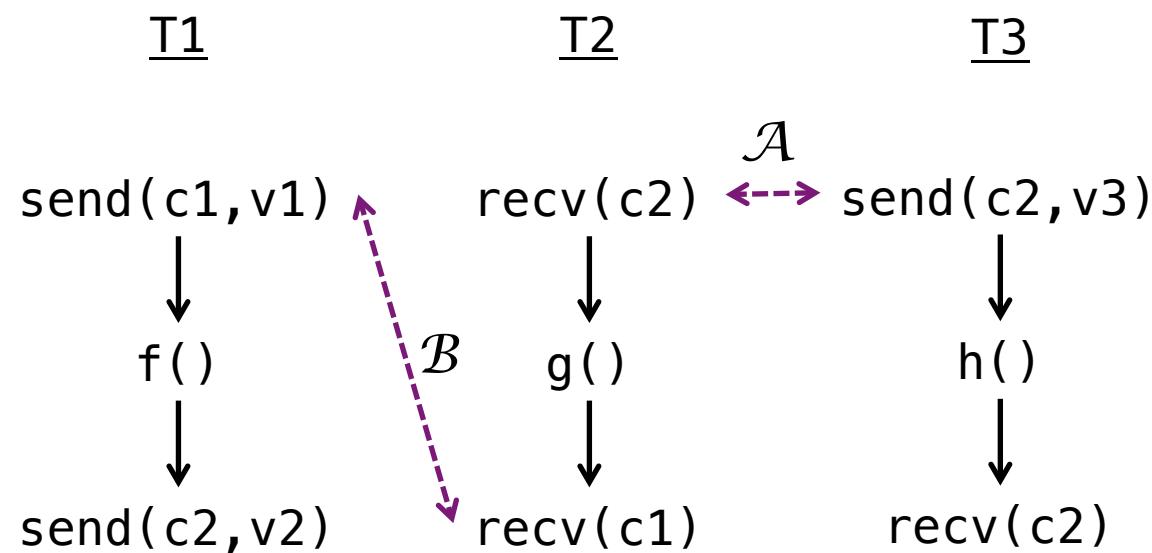
T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



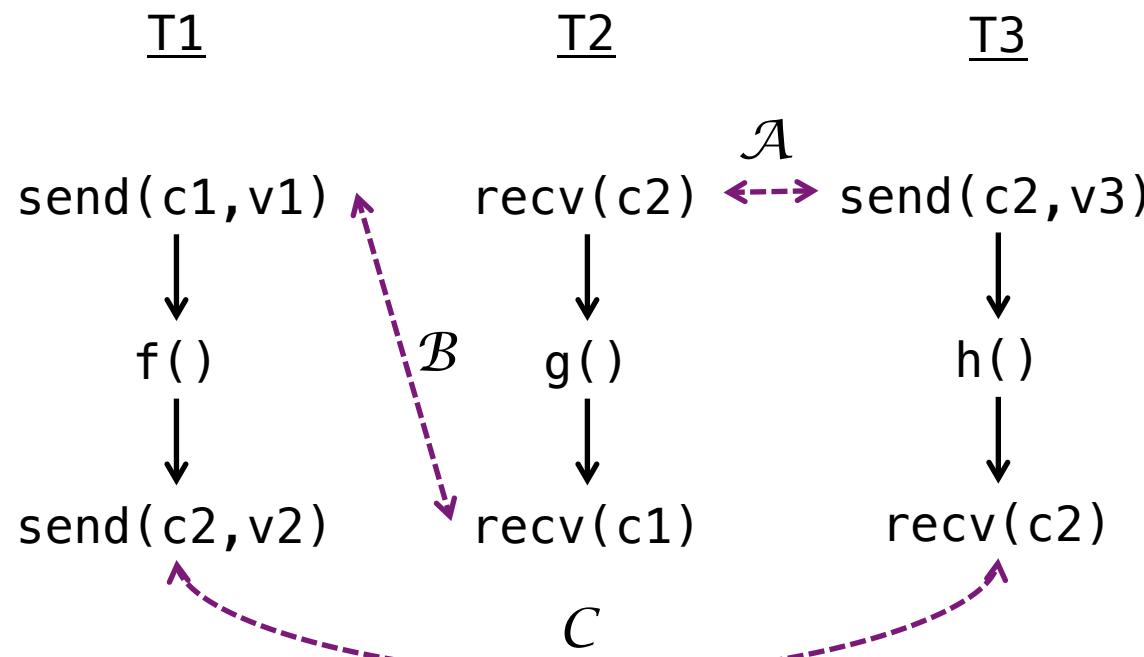
T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



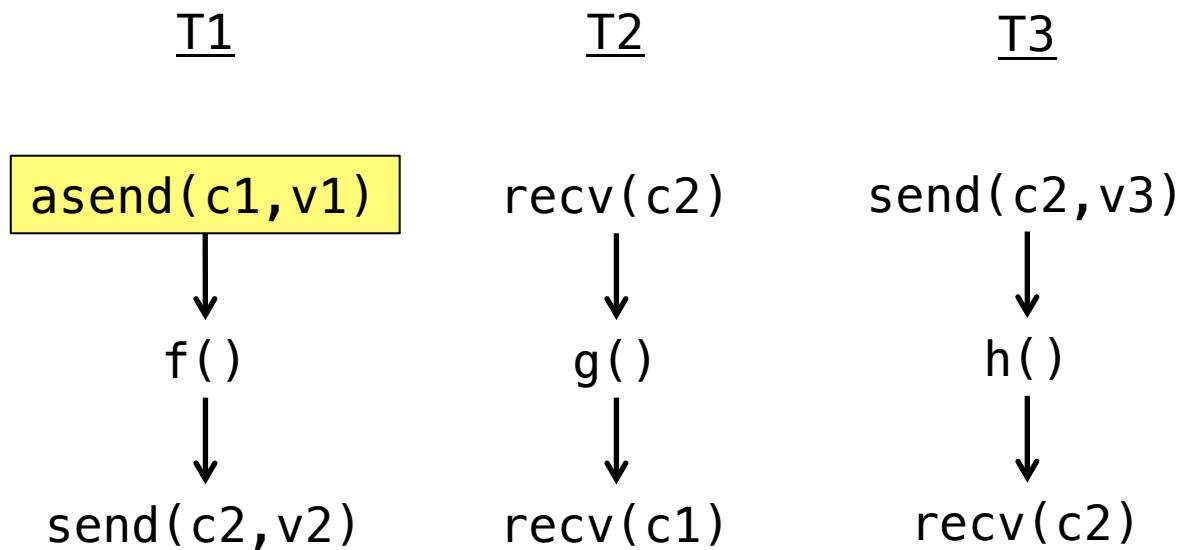
T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



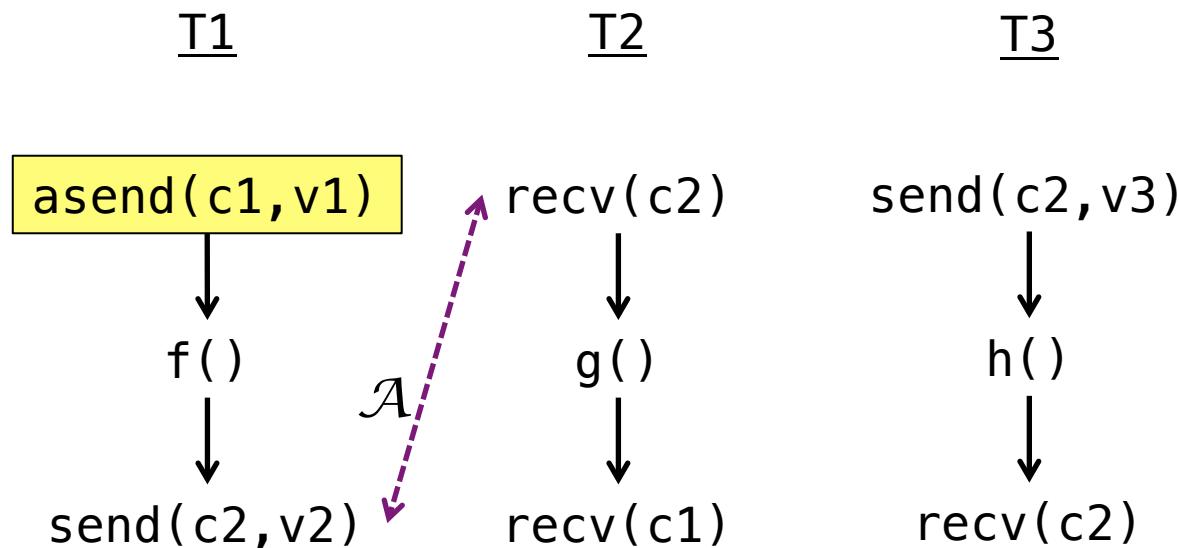
T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



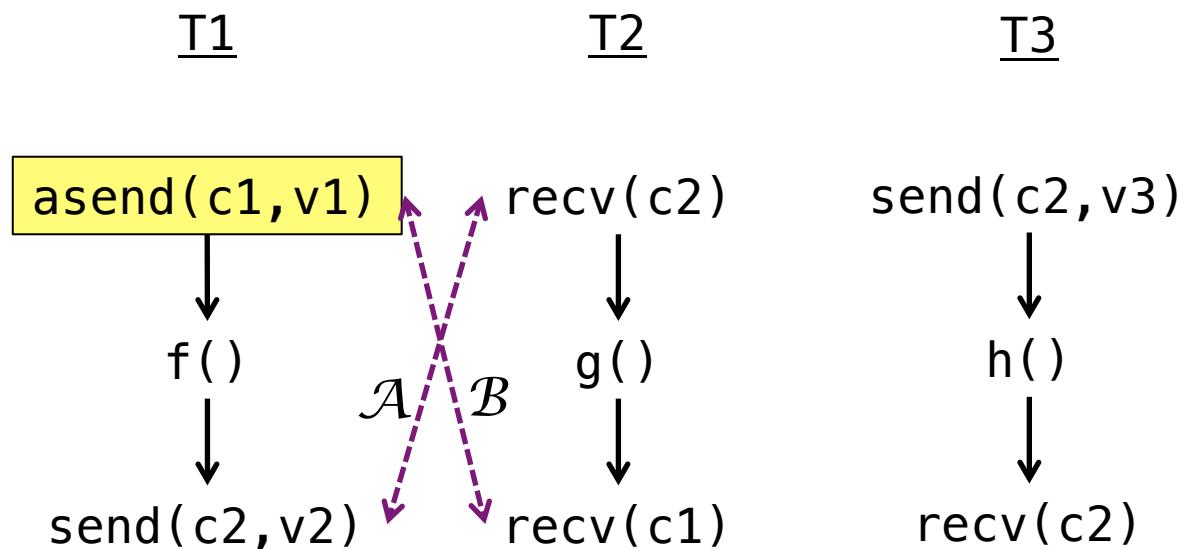
T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



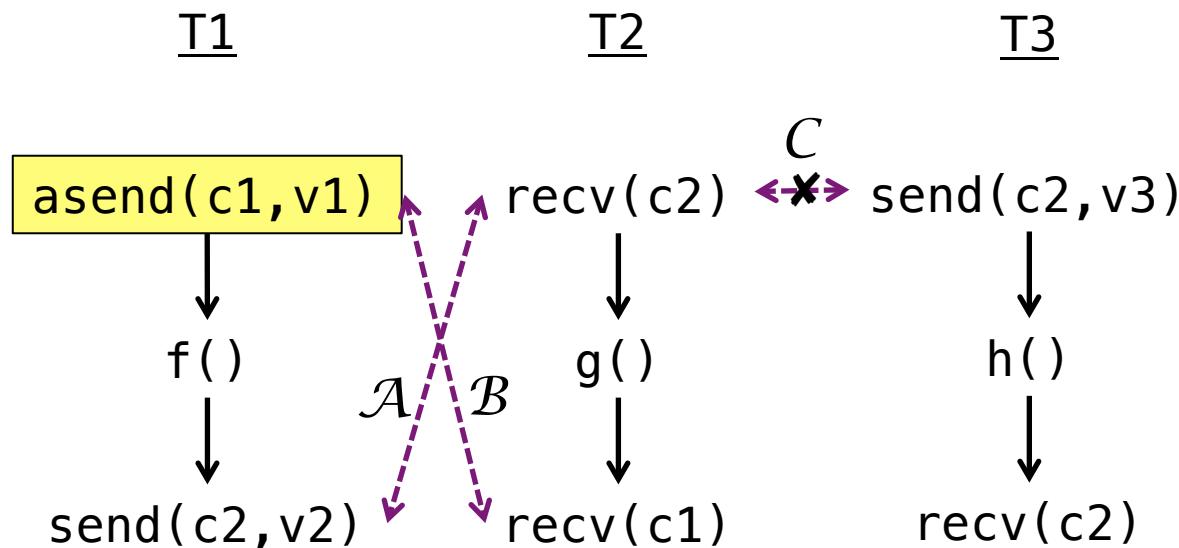
T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



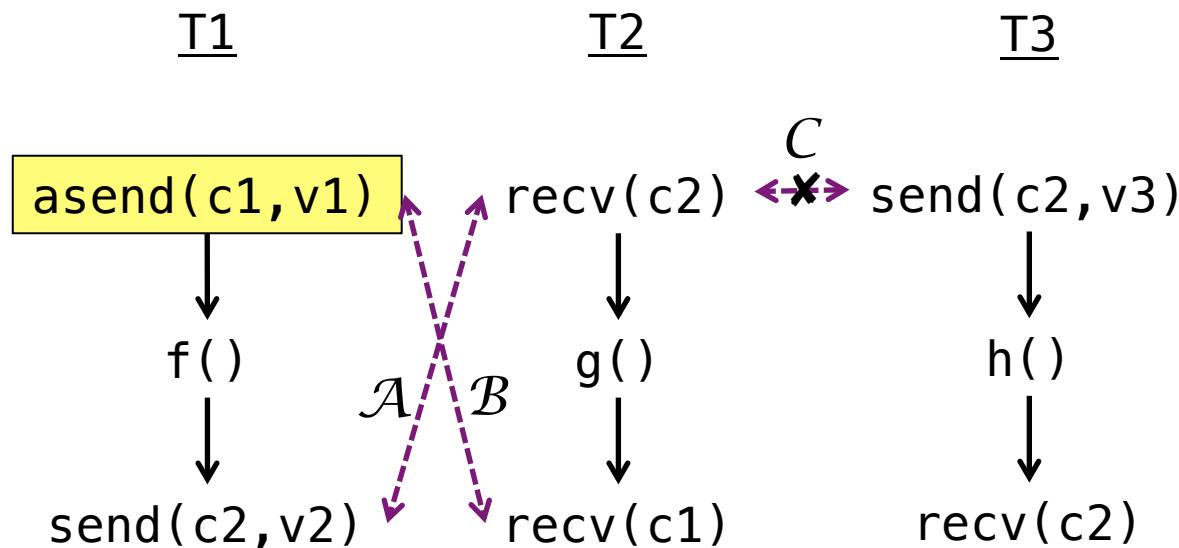
T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)



Cyclic dependence \Rightarrow divergent behavior

Distributed group chat app

No central server & causal dependence → *causal broadcast*

Distributed group chat app

No central server & causal dependence → *causal broadcast*

```
fun bsend (BCHAN (vcList, acList), v: 'a, id: int) : unit =
let
  val _ = map (fn vc => if (vc = nth (vcList, id)) then () else send (vc, v))
    vcList (* phase 1 -- Value distribution *)
  val _ = map (fn ac => if (ac = nth (acList, id)) then () else recv ac)
    acList (* phase 2 -- Acknowledgments *)
in ()
end
```

synchronously send values

*prevent receivers from proceeding until
all members have received the value*

Distributed group chat app: Performance

Distributed group chat app: Performance

Site A – US East

bsend(c,x)

brecv(c)

Site B – EU

brecv(c)

bsend(c,y)

Site C – US West

brecv(c)

brecv(c)

Distributed group chat app: Performance

Site A – US East

bsend(c,x)

brecv(c)

Site B – EU

brecv(c)

bsend(c,y)

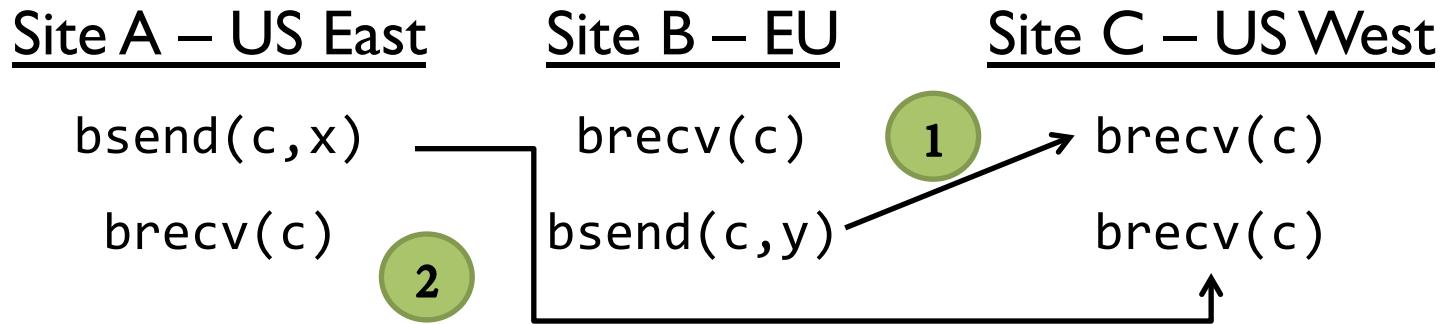
Site C – US West

brecv(c)

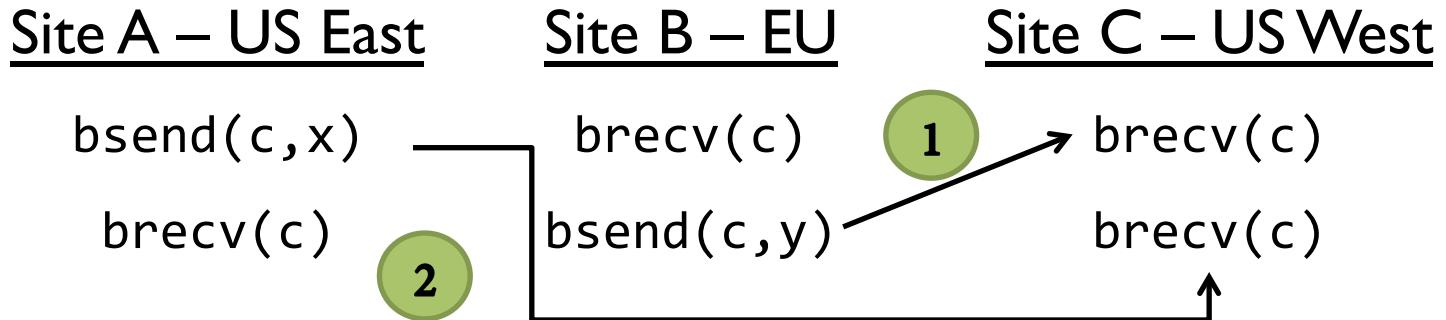
brecv(c)



Distributed group chat app: Performance



Distributed group chat app: Performance



<i>Execution</i>	<i>Avg.time (ms)</i>	<i>Errors</i>
<i>Sync</i>	1540	0
<i>Unsafe Async</i>	520	7
<i>Safe Async (R^{CML})</i>	533	0

Formalization

Reason axiomatically

$$E := \langle P, A, \rightarrow_{po}, \rightarrow_{co} \rangle$$

Happens-before relation

$$\begin{aligned}\rightarrow_{hb} &= (\rightarrow_{po} \cup \rightarrow_{td} \cup \\ &\quad \{(\alpha, \beta) \mid \alpha \rightarrow_{co} \alpha' \wedge \alpha' \rightarrow_{po} \beta\} \cup \\ &\quad \{(\beta, \alpha) \mid \beta \rightarrow_{po} \alpha' \wedge \alpha' \rightarrow_{co} \alpha\})^+\end{aligned}$$

Well-formed execution

$$\text{Obs}(\text{WF_Exec}(P)) \in \{\text{Obs}(\text{Sync_Exec}(P))\}$$

Formalization

Reason axiomatically

$$E := \langle P, A, \rightarrow_{po}, \rightarrow_{co} \rangle$$

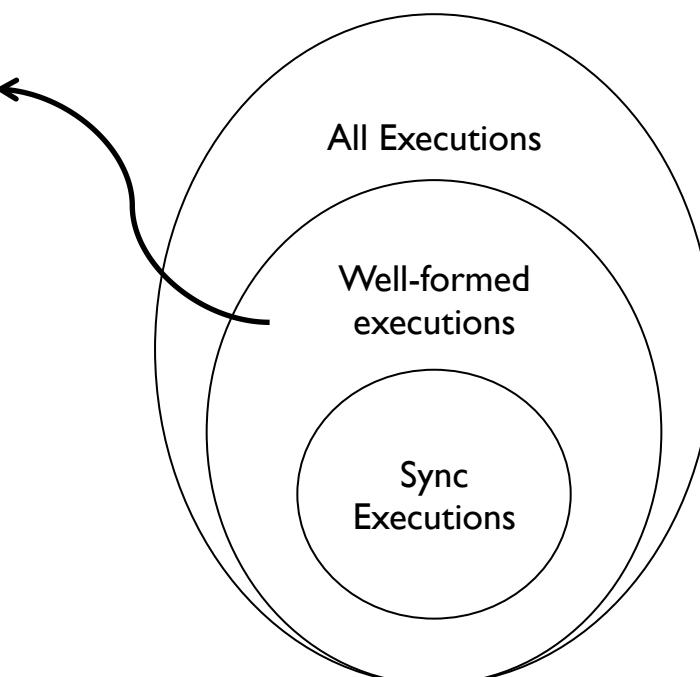
Happens-before relation

$$\rightarrow_{hb} = (\rightarrow_{po} \cup \rightarrow_{td} \cup \{(\alpha, \beta) \mid \alpha \rightarrow_{co} \alpha' \wedge \alpha' \rightarrow_{po} \beta\} \cup \{(\beta, \alpha) \mid \beta \rightarrow_{po} \alpha' \wedge \alpha' \rightarrow_{co} \alpha\})^+$$

Well-formed execution

$$\text{Obs}(\text{WF_Exec}(P)) \in \{\text{Obs}(\text{Sync_Exec}(P))\}$$

- No happens before cycle
- Sensible intra-thread semantics
- No outstanding speculative actions



Formalization

Reason axiomatically

$$E := \langle P, A, \rightarrow_{po}, \rightarrow_{co} \rangle$$

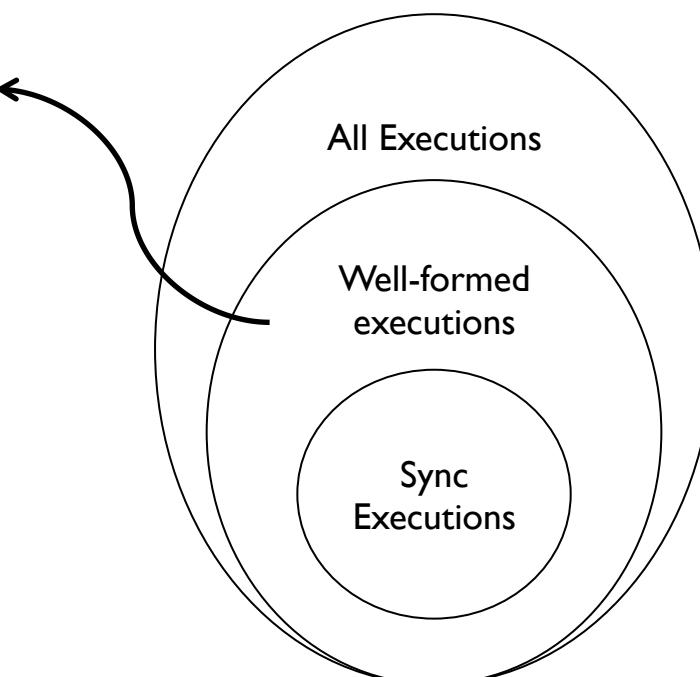
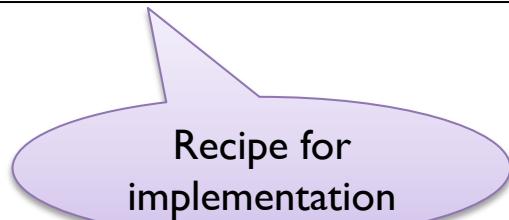
Happens-before relation

$$\rightarrow_{hb} = (\rightarrow_{po} \cup \rightarrow_{td} \cup \{(\alpha, \beta) \mid \alpha \rightarrow_{co} \alpha' \wedge \alpha' \rightarrow_{po} \beta\} \cup \{(\beta, \alpha) \mid \beta \rightarrow_{po} \alpha' \wedge \alpha' \rightarrow_{co} \alpha\})^+$$

Well-formed execution

$$\text{Obs}(\text{WF_Exec}(P)) \in \{\text{Obs}(\text{Sync_Exec}(P))\}$$

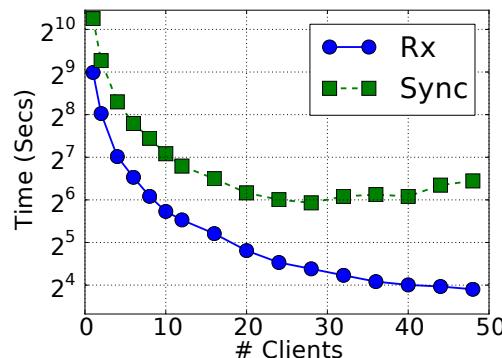
- No happens before cycle
- Sensible intra-thread semantics
- No outstanding speculative actions



Implementation & Results

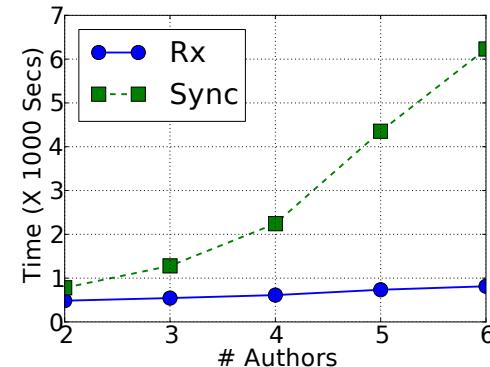
Implementation & Results

- Dependence graph \equiv Axiomatic execution
 - WF Check before observable actions
 - Ill-formed? Rollback and re-execute non-speculatively – Progress!
- WF Check, checkpoint, rollback are uncoordinated!
- Replicated channel consistency through speculative execution
- Benchmark: Optimistic OLTP & P2P Collaborative editing



OLTP

With 48 clients - **5.8X** faster than sync
1.4X slower than async



Collaborative Editing

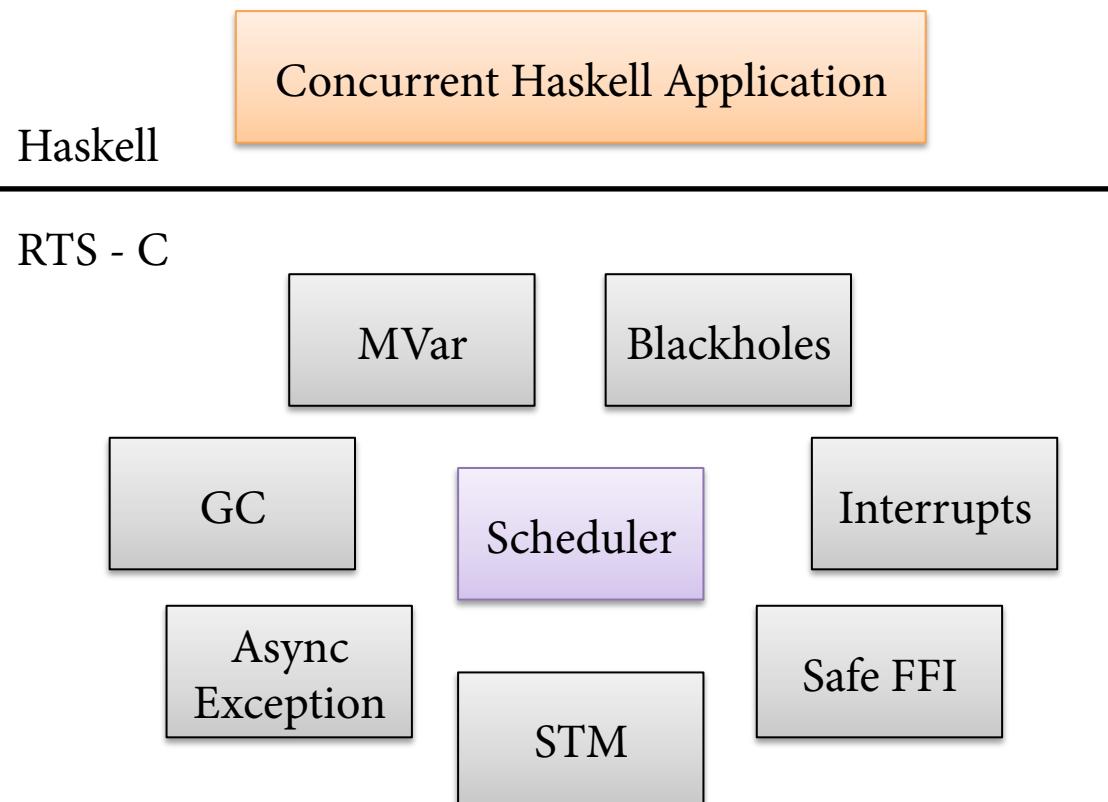
With 6 authors - **7.6X** faster than sync
2.3X slower than async

Scheduler Activations in Haskell

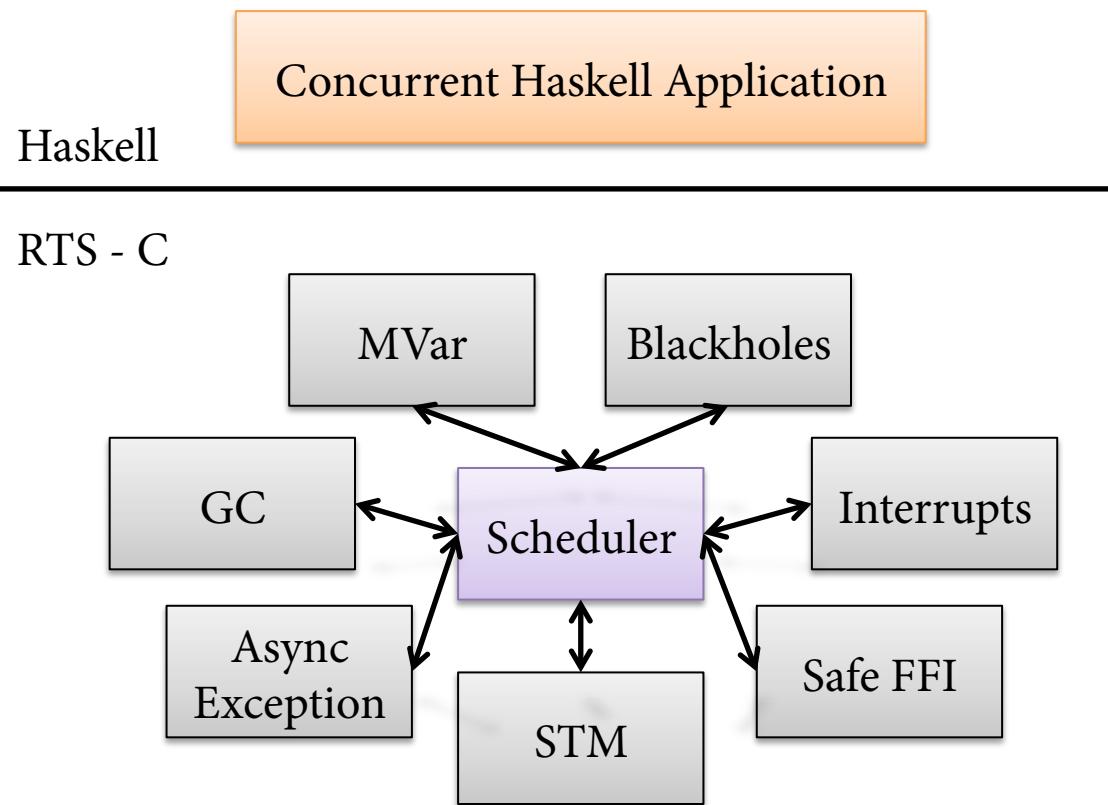
How to write schedulers for Haskell threads as Haskell libraries?

How to write schedulers for Haskell threads as Haskell libraries **in GHC?**

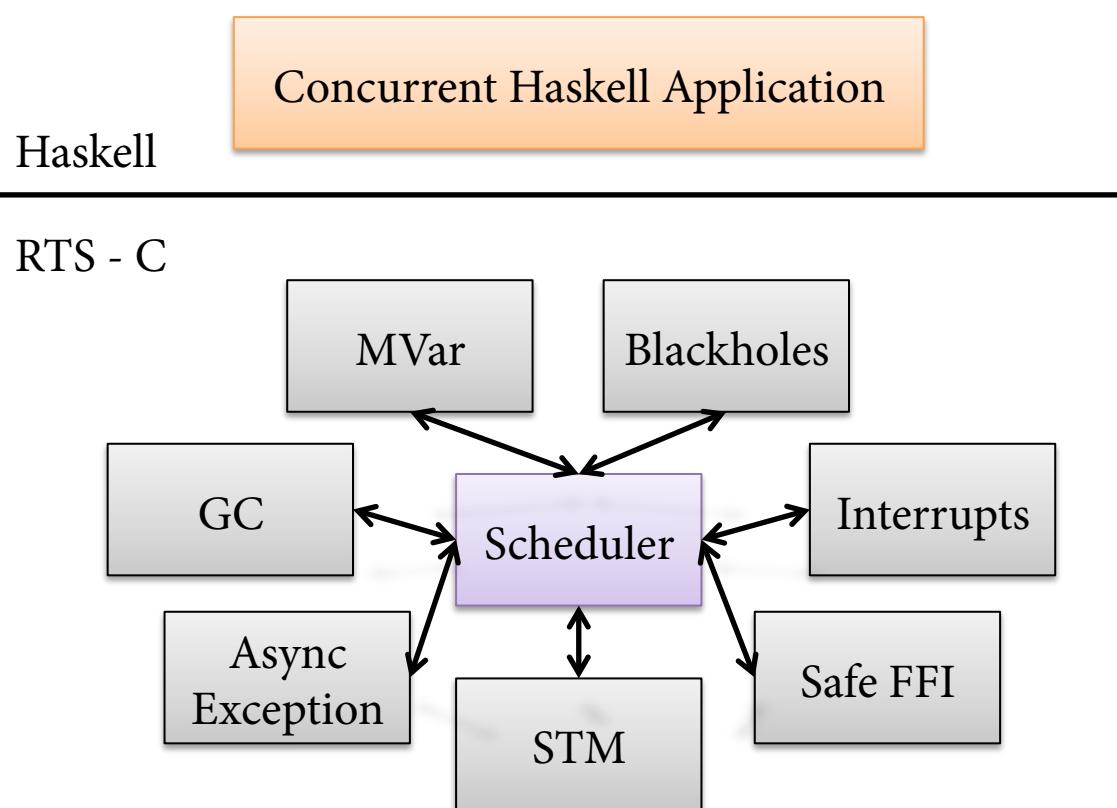
How to write schedulers for Haskell threads as Haskell libraries **in GHC?**



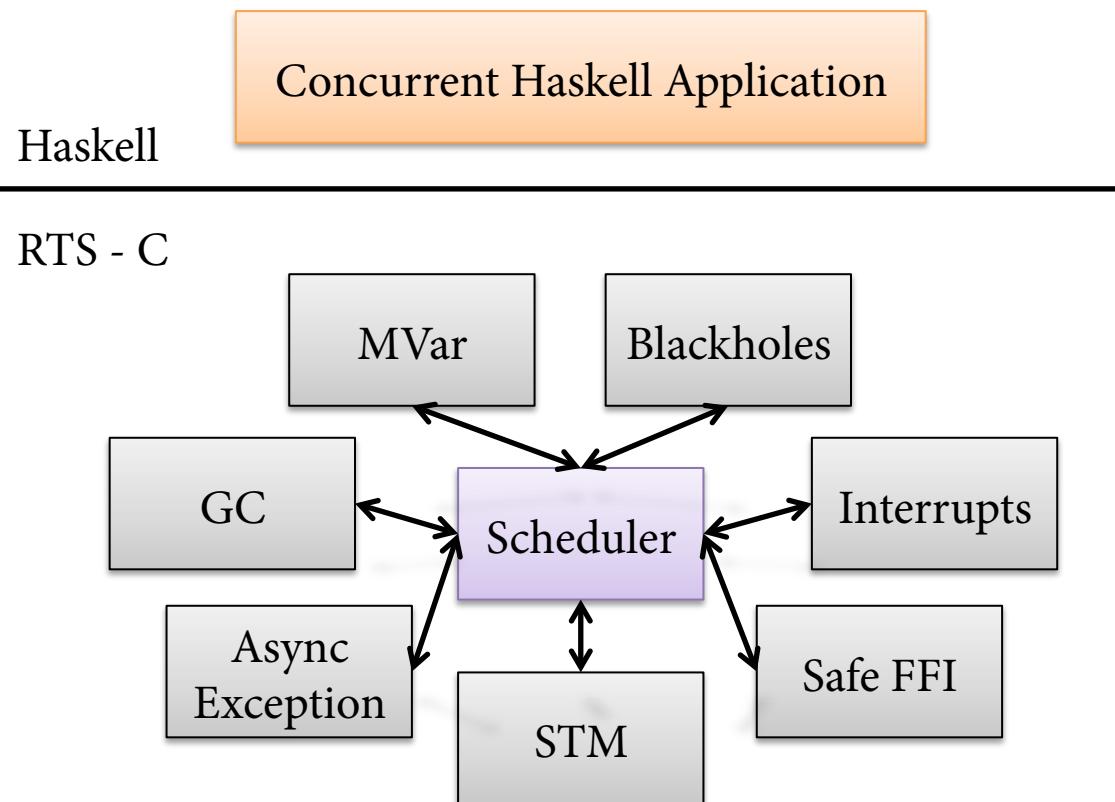
How to write schedulers for Haskell threads as Haskell libraries **in GHC?**



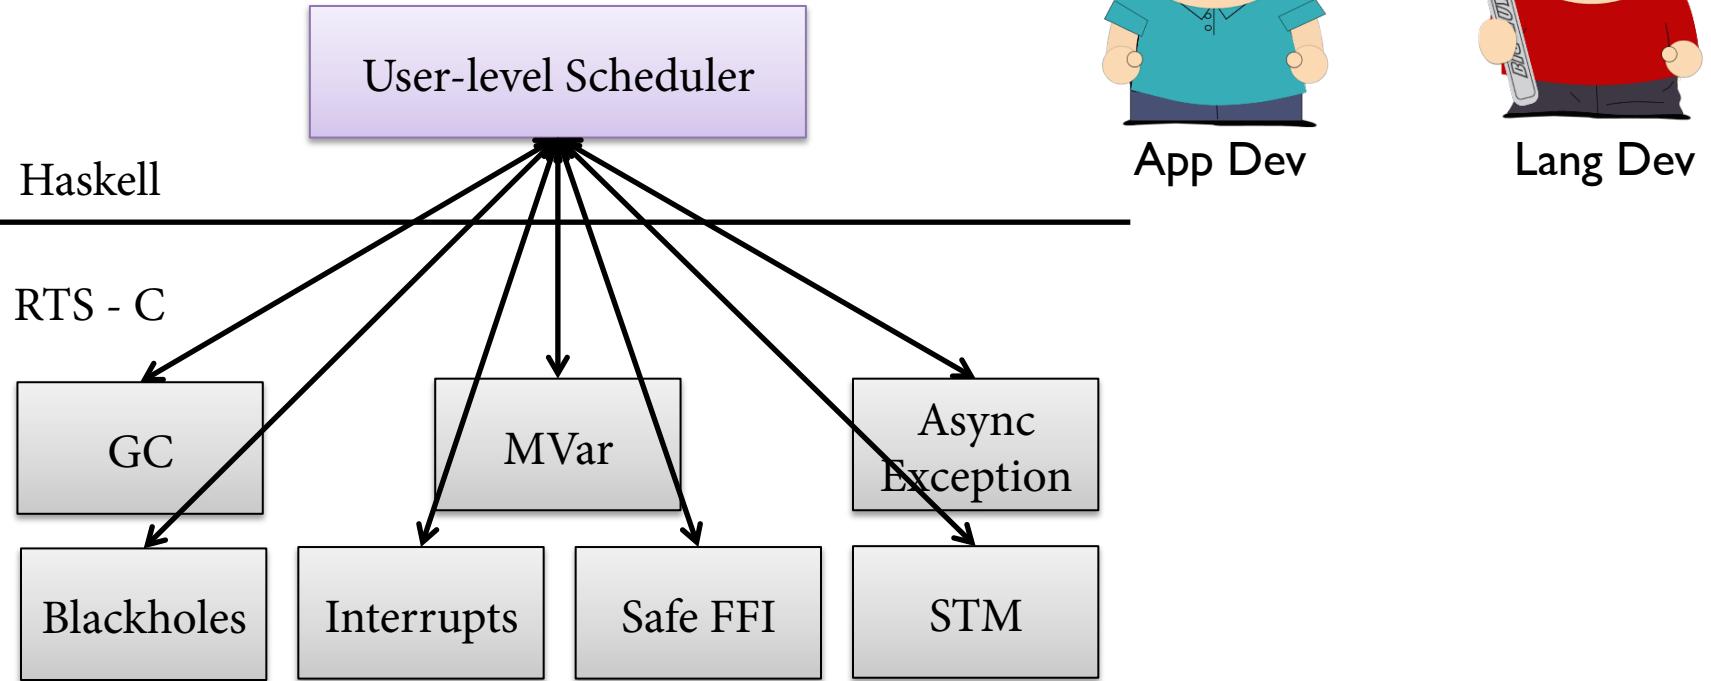
How to write schedulers for Haskell threads as Haskell libraries **in GHC?**



How to write schedulers for Haskell threads as Haskell libraries **in GHC?**



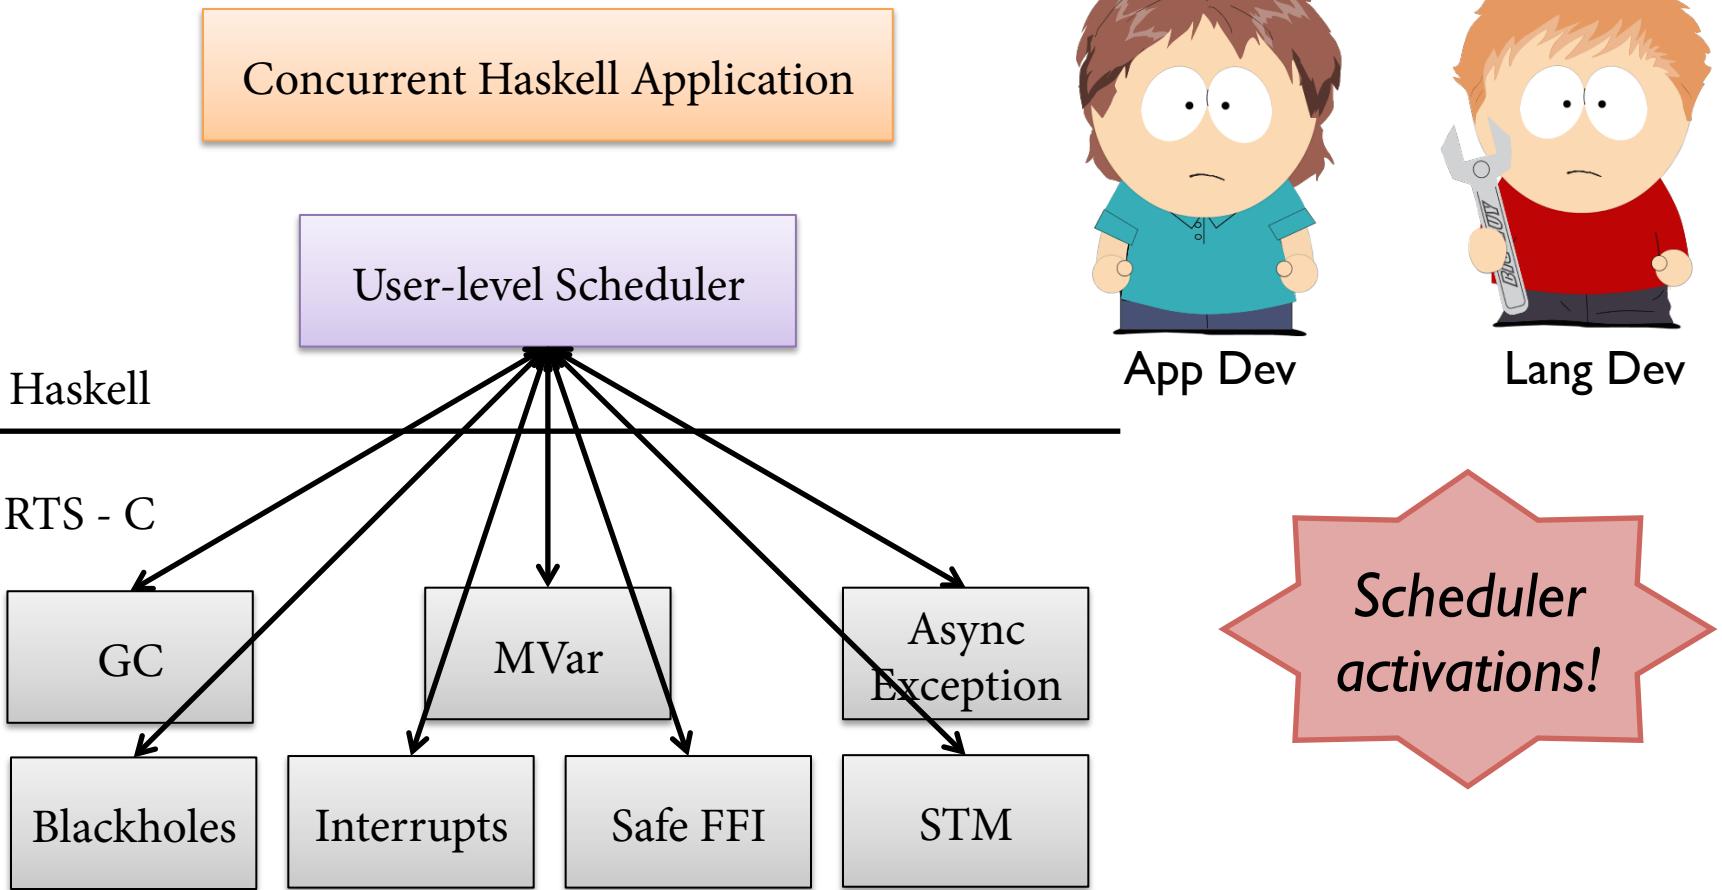
Concurrent Haskell Application



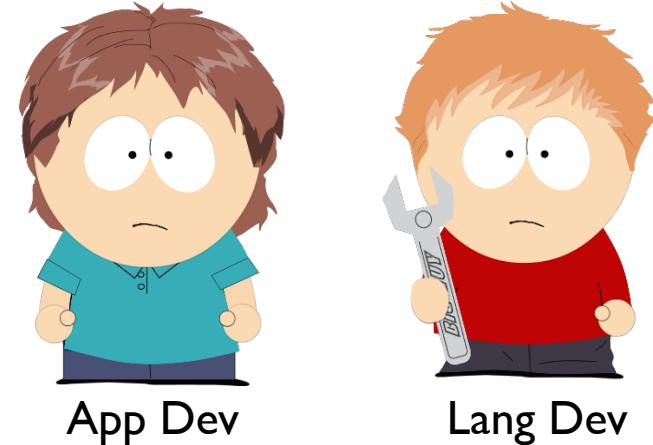
App Dev



Lang Dev



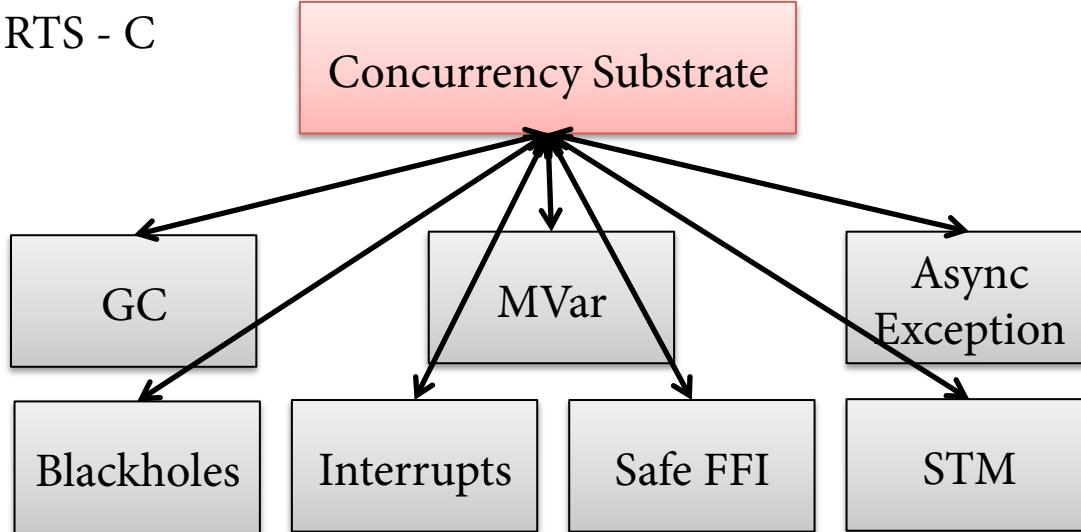
Concurrent Haskell Application



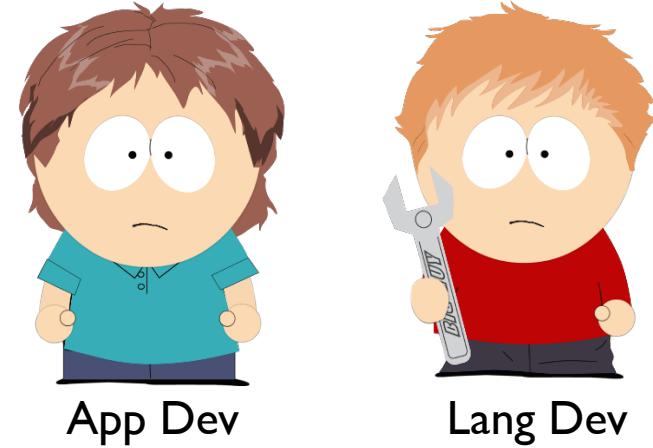
User-level Scheduler

Haskell

RTS - C



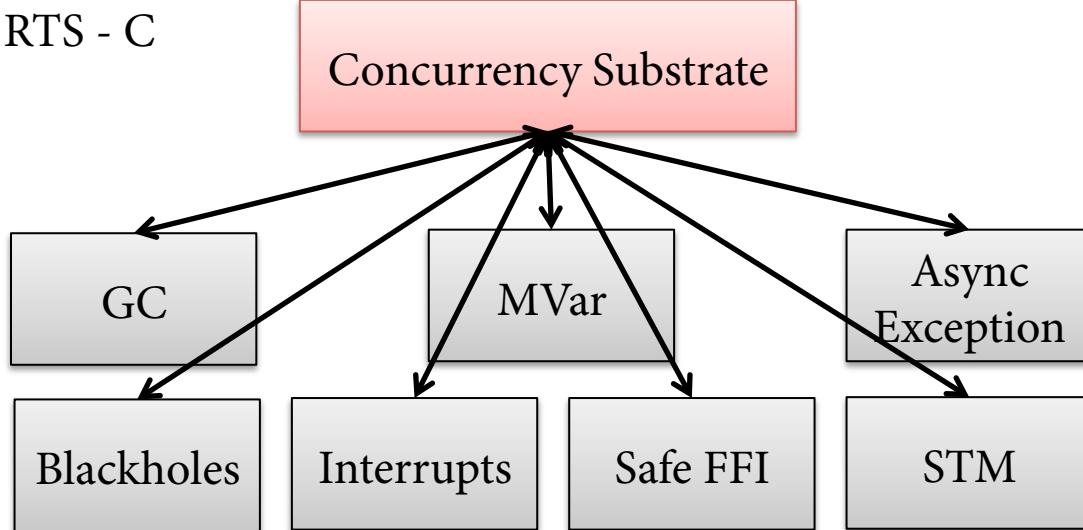
Concurrent Haskell Application



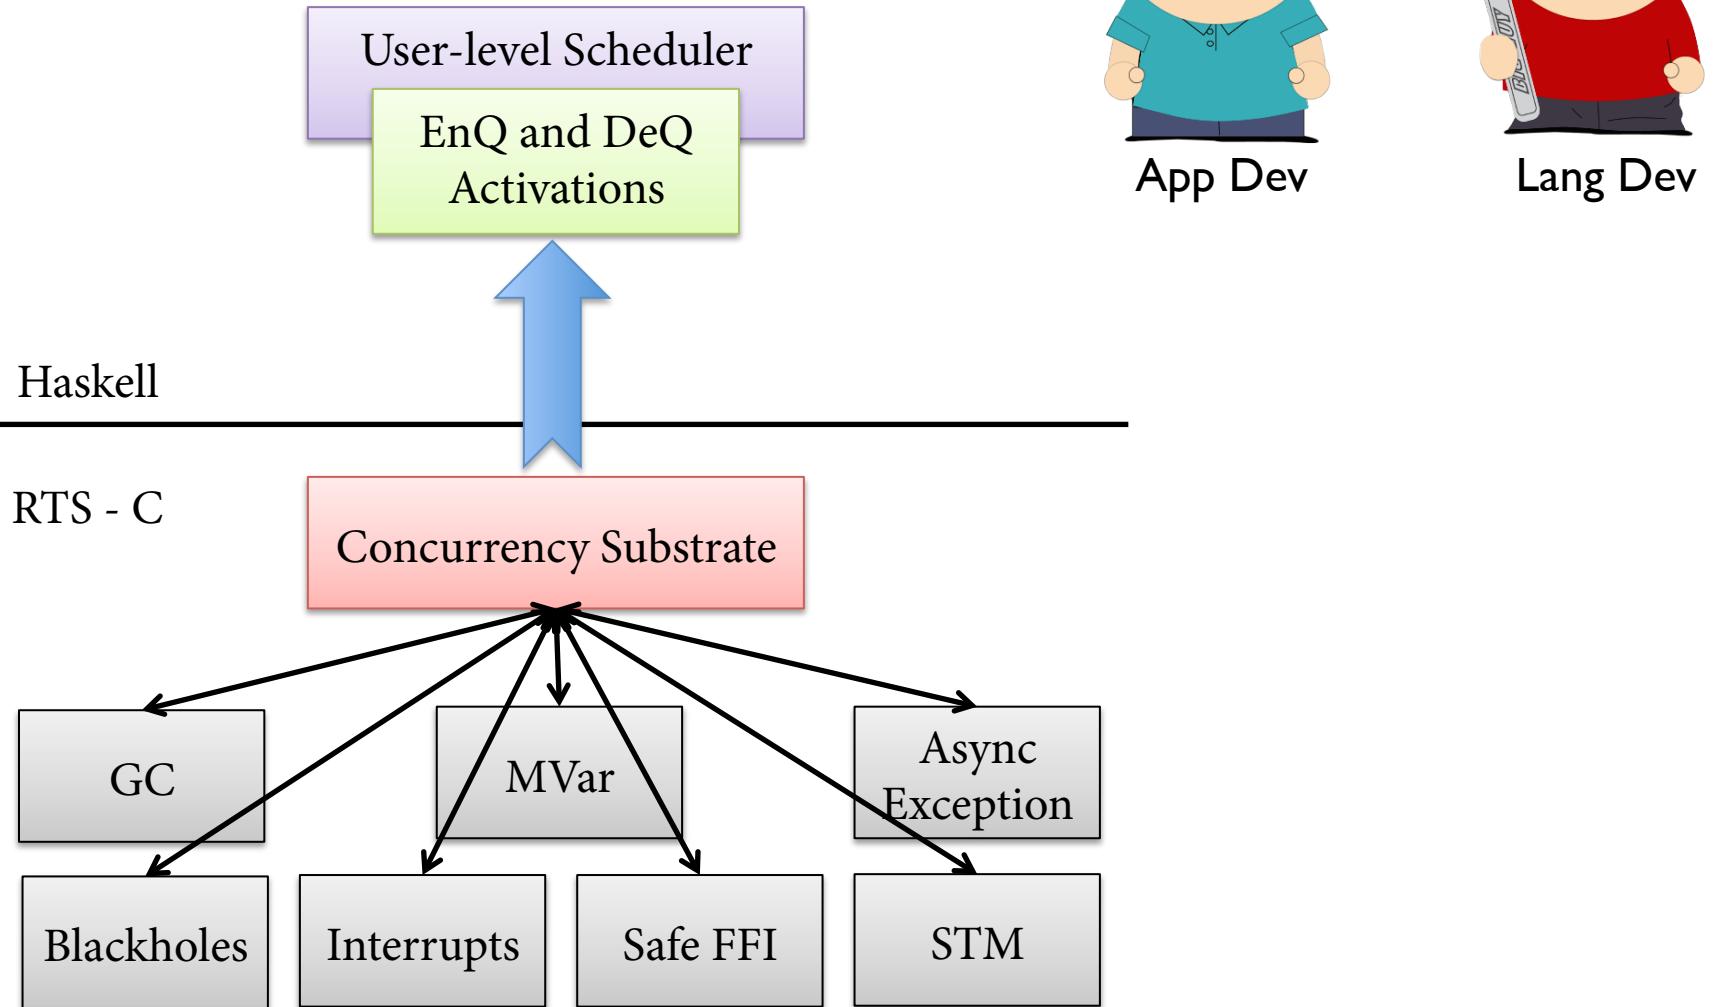
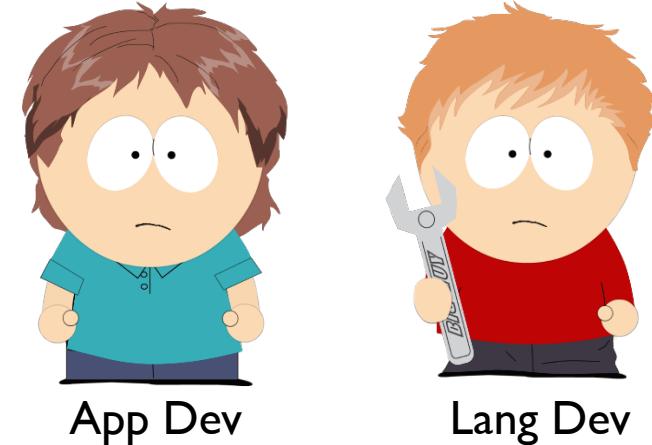
User-level Scheduler
EnQ and DeQ
Activations

Haskell

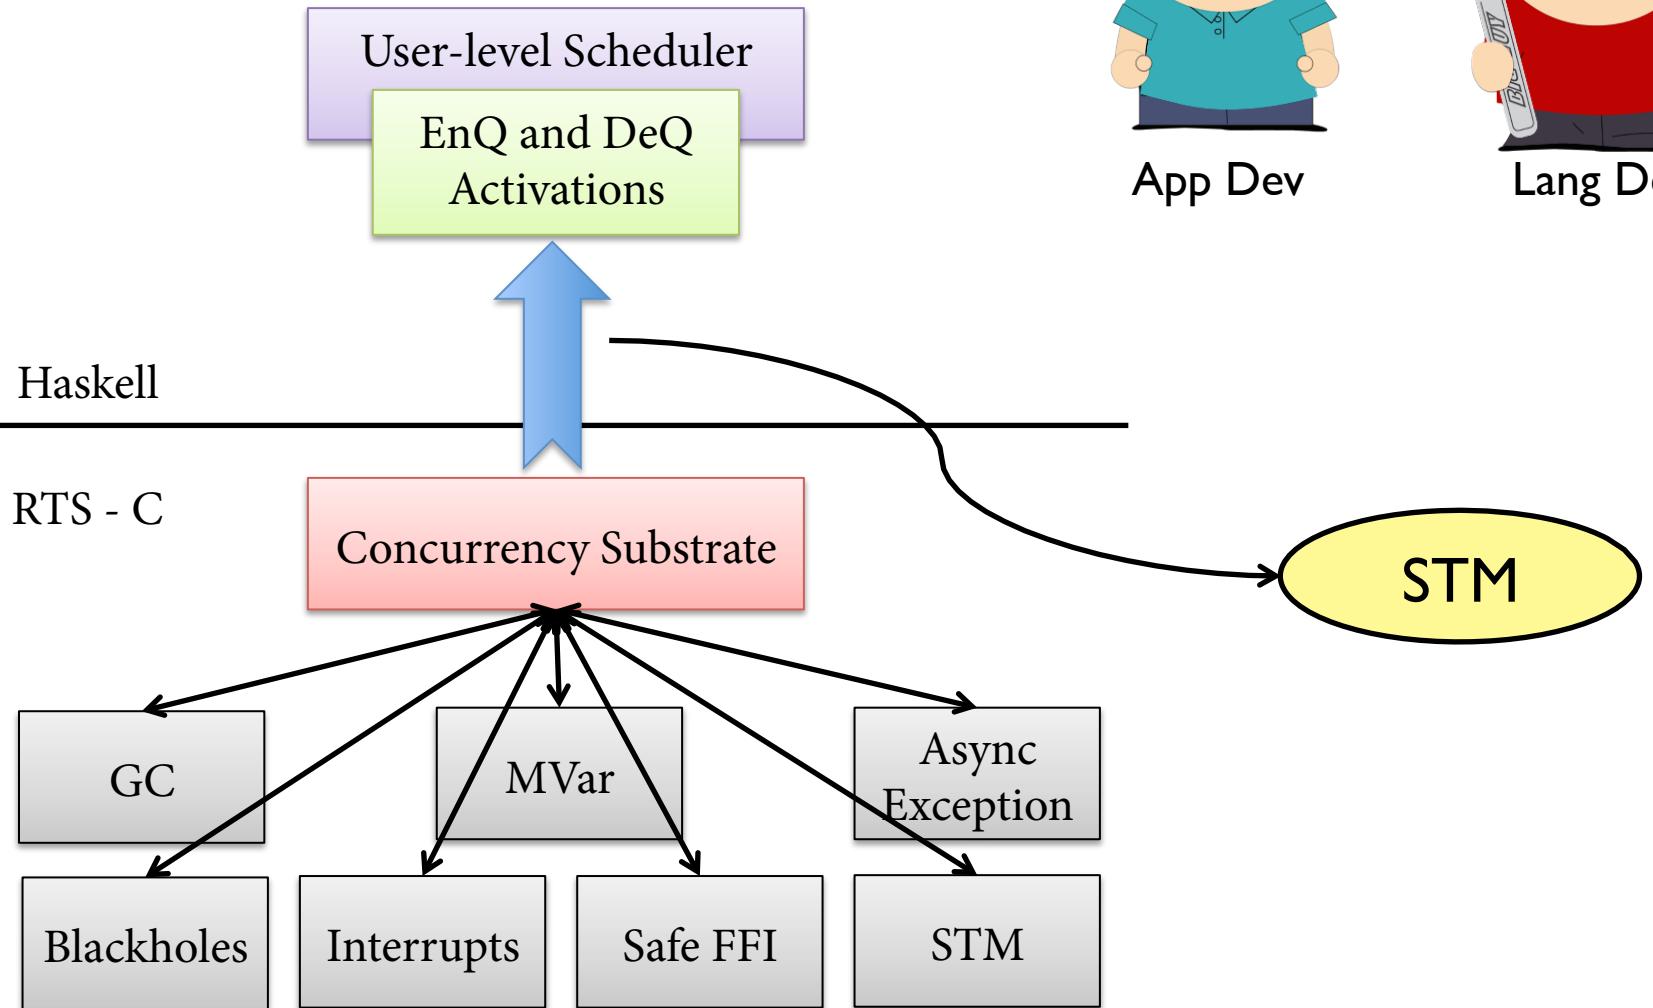
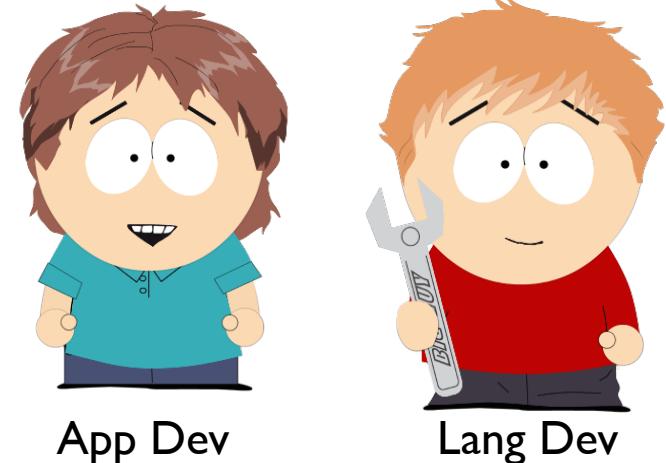
RTS - C



Concurrent Haskell Application



Concurrent Haskell Application



Concurrency Substrate API

```
data SCont
newSCont      :: IO () -> IO SCont
switch        :: (SCont -> STM SCont) -> IO ()
runOnIdleHEC :: SCont -> IO ()
```

Concurrency Substrate API

```
data SCont
newSCont      :: IO () -> IO SCont
switch        :: (SCont -> STM SCont) -> IO ()
runOnIdleHEC :: SCont -> IO ()

type DequeueAct = SCont -> STM SCont
type EnqueueAct = SCont -> STM ()

-- read activations
dequeueAct    :: DequeueAct
enqueueAct    :: EnqueueAct
-- update activations
setDequeueAct :: DequeueAct -> IO ()
setEnqueueAct :: EnqueueAct -> IO ()
```

Concurrency Substrate API

```
data SCont
newSCont      :: IO () -> IO SCont
switch        :: (SCont -> STM SCont) -> IO ()
runOnIdleHEC :: SCont -> IO ()
```

```
type DequeueAct = SCont -> STM SCont
type EnqueueAct = SCont -> STM ()
```

Scheduler access
is under STM

```
-- read activations
dequeueAct     :: DequeueAct
enqueueAct     :: EnqueueAct
-- update activations
setDequeueAct :: DequeueAct -> IO ()
setEnqueueAct :: EnqueueAct -> IO ()
```

Concurrency Substrate API

```
data SCont
newSCont      :: IO () -> IO SCont
switch        :: (SCont -> STM SCont) -> IO ()
runOnIdleHEC :: SCont -> IO ()

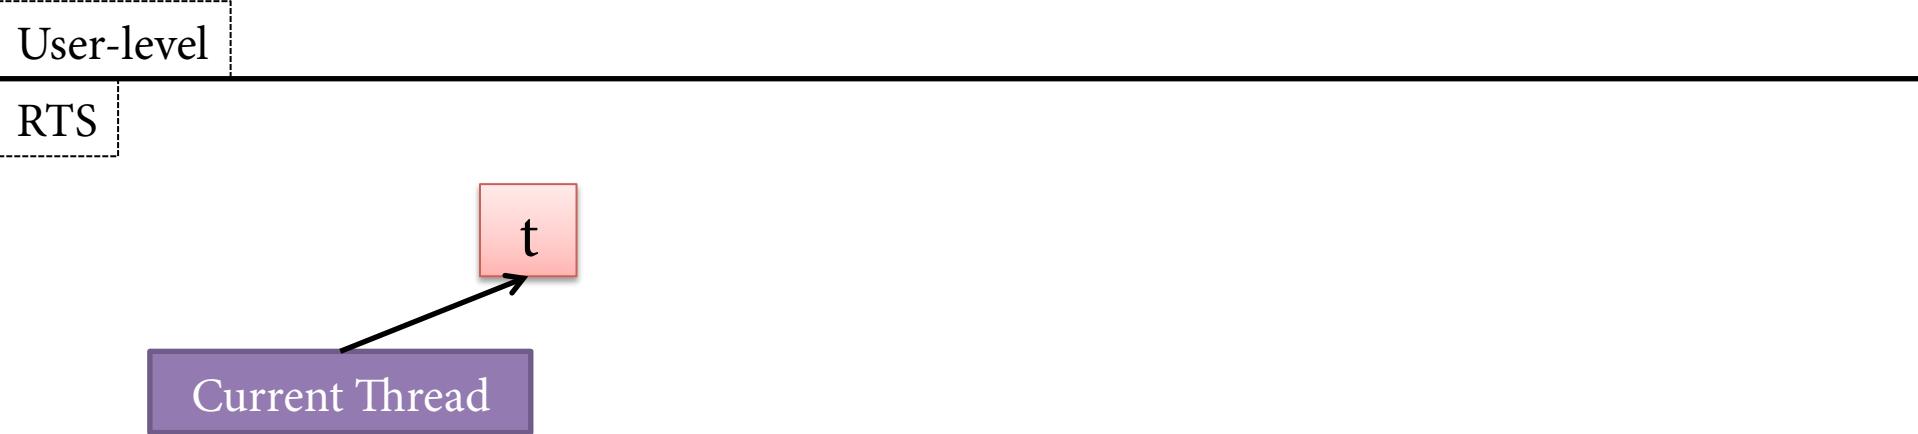
type DequeueAct = SCont -> STM SCont
type EnqueueAct = SCont -> STM ()
```

Scheduler access
is under STM

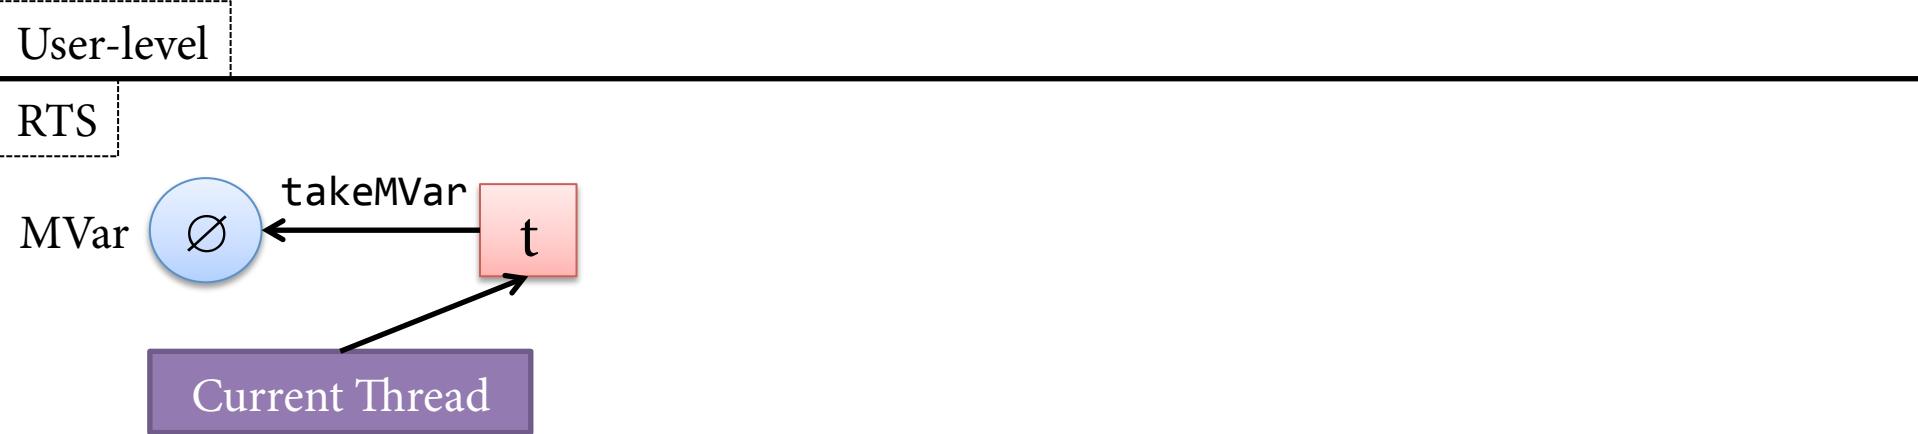
```
-- read activations
dequeueAct    :: DequeueAct
enqueueAct    :: EnqueueAct
-- update activations
setDequeueAct :: DequeueAct -> IO ()
setEnqueueAct :: EnqueueAct -> IO ()

-- update activations
getAux         :: SCont -> STM Dynamic
setAux         :: SCont -> Dynamic -> STM ()
```

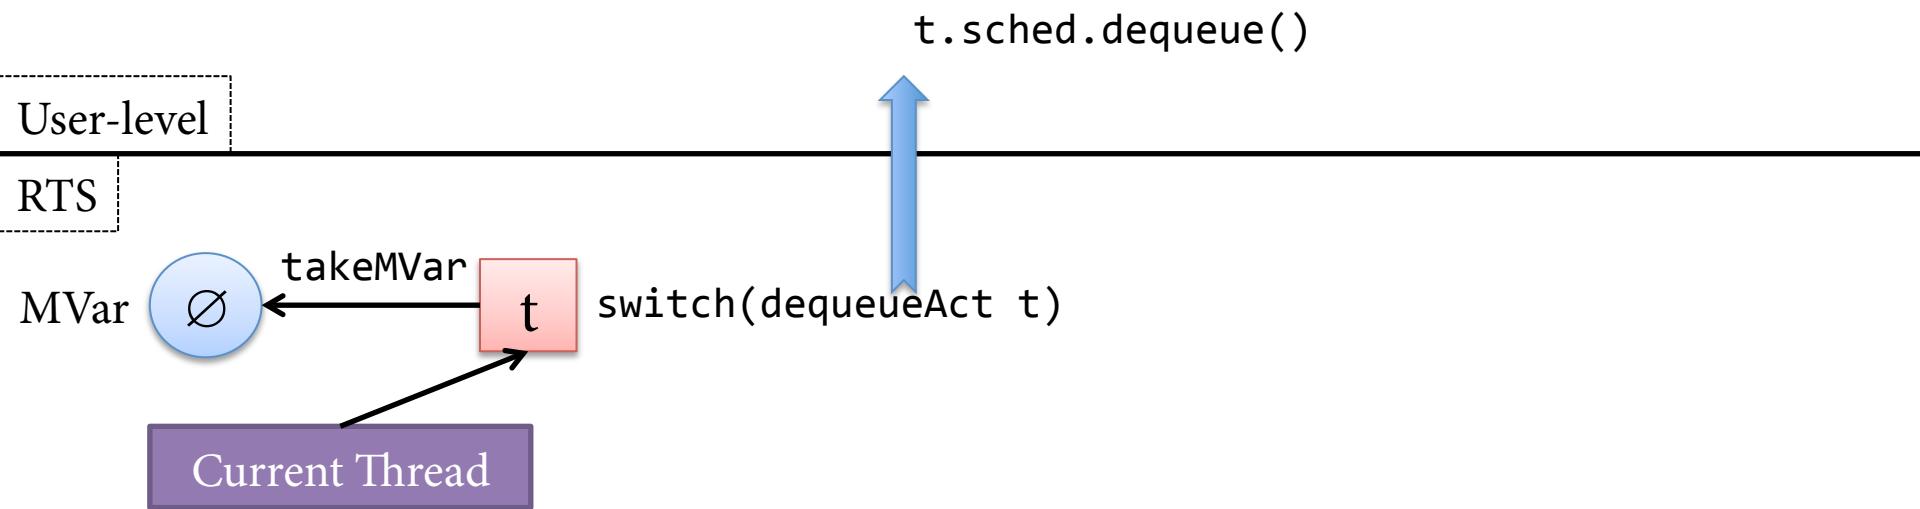
RTS Interaction - Blocking



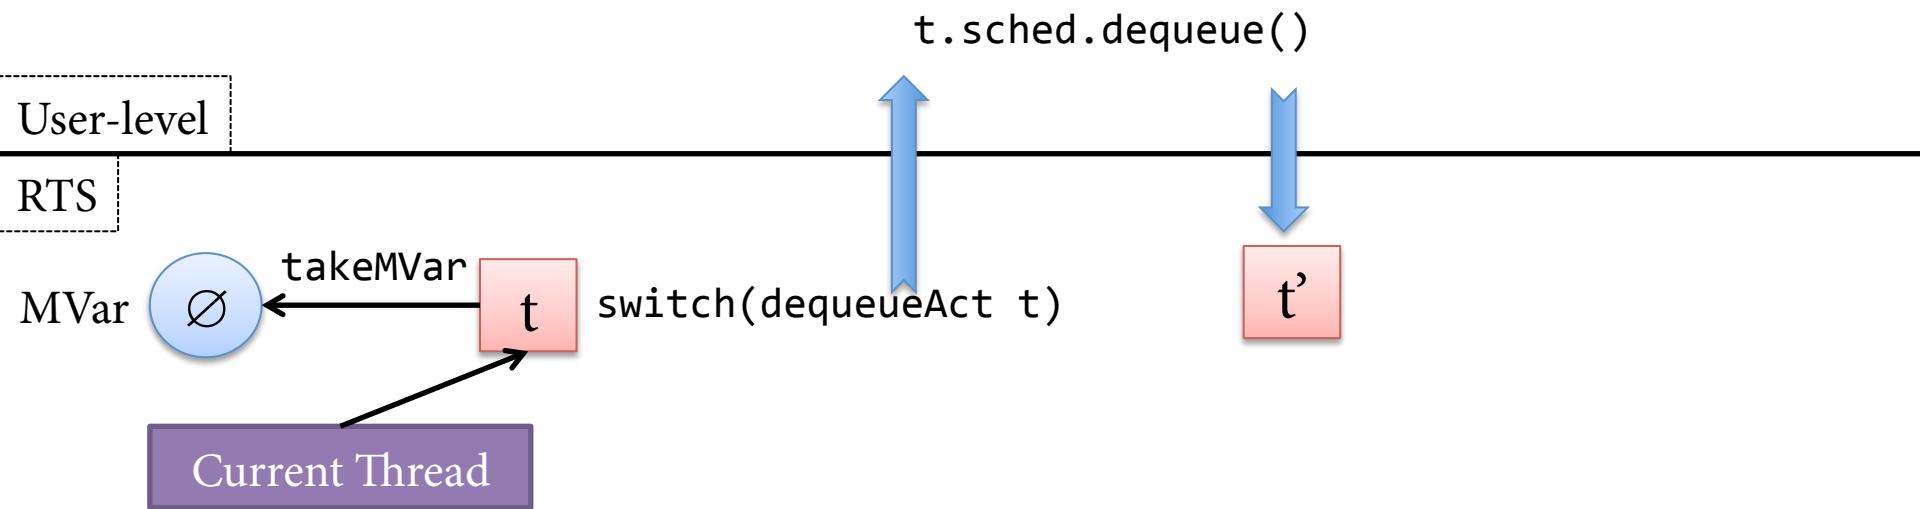
RTS Interaction - Blocking



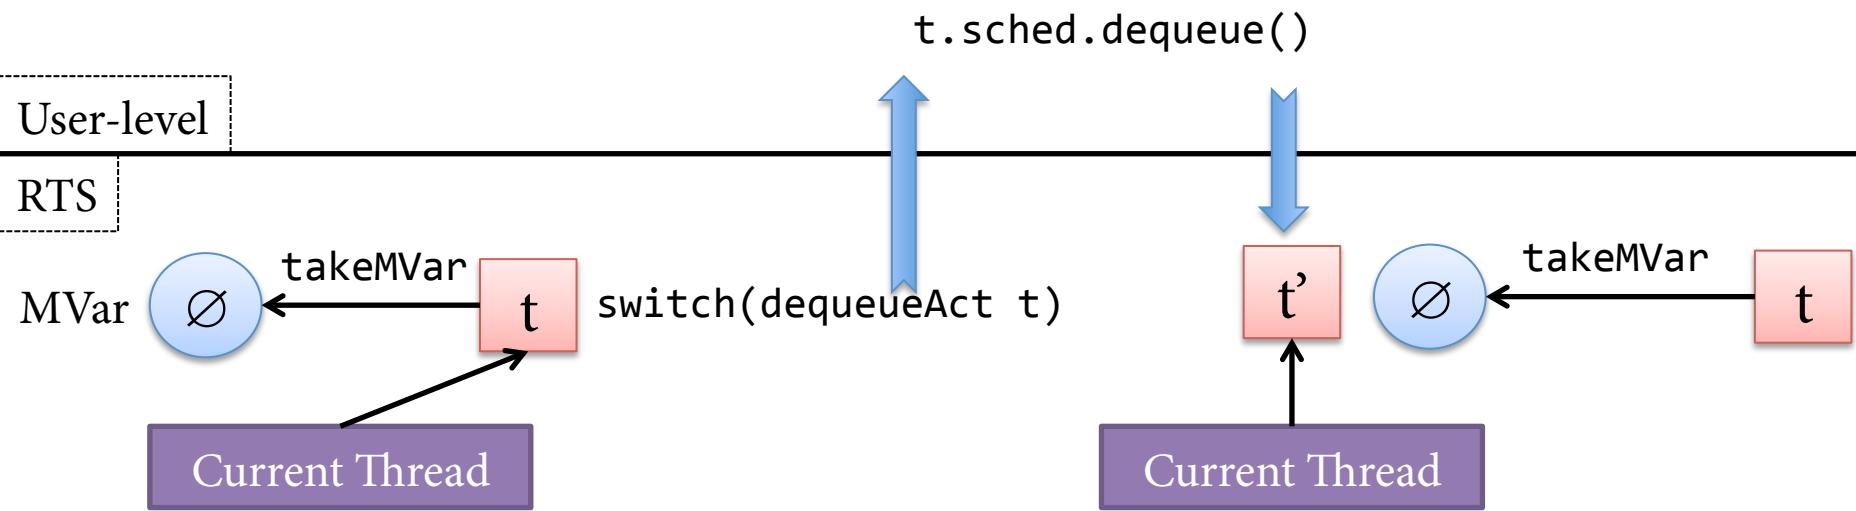
RTS Interaction - Blocking



RTS Interaction - Blocking



RTS Interaction - Blocking



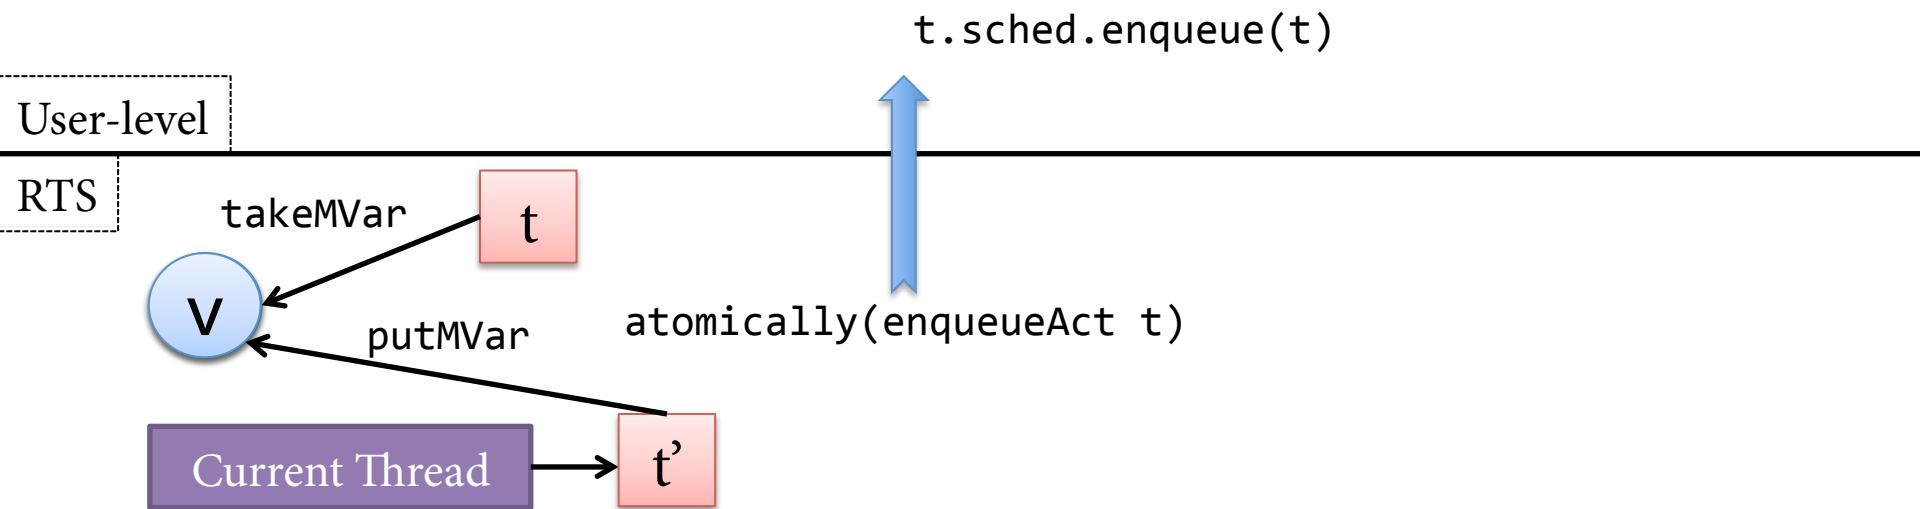
RTS Interaction - Unblocking



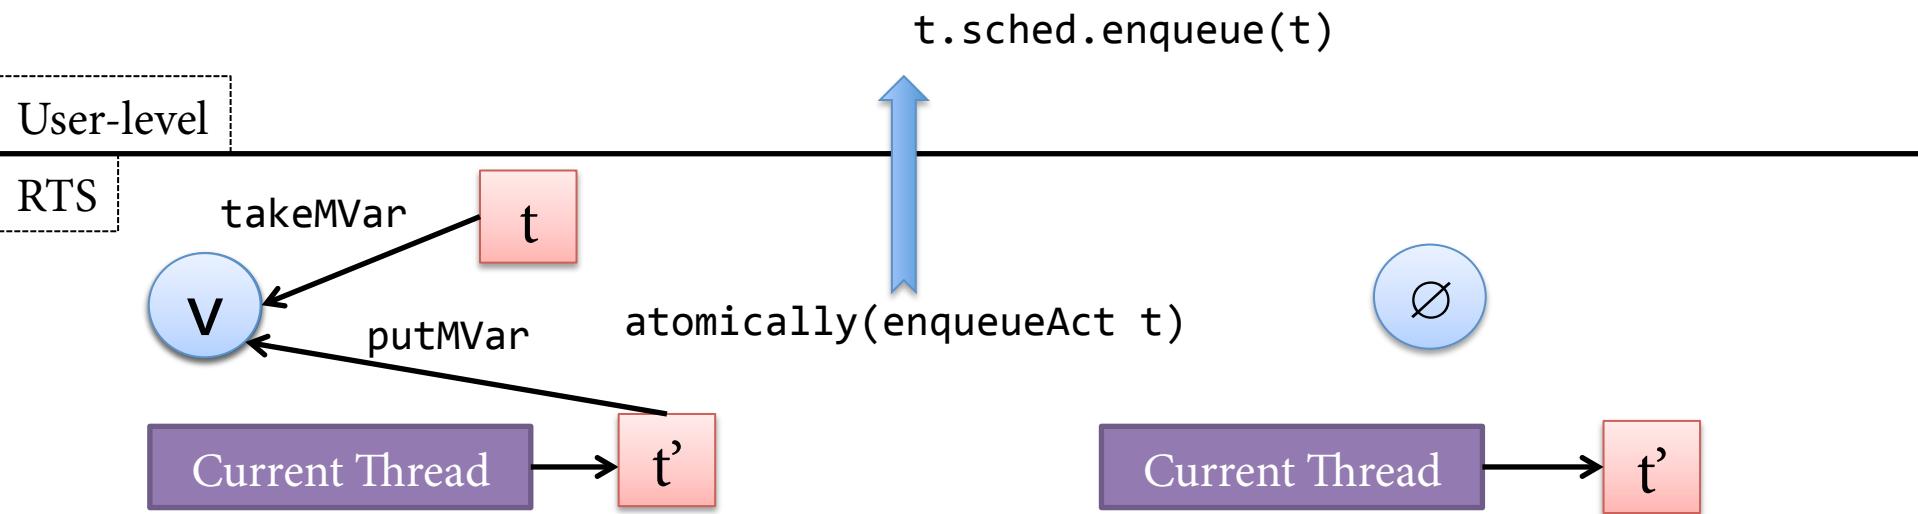
RTS Interaction - Unblocking



RTS Interaction - Unblocking



RTS Interaction - Unblocking



Multicore capable, preemptive, round-robin work-sharing scheduler

```

newtype Sched = Sched (Array Int (TVar[SCont]))

dequeueActivation:: Sched -> SCont -> STM SCont
dequeueActivation (Sched pa) _ = do
  cc <- getCurrentHEC -- get current HEC number
  l <- readTVar $ pa!cc
  case l of
    [] -> retry
    x:tl -> do
      writeTVar (pa!cc) tl
      return x

enqueueActivation:: Sched -> SCont -> STM ()
enqueueActivation (Sched pa) sc = do
  dyn <- getAux sc
  let (hec:Int, _::TVar Int) = fromJust $
    fromDynamic dyn
  l <- readTVar $ pa!hec
  writeTVar (pa!hec) $ l++[sc]

```

```

newScheduler :: IO ()
newScheduler = do
  -- Initialise Auxiliary state
  switch $ `s` -> do
    counter <- newTVar (0:Int)
    setAux s $ toDyn $ (0:Int,counter)
    return s
  -- Allocate scheduler
  nc <-getNumHECs
  sched <- (Sched . listArray (0,nc-1)) <$>
    replicateM n (newTVar [])
  -- Initialise activations
  setDequeueAct s $ dequeueActivation sched
  setEnqueueAct s $ enqueueActivation sched

newHEC :: IO ()
newHEC = do
  -- Initial task
  s <- newSCont $ switch dequeueAct
  -- Run in parallel
  runOnIdleHEC s

```

```

forkIO :: IO () -> IO SCont
forkIO task = do
  numHECs <- getNumHECs
  -- epilogue: Switch to next thread
  newSC <- newSCont (task >> switch dequeueAct)
  -- Create and initialise new Aux state
  switch $ `s` -> do
    dyn <- getAux s
    let (_::Int, t::TVar Int) = fromJust $
      fromDynamic dyn
    nextHEC <- readTVar t
    writeTVar t $ (nextHEC + 1) `mod` numHECs
    setAux newSC $ toDyn (nextHEC, t)
    return s
  -- Add new thread to scheduler
  atomically $ enqueueAct newSC
  return newSC

yield :: IO ()
yield = switch (\s -> enqueueAct s >> dequeueAct s)

```

Multicore capable, preemptive, round-robin work-sharing scheduler

```

newtype Sched = Sched (Array Int (TVar[SCont]))
dequeueActivation:: Sched -> SCont -> STM SCont
dequeueActivation (Sched pa) _ = do
  cc <- getCurrentHEC -- get current HEC number
  l <- readTVar $ pa!cc
  case l of
    [] -> retry
    x:tl -> do
      writeTVar (pa!cc) tl
      return x
enqueueActivation:: Sched -> SCont -> STM ()
enqueueActivation (Sched pa) sc = do
  dyn <- getAux sc
  let (hec:Int, _::TVar Int) = fromJust $ fromDynamic dyn
  l <- readTVar $ pa!hec
  writeTVar (pa!hec) $ l++[sc]

```

```

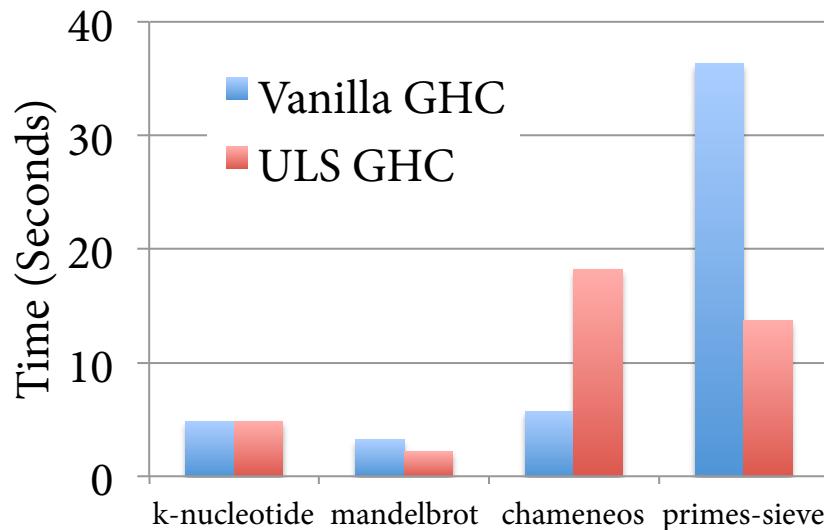
newScheduler :: IO ()
newScheduler = do
  -- Initialise Auxiliary state
  switch $ `s` -> do
    counter <- newTVar (0:Int)
    setAux s $ toDyn $ (0:Int,counter)
    return s
  -- Allocate scheduler
  nc <-getNumHECs
  sched <- (Sched . listArray (0,nc-1)) <$>
    replicateM n (newTVar [])
  -- Initialise activations
  setDequeueAct s $ dequeueActivation sched
  setEnqueueAct s $ enqueueActivation sched
  newHEC :: IO ()
  newHEC = do
    -- Initial task
    s <- newSCont $ switch dequeueAct
    -- Run in parallel
    runOnIdleHEC s

```

```

forkIO :: IO () -> IO SCont
forkIO task = do
  numHECs <- getNumHECs
  -- epilogue: Switch to next thread
  newSC <- newSCont (task >> switch dequeueAct)
  -- Create and initialise new Aux state
  switch $ `s` -> do
    dyn <- getAux s
    let (_::Int, t::TVar Int) = fromJust $ fromDynamic dyn
    nextHEC <- readTVar t
    writeTVar t $ (nextHEC + 1) `mod` numHECs
    setAux newSC $ toDyn (nextHEC, t)
    return s
  -- Add new thread to scheduler
  atomically $ enqueueAct newSC
  return newSC
yield :: IO ()
yield = switch (\s -> enqueueAct s >> dequeueAct s)

```



Multicore capable, preemptive, round-robin work-sharing scheduler

```

newtype Sched = Sched (Array Int (TVar[SCont]))
dequeueActivation:: Sched -> SCont -> STM SCont
dequeueActivation (Sched pa) _ = do
  cc <- getCurrentHEC -- get current HEC number
  l <- readTVar $ pa!cc
  case l of
    [] -> retry
    x:tl -> do
      writeTVar (pa!cc) tl
      return x
enqueueActivation:: Sched -> SCont -> STM ()
enqueueActivation (Sched pa) sc = do
  dyn <- getAux sc
  let (hec:Int, _::TVar Int) = fromJust $ fromDynamic dyn
  l <- readTVar $ pa!hec
  writeTVar (pa!hec) $ l++[sc]

```

```

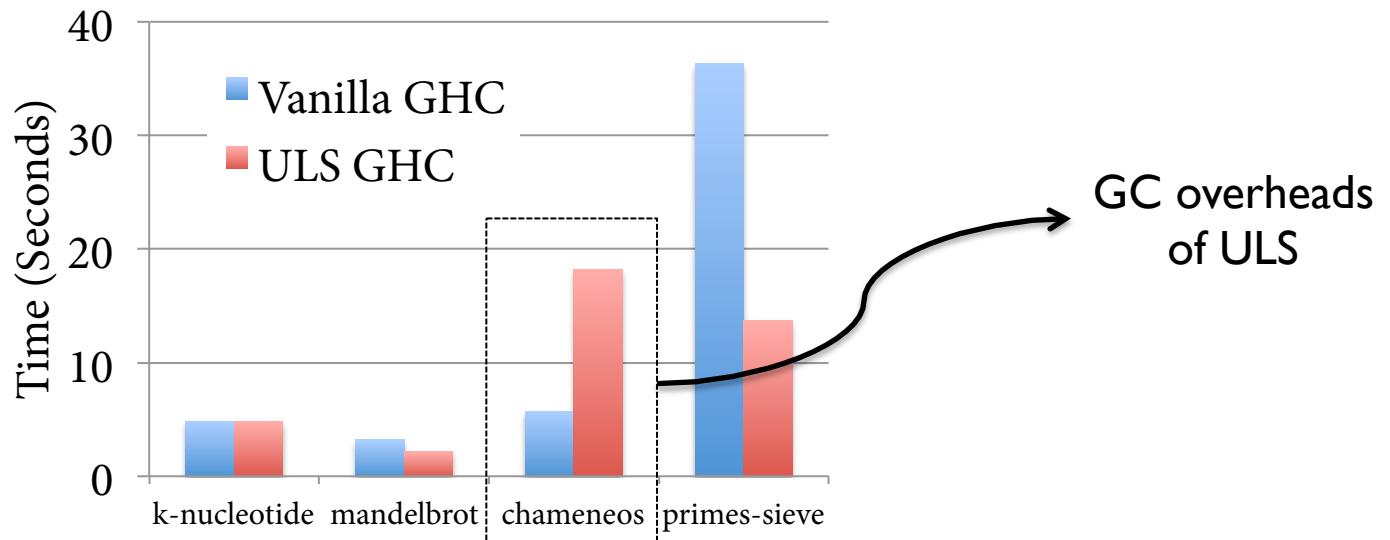
newScheduler :: IO ()
newScheduler = do
  -- Initialise Auxiliary state
  switch $ `s` -> do
    counter <- newTVar (0:Int)
    setAux s $ toDyn $ (0:Int,counter)
    return s
  -- Allocate scheduler
  nc <-getNumHECs
  sched <- (Sched . listArray (0,nc-1)) <$>
    replicateM n (newTVar [])
  -- Initialise activations
  setDequeueAct s $ dequeueActivation sched
  setEnqueueAct s $ enqueueActivation sched
  newHEC :: IO ()
  newHEC = do
    -- Initial task
    s <- newSCont $ switch dequeueAct
    -- Run in parallel
    runOnIdleHEC s

```

```

forkIO :: IO () -> IO SCont
forkIO task = do
  numHECs <- getNumHECs
  -- epilogue: Switch to next thread
  newSC <- newSCont (task >> switch dequeueAct)
  -- Create and initialise new Aux state
  switch $ `s` -> do
    dyn <- getAux s
    let (_::Int, t::TVar Int) = fromJust $ fromDynamic dyn
    nextHEC <- readTVar t
    writeTVar t $ (nextHEC + 1) `mod` numHECs
    setAux newSC $ toDyn (nextHEC, t)
    return s
  -- Add new thread to scheduler
  atomically $ enqueueAct newSC
  return newSC
yield :: IO ()
yield = switch (\s -> enqueueAct s >> dequeueAct s)

```



Multicore capable, preemptive, round-robin work-sharing scheduler

```

newtype Sched = Sched (Array Int (TVar[SCont]))
dequeueActivation:: Sched -> SCont -> STM SCont
dequeueActivation (Sched pa) _ = do
  cc <- getCurrentHEC -- get current HEC number
  l <- readTVar $ pa!cc
  case l of
    [] -> retry
    x:tl -> do
      writeTVar (pa!cc) tl
      return x
enqueueActivation:: Sched -> SCont -> STM ()
enqueueActivation (Sched pa) sc = do
  dyn <- getAux sc
  let (hec:Int, _::TVar Int) = fromJust $ fromDynamic dyn
  l <- readTVar $ pa!hec
  writeTVar (pa!hec) $ l++[sc]

```

```

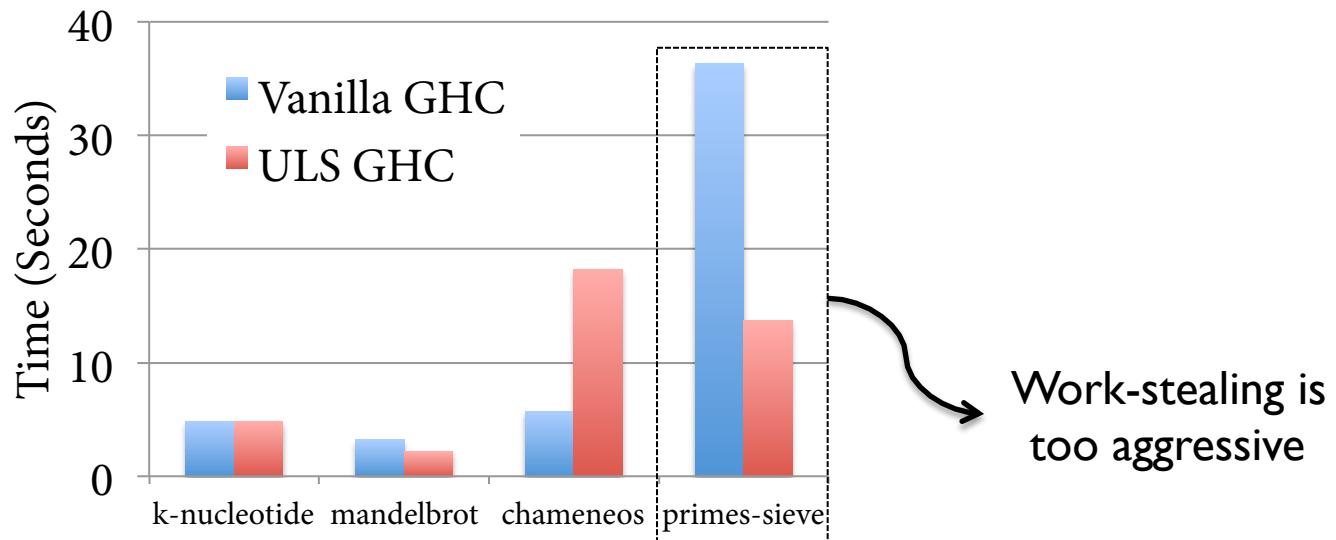
newScheduler :: IO ()
newScheduler = do
  -- Initialise Auxiliary state
  switch $ `s` -> do
    counter <- newTVar (0:Int)
    setAux s $ toDyn $ (0:Int,counter)
    return s
  -- Allocate scheduler
  nc <-getNumHECs
  sched <- (Sched . listArray (0,nc-1)) <$>
    replicateM n (newTVar [])
  -- Initialise activations
  setDequeueAct s $ dequeueActivation sched
  setEnqueueAct s $ enqueueActivation sched
  newHEC :: IO ()
  newHEC = do
    -- Initial task
    s <- newSCont $ switch dequeueAct
    -- Run in parallel
    runOnIdleHEC s

```

```

forkIO :: IO () -> IO SCont
forkIO task = do
  numHECs <- getNumHECs
  -- epilogue: Switch to next thread
  newSC <- newSCont (task >> switch dequeueAct)
  -- Create and initialise new Aux state
  switch $ `s` -> do
    dyn <- getAux s
    let (_::Int, t::TVar Int) = fromJust $ fromDynamic dyn
    nextHEC <- readTVar t
    writeTVar t $ (nextHEC + 1) `mod` numHECs
    setAux newSC $ toDyn (nextHEC, t)
    return s
  -- Add new thread to scheduler
  atomically $ enqueueAct newSC
  return newSC
yield :: IO ()
yield = switch (\s -> enqueueAct s >> dequeueAct s)

```



Formalization

Concurrency Substrate

	$x, y \in Variable$	$r, s, \in Name$
Md	$\text{return } M \mid M \gg= N$	
Ex	$\text{throw } M \mid \text{catch } M\ N \mid \text{catchSTM } M\ N$	
Stm	$\text{newTVar } M \mid \text{readTVar } r \mid \text{writeTVar } r\ M$	
	\mid	
	$\text{atomically } M \mid \text{retry}$	
Sc	$\text{newSCont } M \mid \text{switch } M \mid \text{runOnIdleHEC } s$	
Sls	$\text{getAux } s \mid \text{setAux } s\ M$	
Act	$\text{dequeueAct } s \mid \text{enqueueAct } s$	
	\mid	
	$\text{setDequeueAct } M \mid \text{setEnqueueAct } M$	
Term		
M, N	$\text{::= } r \mid x \mid \lambda.x \rightarrow M \mid M\ N \mid \dots$	
	\mid	
	$Md \mid Ex \mid Stm \mid Sc \mid Sls \mid Act$	
Program state	$P \text{ ::= } S; \Theta$	
HEC soup	$S \text{ ::= } \emptyset \mid H \parallel S$	
HEC	$H \text{ ::= } \langle s, M, D \rangle \mid \langle s, M, D \rangle_{Sleeping}$	
	$\mid \langle s, M, D \rangle_{Outcall} \mid \text{Idle}$	
Heap	$\Theta \text{ ::= } r \mapsto M \oplus s \mapsto (M, D)$	
SLS Store	$D \text{ ::= } (M, N, r)$	
IO Context	$\mathbb{E} \text{ ::= } \bullet \mid \mathbb{E} \gg= M \mid \text{catch } \mathbb{E}\ M$	
STM Context	$\mathbb{P} \text{ ::= } \bullet \mid \mathbb{P} \gg= M$	

Upcalls from the RTS

	$\text{Dequeue upcall instantiation } H; \Theta \xrightarrow{\text{deq}} H'; \Theta'$
(UPDEQUEUE)	
	$s' \text{ fresh } r \text{ fresh } D' = (\text{deq}(D), \text{enq}(D), r)$
	$M' = \text{switch } (\lambda x. \text{deq}(D) s)$
	$\Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]$
	$\underline{(s, M, D); \Theta \xrightarrow{\text{deq}} \langle s', M', D' \rangle; \Theta'}}$
	$\text{Enqueue upcall instantiation } H; \Theta \xrightarrow{\text{enq } s} H'; \Theta'$
(UPENQUEUEIDLE)	
	$s' \text{ fresh } r \text{ fresh } D' = (\text{deq}(D), \text{enq}(D), r)$
	$M' = \text{atomically } (\text{enq}(D) s)$
	$\Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]$
	$\underline{\text{Idle}; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M', D' \rangle; \Theta'}$
(UPENQUEUERUNNING)	
	$M'' = \text{atomically } (\text{enq}(D) s) \gg M'$
	$\underline{\langle s', M', D' \rangle; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M'', D' \rangle; \Theta[s \mapsto (M, D)]}$

Formalization

Concurrency Substrate

	$x, y \in Variable$	$r, s, \in Name$
Md	$\text{return } M \mid M \gg= N$	
Ex	$\text{throw } M \mid \text{catch } M\ N \mid \text{catchSTM } M\ N$	
Stm	$\text{newTVar } M \mid \text{readTVar } r \mid \text{writeTVar } r\ M$	
	\mid	
	$\text{atomically } M \mid \text{retry}$	
Sc	$\text{newSCont } M \mid \text{switch } M \mid \text{runOnIdleHEC } s$	
Sls	$\text{getAux } s \mid \text{setAux } s\ M$	
Act	$\text{dequeueAct } s \mid \text{enqueueAct } s$	
	\mid	
	$\text{setDequeueAct } M \mid \text{setEnqueueAct } M$	
Term		
M, N	$\text{ ::= } r \mid x \mid \lambda.x \rightarrow M \mid M\ N \mid \dots$	
	\mid	
	$Md \mid Ex \mid Stm \mid Sc \mid Sls \mid Act$	
Program state	$P \text{ ::= } S; \Theta$	
HEC soup	$S \text{ ::= } \emptyset \mid H \parallel S$	
HEC	$H \text{ ::= } \langle s, M, D \rangle \mid \langle s, M, D \rangle_{\text{Sleeping}}$	
	$\mid \langle s, M, D \rangle_{\text{Outcall}} \mid \text{Idle}$	
Heap	$\Theta \text{ ::= } r \mapsto M \oplus s \mapsto (M, D)$	
SLS Store	$D \text{ ::= } (M, N, r)$	
IO Context	$\mathbb{E} \text{ ::= } \bullet \mid \mathbb{E} \gg= M \mid \text{catch } \mathbb{E}\ M$	
STM Context	$\mathbb{P} \text{ ::= } \bullet \mid \mathbb{P} \gg= M$	

Upcalls from the RTS

	Dequeue upcall instantiation	$H; \Theta \xrightarrow{\text{deq}} H'; \Theta'$
(UPDEQUEUE)		
$s' \text{ fresh }$	$r \text{ fresh }$	$D' = (\text{deq}(D), \text{enq}(D), r)$
		$M' = \text{switch } (\lambda x. \text{deq}(D))\ s$
		$\Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]$
		$\underline{(s, M, D); \Theta \xrightarrow{\text{deq}} \langle s', M', D' \rangle; \Theta'}}$
Enqueue upcall instantiation		
		$H; \Theta \xrightarrow{\text{enq } s} H'; \Theta'$
(UPENQUEUEIDLE)		
$s' \text{ fresh }$	$r \text{ fresh }$	$D' = (\text{deq}(D), \text{enq}(D), r)$
		$M' = \text{atomically } (\text{enq}(D))\ s$
		$\Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]$
		$\underline{\text{Idle}; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M', D' \rangle; \Theta'}$
(UPENQUEUERUNNING)		
		$M'' = \text{atomically } (\text{enq}(D))\ s \gg M'$
		$\underline{\langle s', M', D' \rangle; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M'', D' \rangle; \Theta[s \mapsto (M, D)]}$



Formalization

Concurrency Substrate

$x, y \in Variable \quad r, s, \in Name$

$Md ::= \text{return } M \mid M \gg= N$
 $Ex ::= \text{throw } M \mid \text{catch } M \ N \mid \text{catchSTM } M \ N$
 $Stm ::= \text{newTVar } M \mid \text{readTVar } r \mid \text{writeTVar } r \ M$
 $\mid \text{atomically } M \mid \text{retry}$
 $Sc ::= \text{newSCont } M \mid \text{switch } M \mid \text{runOnIdleHEC } s$
 $Sls ::= \text{getAux } s \mid \text{setAux } s \ M$
 $Act ::= \text{dequeueAct } s \mid \text{enqueueAct } s$
 $\mid \text{setDequeueAct } M \mid \text{setEnqueueAct } M$

Term

$M, N ::= r \mid x \mid \lambda.x \rightarrow M \mid M \ N \mid \dots$
 $\mid Md \mid Ex \mid Stm \mid Sc \mid Sls \mid Act$

Program state

P	$::= S; \Theta$
HEC soup	$S ::= \emptyset \mid H \parallel S$
HEC	$H ::= \langle s, M, D \rangle \mid \langle s, M, D \rangle_{\text{Sleeping}}$ $\mid \langle s, M, D \rangle_{\text{Outcall}} \mid \text{Idle}$
Heap	$\Theta ::= r \mapsto M \oplus s \mapsto (M, D)$
SLS Store	$D ::= (M, N, r)$
IO Context	$\mathbb{E} ::= \bullet \mid \mathbb{E} \gg= M \mid \text{catch } \mathbb{E} \ M$
STM Context	$\mathbb{P} ::= \bullet \mid \mathbb{P} \gg= M$

Upcalls from the RTS

Dequeue upcall instantiation $H; \Theta \xrightarrow{\text{deq}} H'; \Theta'$

(UPDEQUEUE)

$$\frac{s' \text{ fresh } r \text{ fresh } D' = (\text{deq}(D), \text{enq}(D), r) \quad M' = \text{switch } (\lambda x. \text{deq}(D) \ s) \quad \Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]}{\langle s, M, D \rangle; \Theta \xrightarrow{\text{deq}} \langle s', M', D' \rangle; \Theta'}$$

Enqueue upcall instantiation $H; \Theta \xrightarrow{\text{enq } s} H'; \Theta'$

(UPENQUEUEIDLE)

$$\frac{s' \text{ fresh } r \text{ fresh } D' = (\text{deq}(D), \text{enq}(D), r) \quad M' = \text{atomically } (\text{enq}(D) \ s) \quad \Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()] \quad \text{Idle}; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M', D' \rangle; \Theta'}{\text{Idle}; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M', D' \rangle; \Theta'}$$

(UPENQUEUERUNNING)

$$\frac{M'' = \text{atomically } (\text{enq}(D) \ s) \gg M'}{\langle s', M', D' \rangle; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M'', D' \rangle; \Theta[s \mapsto (M, D)]}$$

IO

$E[M] \xrightarrow{\text{tick}} E[\text{yield} \gg M]$

atomically
STM

$E[\text{atomically}(P[M])] \xrightarrow{\text{tick}} \text{Suspend transaction and switch}$

Formalization

Concurrency Substrate

	$x, y \in Variable$	$r, s, \in Name$
Md	$\text{return } M \mid M \gg= N$	
Ex	$\text{throw } M \mid \text{catch } M\ N \mid \text{catchSTM } M\ N$	
Stm	$\text{newTVar } M \mid \text{readTVar } r \mid \text{writeTVar } r\ M$	
	\mid	
	$\text{atomically } M \mid \text{retry}$	
Sc	$\text{newSCont } M \mid \text{switch } M \mid \text{runOnIdleHEC } s$	
Sls	$\text{getAux } s \mid \text{setAux } s\ M$	
Act	$\text{dequeueAct } s \mid \text{enqueueAct } s$	
	\mid	
	$\text{setDequeueAct } M \mid \text{setEnqueueAct } M$	
Term		
M, N	$\text{ ::= } r \mid x \mid \lambda.x \rightarrow M \mid M\ N \mid \dots$	
	\mid	
	$Md \mid Ex \mid Stm \mid Sc \mid Sls \mid Act$	
Program state	$P \text{ ::= } S; \Theta$	
HEC soup	$S \text{ ::= } \emptyset \mid H \parallel S$	
HEC	$H \text{ ::= } \langle s, M, D \rangle \mid \langle s, M, D \rangle_{Sleeping}$	
	$\mid \langle s, M, D \rangle_{Outcall} \mid \text{Idle}$	
Heap	$\Theta \text{ ::= } r \mapsto M \oplus s \mapsto (M, D)$	
SLS Store	$D \text{ ::= } (M, N, r)$	
IO Context	$\mathbb{E} \text{ ::= } \bullet \mid \mathbb{E} \gg= M \mid \text{catch } \mathbb{E}\ M$	
STM Context	$\mathbb{P} \text{ ::= } \bullet \mid \mathbb{P} \gg= M$	

Upcalls from the RTS

	Dequeue upcall instantiation	$H; \Theta \xrightarrow{\text{deq}} H'; \Theta'$
(UPDEQUEUE)		
	$s' \text{ fresh } r \text{ fresh } D' = (\text{deq}(D), \text{enq}(D), r)$	
	$M' = \text{switch } (\lambda x. \text{deq}(D) s)$	
	$\Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]$	
	$\underline{(s, M, D); \Theta \xrightarrow{\text{deq}} \langle s', M', D' \rangle; \Theta'}}$	
Enqueue upcall instantiation		
	$H; \Theta \xrightarrow{\text{enq } s} H'; \Theta'$	
(UPENQUEUEIDLE)		
	$s' \text{ fresh } r \text{ fresh } D' = (\text{deq}(D), \text{enq}(D), r)$	
	$M' = \text{atomically } (\text{enq}(D) s)$	
	$\Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]$	
	$\text{Idle}; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M', D' \rangle; \Theta'$	
(UPENQUEUERUNNING)		
	$M'' = \text{atomically } (\text{enq}(D) s) \gg M'$	
	$\underline{\langle s', M', D' \rangle; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M'', D' \rangle; \Theta[s \mapsto (M, D)]}$	

IO

$E[M]$

tick

$E[\text{yield} \gg M]$

atomically
STM

$E[\text{atomically}(P[M])]$

tick

Suspend transaction
and switch

Switch
STM

$E[\text{switch}(P[M])]$

tick

Formalization

Concurrency Substrate

	$x, y \in Variable$	$r, s, \in Name$
Md	$\text{return } M \mid M \ggg N$	
Ex	$\text{throw } M \mid \text{catch } M\ N \mid \text{catchSTM } M\ N$	
Stm	$\text{newTVar } M \mid \text{readTVar } r \mid \text{writeTVar } r\ M$	
	\mid	
	$\text{atomically } M \mid \text{retry}$	
Sc	$\text{newSCont } M \mid \text{switch } M \mid \text{runOnIdleHEC } s$	
Sls	$\text{getAux } s \mid \text{setAux } s\ M$	
Act	$\text{dequeueAct } s \mid \text{enqueueAct } s$	
	\mid	
	$\text{setDequeueAct } M \mid \text{setEnqueueAct } M$	
Term		
M, N	$r \mid x \mid \lambda.x \rightarrow M \mid M\ N \mid \dots$	
	\mid	
	$Md \mid Ex \mid Stm \mid Sc \mid Sls \mid Act$	
Program state	$P ::= S; \Theta$	
HEC soup	$S ::= \emptyset \mid H \parallel S$	
HEC	$H ::= \langle s, M, D \rangle \mid \langle s, M, D \rangle_{Sleeping}$	
	$\mid \langle s, M, D \rangle_{Outcall} \mid \text{Idle}$	
Heap	$\Theta ::= r \mapsto M \oplus s \mapsto (M, D)$	
SLS Store	$D ::= (M, N, r)$	
IO Context	$\mathbb{E} ::= \bullet \mid \mathbb{E} \ggg M \mid \text{catch } \mathbb{E}\ M$	
STM Context	$\mathbb{P} ::= \bullet \mid \mathbb{P} \ggg M$	

Upcalls from the RTS

Dequeue upcall instantiation	$H; \Theta \xrightarrow{\text{deq}} H'; \Theta'$
(UPDEQUEUE)	
$s' \text{ fresh } r \text{ fresh } D' = (\text{deq}(D), \text{enq}(D), r)$	
$M' = \text{switch } (\lambda x. \text{deq}(D) s)$	
$\Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]$	
$\langle s, M, D \rangle; \Theta \xrightarrow{\text{deq}} \langle s', M', D' \rangle; \Theta'$	
Enqueue upcall instantiation	
$H; \Theta \xrightarrow{\text{enq } s} H'; \Theta'$	
(UPENQUEUEIDLE)	
$s' \text{ fresh } r \text{ fresh } D' = (\text{deq}(D), \text{enq}(D), r)$	
$M' = \text{atomically } (\text{enq}(D) s)$	
$\Theta' = \Theta[s \mapsto (M, D)][r \mapsto \text{toDyn } ()]$	
$\text{Idle}; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M', D' \rangle; \Theta'$	
(UPENQUEUERUNNING)	
$M'' = \text{atomically } (\text{enq}(D) s) \gg M'$	
$\langle s', M', D' \rangle; \Theta[s \mapsto (M, D)] \xrightarrow{\text{enq } s} \langle s', M'', D' \rangle; \Theta[s \mapsto (M, D)]$	

IO

$E[M] \xrightarrow{\text{tick}} E[\text{yield} \gg M]$

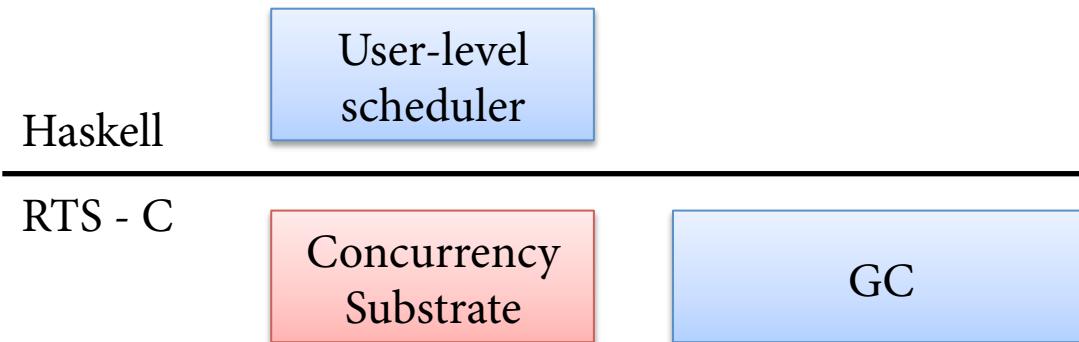
atomically
STM

$E[\text{atomically}(P[M])] \xrightarrow{\text{tick}}$ Suspend transaction
and switch

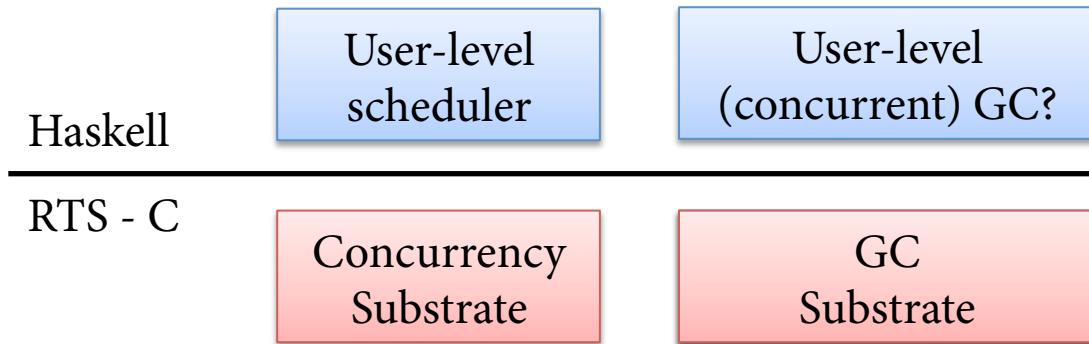
Switch
STM

$E[\text{switch}(P[M])] \xrightarrow{\text{tick}}$ Disable Interrupts!

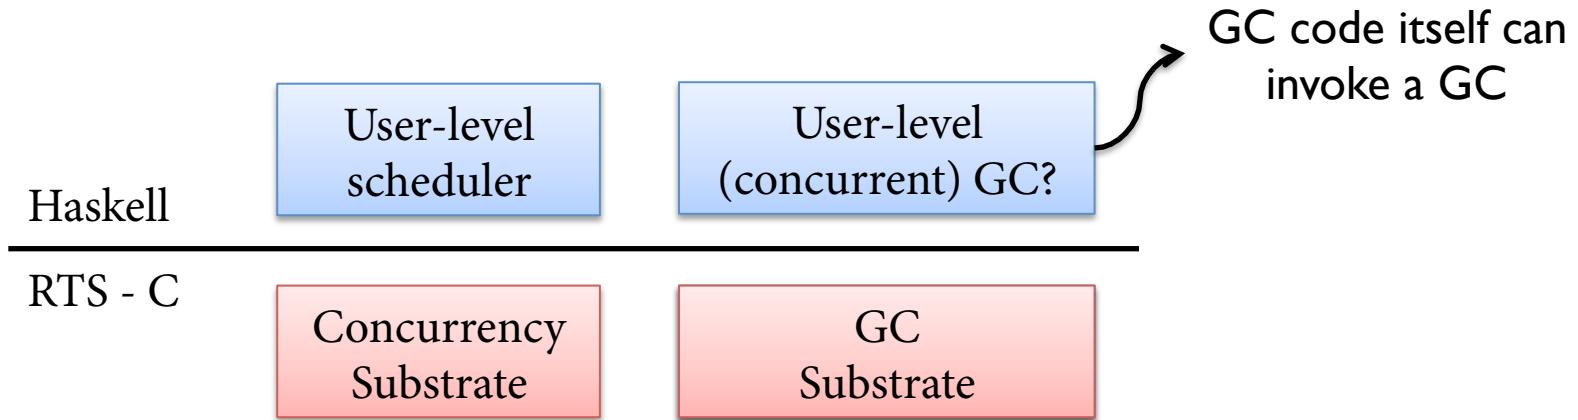
Future directions



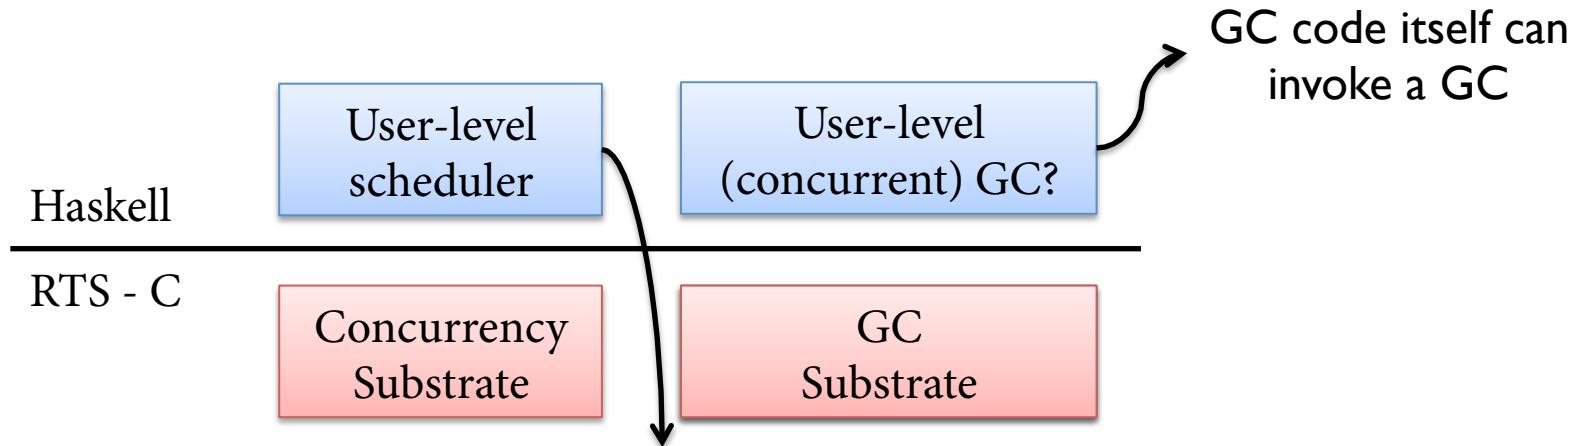
Future directions



Future directions

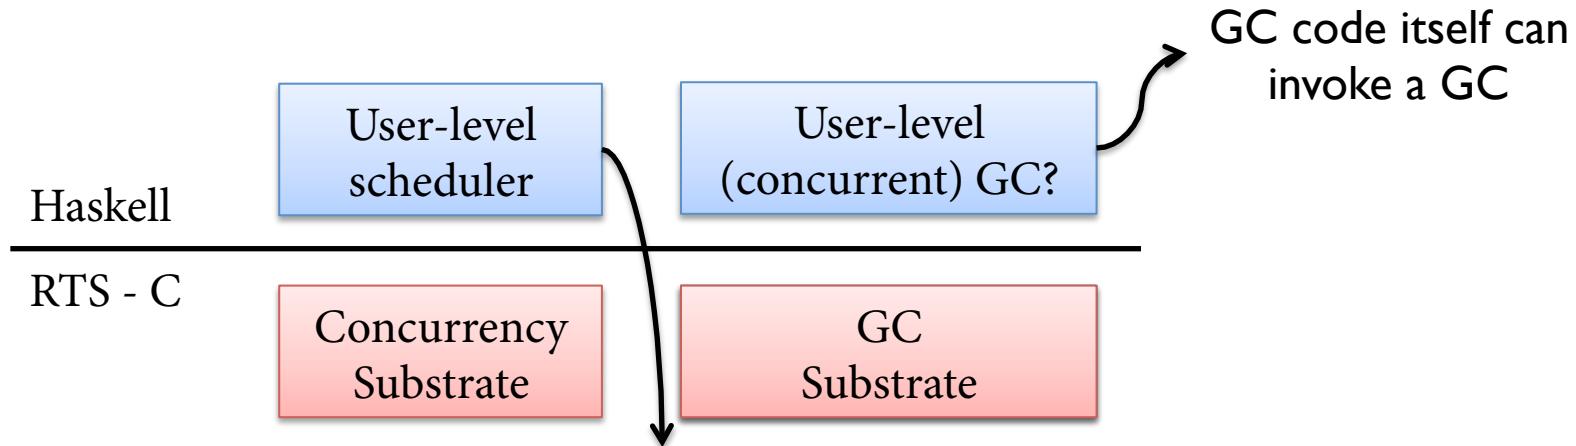


Future directions



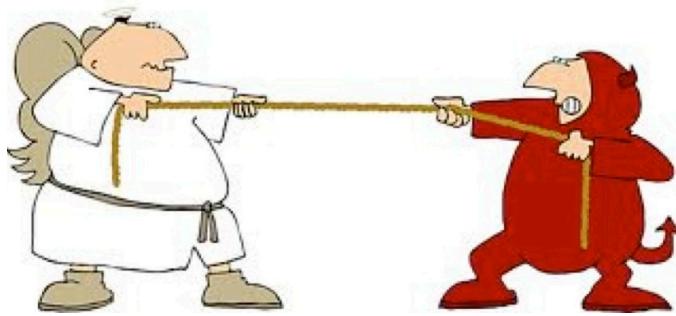
- Verifying the functional correctness of schedulers
 - Correct Scheduler \Rightarrow Each thread runs as if given its **own processor** and **register set**
 - Scheduler access under STM \Rightarrow Treat scheduler as **sequential process**

Future directions



- Verifying the functional correctness of schedulers
 - Correct Scheduler \Rightarrow Each thread runs as if given its **own processor** and **register set**
 - Scheduler access under STM \Rightarrow Treat scheduler as **sequential process**
- FP abstractions for eventually consistent systems
 - Operations on ECDTs described as pure functions over axiomatic executions
 - A relational specification language for specifying consistency assertions over axiomatic executions

Conclusion



- Functional programming abstractions simplify concurrent programming
 - Rx-CML: Synchronous communication over geo-distributed systems
 - Concurrency Substrate: Scheduler activation + STM for writing schedulers
- Abstractions introduce indirection resulting in overheads
 - Rx-CML slower than explicit async under contention
 - Conc. Subs: Scheduler allocations increase GC overheads