



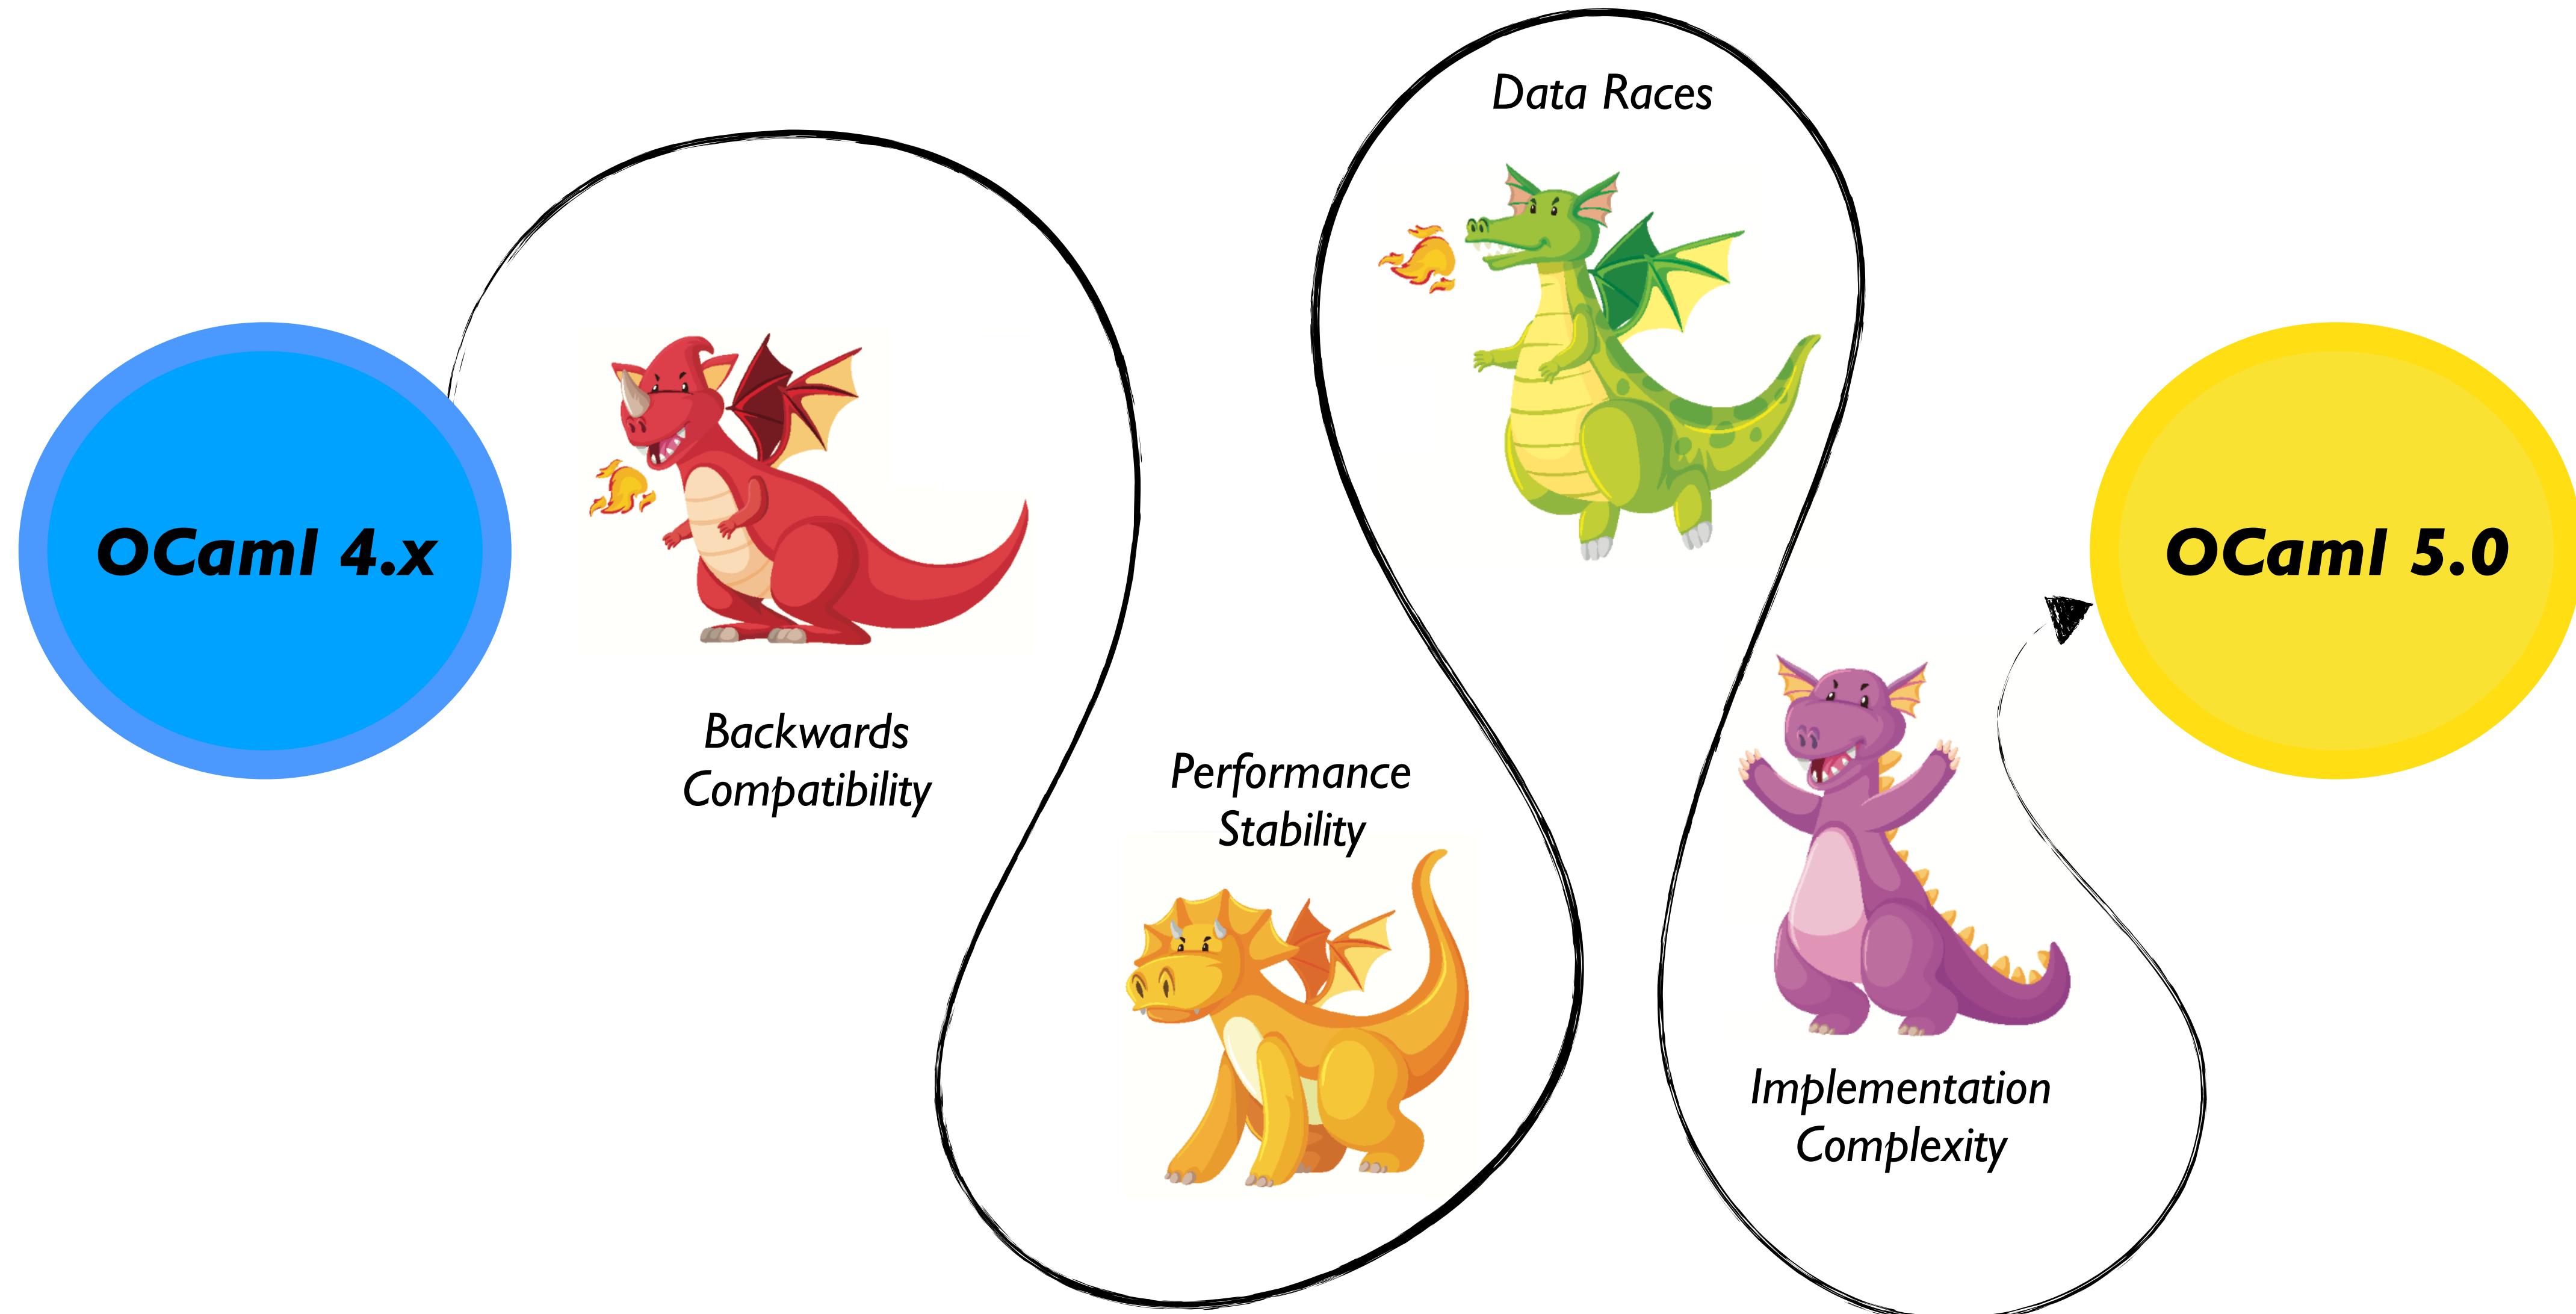
# OCaml 5.0

“KC” Sivaramakrishnan

IIT  
MADRAS  
SADAM



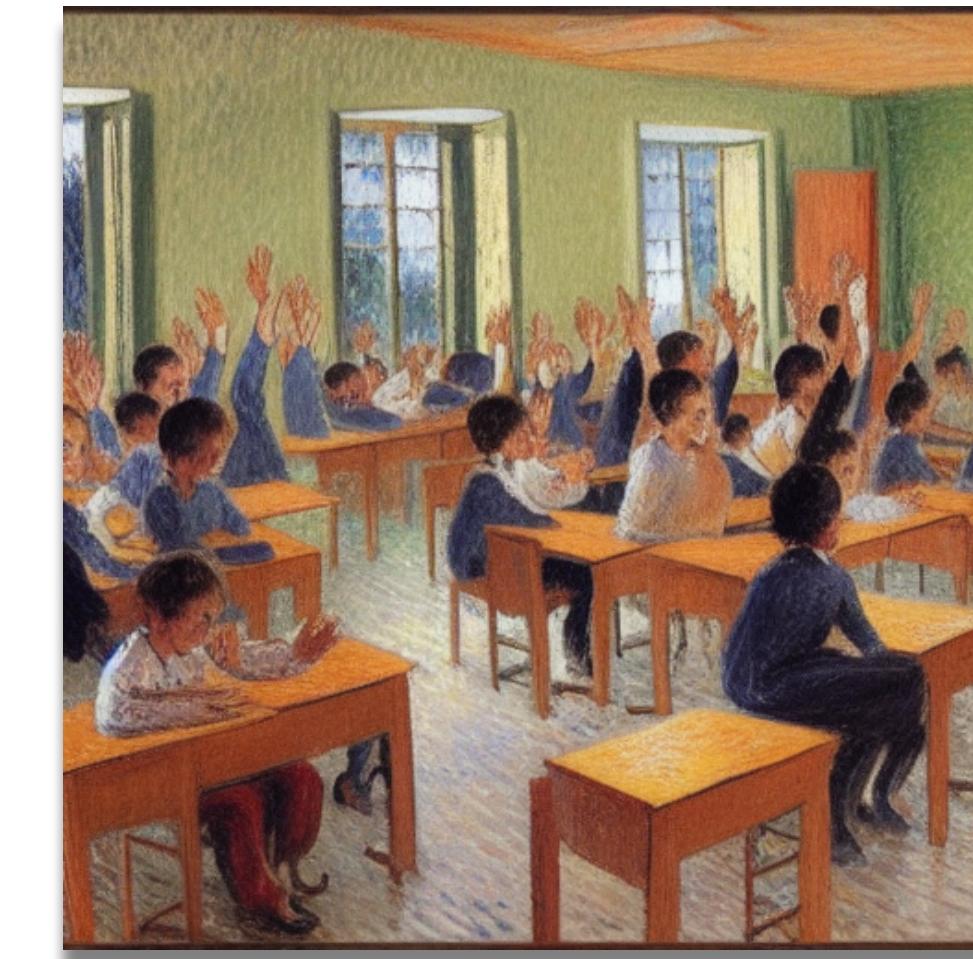
# ICFP Keynote



# This talk...



*What's in the can?*



FAQs

*Merge Process*



*Moving to OCaml 5.0*

# Concurrency and Parallelism

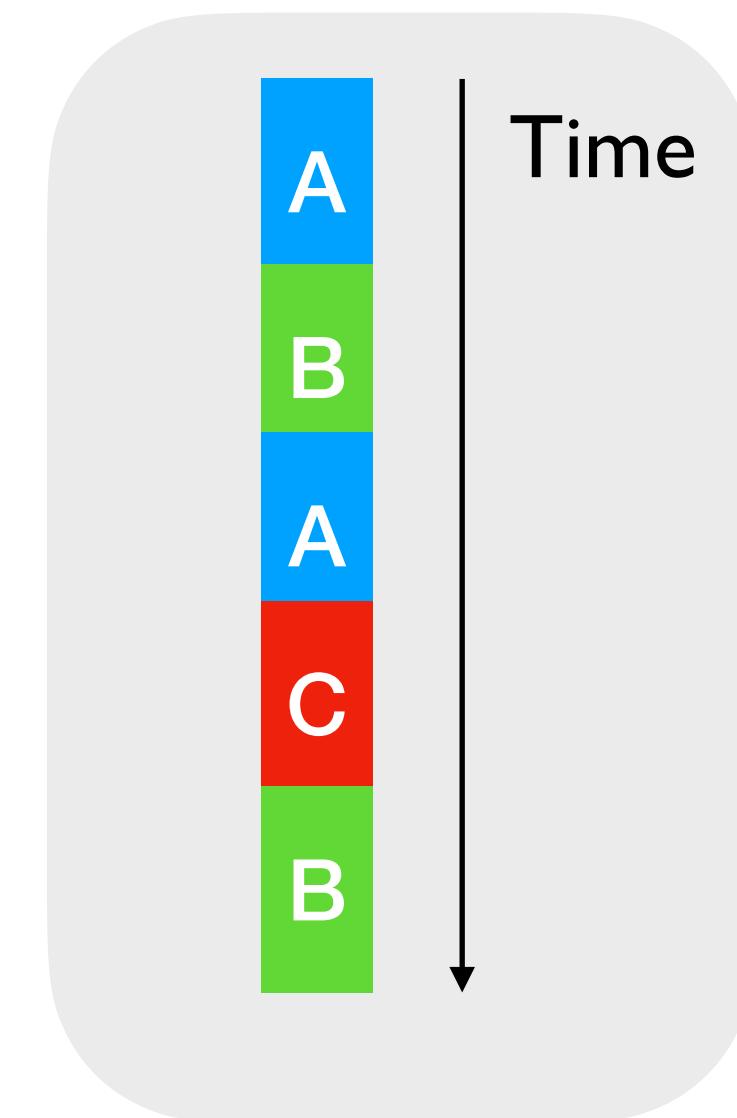
# Concurrency and Parallelism

*Concurrency*

*Parallelism*

# Concurrency and Parallelism

*Concurrency*



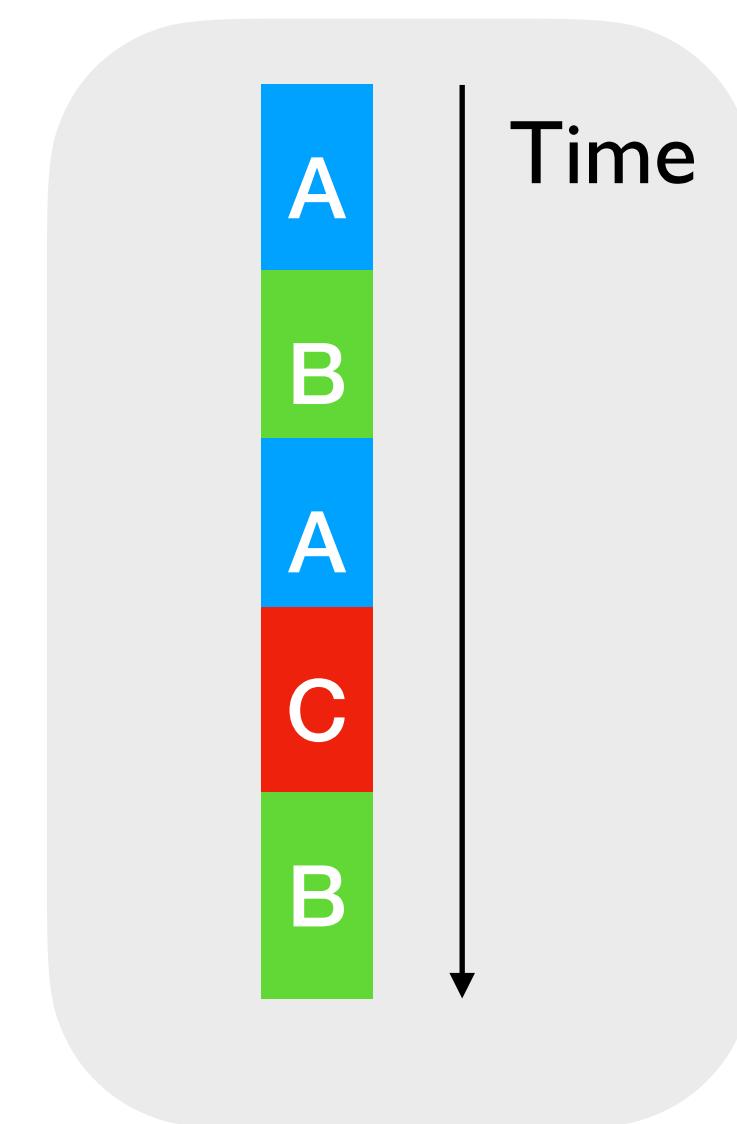
*Overlapped  
execution*

*Effect Handlers*

*Parallelism*

# Concurrency and Parallelism

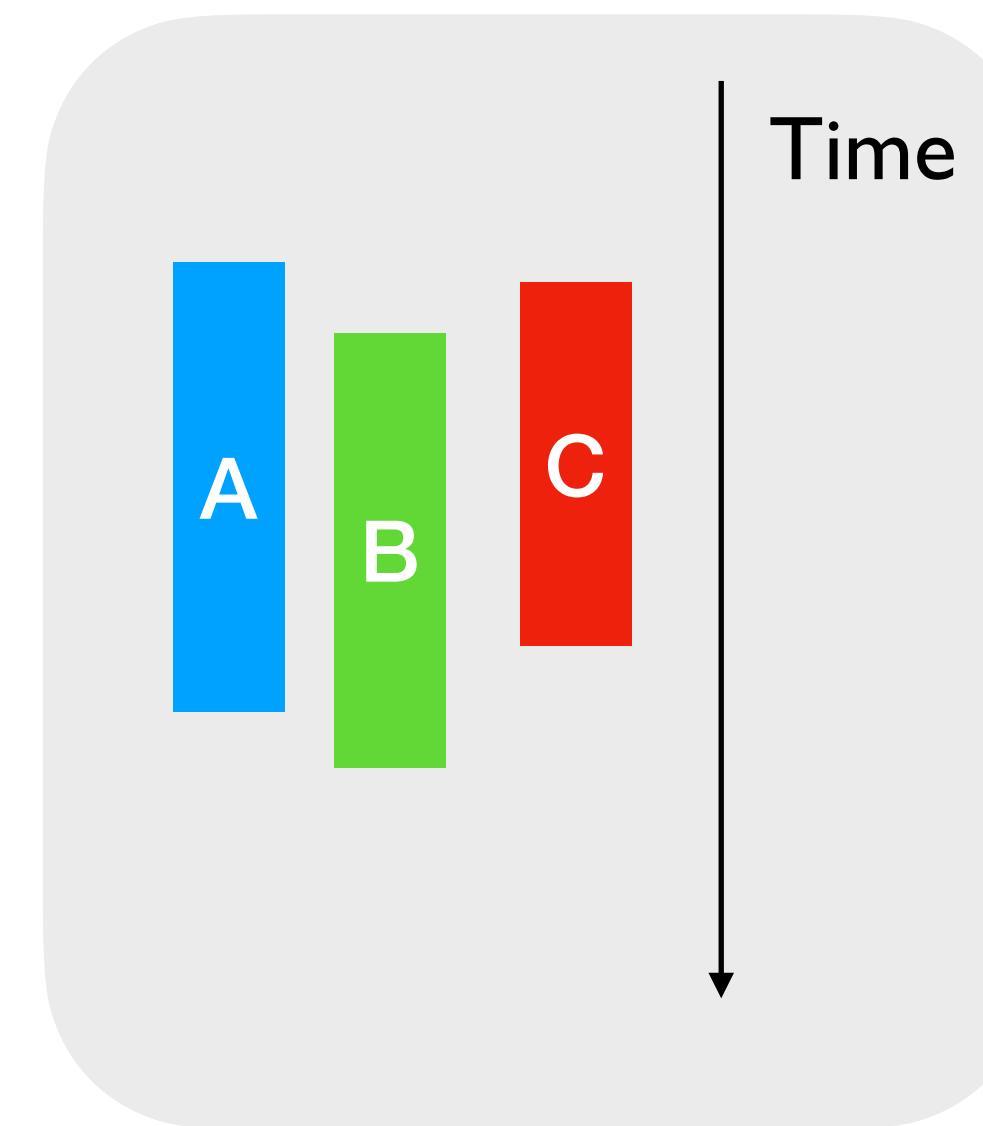
## *Concurrency*



*Overlapped  
execution*

## *Effect Handlers*

## *Parallelism*

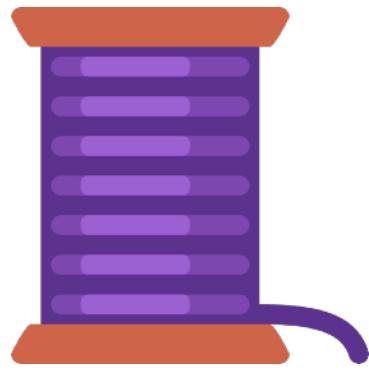


*Simultaneous  
execution*

## *Domains*

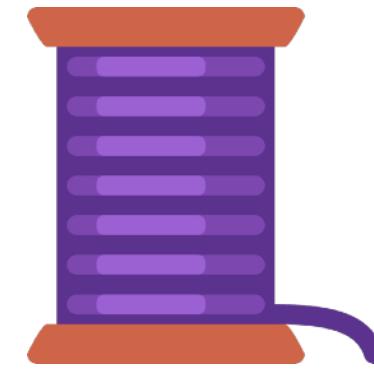
# Domains

- Units of parallelism
- *Heavy-weight* entities
  - ◆ Recommended to have 1 domain per core



OCaml

*Domain 0*

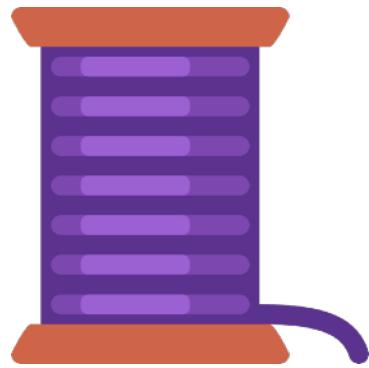


OCaml

*Domain 1*

# Domains

- Units of parallelism
- *Heavy-weight* entities
  - ◆ Recommended to have 1 domain per core
- API
  - ◆ *Create and destroy* — Spawn and Join
  - ◆ *Blocking synchronisation* — Mutex, Condition and Semaphore
  - ◆ *Non-blocking synchronisation* — Atomic
  - ◆ Domain-local state



OCaml

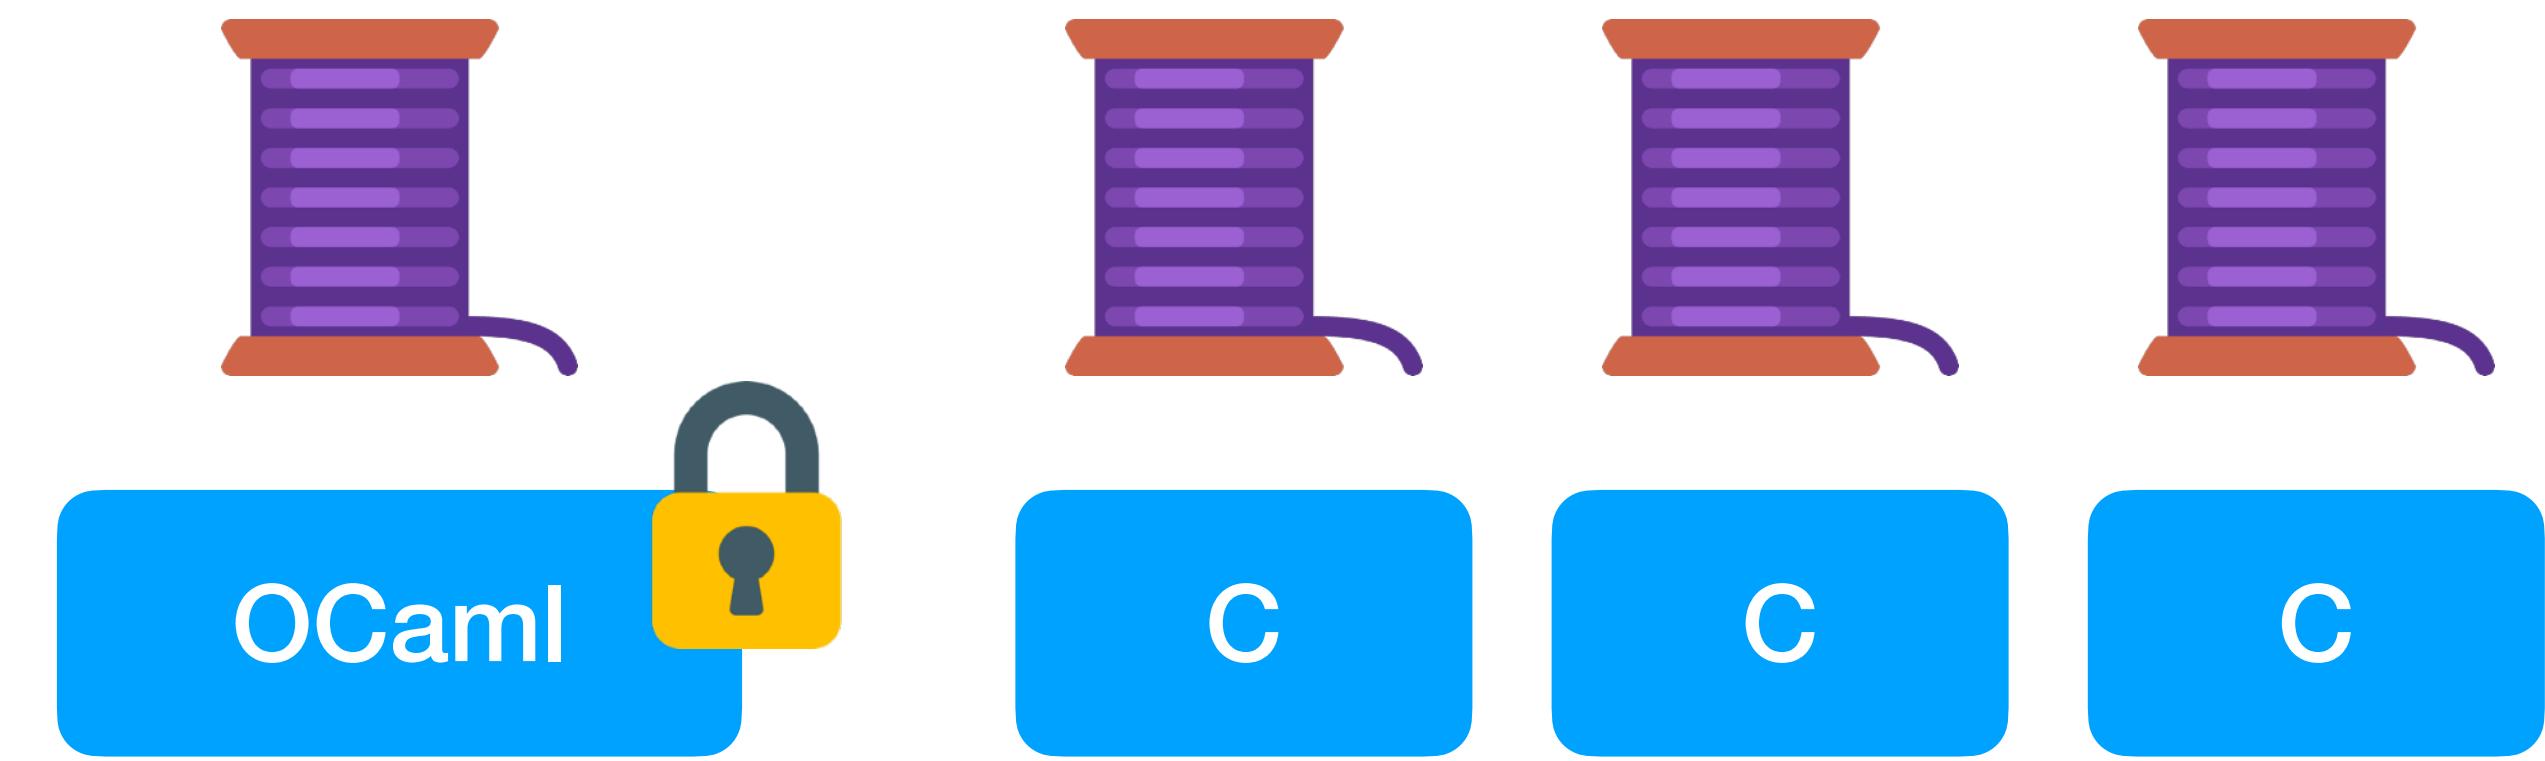
Domain 0



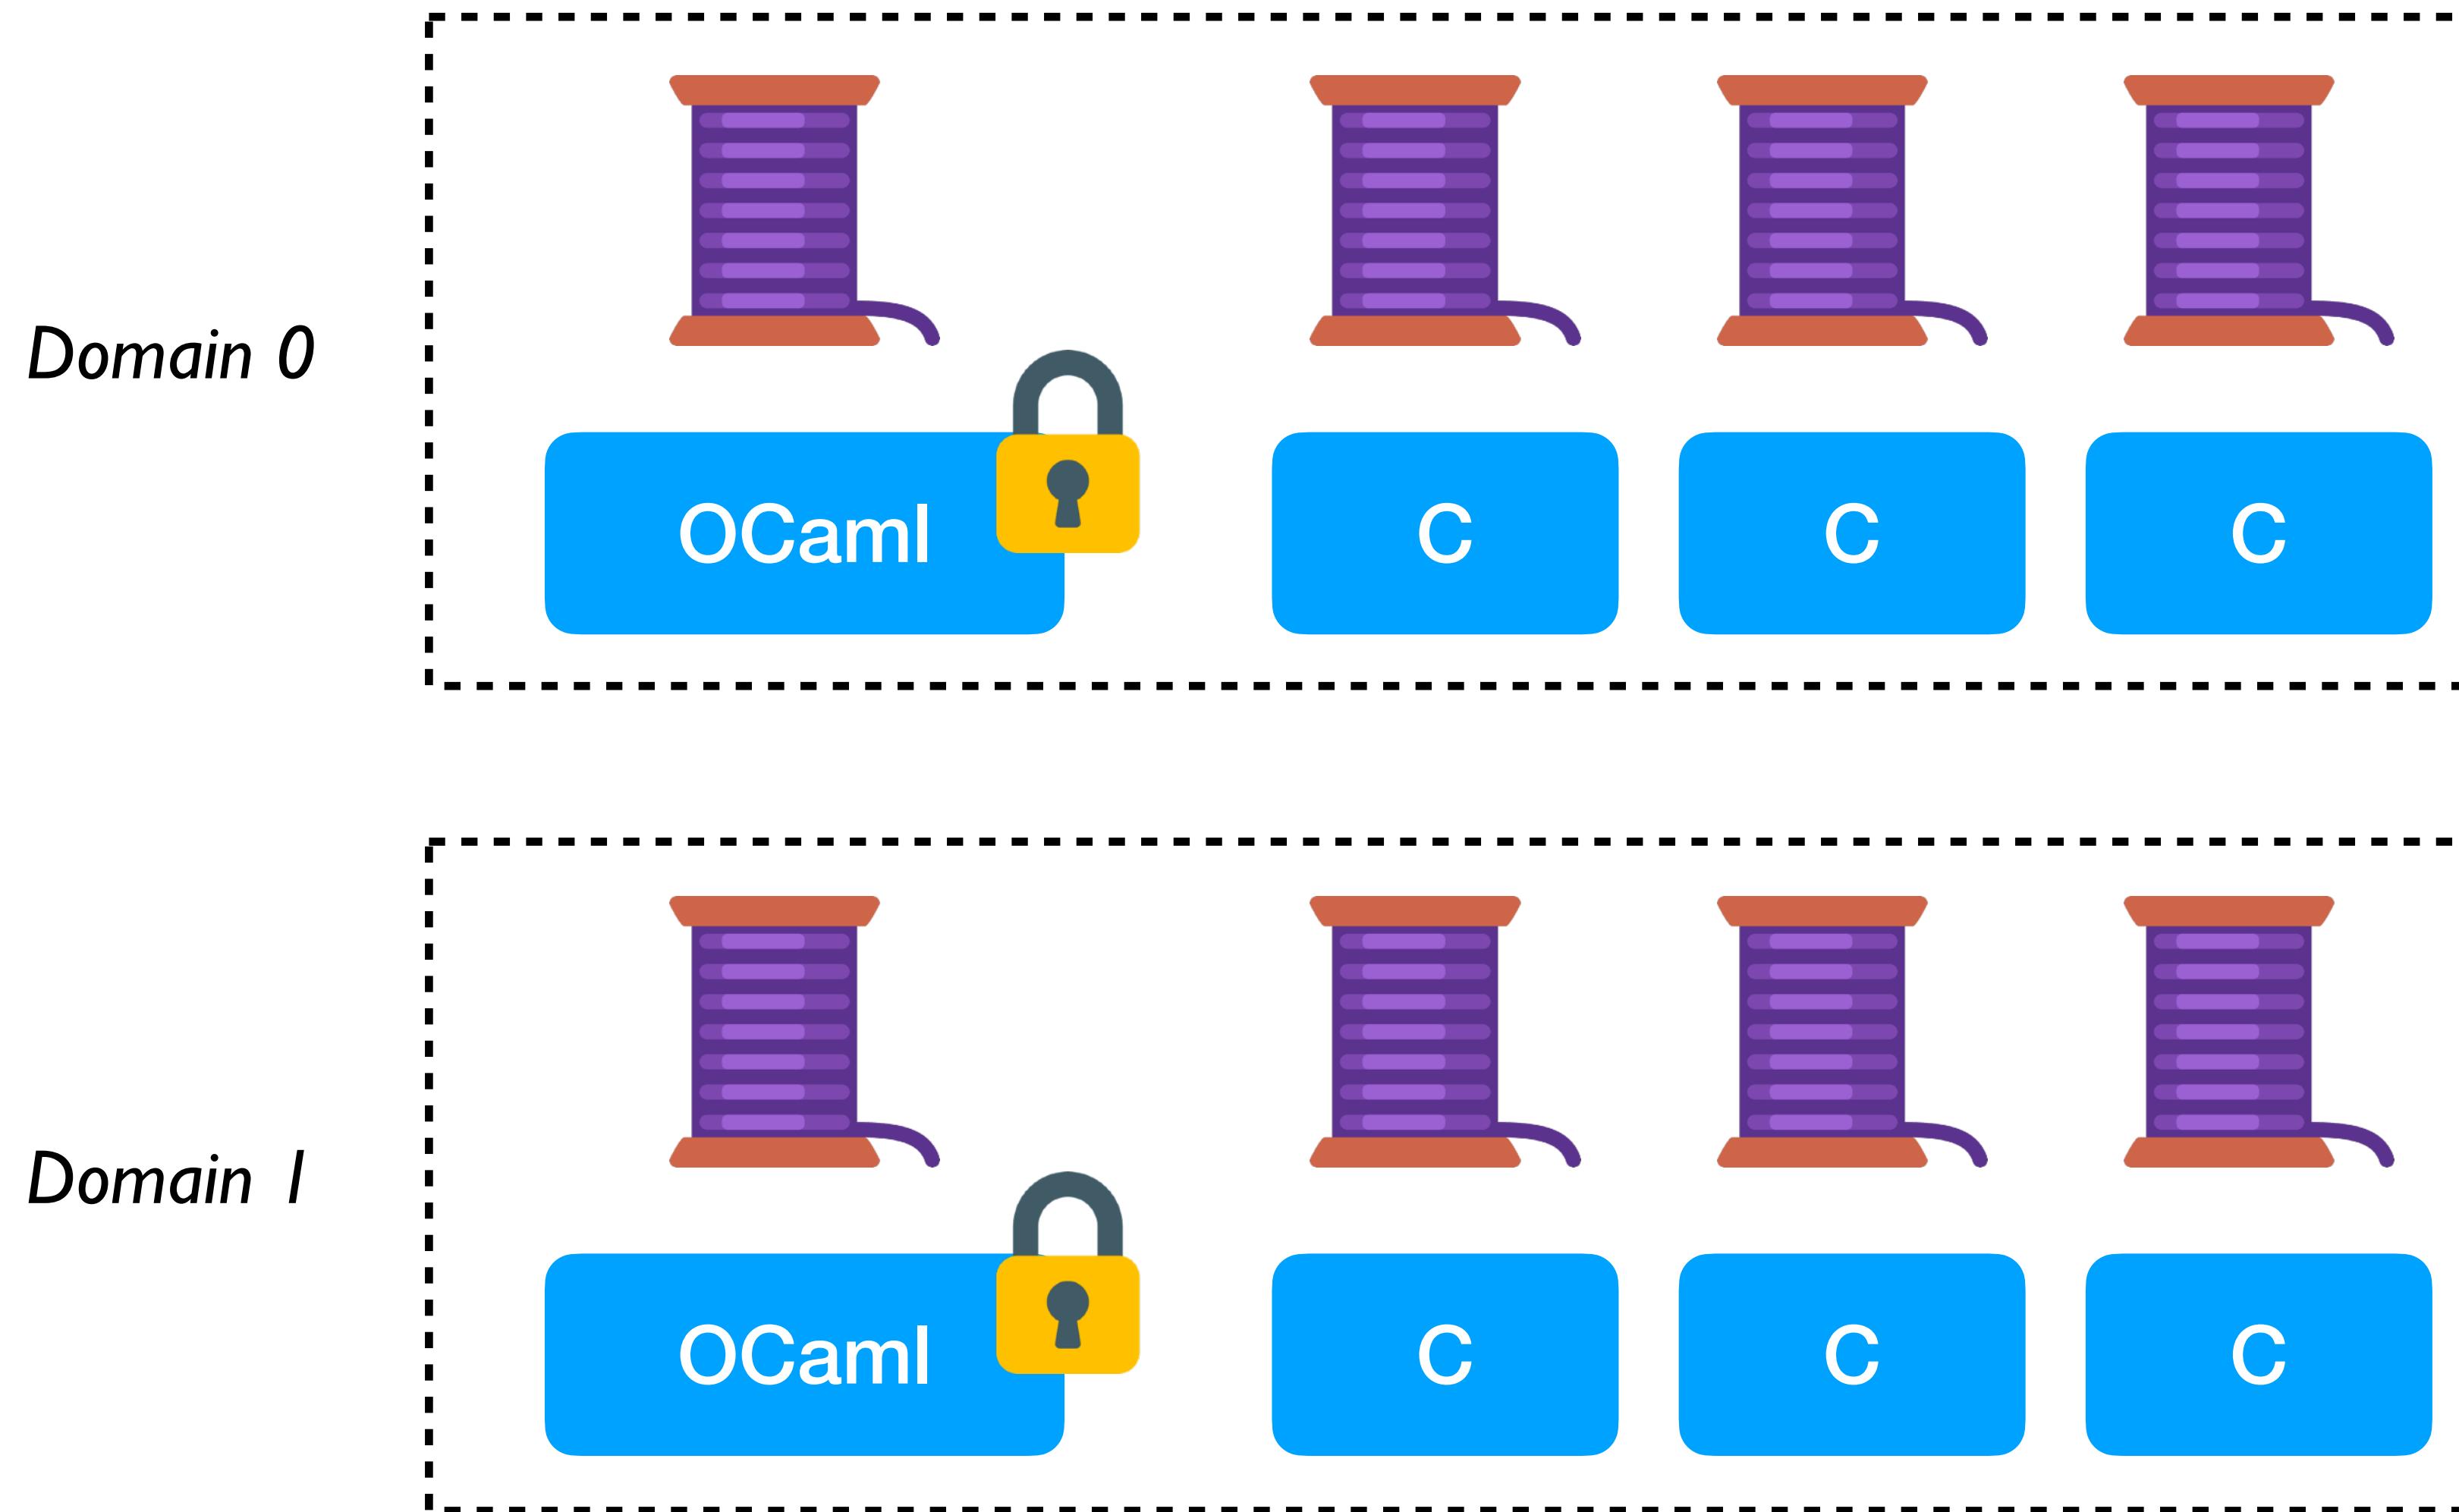
OCaml

Domain 1

# Threads



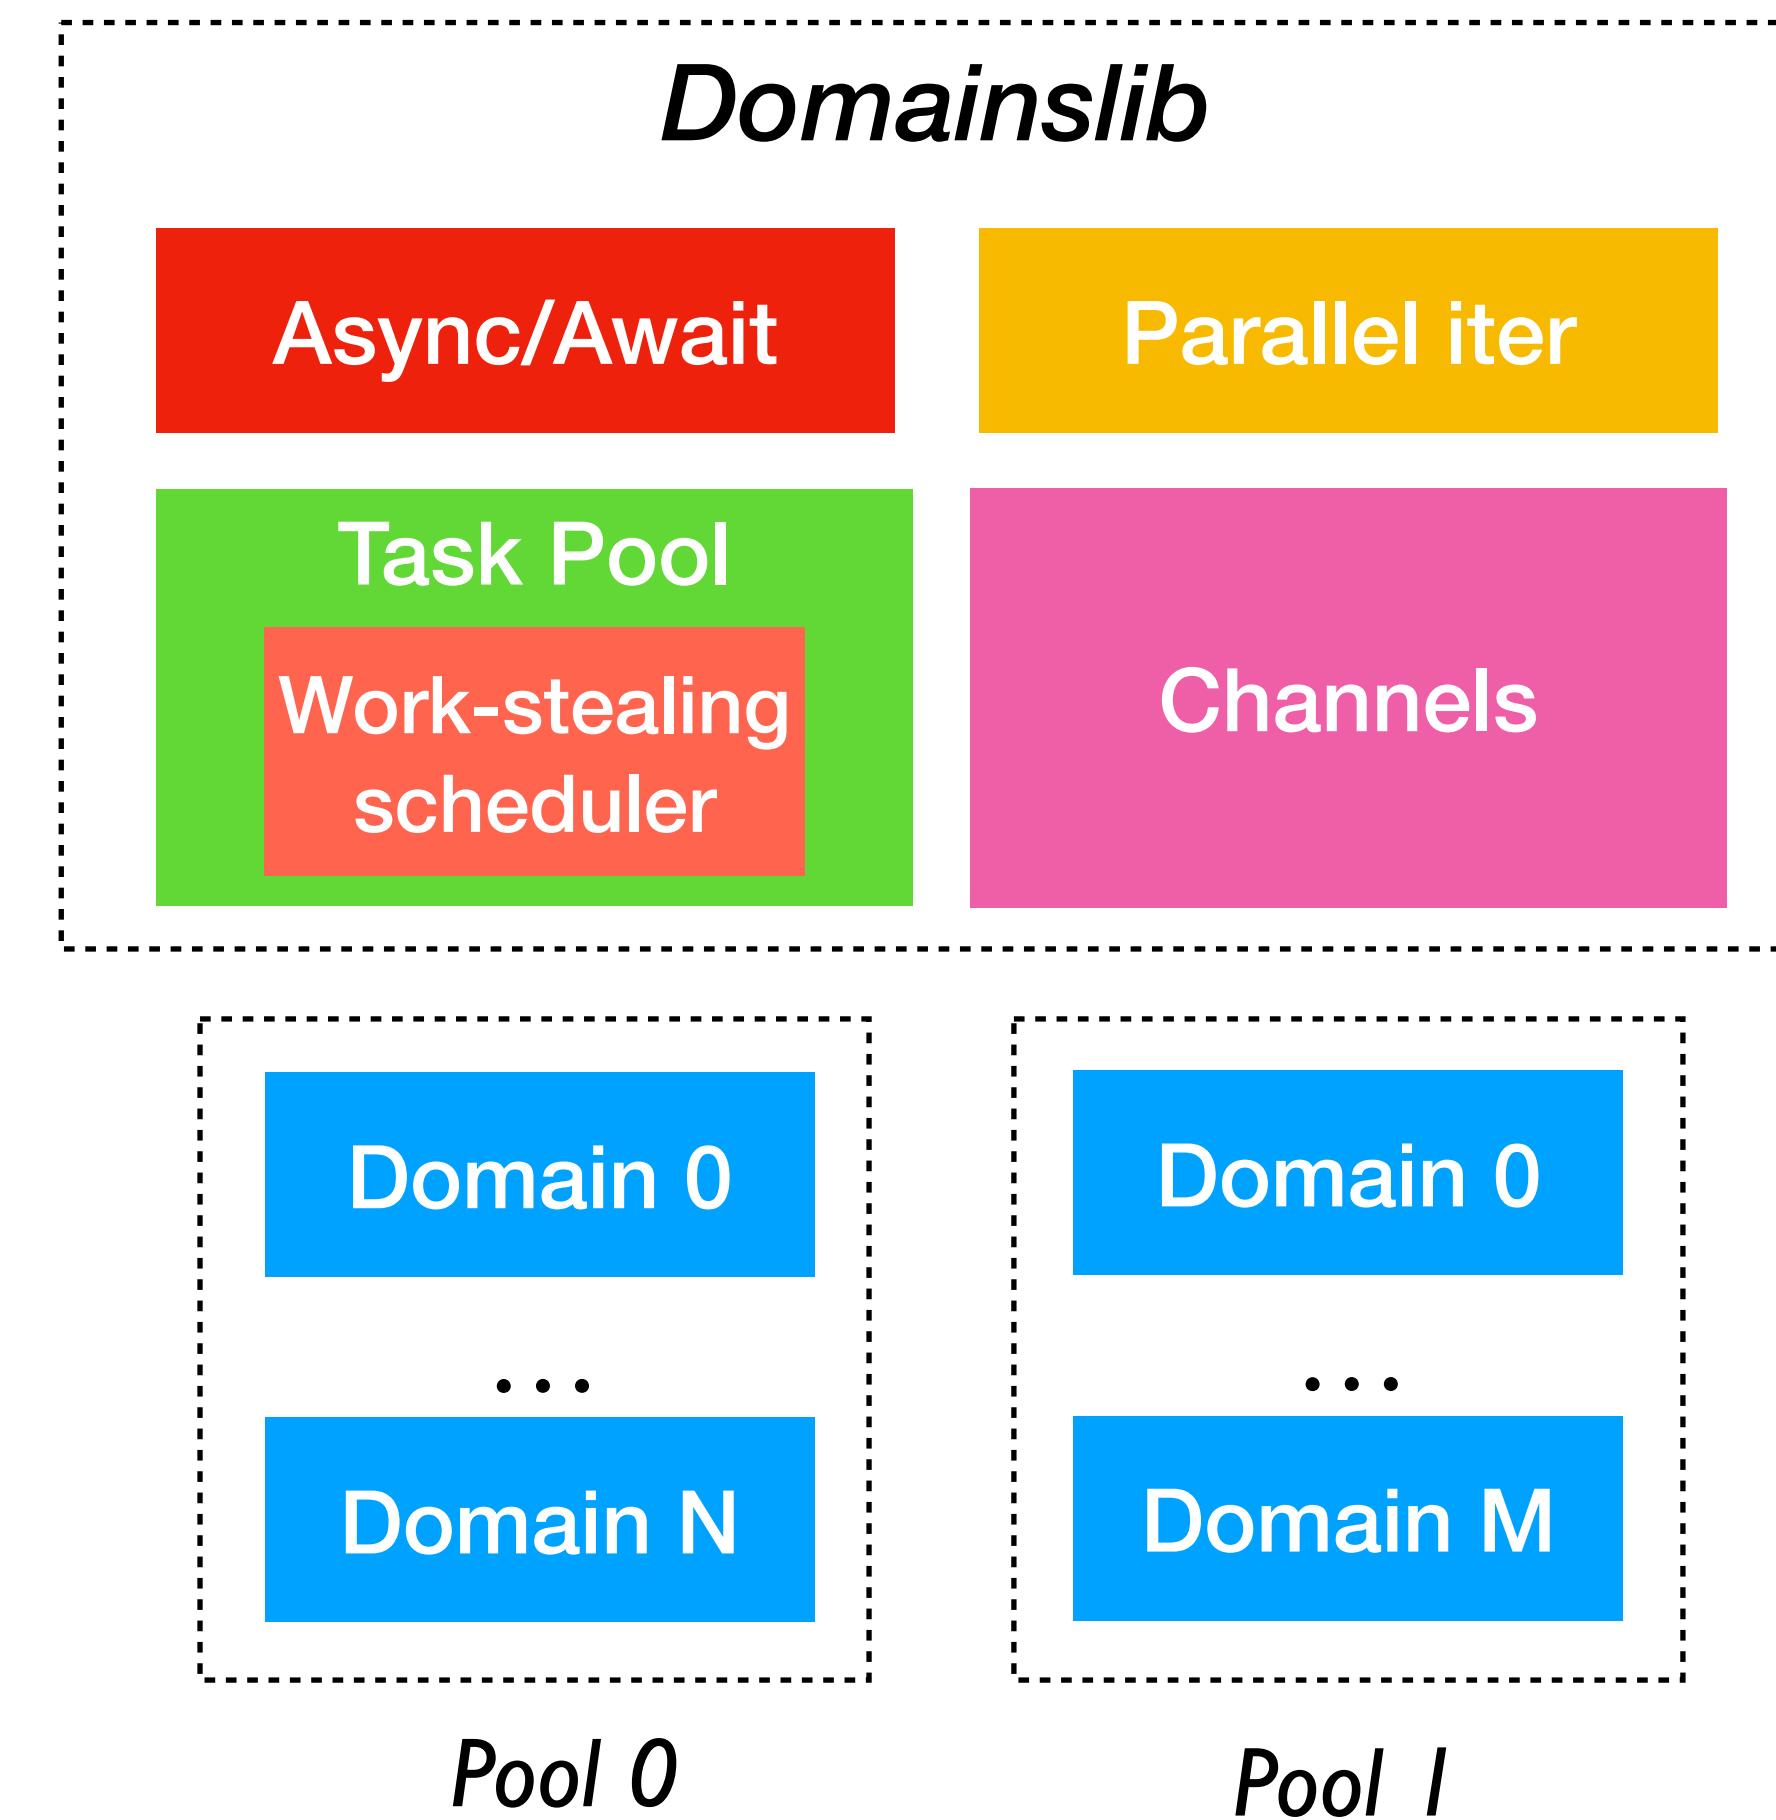
# Domains with Threads



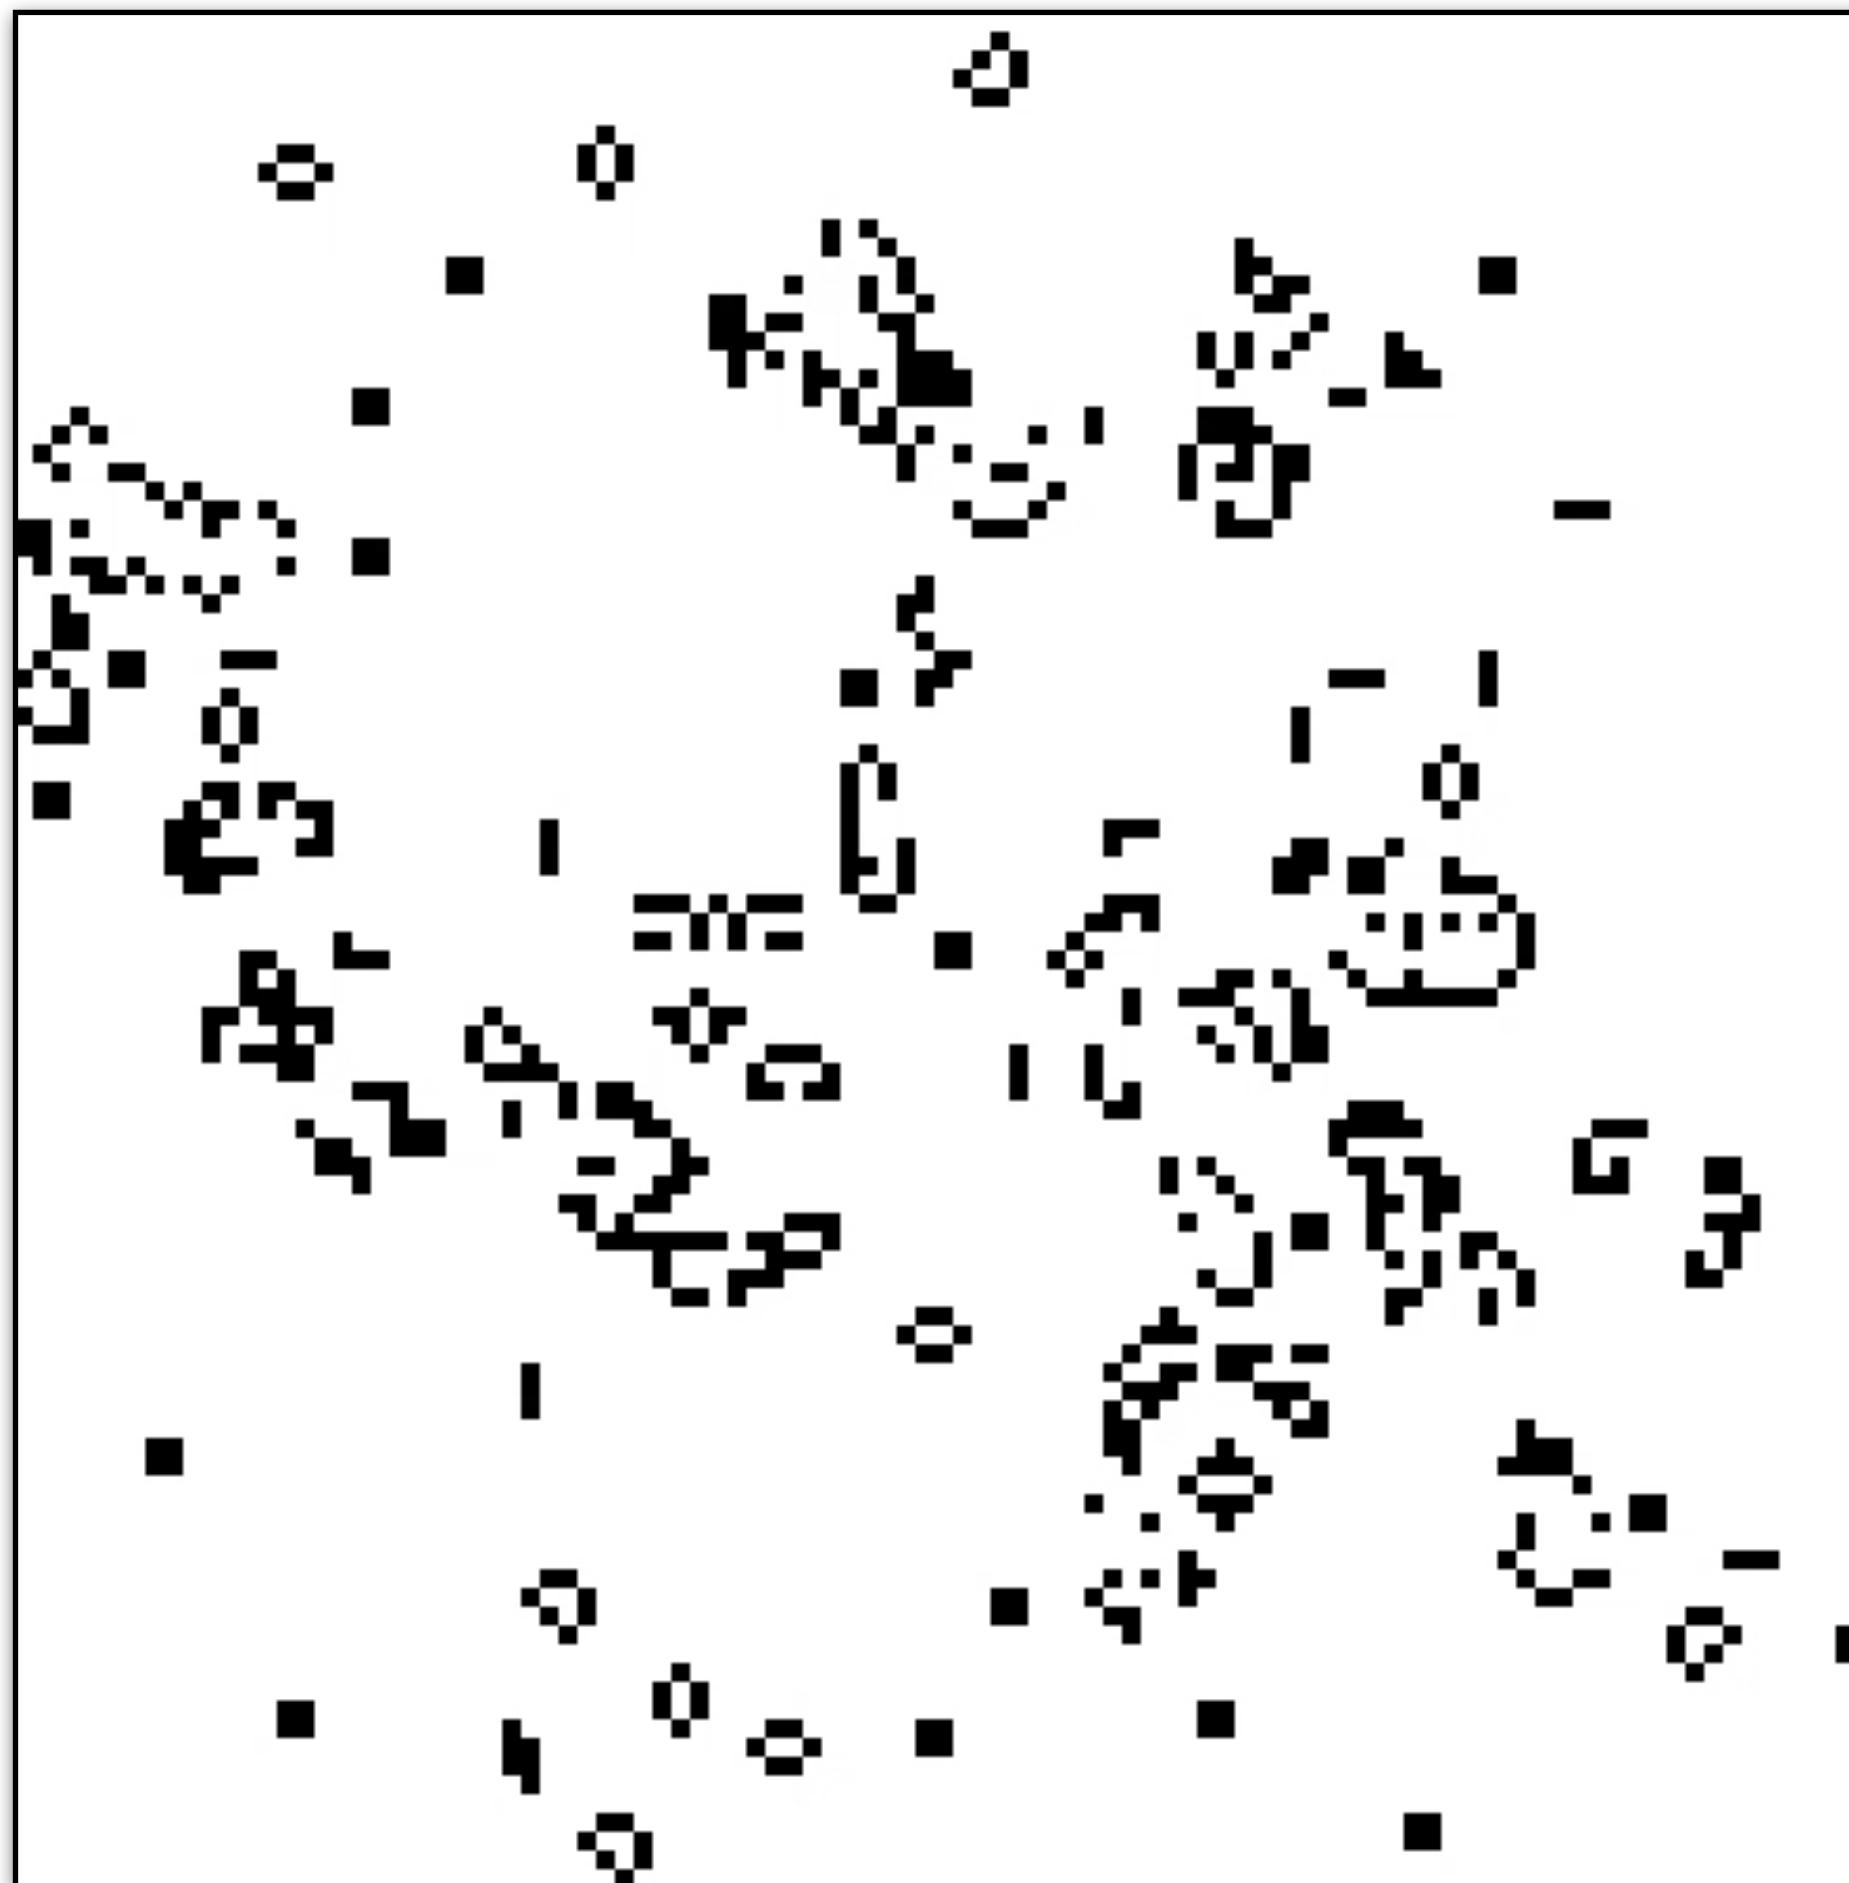
*Blocking and non-blocking  
synchronisation works  
uniformly across threads  
and domains*

# Domainslib

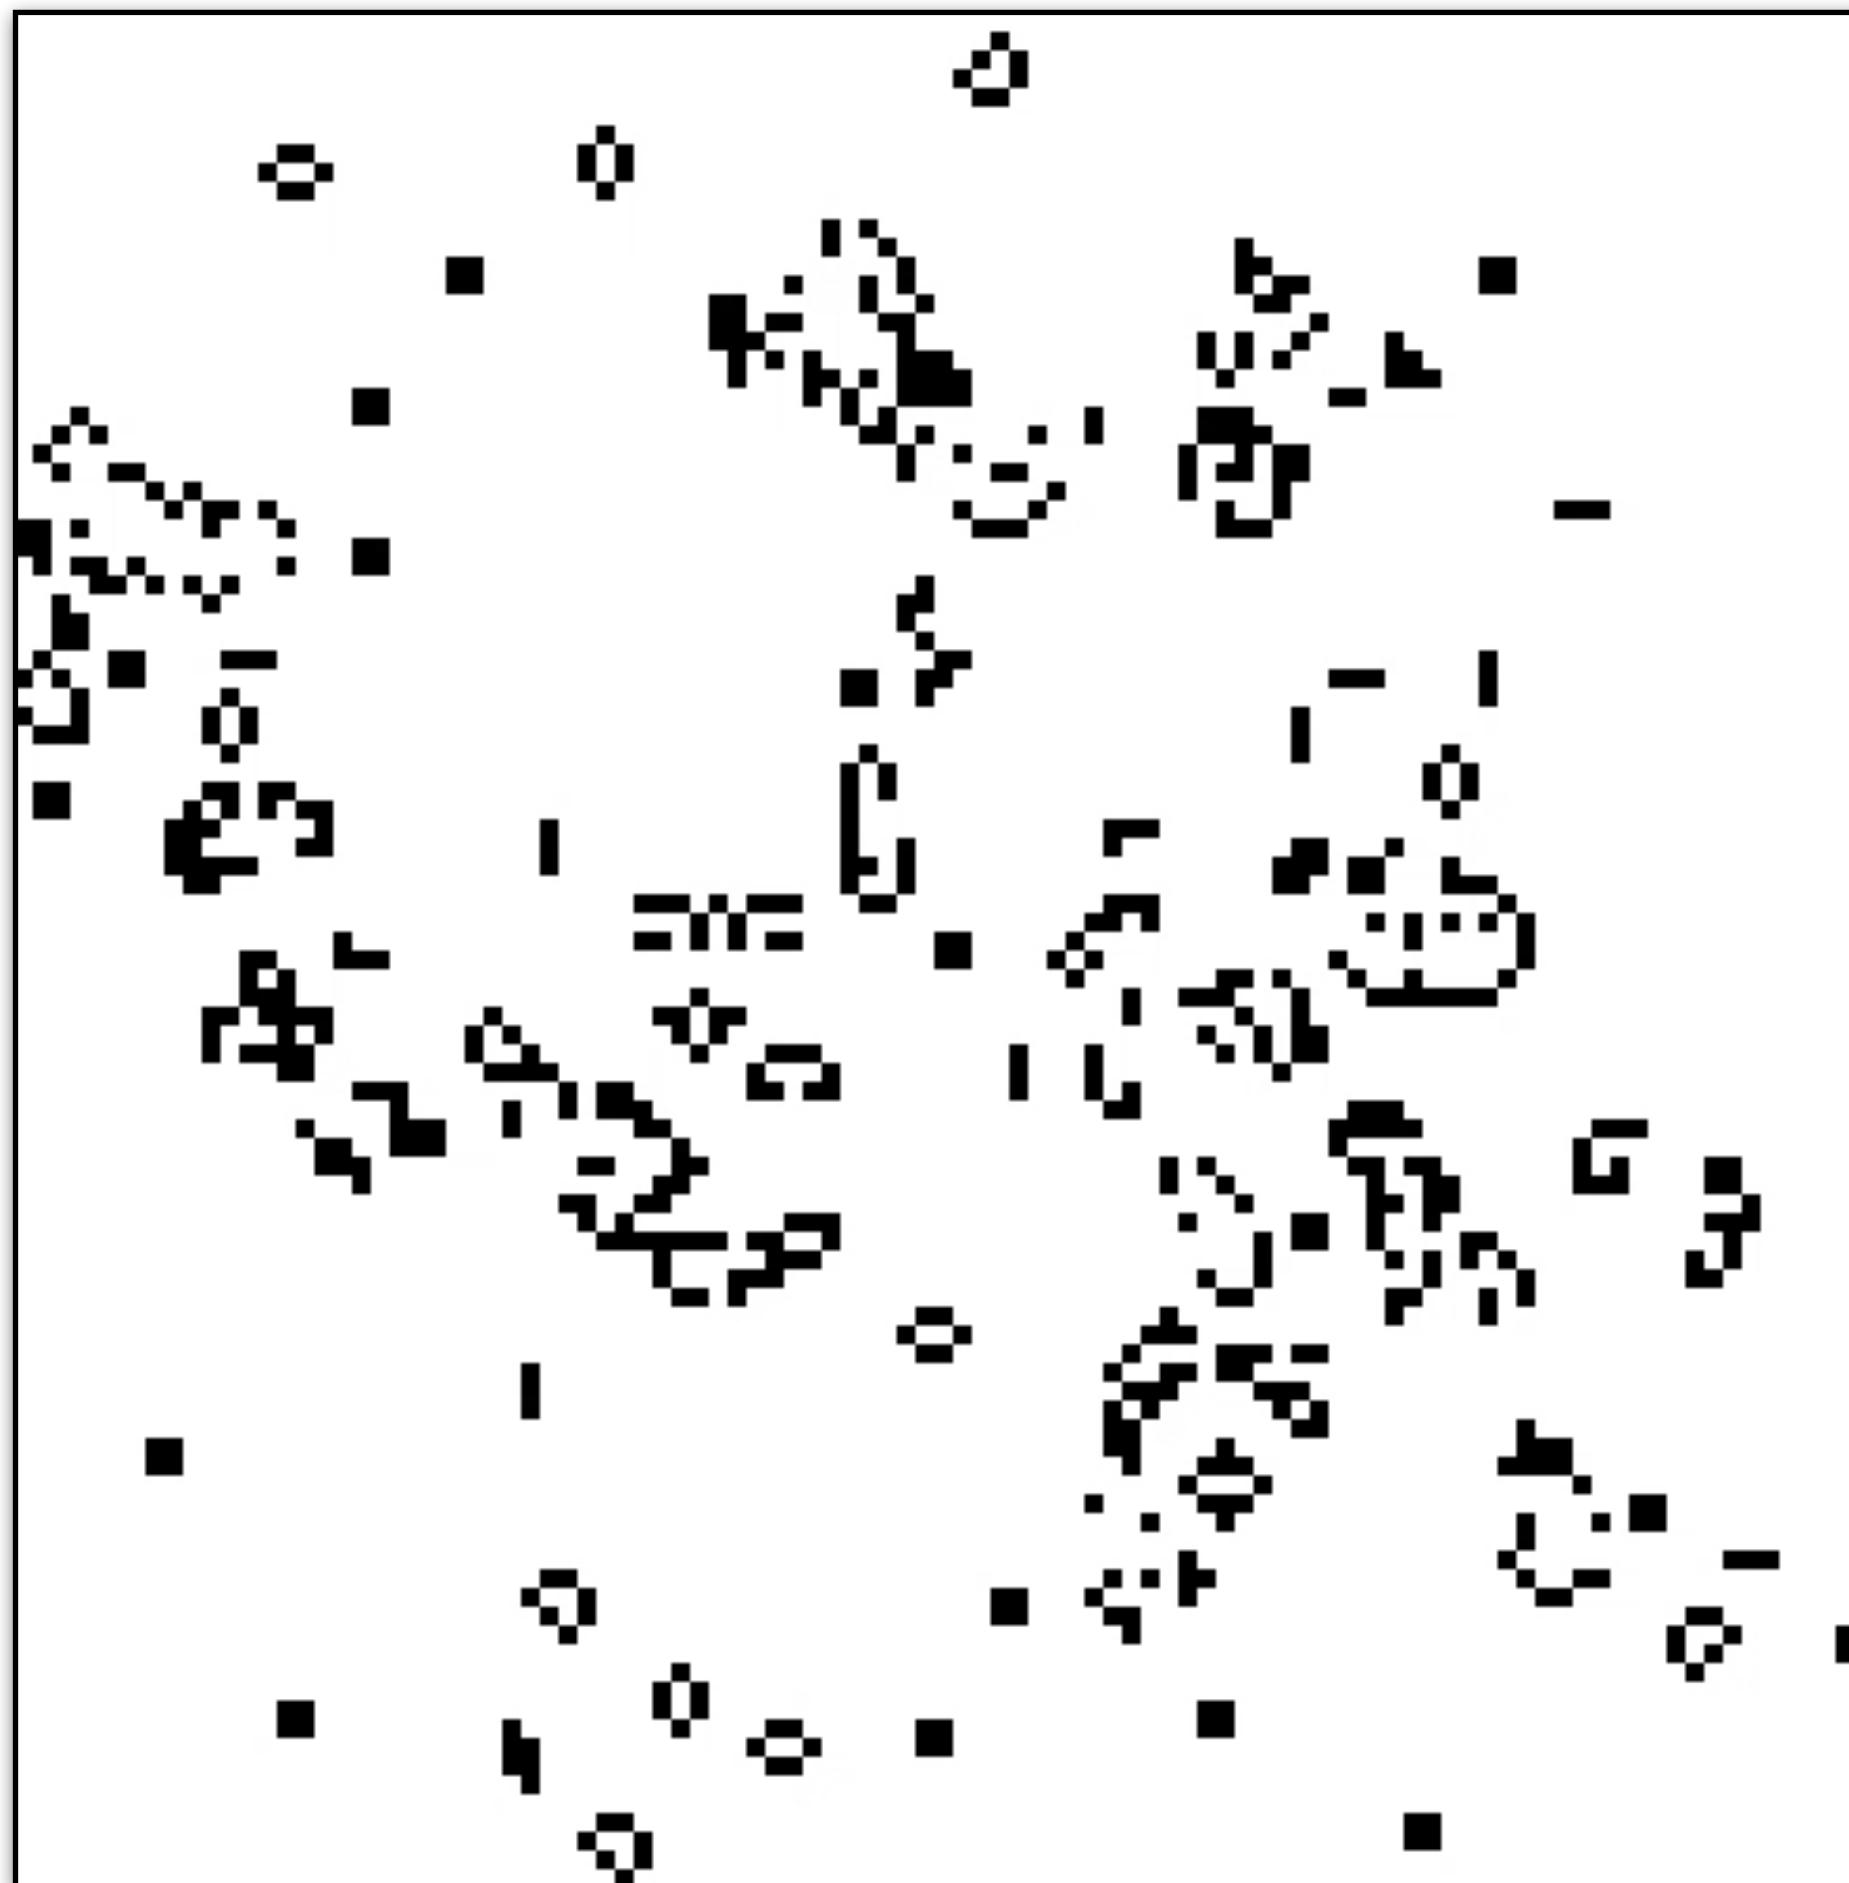
- A library for nested-parallel programming (OpenMP, Cilk, NESL,...)



# Conway's Game of Life



# Conway's Game of Life

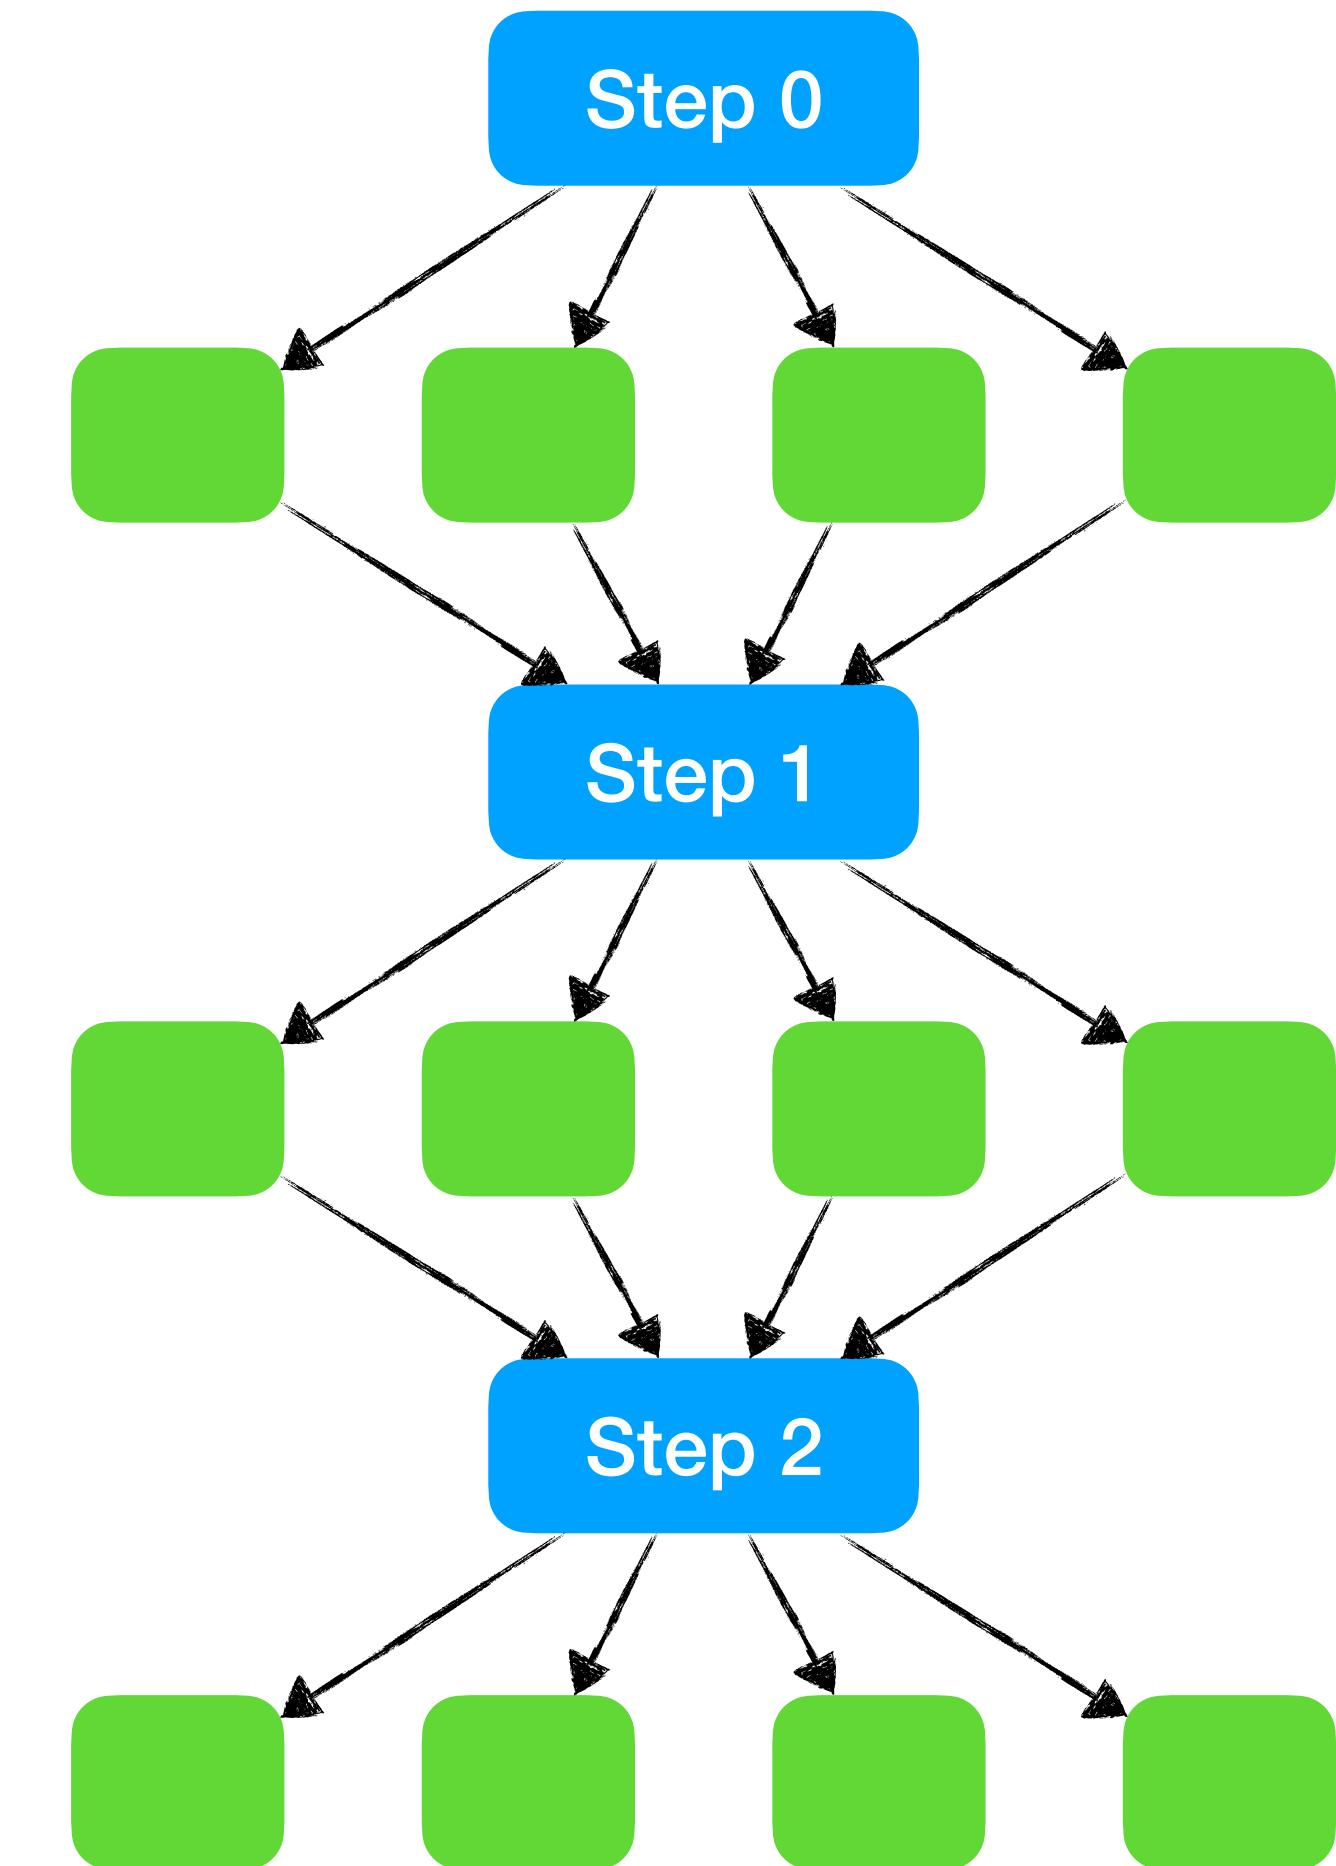


# Conway's Game of Life

```
let next () =  
  ...  
  for x = 0 to board_size - 1 do  
    for y = 0 to board_size - 1 do  
      next_board.(x).(y) <- next_cell cur_board x y  
    done  
  done;  
  ...
```

# Conway's Game of Life

```
let next () =  
  ...  
  for x = 0 to board_size - 1 do  
    for y = 0 to board_size - 1 do  
      next_board.(x).(y) <- next_cell cur_board x y  
    done  
  done;  
  ...  
  
let next () =  
  ...  
  T.parallel_for pool ~start:0 ~finish:(board_size - 1)  
    ~body:(fun x ->  
      for y = 0 to board_size - 1 do  
        next_board.(x).(y) <- next_cell cur_board x y  
      done);  
  ...
```

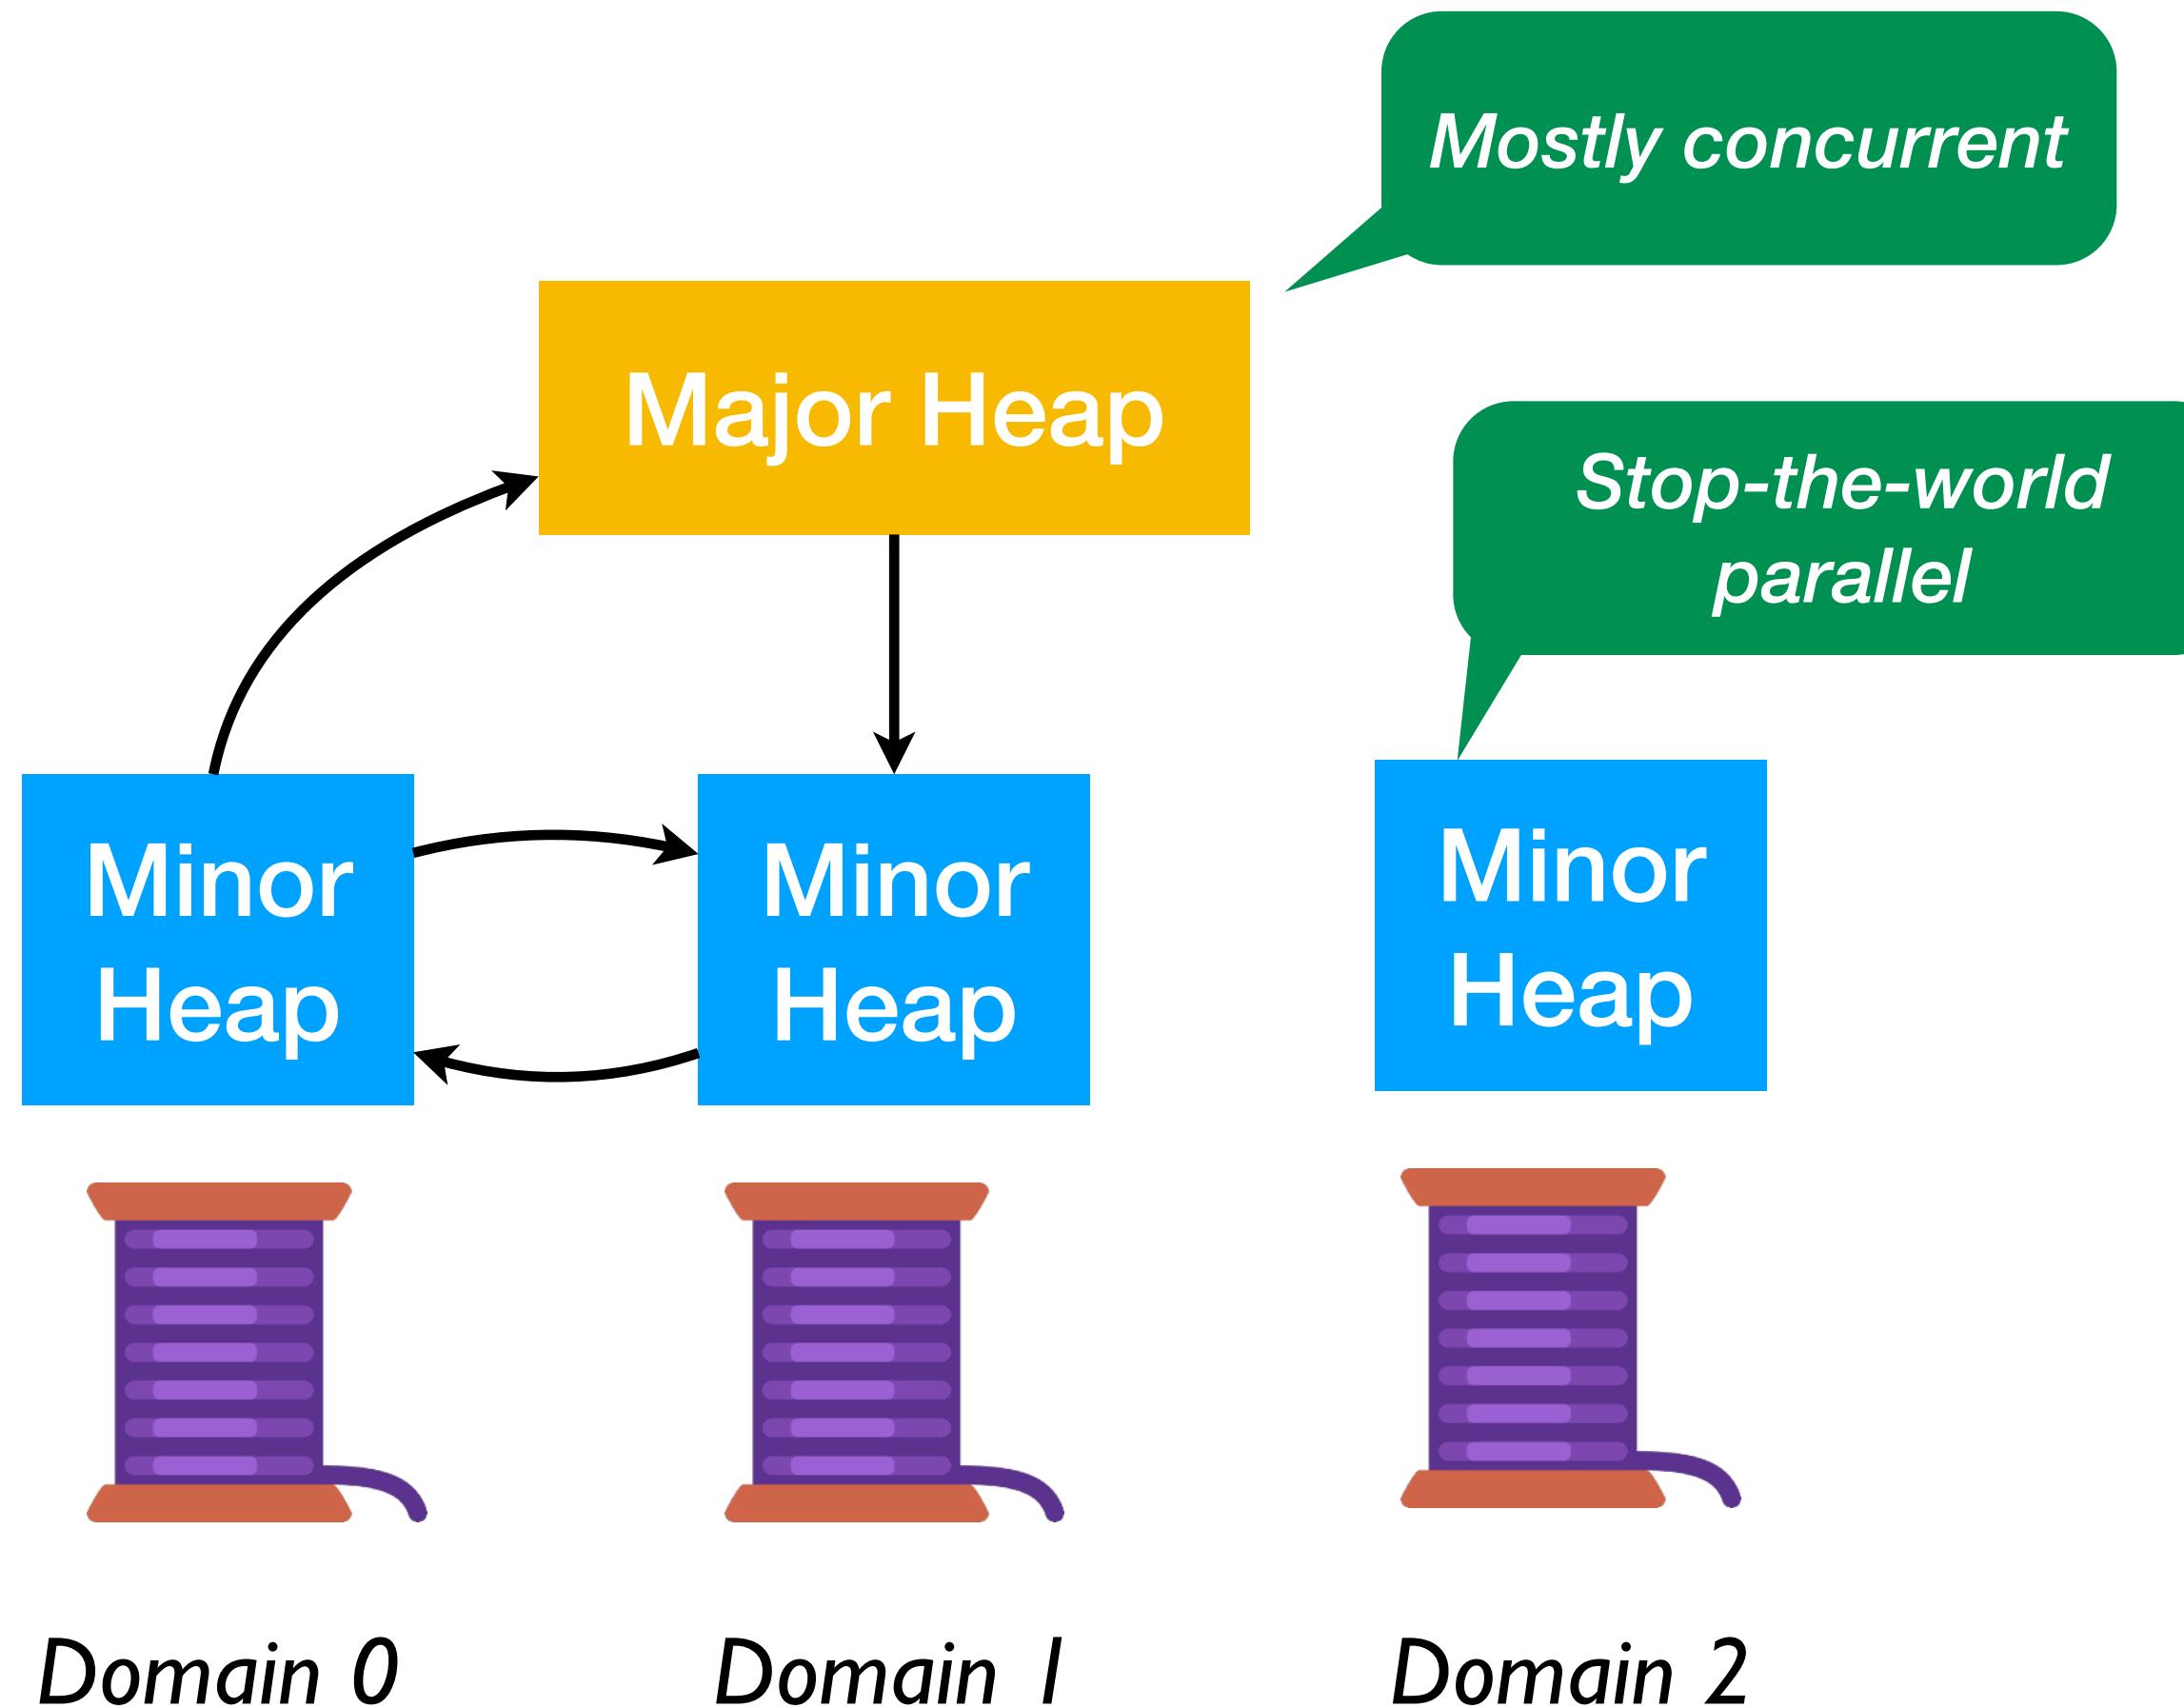


# Performance: Game of Life

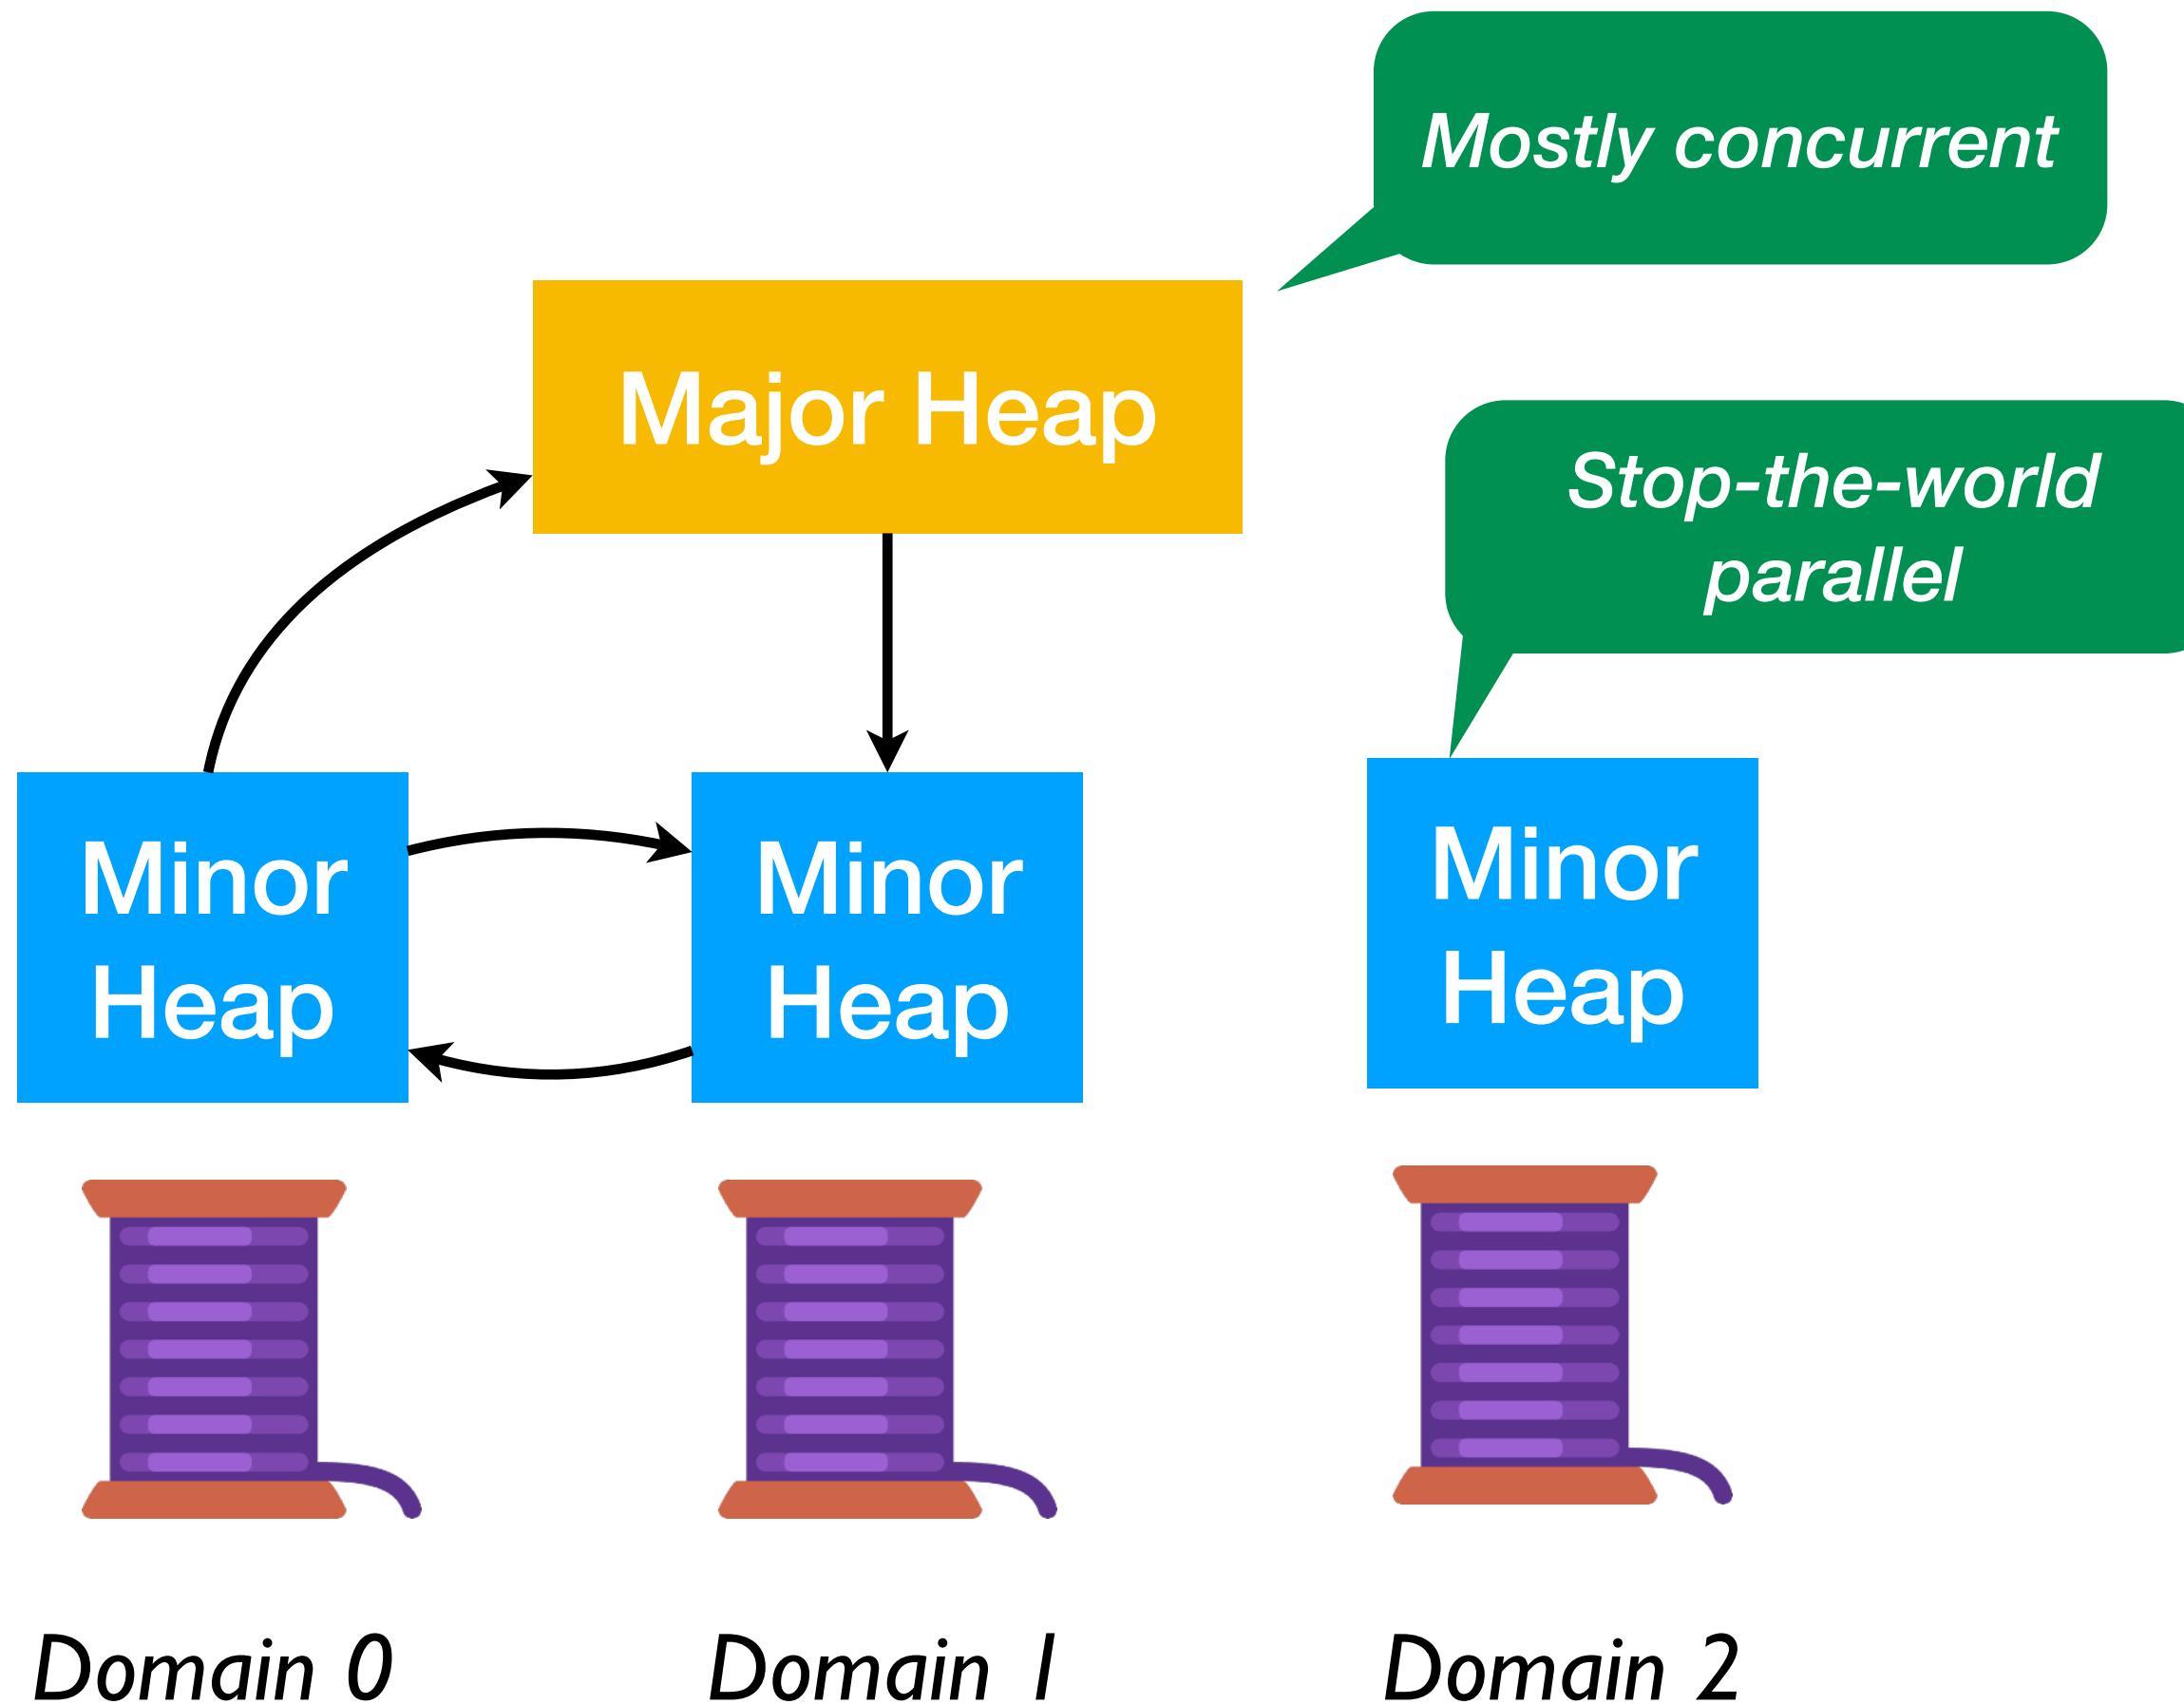
Board size = 1024, Iterations = 512

Cores	Time (Seconds)	Vs Serial
1	24.326	1
2	12.290	1.980
4	6.260	3.890
8	3.238	7.51
16	1.726	14.09
<b>24</b>	<b>1.212</b>	<b>20.07</b>

# Allocation and Collection

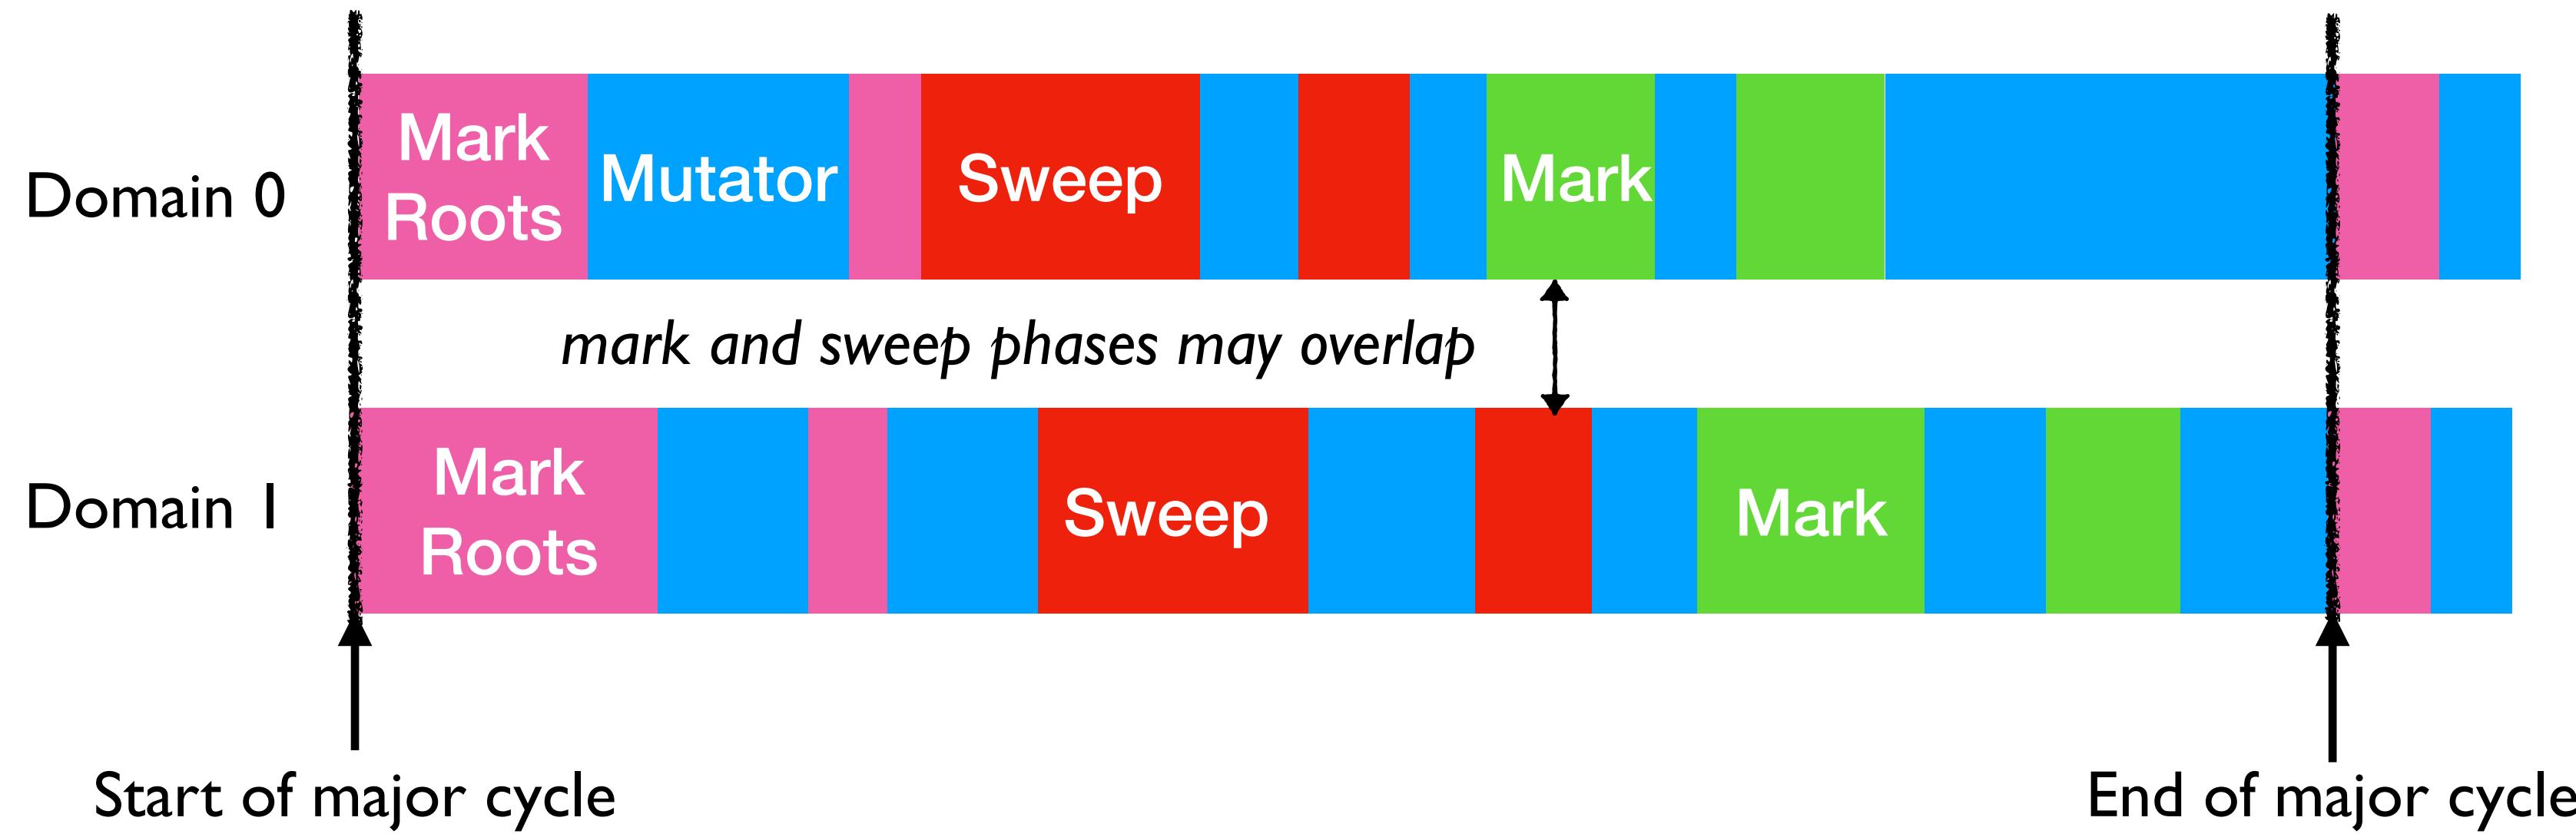


# Allocation and Collection

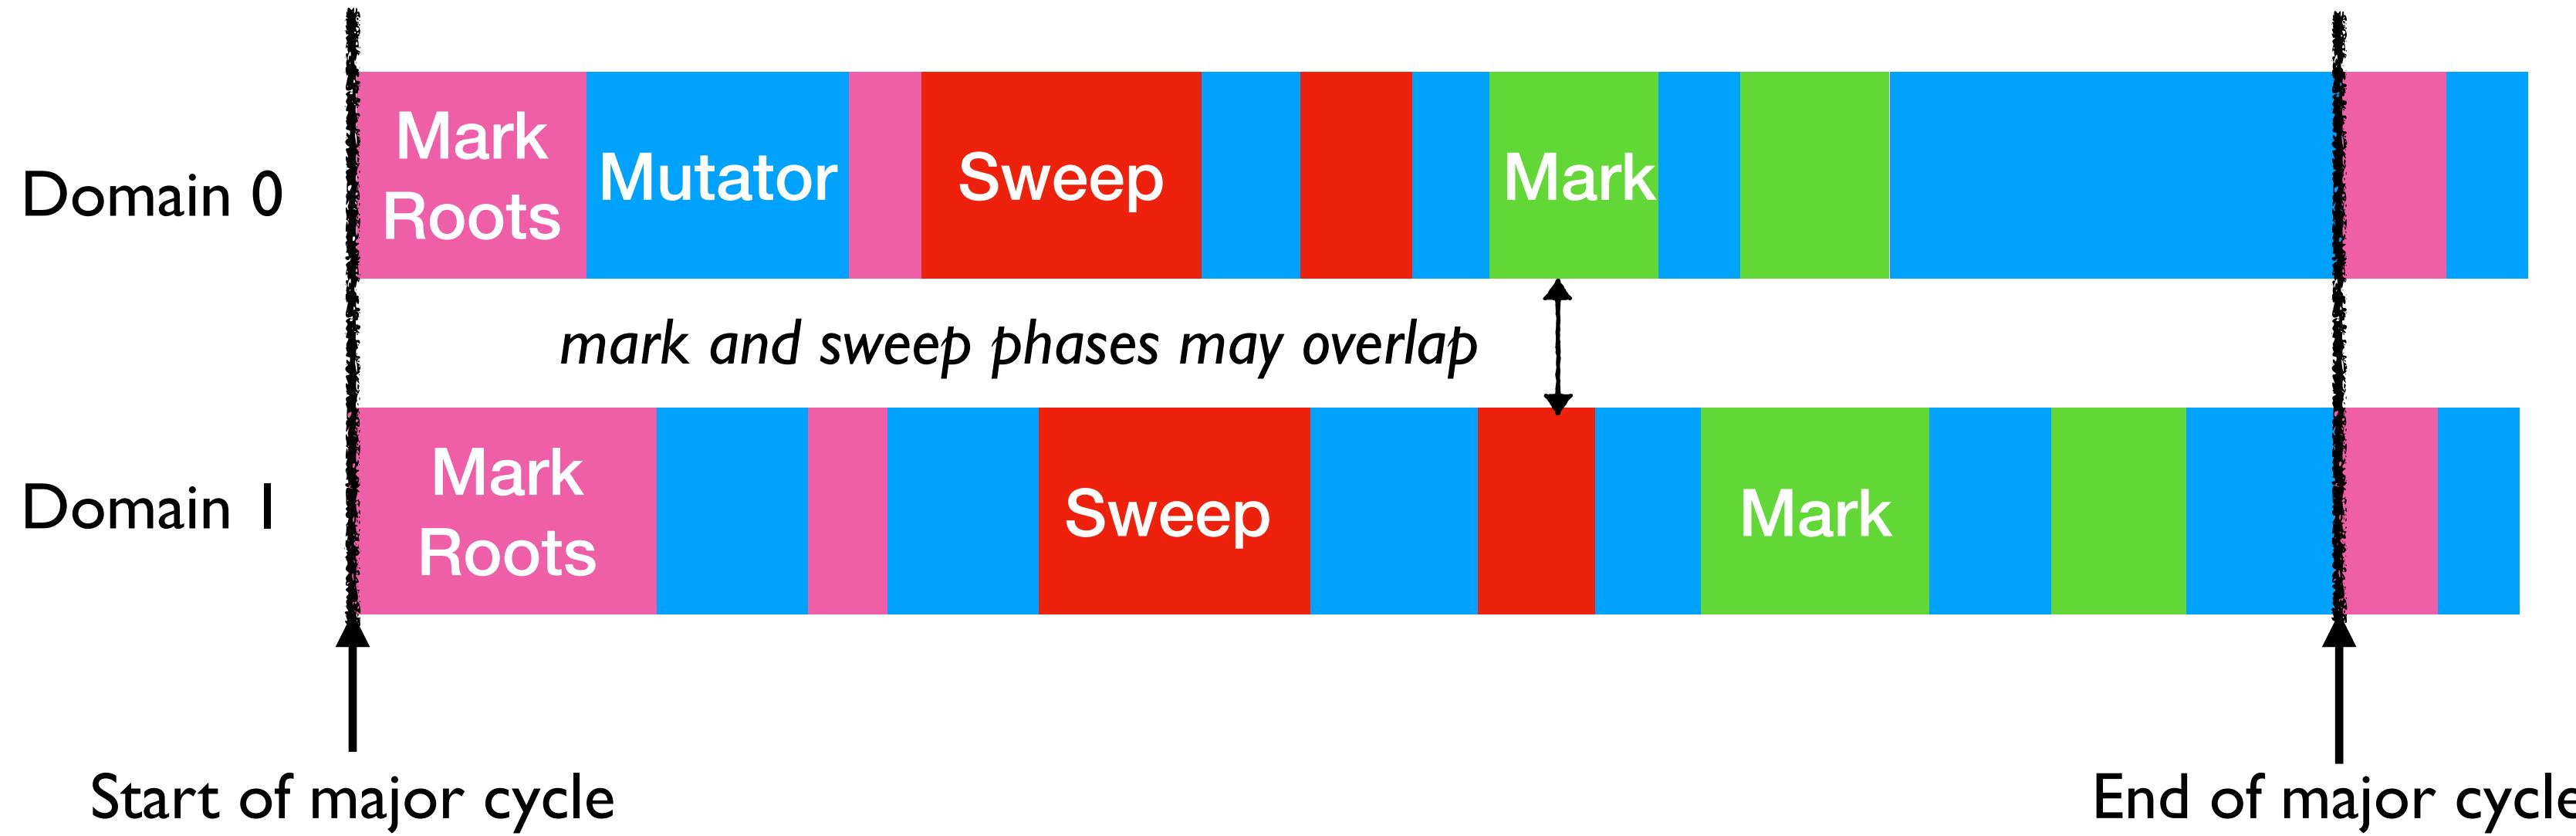


- Minor heap allocations require *no synchronization*
- Major heap allocator is
  - ◆ **Small:** Thread-local, size-segmented free list
  - ◆ **Large:** malloc
- Goal is to match *best-fit* for sequential programs
  - ◆ If we're slower than best-fit, then it is a *performance regression*

# Concurrent GC

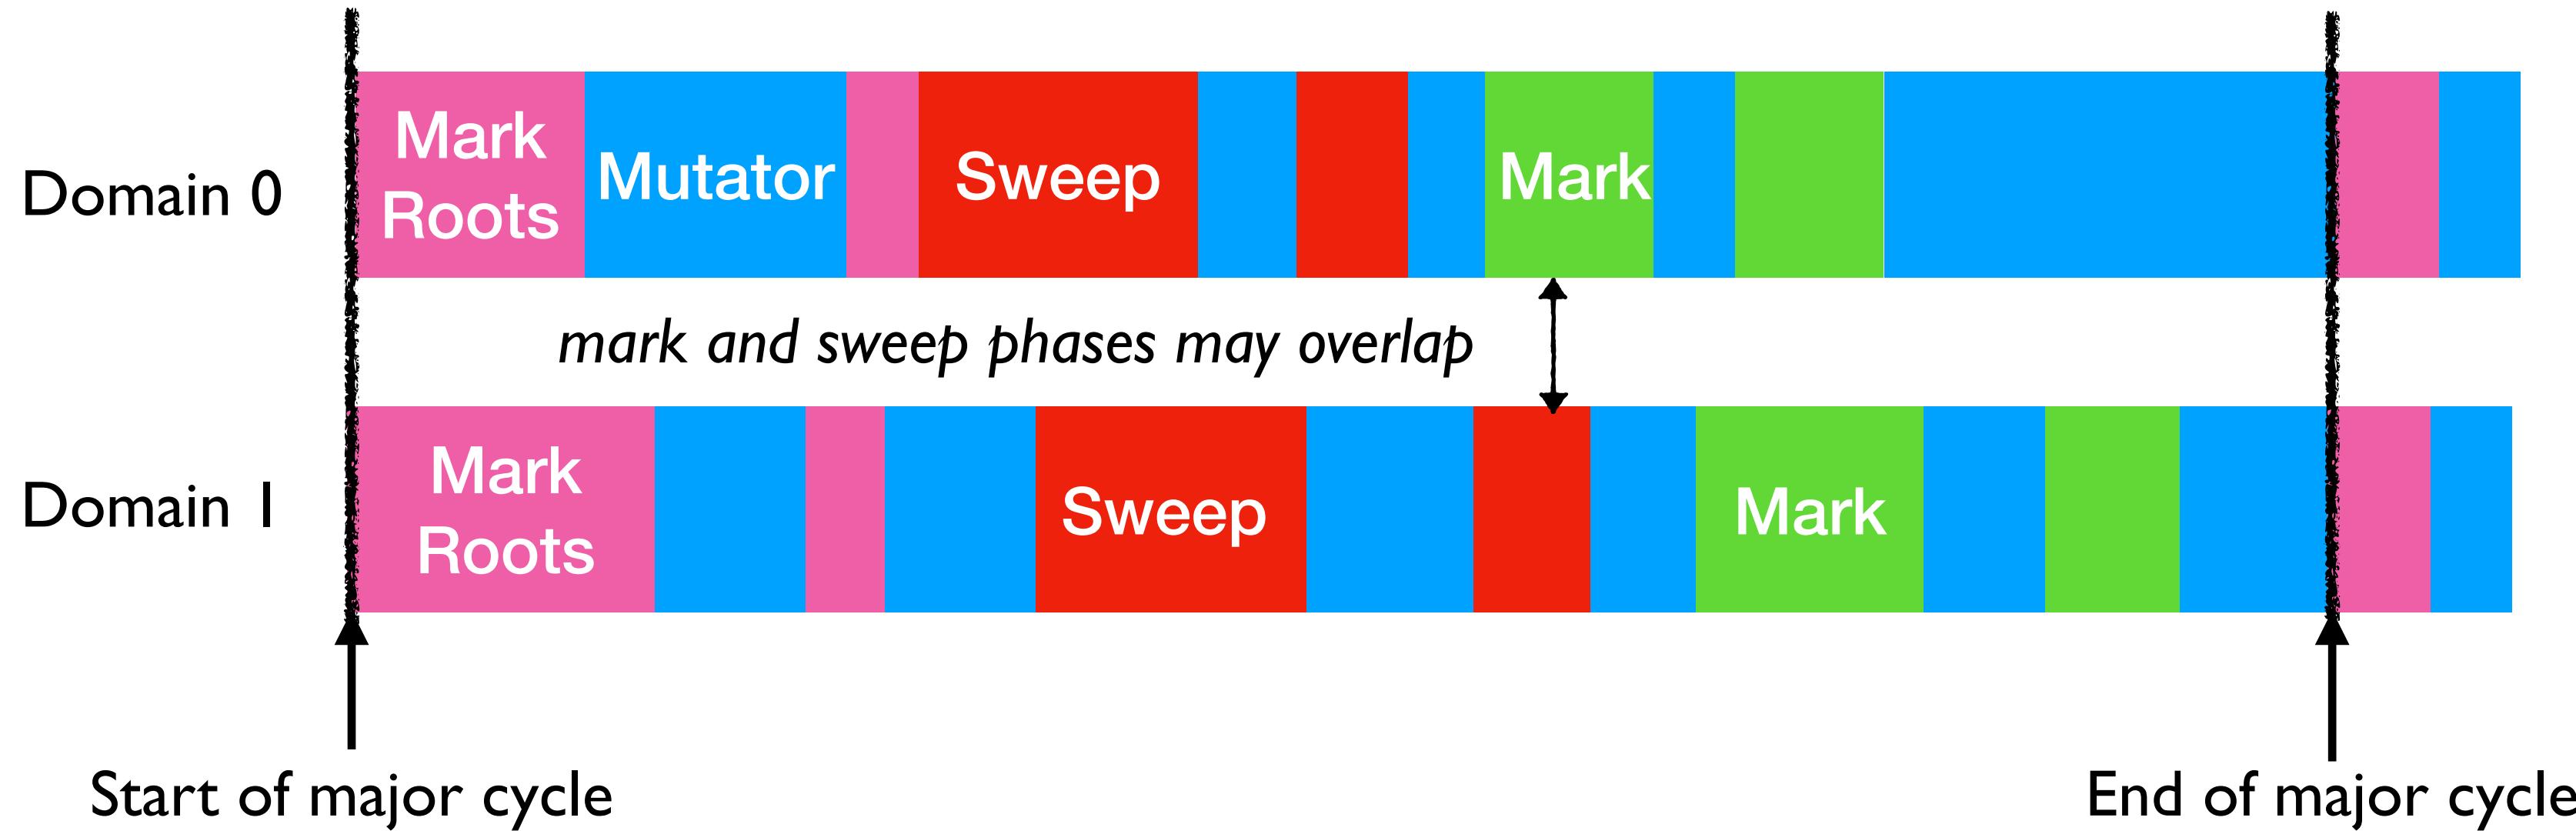


# Concurrent GC



- Stop-the-world parallel *minor* GC + non-moving *major* GC
  - ◆ *Objects don't move while the mutator is running!*

# Concurrent GC



- Stop-the-world parallel *minor* GC + non-moving *major* GC
  - ◆ *Objects don't move while the mutator is running!*
- **No additional rules for the C FFI in OCaml 5.0**
  - ◆ Same rules as OCaml 4.x hold even for parallel programs!

# OCaml memory model

- Simple (*comprehensible!*) operational memory model
  - ◆ Only atomic and non-atomic locations
  - ◆ DRF-SC
  - ◆ No “out of thin air” values
  - ◆ Squeeze at most perf ⇒ write that module in C, C++ or Rust.



# OCaml memory model



- Simple (*comprehensible!*) operational memory model
  - ◆ Only atomic and non-atomic locations
  - ◆ DRF-SC
  - ◆ No “out of thin air” values
  - ◆ Squeeze at most perf  $\Rightarrow$  write that module in C, C++ or Rust.
- **Key innovation:** *Local data race freedom*
  - ◆ Permits compositional reasoning

# OCaml memory model



- Simple (*comprehensible!*) operational memory model
  - ◆ Only atomic and non-atomic locations
  - ◆ DRF-SC
  - ◆ No “out of thin air” values
  - ◆ Squeeze at most perf  $\Rightarrow$  write that module in C, C++ or Rust.
- **Key innovation:** *Local data race freedom*
  - ◆ Permits compositional reasoning
- Performance impact
  - ◆ Free on x86 and < 1% on ARM

# OCaml memory model

- Simple (*comprehensible!*) operational memory model

## Bounding Data Races in Space and Time

(Extended version, with appendices)

PLDI '18

Stephen Dolan  
University of Cambridge, UK

KC Sivaramakrishnan  
University of Cambridge, UK

Anil Madhavapeddy  
University of Cambridge, UK

### Abstract

We propose a new semantics for shared-memory parallel programs that gives strong guarantees even in the presence of data races. Our *local data race freedom* property guarantees that all data-race-free portions of programs exhibit

The primary reasoning tools provided to programmers by these models are the *data-race-freedom (DRF) theorems*. Programmers are required to mark as *atomic* all variables used for synchronisation between threads, and to avoid *data races*, which are concurrent accesses (except concurrent reads) to

- Performance impact
  - ◆ Free on x86 and < 1% on ARM

19

# OCaml memory model

- PLDI '18 paper only formalised compilation to *hardware memory models*
  - ◆ Omitted object initialisation

# OCaml memory model

- PLDI '18 paper only formalised compilation to *hardware memory models*
  - ◆ Omitted object initialisation
- OCaml 5.0 extended the work to cover
  - ◆ Object initialisation
  - ◆ Compilation to CII memory model

# OCaml memory model

- PLDI '18 paper only formalised compilation to *hardware memory models*
  - ◆ Omitted object initialisation
- OCaml 5.0 extended the work to cover
  - ◆ Object initialisation
  - ◆ Compilation to CII memory model
- C FFI has been made stronger (by making the access `volatile`)

```
#define Field(x, i) (((volatile value *)(x)) [I])  
  
void caml_modify (volatile value *, value);  
  
void caml_initialize (volatile value *, value);
```

- ◆ Assumes Linux Kernel Memory Model (LKMM)
- ◆ Does not break code

# OCaml memory model

- **C FFI also respects LDRF!**

# OCaml memory model

- **C FFI also respects LDRF!**

```
let msg = ref 0
let flag = Atomic.make false

let t1 =
  msg := 1;
  Atomic.set flag true

let t2 =
  let rf = Atomic.get flag in
  let rm = !msg in
  assert (not (rf = true && rm = 0))
```

# OCaml memory model

- **C FFI also respects LDRF!**

```
let msg = ref 0
let flag = Atomic.make false

let t1 =
  msg := 1;
  Atomic.set flag true
```

```
let t2 =
  let rf = Atomic.get flag in
  let rm = !msg in
  assert (not (rf = true && rm = 0))
```

```
/* t1 implemented in C */
void t1 (value msg, value flag) {
  caml_modify (&Field(msg,0), Val_int(1));
  caml_atomic_exchange (flag, Val_true);
}
```

# ThreadSanitizer

```
1  type t = { mutable x : int }
2
3  let v = { x = 0 }
4
5  let () =
6    let t1 = Domain.spawn (fun () -> v.x <- 10) in
7    let t2 = Domain.spawn (fun () -> Unix.sleep v.x) in
8    Domain.join t1;
9    Domain.join t2
```

# ThreadSanitizer

```
1  type t = { mutable x : int }
2
3  let v = { x = 0 }
4
5  let () =
6    let t1 = Domain.spawn (fun () -> v.x <- 10) in
7    let t2 = Domain.spawn (fun () -> Unix.sleep v.x) in
8    Domain.join t1;
9    Domain.join t2
```

WARNING: ThreadSanitizer: data race (pid=502344)

Read of size 8 at 0x7fc0b15fe458 by thread T4 (mutexes: write M0):

```
#0 camlDune__exe__Simple_race__fun_600 /workspace_root/simple_race.ml:7 (simple_race.exe+0x51e9b1)
#1 caml_callback ???:? (simple_race.exe+0x5777f0)
#2 domain_thread_func domain.c:? (simple_race.exe+0x57b8fc)
```

Previous write of size 8 at 0x7fc0b15fe458 by thread T1 (mutexes: write M1):

```
#0 camlDune__exe__Simple_race__fun_596 /workspace_root/simple_race.ml:6 (simple_race.exe+0x51e971)
#1 caml_callback ???:? (simple_race.exe+0x5777f0)
#2 domain_thread_func domain.c:? (simple_race.exe+0x57b8fc)
```

# Effect handlers

- Structured programming with delimited continuations
- No effect system, no dedicated syntax
- Provides both *deep* and *shallow* handlers

# Effect handlers

- Structured programming with delimited continuations
- No effect system, no dedicated syntax
- Provides both *deep* and *shallow* handlers

**Example prints “0 1 2 3 4”**

```
1  open Effect
2  open Effect.Deep
3
4  type _ Effect.t += E : string t
5
6  let comp () =
7    print_string "0 ";
8    print_string (perform E);
9    print_string "3 "
10
11 let main () =
12  try_with comp ()
13  { effc = fun (type a) (e : a t) ->
14    match e with
15    | E -> Some (fun (k : (a,_) continuation) ->
16      print_string "1 ";
17      continue k "2 ";
18      print_string "4 ")
19    | e -> None }
```

# Effect handlers

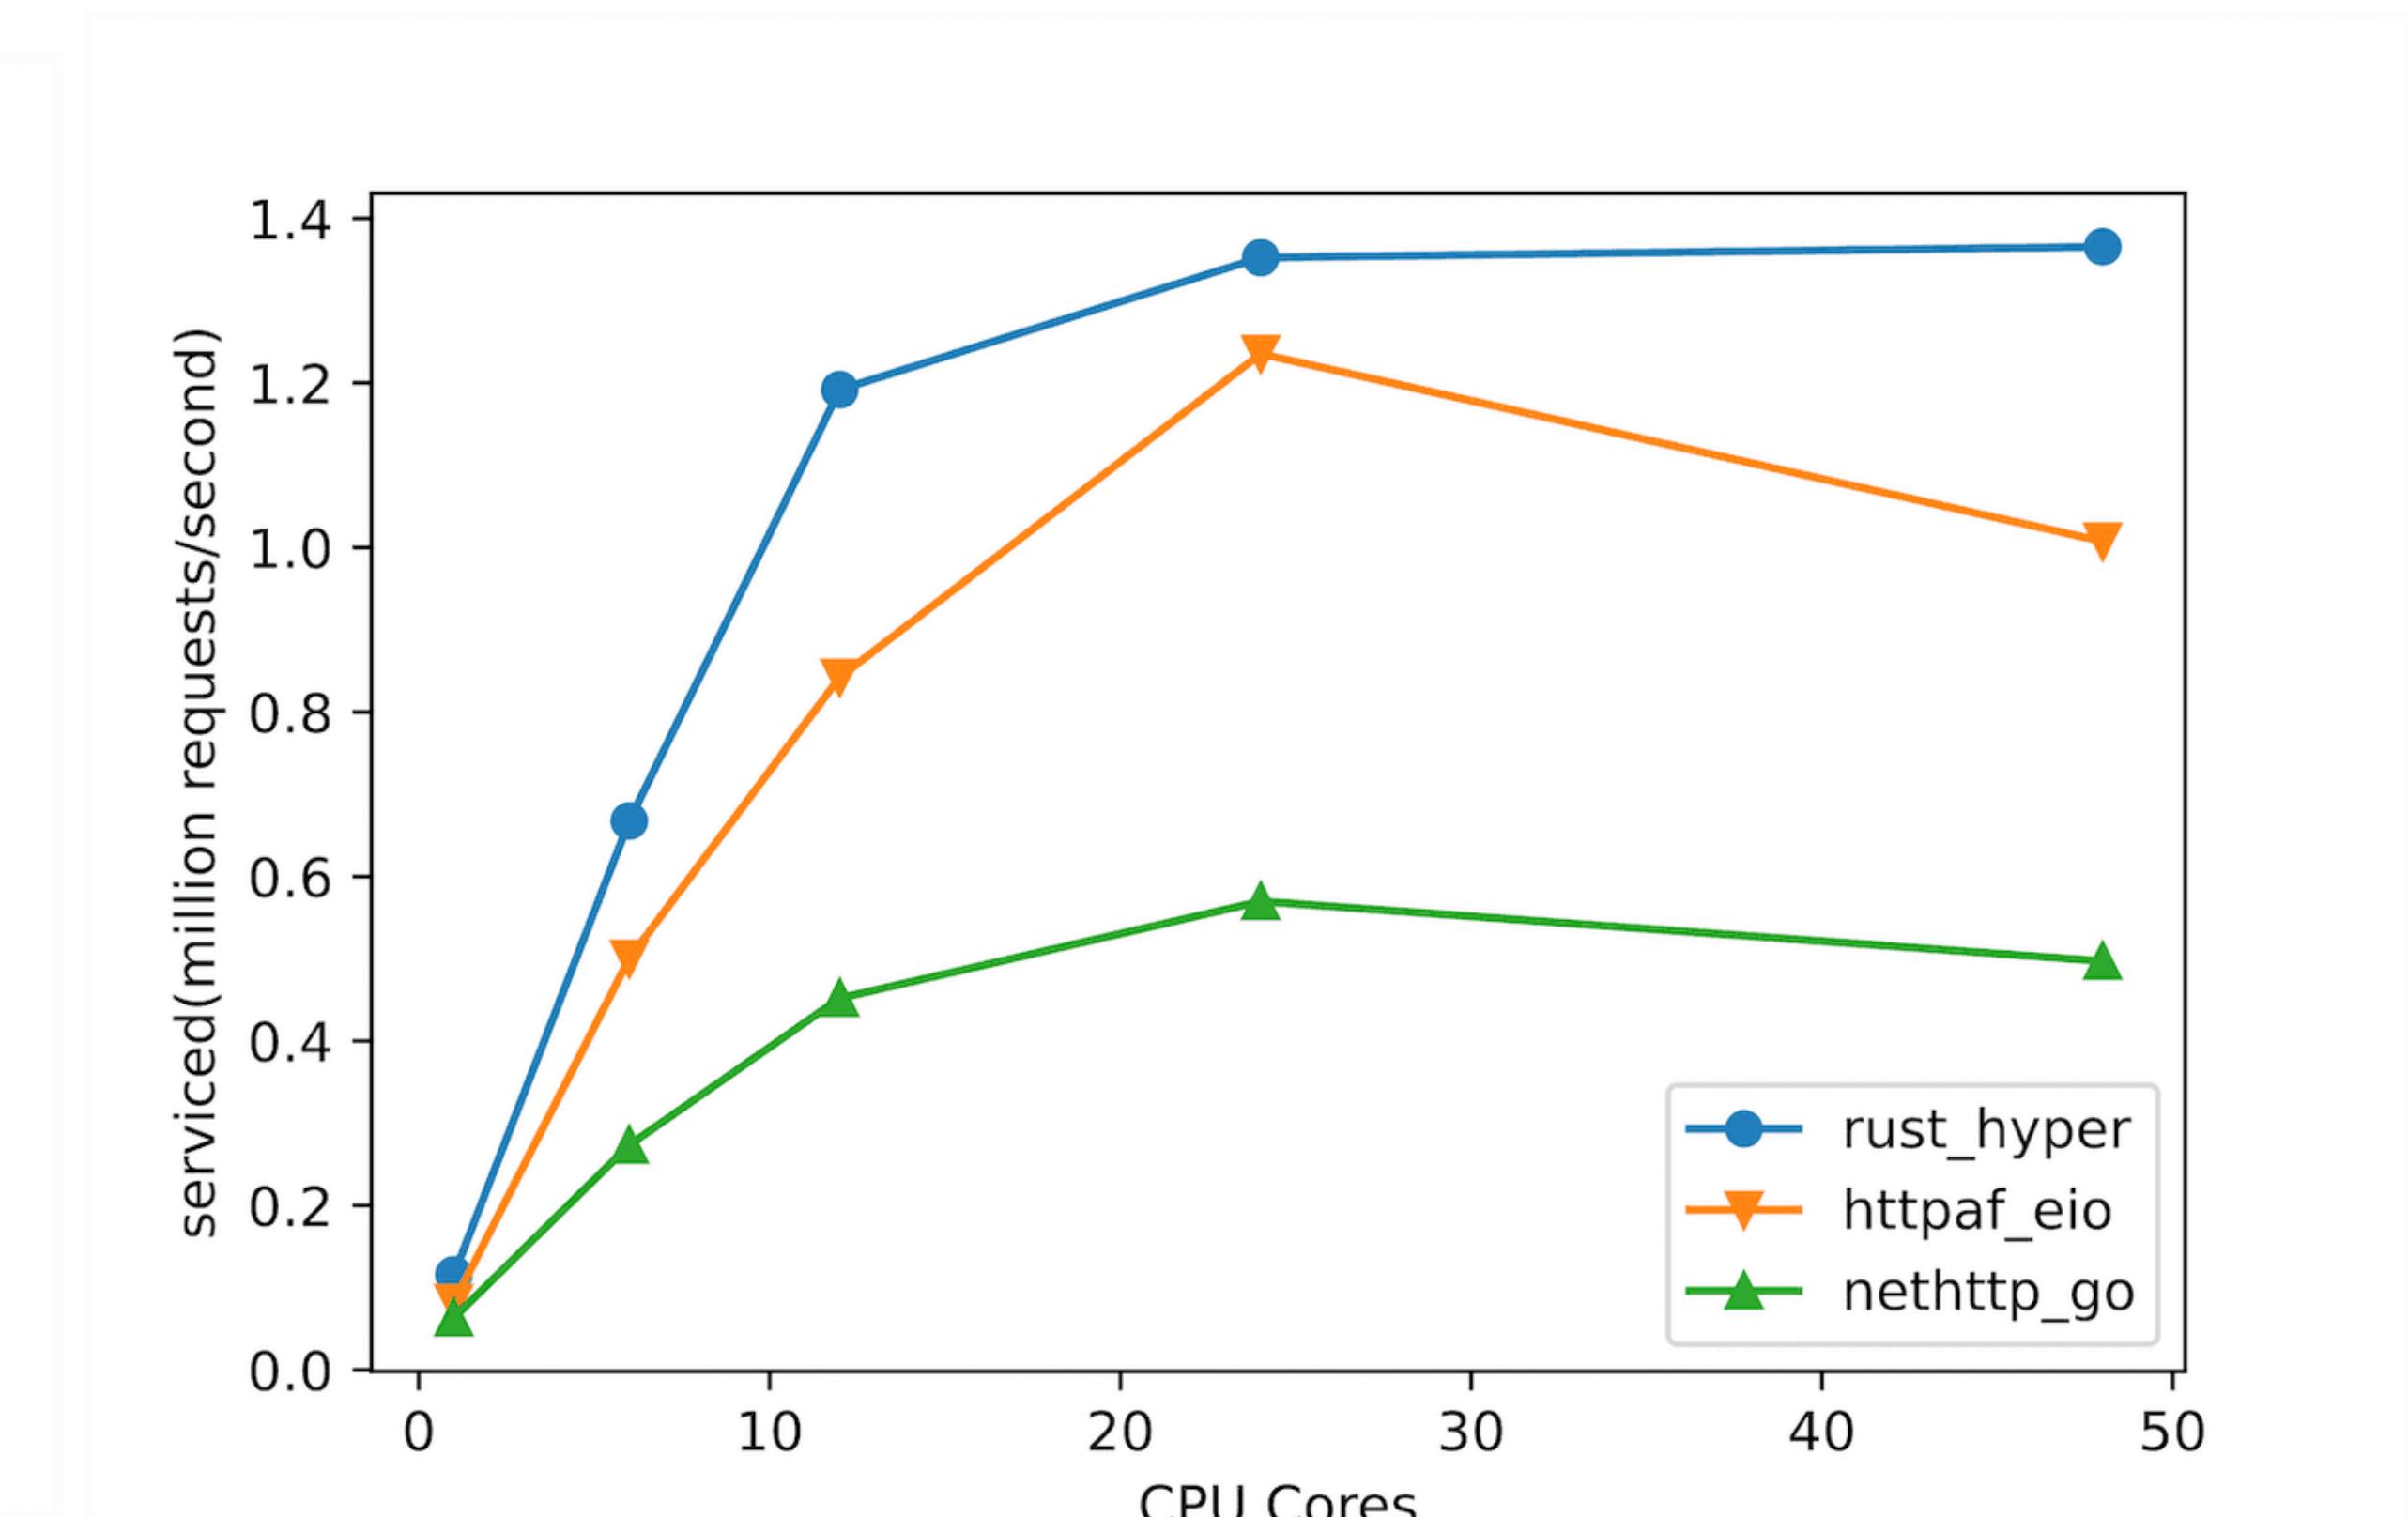
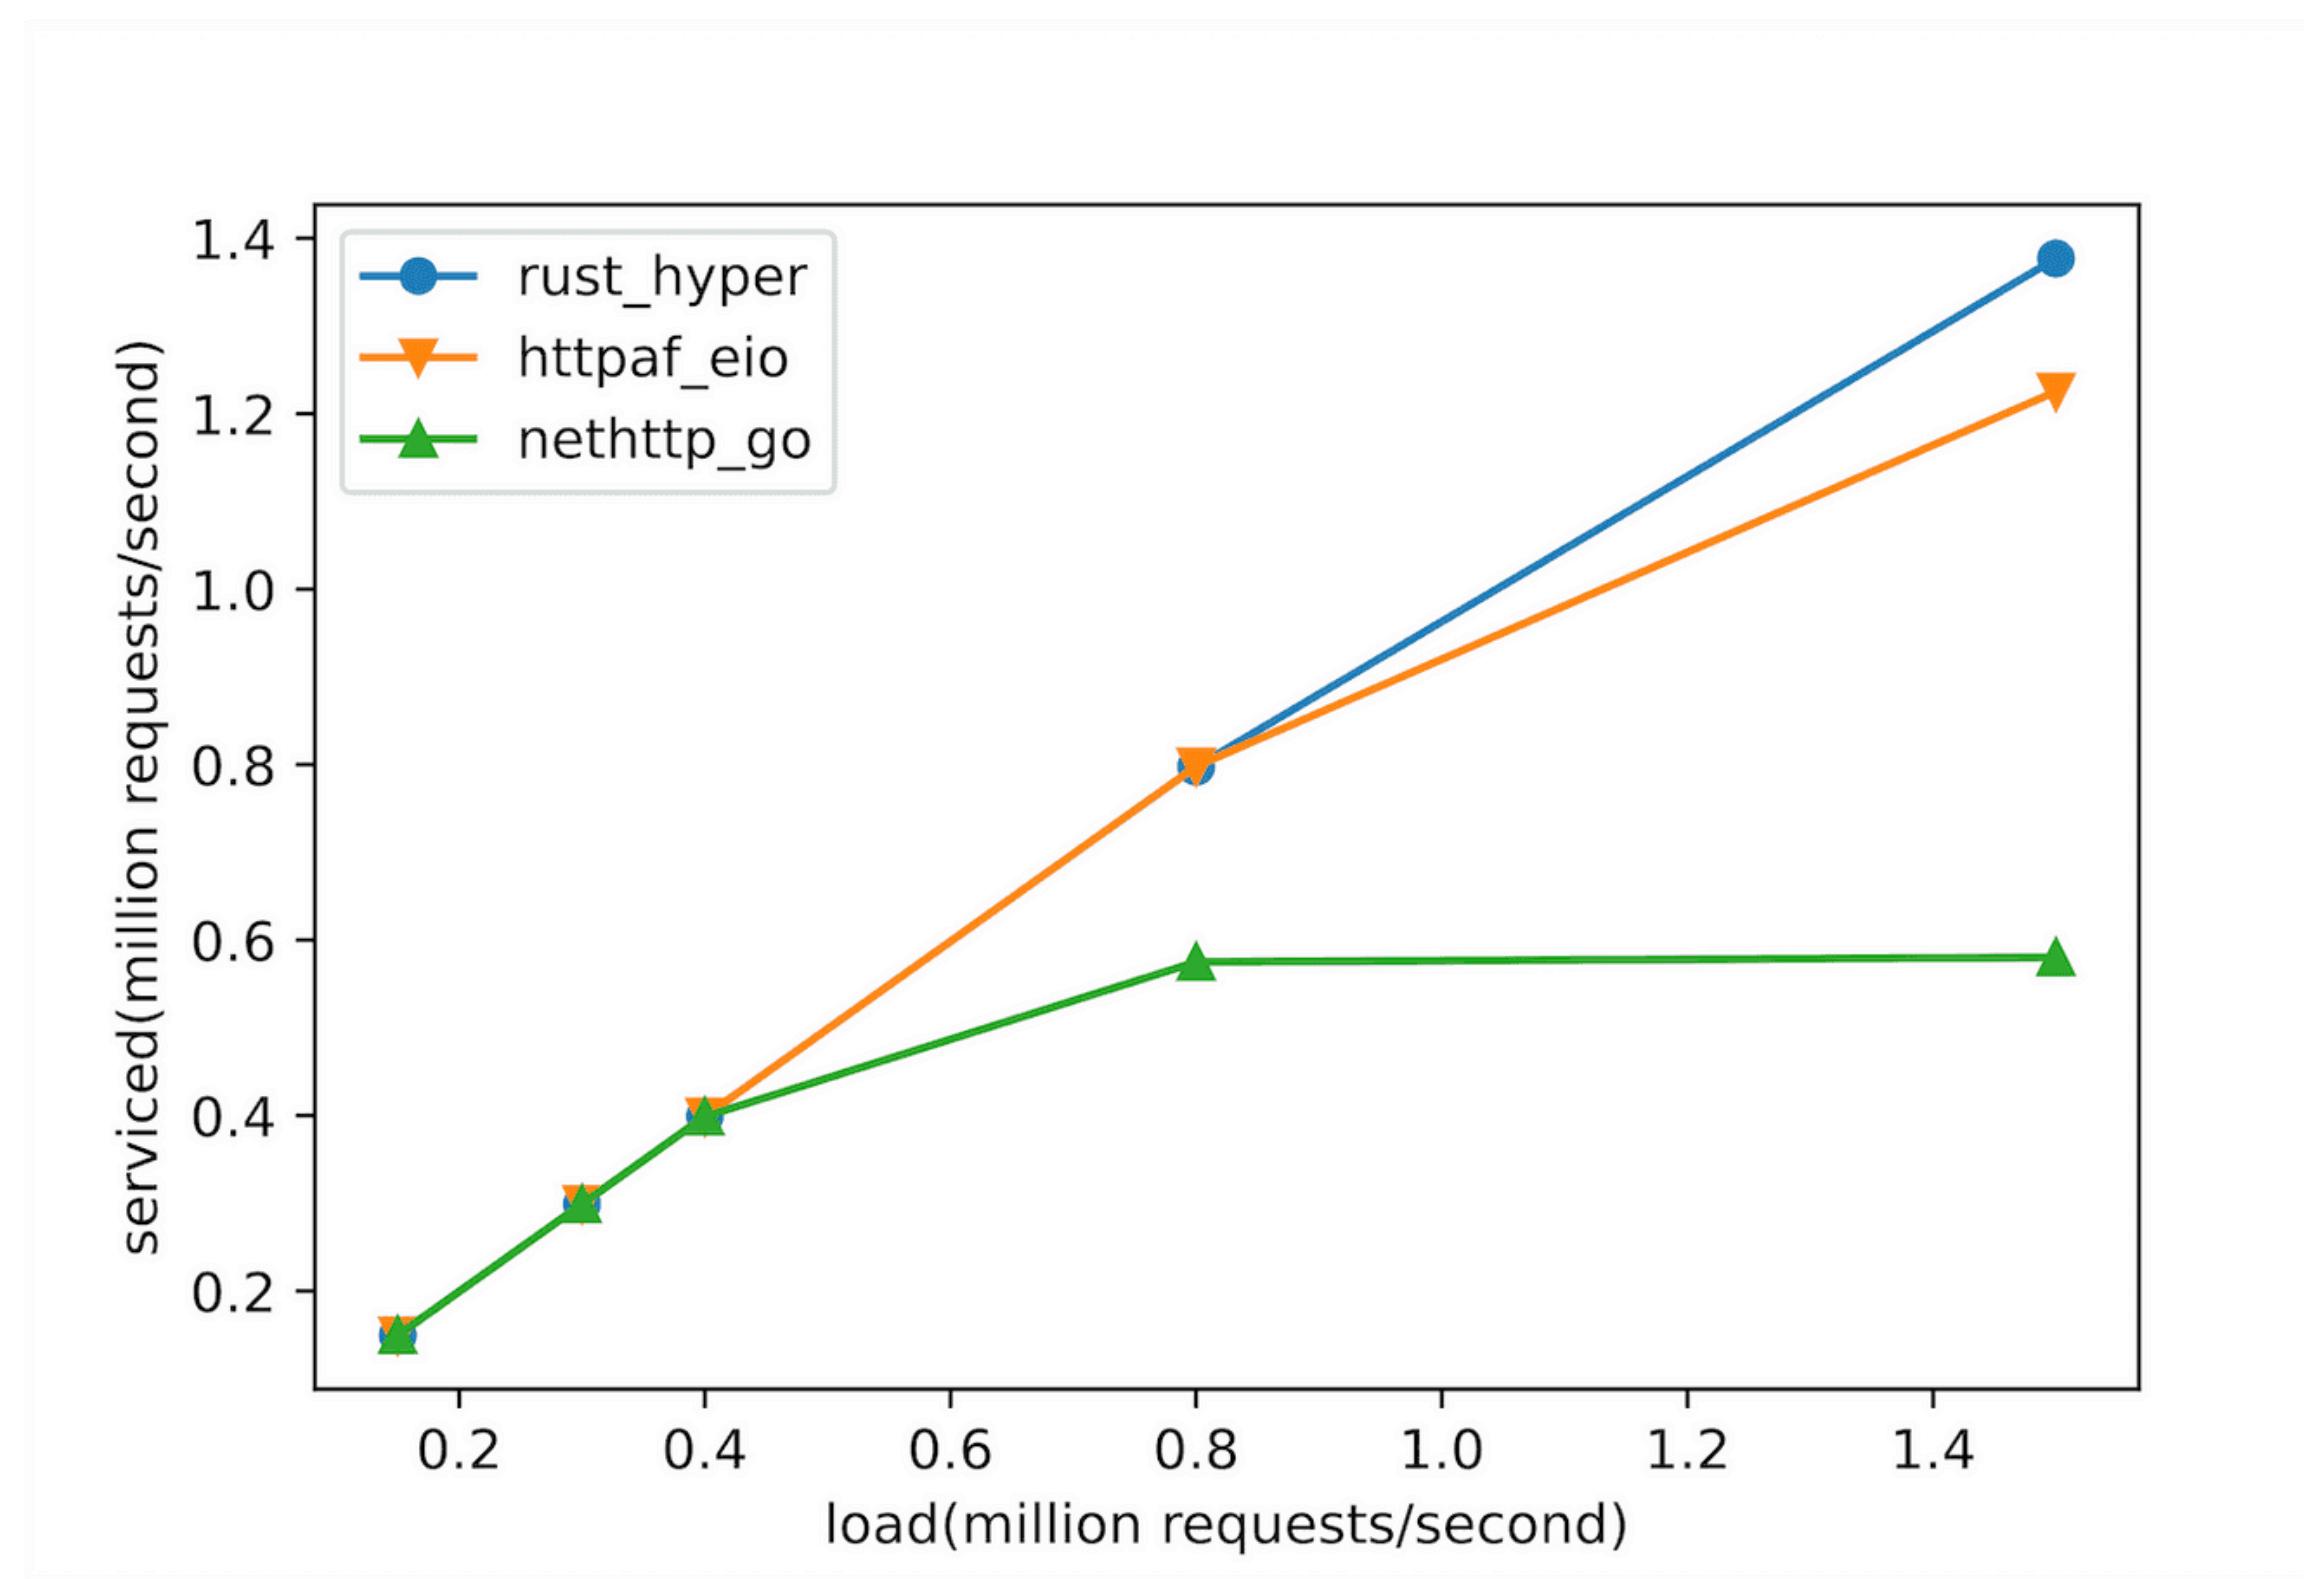
- Structured programming with delimited continuations
- No effect system, no dedicated syntax
- Provides both *deep* and *shallow* handlers

**Example prints “0 1 2 3 4”**

- *Same type safety* as the earlier syntactic version

```
1  open Effect
2  open Effect.Deep
3
4  type _ Effect.t += E : string t
5
6  let comp () =
7    print_string "0 ";
8    print_string (perform E);
9    print_string "3 "
10
11 let main () =
12  try_with comp ()
13  { effc = fun (type a) (e : a t) ->
14    match e with
15    | E -> Some (fun (k : (a,_) continuation) ->
16                  print_string "1 ";
17                  continue k "2 ";
18                  print_string "4 ")
19    | e -> None }
```

# Eio — Direct-style effect-based concurrency



# Integration with Lwt & Async



- Lwt\_eio allows running Lwt and Eio code together
  - ◆ Only sequential
  - ◆ Cancellation semantics is also integrated
  - ◆ Incrementally port Lwt applications to Eio

# Integration with Lwt & Async



- Lwt\_eio allows running Lwt and Eio code together
  - ◆ Only sequential
  - ◆ Cancellation semantics is also integrated
  - ◆ Incrementally port Lwt applications to Eio
- *Very experimental* Async\_eio running Async and Eio code together
  - ◆ Required changes to Async

# Merge Process

- Multicore OCaml was maintained as a separate *fork* of the compiler
  - ◆ Multiple *tricky* rebases to keep the fork up to date with trunk

# Merge Process

- Multicore OCaml was maintained as a separate *fork* of the compiler
  - ◆ Multiple *tricky* rebases to keep the fork up to date with trunk
- *Single PR* to merge multicore change
  - ◆ Not worth splitting into multiple PR — context loss

# Merge Process

- Multicore OCaml was maintained as a separate *fork* of the compiler
  - ◆ Multiple *tricky* rebases to keep the fork up to date with trunk
- *Single PR* to merge multicore change
  - ◆ Not worth splitting into multiple PR — context loss
- Asynchronous & Synchronous review phases (Nov 2021)

Working Group	Multicore Lead	Dev Team Lead
GC	Tom Kelly & Sadiq Jaffer	Damien Doligez
Domains	Tom Kelly	Luc Maranget
Runtime multi-domain safety	Enguerrand Decorne	Xavier Leroy
Stdlib changes	KC Sivaramakrishnan	Florian Angeletti (and Gabriel Scherer)
Fibers	KC Sivaramakrishnan	Damien Doligez (and Xavier Leroy)

# Merge Process

The screenshot shows a GitHub pull request page for the repository `ocaml/ocaml`. The pull request is titled `Multicore OCaml #10831`. The status is `Merged`, and it was merged by `xavierleroy` on 10 Jan. The pull request has 250 commits, 9 checks, and 573 files changed, with a total of `+22,955 -14,062` lines of code. A comment from `kayceesrk` on 21 Dec 2021 describes the changes, mentioning support for shared-memory parallelism and direct-style concurrency. The pull request has been reviewed by `damiendoligez`, `avsm`, `Engil`, `gasche`, and `sadiqj`. The page also shows the repository's main statistics: 268 issues, 258 pull requests, and 4,103 commits in the `ocaml:trunk` branch.

ocaml / ocaml Public Edit Pins Unwatch 199 Fork 920 Starred 4k

Code Issues 268 Pull requests 258 Actions Projects Security Insights

## Multicore OCaml #10831

Merged xavierleroy merged 4,103 commits into `ocaml:trunk` from `ocaml-multicore:multicore-pr` on 10 Jan

Conversation 393 Commits 250 Checks 9 Files changed 573 +22,955 -14,062

**kayceesrk** commented on 21 Dec 2021 · edited Member

This PR adds support for shared-memory parallelism through domains and direct-style concurrency through effect handlers (without syntactic support). It intends to have backwards compatibility in terms of language features, C API, and also the performance of single-threaded code.

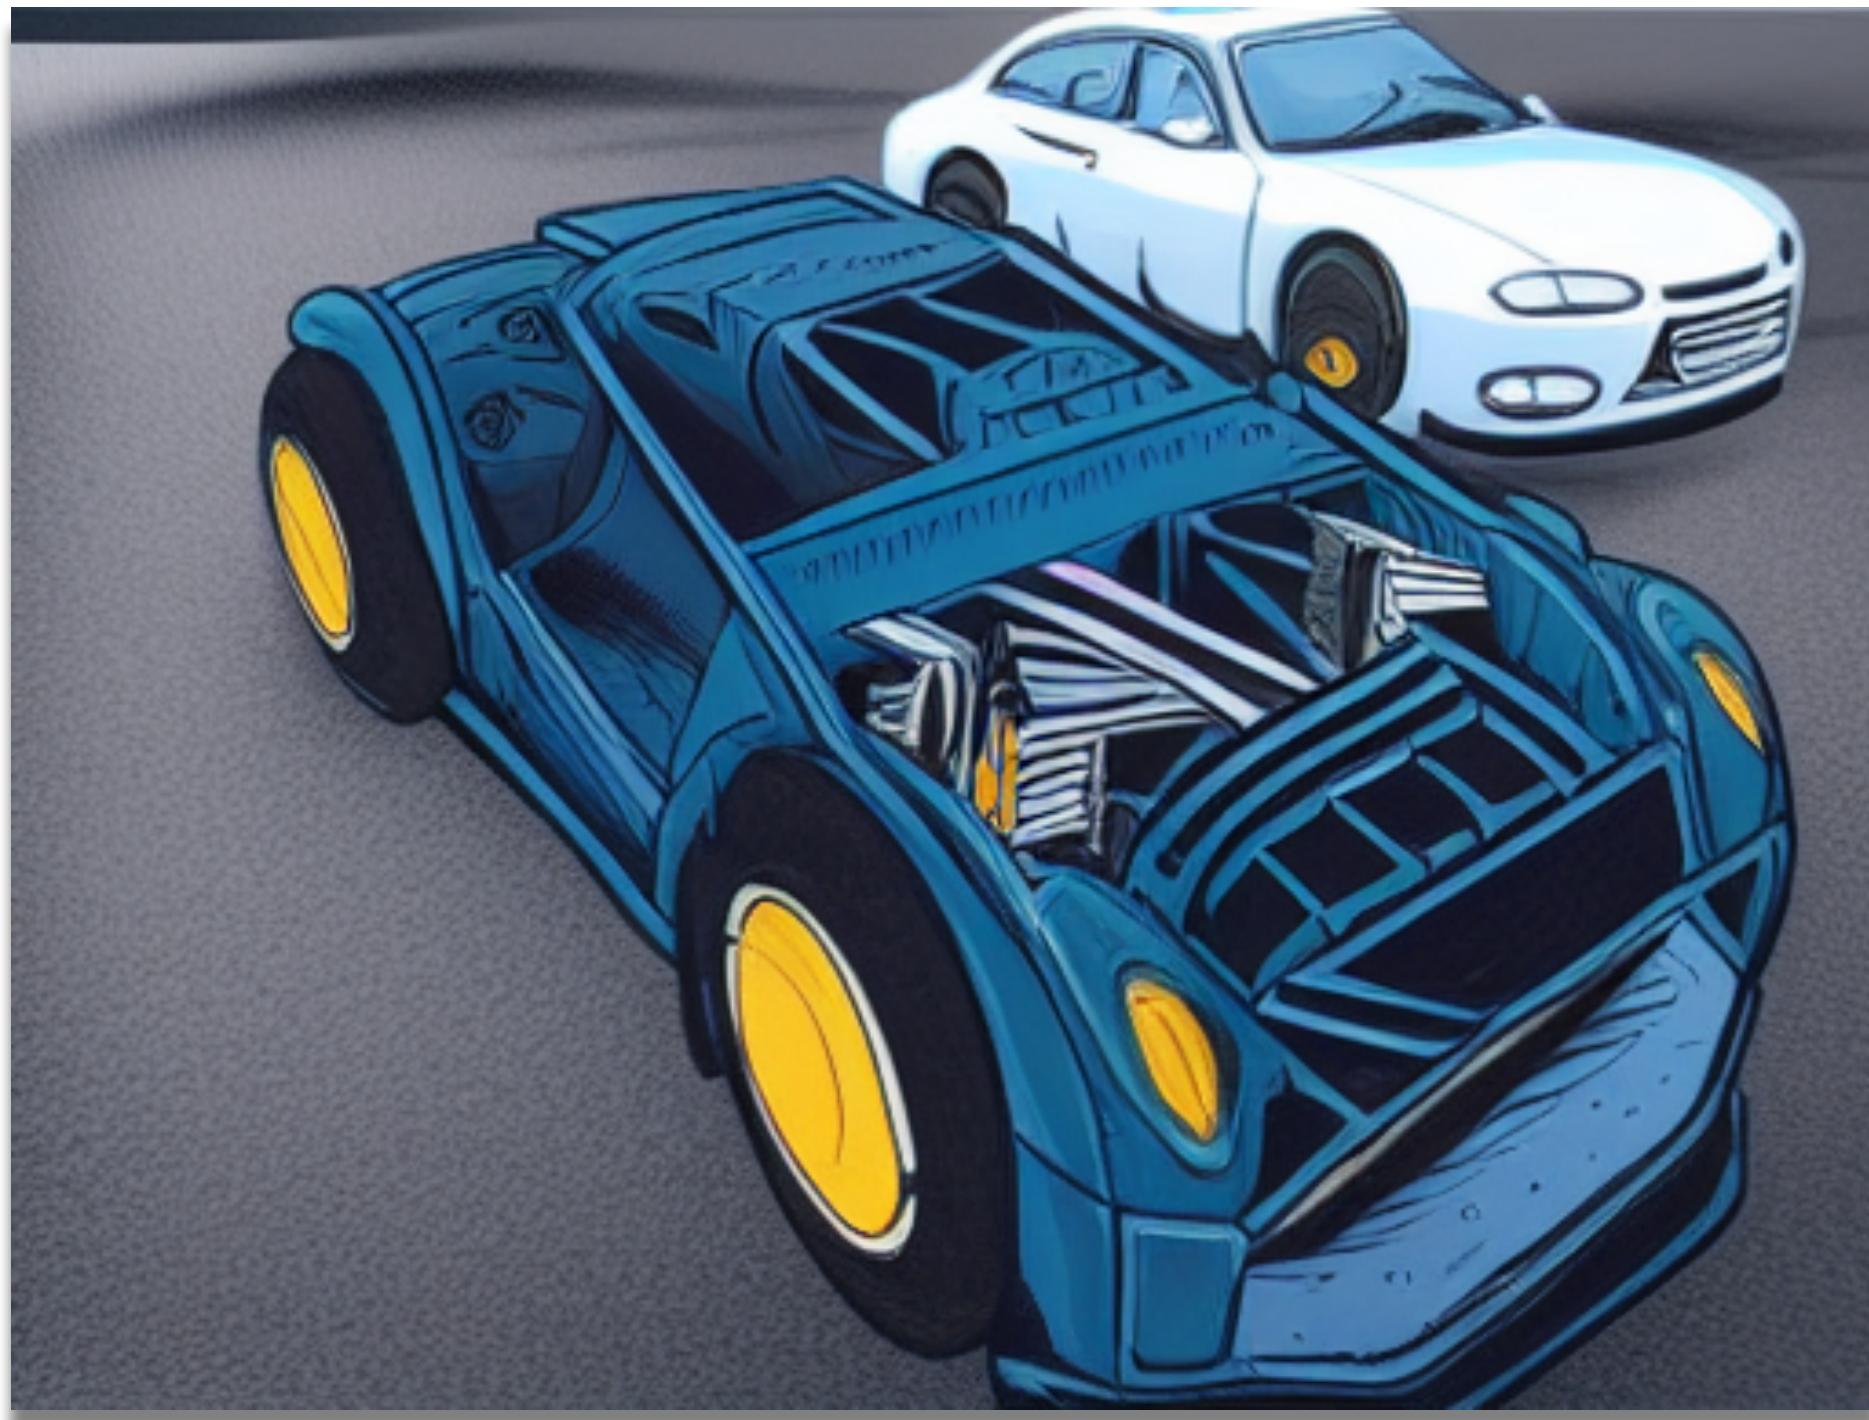
**For users**

Reviewers

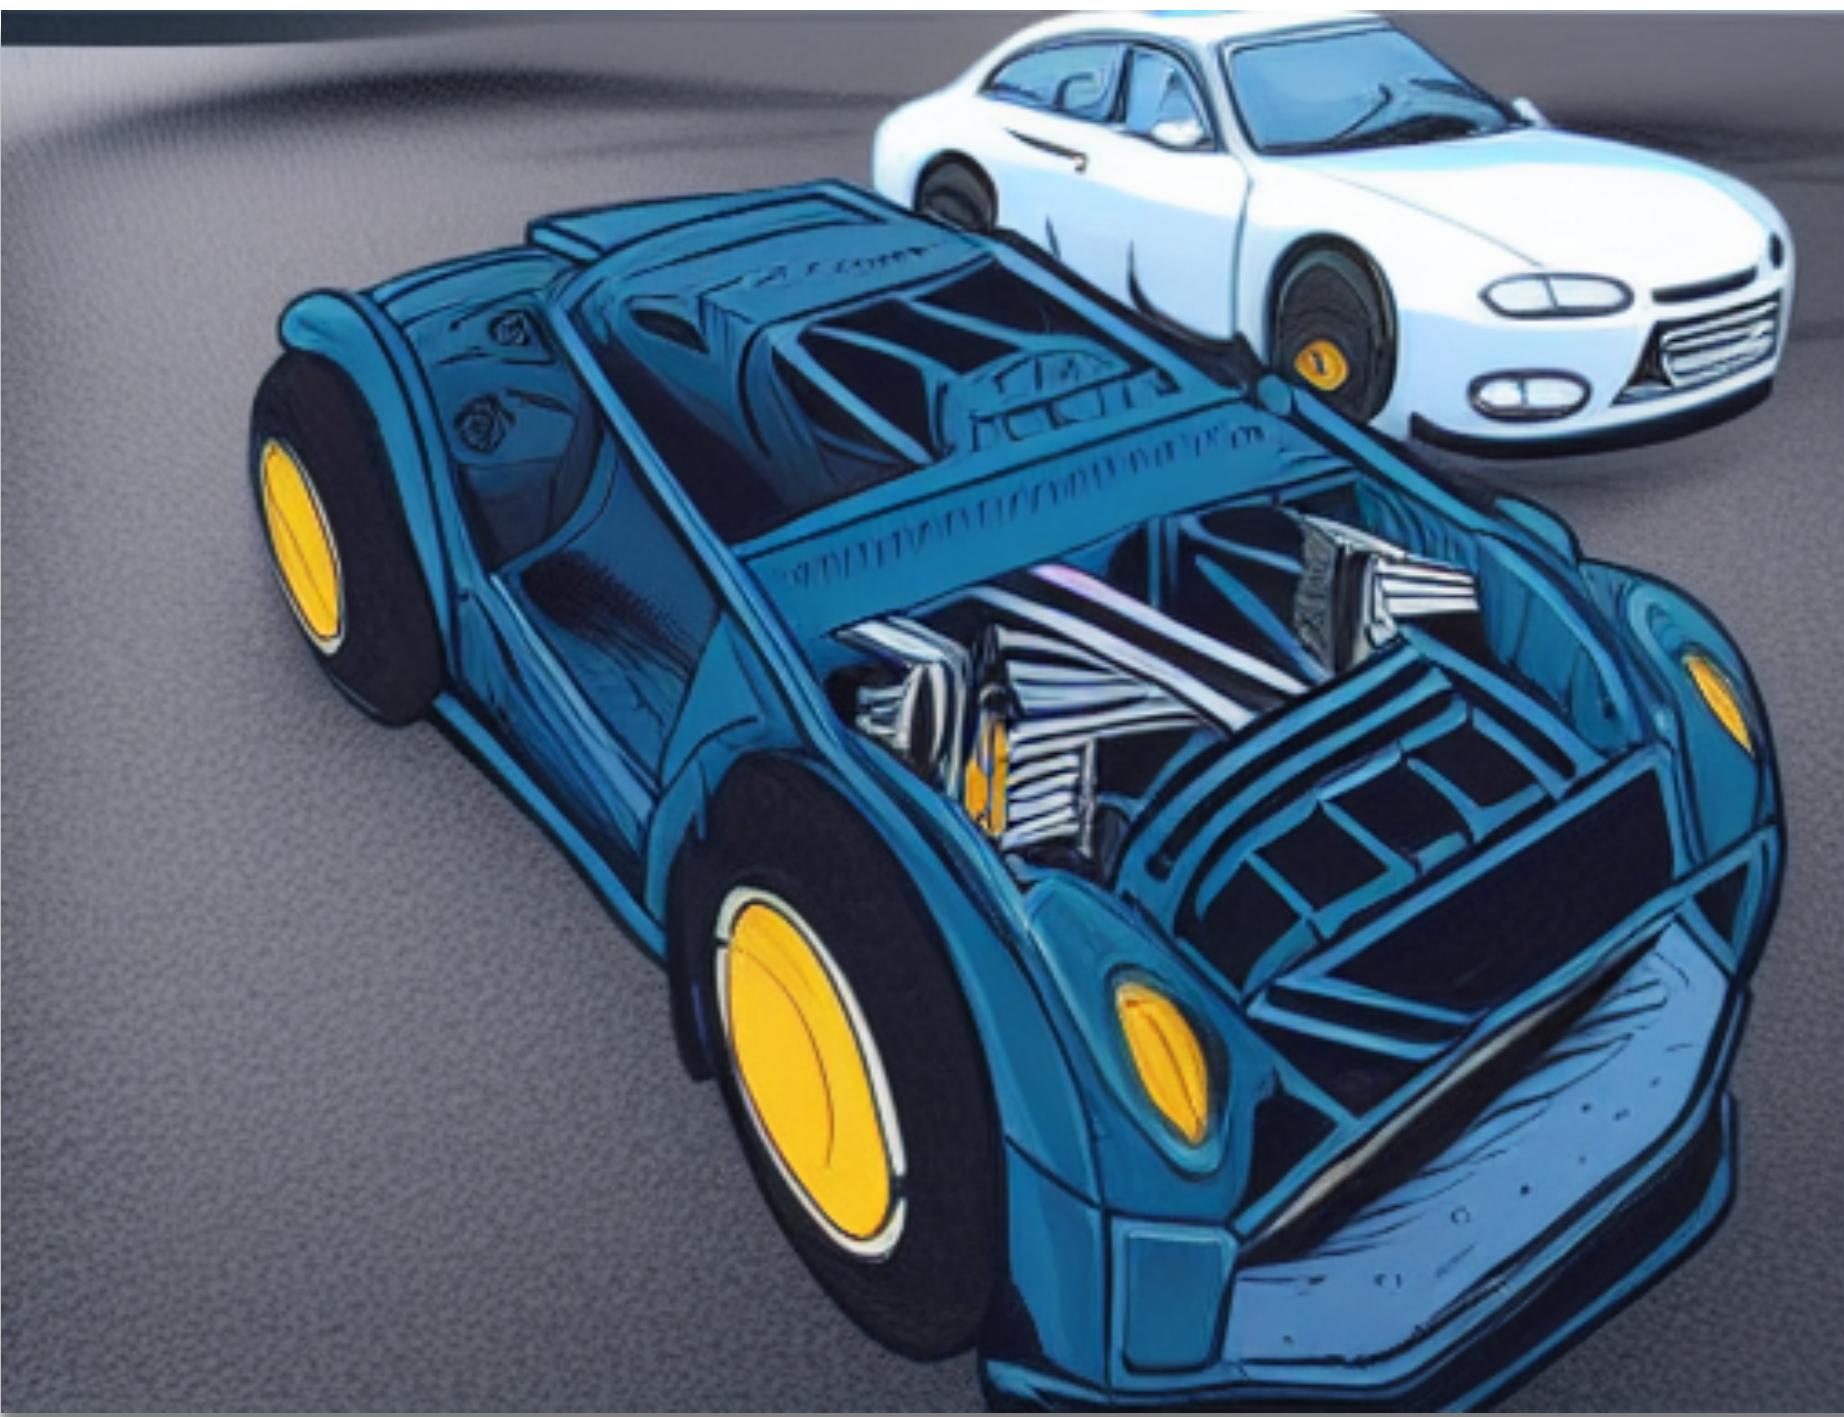
- damiendoligez
- avsm
- Engil
- gasche
- sadiqj

# OCaml 5.0 — an MVP release

- Many features broken and are being added back
  - ◆ This will continue after 5.0 gets released



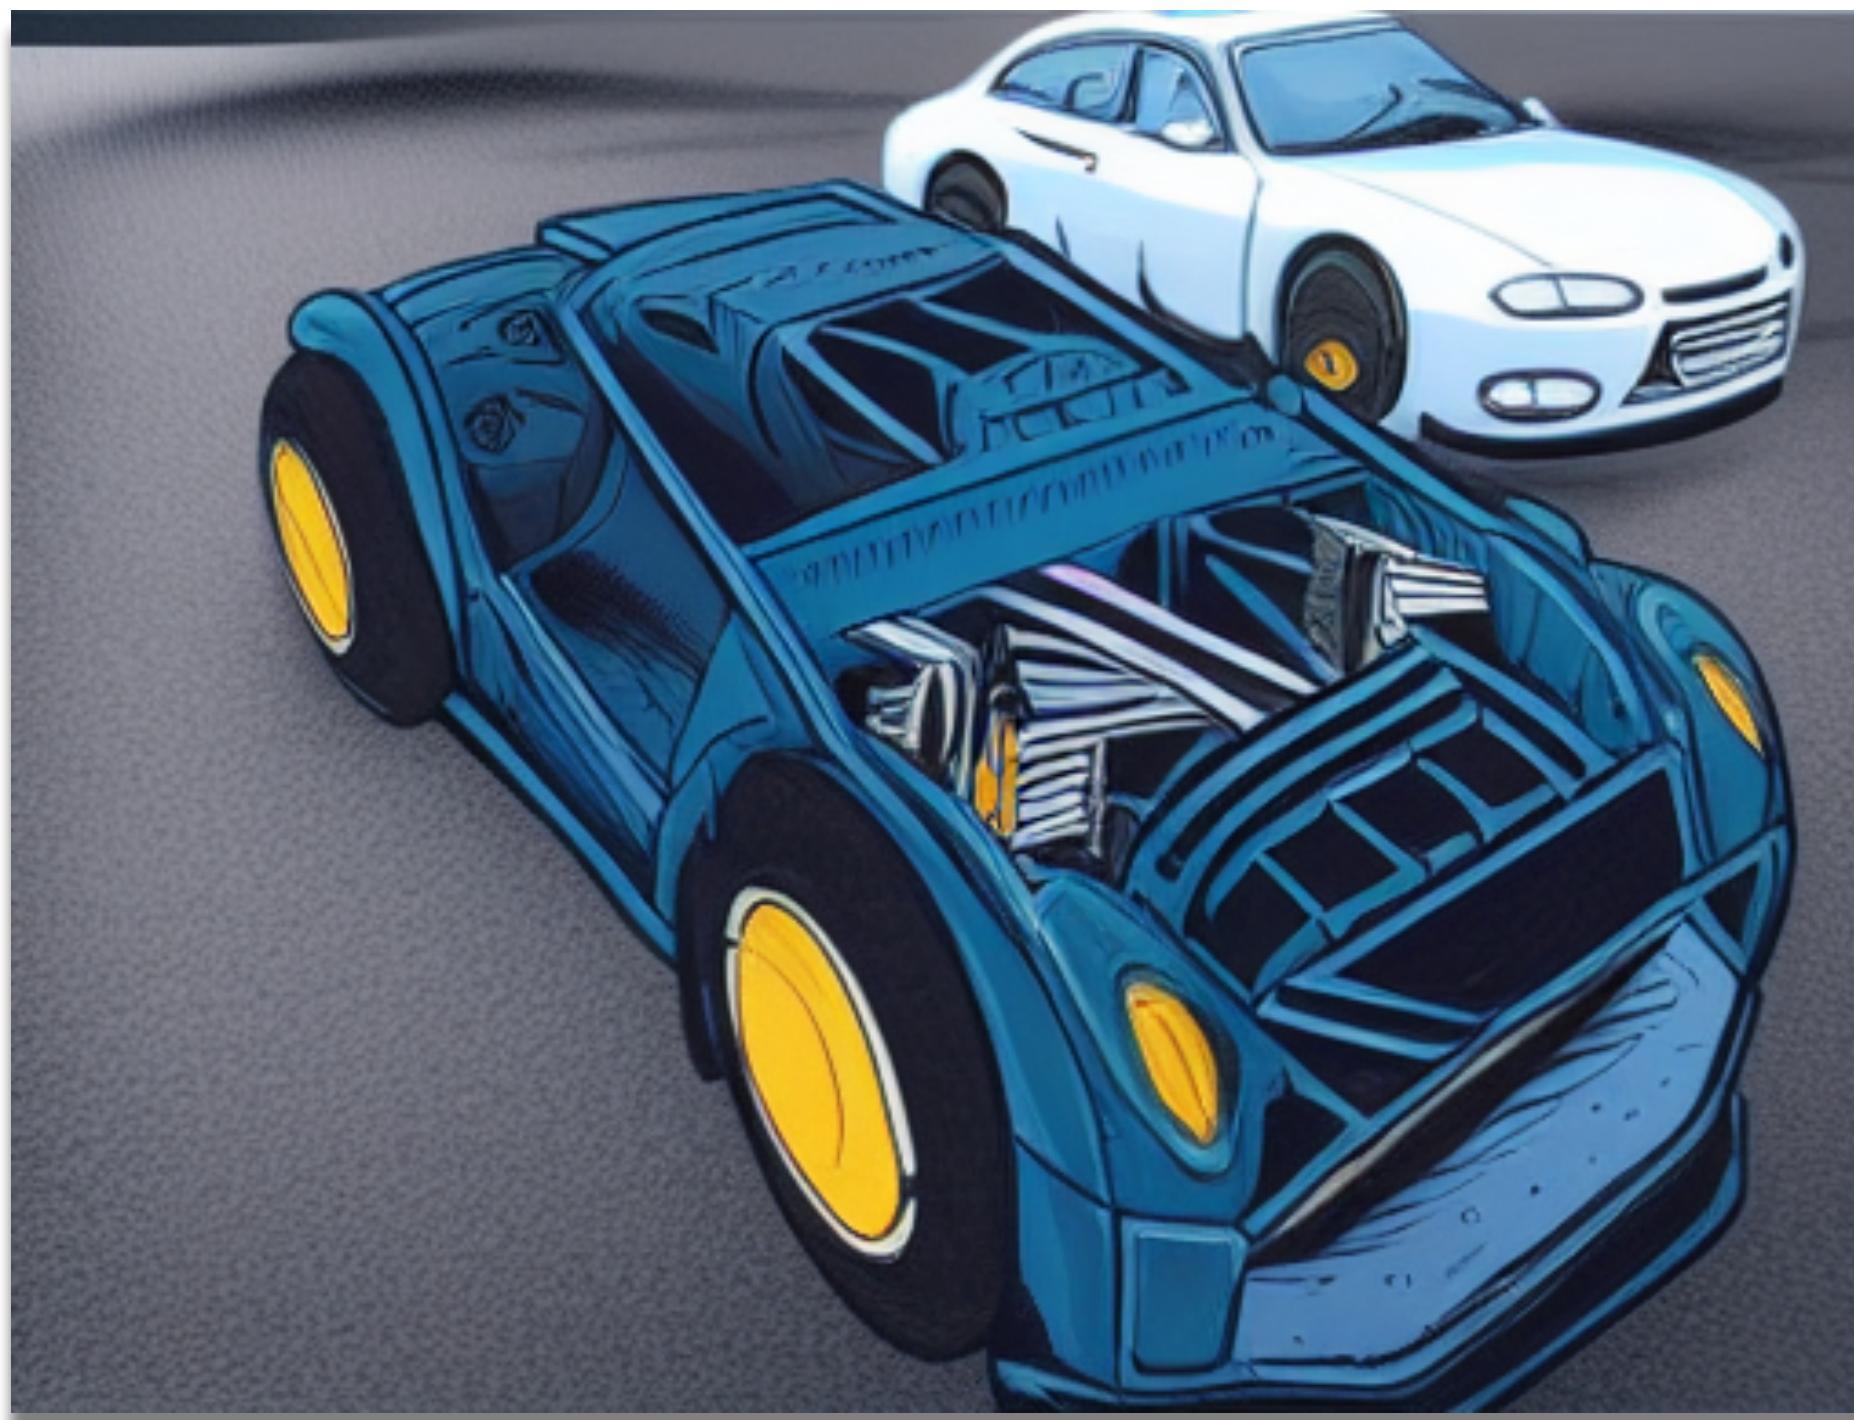
# OCaml 5.0 — an MVP release



- Many features broken and are being added back
  - ◆ This will continue after 5.0 gets released
- Platform support
  - ◆ 32-bit will be *bytecode only*
  - ◆ On 64-bit,
    - ◆ x86-64 + Linux, macOS, Windows, OpenBSD, FreeBSD
    - ◆ Arm64 + Linux, macOS (*Apple Silicon*)
    - ◆ RISC-V (*PR open*)
  - ◆ JavaScript (js0o) — effect handlers are not supported yet!

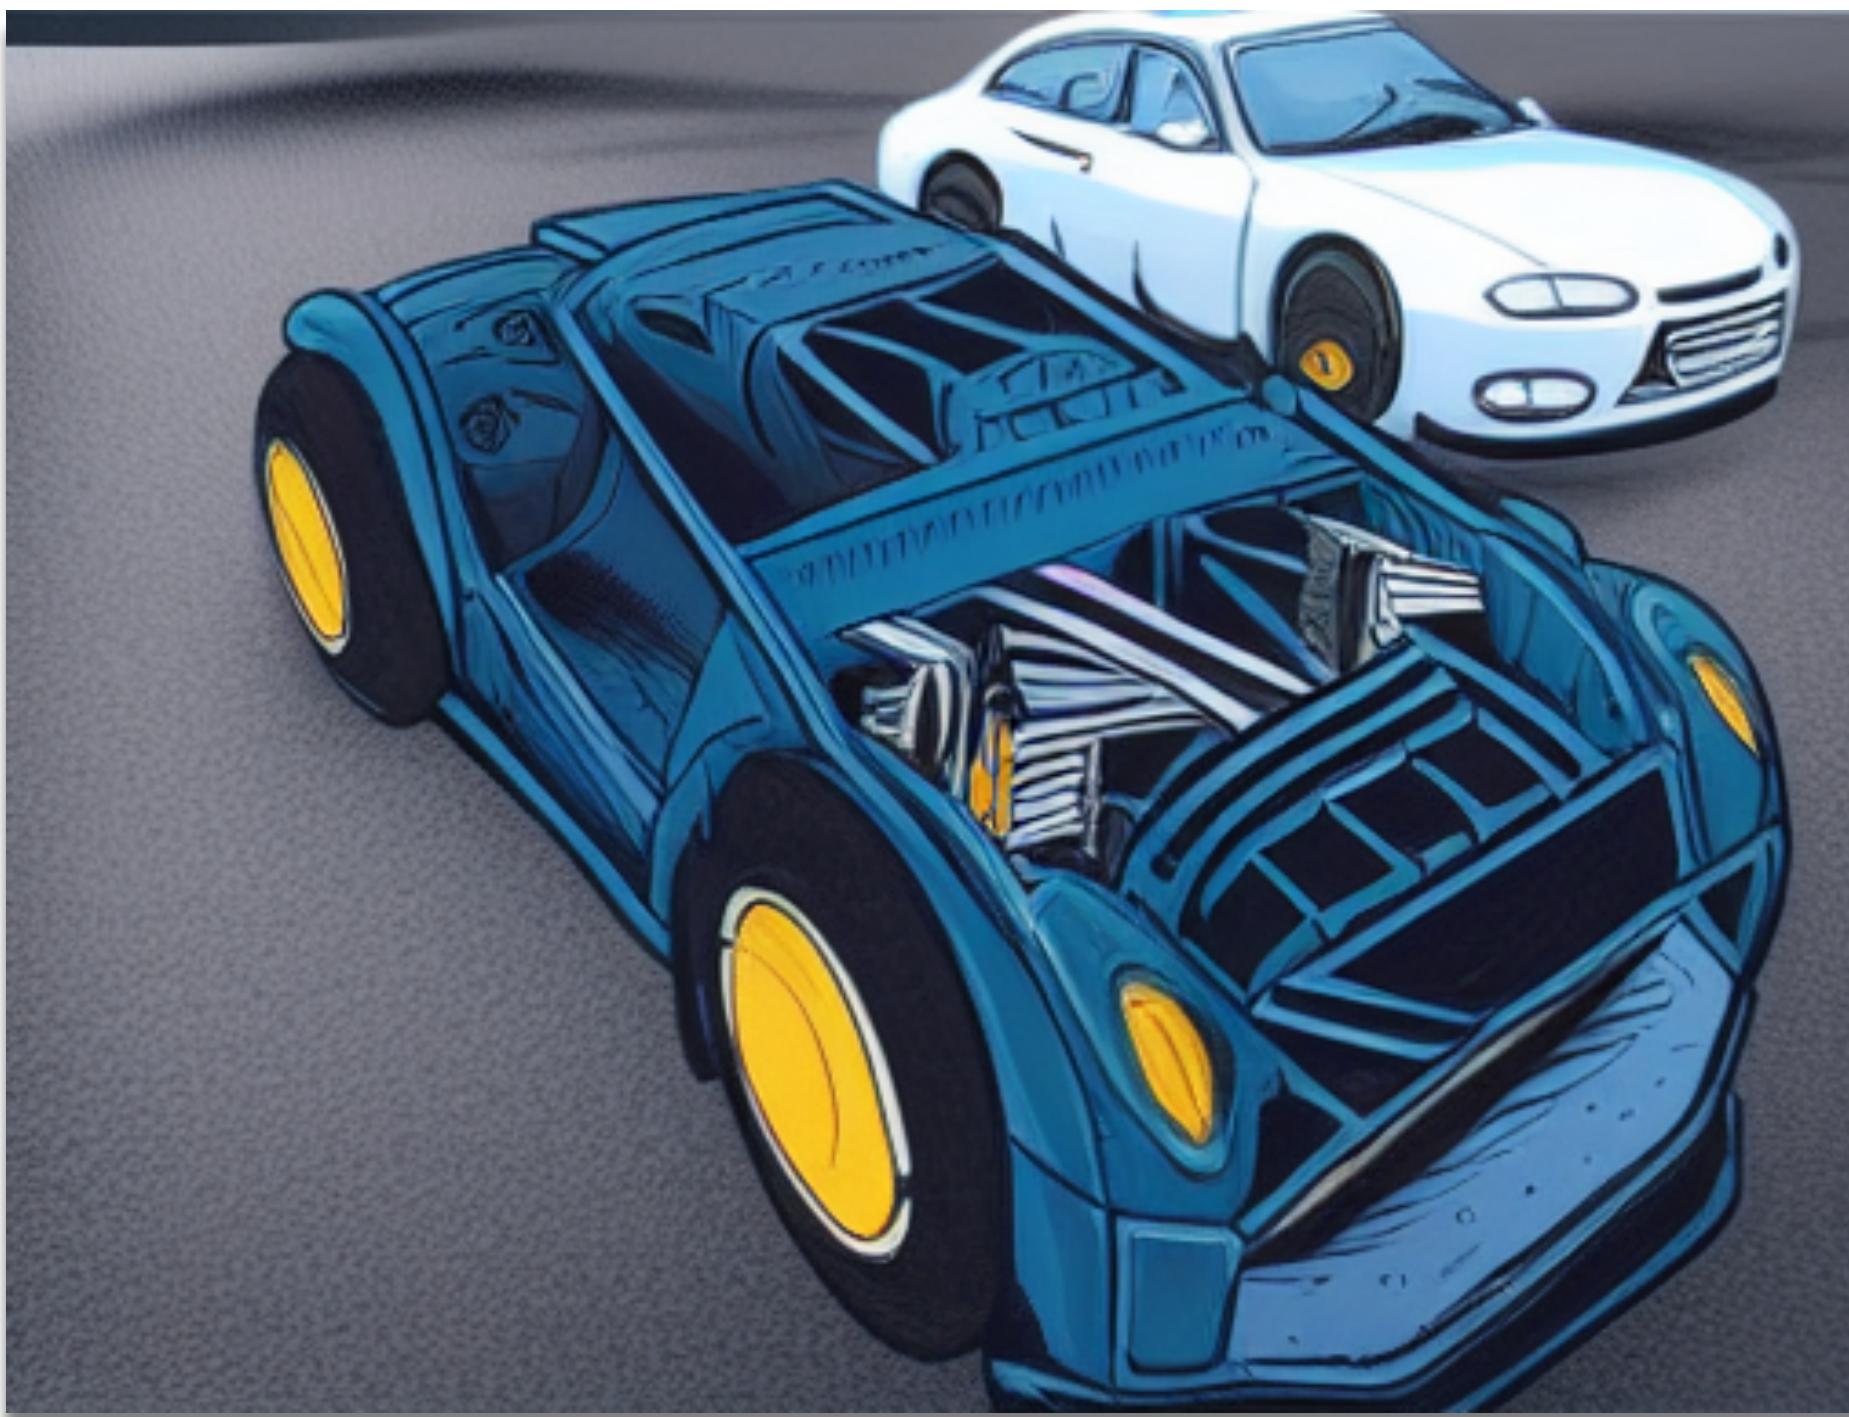
# OCaml 5.0 — an MVP release

- GC performance improvements TBD
  - ◆ Decoupling major slice from minor GC
  - ◆ Mark stack prefetching
  - ◆ Best-fit vs multicore allocator



# OCaml 5.0 — an MVP release

- GC performance improvements TBD
  - ◆ Decoupling major slice from minor GC
  - ◆ Mark stack prefetching
  - ◆ Best-fit vs multicore allocator
- Statmemprof
  - ◆ Work in progress for reinstating asynchronous action safety



# Tidying

- We tidied up accumulated deprecations
  - ◆ `String.uppercase`, `lowercase`, `capitalize`, `uncapitalize`
  - ◆ `Stream`, `Genlex` ~> `camlp-streams`
  - ◆ `Pervasives`, `ThreadUnix` modules deleted
- Major version jump to make good changes
  - ◆ C function names are all prefixed uniformly
  - ◆ Additional libraries `Unix`, `Str` installed as findlib packages

# OPAM Health Check

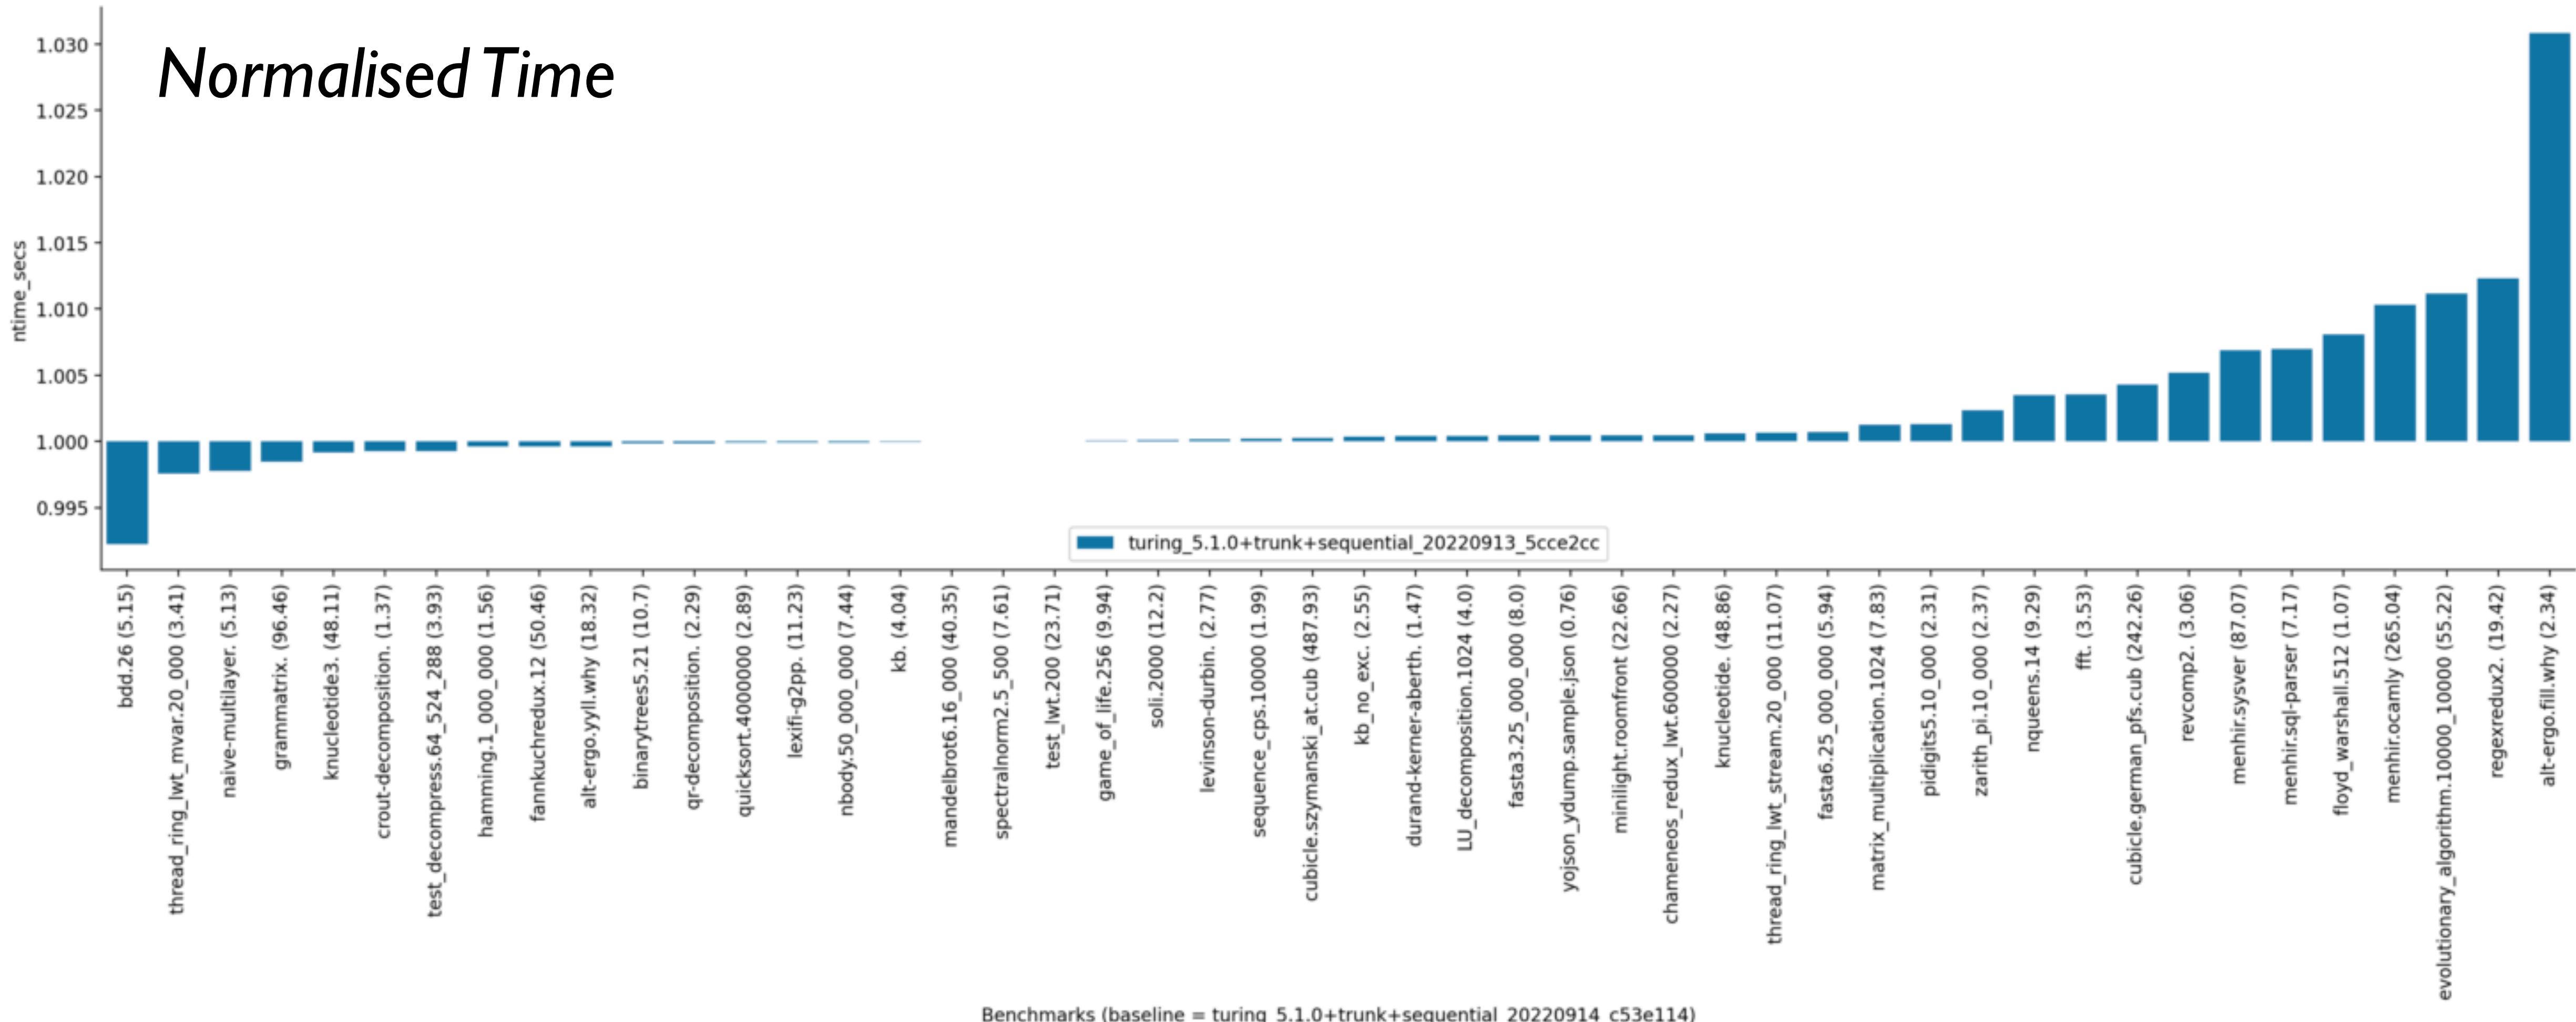
# OPAM Health Check

	4.14	5.0+alpha-repo	number of revdeps
0install.2.18	✓	✗	1
BetterErrors.0.0.1	✓	✗	7
TCSLib.0.3	✓	✗	1
absolute.0.1	✓	✗	0
acgtk.1.5.3	✓	✗	0
advi.2.0.0	✓	✗	0
aez.0.3	✓	✗	0
ahrocksdb.0.2.2	✗	✗	0
aio.0.0.3	✓	✗	0
alt-ergo-free.2.2.0	✓	✗	7
amqp-client-async.2.2.2	✓	✗	0
amqp-client-lwt.2.2.2	✓	✗	0
ancient.0.9.1	✓	✗	0
apron.v0.9.13	✓	✗	17

<http://check.ocamlabs.io/>

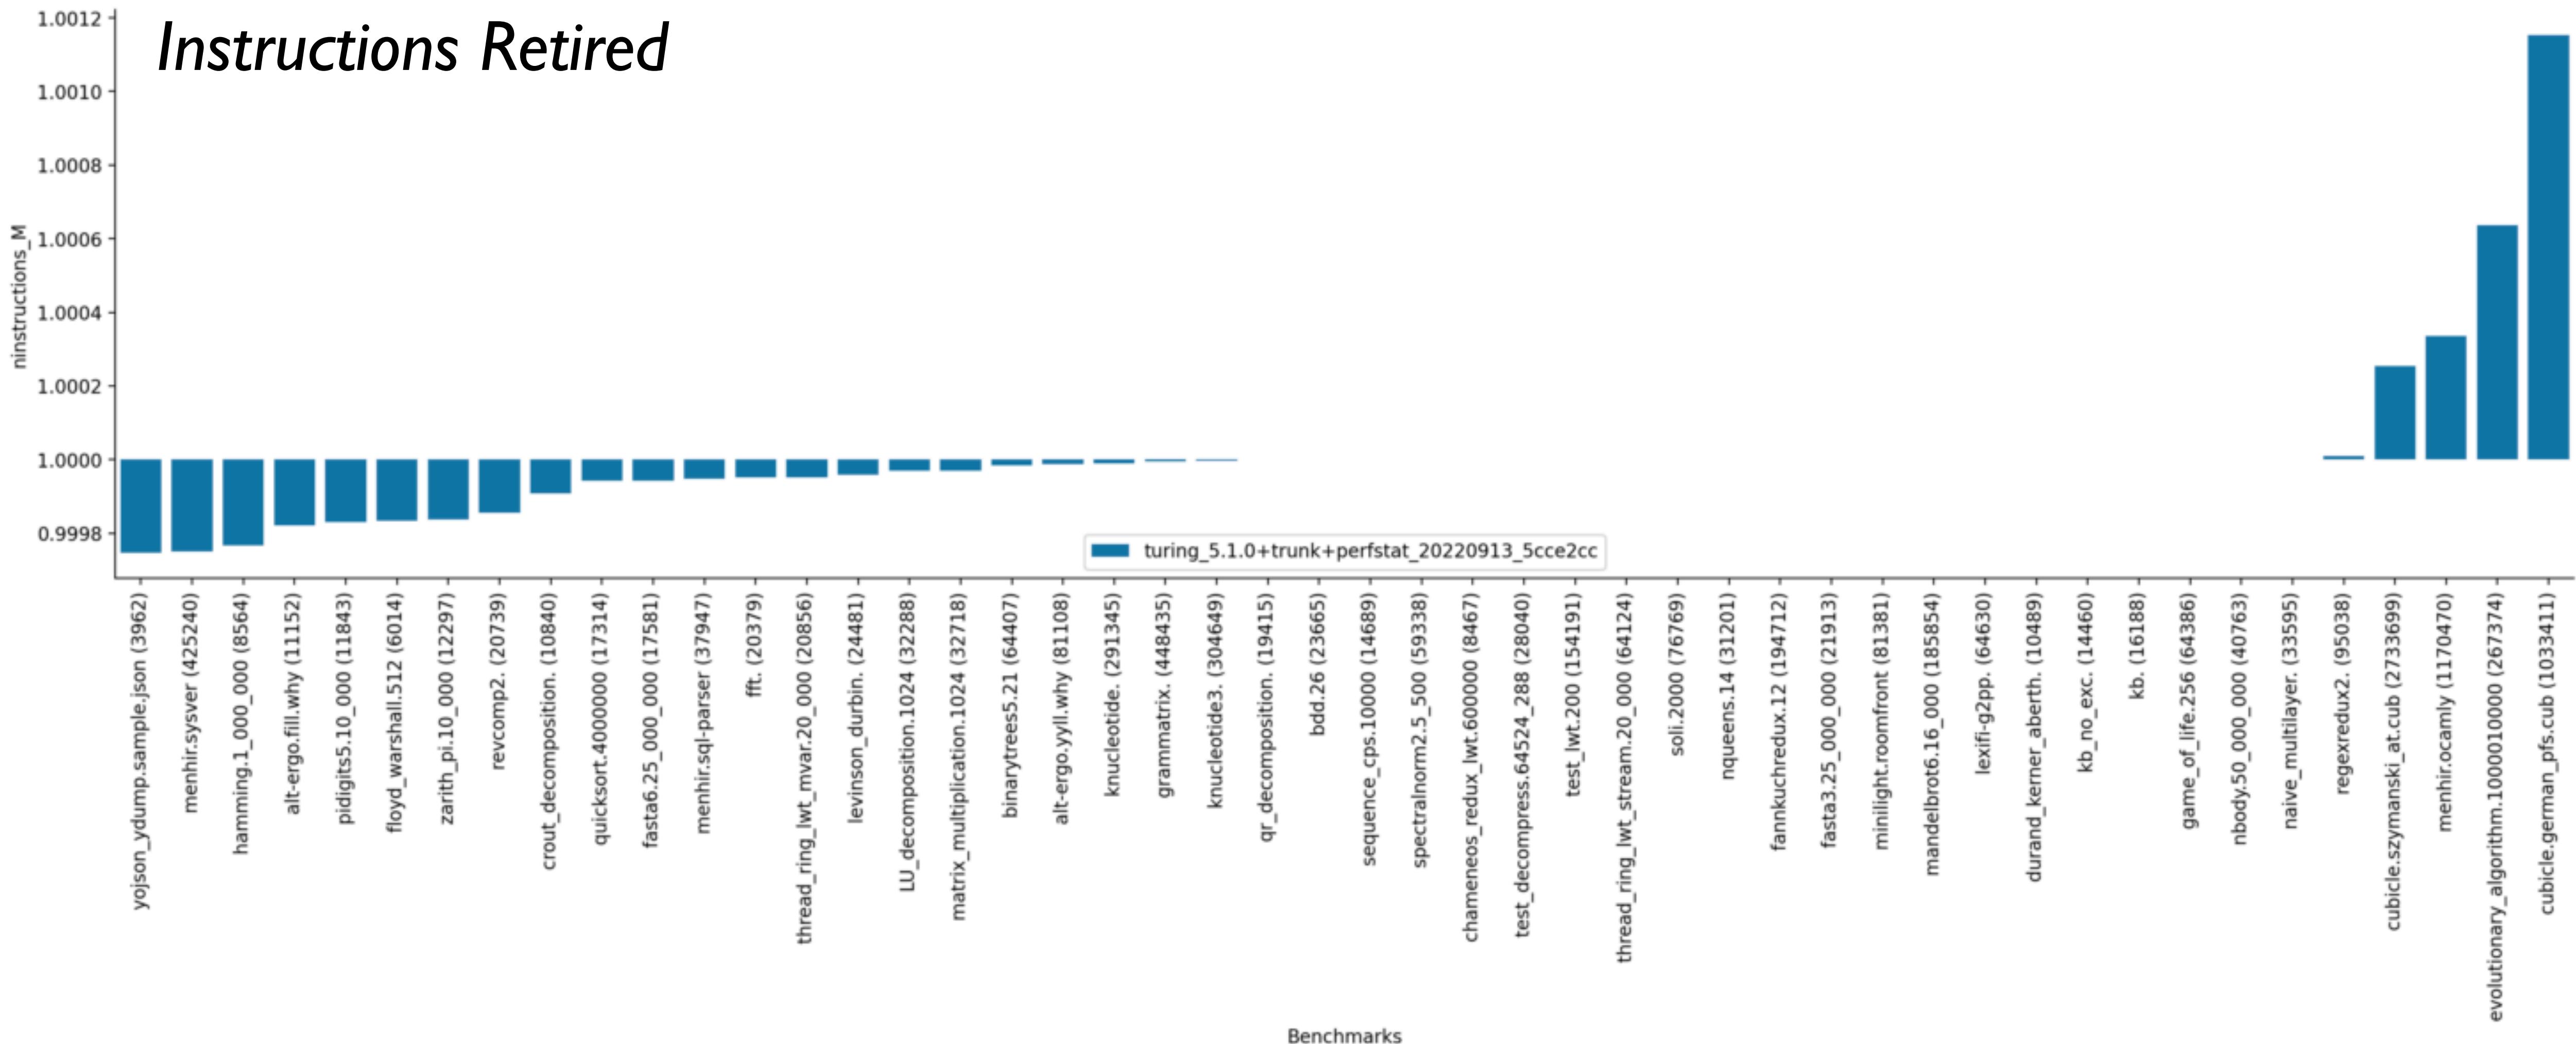
# Sandmark Nightly Service

[sandmark.tarides.com](http://sandmark.tarides.com)



# Sandmark Nightly Service

## *Instructions Retired*



# OCaml 5.0 needs you!



- OCaml 4 will have longer term support than usual
- Even if you don't plan to use concurrency and parallelism, *switch to OCaml 5.0*
  - ♦ Only then can we move away from OCaml 4.x

# OCaml 5.0 needs you!



- OCaml 4 will have longer term support than usual
- Even if you don't plan to use concurrency and parallelism, *switch to OCaml 5.0*
  - ◆ Only then can we move away from OCaml 4.x
- Sequential programs *must* work with same perf on 5.0
  - ◆ Test, deploy, evaluate, benchmark sequential programs in 5.0
  - ◆ Report bugs & performance regressions

# OCaml 5.0 needs you!



- OCaml 4 will have longer term support than usual
- Even if you don't plan to use concurrency and parallelism, *switch to OCaml 5.0*
  - ◆ Only then can we move away from OCaml 4.x
- Sequential programs *must* work with same perf on 5.0
  - ◆ Test, deploy, evaluate, benchmark sequential programs in 5.0
  - ◆ Report bugs & performance regressions
- **What is stopping you from switching to OCaml 5.0?**
  - ◆ Let us know so that we can work on it!