

# Rx-CML

## : A Prescription for Safely Relaxing Synchrony

KC Sivaramakrishnan

Lukasz Ziarek  
SUNY Buffalo

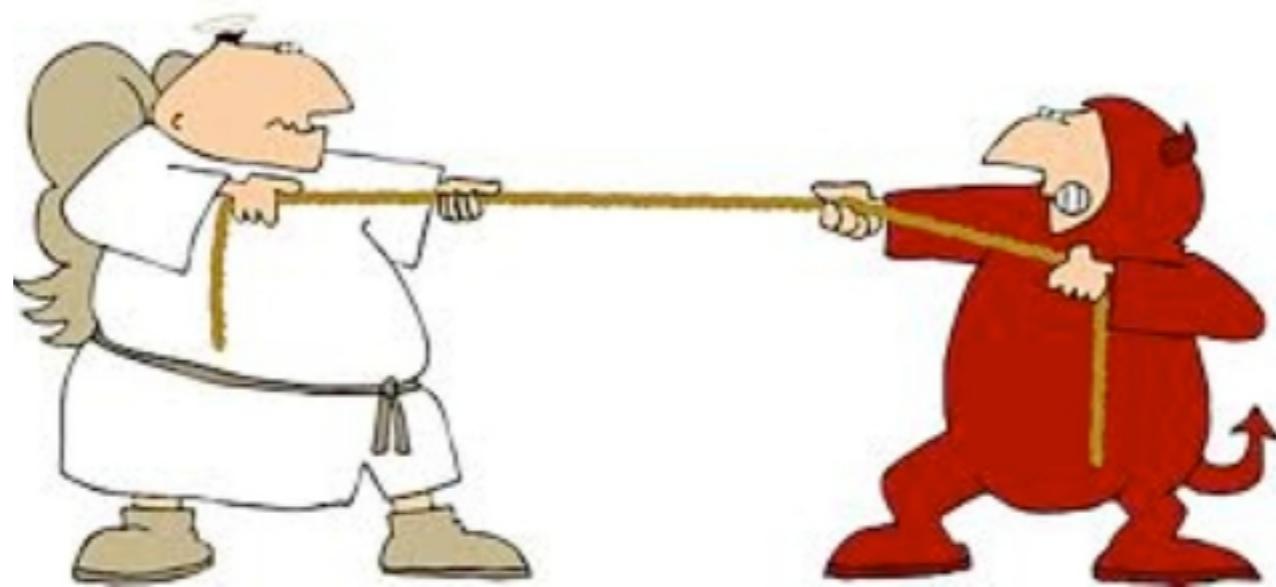
Suresh Jagannathan  
Purdue University

# Introduction

Two often competing goals when *designing* and  
*implementing* concurrency abstractions

# Introduction

Two often competing goals when *designing* and *implementing* concurrency abstractions




*Simplicity*  
*Safety*

*Performance*  
*Functionality*

# Introduction

Two often competing goals when *designing* and *implementing* concurrency abstractions



*Simplicity*  
*Safety*

*Performance*  
*Functionality*



Always desirable to marry the two whenever possible


# Big Picture

# Big Picture

- Functional language + Synchronous message passing
  - ★ Communication = Data transfer + Synchronization

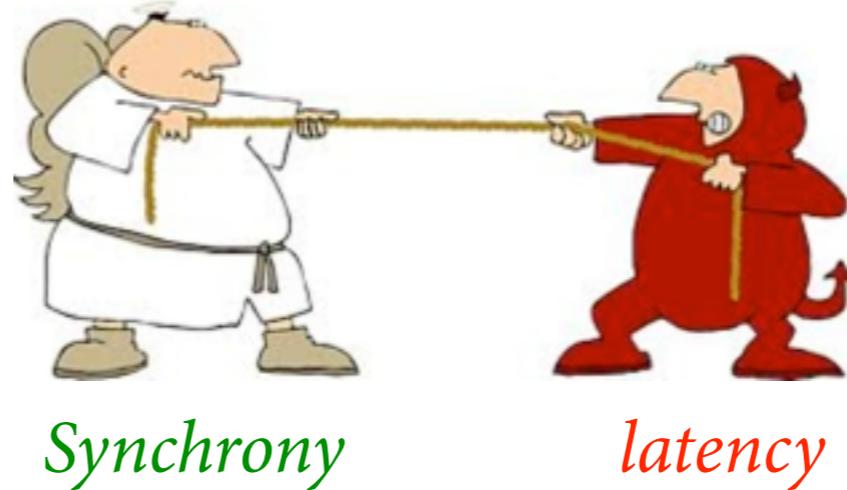
# Big Picture

- Functional language + Synchronous message passing
  - ★ Communication = Data transfer + Synchronization
- However, in the cloud,



# Big Picture

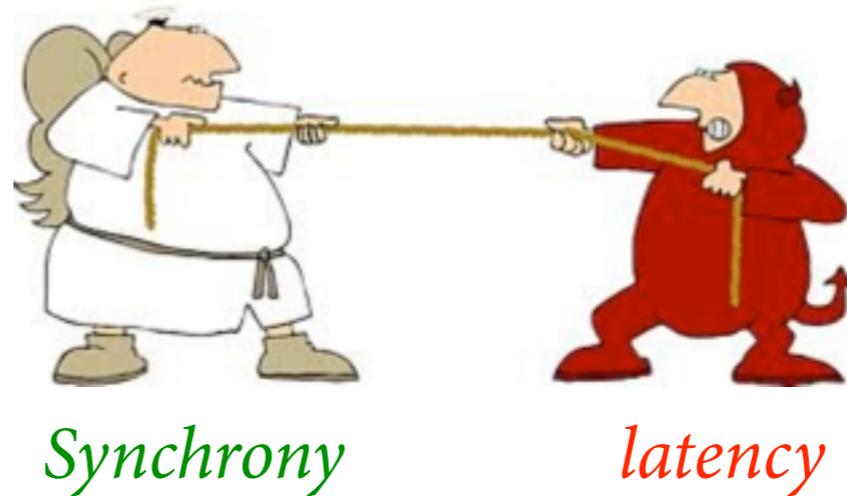
- Functional language + Synchronous message passing
  - ★ Communication = Data transfer + Synchronization
- However, in the cloud,



- ★ Explicit asynchrony **complicates reasoning**

# Big Picture

- Functional language + Synchronous message passing
  - ★ Communication = Data transfer + Synchronization
- However, in the cloud,



- ★ Explicit asynchrony **complicates reasoning**

*Can we discharge synchronous communications **asynchronously** while ensuring **observable equivalence**?*

# Goal

# Goal

1. Formalize the conditions under which the following equivalence holds:

$$[\![ \text{send} (c, v) ]\!]_k \equiv [\![ \text{asend} (c, v) ]\!]_k$$

# Goal

1. Formalize the conditions under which the following equivalence holds:

$$[\![ \text{send} (c, v) ]\!]_k \equiv [\![ \text{asend} (c, v) ]\!]_k$$

2. A cloud infrastructure + speculative execution framework
  - a. discharges synchronous sends asynchronously
  - b. detects when the equivalence fails, and
  - c. repairs failed executions

# Context

# Context

- A distributed extension of MultiMLton - MLton for scalable architectures

# Context

- A distributed extension of MultiMLton - MLton for scalable architectures
- Parallel extension of *Concurrent ML*
  - ★ Dynamic lightweight threads
  - ★ *Synchronous* message passing
  - ★ First-class events
    - ◆ *Composable synchronous protocols*

# Context

- A distributed extension of MultiMLton - MLton for scalable architectures
- Parallel extension of *Concurrent ML*
  - ★ Dynamic lightweight threads
  - ★ *Synchronous* message passing
  - ★ First-class events
    - ◆ *Composable synchronous protocols*


```
val channel : unit -> 'a chan
val spawn   : (unit -> unit) -> thread_id
val send    : 'a chan * 'a -> unit
val recv    : 'a chan -> 'a
val sendEvt : 'a chan * 'a -> unit event
val recvEvt : 'a chan -> 'a event
val sync    : 'a event -> 'a
```

```
val never     : 'a event
val alwaysEvt : 'a -> 'a event
val wrap      : 'a event -> ('a -> 'b) -> 'b event
val guard    : (unit -> 'a event) -> 'a event
val choose   : 'a event list -> 'a event
...
```

# Basic Idea (1)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)


*Synchronous  
Execution*



# Basic Idea (1)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)

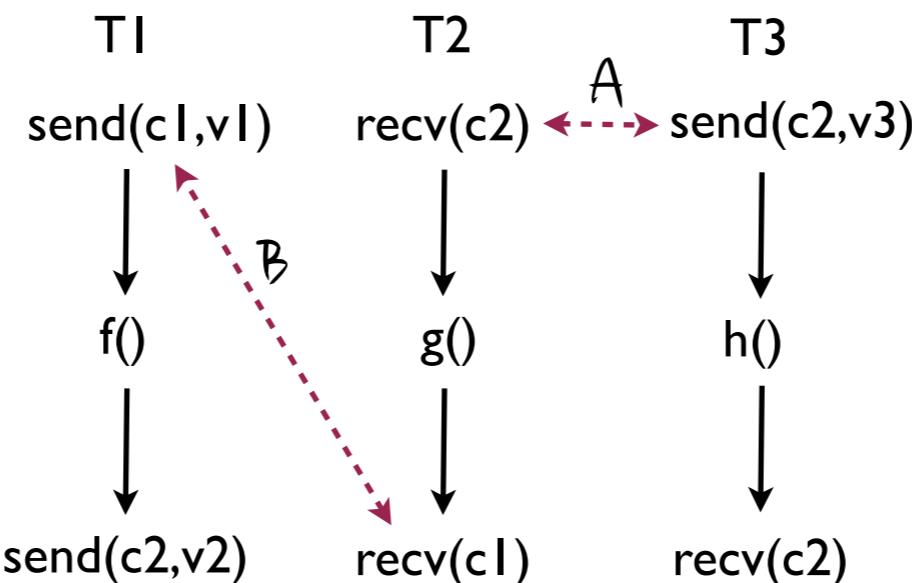
*Synchronous  
Execution*



# Basic Idea (1)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)


*Synchronous  
Execution*



# Basic Idea (1)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)


*Synchronous  
Execution*



# Basic Idea (2)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)

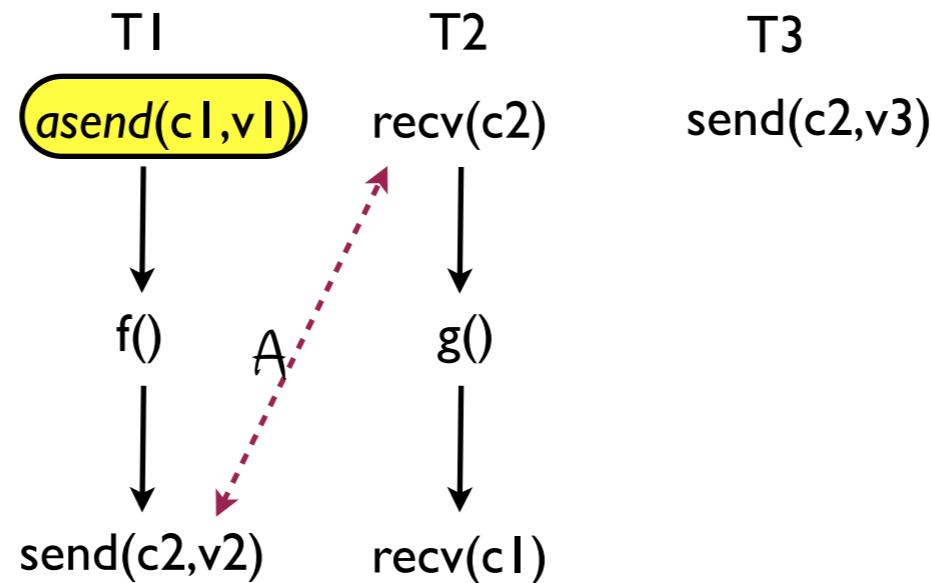
*Asynchronous  
Execution*



# Basic Idea (2)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)

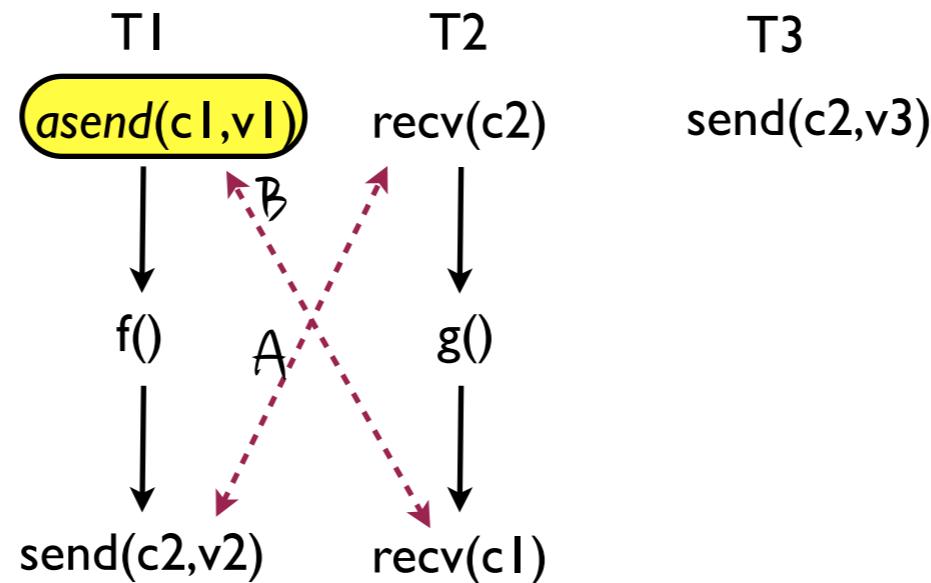
*Asynchronous  
Execution*



# Basic Idea (2)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)


*Asynchronous  
Execution*



# Basic Idea (2)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)


*Asynchronous  
Execution*



# Basic Idea (2)

T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)

*Asynchronous  
Execution*



- Synchronous evaluation **never** results in cyclic dependence
  - ★ Cyclic dependence => divergent behavior w.r.t synchronous evaluation

# Example: Distributed Group Chat

# Example: Distributed Group Chat

- No central server + Preserve causal dependence

# Example: Distributed Group Chat

- No central server + Preserve causal dependence
  - ★ **Causal broadcast primitive:** If message A is generated **as a response** to message B, then A is **delivered after** B at every site

# Example: Distributed Group Chat

- No central server + Preserve causal dependence
  - ★ **Causal broadcast primitive:** If message A is generated **as a response** to message B, then A is **delivered after** B at every site
  - ★ Asynchronous messaging => **explicitly** manage vector clocks and buffering on receiver side

# Example: Distributed Group Chat

- No central server + Preserve causal dependence
  - ★ **Causal broadcast primitive:** If message A is generated **as a response** to message B, then A is **delivered after** B at every site
  - ★ Asynchronous messaging => **explicitly** manage vector clocks and buffering on receiver side
  - ★ Synchronous messaging => **directly** using point-to-point messaging

# Example: Distributed Group Chat

- No central server + Preserve causal dependence
  - ★ **Causal broadcast primitive:** If message A is generated **as a response** to message B, then A is **delivered after** B at every site
  - ★ Asynchronous messaging => **explicitly** manage vector clocks and buffering on receiver side
  - ★ Synchronous messaging => **directly** using point-to-point messaging

```
fun bsend (BCHAN (vcList, acList), v: 'a, id: int) : unit =
let
  val _ = map (fn vc => if (vc = nth (vcList, id)) then () else send (vc, v))
    vcList (* phase 1 -- Value distribution *)
  val _ = map (fn ac => if (ac = nth (acList, id)) then () else recv ac)
    acList (* phase 2 -- Acknowledgments *)
in ()
end
synchronously send values
```

*prevent receivers from proceeding until all members have received the value*

# Example: Distributed Group Chat

- No central server + Preserve causal dependence
  - ★ **Causal broadcast primitive:** If message A is generated **as a response** to message B, then A is **delivered after** B at every site
  - ★ Asynchronous messaging => **explicitly** manage vector clocks and buffering on receiver side
  - ★ Synchronous messaging => **directly** using point-to-point messaging

```
fun bsend (BCHAN (vcList, acList), v: 'a, id: int) : unit =
let
  val _ = map (fn vc => if (vc = nth (vcList, id)) then () else send (vc, v))
    vcList (* phase 1 -- Value distribution *)
  val _ = map (fn ac => if (ac = nth (acList, id)) then () else recv ac)
    acList (* phase 2 -- Acknowledgments *)
in ()
end
  synchronously send values
  prevent receivers from proceeding until
  all members have received the value
```

- Simple but likely to be inefficient - **phase 2 is a global barrier!**

# Example: Distributed Group Chat

- No central server + Preserve causal dependence
  - ★ **Causal broadcast primitive:** If message A is generated **as a response** to message B, then A is **delivered after** B at every site
  - ★ Asynchronous messaging => **explicitly** manage vector clocks and buffering on receiver side
  - ★ Synchronous messaging => **directly** using point-to-point messaging

```
fun bsend (BCHAN (vcList, acList), v: 'a, id: int) : unit =
let
  val _ = map (fn vc => if (vc = nth (vcList, id)) then () else send (vc, v))
    vcList (* phase 1 -- Value distribution *)
  val _ = map (fn ac => if (ac = nth (acList, id)) then () else recv ac)
    acList (* phase 2 -- Acknowledgments *)
in ()
end
synchronously send values
prevent receivers from proceeding until all members have received the value
```

- Simple but likely to be inefficient - **phase 2 is a global barrier!**
  - ★ Discharging asynchronously **breaks causal ordering**

# Example: Distributed Group Chat


- No central server + Preserve causal dependence
  - ★ **Causal broadcast primitive:** If message A is generated **as a response** to message B, then A is **delivered after** B at every site
  - ★ Asynchronous messaging => **explicitly** manage vector clocks and buffering on receiver side
  - ★ Synchronous messaging => **directly** using point-to-point messaging

```
fun bsend (BCHAN (vcList, acList), v: 'a, id: int) : unit =
let
  val _ = map (fn vc => if (vc = nth (vcList, id)) then () else send (vc, v))
    vcList (* phase 1 -- Value distribution *)
  val _ = map (fn ac => if (ac = nth (acList, id)) then () else recv ac)
    acList (* phase 2 -- Acknowledgments *)
in ()
end
  synchronously send values
  prevent receivers from proceeding until
  all members have received the value
```

- Simple but likely to be inefficient - **phase 2 is a global barrier!**
  - ★ Discharging asynchronously **breaks causal ordering**
  - ★ *Our idea: program synchronously, discharge asynchronously, detect and remediate causal ordering violations*

# Example: Distributed Group Chat

- A distributed group chat program = {Node}
- Node = MultiMLton process = {CML threads}



# Distributed Group Chat - Run 1



Display	IP	Daemon
---------	----	--------




Display	IP	Daemon
---------	----	--------


# Distributed Group Chat - Run 1




# Distributed Group Chat - Run 1




# Distributed Group Chat - Run 1




# Distributed Group Chat - Run 1




# Distributed Group Chat - Run 1



# Distributed Group Chat - Run 1



# Distributed Group Chat - Run 1



- Observations
  - ★ X and Y independently generated => No causal dependence between bcast (X) and bcast (Y)
- No Cycles => Correct execution!

# Distributed Group Chat - Run 2



Display	IP	Daemon
---------	----	--------




Display	IP	Daemon
---------	----	--------




Display	IP	Daemon
---------	----	--------


# Distributed Group Chat - Run 2




# Distributed Group Chat - Run 2




# Distributed Group Chat - Run 2




# Distributed Group Chat - Run 2




# Distributed Group Chat - Run 2




# Distributed Group Chat - Run 2




# Distributed Group Chat - Run 2



# Distributed Group Chat - Run 2



# Distributed Group Chat - Run 2



# Distributed Group Chat - Results

- Simulation on 3 geo-distributed Amazon EC2 instances
- Measure time between message initiation and receipt by all parties over 1000 iterations

<i>Execution</i>	<i>Avg.time (ms)</i>	<i>Errors</i>
<i>Sync</i>	1540	0
<i>Unsafe Async</i>	520	7
<i>Safe Async (Rx<sup>CML</sup>)</i>	533	0

# Formalization Overview

# Formalization Overview

- Reason *axiomatically* about executions (relaxed or otherwise)
  - ★ Similar to formalizations used in relaxed memory models
  - ★ Declarative characterization of (relaxed) CML behavior

# Formalization Overview

- Reason *axiomatically* about executions (relaxed or otherwise)
  - ★ Similar to formalizations used in relaxed memory models
  - ★ Declarative characterization of (relaxed) CML behavior
- *Actions + happens-before relation*
  - ★ Captures visibility and dependence properties

# Formalization Overview

- Reason *axiomatically* about executions (relaxed or otherwise)
  - ★ Similar to formalizations used in relaxed memory models
  - ★ Declarative characterization of (relaxed) CML behavior
- *Actions + happens-before relation*
  - ★ Captures visibility and dependence properties
- Happens-before is intentionally *relaxed*: may define more behaviors than possible in CML
  - ★ Strengthen the relation with *well-formedness* conditions


# Actions and Execution

- Actions:

$\mathbb{A} :=$	$b_t$	(thread t starts)
	$e_t$	(thread t ends)
	$j_t^m t'$	(thread t detects thread t' has terminated)
	$f_t^m t'$	(thread t creates a new thread t')
	$s_t^m c, v$	(thread t sends value v on channel c)
	$r_t^m c$	(thread t receives a value on channel c)
	$p_t^m v$	(thread t outputs an observable value v)

$$c \in \mathbb{C} \quad t, t' \in \mathbb{T} \quad v \in \mathbb{V} \quad m, n \in \mathbb{N}$$

- Execution:



# Communication and Thread Dependence

- Synchronous communication  $\rightarrow$  communication order is symmetric:

$$a \rightarrow_{co} b \implies b \rightarrow_{co} a$$

- Thread dependence order:

$\alpha \rightarrow_{td} \beta$  if:

- (1)  $\alpha = f_t^m t'$  and  $\beta = b_{t'}$  or
- (2)  $\alpha = e_t$  and  $\beta = j_{t'}^m t$

# Happens-before Relation

- Establishes both intra- and inter-thread dependences:

$$\rightarrow_{hb} = (\rightarrow_{po} \cup \rightarrow_{td} \cup \\ \{(\alpha, \beta) \mid \alpha \rightarrow_{co} \alpha' \wedge \alpha' \rightarrow_{po} \beta\} \cup \\ \{(\beta, \alpha) \mid \beta \rightarrow_{po} \alpha' \wedge \alpha' \rightarrow_{co} \alpha\})^+$$

# Happens-before Relation

- Establishes both intra- and inter-thread dependences:

$$\rightarrow_{hb} = (\rightarrow_{po} \cup \rightarrow_{td} \cup \\ \{(\alpha, \beta) \mid \alpha \rightarrow_{co} \alpha' \wedge \alpha' \rightarrow_{po} \beta\} \cup \\ \{(\beta, \alpha) \mid \beta \rightarrow_{po} \alpha' \wedge \alpha' \rightarrow_{co} \alpha\})^+$$


- Two actions not related by happens-before relation are said to be concurrent
  - ★ A send action and its matching receive action are concurrent!

# Happens-before Relation

- Establishes both intra- and inter-thread dependences:

$$\rightarrow_{hb} = (\rightarrow_{po} \cup \rightarrow_{td} \cup \{(\alpha, \beta) \mid \alpha \rightarrow_{co} \alpha' \wedge \alpha' \rightarrow_{po} \beta\} \cup \{(\beta, \alpha) \mid \beta \rightarrow_{po} \alpha' \wedge \alpha' \rightarrow_{co} \alpha\})^+$$

- Two actions not related by happens-before relation are said to be concurrent
  - ★ A send action and its matching receive action are concurrent!



# Happens-before Relation

- Establishes both intra- and inter-thread dependences:

$$\rightarrow_{hb} = (\rightarrow_{po} \cup \rightarrow_{td} \cup \{(\alpha, \beta) \mid \alpha \rightarrow_{co} \alpha' \wedge \alpha' \rightarrow_{po} \beta\} \cup \{(\beta, \alpha) \mid \beta \rightarrow_{po} \alpha' \wedge \alpha' \rightarrow_{co} \alpha\})^+$$

- Two actions not related by happens-before relation are said to be concurrent
  - ★ A send action and its matching receive action are concurrent!



# Example


T1	T2	T3
send(c1,v1)	recv(c2)	send(c2,v3)
f()	g()	h()
send(c2,v2)	recv(c1)	recv(c2)

- Assume T1 spawns T2 and T3
- Let f, g, h = print 1, print 2, print 3

# Example

T1	T2	T3
send(c1, v1)	recv(c2)	send(c2, v3)
f()	g()	h()
send(c2, v2)	recv(c1)	recv(c2)

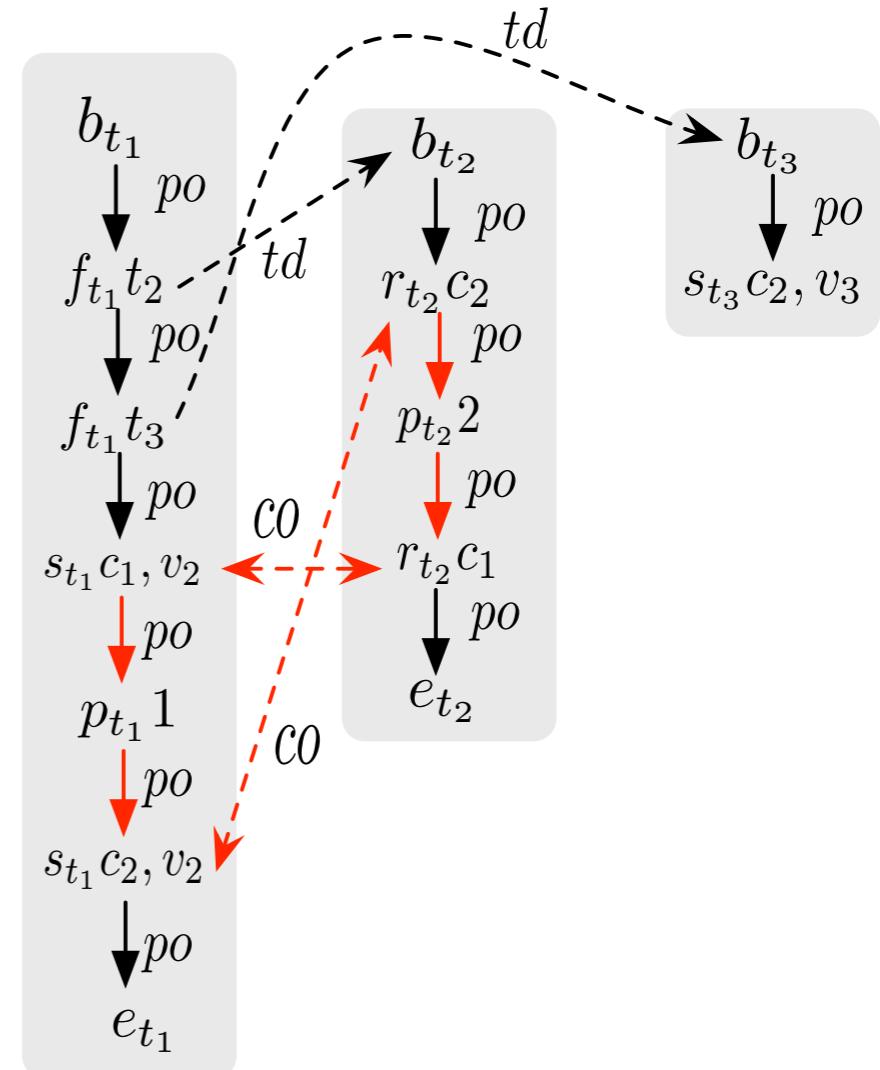
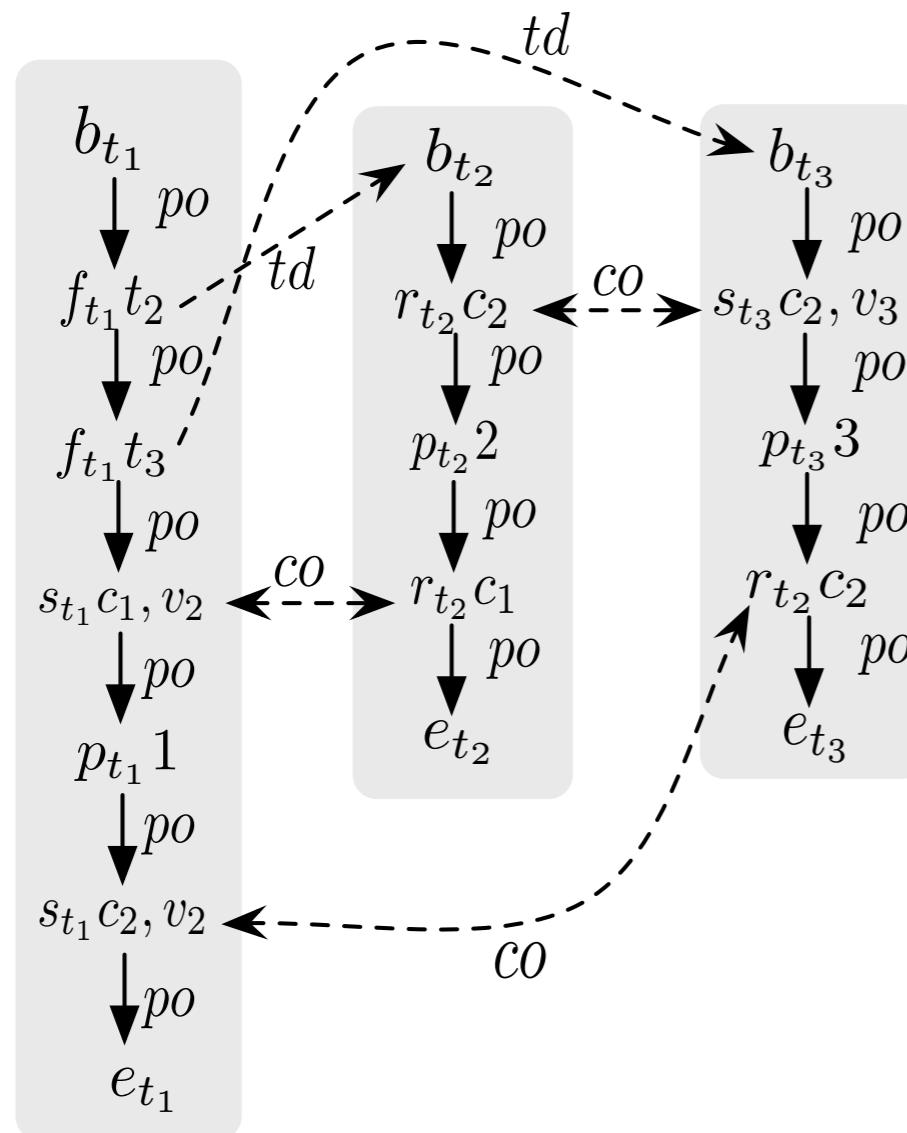
- Assume T1 spawns T2 and T3
- Let f, g, h = print 1, print 2, print 3



# Example


T1	T2	T3
send(c1, v1)	recv(c2)	send(c2, v3)
f()	g()	h()
send(c2, v2)	recv(c1)	recv(c2)

- Assume T1 spawns T2 and T3
- Let f, g, h = print 1, print 2, print 3




# Well-formed Executions

$\text{Obs}(\text{Well-formed Execution of } P) \in \{\text{Obs}(\text{CML Execution of } P)\}$




# Well-formed Executions

$\text{Obs}(\text{Well-formed Execution of } P) \in \{\text{Obs}(\text{CML Execution of } P)\}$



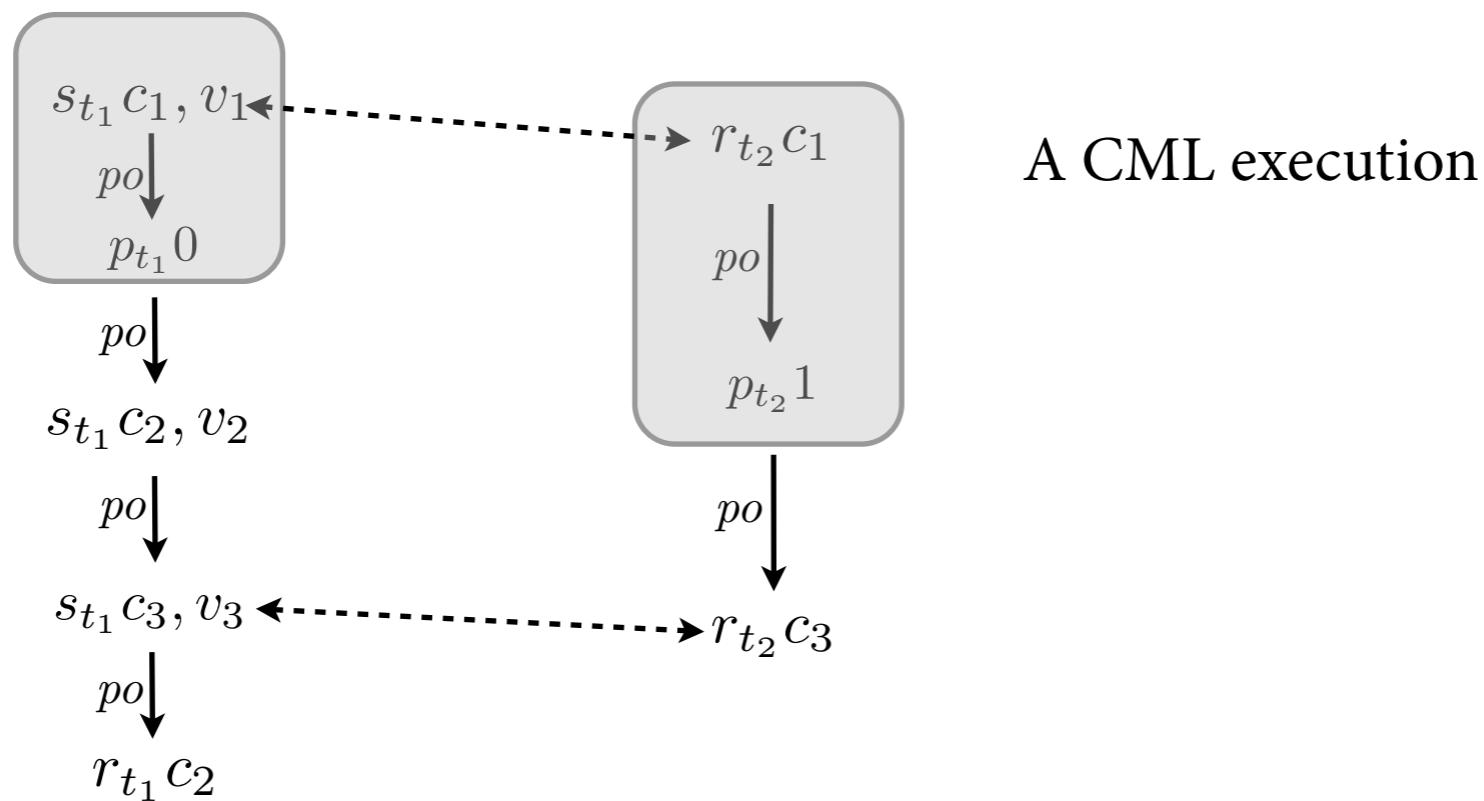
# Well-formedness -> Speculation

*Track executions to see if they become ill-formed (rollback) or turn into CML executions (commit)*




# Well-formedness -> Speculation

*Track executions to see if they become ill-formed (rollback) or turn into CML executions (commit)*



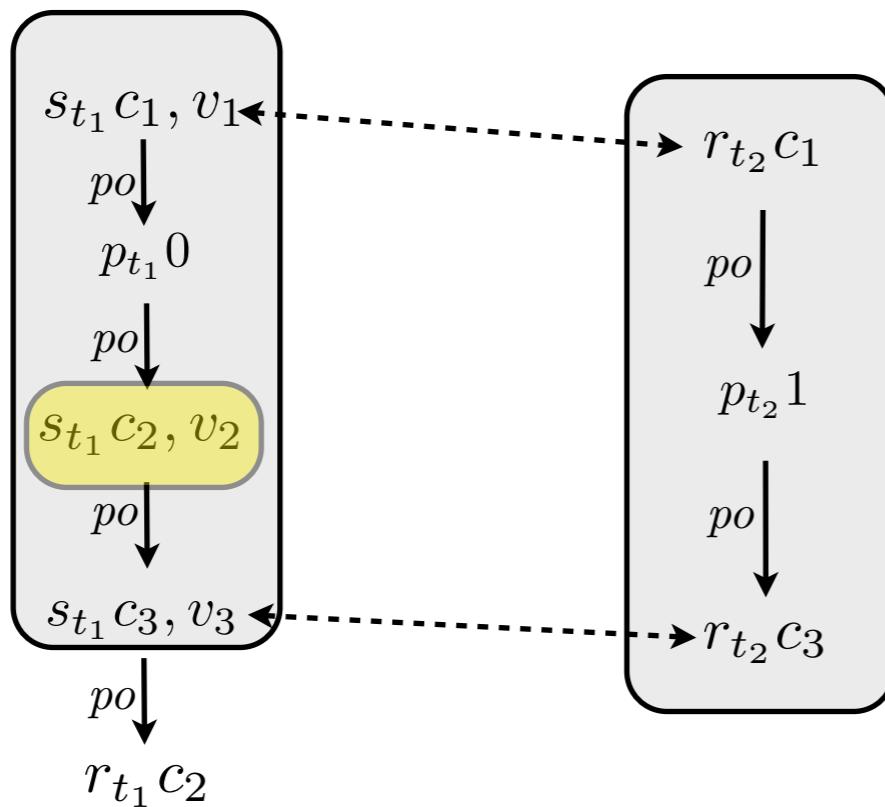
# Well-formedness -> Speculation

*Track executions to see if they become ill-formed (rollback) or turn into CML executions (commit)*



# Well-formedness -> Speculation


*Track executions to see if they become ill-formed (rollback) or turn into CML executions (commit)*



A well-formed execution that *can lead* to a CML execution


# Well-formedness -> Speculation

*Track executions to see if they become ill-formed (rollback) or turn into CML executions (commit)*



# Well-formedness -> Speculation

*Track executions to see if they become ill-formed (rollback) or turn into CML executions (commit)*




# Implementation: Overview

# Implementation: Overview


- Rx-CML: Relaxed CML

- ★ MultiMLton with distribution support
- ★ Rx-CML application = {Instances}
- ★ Supports full CML
- ★ Built-in serialization (immutable values and function closures)
- ★ Transport layer is ZeroMQ




# Implementation: Overview

- Rx-CML: Relaxed CML
  - ★ MultiMLton with distribution support
  - ★ Rx-CML application = {Instances}
  - ★ Supports full CML
  - ★ Built-in serialization (immutable values and function closures)
  - ★ Transport layer is ZeroMQ
- Check the integrity of the speculative actions on-the-fly




# Implementation: Overview

- Rx-CML: Relaxed CML
  - ★ MultiMLton with distribution support
  - ★ Rx-CML application = {Instances}
  - ★ Supports full CML
  - ★ Built-in serialization (immutable values and function closures)
  - ★ Transport layer is ZeroMQ
- Check the integrity of the speculative actions on-the-fly
  - ★ Build a dependence graph that captures happens-before relation




# Implementation: Overview

- Rx-CML: Relaxed CML
  - ★ MultiMLton with distribution support
  - ★ Rx-CML application = {Instances}
  - ★ Supports full CML
  - ★ Built-in serialization (immutable values and function closures)
  - ★ Transport layer is ZeroMQ
- Check the integrity of the speculative actions on-the-fly
  - ★ Build a dependence graph that captures happens-before relation
    - ◆ Same structure as an axiomatic execution




# Implementation: Overview

- Rx-CML: Relaxed CML
  - ★ MultiMLton with distribution support
  - ★ Rx-CML application = {Instances}
  - ★ Supports full CML
  - ★ Built-in serialization (immutable values and function closures)
  - ★ Transport layer is ZeroMQ
- Check the integrity of the speculative actions on-the-fly
  - ★ Build a dependence graph that captures happens-before relation
    - ◆ Same structure as an axiomatic execution
  - ★ Automatically check dependence graph integrity before an observable action (ref cell accesses, system calls, FFI, etc)



# Implementation: Overview

- Rx-CML: Relaxed CML
  - ★ MultiMLton with distribution support
  - ★ Rx-CML application = {Instances}
  - ★ Supports full CML
  - ★ Built-in serialization (immutable values and function closures)
  - ★ Transport layer is ZeroMQ
- Check the integrity of the speculative actions on-the-fly
  - ★ Build a dependence graph that captures happens-before relation
    - ◆ Same structure as an axiomatic execution
  - ★ Automatically check dependence graph integrity before an observable action (ref cell accesses, system calls, FFI, etc)
  - ★ Roll-back ill-formed executions, re-execute non-speculatively



# Channel Consistency (1)

# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?

# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance

# Channel Consistency (1)

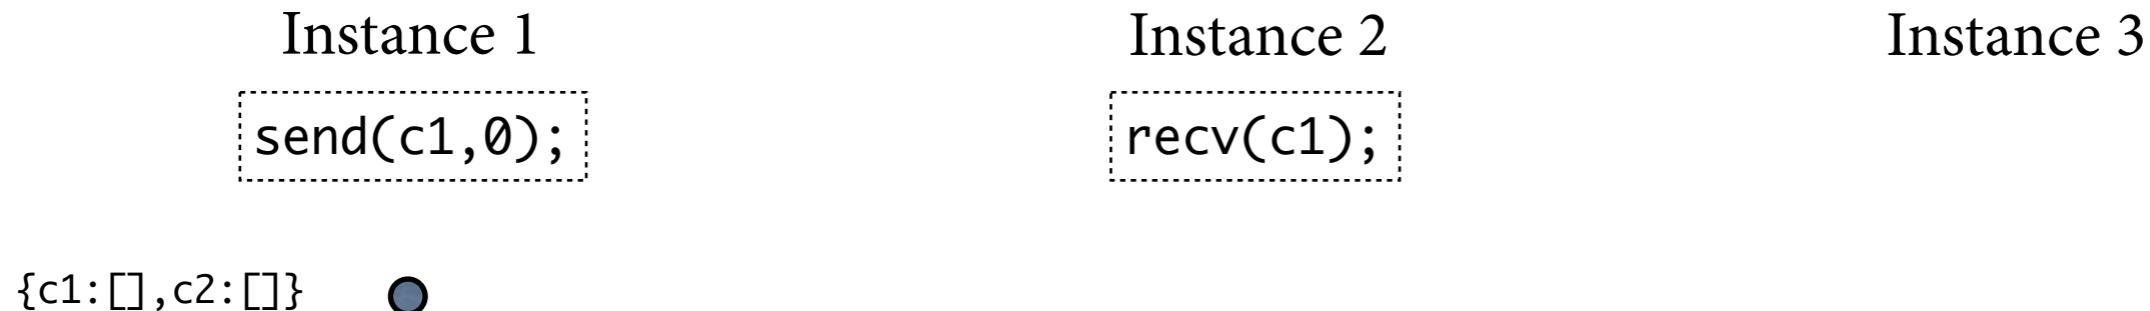
- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance
  - ★ Maintains a replica of CML channel

# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance
  - ★ Maintains a replica of CML channel
  - ★ Utilize speculative execution to recover from inconsistent state


# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance
  - ★ Maintains a replica of CML channel
  - ★ Utilize speculative execution to recover from inconsistent state




# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance
  - ★ Maintains a replica of CML channel
  - ★ Utilize speculative execution to recover from inconsistent state




# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance
  - ★ Maintains a replica of CML channel
  - ★ Utilize speculative execution to recover from inconsistent state




# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance
  - ★ Maintains a replica of CML channel
  - ★ Utilize speculative execution to recover from inconsistent state




# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance
  - ★ Maintains a replica of CML channel
  - ★ Utilize speculative execution to recover from inconsistent state



# Channel Consistency (1)

- Single consistent channel image across all instances without coherence?
  - ★ Communication manager thread @ every instance
  - ★ Maintains a replica of CML channel
  - ★ Utilize speculative execution to recover from inconsistent state



# Channel Consistency (2)

Instance 1

```
send(c2,1);
```

Instance 2

```
recv(c2);
```

Instance 3

```
recv(c2);
```

# Channel Consistency (2)

Instance 1


send(c2,1);

Instance 2


recv(c2);

Instance 3


recv(c2);




# Channel Consistency (2)




# Channel Consistency (2)




# Channel Consistency (2)



# Channel Consistency (2)



# Channel Consistency (2)



# Speculative Execution

# Speculative Execution

- Consistent, replicated dependence graph @ each instance
  - ★ *Snoop on match messages* from communication manager
  - ★ Broadcast thread spawn and thread join messages

# Speculative Execution

- Consistent, replicated dependence graph @ each instance
  - ★ *Snoop on match messages* from communication manager
  - ★ Broadcast thread spawn and thread join messages
- Well-formedness check *is local to the instance!*
  - ★ GC dependence graph on successful well-formedness check

# Speculative Execution

- Consistent, replicated dependence graph @ each instance
  - ★ *Snoop on match messages* from communication manager
  - ★ Broadcast thread spawn and thread join messages
- Well-formedness check *is local to the instance!*
  - ★ GC dependence graph on successful well-formedness check
- Automatic checkpointing
  - ★ 1 continuation per thread
  - ★ *Uncoordinated!* - thread local - does not require barriers

# Speculative Execution

- Consistent, replicated dependence graph @ each instance
  - ★ *Snoop on match messages* from communication manager
  - ★ Broadcast thread spawn and thread join messages
- Well-formedness check *is local to the instance!*
  - ★ GC dependence graph on successful well-formedness check
- Automatic checkpointing
  - ★ 1 continuation per thread
  - ★ *Uncoordinated!* - thread local - does not require barriers
- Remediation
  - ★ *Uncoordinated!* - Transitively inform each mis-speculated thread to rollback
  - ★ *Check-point (Continuation) + Log-based (Dependence graph)* recovery
  - ★ Rollback to last checkpoint, replay correct speculative actions
  - ★ Continues non-speculatively until next observable action = Progress

# Results

# Results


- Optimistic OLTP
  - ★ Distributed version of STAMP Vacation benchmark
  - ★ Database split into 64 shards, with concurrent transaction requests from geo-distributed clients
  - ★ Uses explicit lock servers -> Rx-CML executes transactions optimistically

# Results

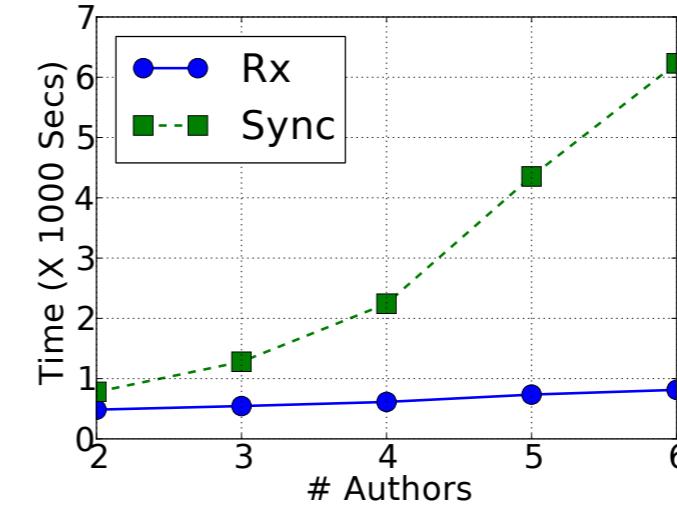
- Optimistic OLTP
  - ★ Distributed version of STAMP Vacation benchmark
  - ★ Database split into 64 shards, with concurrent transaction requests from geo-distributed clients
  - ★ Uses explicit lock servers -> Rx-CML executes transactions optimistically
- P2P Collaborative editing
  - ★ Simulates concurrent document editing (operational transformation)
  - ★ Total order broadcast - built out of synchronous events + choice combinator.

# Results

- Optimistic OLTP
  - ★ Distributed version of STAMP Vacation benchmark
  - ★ Database split into 64 shards, with concurrent transaction requests from geo-distributed clients
  - ★ Uses explicit lock servers -> Rx-CML executes transactions optimistically
- P2P Collaborative editing
  - ★ Simulates concurrent document editing (operational transformation)
  - ★ Total order broadcast - built out of synchronous events + choice combinator.



*OLTP*



*Collaborative Editing*

- Rx-CML was **5.8X to 7.6X** faster than the synchronous version
  - ★ 9-17% of communications were mis-specified.

# Conclusion

# Conclusion

- Composable synchronous events vs. High latency compute cloud

# Conclusion

- Composable synchronous events vs. High latency compute cloud
- Rx-CML (Relaxed CML)
  - ★ optimistic concurrency control for CML
  - ★ reason synchronously, but implement asynchronously
  - ★ retain simplicity and composability, but gain performance

# Conclusion

- Composable synchronous events vs. High latency compute cloud
- Rx-CML (Relaxed CML)
  - ★ optimistic concurrency control for CML
  - ★ reason synchronously, but implement asynchronously
  - ★ retain simplicity and composability, but gain performance
- Distributed implementation of MultiMLton
  - ★ Case studies demonstrate effectiveness of the approach

# Conclusion

- Composable synchronous events vs. High latency compute cloud
- Rx-CML (Relaxed CML)
  - ★ optimistic concurrency control for CML
  - ★ reason synchronously, but implement asynchronously
  - ★ retain simplicity and composability, but gain performance
- Distributed implementation of MultiMLton
  - ★ Case studies demonstrate effectiveness of the approach
- Future Work - Fault tolerance
  - ★ Make checkpoints and dependence graph resilient
  - ★ Treat failures as mis-speculations -> rollback to last saved checkpoint

# Questions?



<http://multimlton.cs.purdue.edu>