

Relational Reasoning for Mergeable Replicated Data Types

KC Sivaramakrishnan

joint work with

Gowtham Kaki, Swarn Priya, Suresh Jagannathan







VISA®

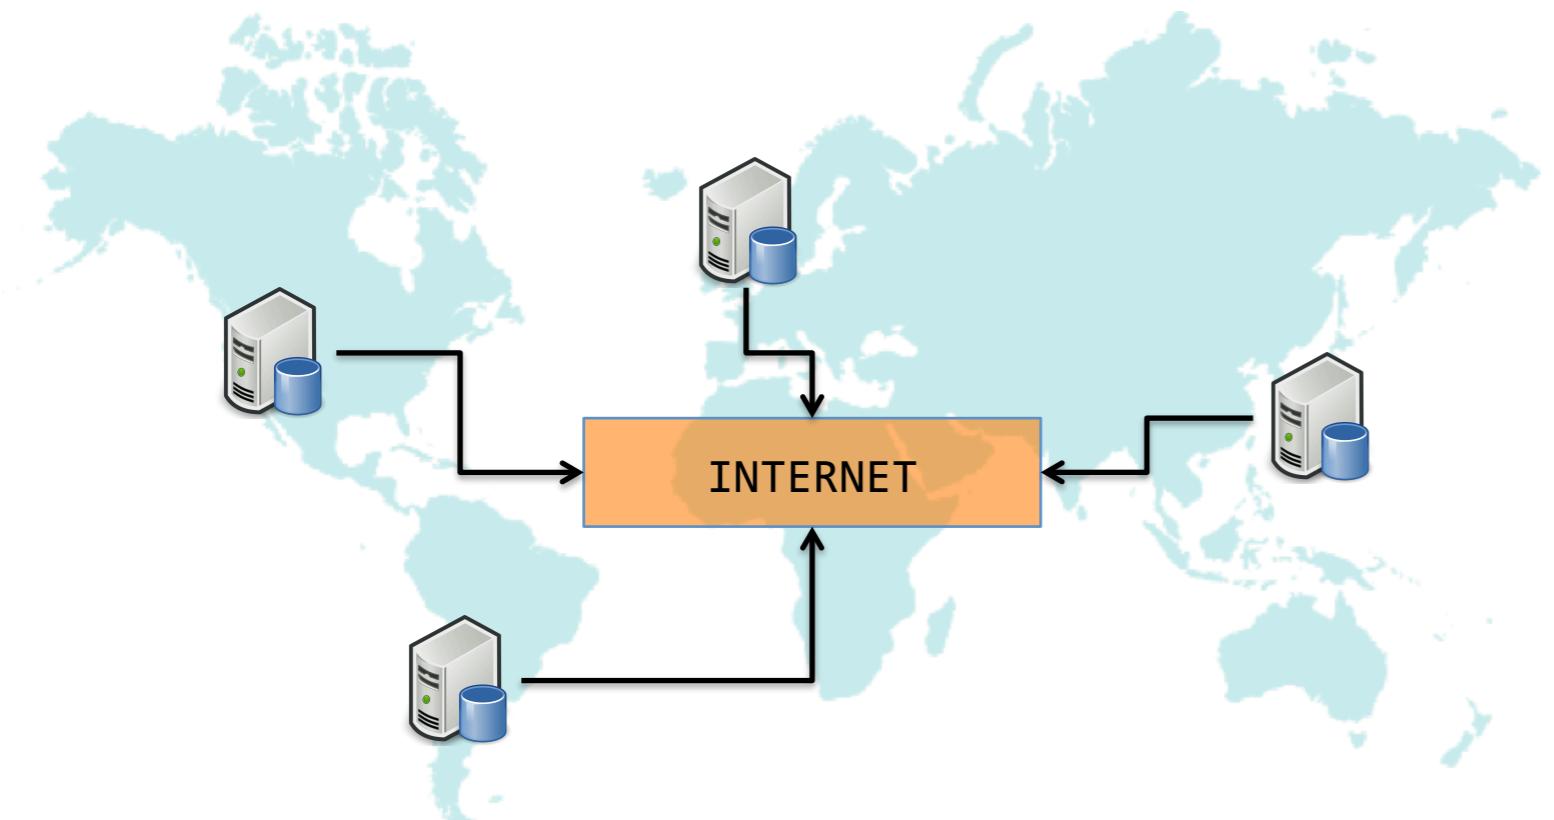
ebay™



YouTube

NETFLIX

HSBC ◀▶





VISA®

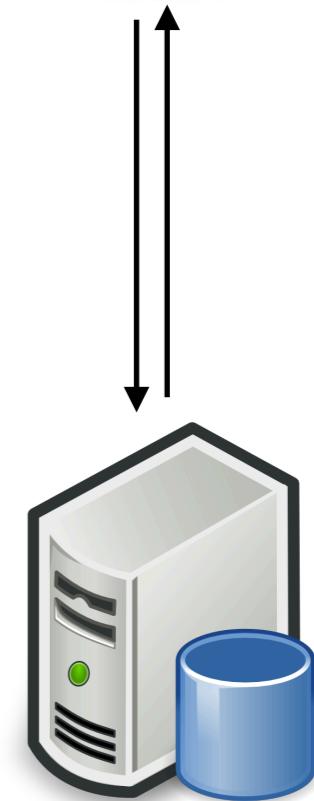
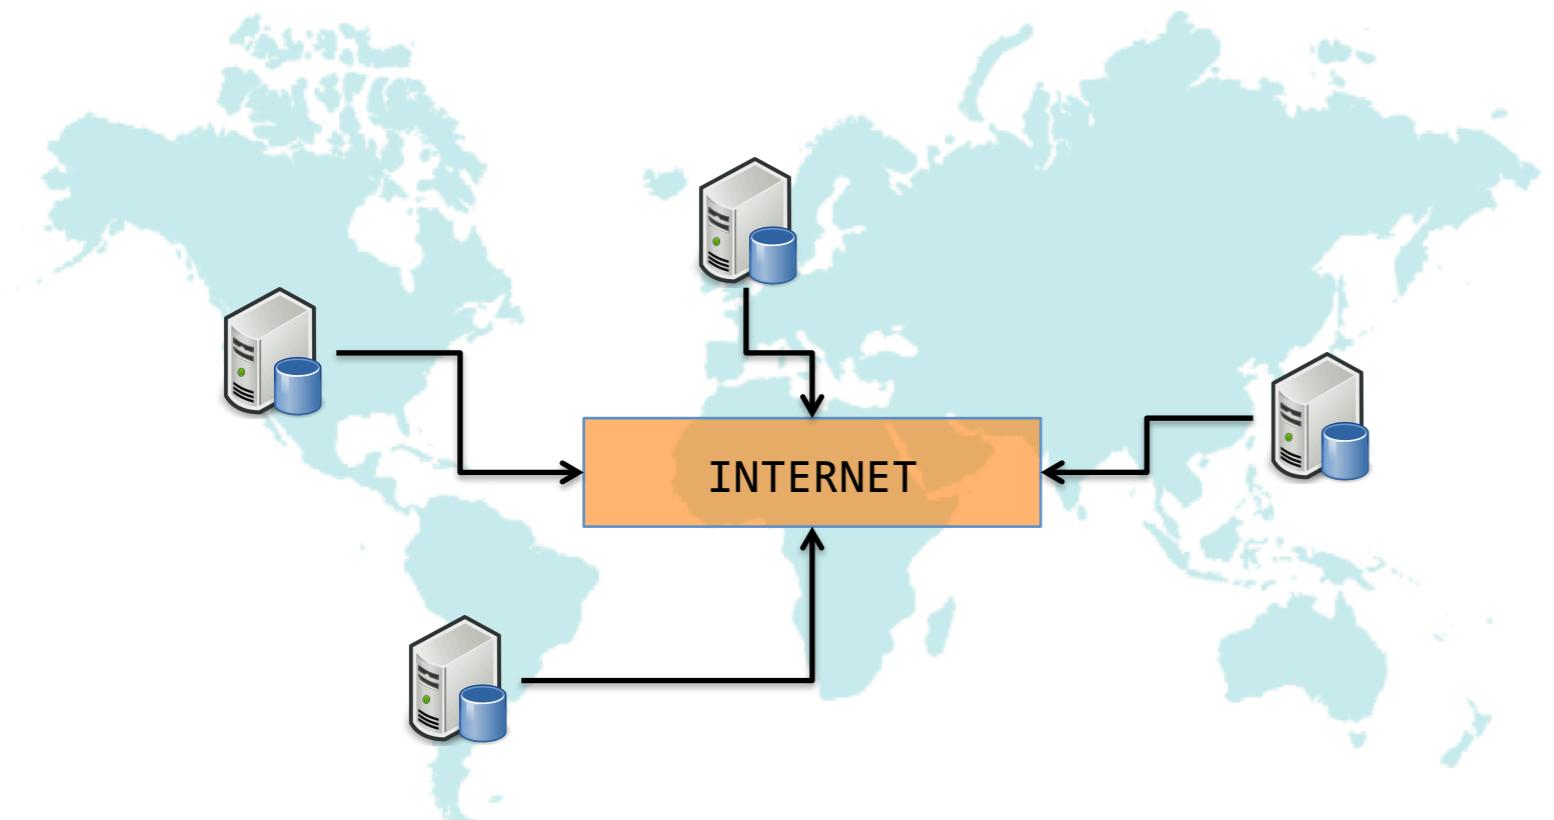
ebay™



NETFLIX

YouTube

HSBC ◀▶





VISA®

ebay™



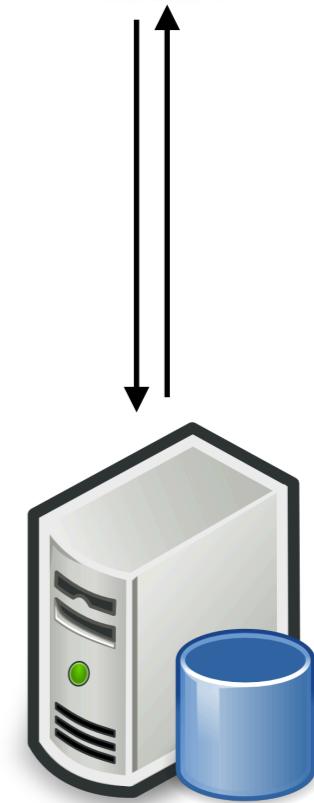
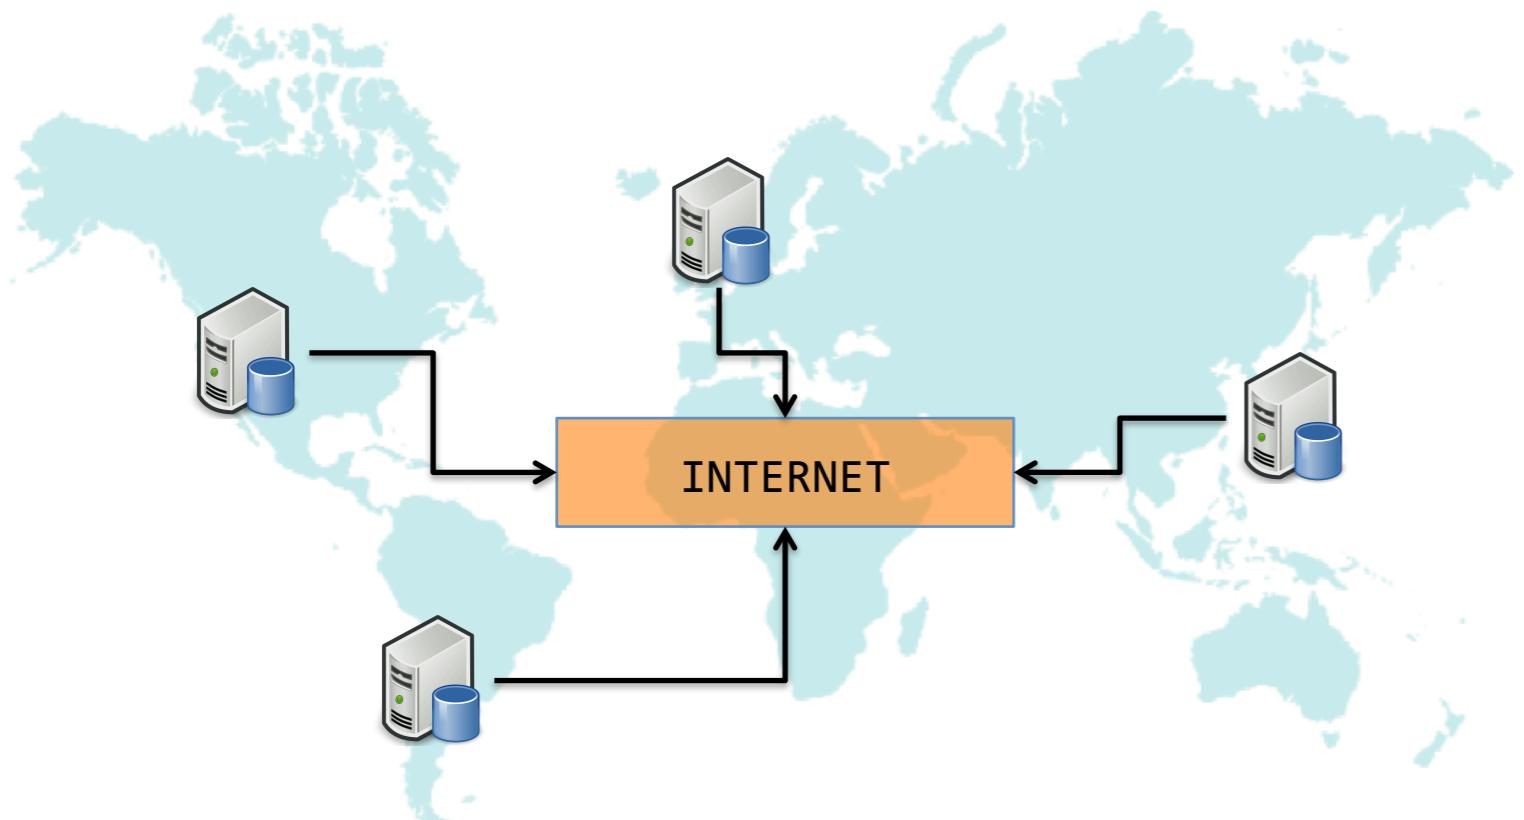
NETFLIX

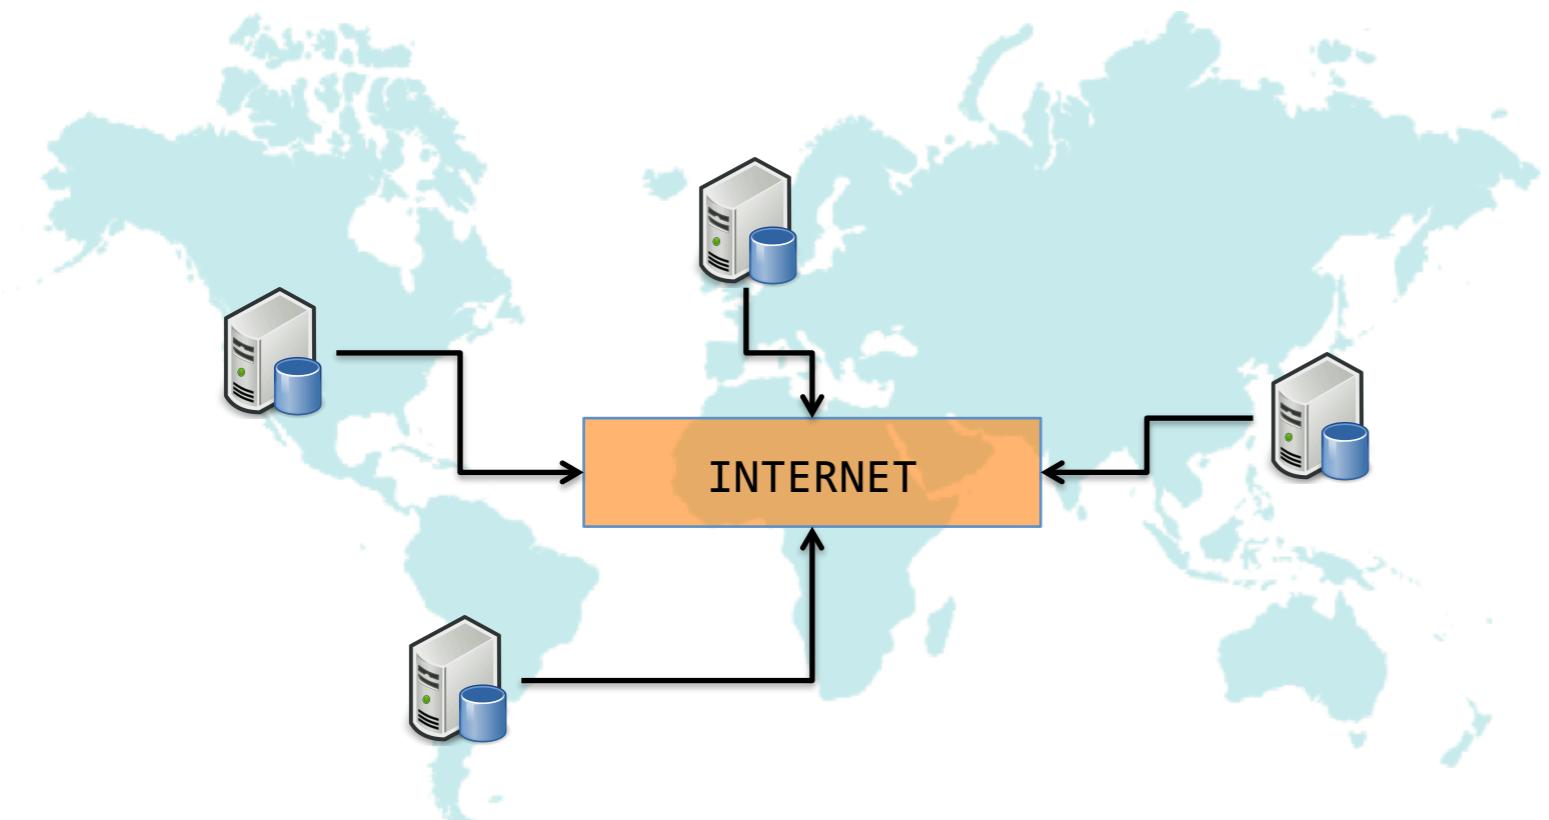
YouTube

HSBC ◀▶

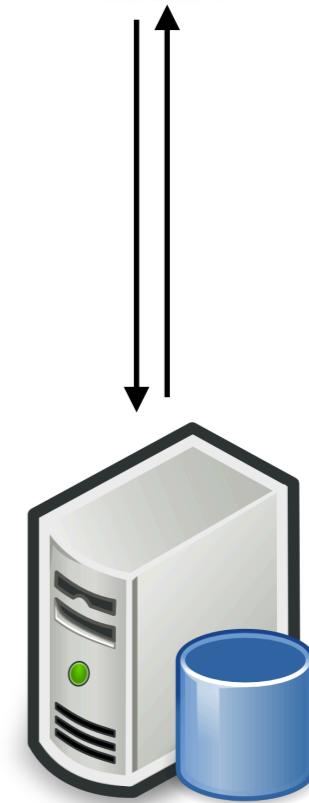


≠





≠



- Weak Consistency & Isolation

- Serializability
- Linearizability

When *system-level concerns* like replication & availability affect *application-level design* decisions, programming becomes *complicated*.



natefoster
@natefoster

v

Seems like Twitter wants me to follow this guy.

Who to follow



Doug Woos
@dougwoos

PhD student, joining @BrownCSDept as a lecturer in Fall 2019. Into programming languages, distributed systems, baseball, and other stuff. he/him

Follow



Doug Woos
@dougwoos

PhD student, joining @BrownCSDept as a lecturer in Fall 2019. Into programming languages, distributed systems, baseball, and other stuff. he/him

Follow



Doug Woos
@dougwoos

PhD student, joining @BrownCSDept as a lecturer in Fall 2019. Into programming languages, distributed systems, baseball, and other stuff. he/him

Follow

5:43 AM · May 14, 2019 · Twitter for iPhone

Sequential Counter

```
module Counter : sig
    type t
    val read : t -> int
    val add  : t -> int -> t
    val sub  : t -> int -> t
end = struct
    type t = int
    let read x = x
    let add x d = x + d
    let sub x d = x - d
end
```

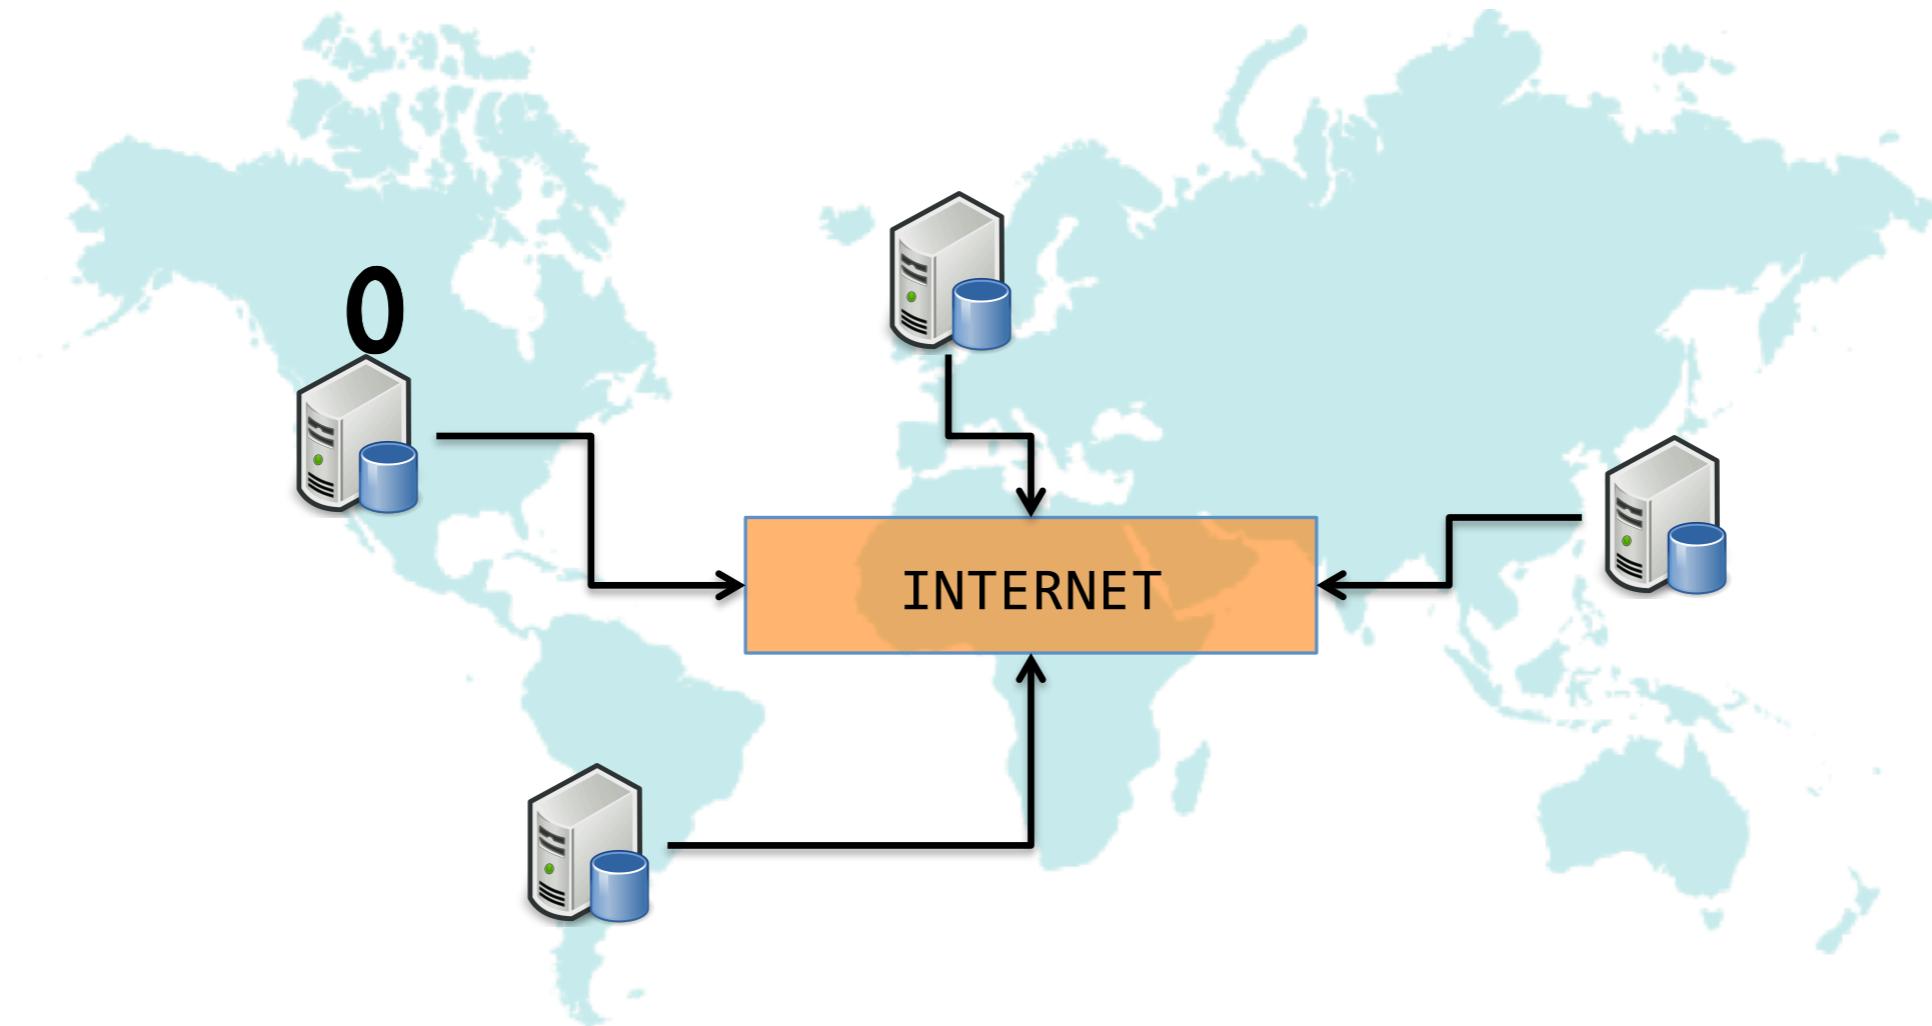
Sequential Counter

```
module Counter : sig
    type t
    val read : t -> int
    val add  : t -> int -> t
    val sub  : t -> int -> t
end = struct
    type t = int
    let read x = x
    let add x d = x + d
    let sub x d = x - d
end
```

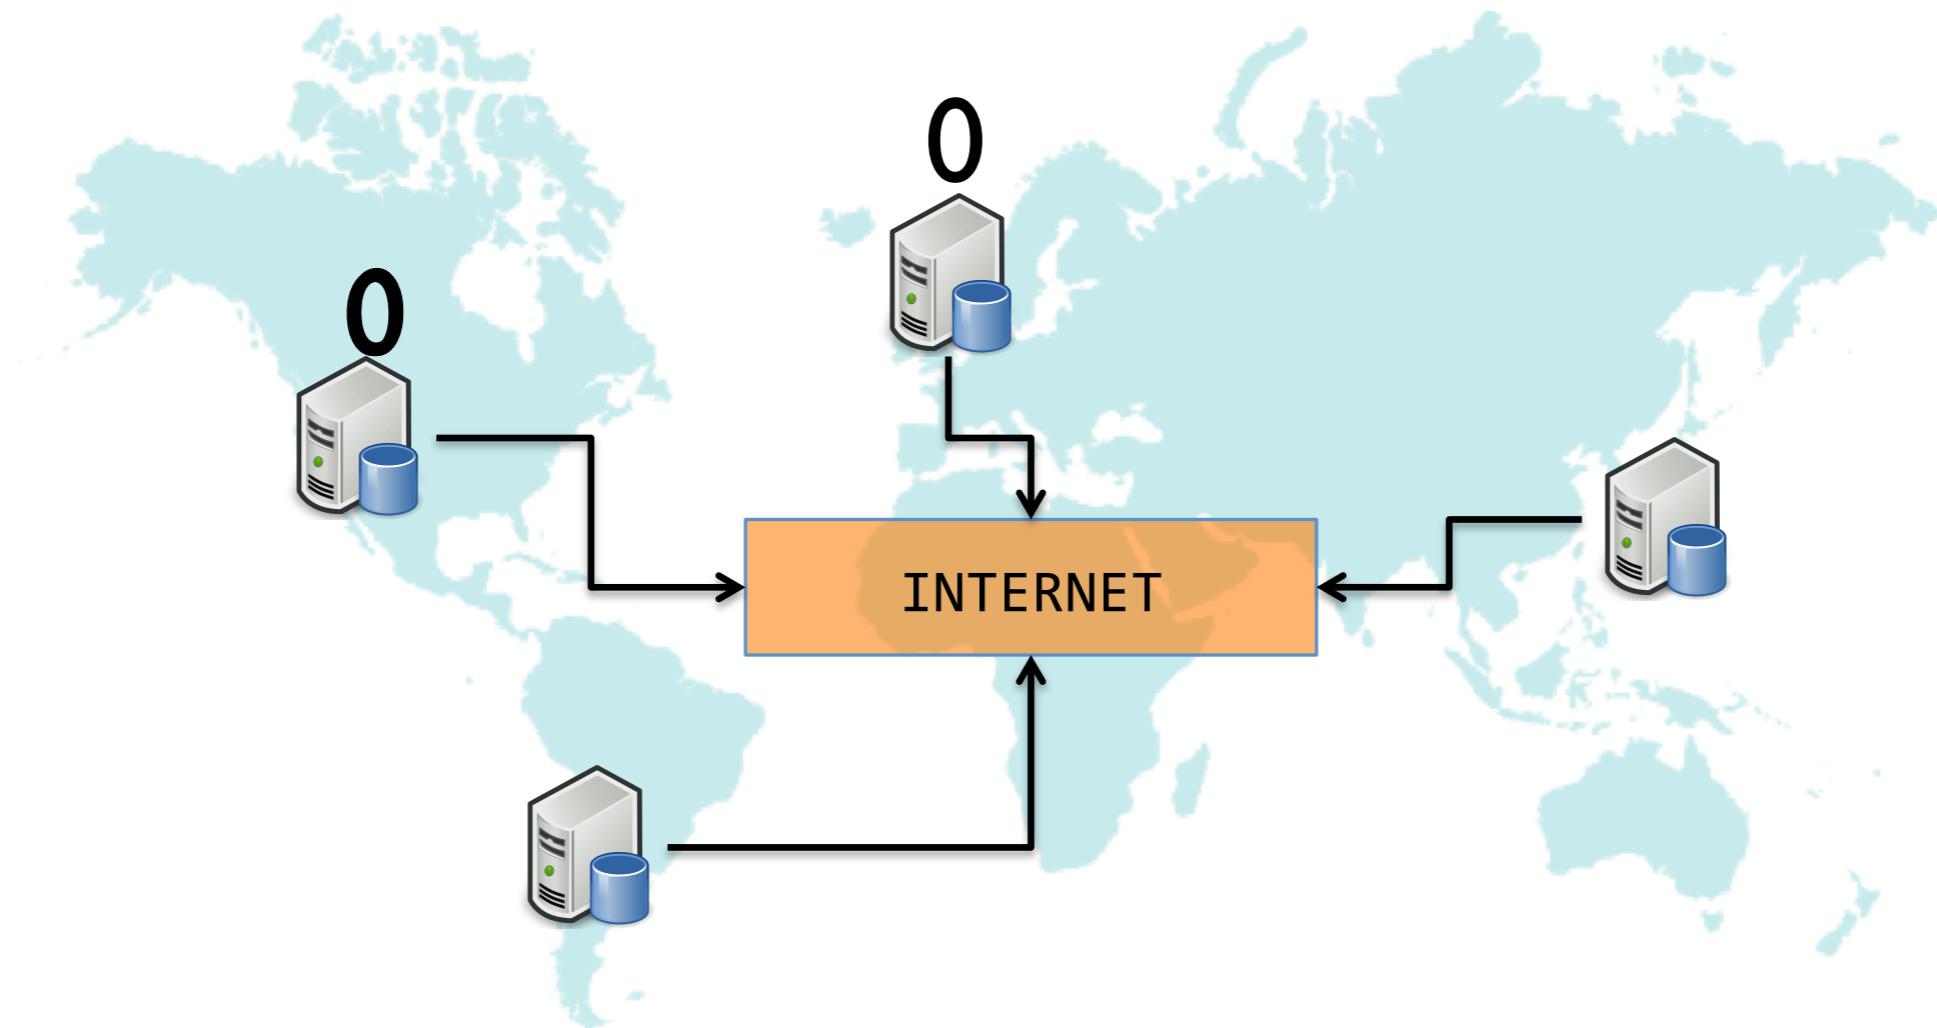
- Written in idiomatic style
- Composable

```
type counter_list = Counter.t list
```

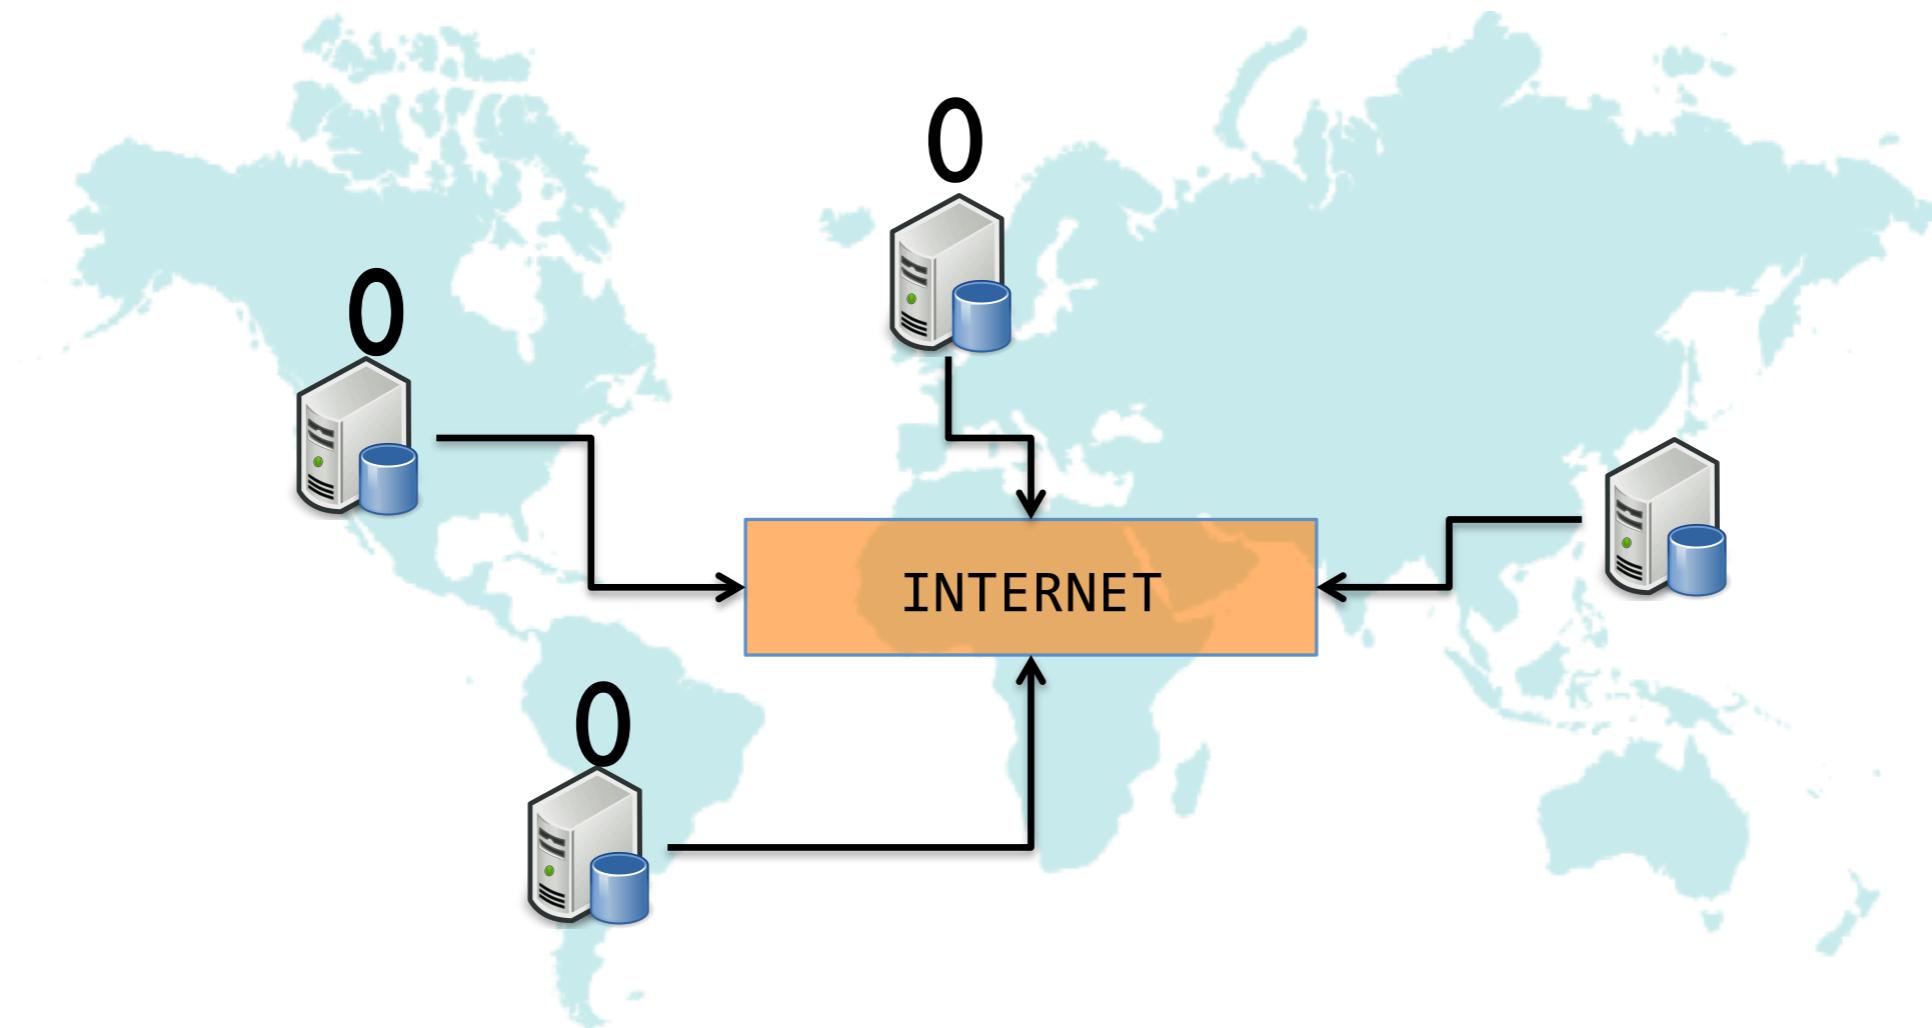
Replicated Counter



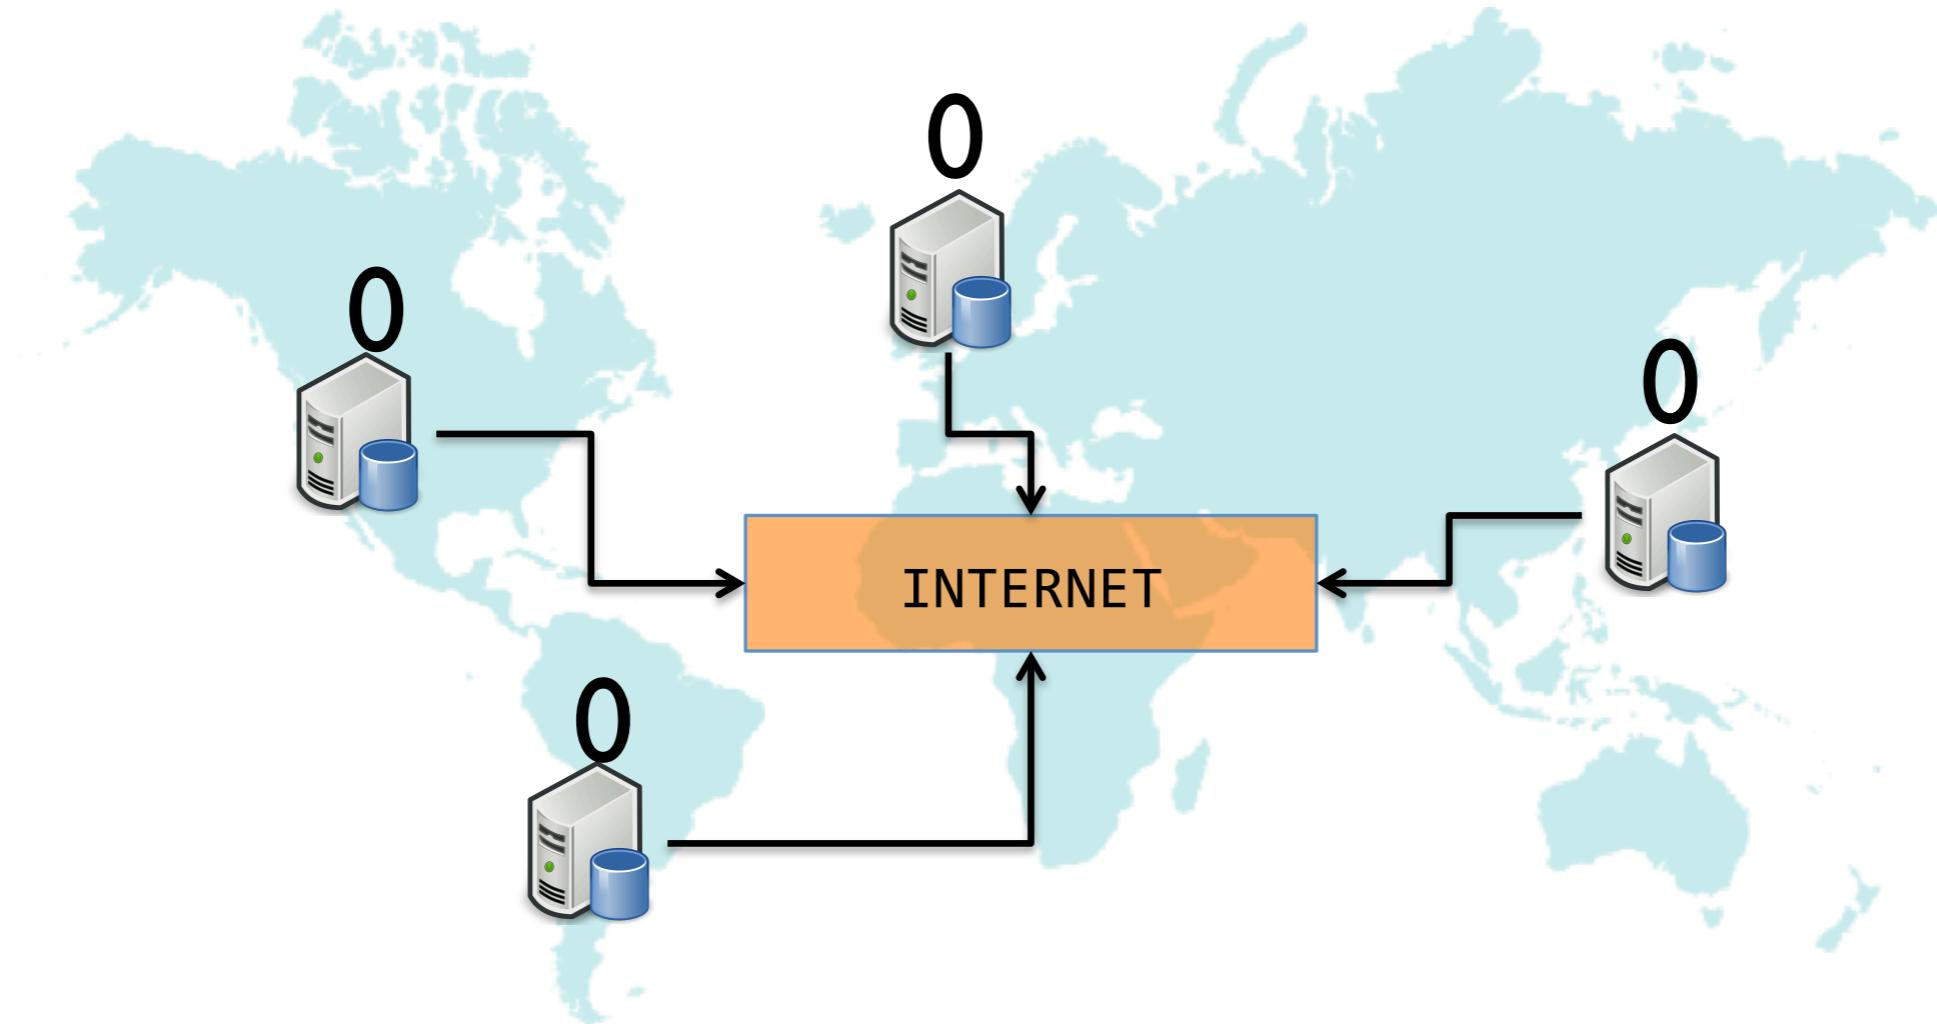
Replicated Counter



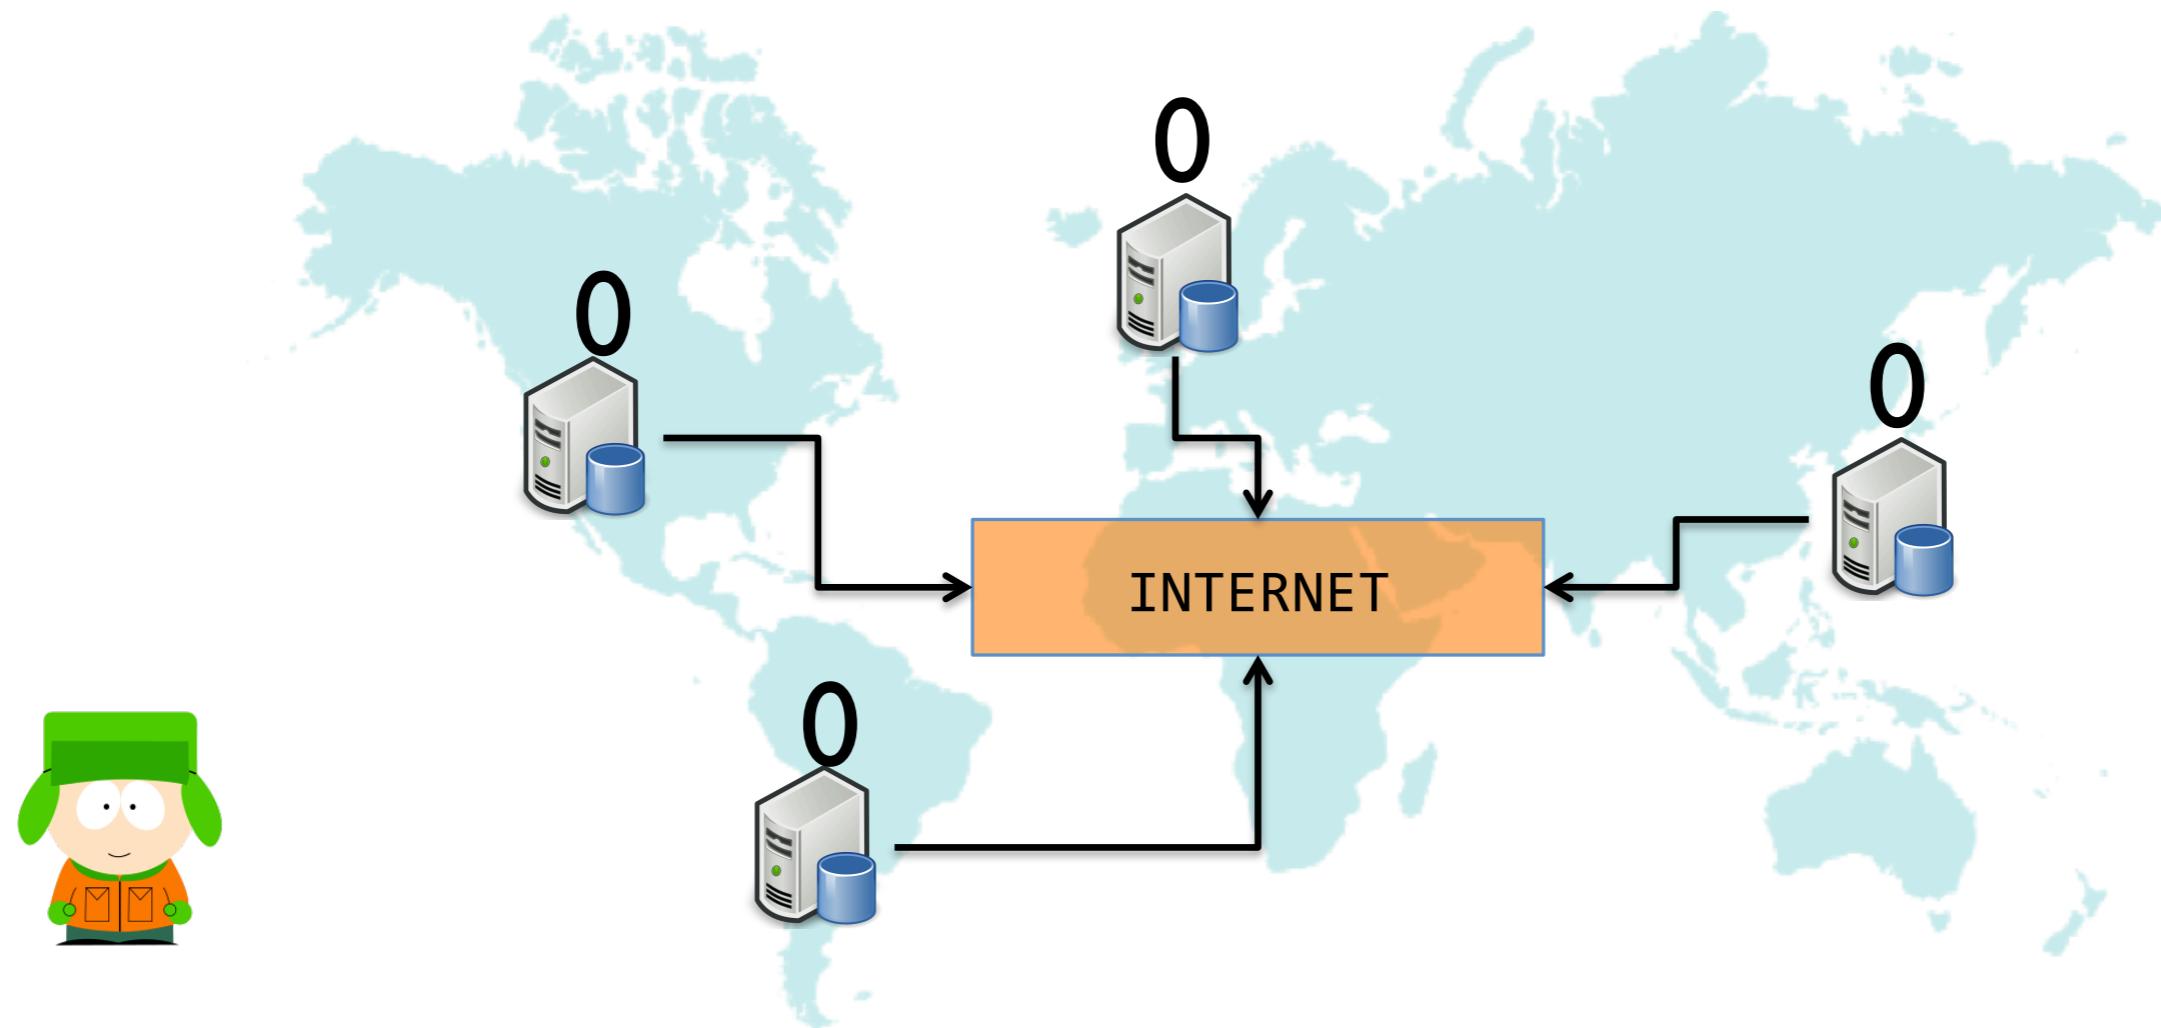
Replicated Counter



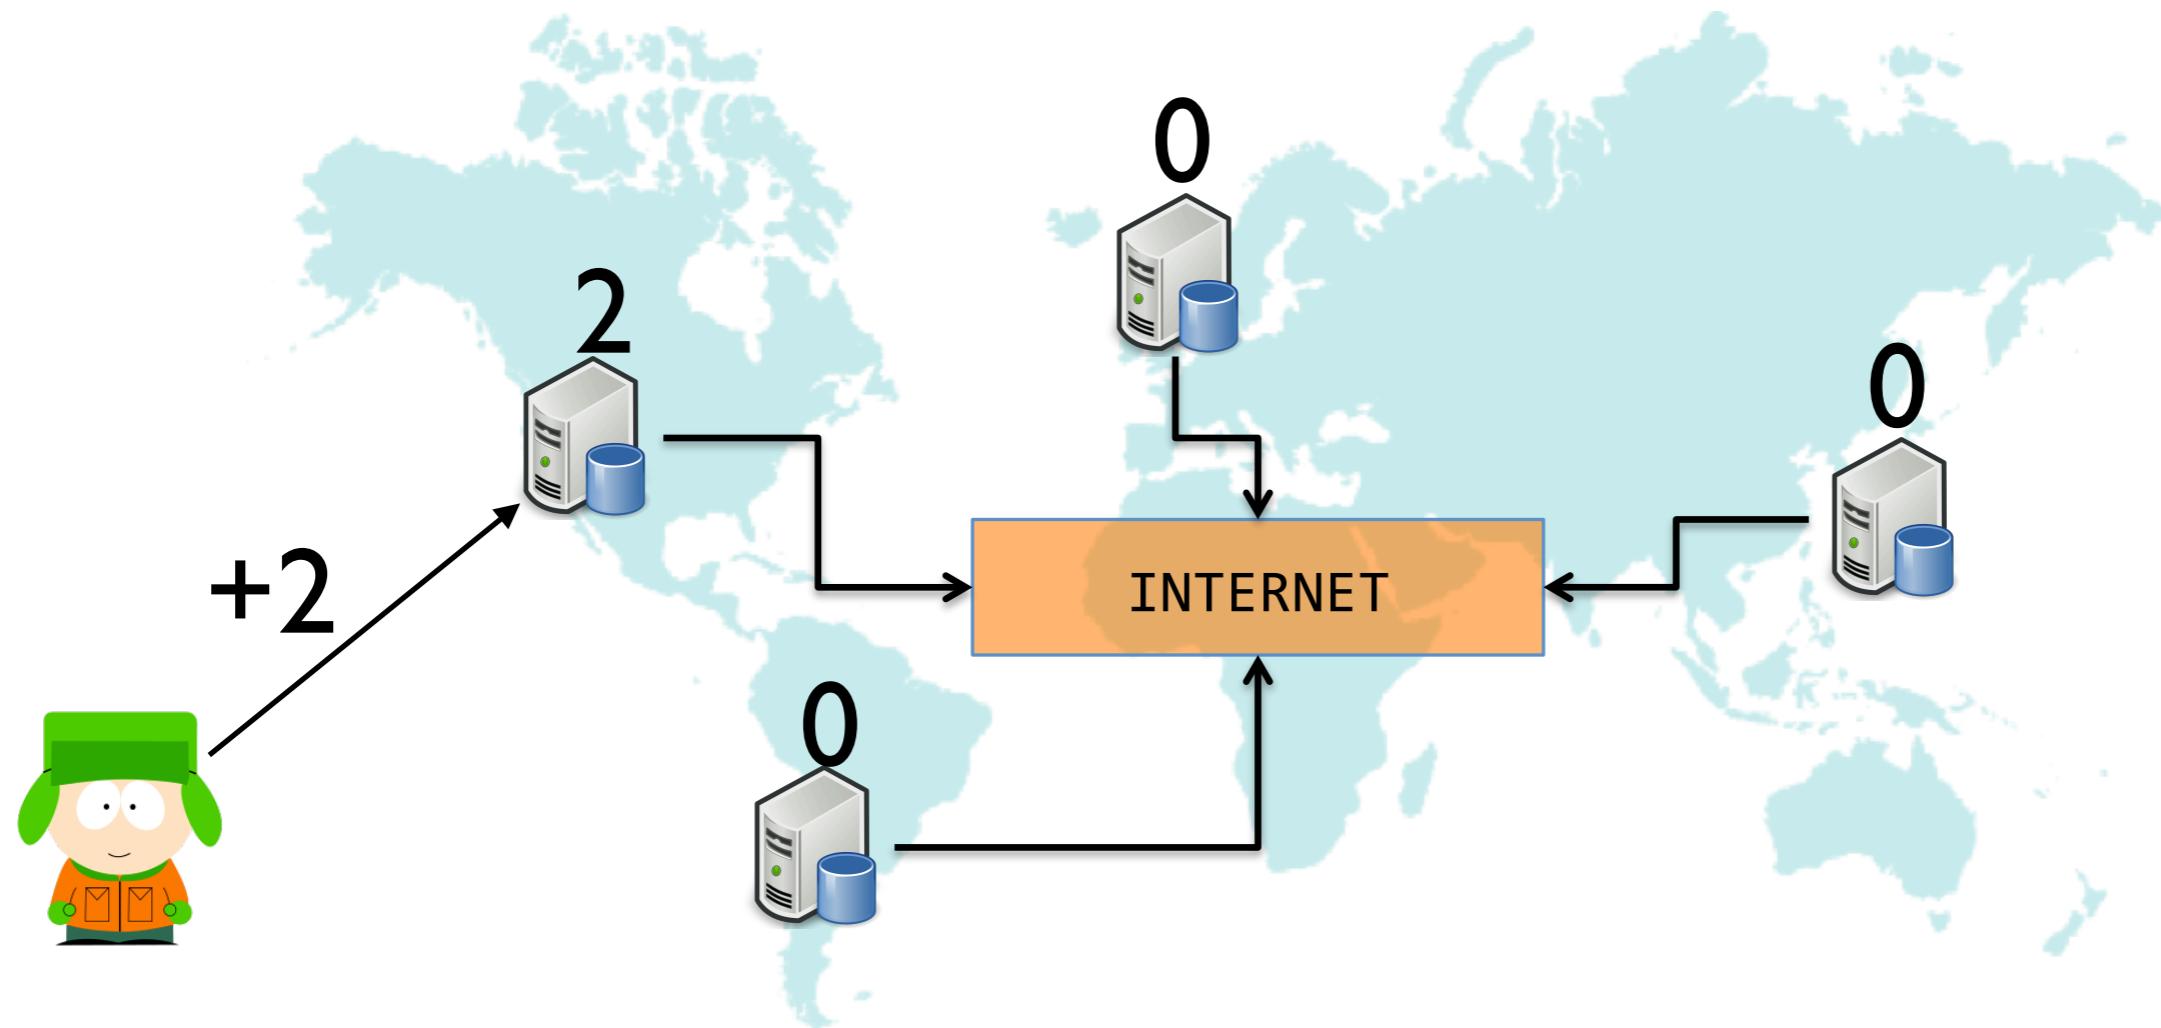
Replicated Counter



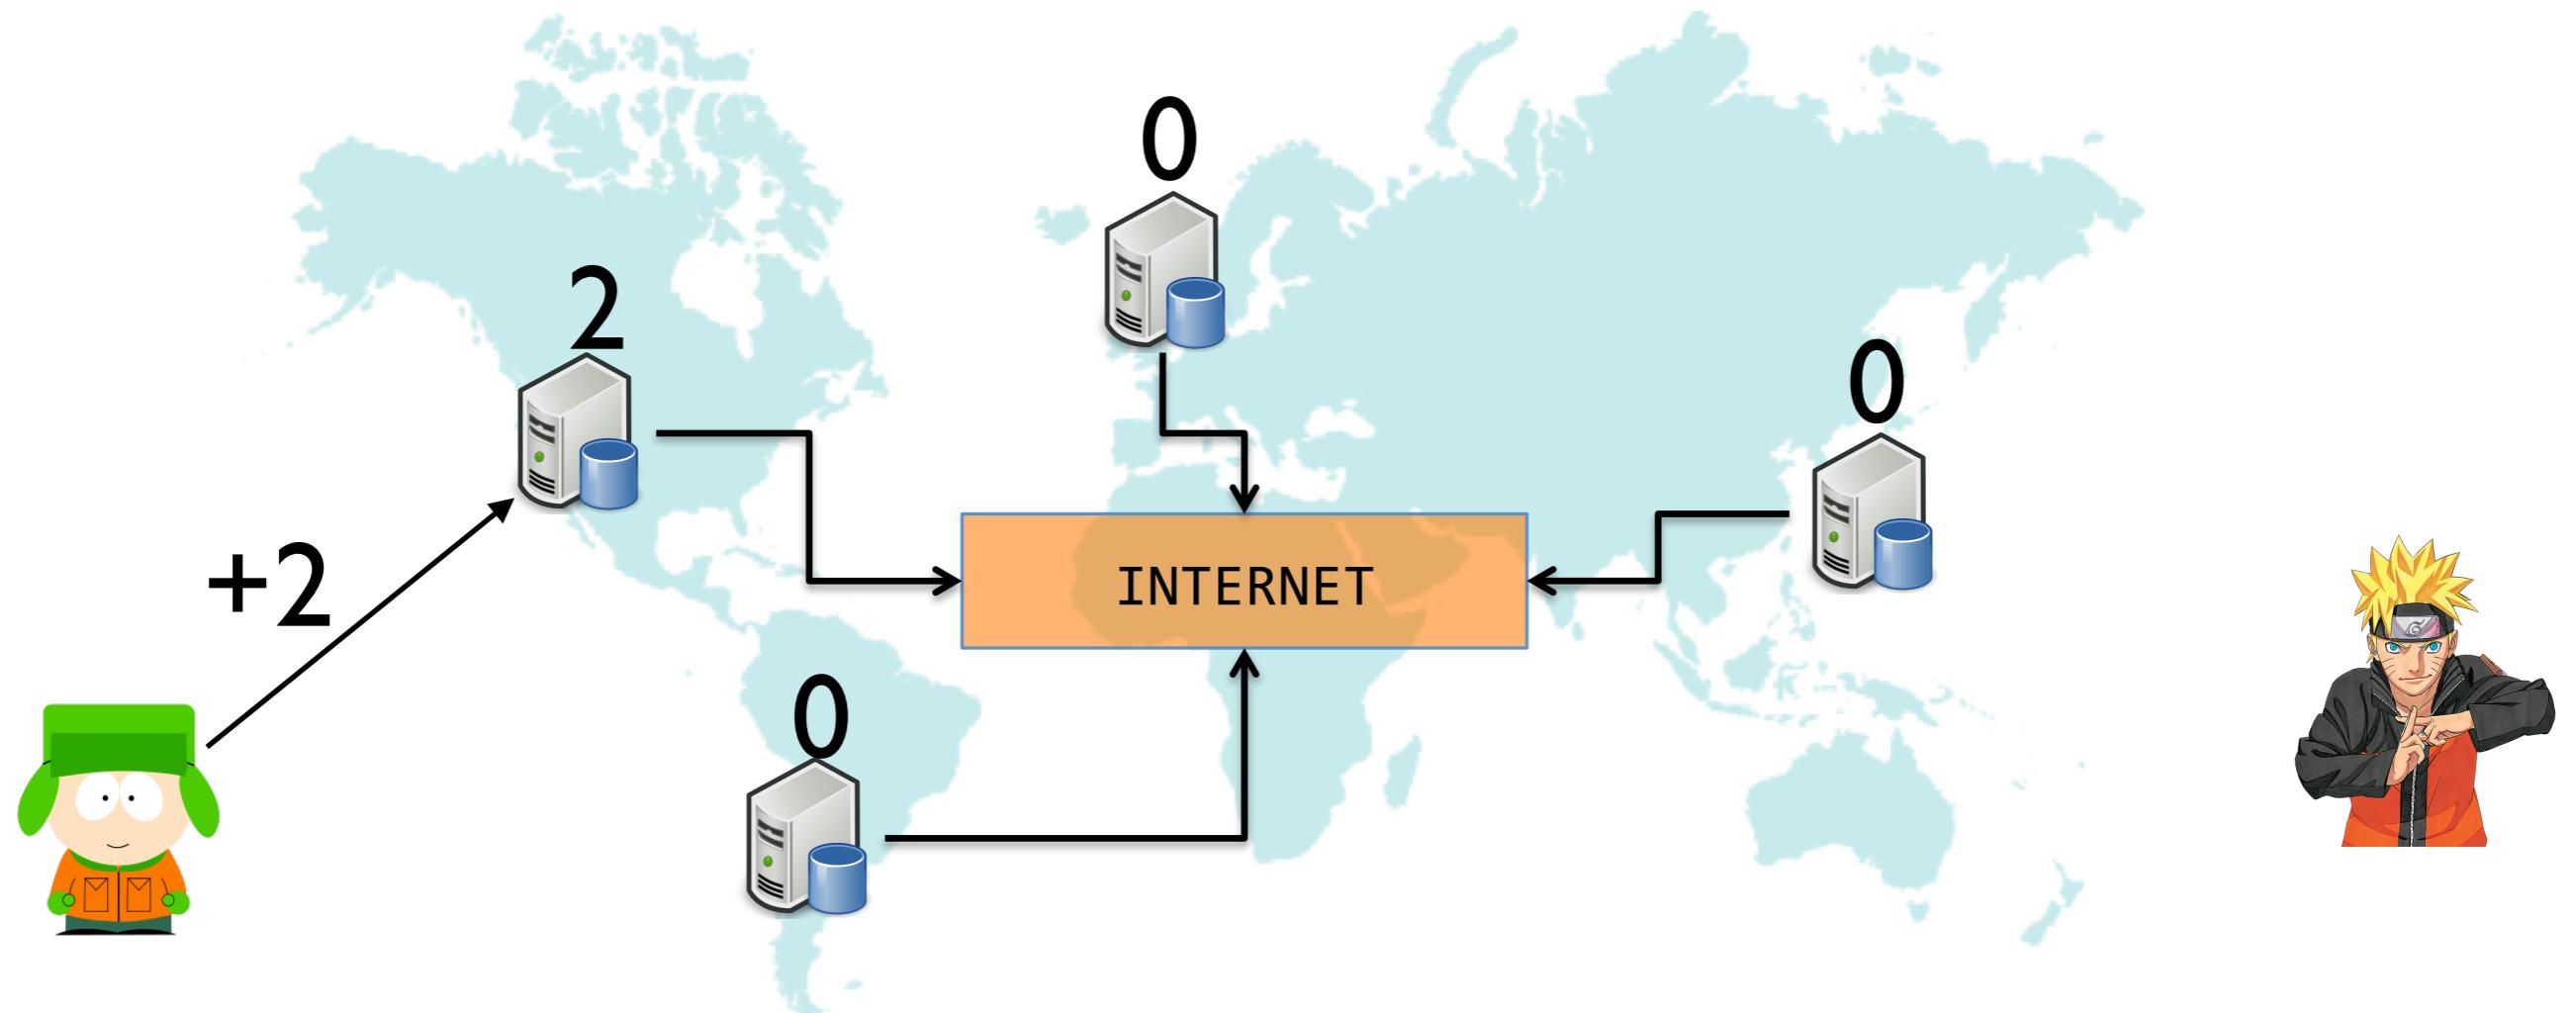
Replicated Counter



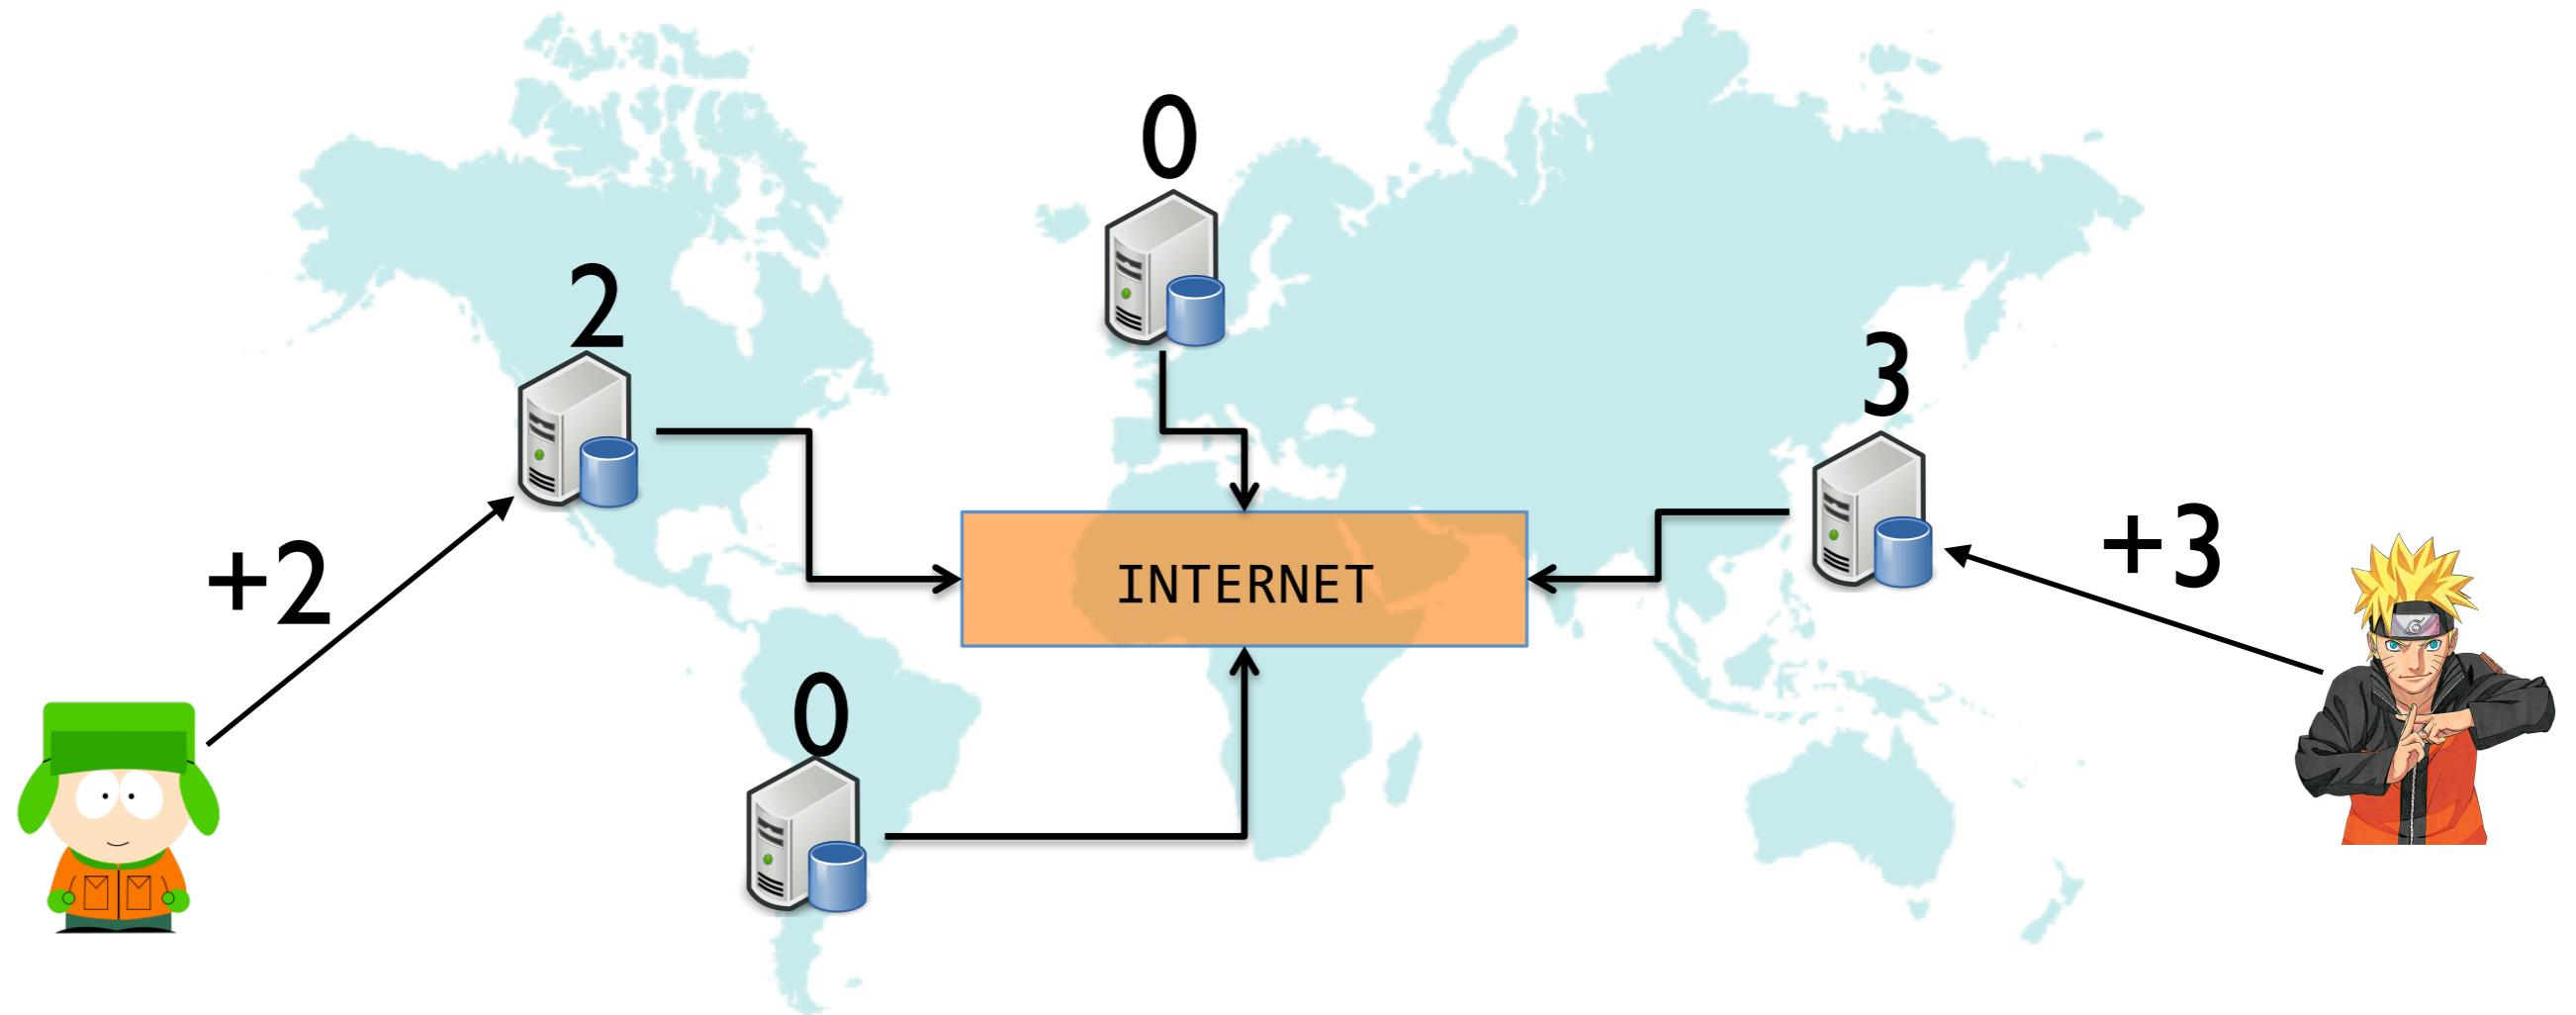
Replicated Counter



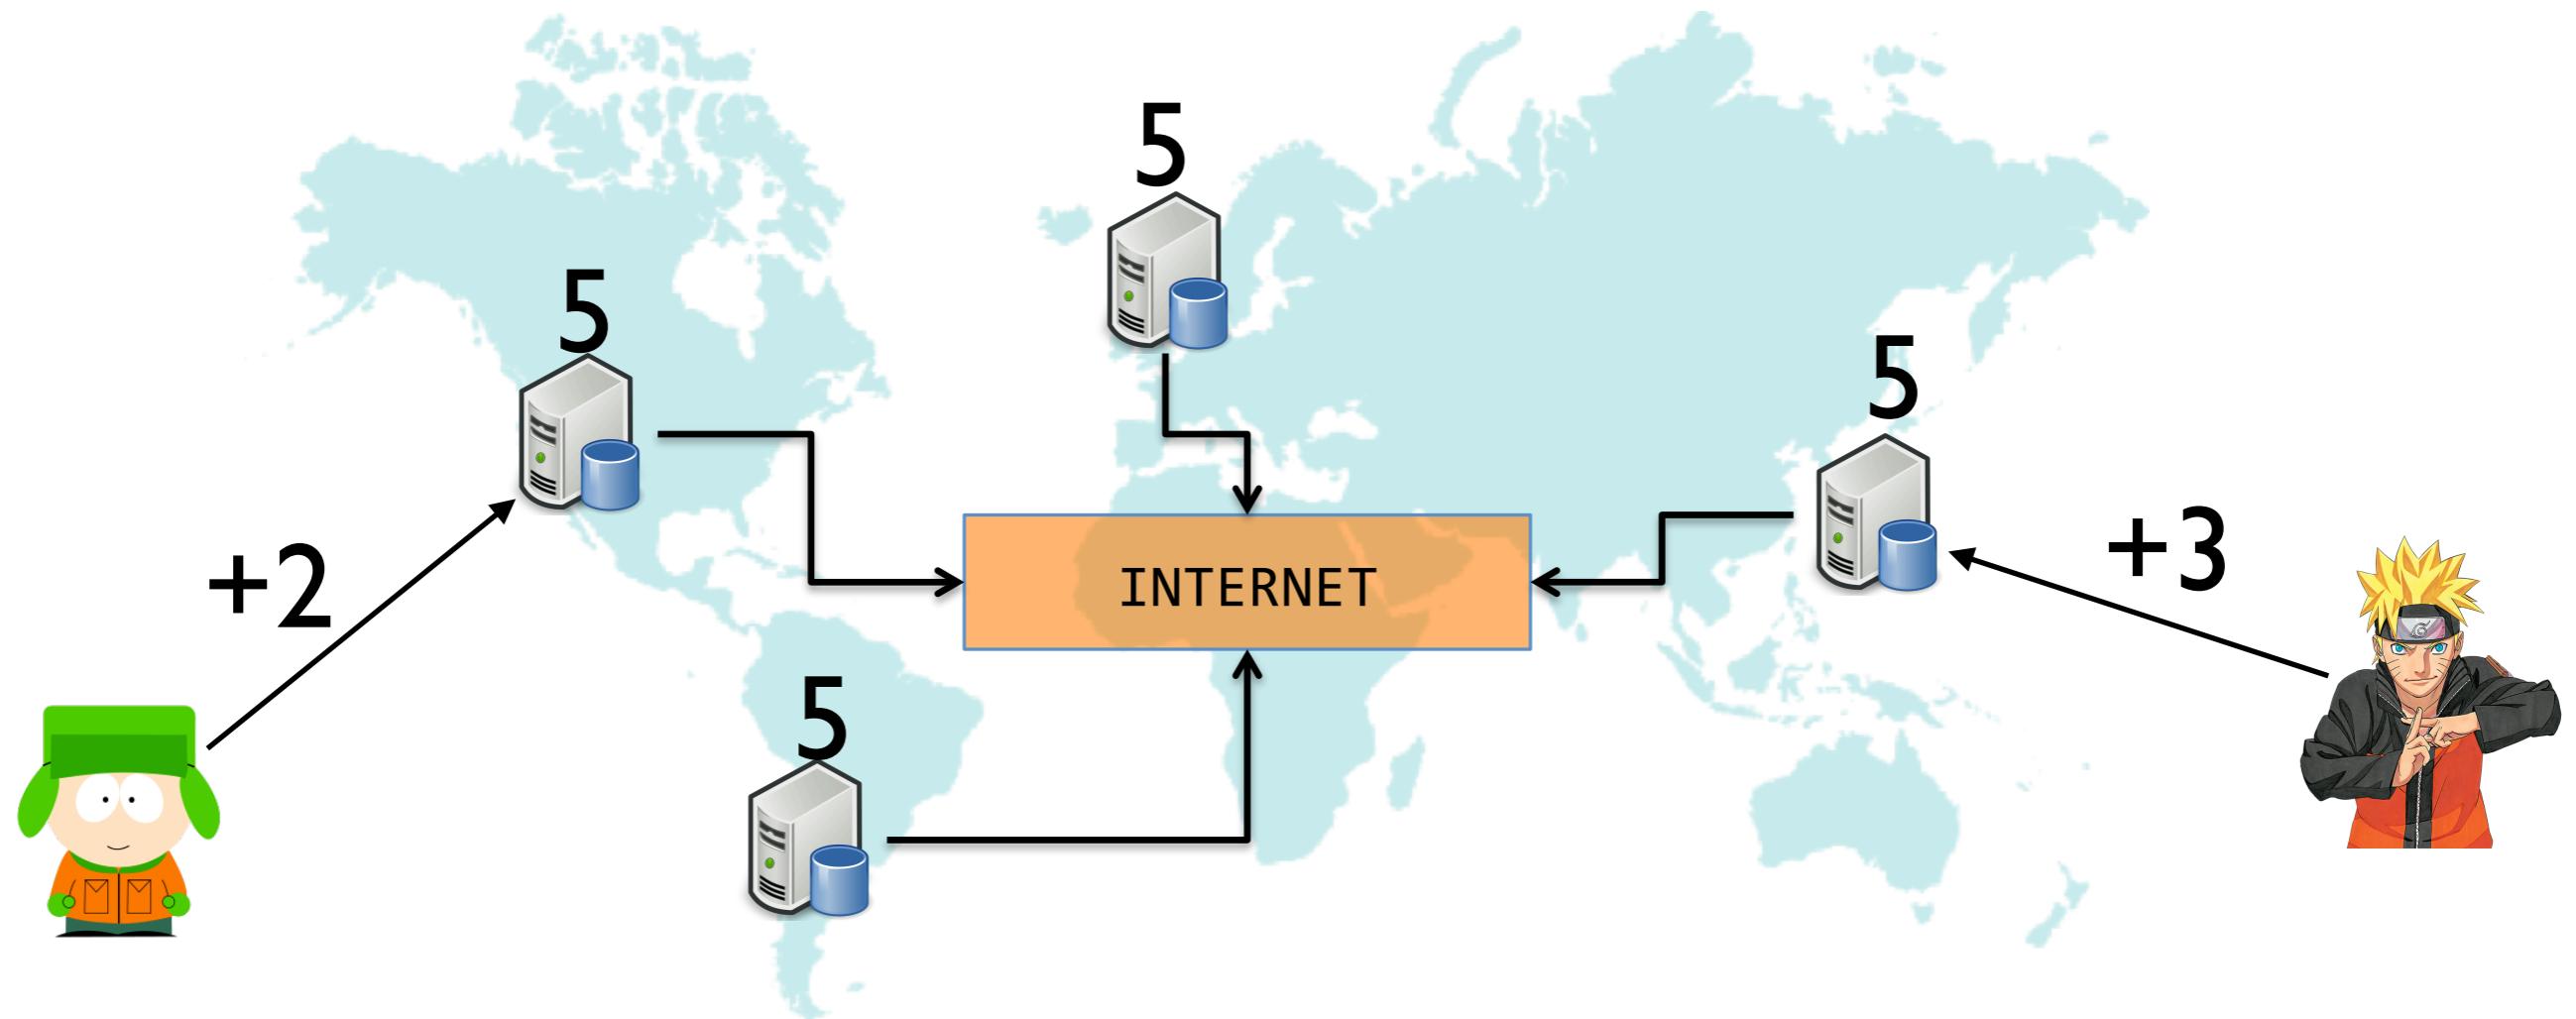
Replicated Counter



Replicated Counter



Replicated Counter

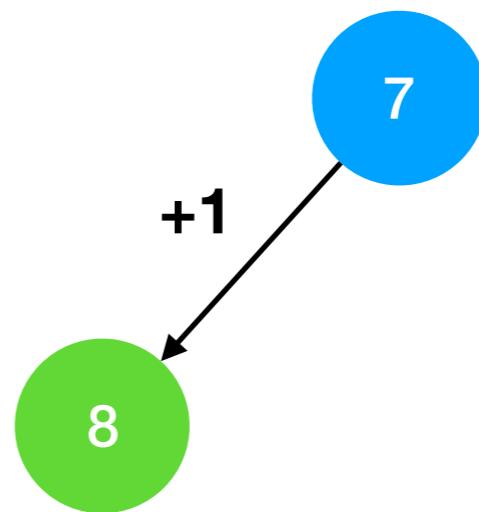


```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```

7

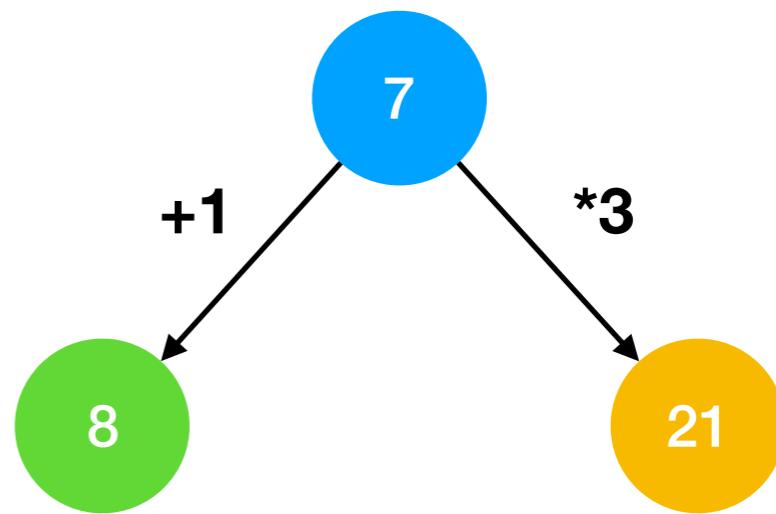
```
module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

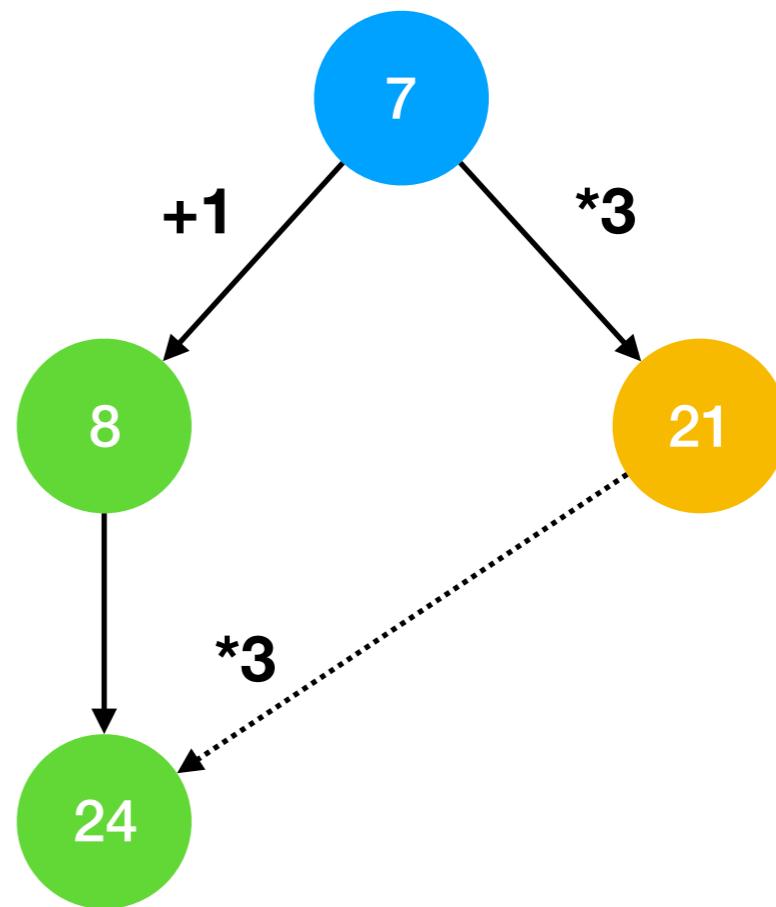
```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

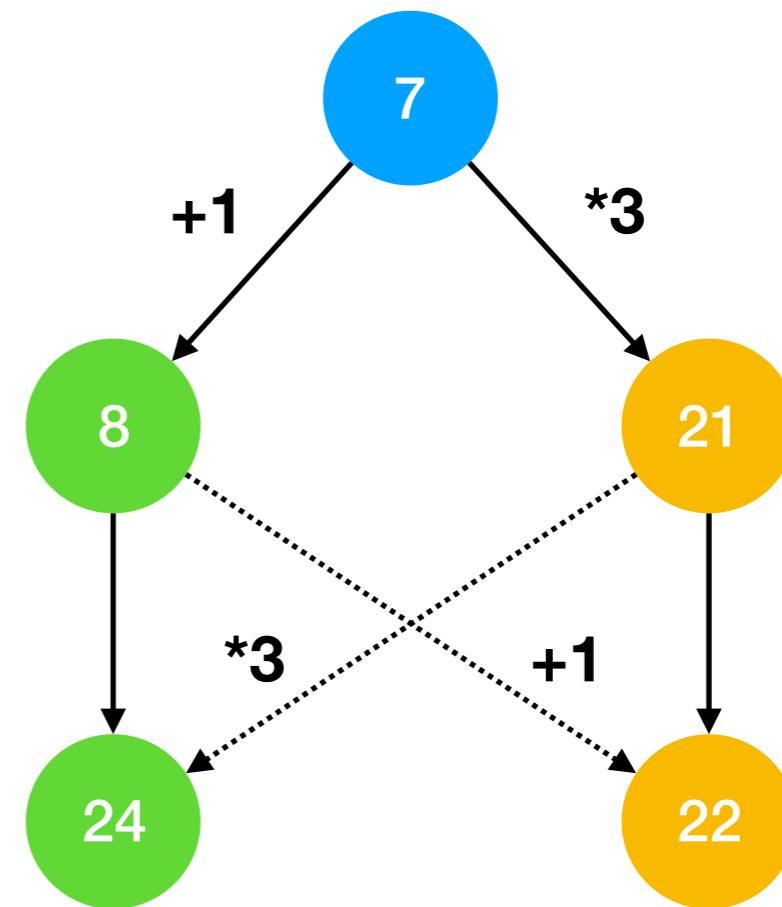
```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

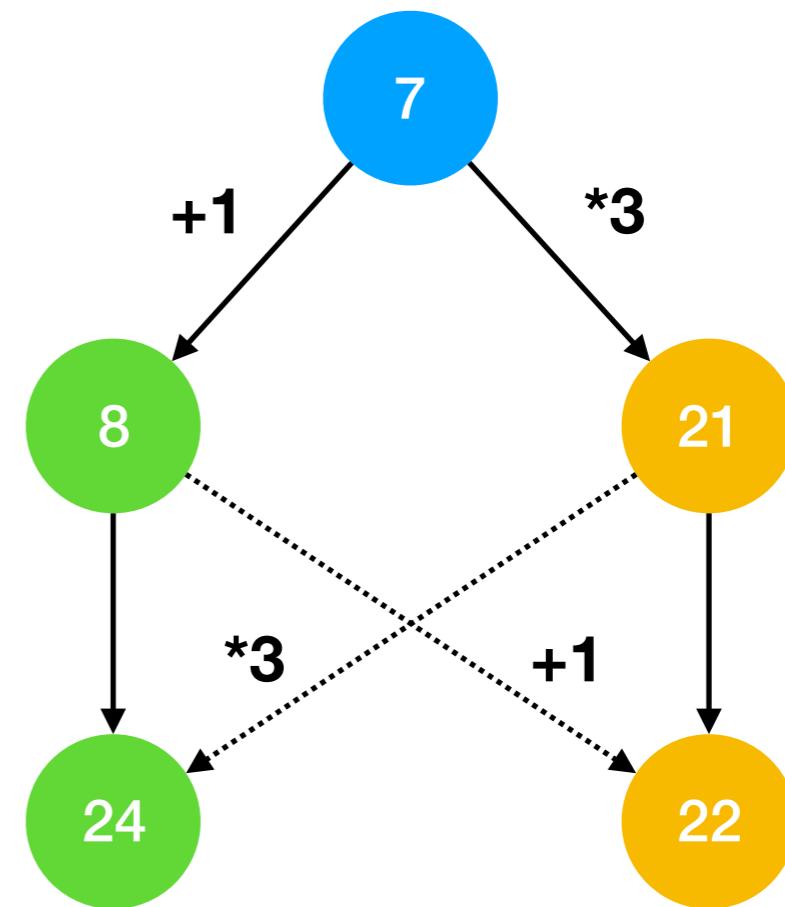
```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```

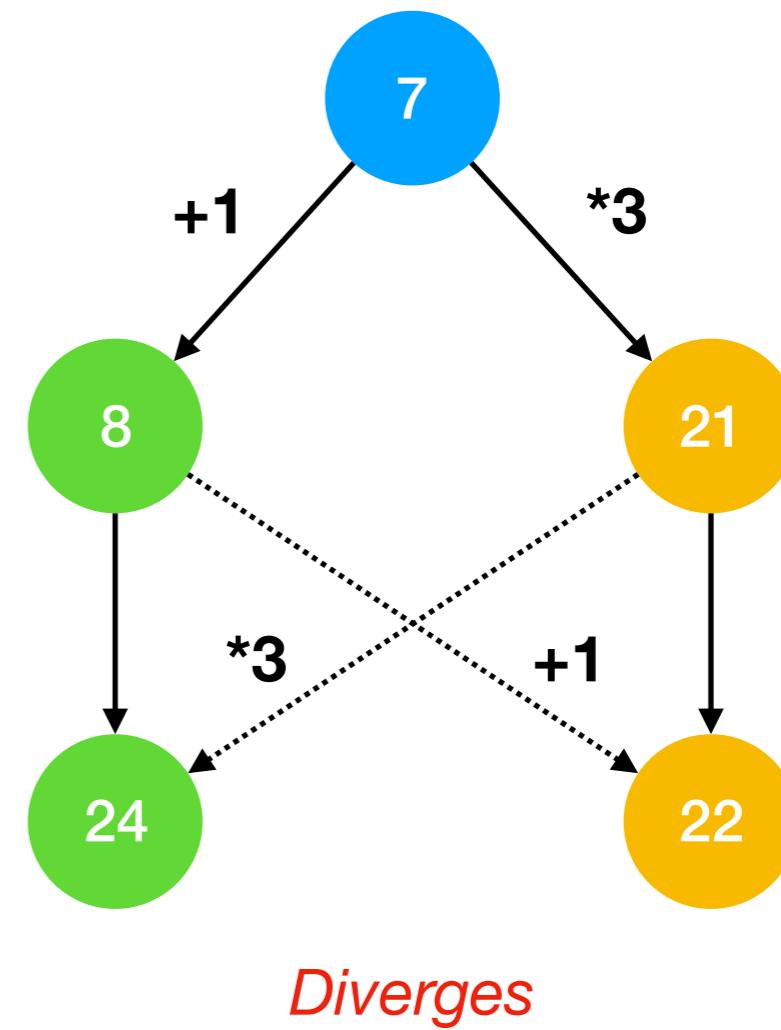


Diverges

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```



Addition and multiplication do not commute

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```

- Capture the effect of multiplication through the commutative addition operation

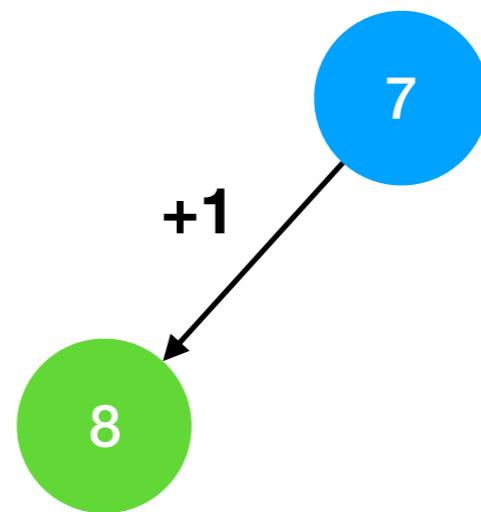
```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end
```

- Capture the effect of multiplication through the commutative addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```

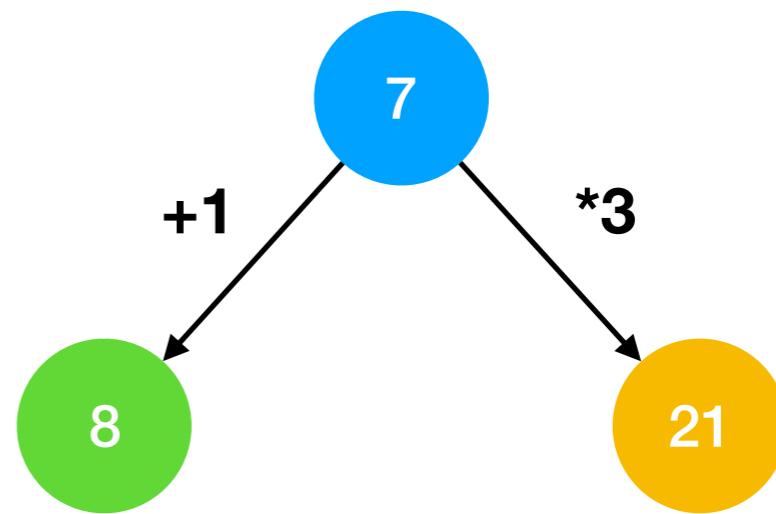


- Capture the effect of multiplication through the commutative addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```

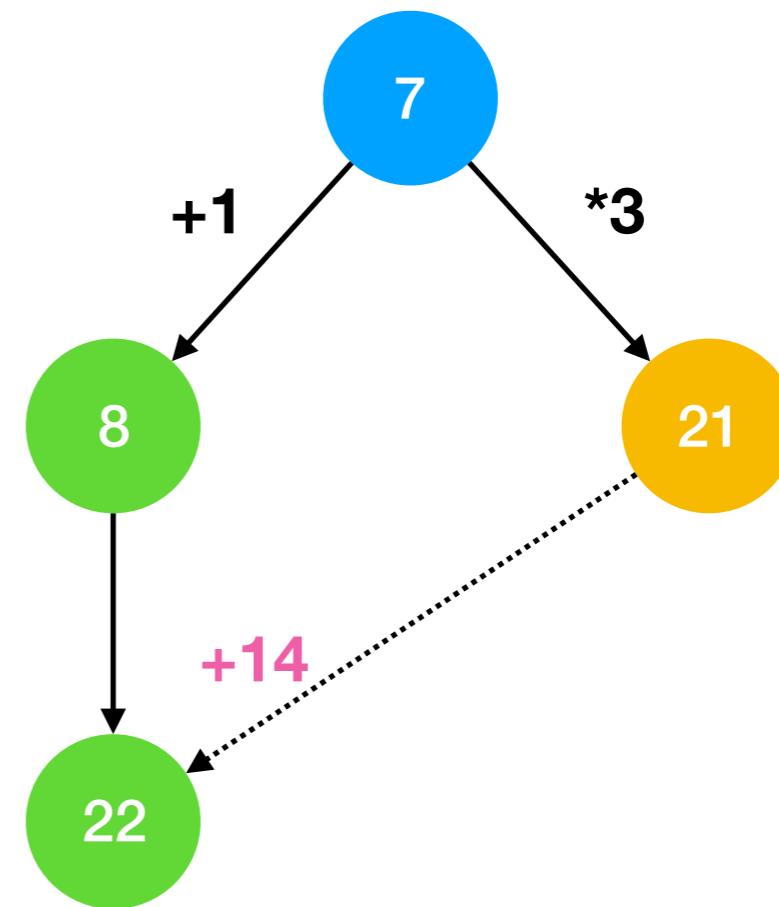


- Capture the effect of multiplication through the commutative addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```

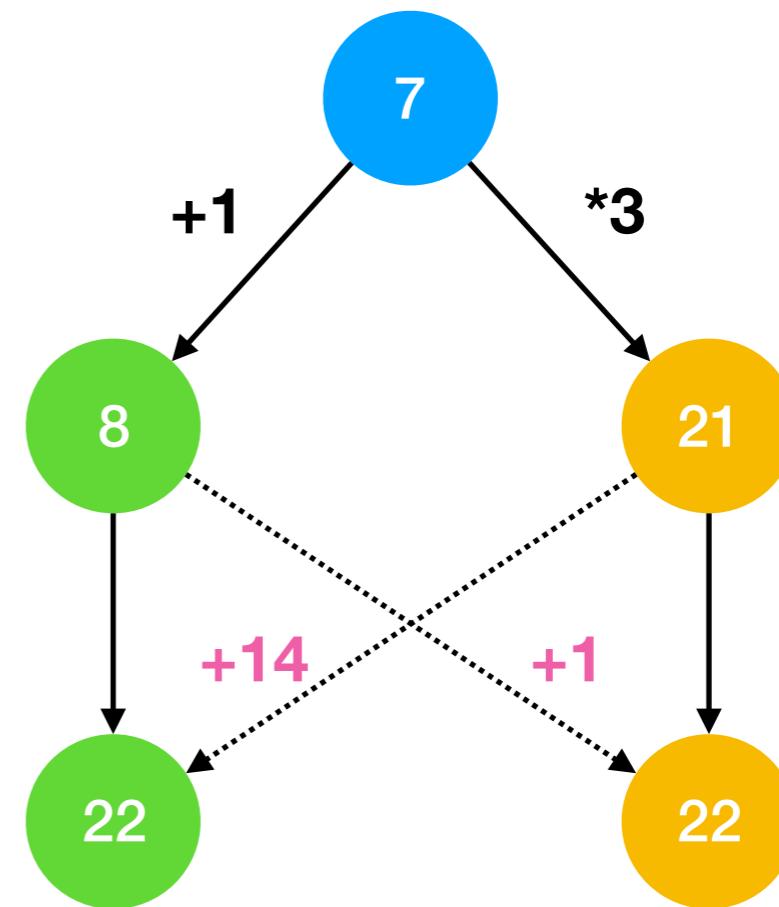


- Capture the effect of multiplication through the commutative addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```

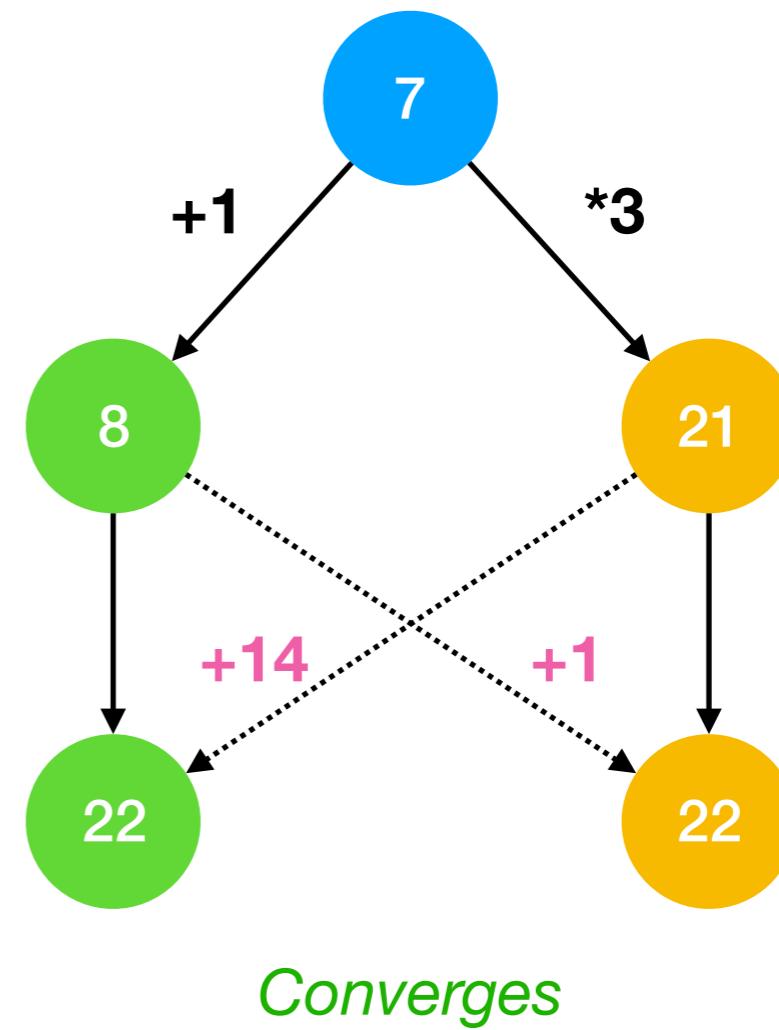


- Capture the effect of multiplication through the commutative addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```

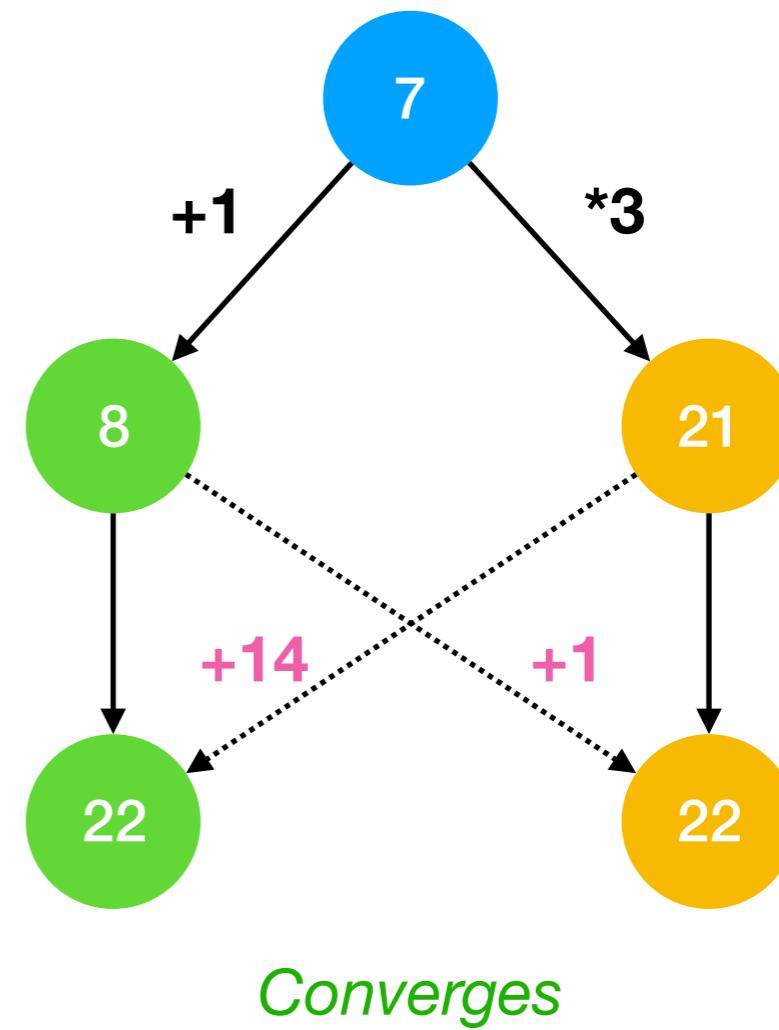


- Capture the effect of multiplication through the commutative addition operation

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
end

```



- Capture the effect of multiplication through the commutative addition operation
- CRDTs

Conflict-free Replicated Data Types (CRDT)

Conflict-free Replicated Data Types (CRDT)

- CRDT is guaranteed to ensure *strong eventual consistency (SEC)*
 - ★ G-counters, PN-counters, OR-Sets, Graphs, Ropes, docs, sheets
 - ★ Simple interface for the clients of CRDTs

Conflict-free Replicated Data Types (CRDT)

- CRDT is guaranteed to ensure *strong eventual consistency (SEC)*
 - ★ G-counters, PN-counters, OR-Sets, Graphs, Ropes, docs, sheets
 - ★ Simple interface for the clients of CRDTs
- Need to reengineer every datatype to ensure SEC (commutativity)
 - ★ Do not mirror sequential counter parts => implementation & proof burden
 - ★ Do not compose!
 - ◆ `counter set` is not a composition of `counter` and `set` CRDTs

Can we *program & reason about* replicated data types
as an extension of their sequential counterparts?

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

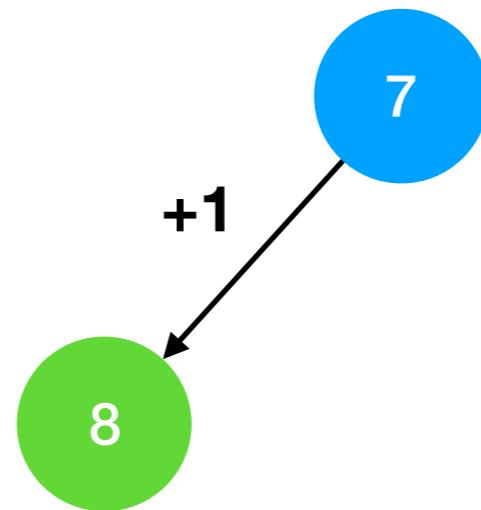
```

```
module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end
```

```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

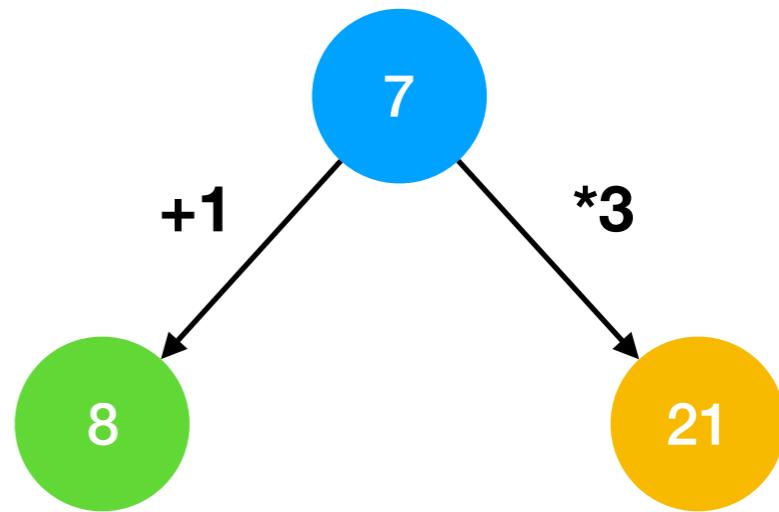
```



```

module Counter : sig
  type t
  val read : t -> int
  val add  : t -> int -> t
  val sub  : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

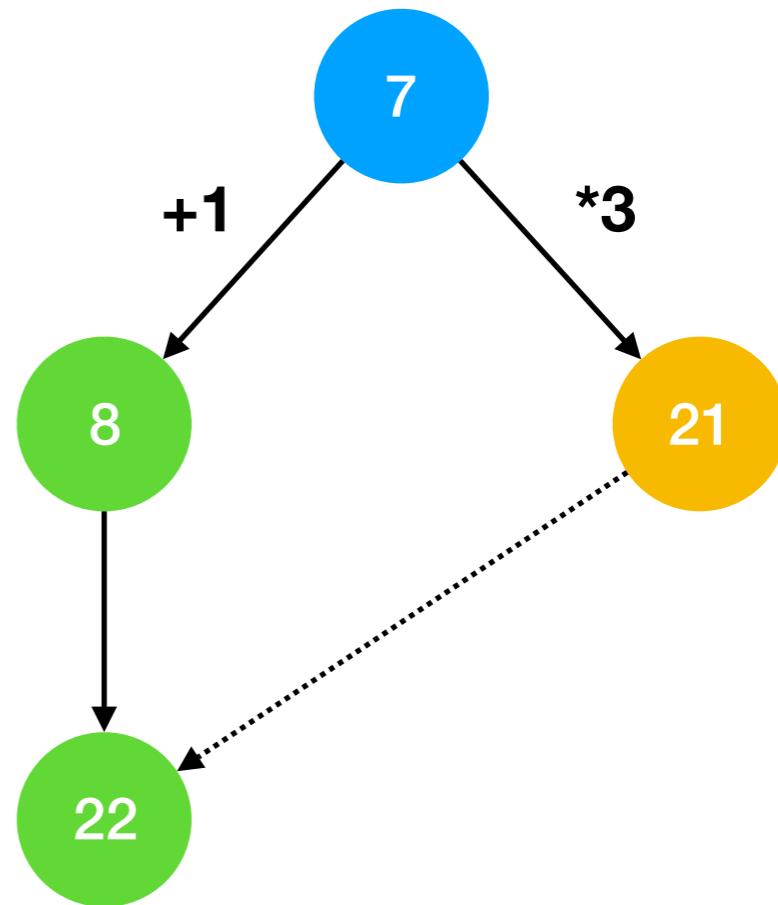
```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

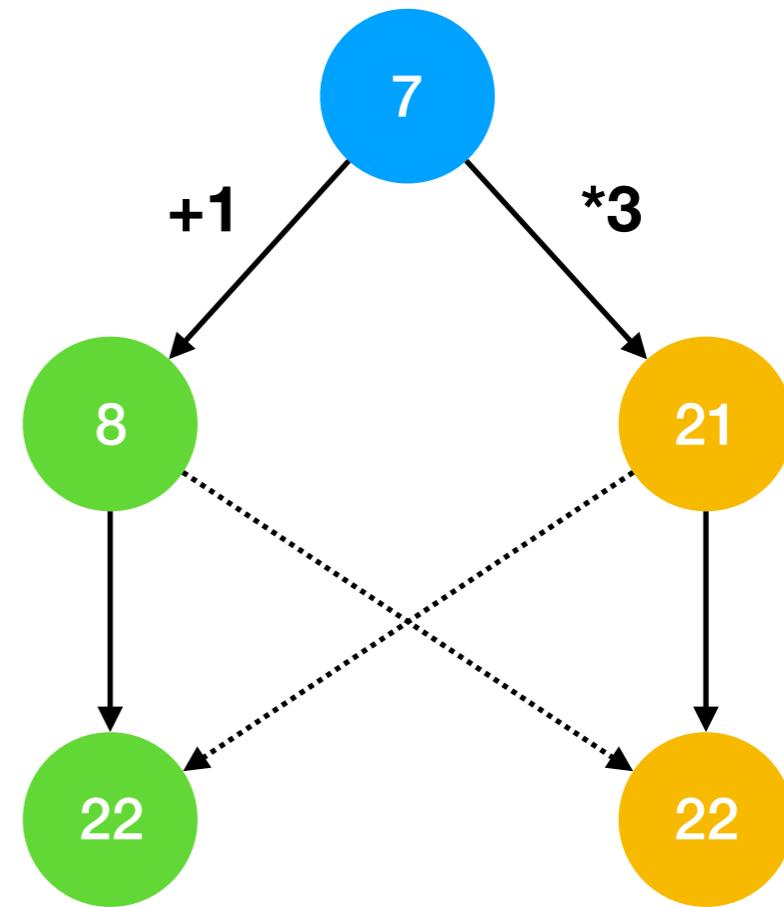
```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

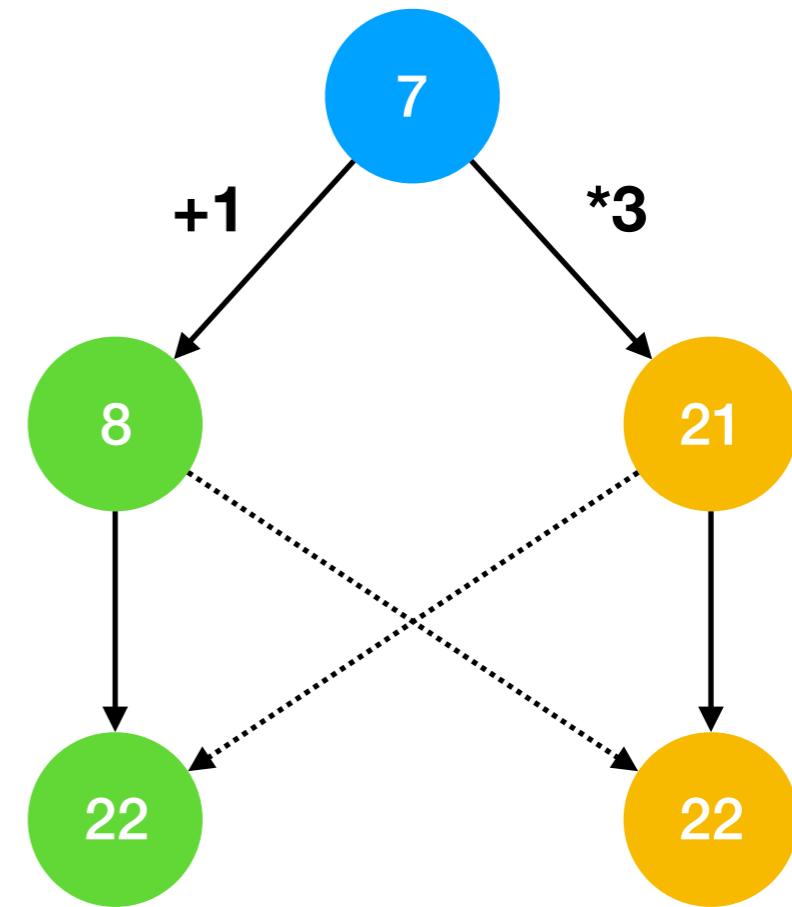
```



```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

```

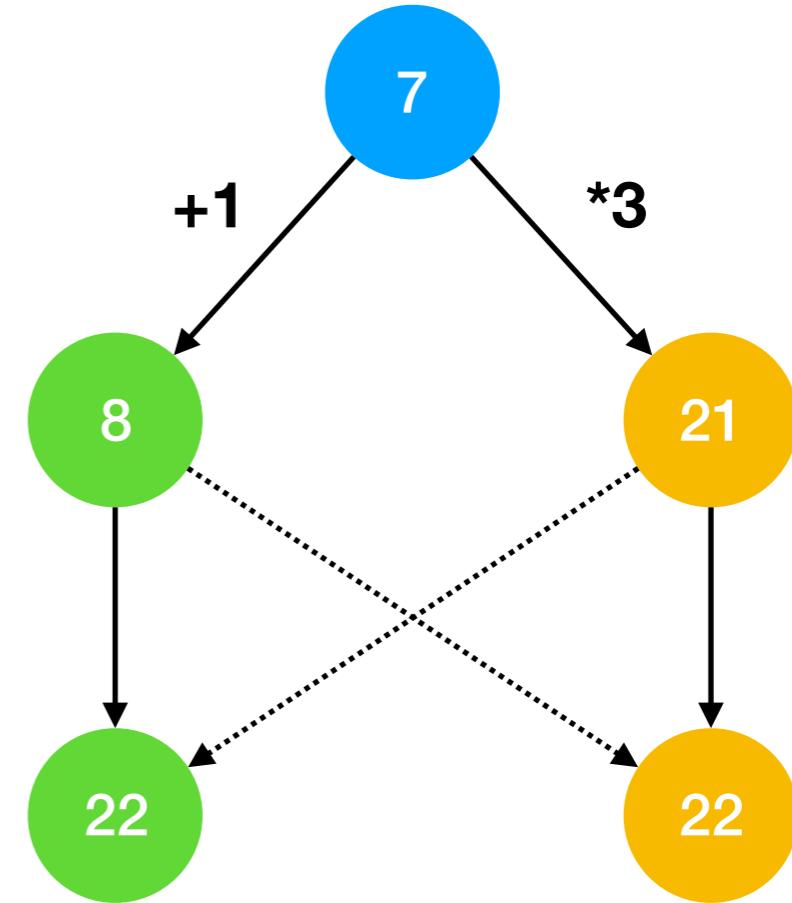


$$22 = 7 + (8-1) + (21 - 7)$$

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

```



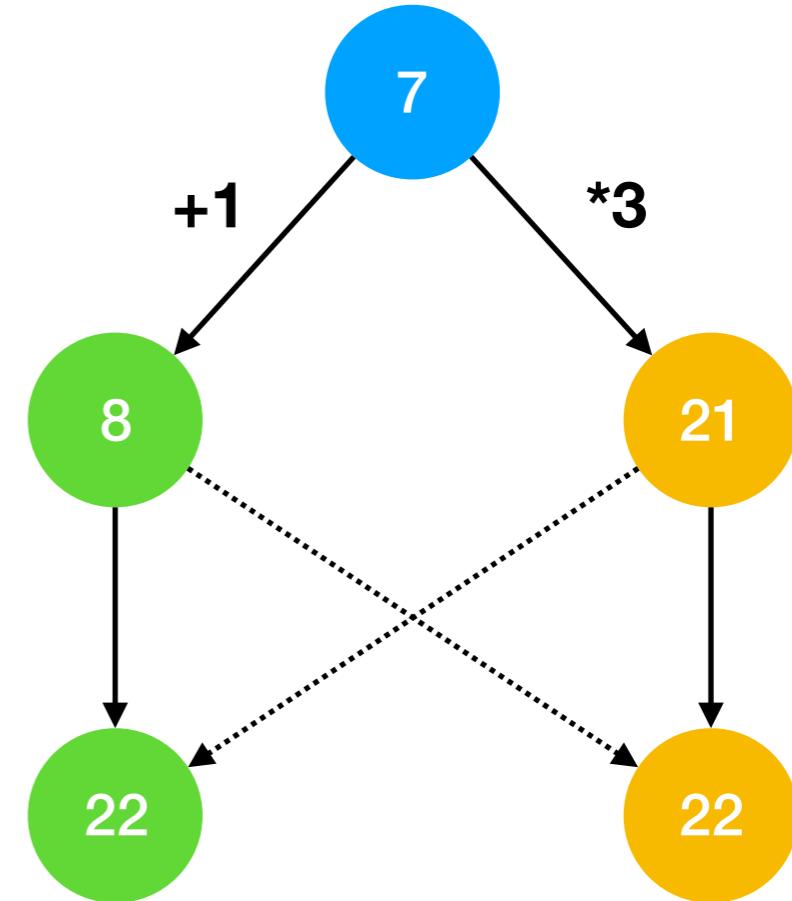
$$22 = 7 + (8-1) + (21 - 7)$$

- 3-way merge function makes the counter suitable for distribution

```

module Counter : sig
  type t
  val read : t -> int
  val add : t -> int -> t
  val sub : t -> int -> t
  val mult : t -> int -> t
  val merge : lca:t -> v1:t -> v2:t -> t
end = struct
  type t = int
  let read x = x
  let add x d = x + d
  let sub x d = x - d
  let mult x n = x * n
  let merge ~lca ~v1 ~v2 =
    lca + (v1 - lca) + (v2 - lca)
end

```



$$22 = 7 + (8-1) + (21 - 7)$$

- 3-way merge function makes the counter suitable for distribution
- Does not appeal to individual operations => independently extend data-type

Systems → PL

Systems → PL

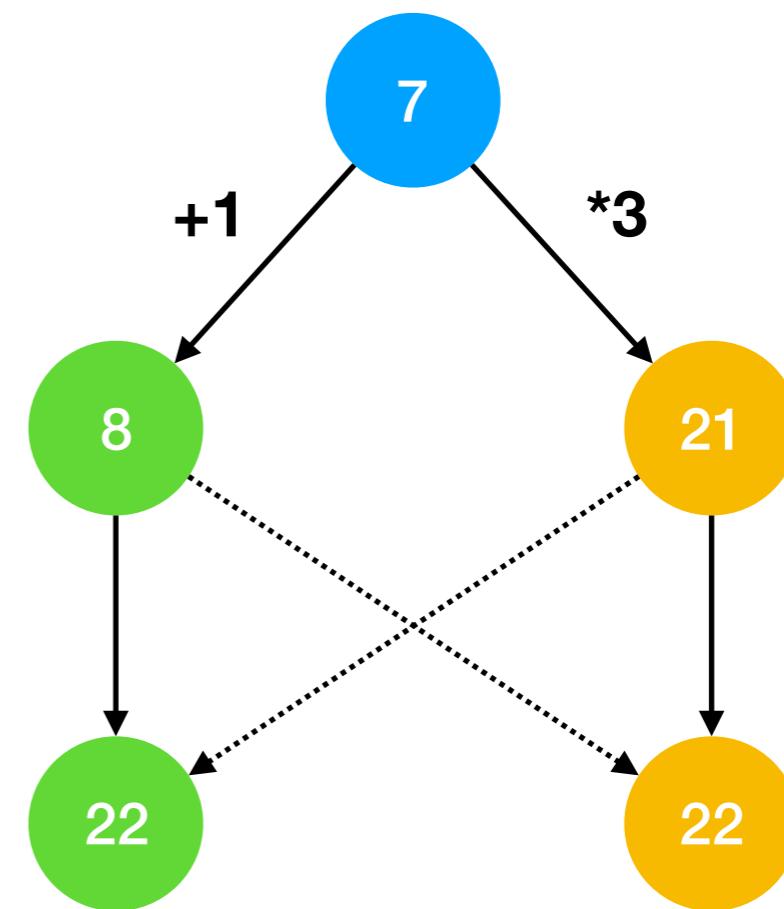
- CRDTs need to take care of systems level concerns such as exactly once delivery

Systems → PL

- CRDTs need to take care of systems level concerns such as exactly once delivery
- 3-way merge handles it automatically

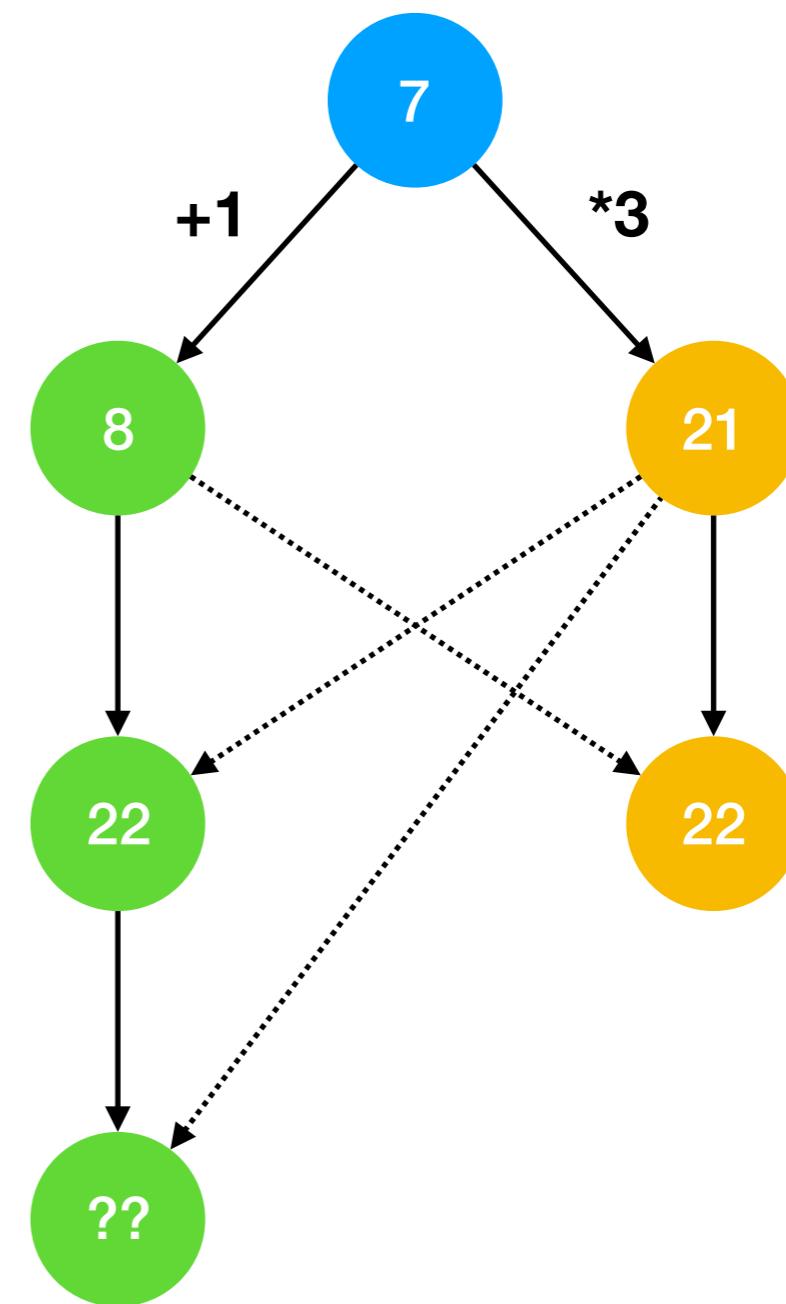
Systems → PL

- CRDTs need to take care of systems level concerns such as exactly once delivery
- 3-way merge handles it automatically



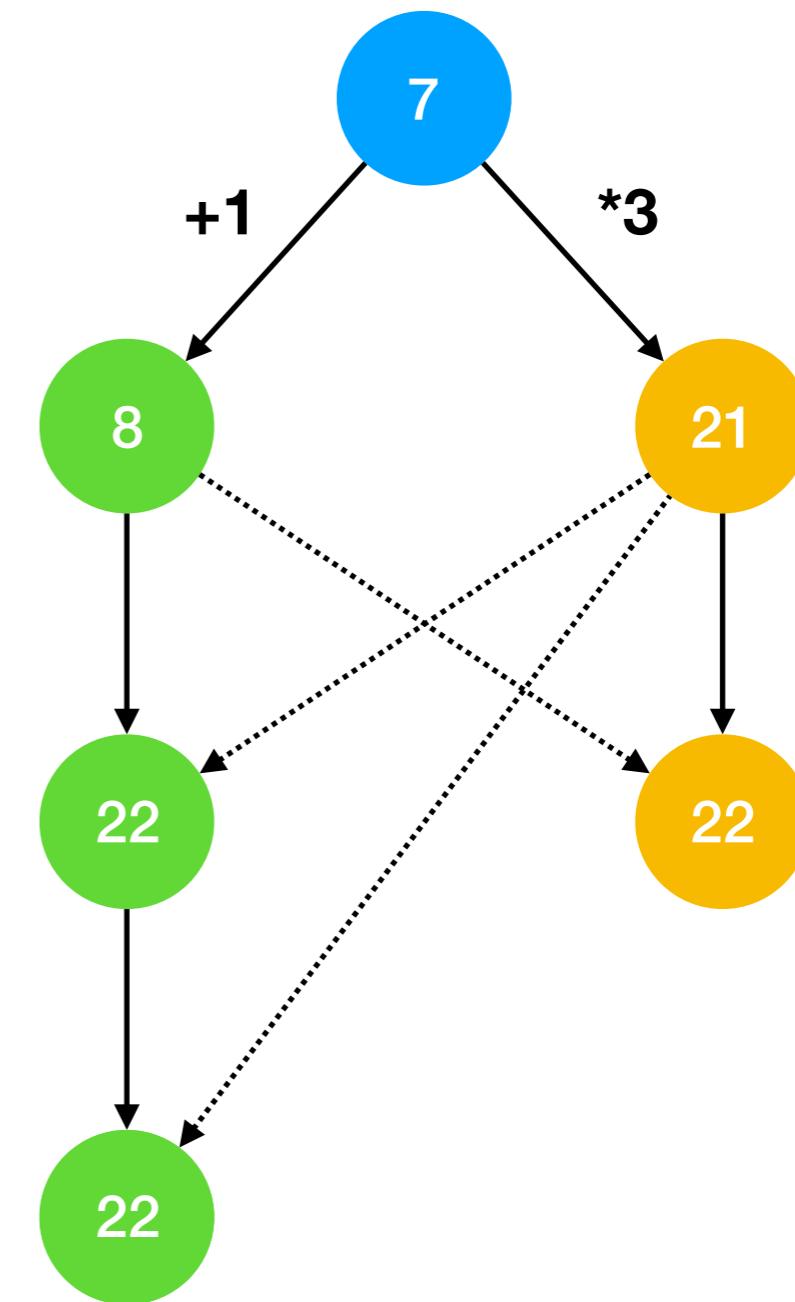
Systems → PL

- CRDTs need to take care of systems level concerns such as exactly once delivery
- 3-way merge handles it automatically



Systems → PL

- CRDTs need to take care of systems level concerns such as exactly once delivery
- 3-way merge handles it automatically



$$22 = 21 + (21-21) + (22 - 21)$$

Does the 3-way merge idea generalise?

```
module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end
```

```
module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end
```

- Try replicating queues by asynchronously transmitting operations

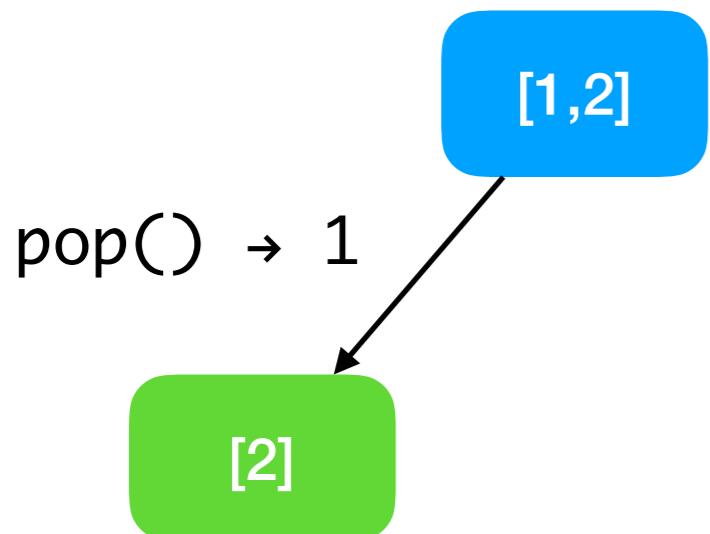
```
module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end
```

- Try replicating queues by asynchronously transmitting operations

[1,2]

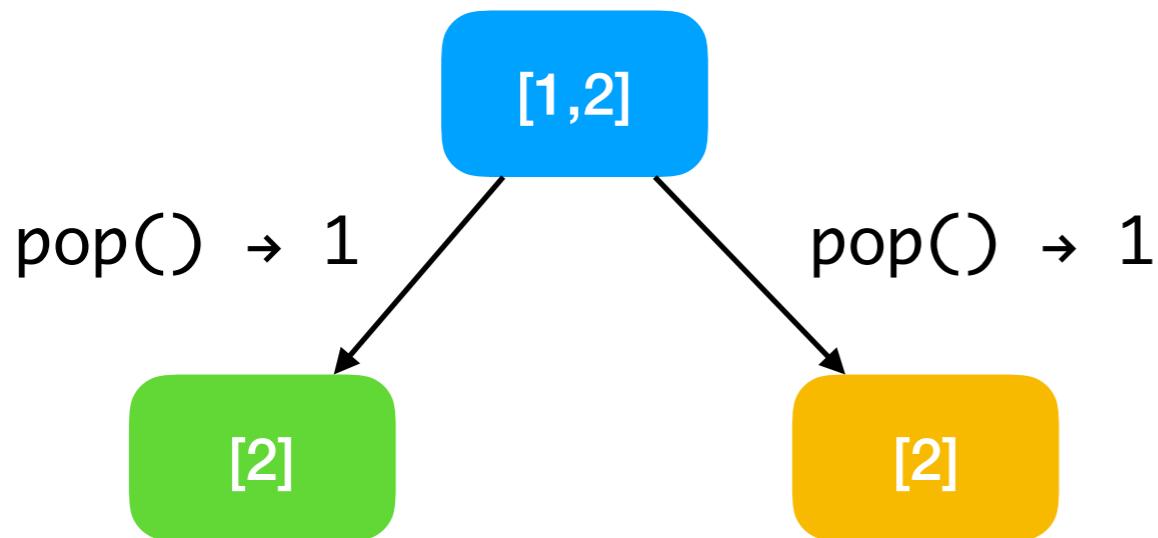
```
module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end
```

- Try replicating queues by asynchronously transmitting operations



```
module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end
```

- Try replicating queues by asynchronously transmitting operations

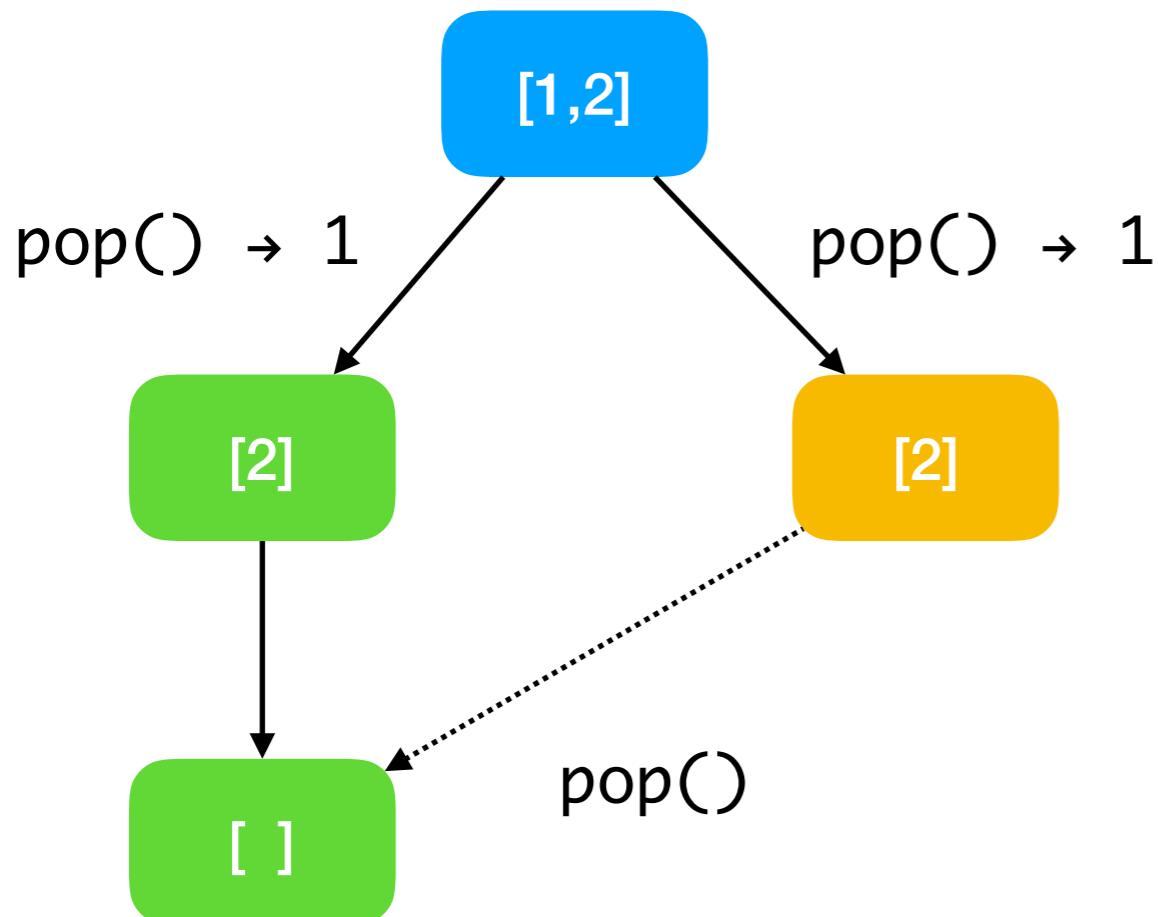


```

module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end

```

- Try replicating queues by asynchronously transmitting operations

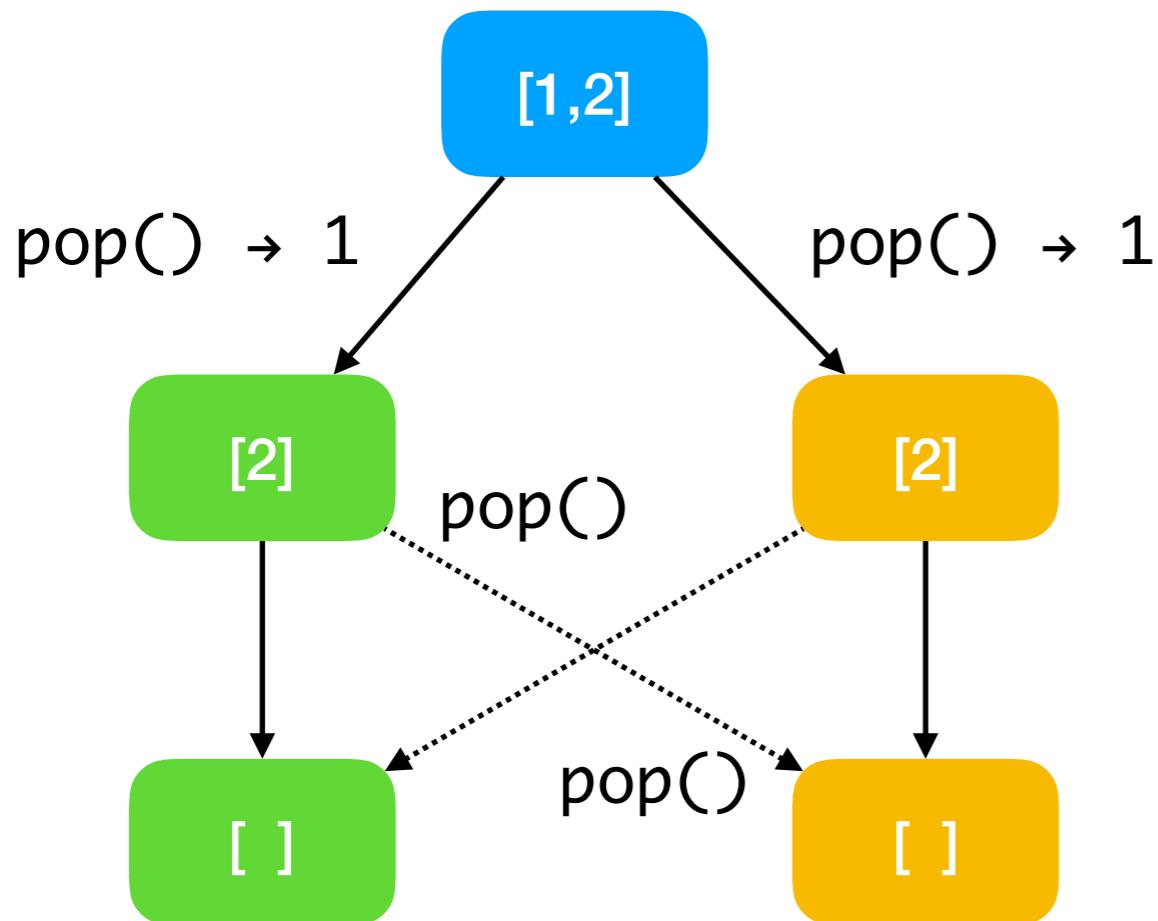


```

module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end

```

- Try replicating queues by asynchronously transmitting operations

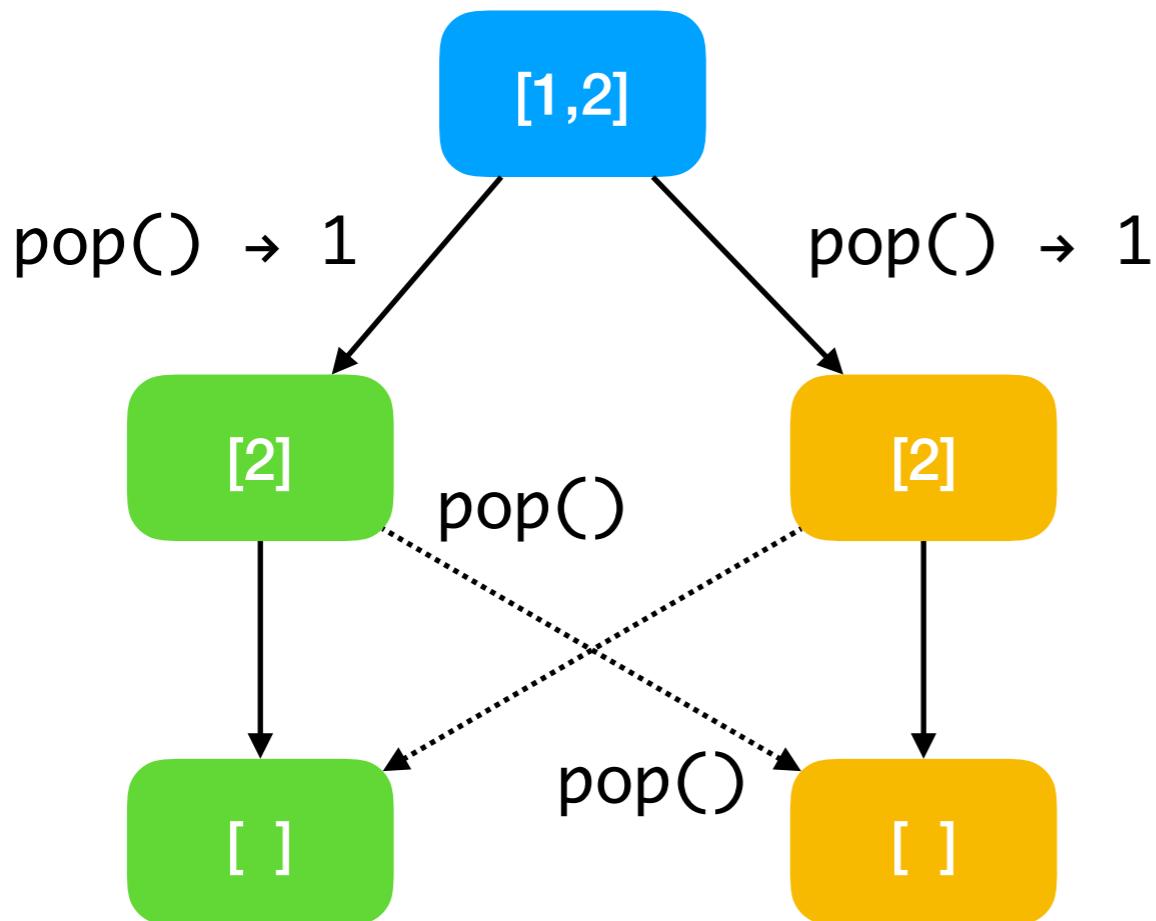


```

module type Queue = sig
    type 'a t
    val push : 'a t -> 'a -> 'a t
    val pop : 'a t -> ('a * 'a t) option
    (* at-least once semantics *)
end

```

- Try replicating queues by asynchronously transmitting operations



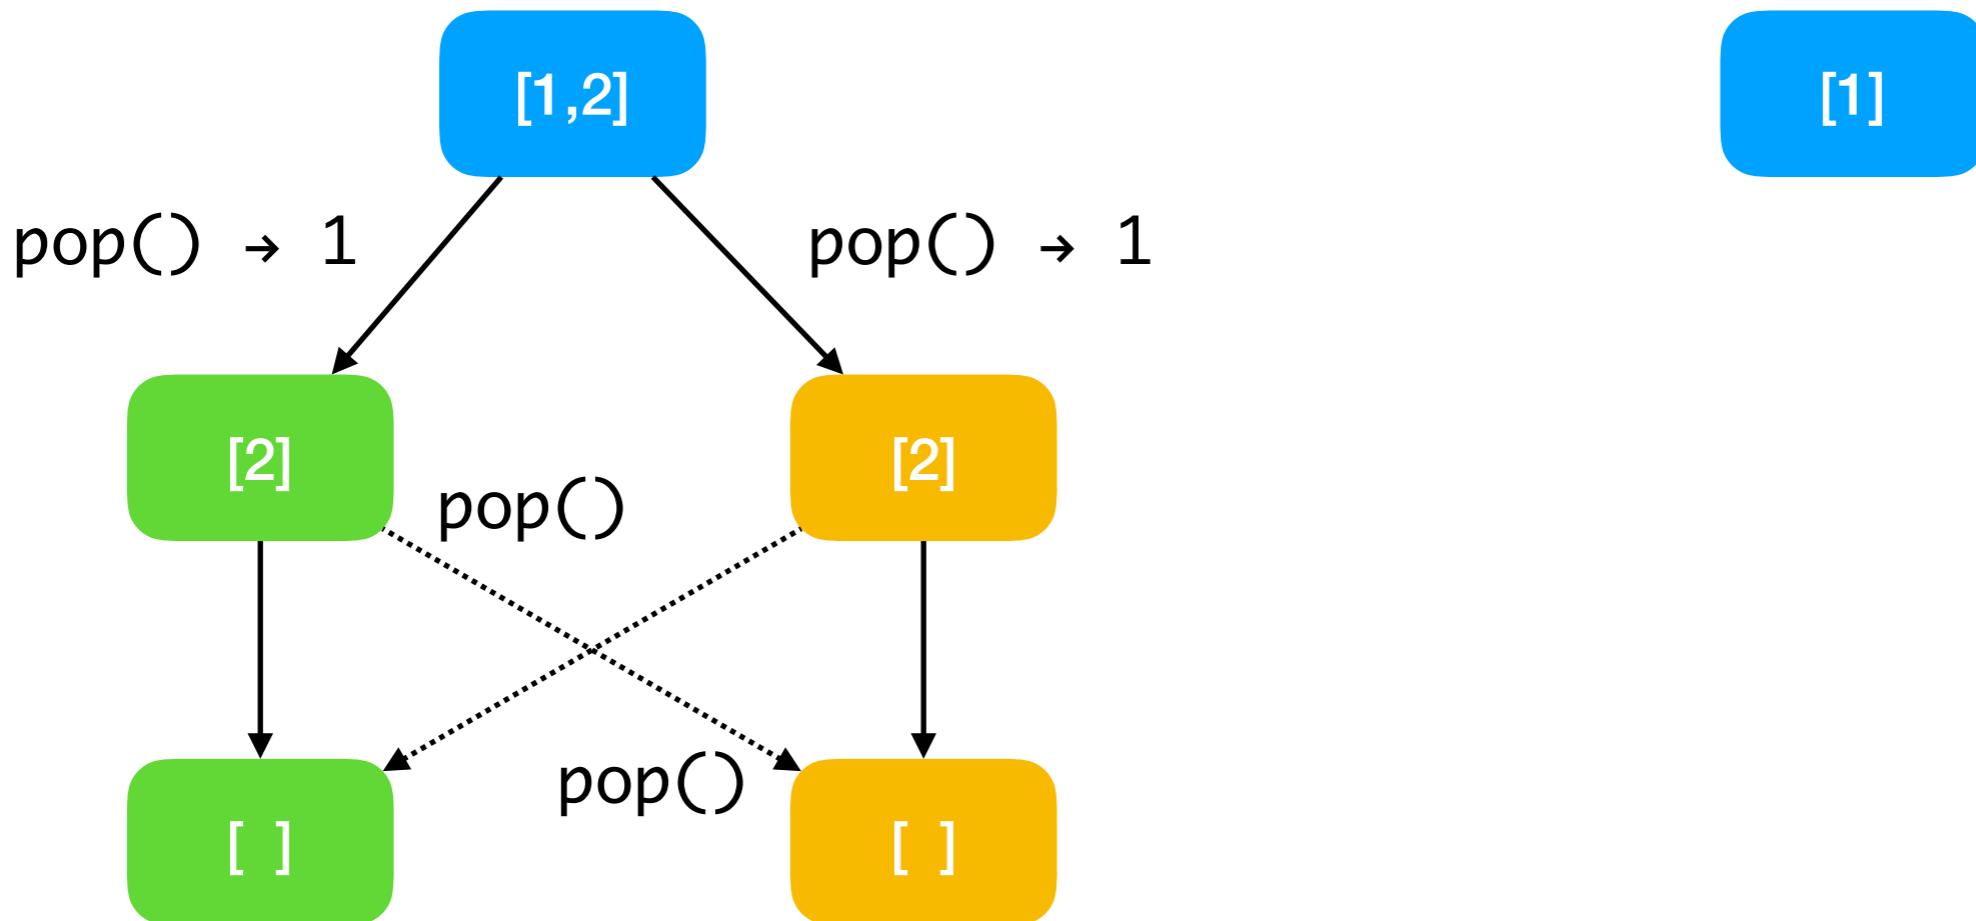
- Convergence is not sufficient; *Intent* is not preserved

```

module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end

```

- Try replicating queues by asynchronously transmitting operations



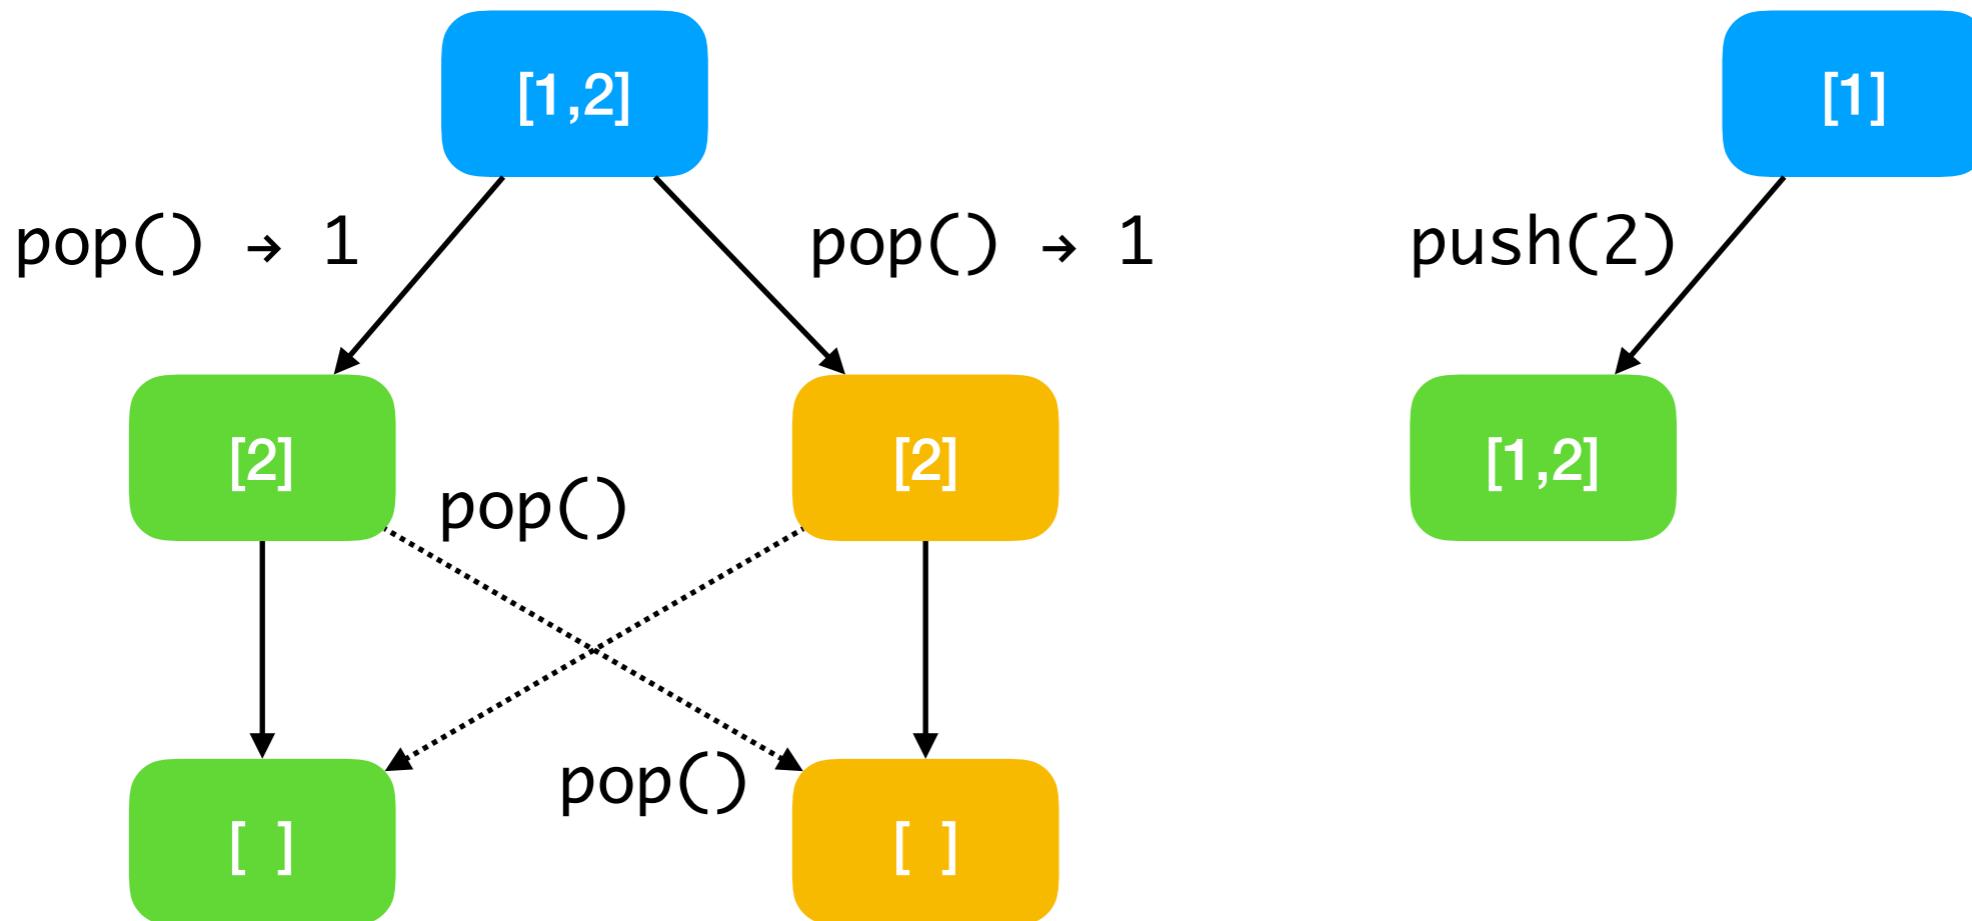
- Convergence is not sufficient; *Intent* is not preserved

```

module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end

```

- Try replicating queues by asynchronously transmitting operations



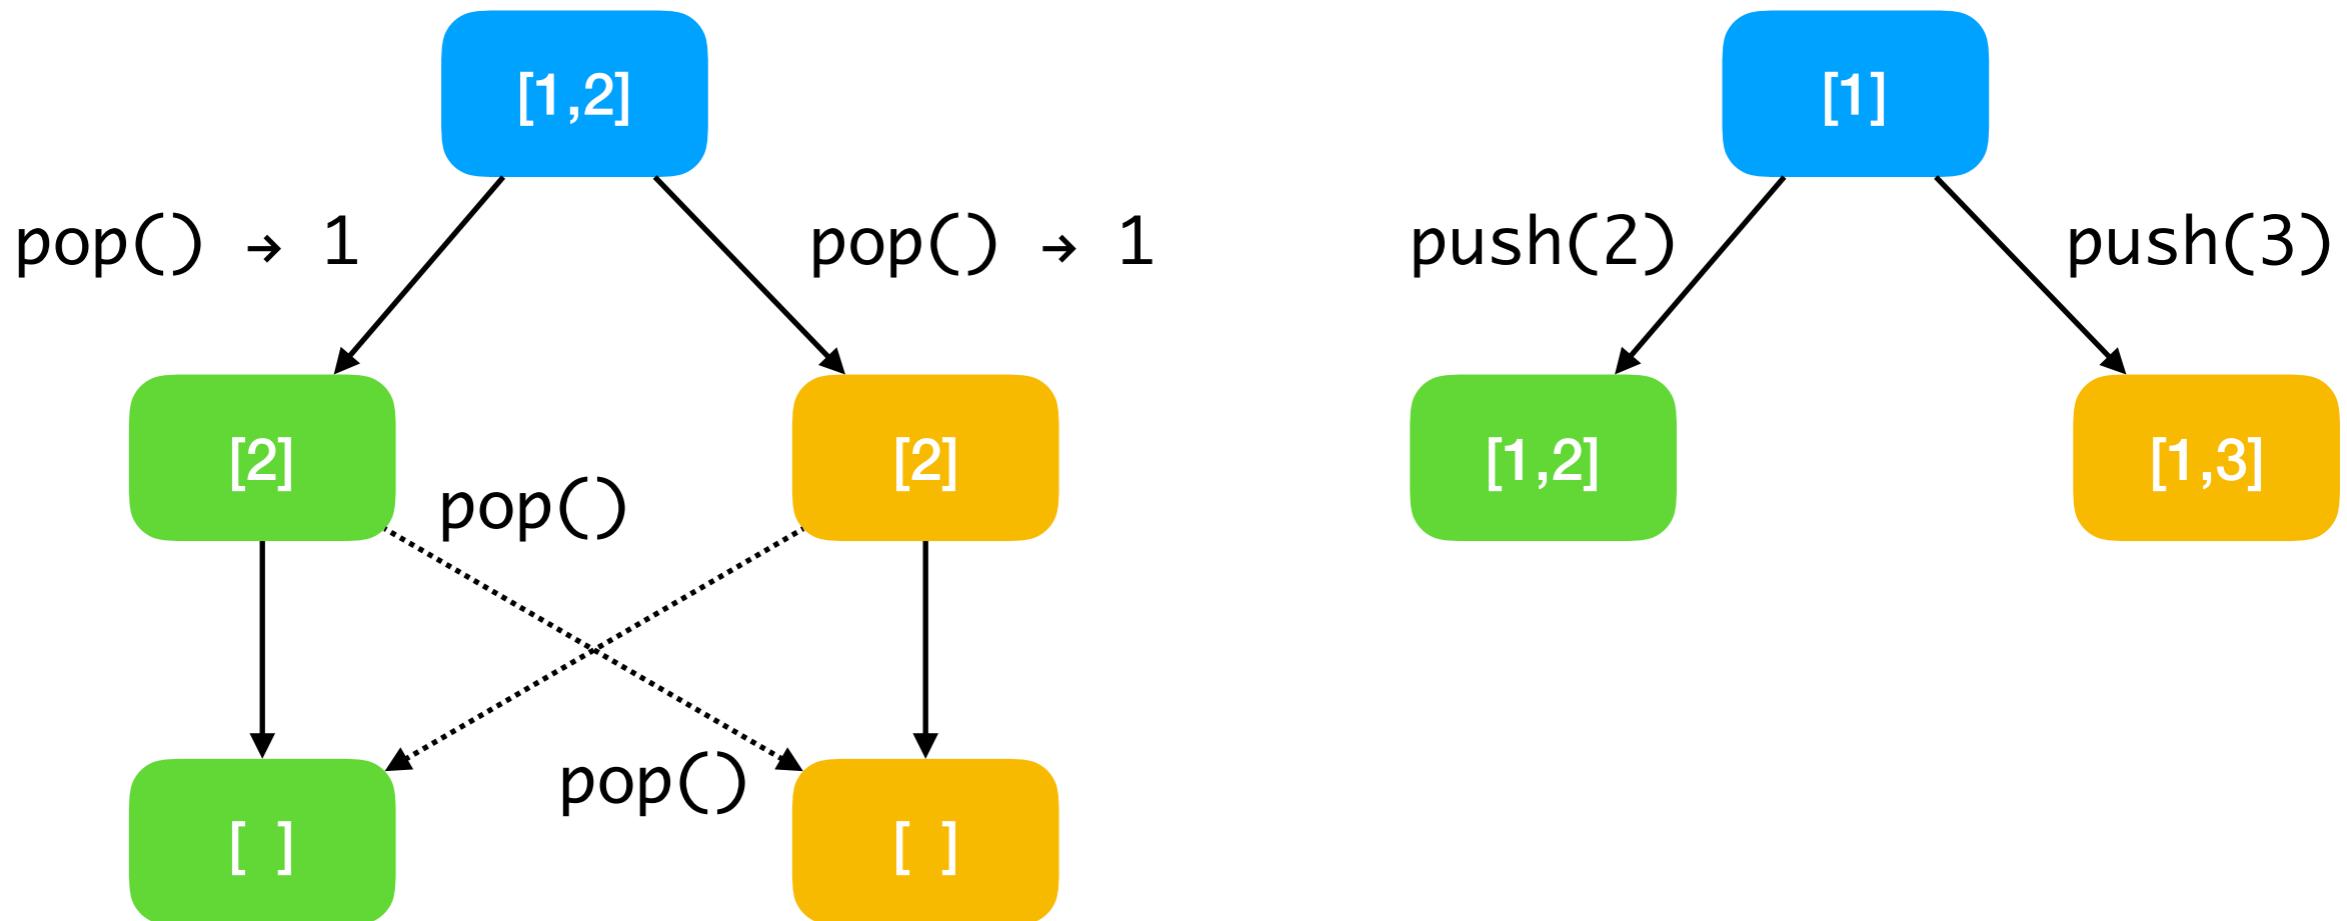
- Convergence is not sufficient; *Intent* is not preserved

```

module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end

```

- Try replicating queues by asynchronously transmitting operations



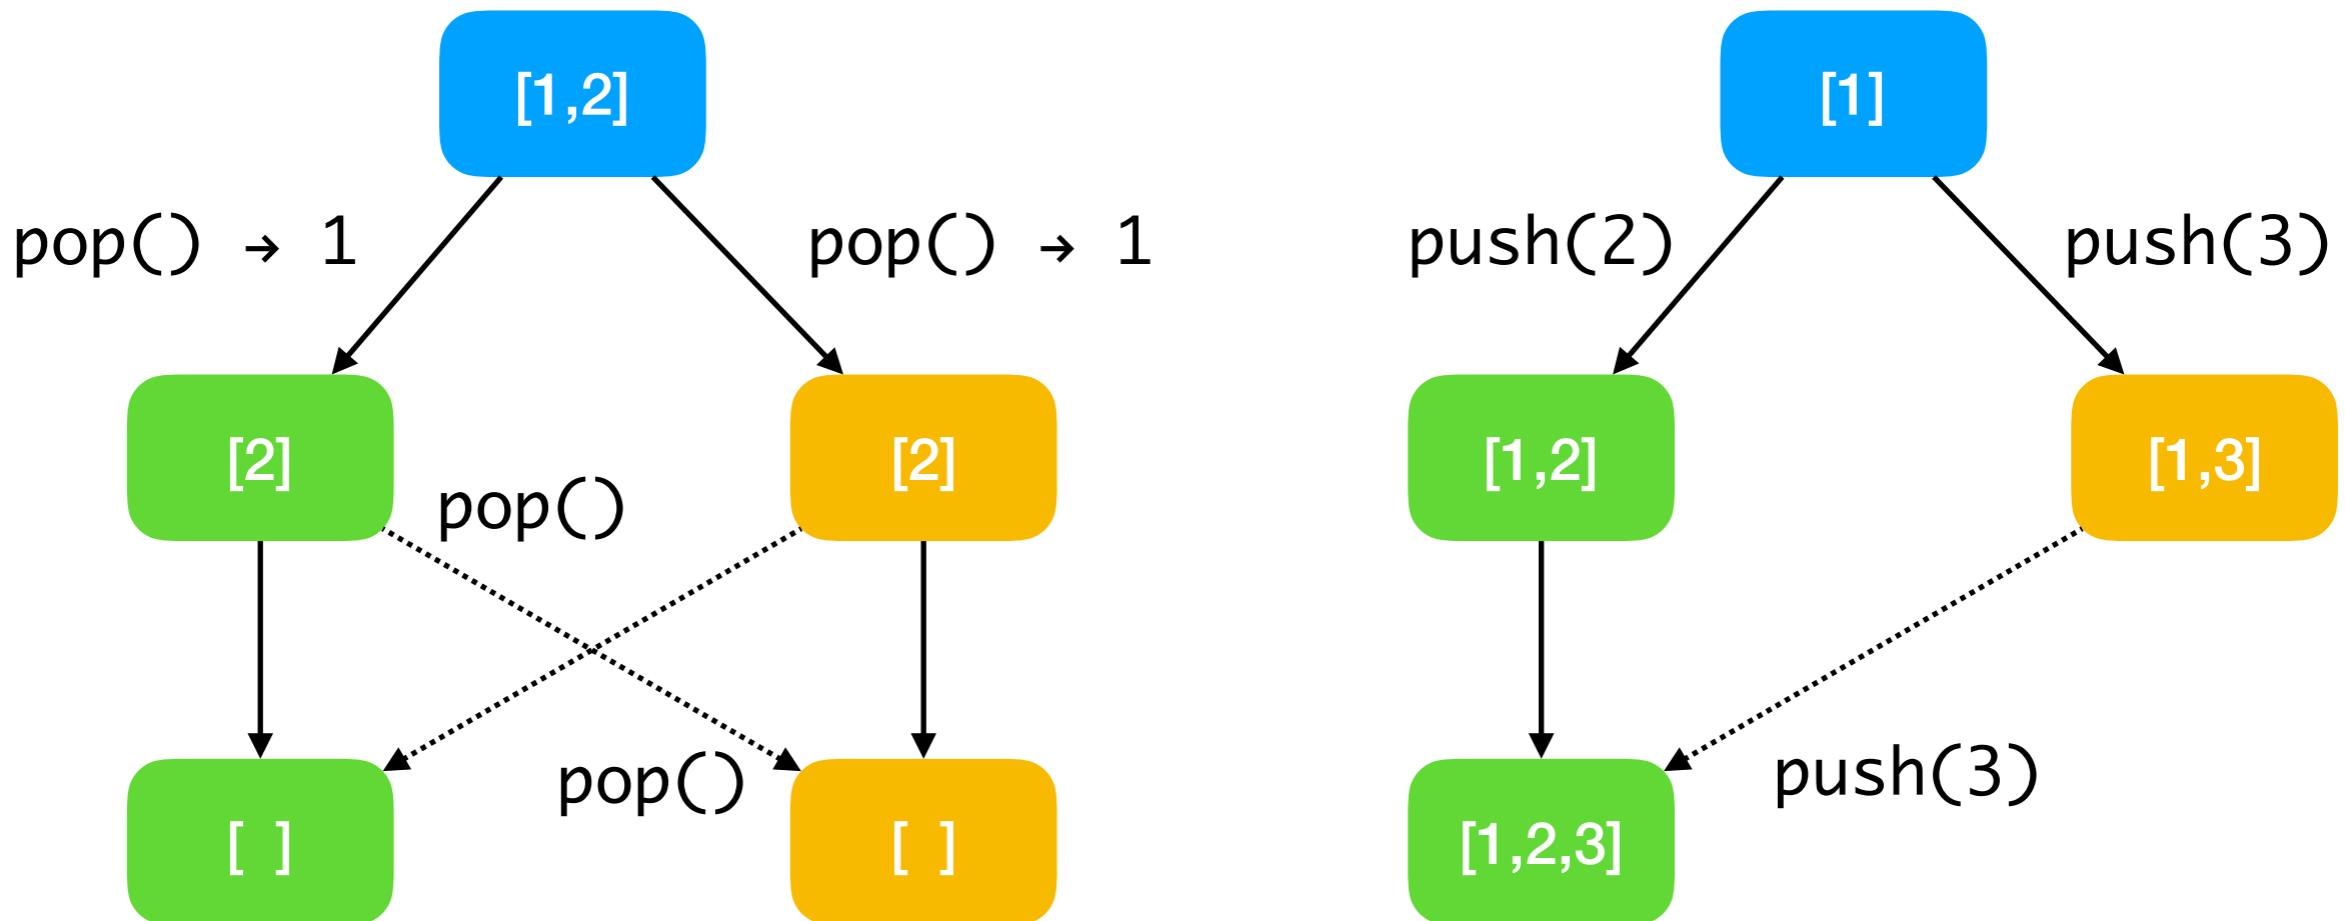
- Convergence is not sufficient; *Intent* is not preserved

```

module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end

```

- Try replicating queues by asynchronously transmitting operations



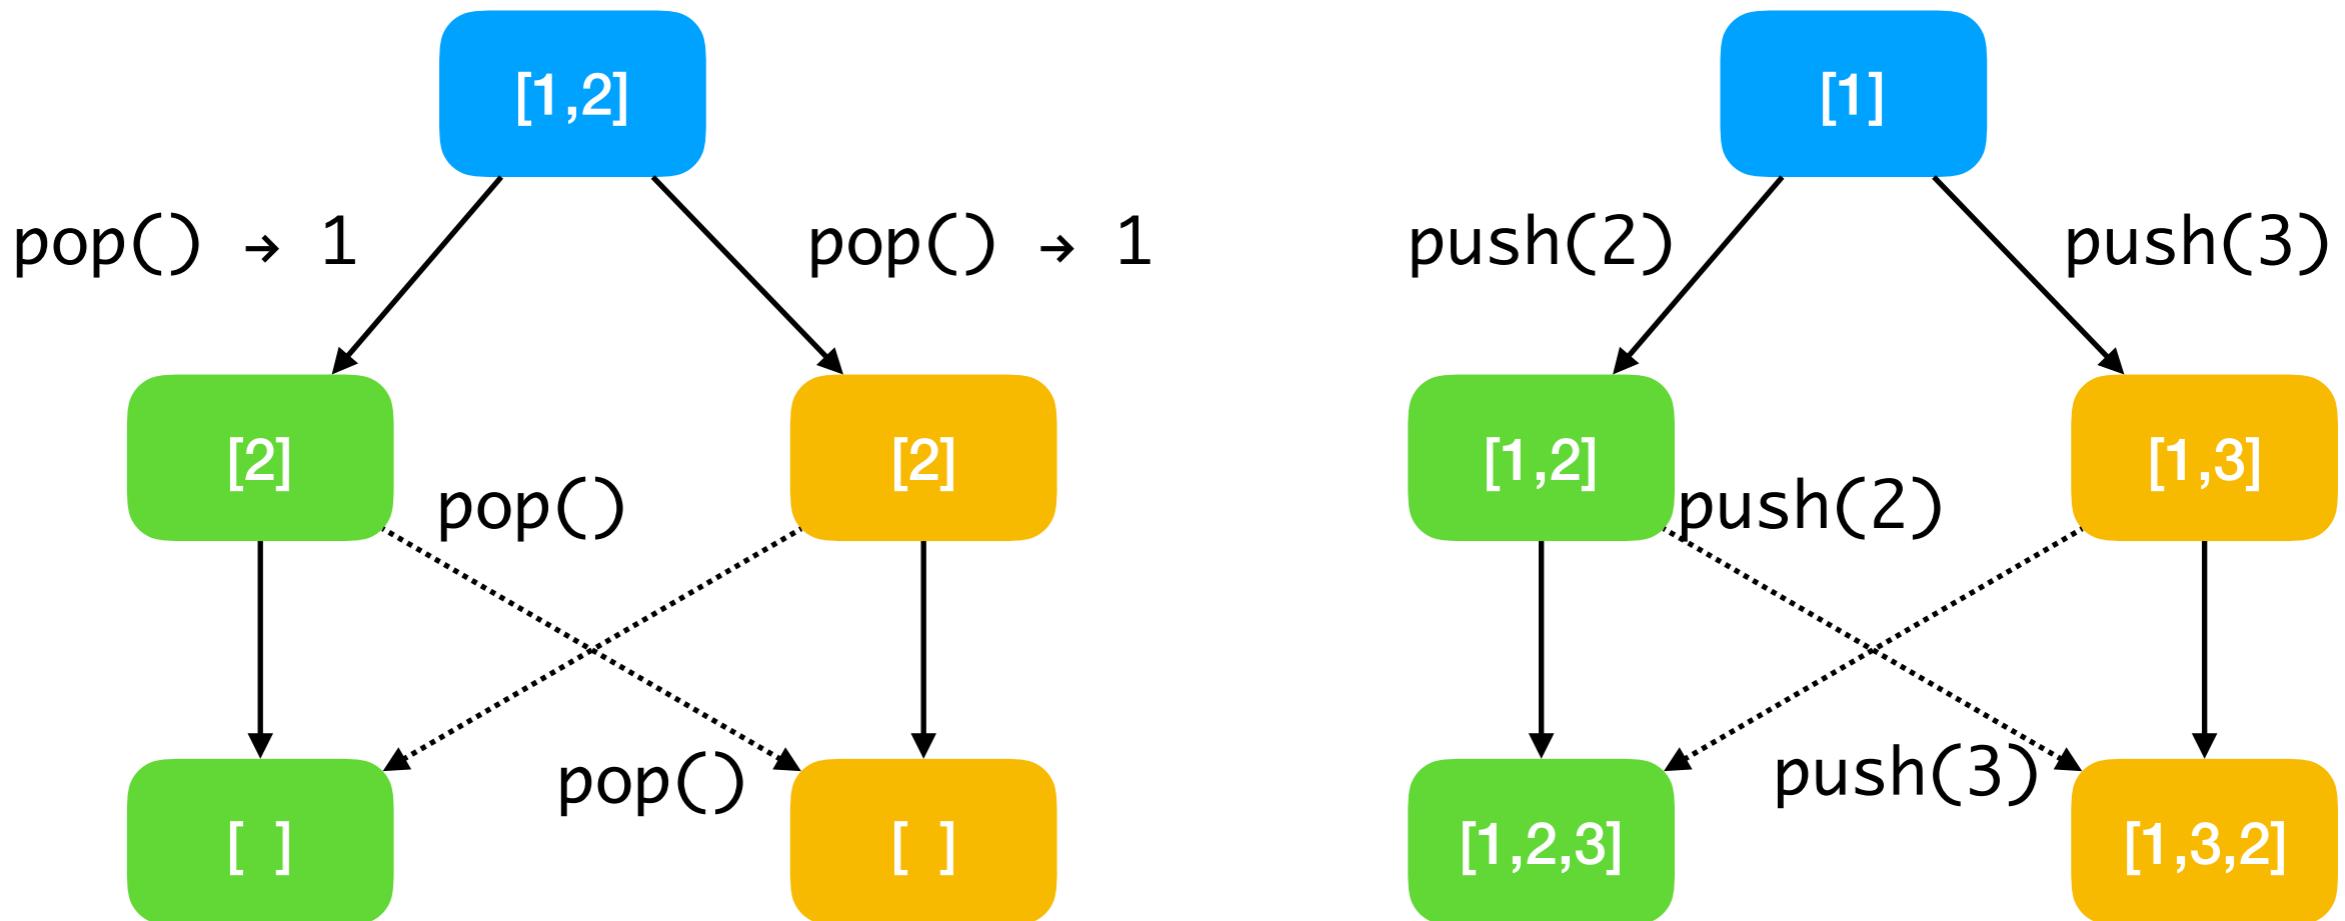
- Convergence is not sufficient; *Intent* is not preserved

```

module type Queue = sig
  type 'a t
  val push : 'a t -> 'a -> 'a t
  val pop : 'a t -> ('a * 'a t) option
  (* at-least once semantics *)
end

```

- Try replicating queues by asynchronously transmitting operations



- Convergence is not sufficient; *Intent* is not preserved

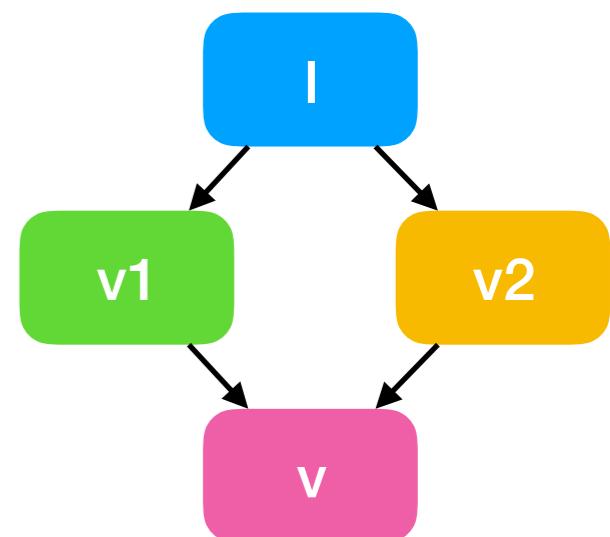
Concretising Intent

Concretising Intent

- Intent is a woolly term
 - ★ *How can we formalise the intent of operations on a data structure?*

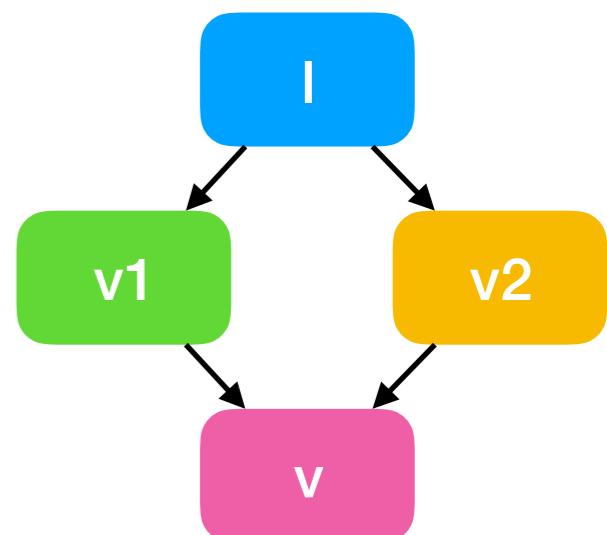
Concretising Intent

- Intent is a woolly term
 - ★ *How can we formalise the intent of operations on a data structure?*



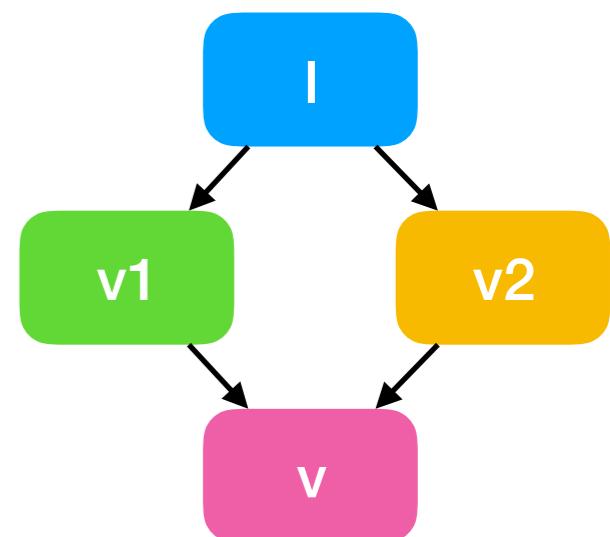
Concretising Intent

- Intent is a woolly term
 - ★ *How can we formalise the intent of operations on a data structure?*
- For a replicated queue,



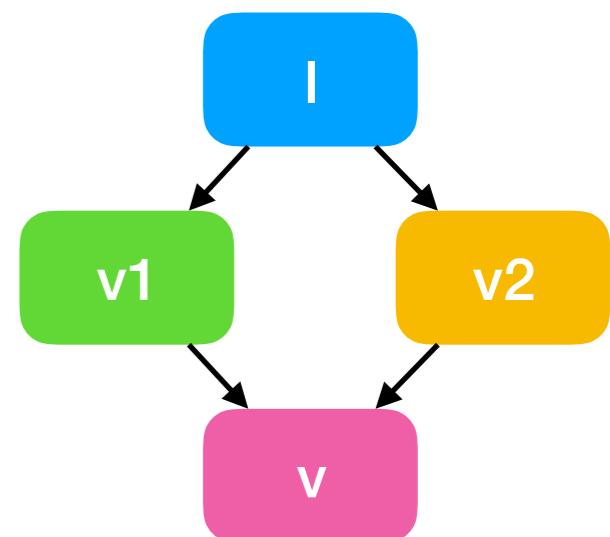
Concretising Intent

- Intent is a woolly term
 - ★ *How can we formalise the intent of operations on a data structure?*
- For a replicated queue,
 - I. Any element popped in either v1 or v2 does not remain in v



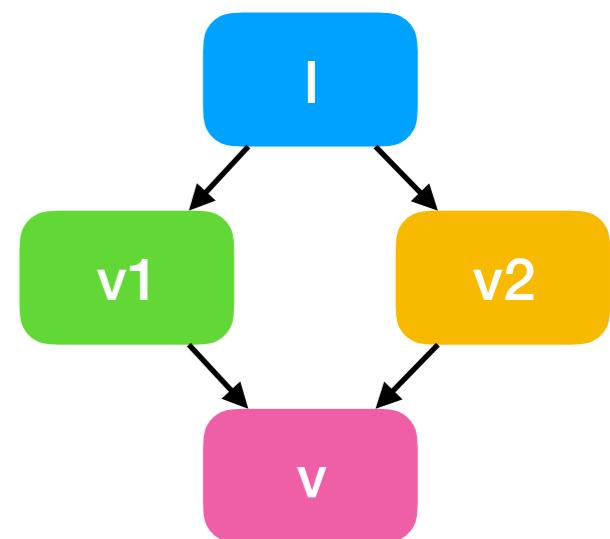
Concretising Intent

- Intent is a woolly term
 - ★ *How can we formalise the intent of operations on a data structure?*
- For a replicated queue,
 1. Any element popped in either v_1 or v_2 does not remain in v
 2. Any element pushed into either v_1 or v_2 appears in v



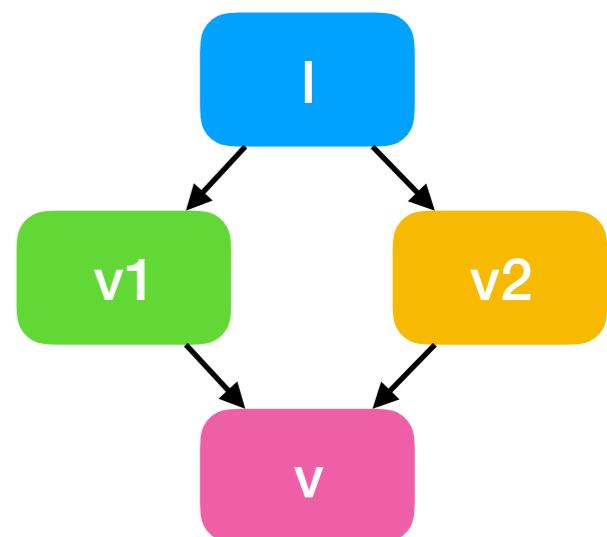
Concretising Intent

- Intent is a woolly term
 - ★ *How can we formalise the intent of operations on a data structure?*
- For a replicated queue,
 - I. Any element popped in either v_1 or v_2 does not remain in v
 2. Any element pushed into either v_1 or v_2 appears in v
 3. An element that remains untouched in I, v_1, v_2 remains in v



Concretising Intent

- Intent is a woolly term
 - ★ *How can we formalise the intent of operations on a data structure?*
- For a replicated queue,
 1. Any element popped in either v_1 or v_2 does not remain in v
 2. Any element pushed into either v_1 or v_2 appears in v
 3. An element that remains untouched in I, v_1, v_2 remains in v
 4. Order of pairs of elements in I, v_1, v_2 must be preserved in m , if those elements are present in v .



Relational Specification

Relational Specification

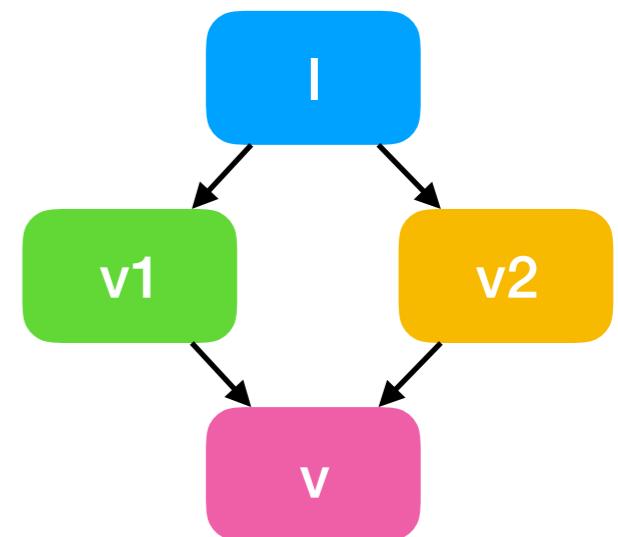
- Let's define relations R_{mem} and R_{ob} to capture membership and ordering
 - ★ $R_{mem} [I,2,3] = \{I,2,3\}$
 - ★ $R_{ob} [I,2,3] = \{ (I,2), (I,3), (2,3) \}$

Relational Specification

- Let's define relations R_{mem} and R_{ob} to capture membership and ordering

★ $R_{mem} [I,2,3] = \{I,2,3\}$

★ $R_{ob} [I,2,3] = \{ (I,2), (I,3), (2,3) \}$

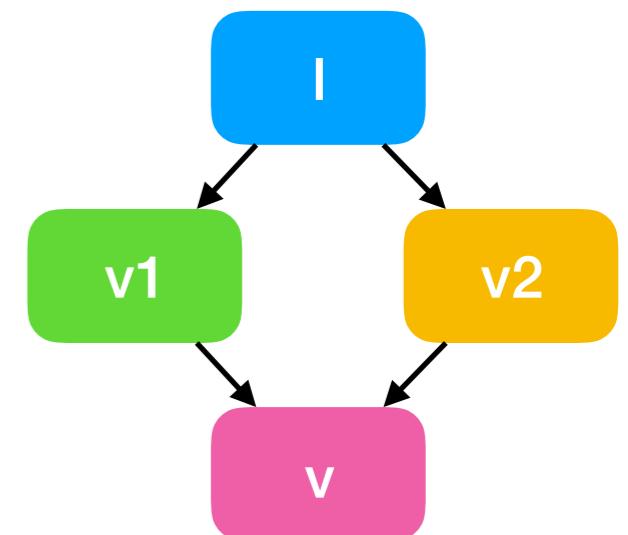


Relational Specification

- Let's define relations R_{mem} and R_{ob} to capture membership and ordering

★ $R_{mem} [l,2,3] = \{l,2,3\}$

★ $R_{ob} [l,2,3] = \{ (l,2), (l,3), (2,3) \}$



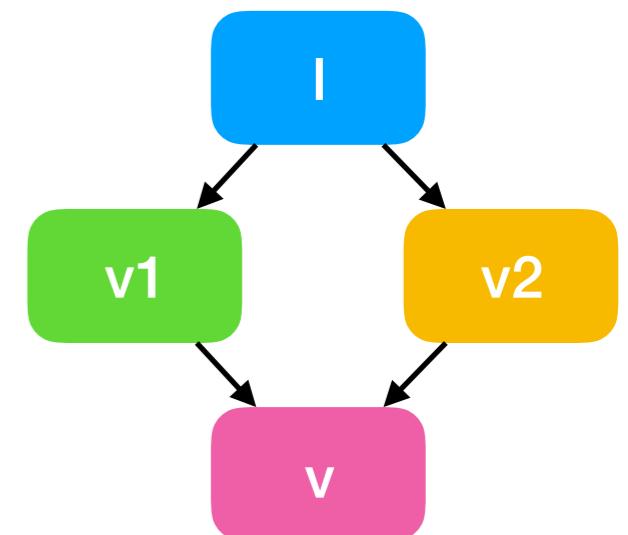
$$\begin{aligned} R_{mem}(v) &= R_{mem}(l) \cap R_{mem}(v_1) \cap R_{mem}(v_2) \\ &\cup R_{mem}(v_1) - R_{mem}(l) \quad \cup \quad R_{mem}(v_2) - R_{mem}(l) \end{aligned}$$

Relational Specification

- Let's define relations R_{mem} and R_{ob} to capture membership and ordering

★ $R_{mem} [l,2,3] = \{l,2,3\}$

★ $R_{ob} [l,2,3] = \{ (l,2), (l,3), (2,3) \}$



$$\begin{aligned} R_{mem}(v) &= R_{mem}(l) \cap R_{mem}(v_1) \cap R_{mem}(v_2) \\ &\cup R_{mem}(v_1) - R_{mem}(l) \quad \cup \quad R_{mem}(v_2) - R_{mem}(l) \end{aligned}$$

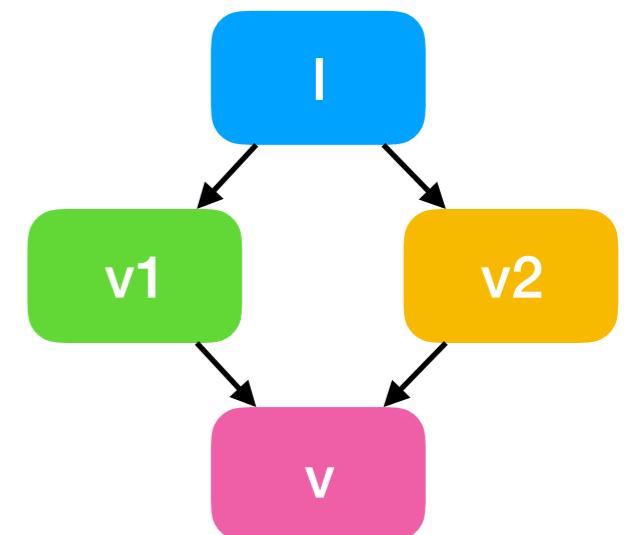
I. Any element popped in either v_1 or v_2 does not remain in v

Relational Specification

- Let's define relations R_{mem} and R_{ob} to capture membership and ordering

★ $R_{mem}[l,2,3] = \{l,2,3\}$

★ $R_{ob}[l,2,3] = \{(l,2), (l,3), (2,3)\}$



$$\begin{aligned} R_{mem}(v) &= R_{mem}(l) \cap R_{mem}(v_1) \cap R_{mem}(v_2) \\ &\cup R_{mem}(v_1) - R_{mem}(l) \quad \cup \quad R_{mem}(v_2) - R_{mem}(l) \end{aligned}$$

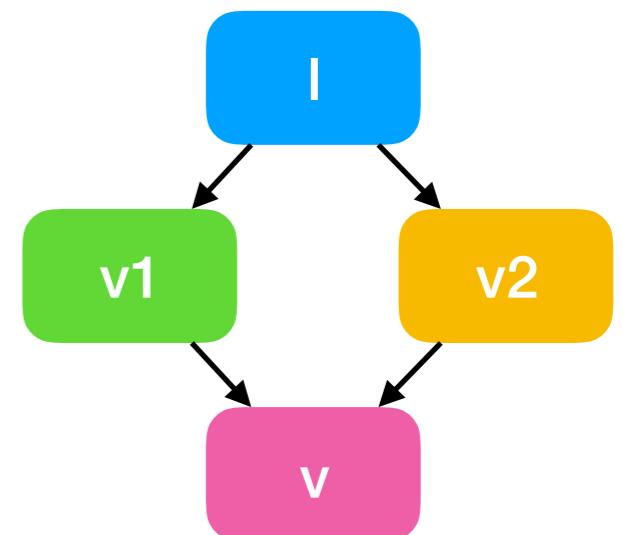
- Any element popped in either v_1 or v_2 does not remain in v
- Any element pushed into either v_1 or v_2 appears in v

Relational Specification

- Let's define relations R_{mem} and R_{ob} to capture membership and ordering

- $\star R_{mem}[l,2,3] = \{l,2,3\}$

- $\star R_{ob}[l,2,3] = \{ (l,2), (l,3), (2,3) \}$

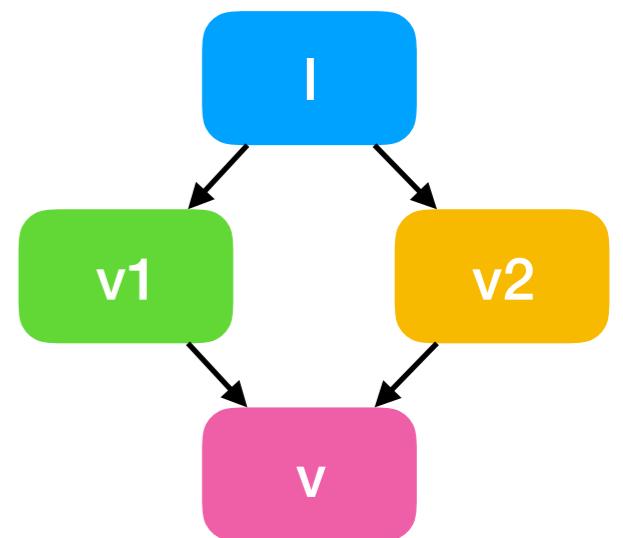


$$R_{mem}(v) = R_{mem}(l) \cap R_{mem}(v_1) \cap R_{mem}(v_2) \\ \cup R_{mem}(v_1) - R_{mem}(l) \cup R_{mem}(v_2) - R_{mem}(l)$$

- Any element popped in either v_1 or v_2 does not remain in v
- Any element pushed into either v_1 or v_2 appears in v
- An element that remains untouched in l, v_1, v_2 remains in v

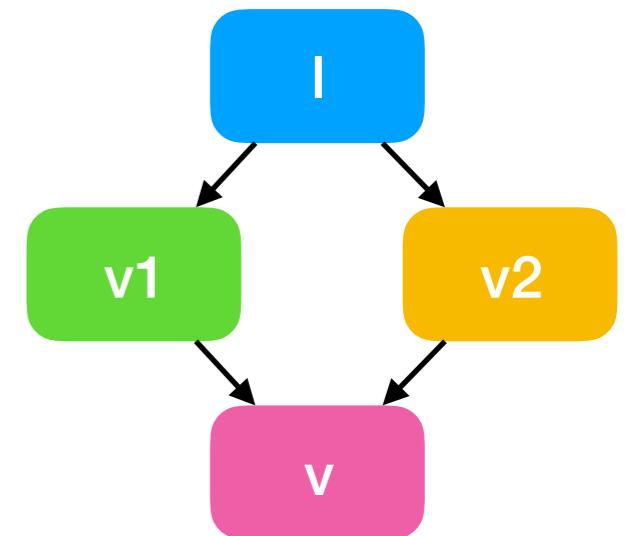
Relational Specification

$$\begin{aligned} R_{ob}(v) \supseteq & (R_{ob}(l) \cap R_{ob}(v_1) \cap R_{ob}(v_2)) \\ & \cup R_{ob}(v_1) - R_{ob}(l) \cup R_{ob}(v_2) - R_{ob}(l)) \\ \cap & (R_{mem}(v) \times R_{mem}(v)) \end{aligned}$$



Relational Specification

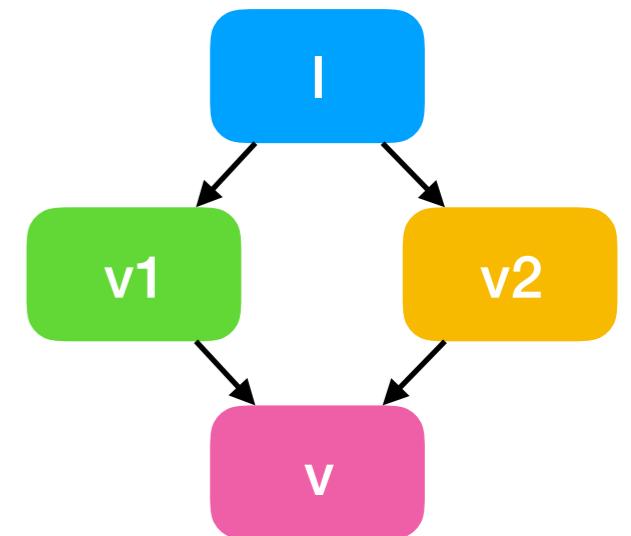
$$\begin{aligned} R_{ob}(v) \supseteq & (R_{ob}(l) \cap R_{ob}(v_1) \cap R_{ob}(v_2)) \\ & \cup (R_{ob}(v_1) - R_{ob}(l)) \cup (R_{ob}(v_2) - R_{ob}(l)) \\ & \cap (R_{mem}(v) \times R_{mem}(v)) \end{aligned}$$



- RHS has to be confined to $R_{mem}(v) \times R_{mem}(v)$ since certain orders might be missing
 - ★ Consider $l = [0], v1 = [0, l], v2 = [], v = [l]$

Relational Specification

$$\begin{aligned} R_{ob}(v) \supseteq & (R_{ob}(l) \cap R_{ob}(v_1) \cap R_{ob}(v_2) \\ & \cup R_{ob}(v_1) - R_{ob}(l) \cup R_{ob}(v_2) - R_{ob}(l)) \\ \cap & (R_{mem}(v) \times R_{mem}(v)) \end{aligned}$$



- RHS has to be confined to $R_{mem}(v) \times R_{mem}(v)$ since certain orders might be missing
 - ★ Consider $l = [0], v1 = [0, l], v2 = [], v = [l]$
- RHS is an underspecification since orders between concurrent insertions will only be present in $R_{ob}(v)$
 - ★ Consider $l = [], v1 = [0], v2 = [l], v = [0, l]$

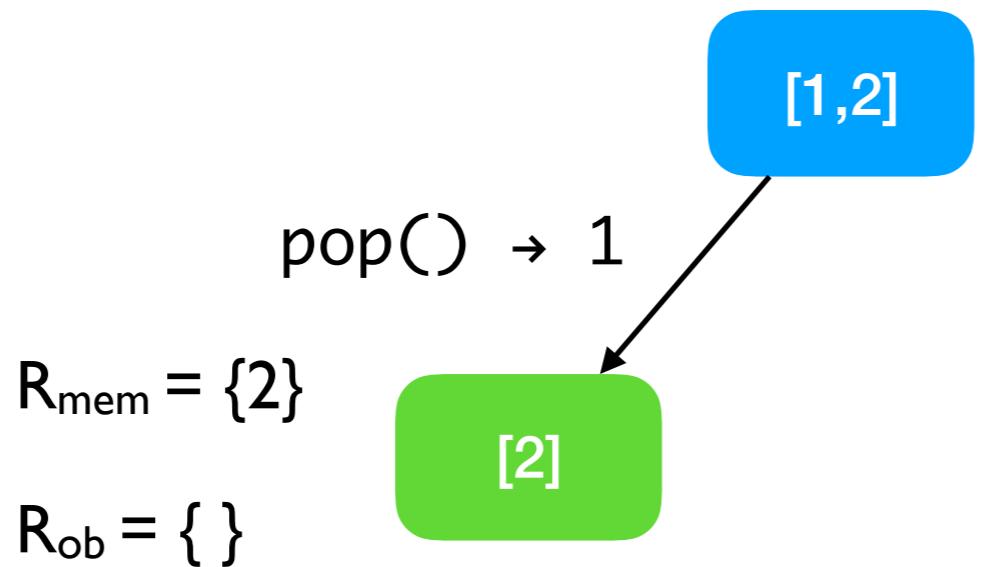
$$R_{mem} = \{1,2\}$$

$$R_{ob} = \{ (1,2) \}$$

[1,2]

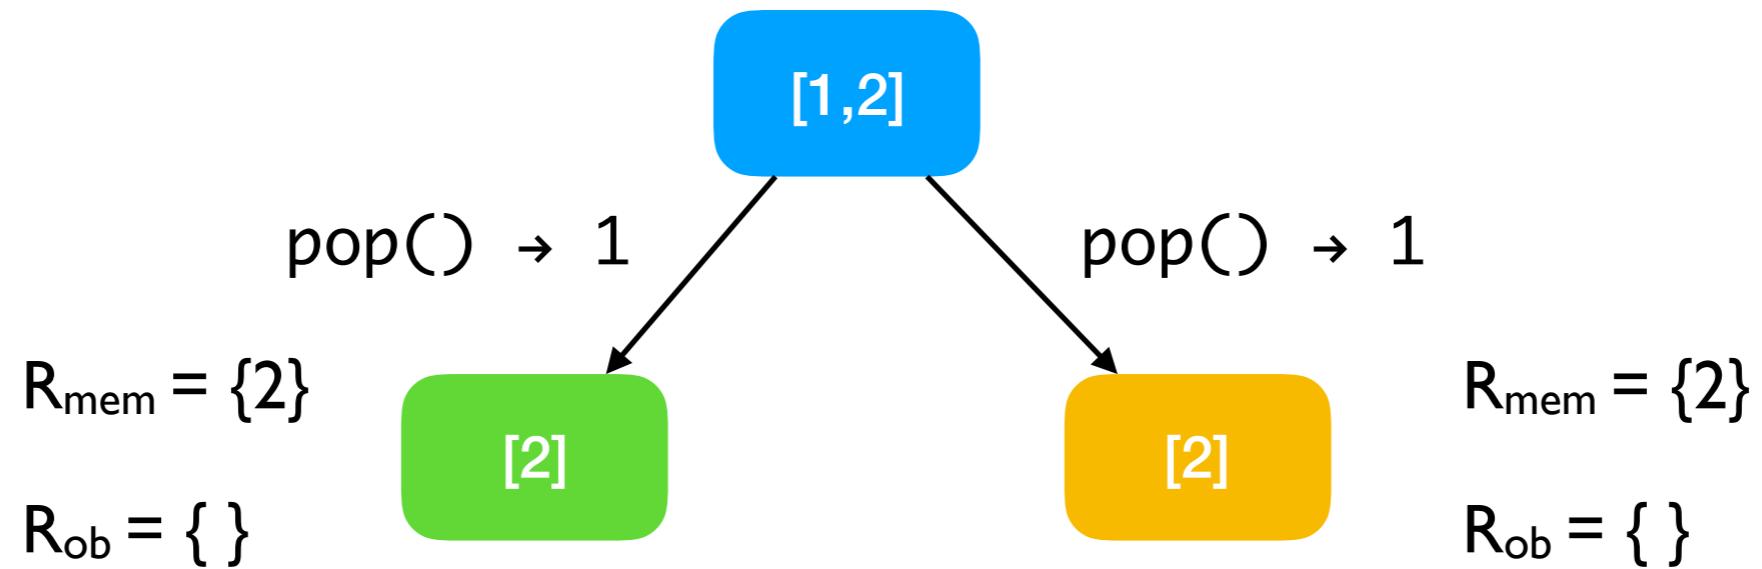
$$R_{mem} = \{1,2\}$$

$$R_{ob} = \{ (1,2) \}$$



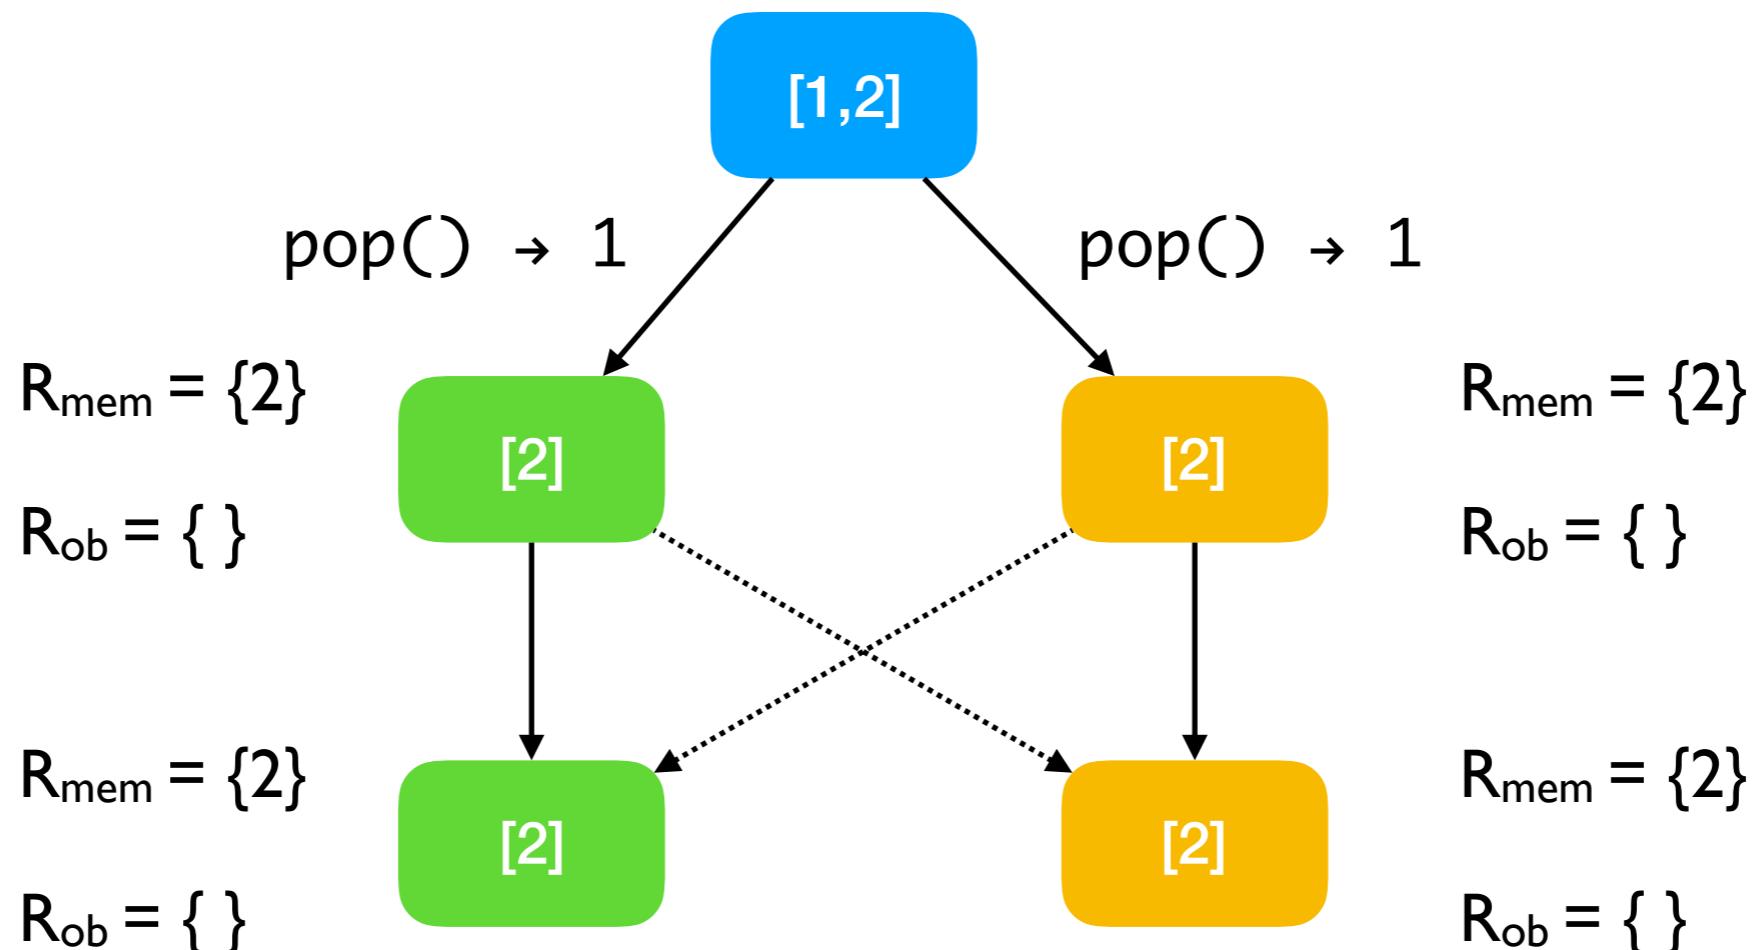
$$R_{mem} = \{1,2\}$$

$$R_{ob} = \{ (1,2) \}$$



$R_{mem} = \{1,2\}$

$R_{ob} = \{ (1,2) \}$



$$R_{mem} = \{ | \}$$

$$R_{ob} = \{ \}$$

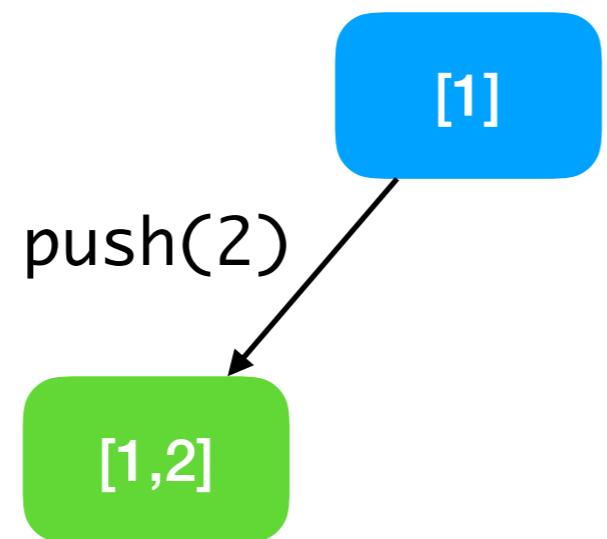
[1]

$$R_{mem} = \{1\}$$

$$R_{ob} = \{ \}$$

$$R_{mem} = \{1,2\}$$

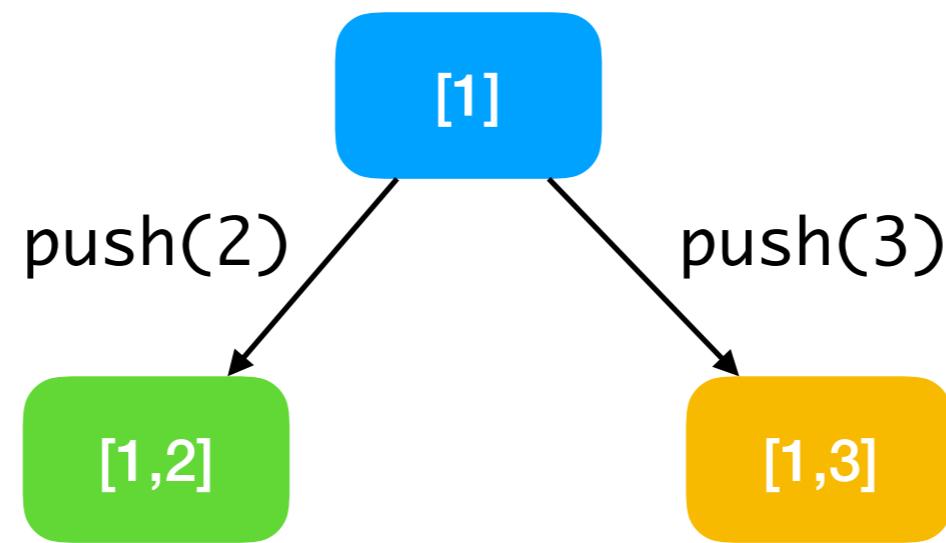
$$R_{ob} = \{ (1,2) \}$$



$$R_{mem} = \{1\}$$

$$R_{ob} = \{ \}$$

$$\begin{aligned} R_{mem} &= \{1,2\} \\ R_{ob} &= \{ (1,2) \} \end{aligned}$$

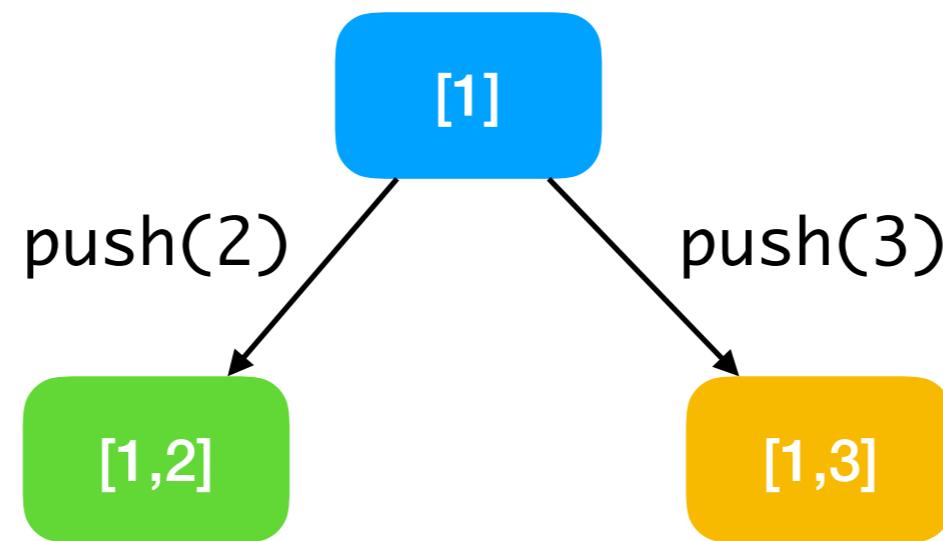


$$\begin{aligned} R_{mem} &= \{1,3\} \\ R_{ob} &= \{ (1,3) \} \end{aligned}$$

$$R_{mem} = \{1\}$$

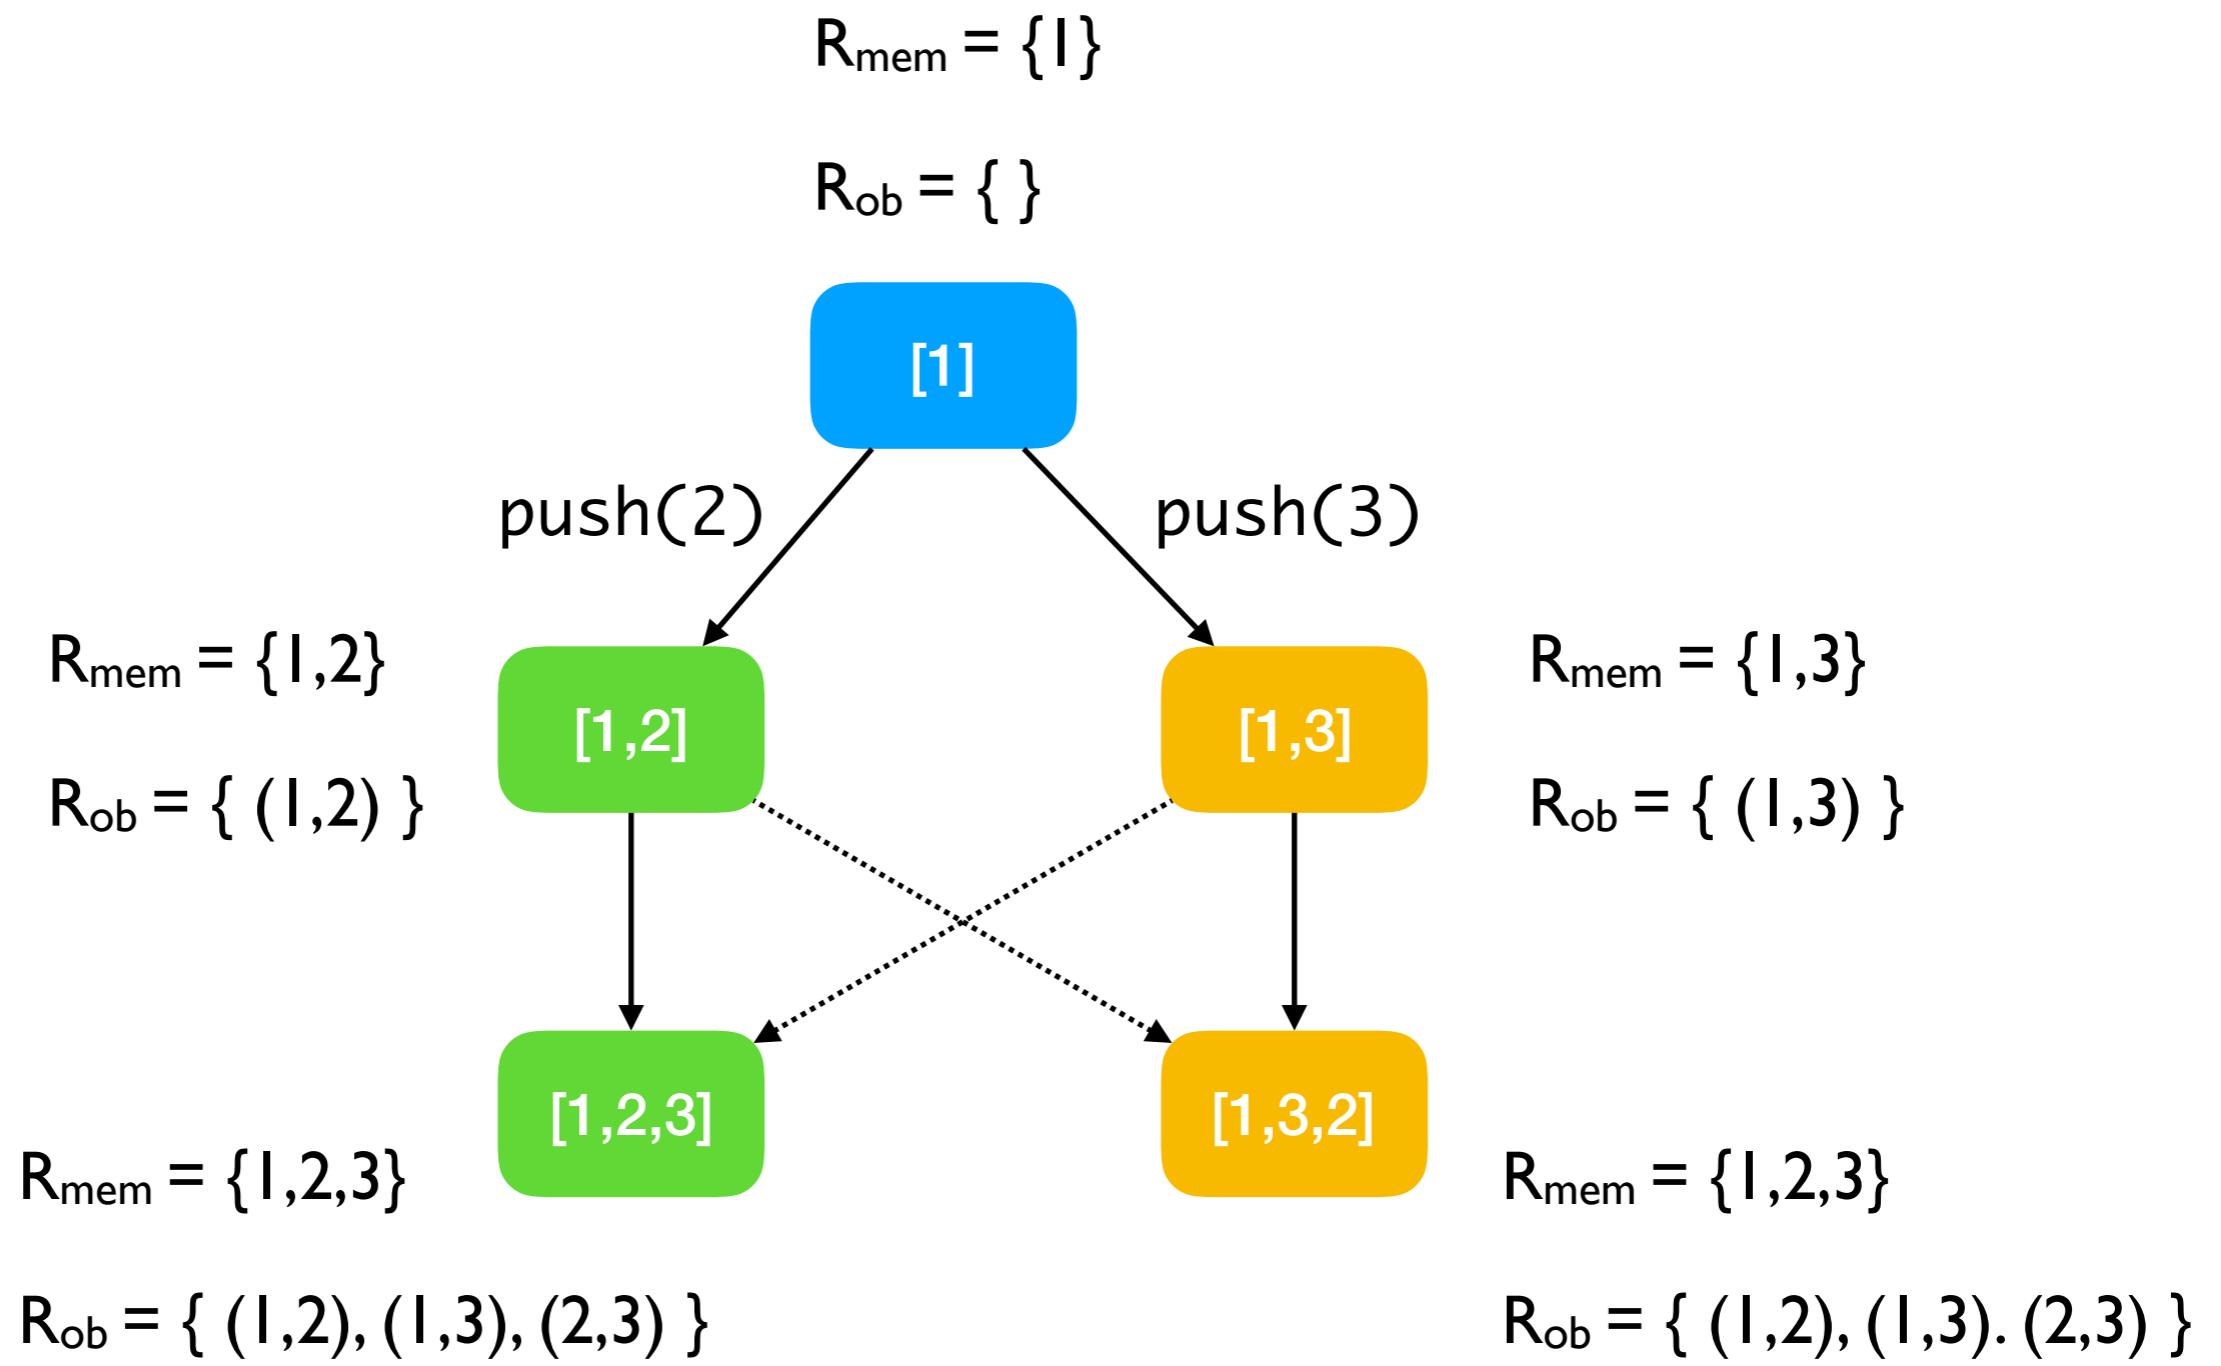
$$R_{ob} = \{ \}$$

$$\begin{aligned} R_{mem} &= \{1,2\} \\ R_{ob} &= \{ (1,2) \} \end{aligned}$$



$$\begin{aligned} R_{mem} &= \{1,3\} \\ R_{ob} &= \{ (1,3) \} \end{aligned}$$

Use < as an arbitration function between concurrent insertions



Use < as an arbitration function between concurrent insertions

Characteristic Relations

A sequence of relations \overline{R}_T is called a characteristic relation of a data type T , if for every $x : T$ and $y : T$, $\overline{R}_T(x) = \overline{R}_T(y)$ iff x and y are extensionally equal as interpreted under T .

Characteristic Relations

A sequence of relations \overline{R}_T is called a characteristic relation of a data type T , if for every $x : T$ and $y : T$, $\overline{R}_T(x) = \overline{R}_T(y)$ iff x and y are extensionally equal as interpreted under T .

- R_{mem} and R_{ob} are the characteristic relations of queue

Characteristic Relations

A sequence of relations \overline{R}_T is called a characteristic relation of a data type T , if for every $x : T$ and $y : T$, $\overline{R}_T(x) = \overline{R}_T(y)$ iff x and y are extensionally equal as interpreted under T .

- R_{mem} and R_{ob} are the characteristic relations of queue
- Appeals only to the sequential properties of the data type
 - ★ Ignore distribution when defining characteristic relations.

Synthesizing Merge

Synthesizing Merge

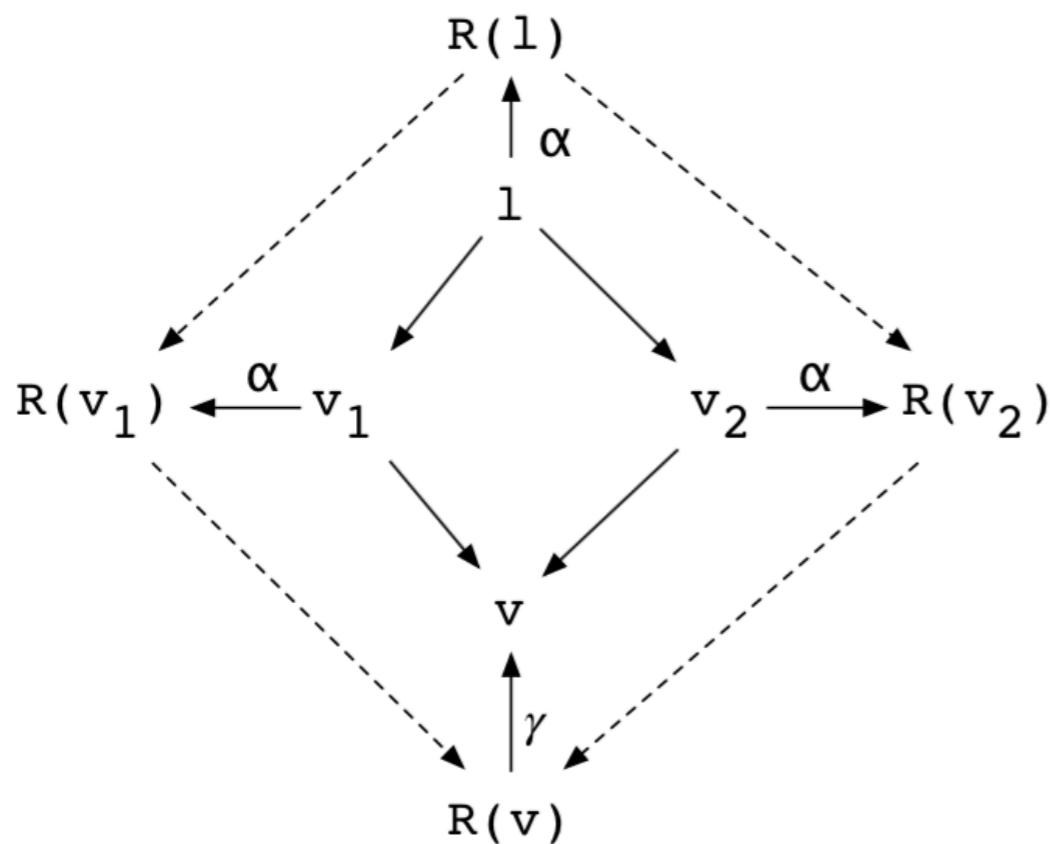
- Semantics of merge in relational domain is quite standard across data types

Synthesizing Merge

- Semantics of merge in relational domain is quite standard across data types
 - ★ Can we synthesise merge functions for arbitrary data type?

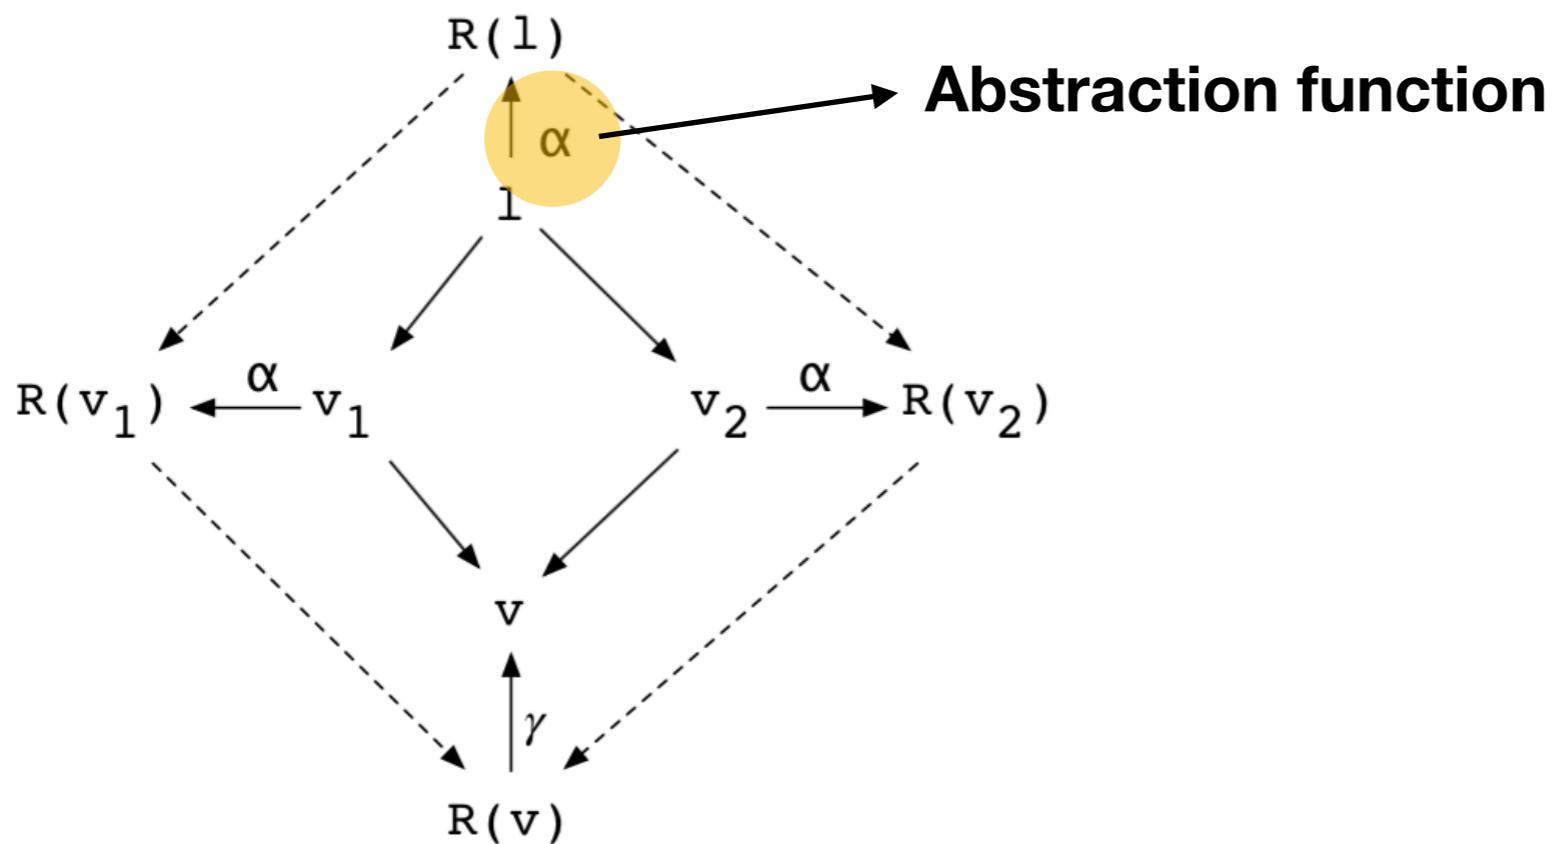
Synthesizing Merge

- Semantics of merge in relational domain is quite standard across data types
 - ★ Can we synthesise merge functions for arbitrary data type?



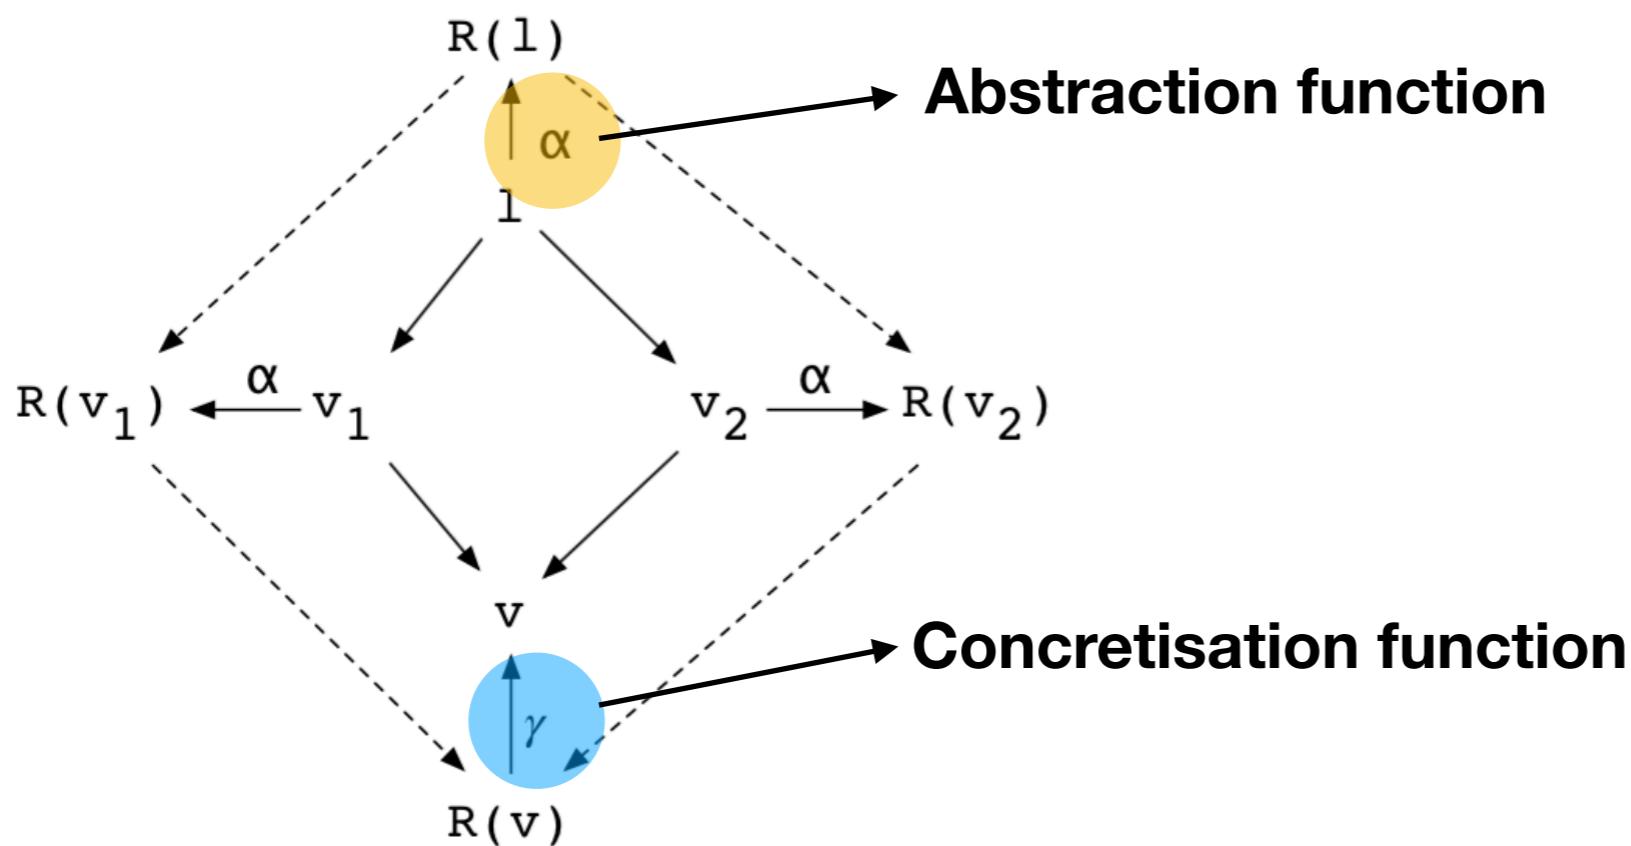
Synthesizing Merge

- Semantics of merge in relational domain is quite standard across data types
 - ★ Can we synthesise merge functions for arbitrary data type?



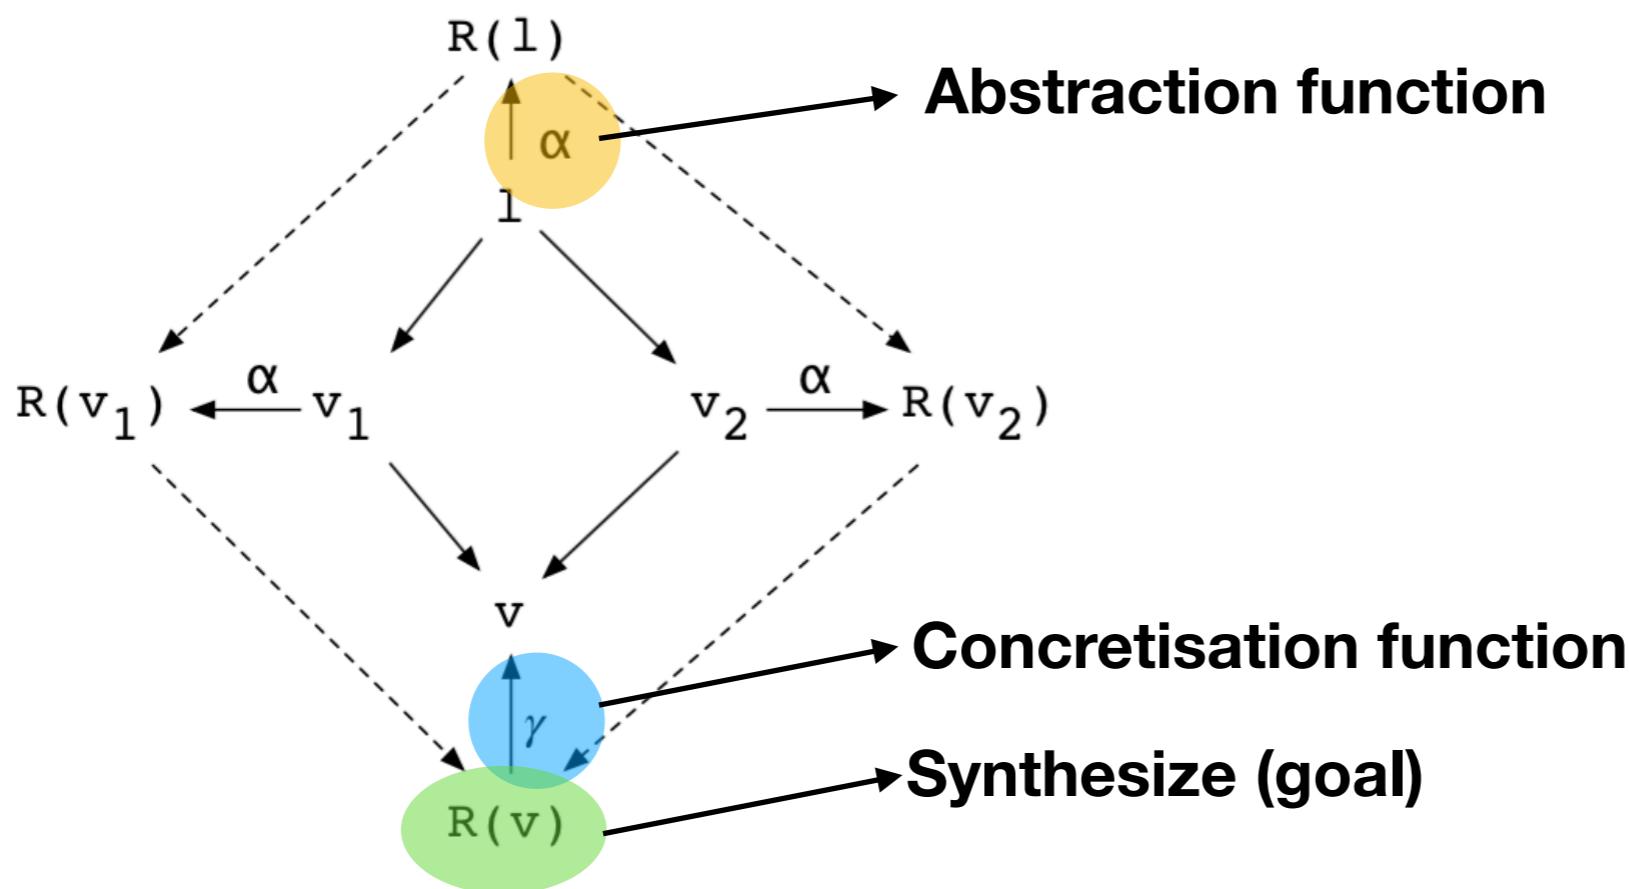
Synthesizing Merge

- Semantics of merge in relational domain is quite standard across data types
 - ★ Can we synthesise merge functions for arbitrary data type?



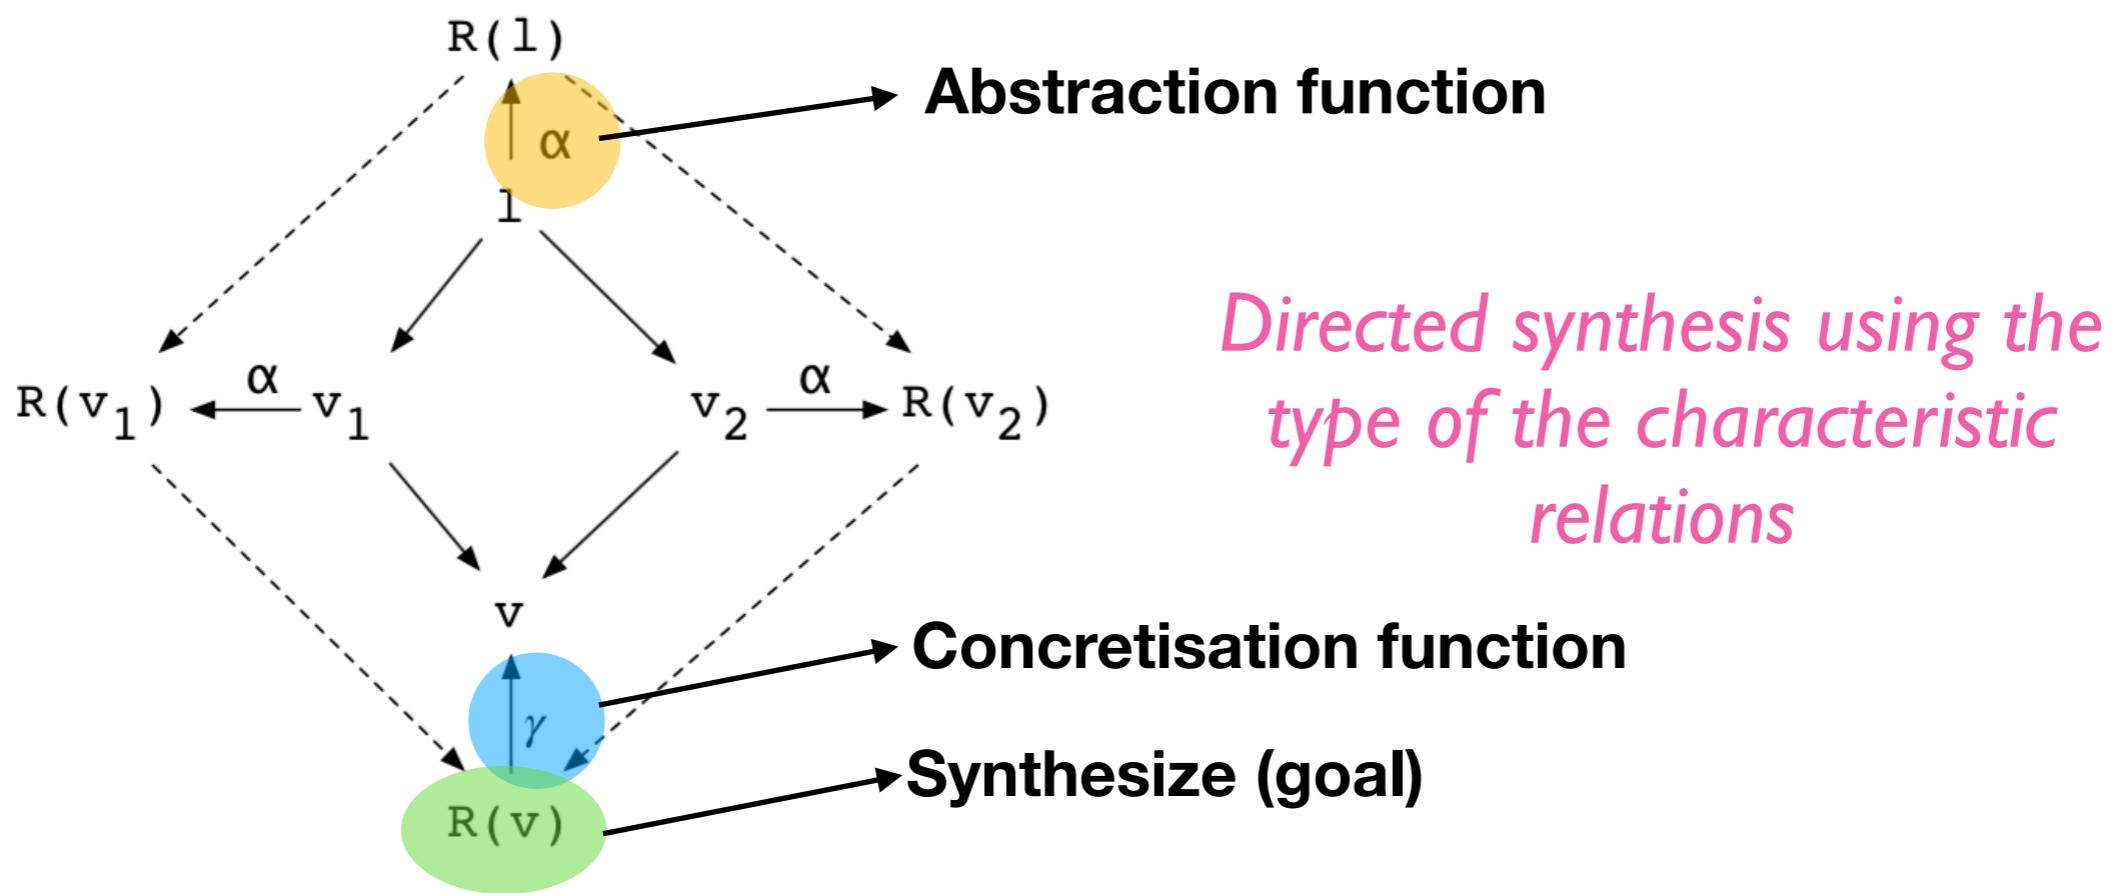
Synthesizing Merge

- Semantics of merge in relational domain is quite standard across data types
 - ★ Can we synthesise merge functions for arbitrary data type?



Synthesizing Merge

- Semantics of merge in relational domain is quite standard across data types
 - ★ Can we synthesise merge functions for arbitrary data type?



Abstraction Function: Queue

Abstraction Function: Queue

```
let rec  $R_{mem}$  = function
| [] ->  $\emptyset$ 
| x :: xs -> {x}  $\cup$   $R_{mem}(xs)$ 
```

Abstraction Function: Queue

```
let rec  $R_{mem}$  = function  
| [] ->  $\emptyset$   
| x :: xs -> {x}  $\cup$   $R_{mem}(xs)$ 
```

$$R_{mem} : \{v : \text{int list}\} \rightarrow \mathcal{P}(\text{int})$$

Abstraction Function: Queue

```
let rec  $R_{mem}$  = function  
| [] ->  $\emptyset$   
| x :: xs -> {x}  $\cup$   $R_{mem}(xs)$ 
```

$$R_{mem} : \{v : \text{int list}\} \rightarrow \mathcal{P}(\text{int})$$

```
let rec  $R_{ob}$  = function  
| [] ->  $\emptyset$   
| x :: xs -> ({x}  $\times$   $R_{mem}(xs)$ )  $\cup$   $R_{ob}(xs)$ 
```

Abstraction Function: Queue

```
let rec  $R_{mem}$  = function  
| [] ->  $\emptyset$   
| x :: xs -> {x}  $\cup$   $R_{mem}(xs)$ 
```

$$R_{mem} : \{v : \text{int list}\} \rightarrow \mathcal{P}(\text{int})$$

```
let rec  $R_{ob}$  = function  
| [] ->  $\emptyset$   
| x :: xs -> ({x}  $\times$   $R_{mem}(xs)$ )  $\cup$   $R_{ob}(xs)$ 
```

$$R_{ob} : \{v : \text{int list}\} \rightarrow \mathcal{P}(R_{mem}(v) \times R_{mem}(v))$$

Abstraction Function: Binary Tree

```
type 'a tree = | E  
              | N of 'a tree * 'a * 'a tree
```

Abstraction Function: Binary Tree

```
type 'a tree = | E
               | N of 'a tree * 'a * 'a tree

let rec Rmem = function
| E -> []
| N(l,x,r) -> Rmem(l) ∪ {x} ∪ Rmem(r)
```

Abstraction Function: Binary Tree

```
type 'a tree = | E  
              | N of 'a tree * 'a * 'a tree
```

```
let rec Rmem = function  
| E -> ∅  
| N(l,x,r) -> Rmem(l) ∪ {x} ∪ Rmem(r)
```

```
type label = L | R  
let rec Rto = function  
| E -> ∅  
| N(l,x,r) ->  
  let l_des = {x} × {L} × Rmem(l) in  
  let r_des = {x} × {R} × Rmem(r) in  
  Rto(l) ∪ l_des ∪ r_des ∪ Rto(r)
```

Abstraction Function: Binary Heap

```
type 'a tree = | E  
              | N of 'a tree * 'a * 'a tree  
  
let rec Rmem = function  
| E -> []  
| N(l,x,r) -> Rmem(l) ∪ {x} ∪ Rmem(r)
```

Abstraction Function: Binary Heap

```
type 'a tree = | E  
              | N of 'a tree * 'a * 'a tree
```

```
let rec Rmem = function  
| E -> ∅  
| N(l, x, r) -> Rmem(l) ∪ {x} ∪ Rmem(r)
```

```
let rec Rans = function  
| E -> ∅  
| N(l, x, r) ->  
  let des_x = Rmem(l) ∪ Rmem(r) in  
  let rans = {x} × des_x in  
  Rans(l) ∪ rans ∪ Rans(r)
```

Data Type	Characteristic Relations
Binary Heap	Membership (R_{mem}), Ancestor ($R_{ans} \subseteq R_{mem} \times R_{mem}$)
Priority Queue	Membership (R_{mem})
Set	Membership (R_{mem})
Graph	Vertex (R_V), Edge (R_E)
Functional Map	Key-Value (R_{kv})
List	Membership (R_{mem}), Order (R_{ob})
Binary Tree	Membership (R_{mem}), Tree-order ($R_{to} \subseteq R_{mem} \times \text{label} \times R_{mem}$)
Binary Search Tree	Membership (R_{mem})

Table 1. Characteristic relations for various data types

Compositionality: Pair

Compositionality: Pair

- The merge of a pair is the merge of the corresponding constituents

Compositionality: Pair

- The merge of a pair is the merge of the corresponding constituents
- A pair data type is defined by the relations:

```
let Rfst = fun (x,_) -> {x}      let Rsnd = fun (_,y) -> {y}
```

Compositionality: Pair

- The merge of a pair is the merge of the corresponding constituents
- A pair data type is defined by the relations:

```
let Rfst = fun (x,_) -> {x}    let Rsnd = fun (_,y) -> {y}
```

- Assume that the pair is composed of 2 counters. The counter merge spec is: $\phi_c(l, v_1, v_2, v) \Leftrightarrow v = l + (v_1 - l) + (v_2 - l)$

Compositionality: Pair

- The merge of a pair is the merge of the corresponding constituents
- A pair data type is defined by the relations:

```
let Rfst = fun (x,_) -> {x}      let Rsnd = fun (_,y) -> {y}
```

- Assume that the pair is composed of 2 counters. The counter merge spec is: $\phi_c(l, v_1, v_2, v) \Leftrightarrow v = l + (v_1 - l) + (v_2 - l)$
- Then, pair merge spec is:

$$\begin{aligned}\phi_{c \times c}(l, v_1, v_2, v) \Leftrightarrow & \forall x, y, z, s. x \in R_{fst}(l) \wedge y \in R_{fst}(v_1) \wedge z \in R_{fst}(v_2) \\ & \wedge \phi_c(x, y, z, s) \Rightarrow s \in R_{fst}(v) \\ & \wedge \forall s. s \in R_{fst}(v) \Rightarrow \exists x, y, z. x \in R_{fst}(l) \wedge y \in R_{fst}(v_1) \\ & \wedge z \in R_{fst}(v_2) \wedge \phi_c(x, y, z, s) \\ & \wedge \dots (\text{respectively for } R_{snd})\end{aligned}$$

Generalising Pairs to Ordinates

Generalising Pairs to Ordinates

- An alternative characteristic relation for a pair is:

```
let Rpair (x,y) = {(1,x), (2,y)}
```

Generalising Pairs to Ordinates

- An alternative characteristic relation for a pair is:

```
let Rpair (x, y) = {(1, x), (2, y)}
```

$$R_{pair} : \{v : \text{counter} * \text{counter}\} \rightarrow \mathcal{P}(\text{int} \times \text{counter})$$

Generalising Pairs to Ordinates

- An alternative characteristic relation for a pair is:

```
let Rpair (x, y) = {(1, x), (2, y)}
```

$$R_{pair} : \{v : \text{counter} * \text{counter}\} \rightarrow \mathcal{P}(\text{int} \times \text{counter})$$

- Corresponding merge specification is:

Generalising Pairs to Ordinates

- An alternative characteristic relation for a pair is:

```
let Rpair (x, y) = {(1, x), (2, y)}
```

$$R_{pair} : \{v : \text{counter} * \text{counter}\} \rightarrow \mathcal{P}(\text{int} \times \text{counter})$$

- Corresponding merge specification is:

$$\begin{aligned} \phi_{c \times c} &= \forall(k : \text{int}). \forall(x, y, z, s : \text{counter}). (k, x) \in R_{pair}(l) \wedge (k, y) \in R_{pair}(v_1) \\ &\quad \wedge (k, z) \in R_{pair}(v_2) \wedge \phi_c(x, y, z, s) \Rightarrow (k, s) \in R(v) \\ &\wedge \forall(k : \text{int}). \forall(s : \text{counter}). (k, s) \in R_{pair}(v) \Rightarrow \exists(x, y, z : \text{counter}). (k, x) \in R_{pair}(l) \\ &\quad \wedge (k, y) \in R_{pair}(v_1) \wedge (k, z) \in R_{pair}(v_2) \wedge \phi_c(x, y, z, s) \end{aligned}$$

Generalising Pairs to Ordinates

- An alternative characteristic relation for a pair is:

```
let Rpair (x, y) = {(1, x), (2, y)}
```

$$R_{pair} : \{v : \text{counter} * \text{counter}\} \rightarrow \mathcal{P}(\text{int} \times \text{counter})$$

- Corresponding merge specification is:

$$\begin{aligned} \phi_{c \times c} &= \forall(k : \text{int}). \forall(x, y, z, s : \text{counter}). (k, x) \in R_{pair}(l) \wedge (k, y) \in R_{pair}(v_1) \\ &\quad \wedge (k, z) \in R_{pair}(v_2) \wedge \phi_c(x, y, z, s) \Rightarrow (k, s) \in R(v) \\ &\wedge \forall(k : \text{int}). \forall(s : \text{counter}). (k, s) \in R_{pair}(v) \Rightarrow \exists(x, y, z : \text{counter}). (k, x) \in R_{pair}(l) \\ &\quad \wedge (k, y) \in R_{pair}(v_1) \wedge (k, z) \in R_{pair}(v_2) \wedge \phi_c(x, y, z, s) \end{aligned}$$

- Given appropriate characteristic relation for a n-tuple ($R_{n\text{-tuple}}$), the same merge specification can be used.

Generalising Pairs to Ordinates

- Similar encoding can be given to maps with non-mergeable types as keys and mergeable types as values

$$\begin{aligned} R_k &: \{v : (\text{string}, \text{int}) \text{ map}\} \rightarrow \mathcal{P}(\text{string}), \\ R_{kv} &: \{v : (\text{string}, \text{int}) \text{ map}\} \rightarrow \mathcal{P}(R_k(v) \times \text{counter}) \end{aligned}$$

Types of Characteristic Relations

Types of Characteristic Relations

- Practical characteristic relations fall into 3 types:

Types of Characteristic Relations

- Practical characteristic relations fall into 3 types:

- ★ Membership (R_{mem})

$R : \{v : T\} \rightarrow \mathcal{P}(\overline{T})$, where T is a non-mergeable type

Types of Characteristic Relations

- Practical characteristic relations fall into 3 types:

- ★ Membership (R_{mem})

$R : \{v : T\} \rightarrow \mathcal{P}(\overline{T})$, where T is a non-mergeable type

- ★ Ordering ($R_{\text{ob}}, R_{\text{ans}}, R_{\text{to}}$)

$R : \{v : T\} \rightarrow \mathcal{P}(\rho)$

- ◆ where ρ is a sequence of non-mergeable types and other relations, which flattens to a sequence of non-mergeable types

Types of Characteristic Relations

- Practical characteristic relations fall into 3 types:

- ★ Membership (R_{mem})

$R : \{v : T\} \rightarrow \mathcal{P}(\overline{T})$, where T is a non-mergeable type

- ★ Ordering ($R_{\text{ob}}, R_{\text{ans}}, R_{\text{to}}$)

$R : \{v : T\} \rightarrow \mathcal{P}(\rho)$

- ◆ where ρ is a sequence of non-mergeable types and other relations, which flattens to a sequence of non-mergeable types

- ★ Ordinates ($R_{\text{pair}}, R_{\text{kv}}$)

$R : \{v : T\} \rightarrow \mathcal{P}(\overline{T} \times \bar{\tau})$, where τ is a mergeable type

Deriving merge spec

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\bar{T})}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{x} : \bar{T}). \bar{x} \in (R(l) \diamond R(v_1) \diamond R(v_2)) \Leftrightarrow \bar{x} \in R(v)} \quad [\text{SET-MERGE}]$$

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\rho) \quad \lfloor \rho \rfloor = \bar{T}}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{x} : \bar{T}). \bar{x} \in (R(l) \diamond R(v_1) \diamond R(v_2) \cap \rho) \Rightarrow \bar{x} \in R(v)} \quad [\text{ORDER-MERGE-1}]$$

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\rho) \quad \lfloor \rho \rfloor = \bar{T}}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{x} : \bar{T}). \bar{x} \in R(v) \Rightarrow \bar{x} \in \rho} \quad [\text{ORDER-MERGE-2}]$$

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\rho) \quad \lfloor \rho \rfloor = \bar{T} \times \bar{\tau} \quad \bar{\tau} \neq \emptyset}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{k} : \bar{T}). \forall(\bar{x}, \bar{y}, \bar{z}, \bar{s} : \bar{\tau}). (\bar{k}, \bar{x}) \in R_+(l) \wedge (\bar{k}, \bar{y}) \in R_+(v_1) \wedge (\bar{k}, \bar{z}) \in R_+(v_2) \wedge \bar{k} \in (R_k(l) \diamond R_k(v_1) \diamond R_k(v_2)) \wedge \bigwedge_i \phi_{\tau_i}(x_i, y_i, z_i, s_i) \wedge (\bar{k}, \bar{s}) \in \rho \Rightarrow (\bar{k}, \bar{s}) \in R(v)} \quad [\text{REL-MERGE-1}]$$

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\rho) \quad \lfloor \rho \rfloor = \bar{T} \times \bar{\tau} \quad \bar{\tau} \neq \emptyset}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{k} : \bar{T}). \forall(\bar{s} : \bar{\tau}). (\bar{k}, \bar{s}) \in R(v) \Rightarrow (\bar{k}, \bar{s}) \in \rho \wedge \exists(\bar{x}, \bar{y}, \bar{z} : \bar{\tau}). (\bar{k}, \bar{x}) \in R_+(l) \wedge (\bar{k}, \bar{y}) \in R_+(v_1) \wedge (\bar{k}, \bar{z}) \in R_+(v_2) \wedge \bar{k} \in (R_k(l) \diamond R_k(v_1) \diamond R_k(v_2)) \wedge \bigwedge_i \phi_{\tau_i}(x_i, y_i, z_i, s_i)} \quad [\text{REL-MERGE-2}]$$

Deriving merge spec

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\bar{T})}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{x} : \bar{T}). \bar{x} \in (R(l) \diamond R(v_1) \diamond R(v_2)) \Leftrightarrow \bar{x} \in R(v)} \quad [\text{SET-MERGE}]$$

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\rho) \quad \lfloor \rho \rfloor = \bar{T}}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{x} : \bar{T}). \bar{x} \in (R(l) \diamond R(v_1) \diamond R(v_2) \cap \rho) \Rightarrow \bar{x} \in R(v)} \quad [\text{ORDER-MERGE-1}]$$

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\rho) \quad \lfloor \rho \rfloor = \bar{T}}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{x} : \bar{T}). \bar{x} \in R(v) \Rightarrow \bar{x} \in \rho} \quad [\text{ORDER-MERGE-2}]$$

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\rho) \quad \lfloor \rho \rfloor = \bar{T} \times \bar{\tau} \quad \bar{\tau} \neq \emptyset}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{k} : \bar{T}). \forall(\bar{x}, \bar{y}, \bar{z}, \bar{s} : \bar{\tau}). (\bar{k}, \bar{x}) \in R_+(l) \wedge (\bar{k}, \bar{y}) \in R_+(v_1) \wedge (\bar{k}, \bar{z}) \in R_+(v_2) \\ \wedge \bar{k} \in (R_k(l) \diamond R_k(v_1) \diamond R_k(v_2)) \wedge \bigwedge_i \phi_{\tau_i}(x_i, y_i, z_i, s_i) \wedge (\bar{k}, \bar{s}) \in \rho \Rightarrow (\bar{k}, \bar{s}) \in R(v)} \quad [\text{REL-MERGE-1}]$$

$$\frac{R : \{v : T\} \rightarrow \mathcal{P}(\rho) \quad \lfloor \rho \rfloor = \bar{T} \times \bar{\tau} \quad \bar{\tau} \neq \emptyset}{\phi_T(l, v_1, v_2, v) \supseteq \forall(\bar{k} : \bar{T}). \forall(\bar{s} : \bar{\tau}). (\bar{k}, \bar{s}) \in R(v) \Rightarrow (\bar{k}, \bar{s}) \in \rho \\ \wedge \exists(\bar{x}, \bar{y}, \bar{z} : \bar{\tau}). (\bar{k}, \bar{x}) \in R_+(l) \wedge (\bar{k}, \bar{y}) \in R_+(v_1) \wedge (\bar{k}, \bar{z}) \in R_+(v_2) \\ \wedge \bar{k} \in (R_k(l) \diamond R_k(v_1) \diamond R_k(v_2)) \wedge \bigwedge_i \phi_{\tau_i}(x_i, y_i, z_i, s_i)} \quad [\text{REL-MERGE-2}]$$

- Not complete, but practical
- Can derive merge spec for
 - ★ Data structures: Set, Heap, Graph, Queue, TreeDoc
 - ★ Larger apps: TPC-C, TPC-E, Twissandra, Rubis

Distributed Implementation

Distributed Implementation

- For making this programming model practical, we need to:
 - ★ Quickly compute LCA
 - ★ Optimise storage through sharing
 - ★ Optimise network transmissions (state based merge)

Distributed Implementation

- For making this programming model practical, we need to:
 - ★ Quickly compute LCA
 - ★ Optimise storage through sharing
 - ★ Optimise network transmissions (state based merge)
- Irmin
 - ★ A reimplementation of Git in pure OCaml
 - ★ Arbitrary OCaml objects, not just files + User-defined 3-way merges
 - ★ Only transmit diffs over the network
 - ★ Multiple storage backends including in-memory, filesystems, log-structured-merge database, distributed databases

Performance

Performance

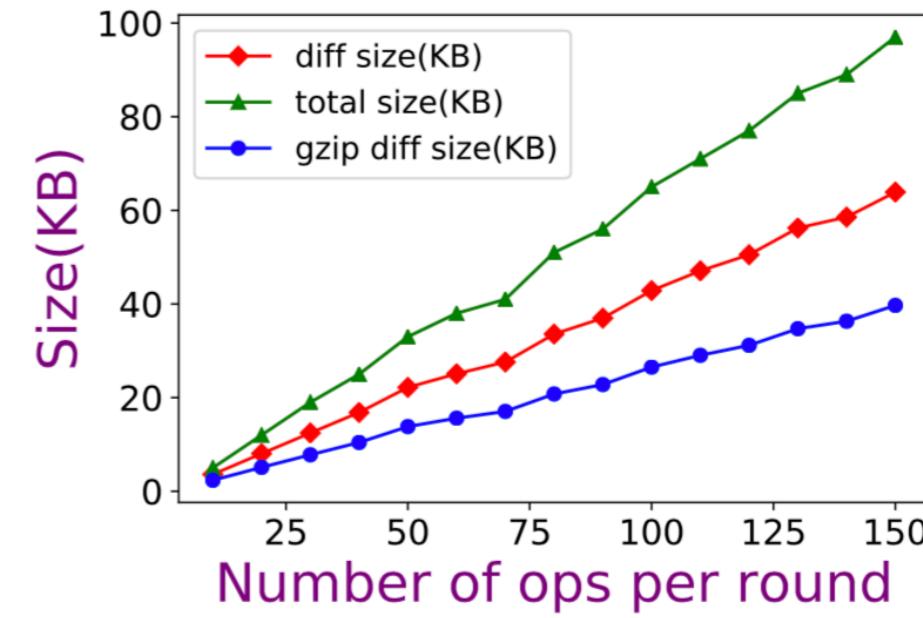
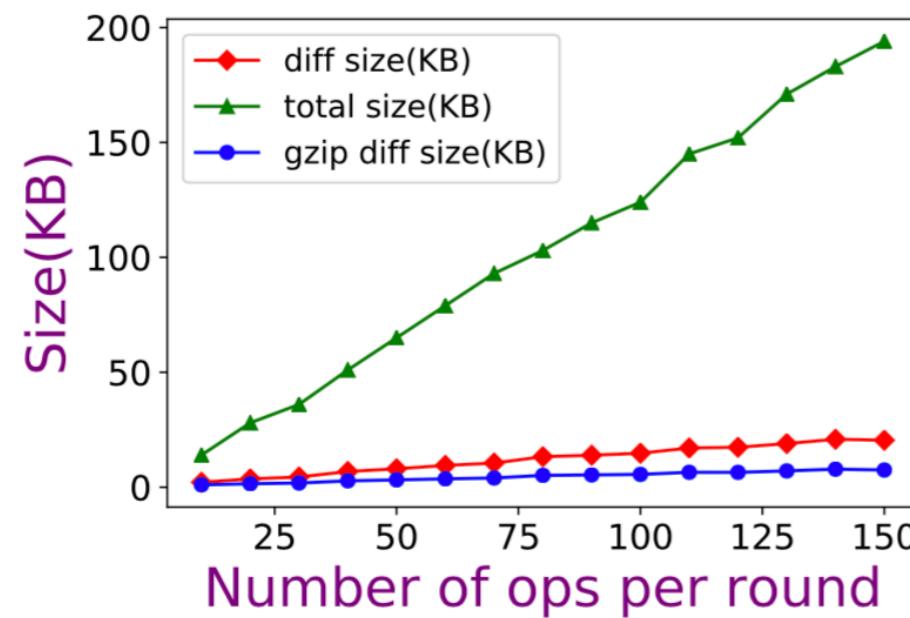
- What is the size of diff compared to the size of data structure?

Performance

- What is the size of diff compared to the size of data structure?
- Setup
 - ★ 2 Replicas, fixed number of rounds, each round has N operations
 - ★ 75% inserts, 25% deletions
 - ★ Synchronise after each round

Performance

- What is the size of diff compared to the size of data structure?
- Setup
 - ★ 2 Replicas, fixed number of rounds, each round has N operations
 - ★ 75% inserts, 25% deletions
 - ★ Synchronise after each round



Thanks for listening!