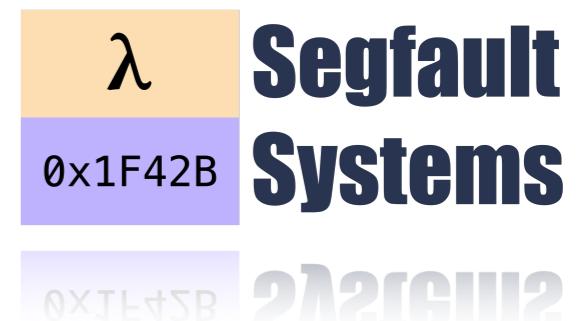


Retrofitting Parallelism onto OCaml

“KC” Sivaramakrishnan

IIT
MADRAS
SADAM



Industry



Projects





Industry



docker



Jane Street

No multicore support!



FACEBOOK



Tarides



Bloomberg

Projects



The *Astrée* Static Analyzer



COMPCERT

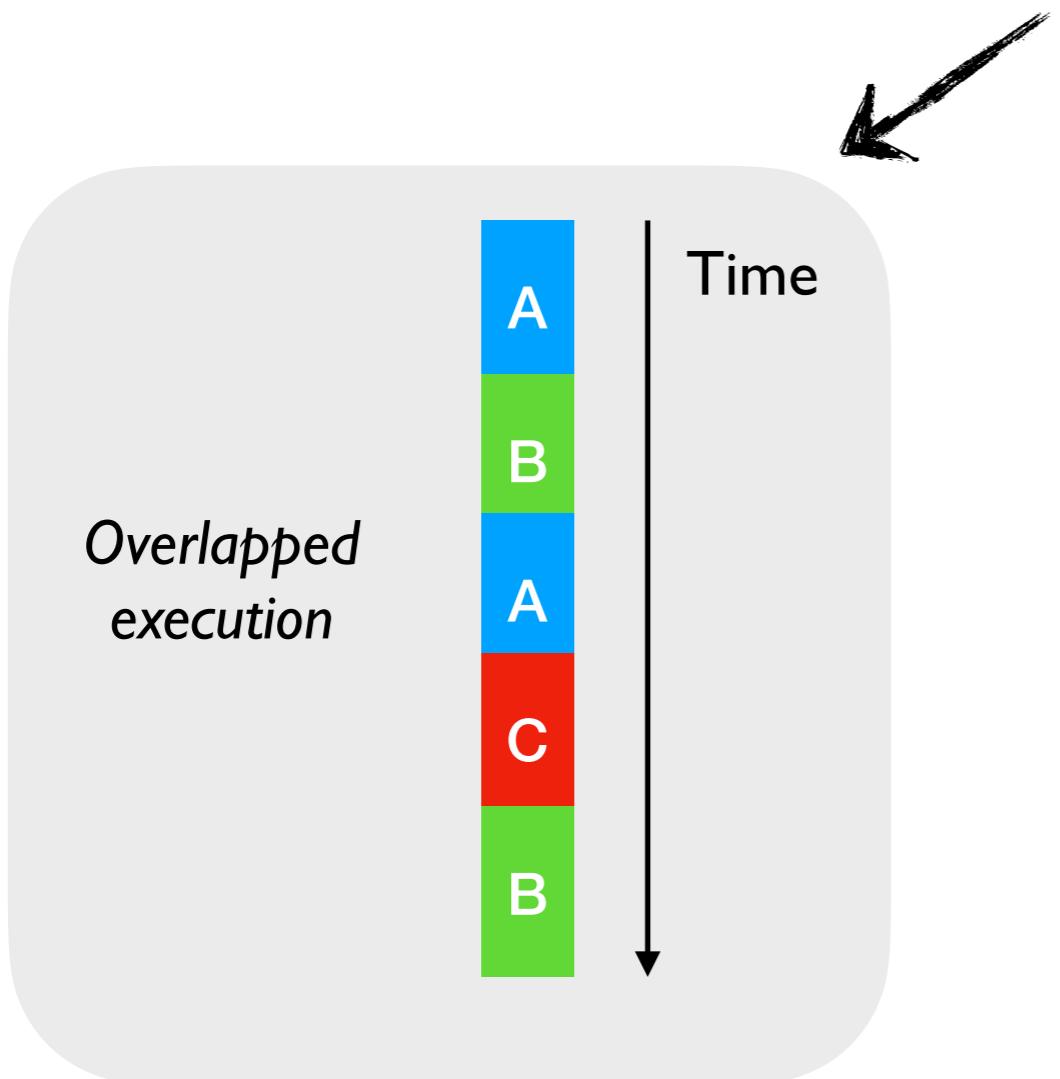


Multicore OCaml

- Adds native support for *concurrency* and *parallelism* to OCaml

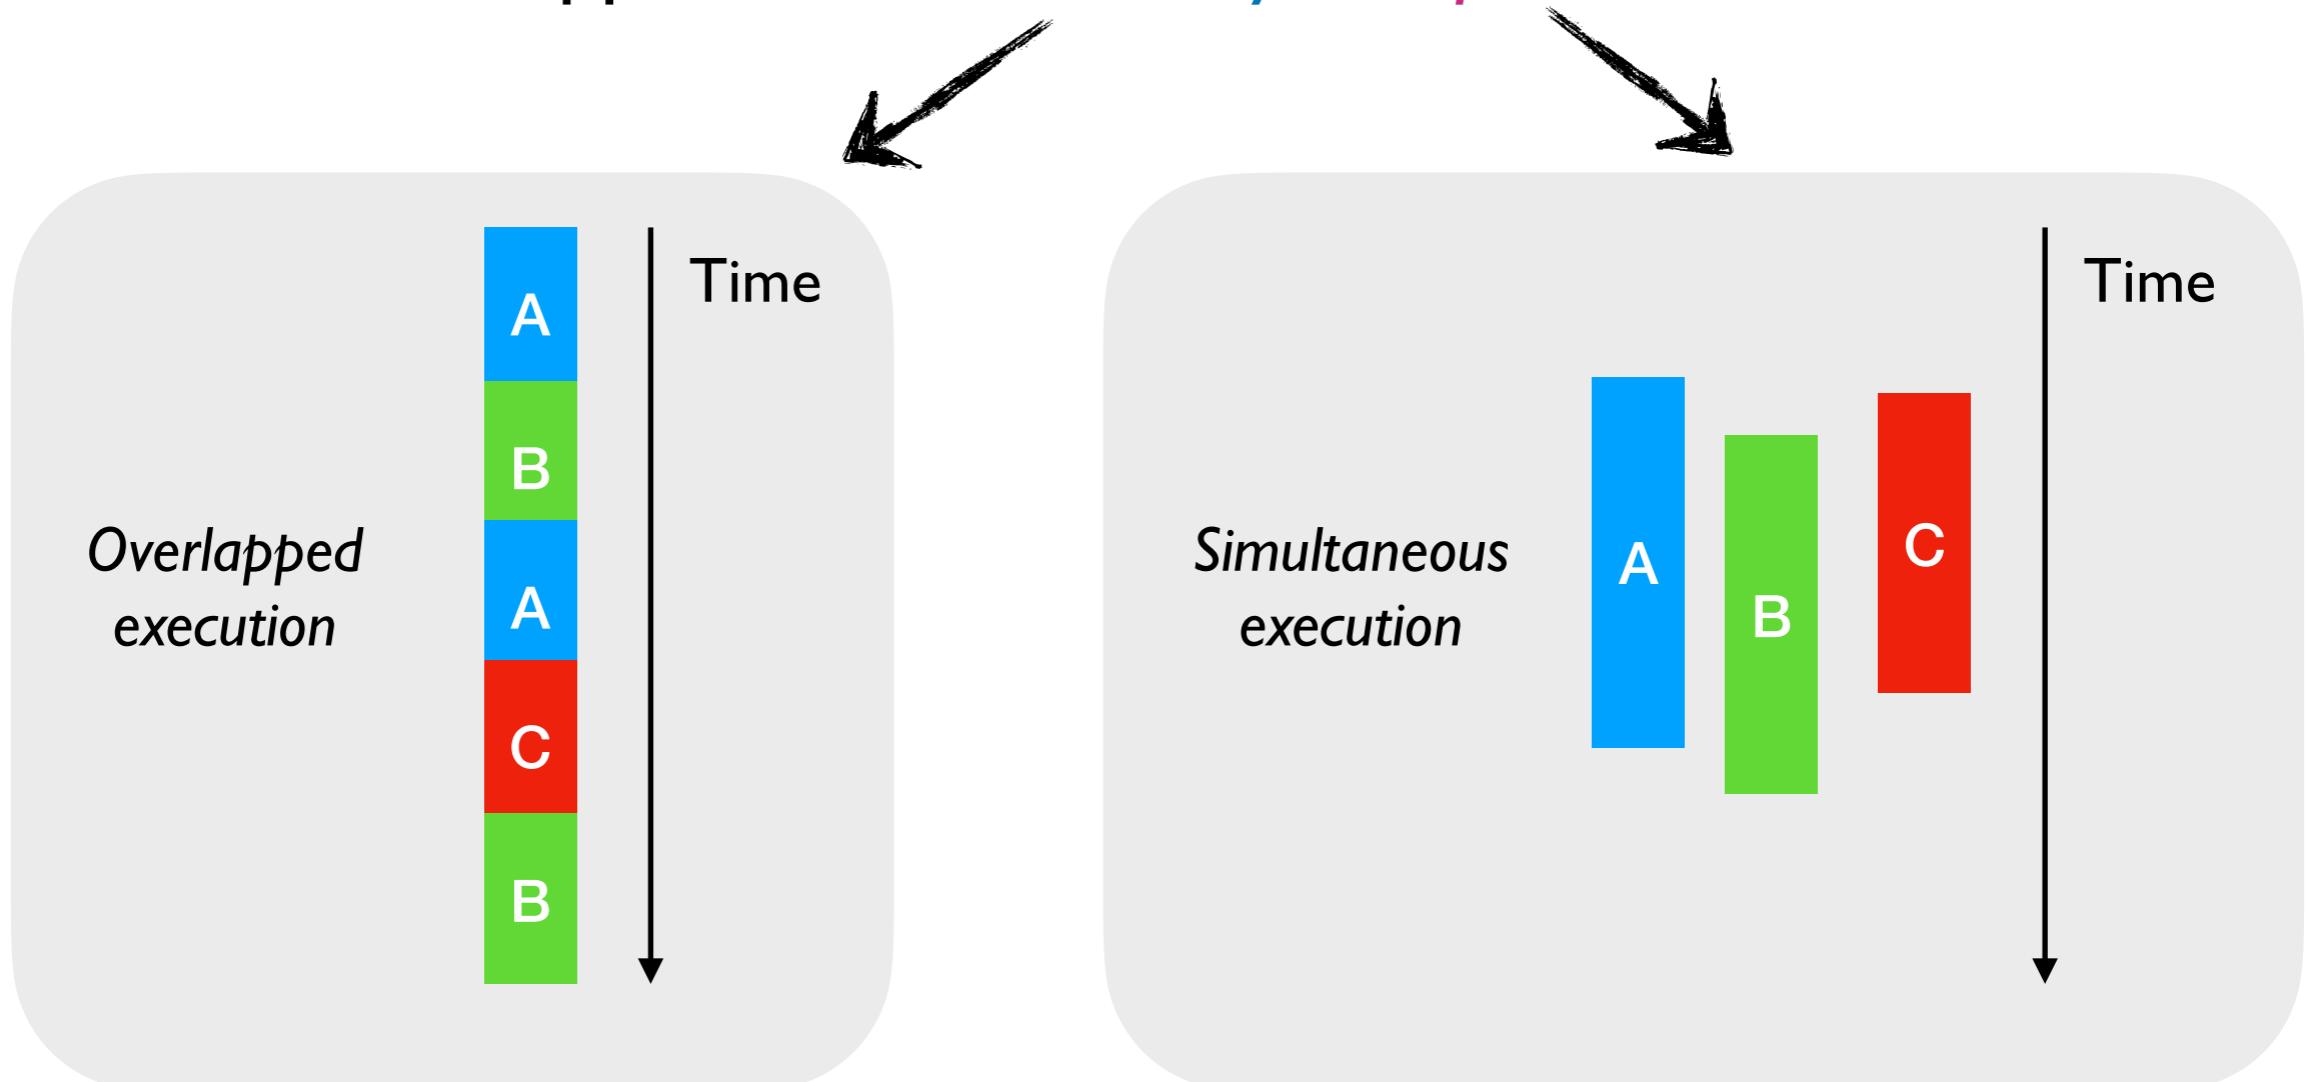
Multicore OCaml

- Adds native support for *concurrency* and *parallelism* to OCaml



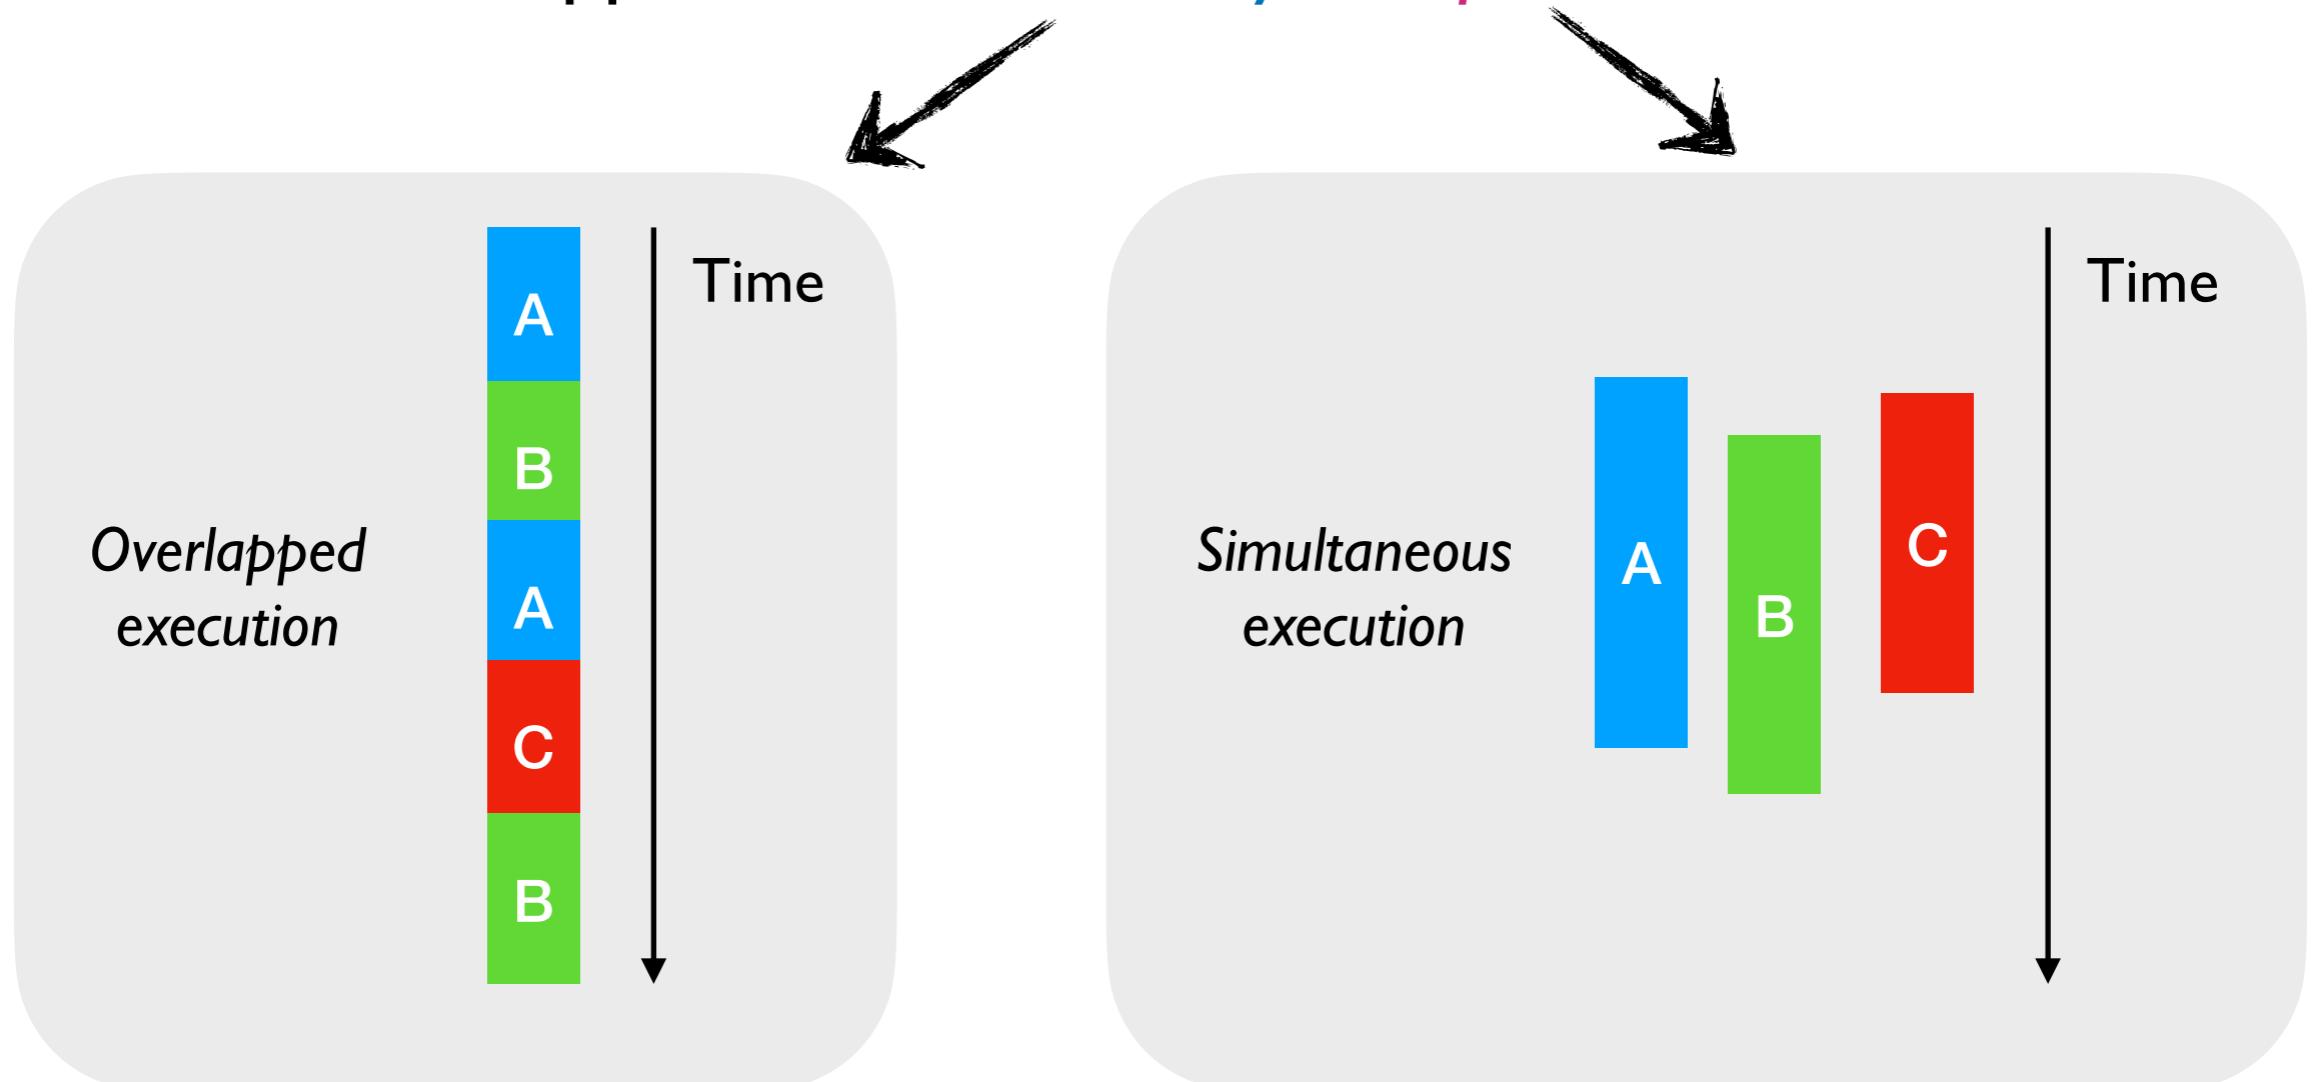
Multicore OCaml

- Adds native support for *concurrency* and *parallelism* to OCaml



Multicore OCaml

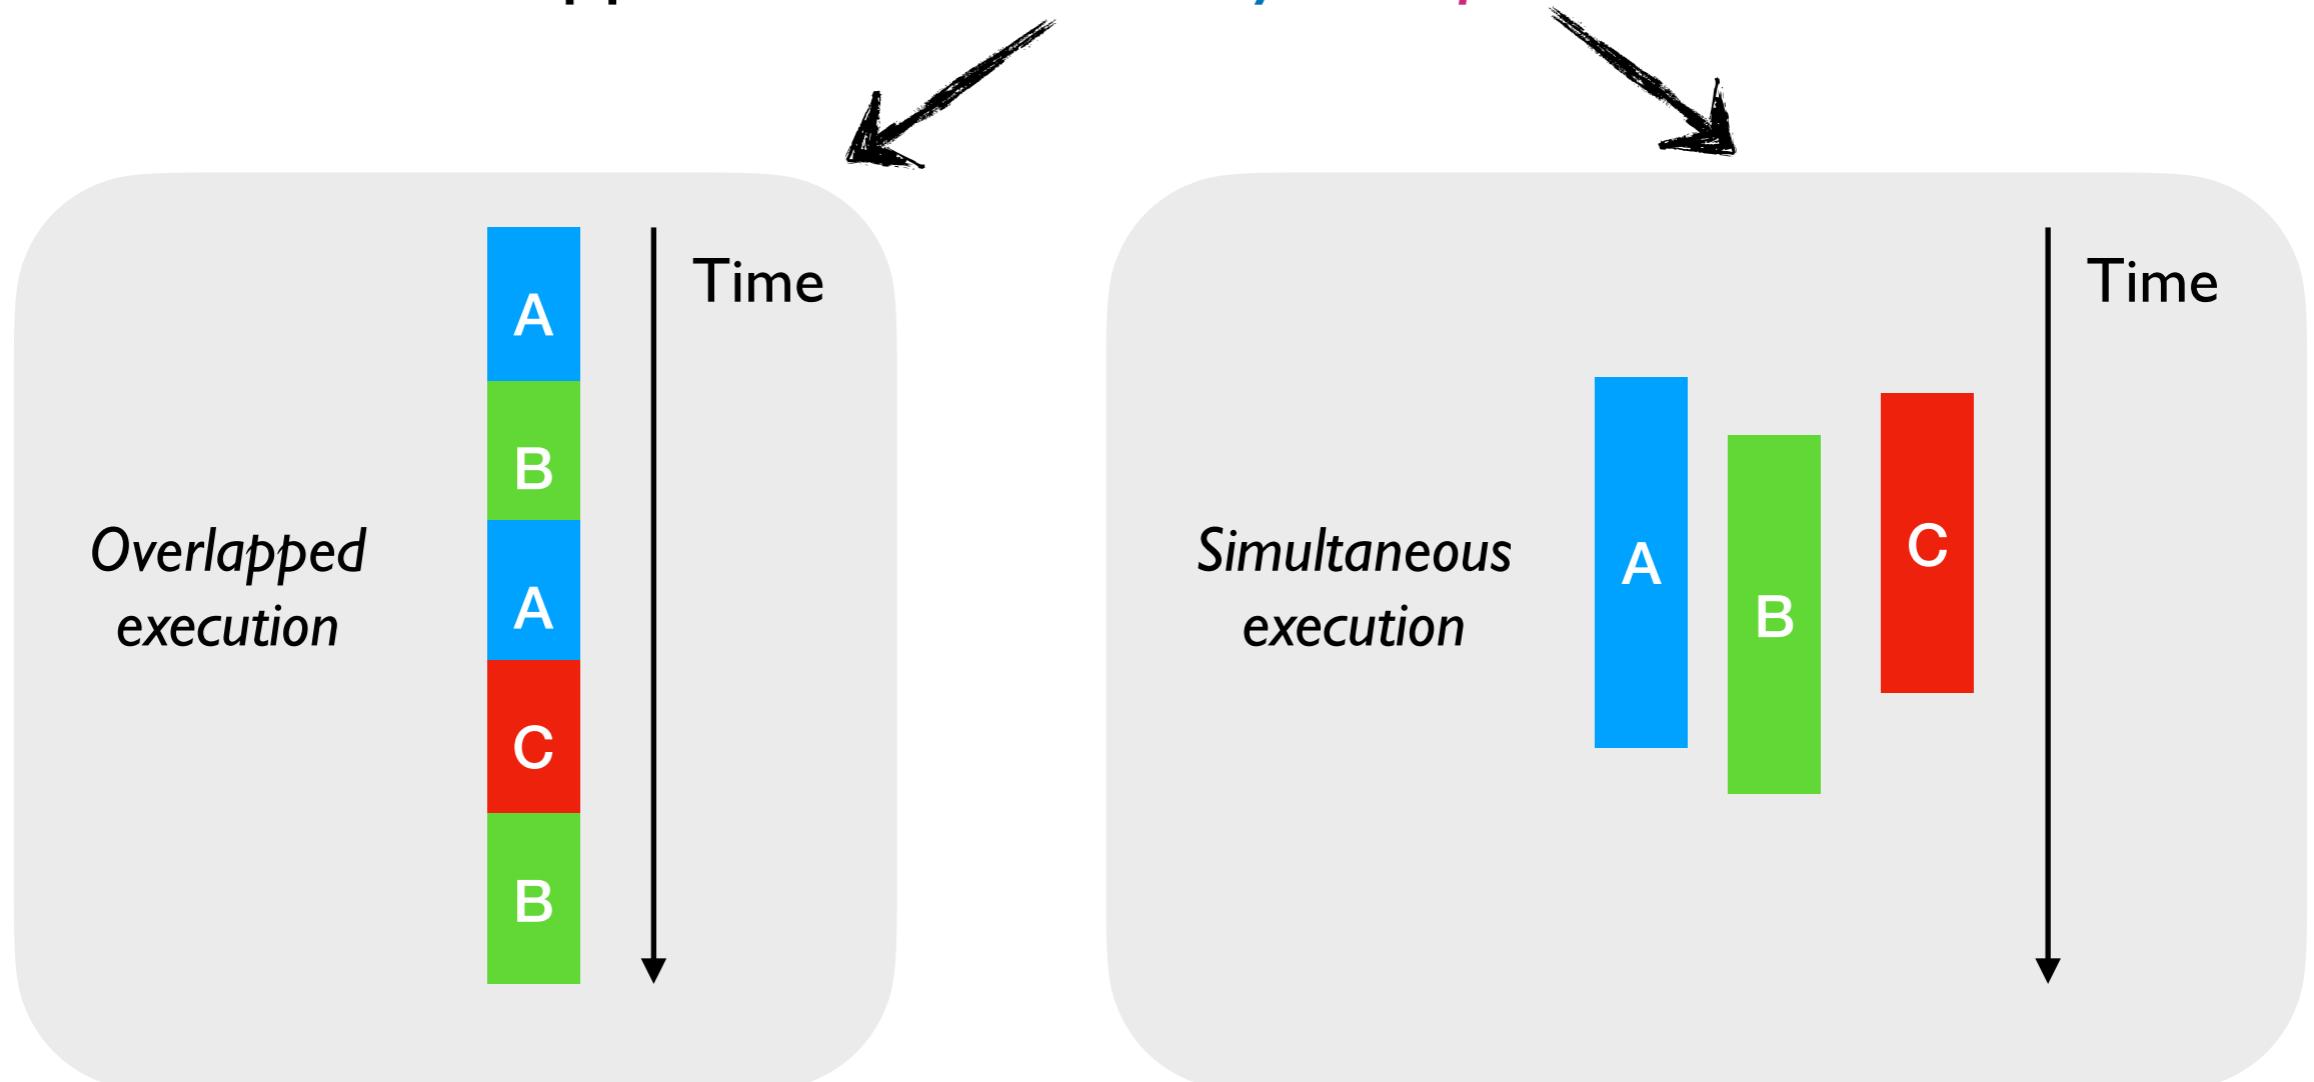
- Adds native support for *concurrency* and *parallelism* to OCaml



Effect Handlers

Multicore OCaml

- Adds native support for *concurrency* and *parallelism* to OCaml



Effect Handlers

Domains

OCaml 5.0

- Domains and *the runtime system support* for effect handlers will land in OCaml 5.0
 - ◆ Expected to be released alongside 4.14 (Q2 2022)
 - ◆ OCaml 4.xx version will have long-term support

OCaml 5.0

- Domains and *the runtime system support* for effect handlers will land in OCaml 5.0
 - ◆ Expected to be released alongside 4.14 (Q2 2022)
 - ◆ OCaml 4.xx version will have long-term support

```
opam install 4.12.0+domains+effects
```

Outline of the talk

1. Challenges of adding parallelism to OCaml
2. Deep dive into the new parallel GC for OCaml
3. High-level parallel programming with *Domainslib*

Challenges

- Millions of lines of legacy code
 - ◆ Written without *concurrency* and *parallelism* in mind
 - ◆ Cost of refactoring sequential code itself is *prohibitive*

Challenges

- Millions of lines of legacy code
 - ◆ Written without *concurrency* and *parallelism* in mind
 - ◆ Cost of refactoring sequential code itself is *prohibitive*
- Low-latency and predictable performance
 - ◆ Great for applications that tolerate $\sim 10ms$ latency

Challenges

- Millions of lines of legacy code
 - ◆ Written without *concurrency* and *parallelism* in mind
 - ◆ Cost of refactoring sequential code itself is *prohibitive*
- Low-latency and predictable performance
 - ◆ Great for applications that tolerate $\sim 10ms$ latency
- Excellent compatibility with debugging and profiling tools
 - ◆ `gdb`, `lldb`, `perf`, `libunwind`, etc.

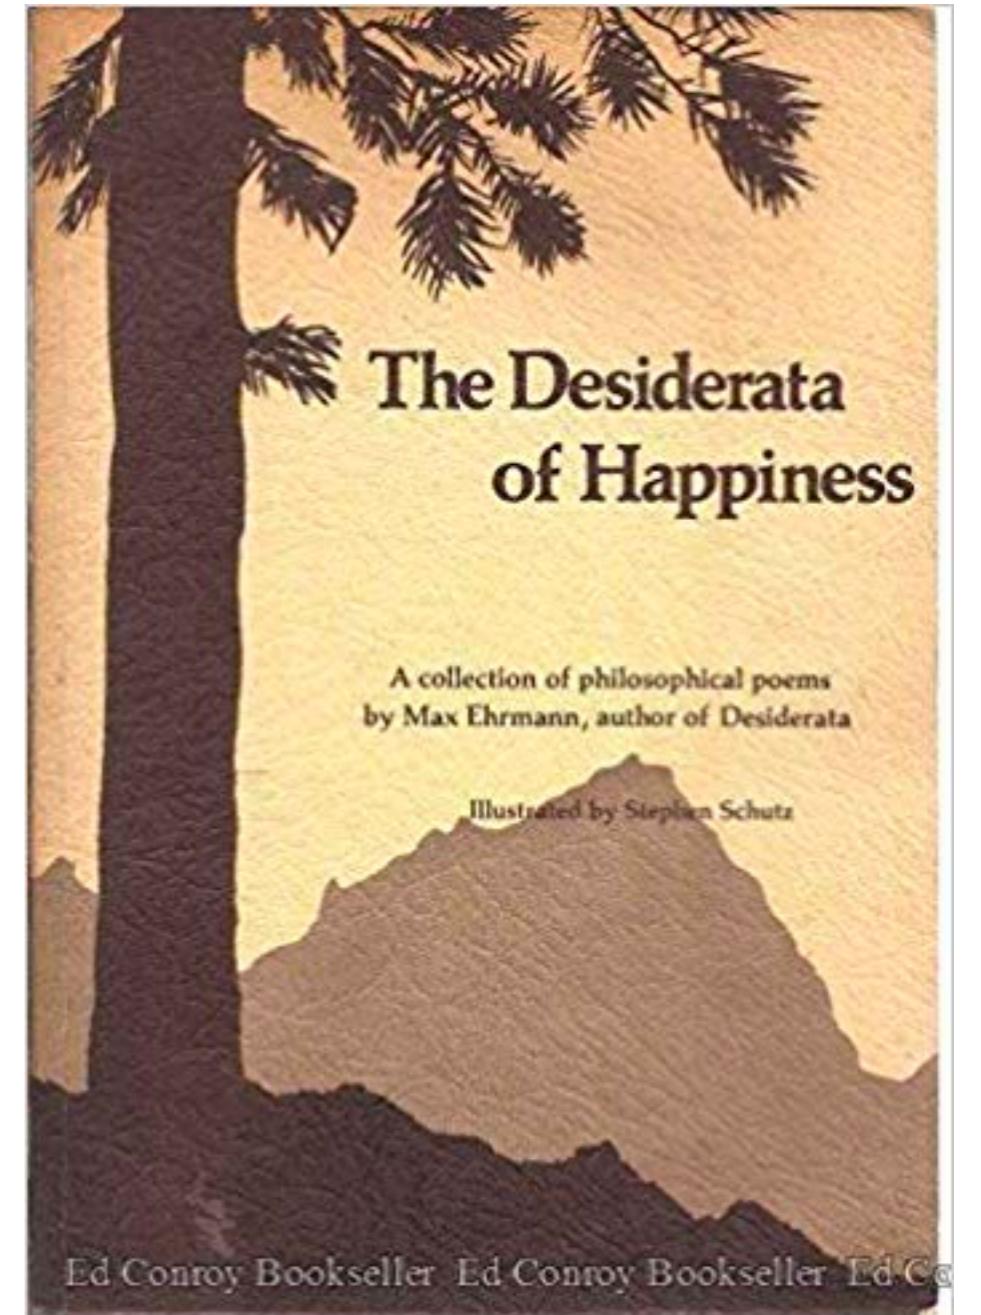
Challenges

- Millions of lines of legacy code
 - ◆ Written without *concurrency* and *parallelism* in mind
 - ◆ Cost of refactoring sequential code itself is *prohibitive*
- Low-latency and predictable performance
 - ◆ Great for applications that tolerate $\sim 10\text{ms}$ latency
- Excellent compatibility with debugging and profiling tools
 - ◆ `gdb`, `lldb`, `perf`, `libunwind`, etc.

Backwards compatibility before scalability

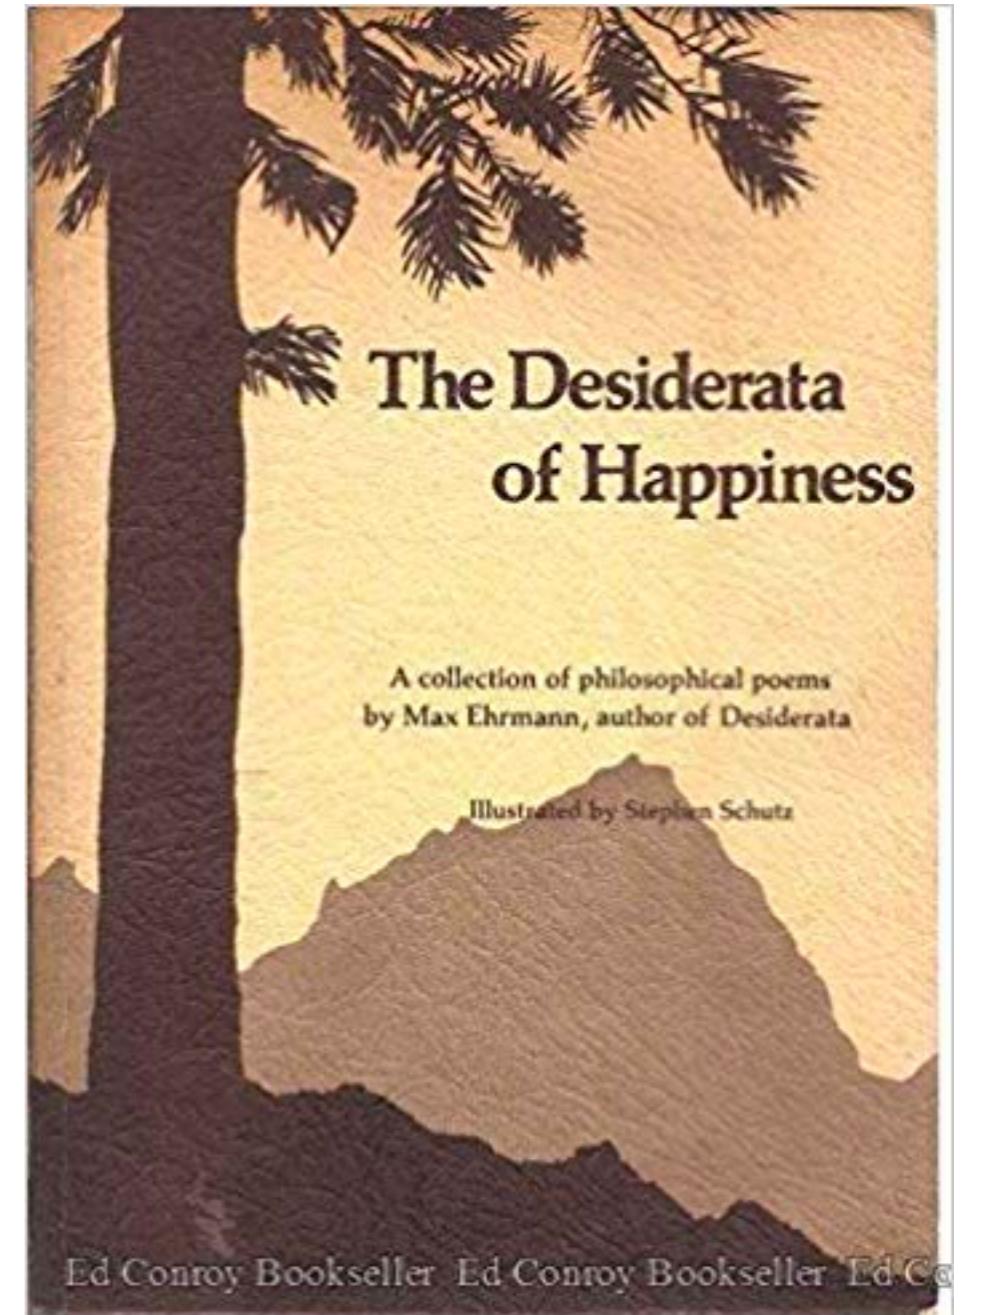
Desiderata

- Feature backwards compatibility
 - ◆ Do not break existing code



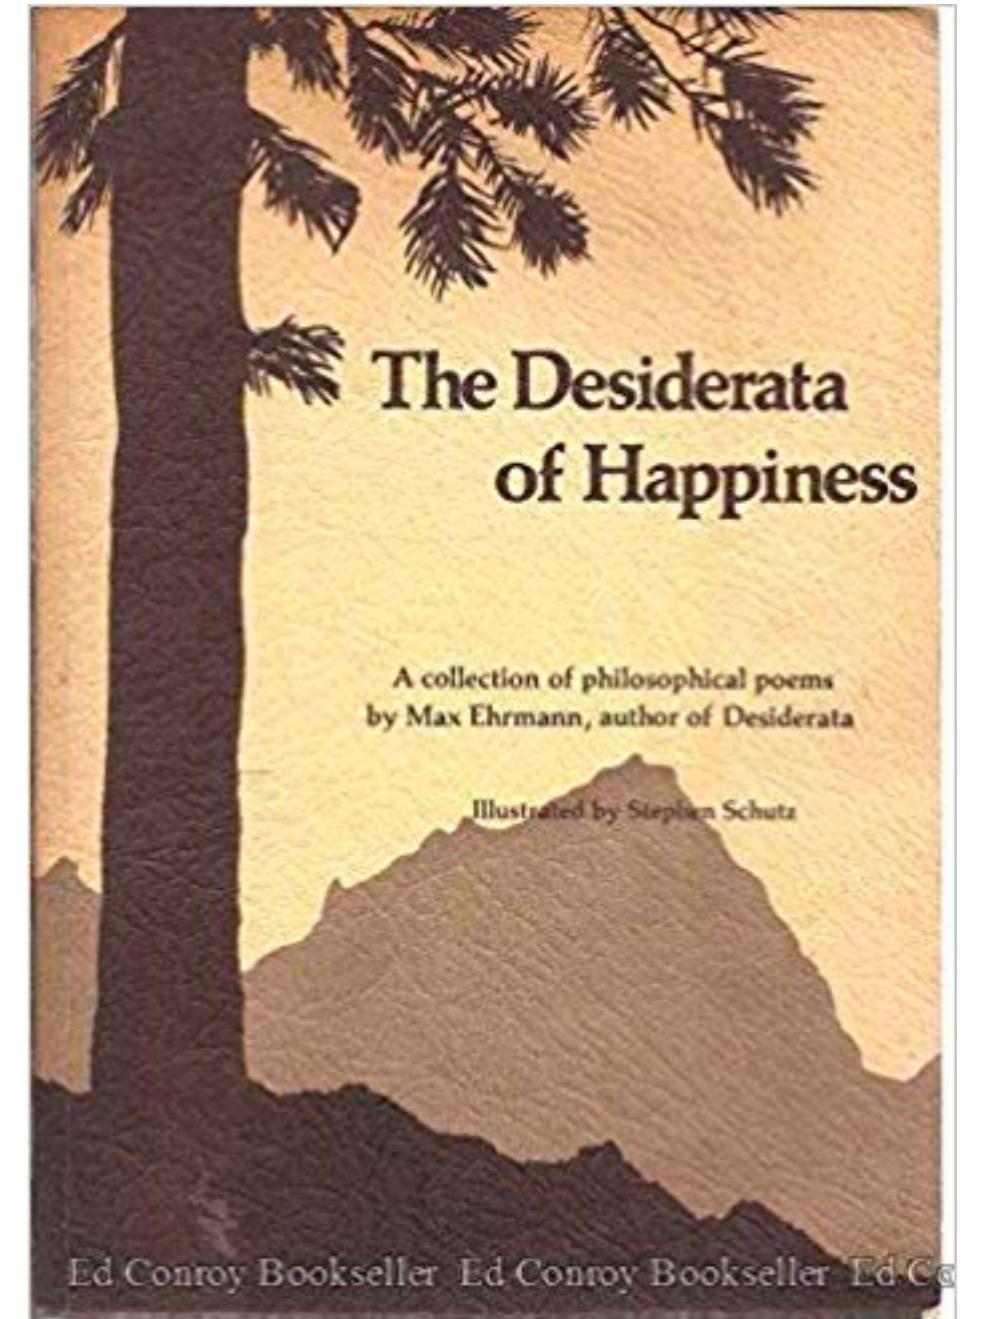
Desiderata

- Feature backwards compatibility
 - ◆ Do not break existing code
- Performance backwards compatibility
 - ◆ Existing programs run just as fast using just the same memory



Desiderata

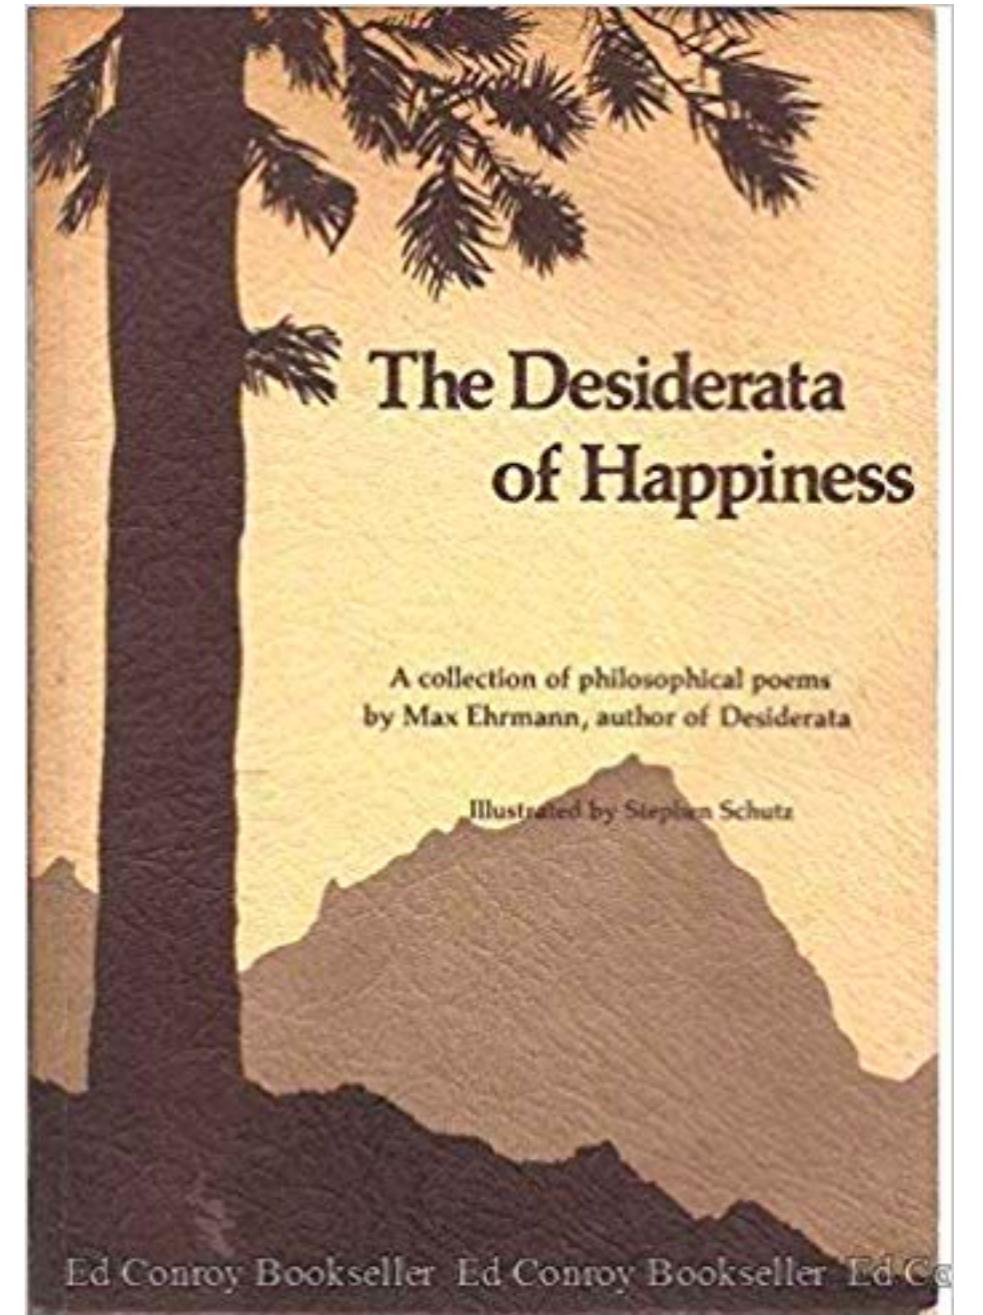
- Feature backwards compatibility
 - ◆ Do not break existing code
- Performance backwards compatibility
 - ◆ Existing programs run just as fast using just the same memory
- GC Latency before multicore scalability



Ed Conroy Bookseller Ed Conroy Bookseller Ed C

Desiderata

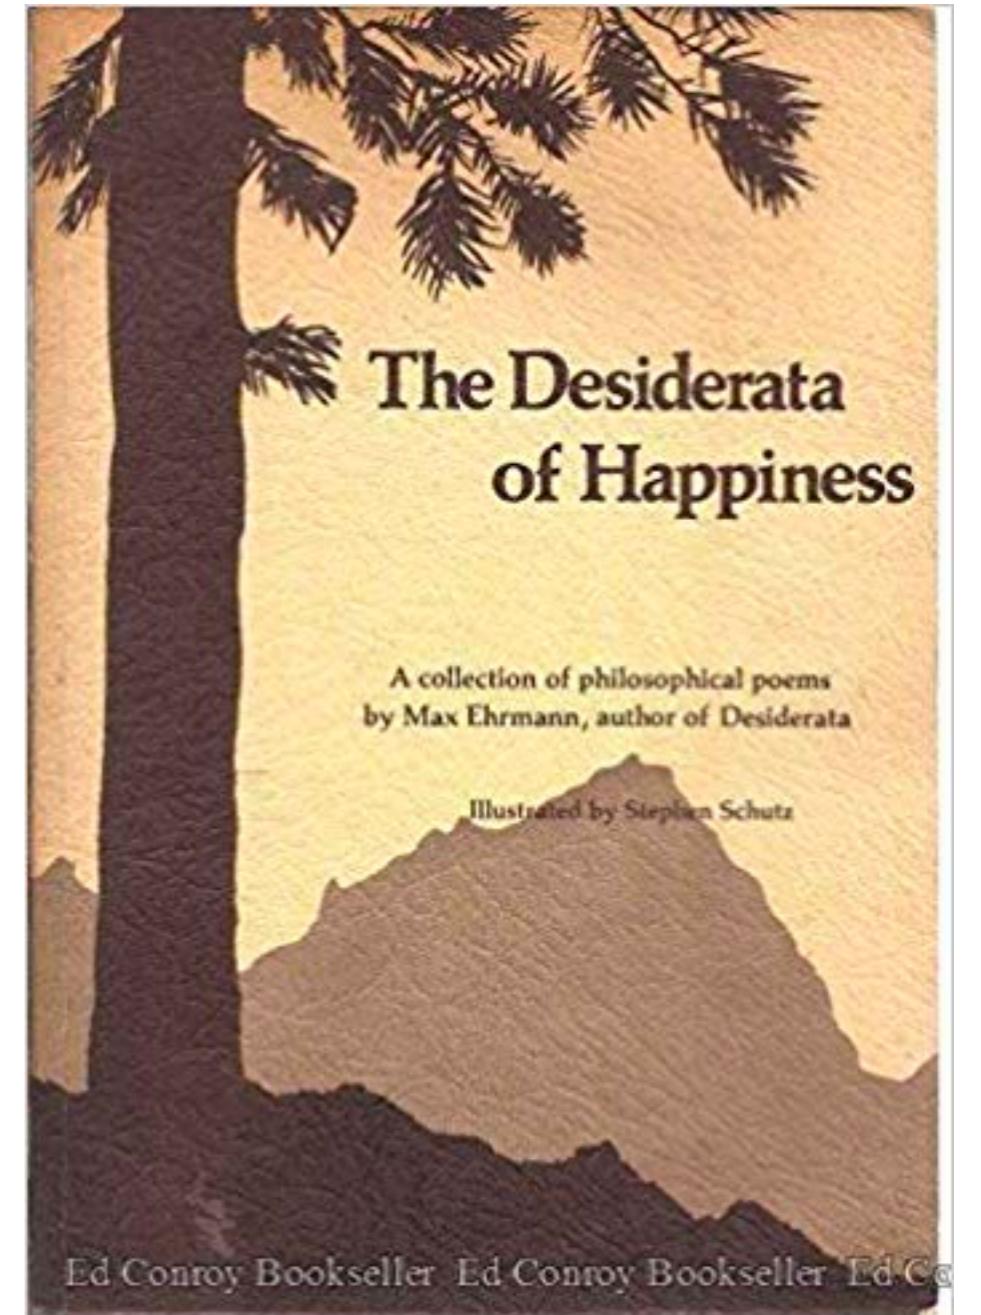
- Feature backwards compatibility
 - ◆ Do not break existing code
- Performance backwards compatibility
 - ◆ Existing programs run just as fast using just the same memory
- GC Latency before multicore scalability
- Compatibility with program inspection tools



Ed Conroy Bookseller Ed Conroy Bookseller Ed C

Desiderata

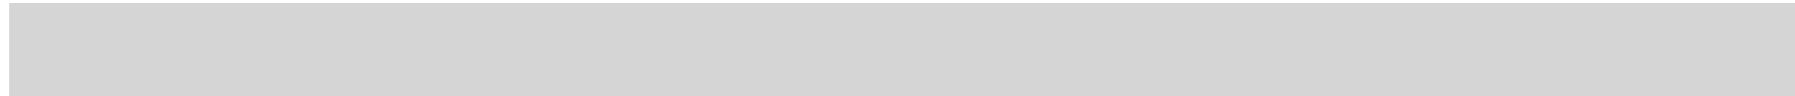
- Feature backwards compatibility
 - ◆ Do not break existing code
- Performance backwards compatibility
 - ◆ Existing programs run just as fast using just the same memory
- GC Latency before multicore scalability
- Compatibility with program inspection tools
- Performant concurrent and parallel programming abstractions



Domains for Parallelism

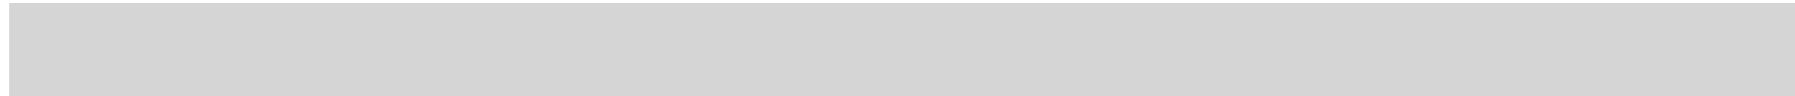
Domains for Parallelism

- A unit of parallelism



Domains for Parallelism

- A unit of parallelism
- **Heavyweight** — maps onto a OS thread
 - ◆ Recommended to have 1 domain per core



Domains for Parallelism

- A unit of parallelism
- **Heavyweight** — maps onto a OS thread
 - ◆ Recommended to have 1 domain per core
- Low-level domain API
 - ◆ Spawn & join, wait & notify

```
Domain.spawn : (unit -> 'a) -> 'a Domain.t
```

Domains for Parallelism

- A unit of parallelism
- **Heavyweight** — maps onto a OS thread
 - ◆ Recommended to have 1 domain per core
- Low-level domain API
 - ◆ Spawn & join, wait & notify

```
Domain.spawn : (unit -> 'a) -> 'a Domain.t
```

- ◆ Domain-local storage
- ◆ Atomic memory operations
- ◆ Dolan et al, “Bounding Data Races in Space and Time”, PLDI’18

Domains for Parallelism

- A unit of parallelism
- **Heavyweight** — maps onto a OS thread

- **Heavyweight** — maps onto a OS thread
 - ◆ Recommended to have 1 domain per core

- Low-level domain API

- Low-level domain API
 - ◆ Spawn & join, wait & notify

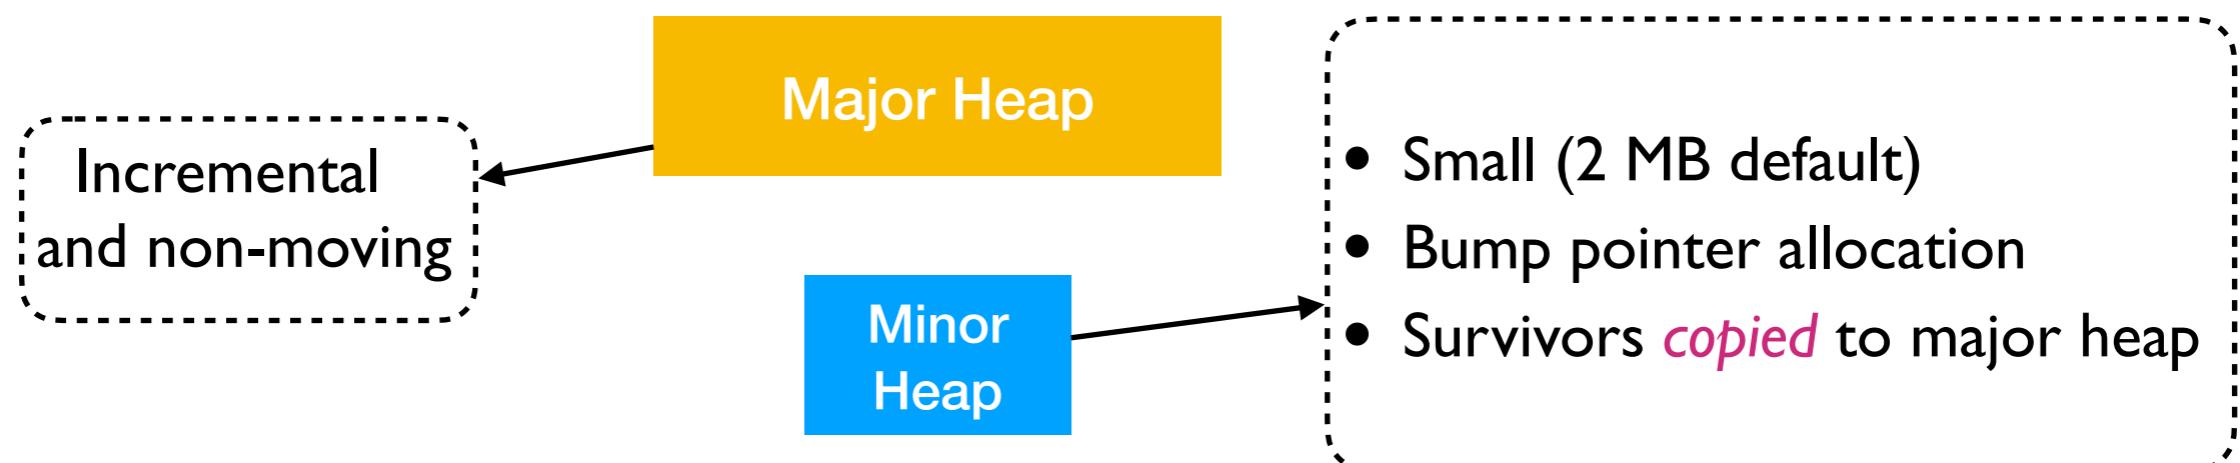
```
Domain.spawn : (unit -> 'a) -> 'a Domain.t
```

- Low-level domain API
 - ◆ Domain-local storage
 - ◆ Atomic memory operations
 - ✿ Dolan et al, “Bounding Data Races in Space and Time”, PLDI’18

- No restrictions on sharing objects between domains
 - ◆ But how does it work?

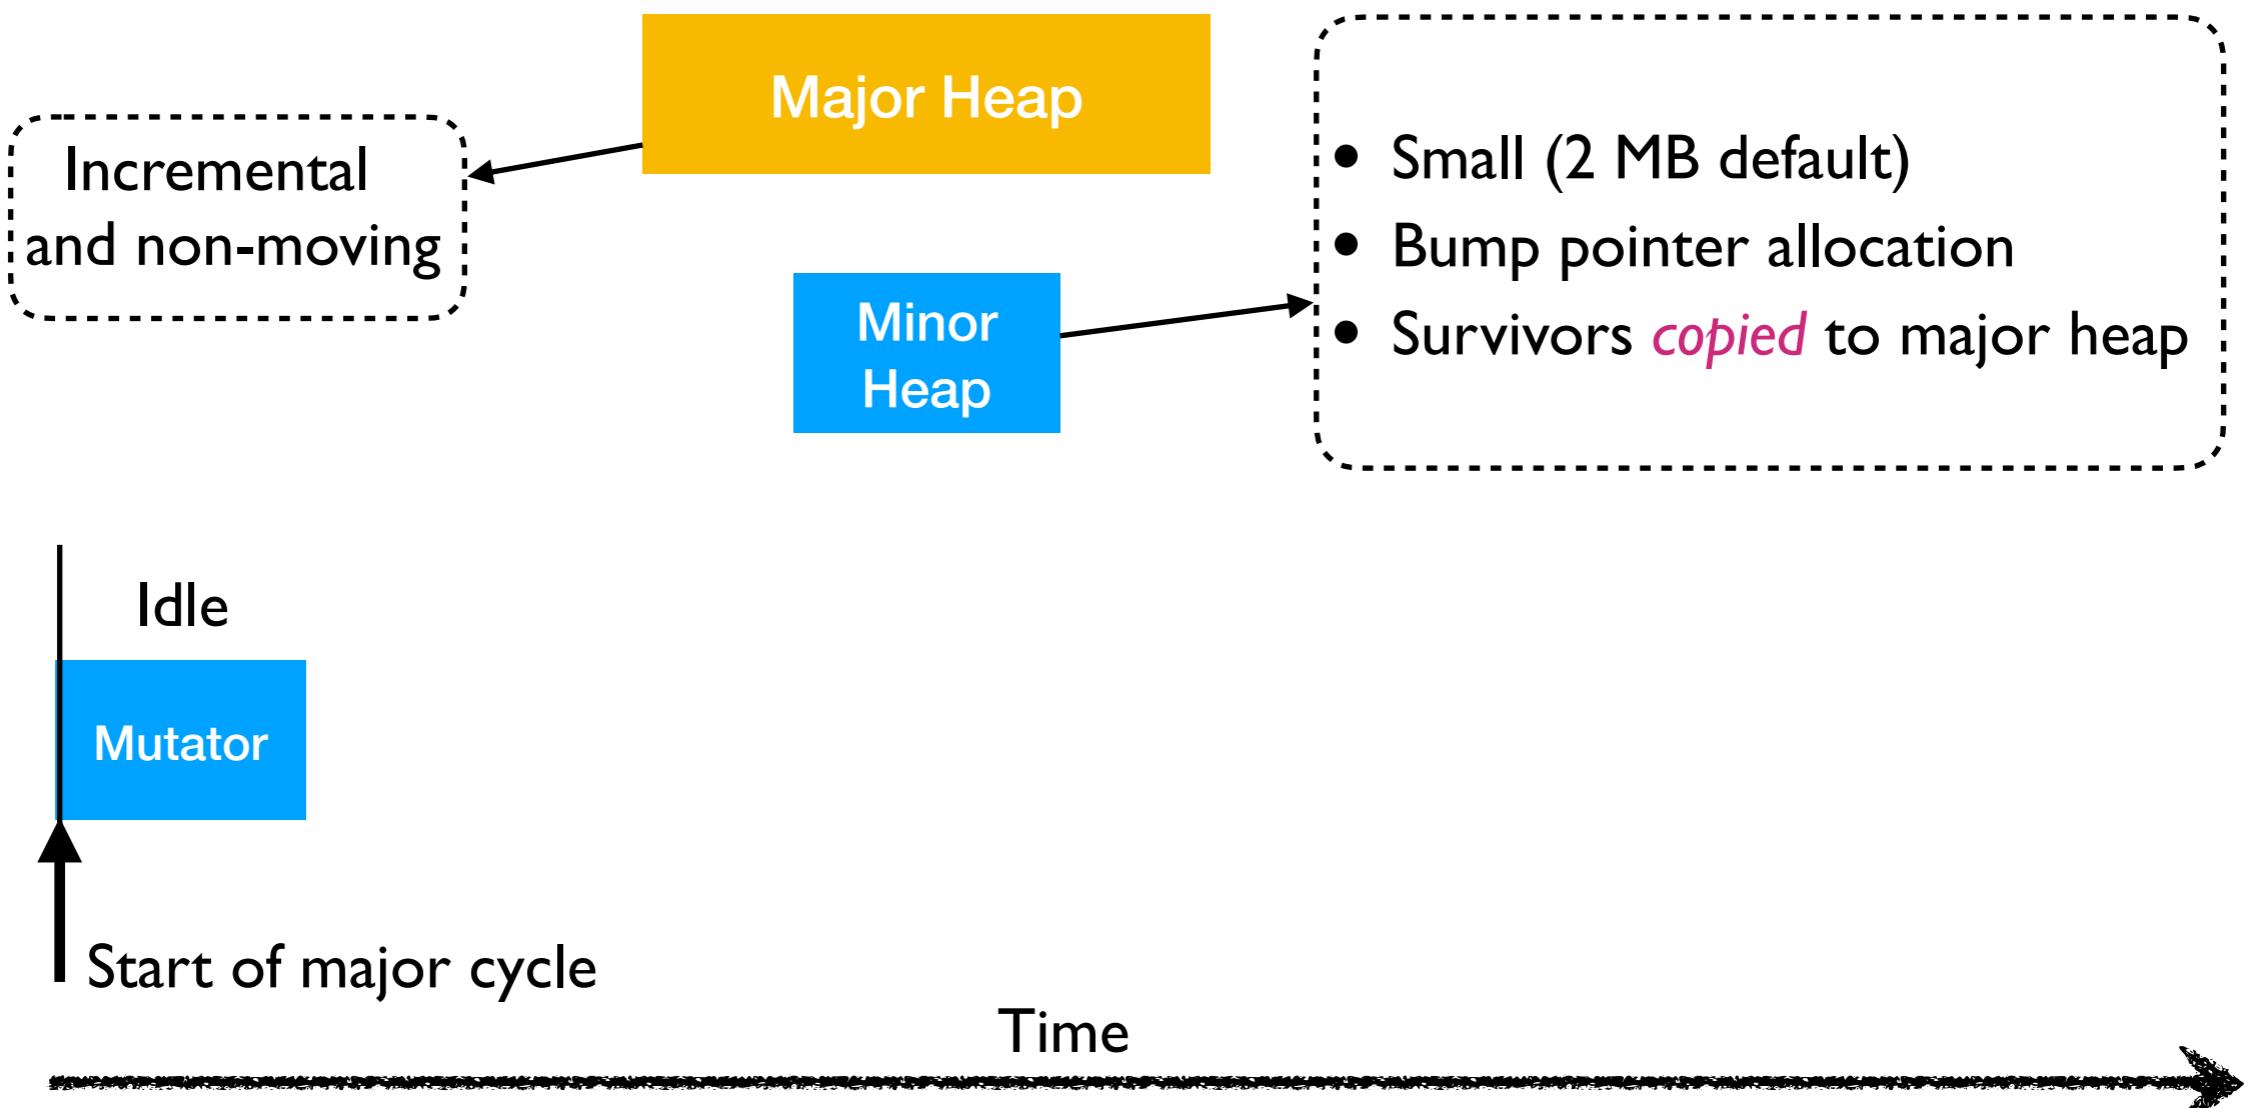
Stock OCaml GC

- A generational, non-moving, incremental, mark-and-sweep GC



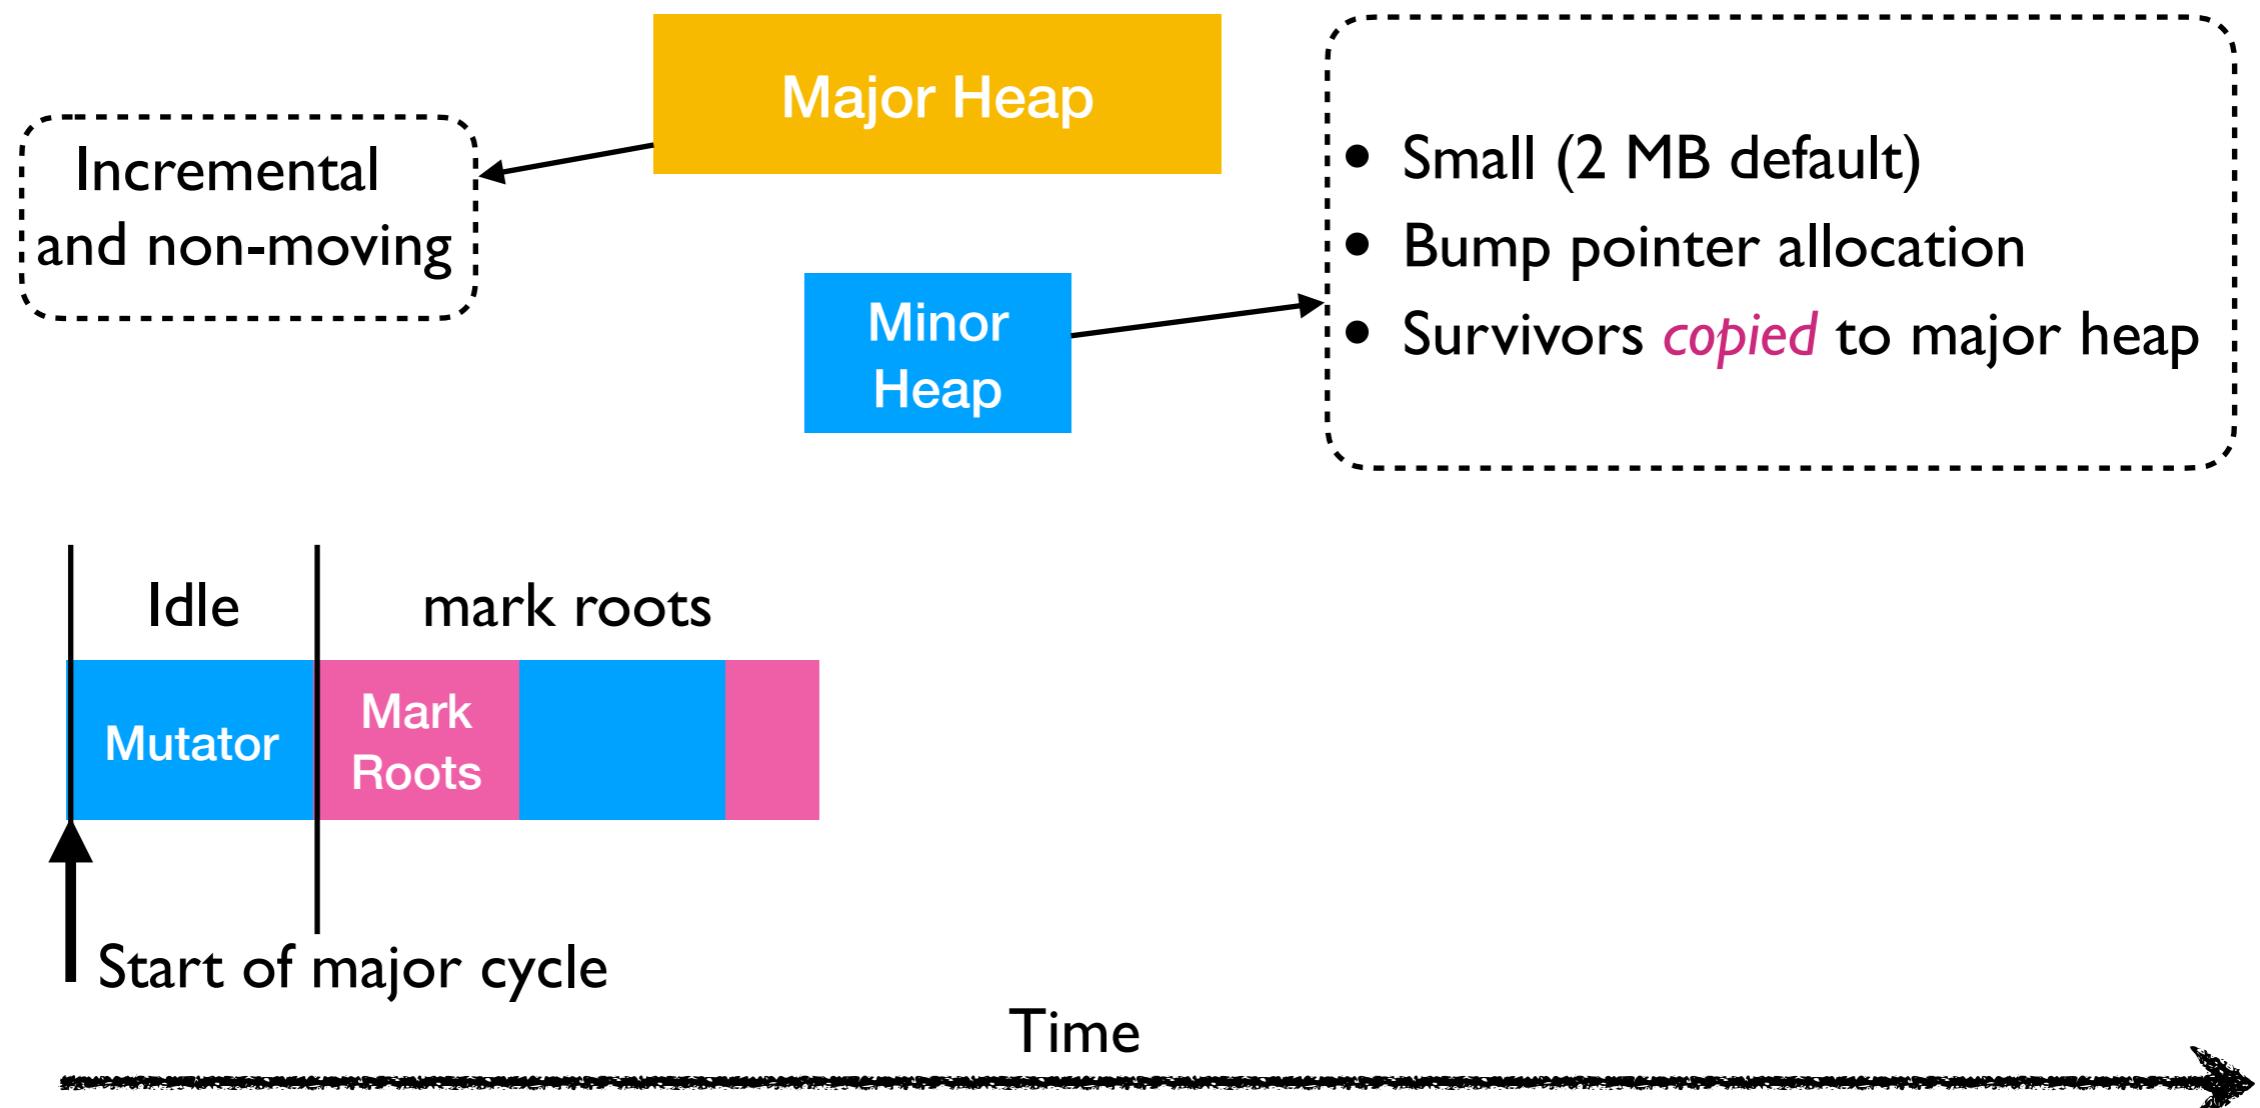
Stock OCaml GC

- A generational, non-moving, incremental, mark-and-sweep GC



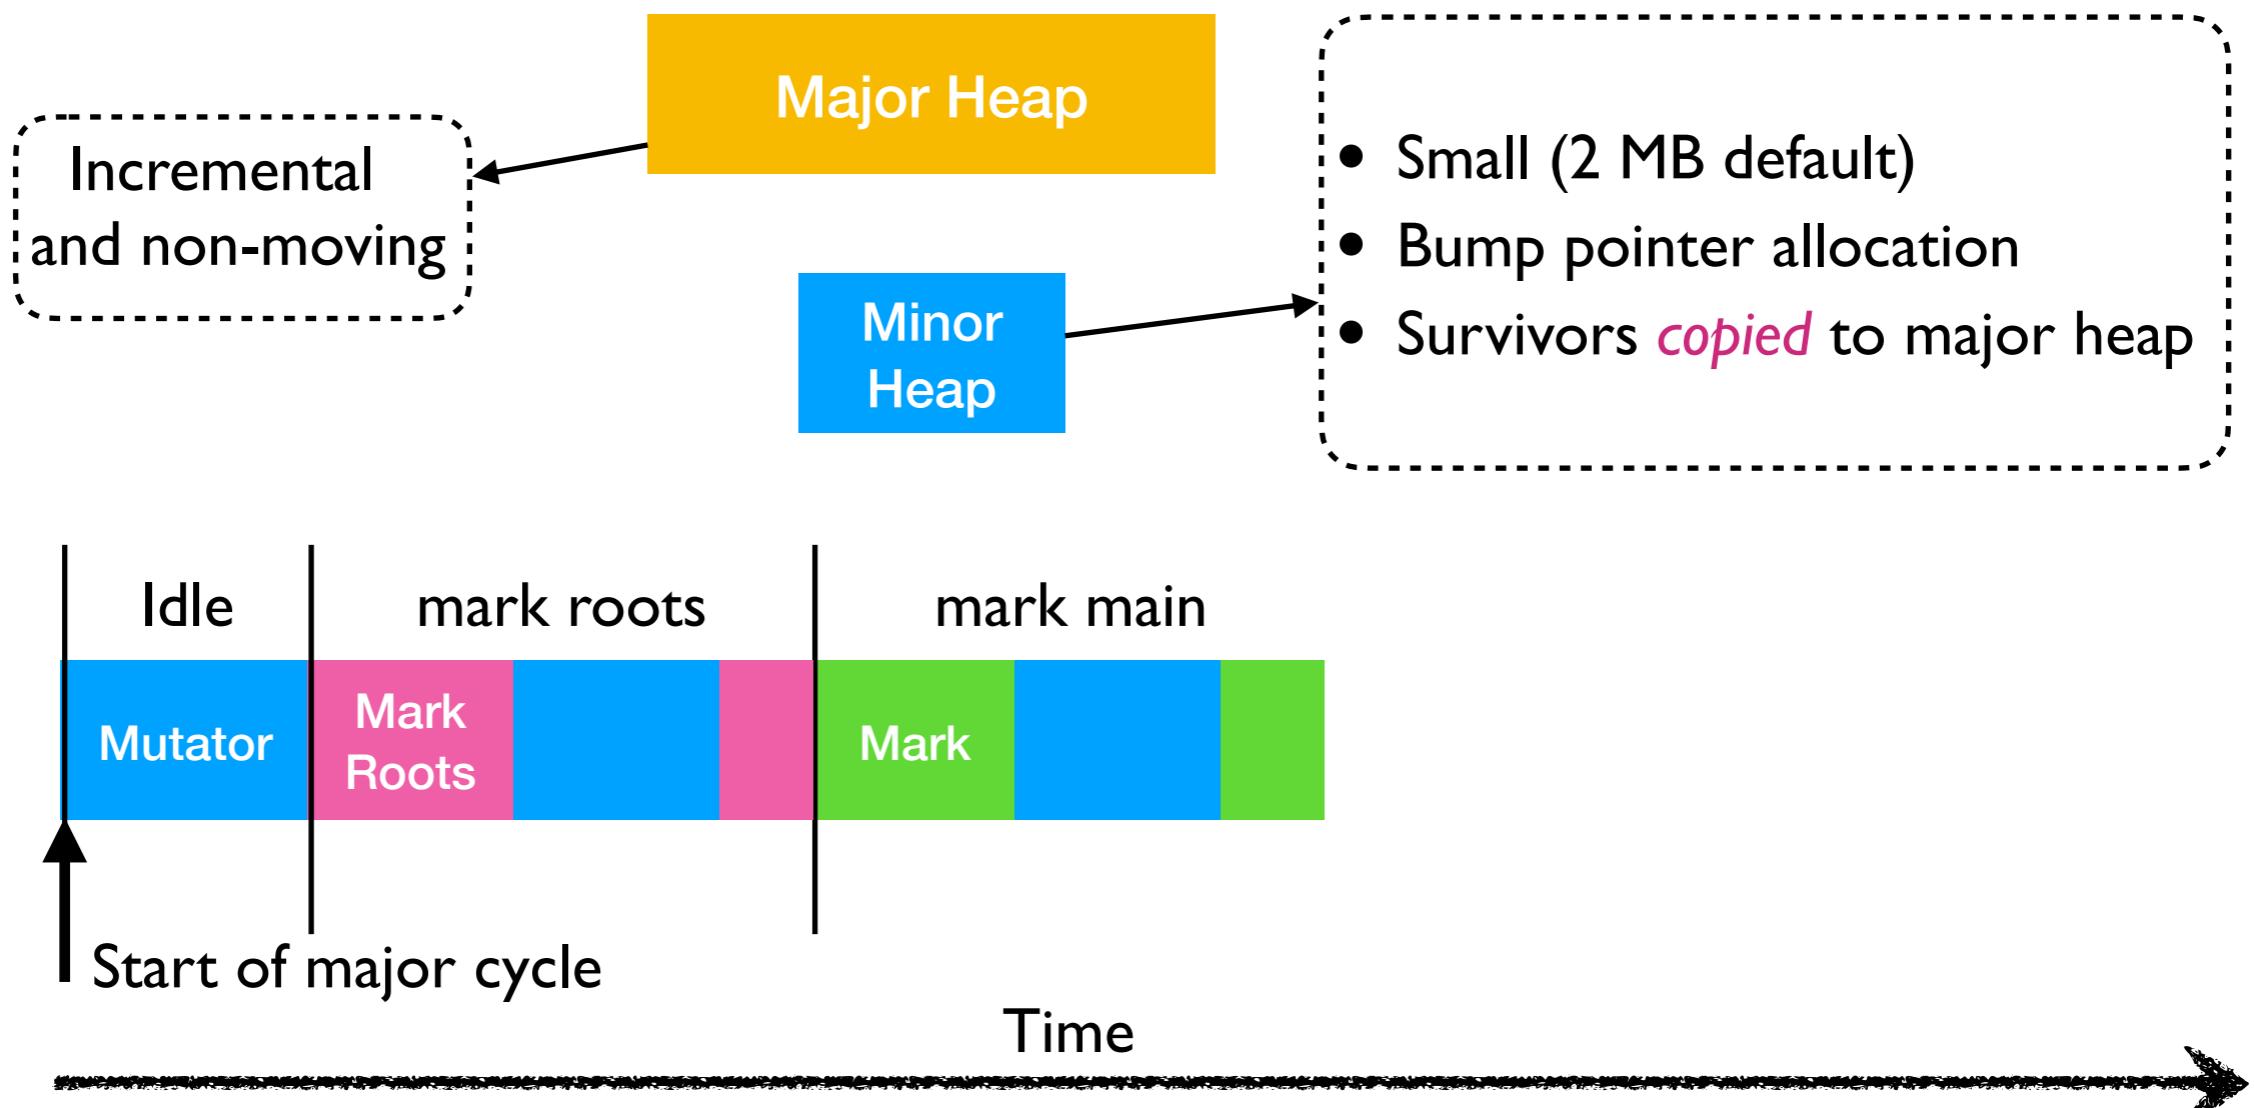
Stock OCaml GC

- A generational, non-moving, incremental, mark-and-sweep GC



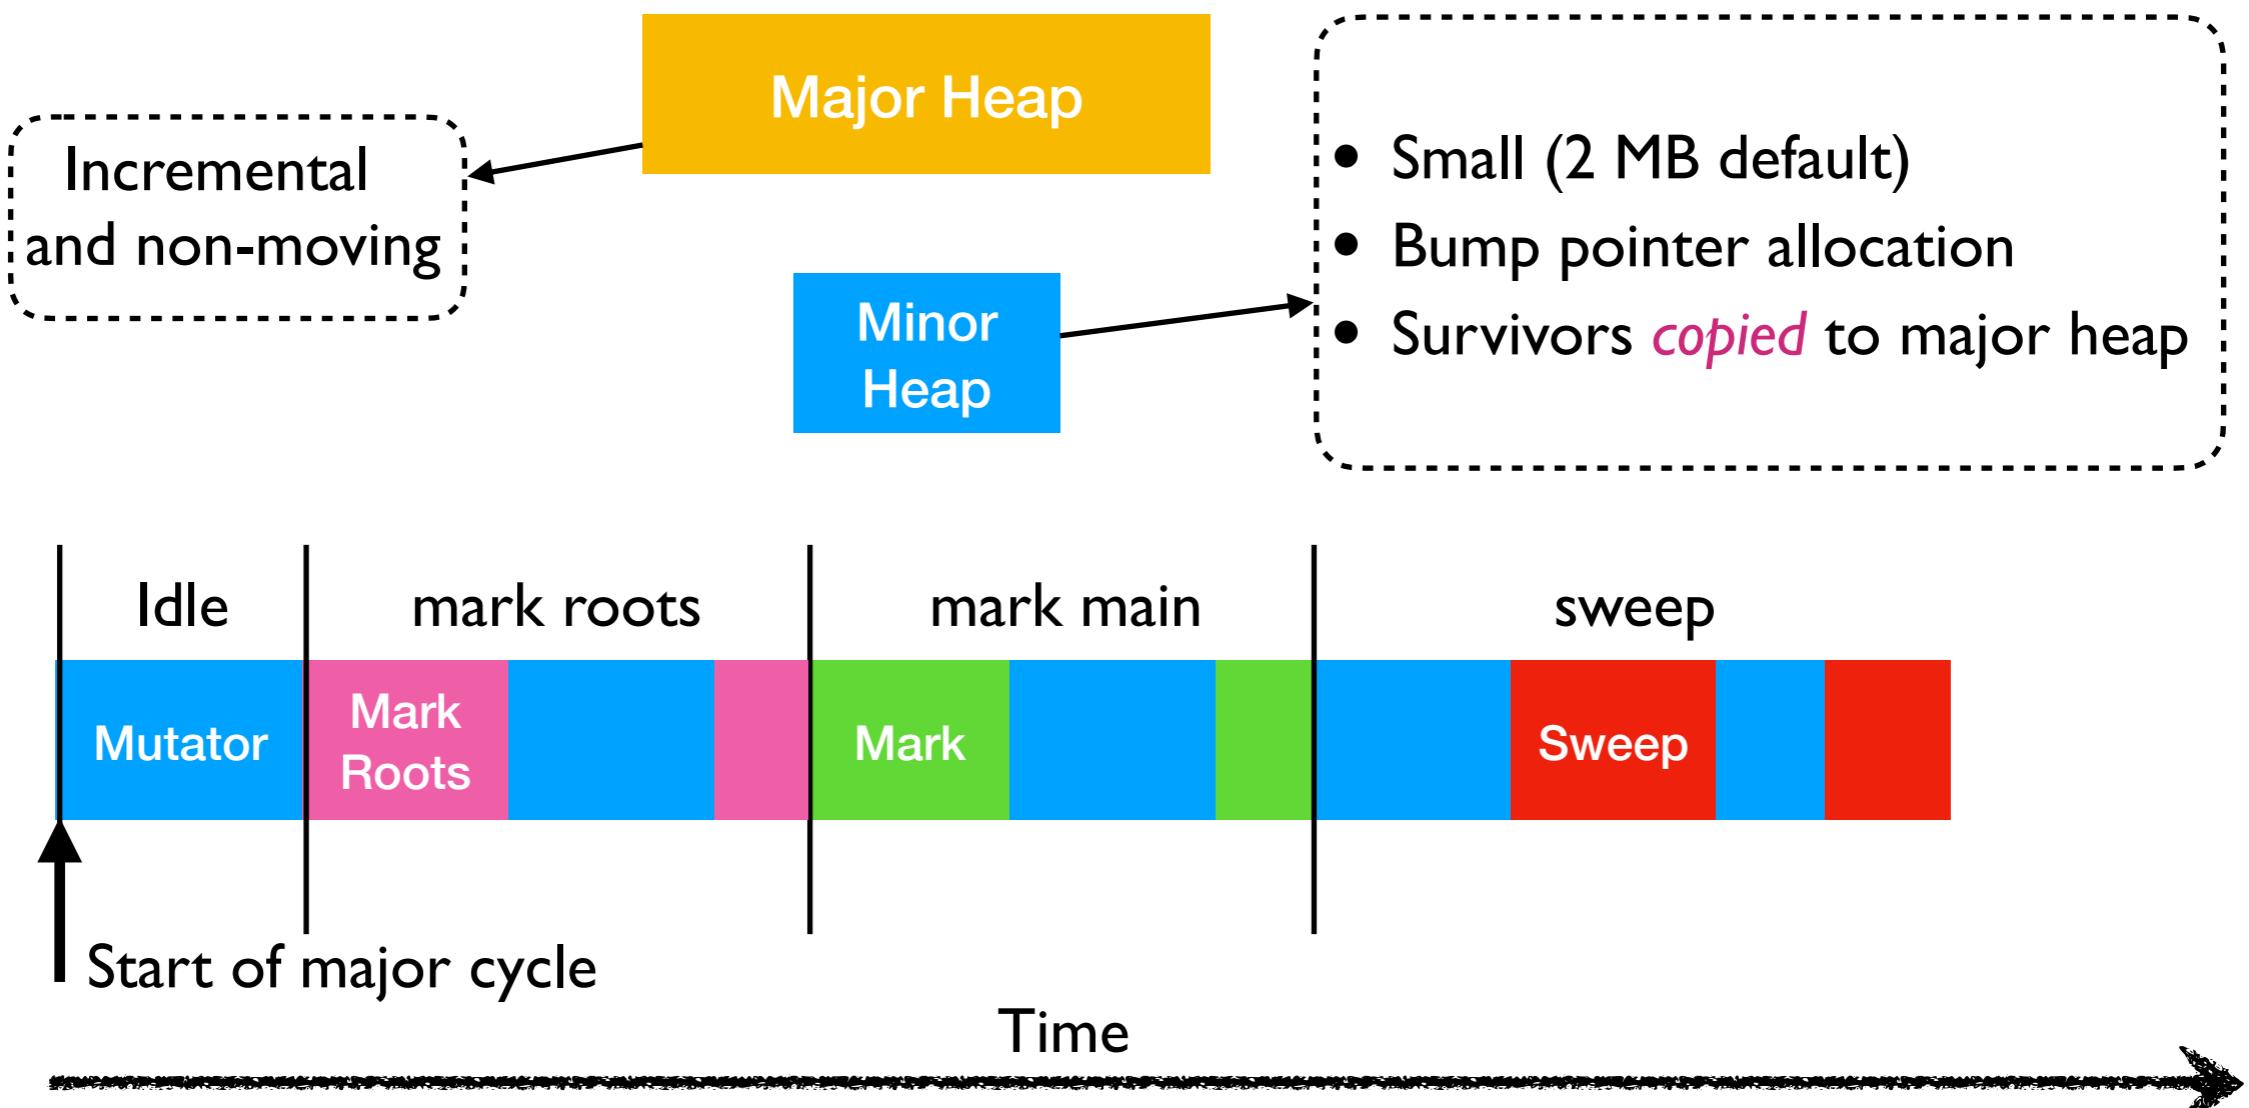
Stock OCaml GC

- A generational, non-moving, incremental, mark-and-sweep GC



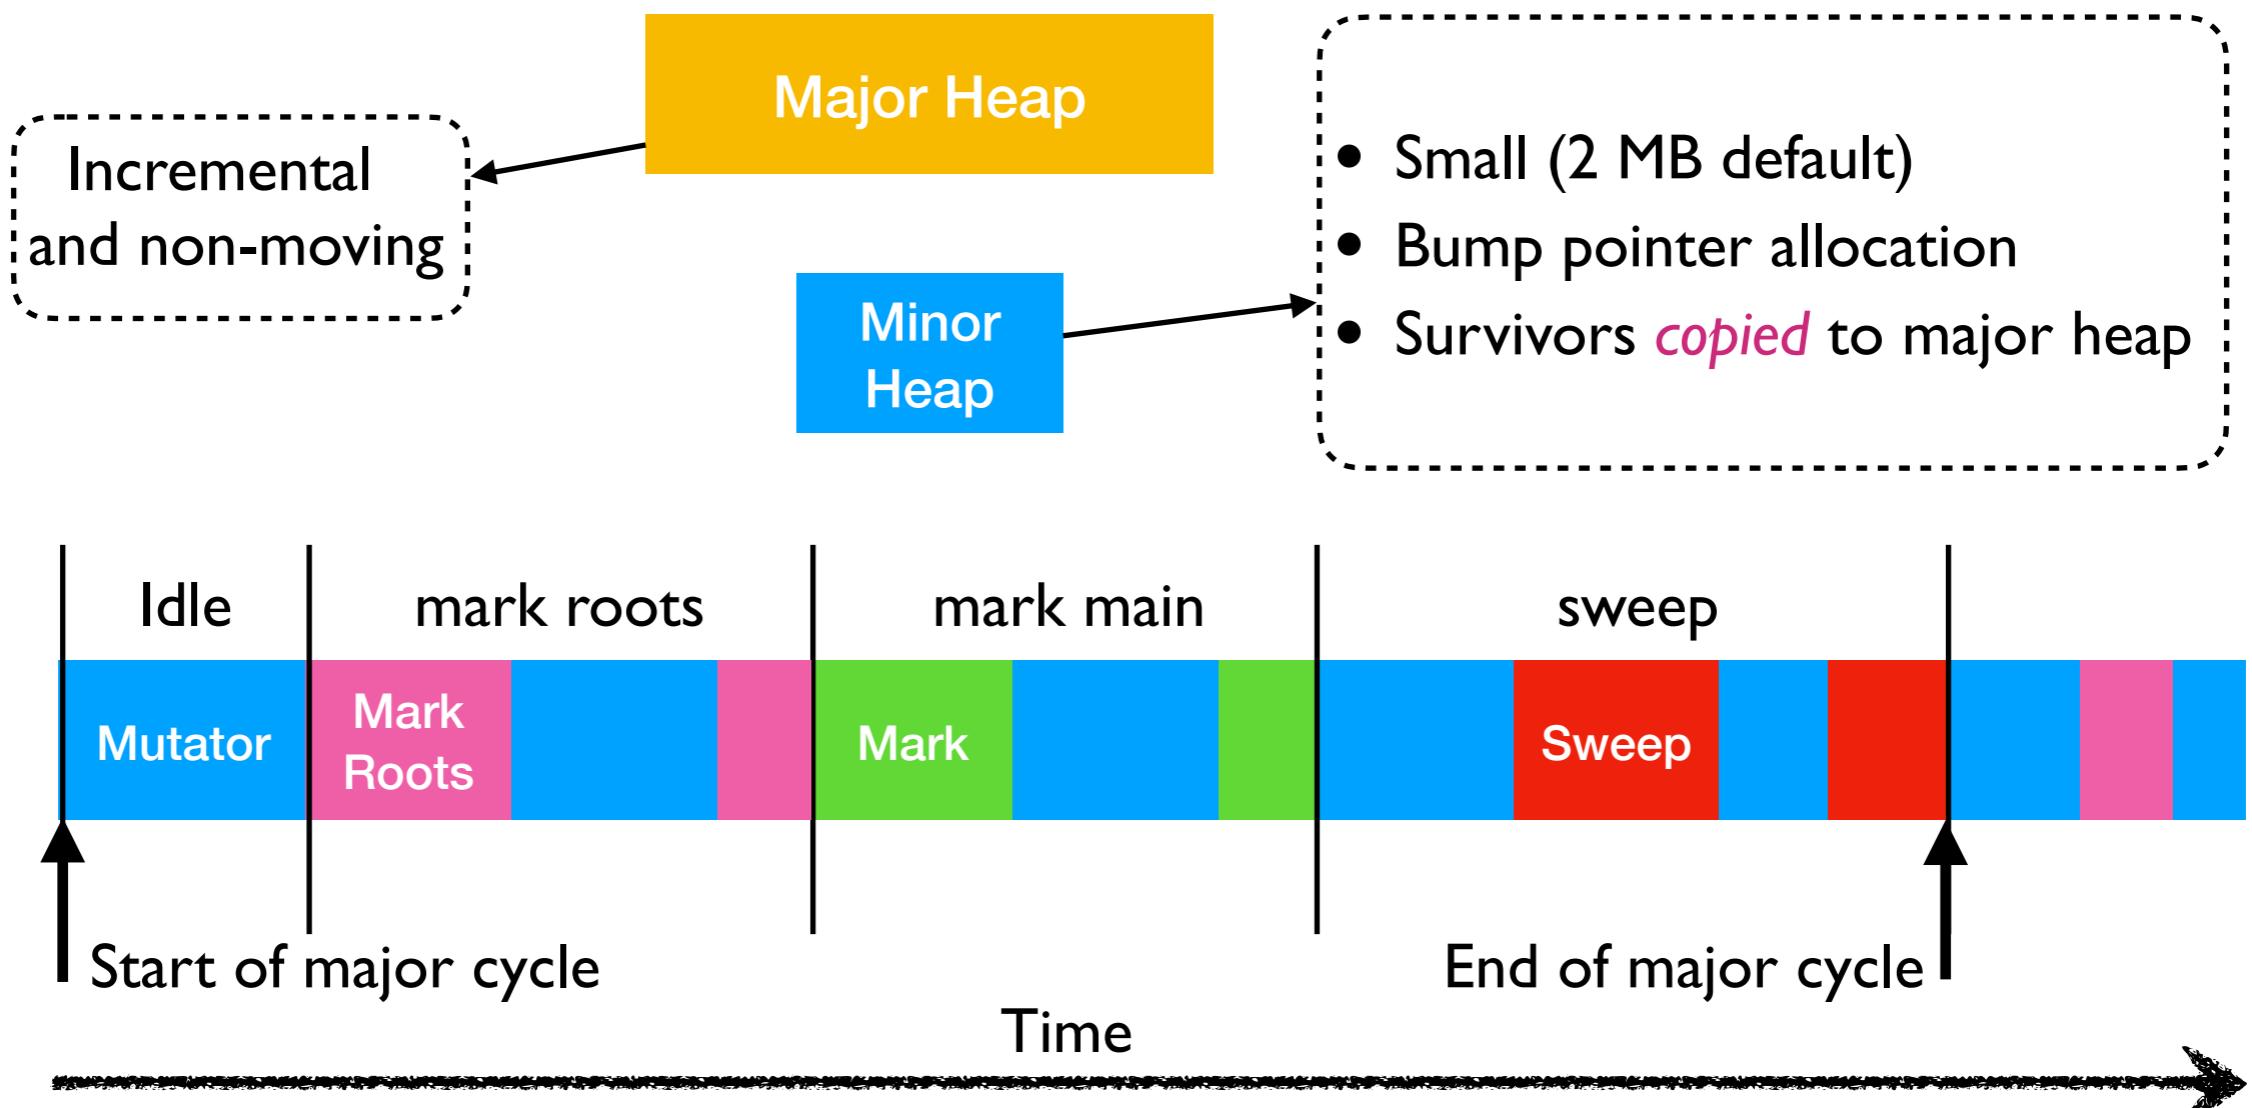
Stock OCaml GC

- A generational, non-moving, incremental, mark-and-sweep GC



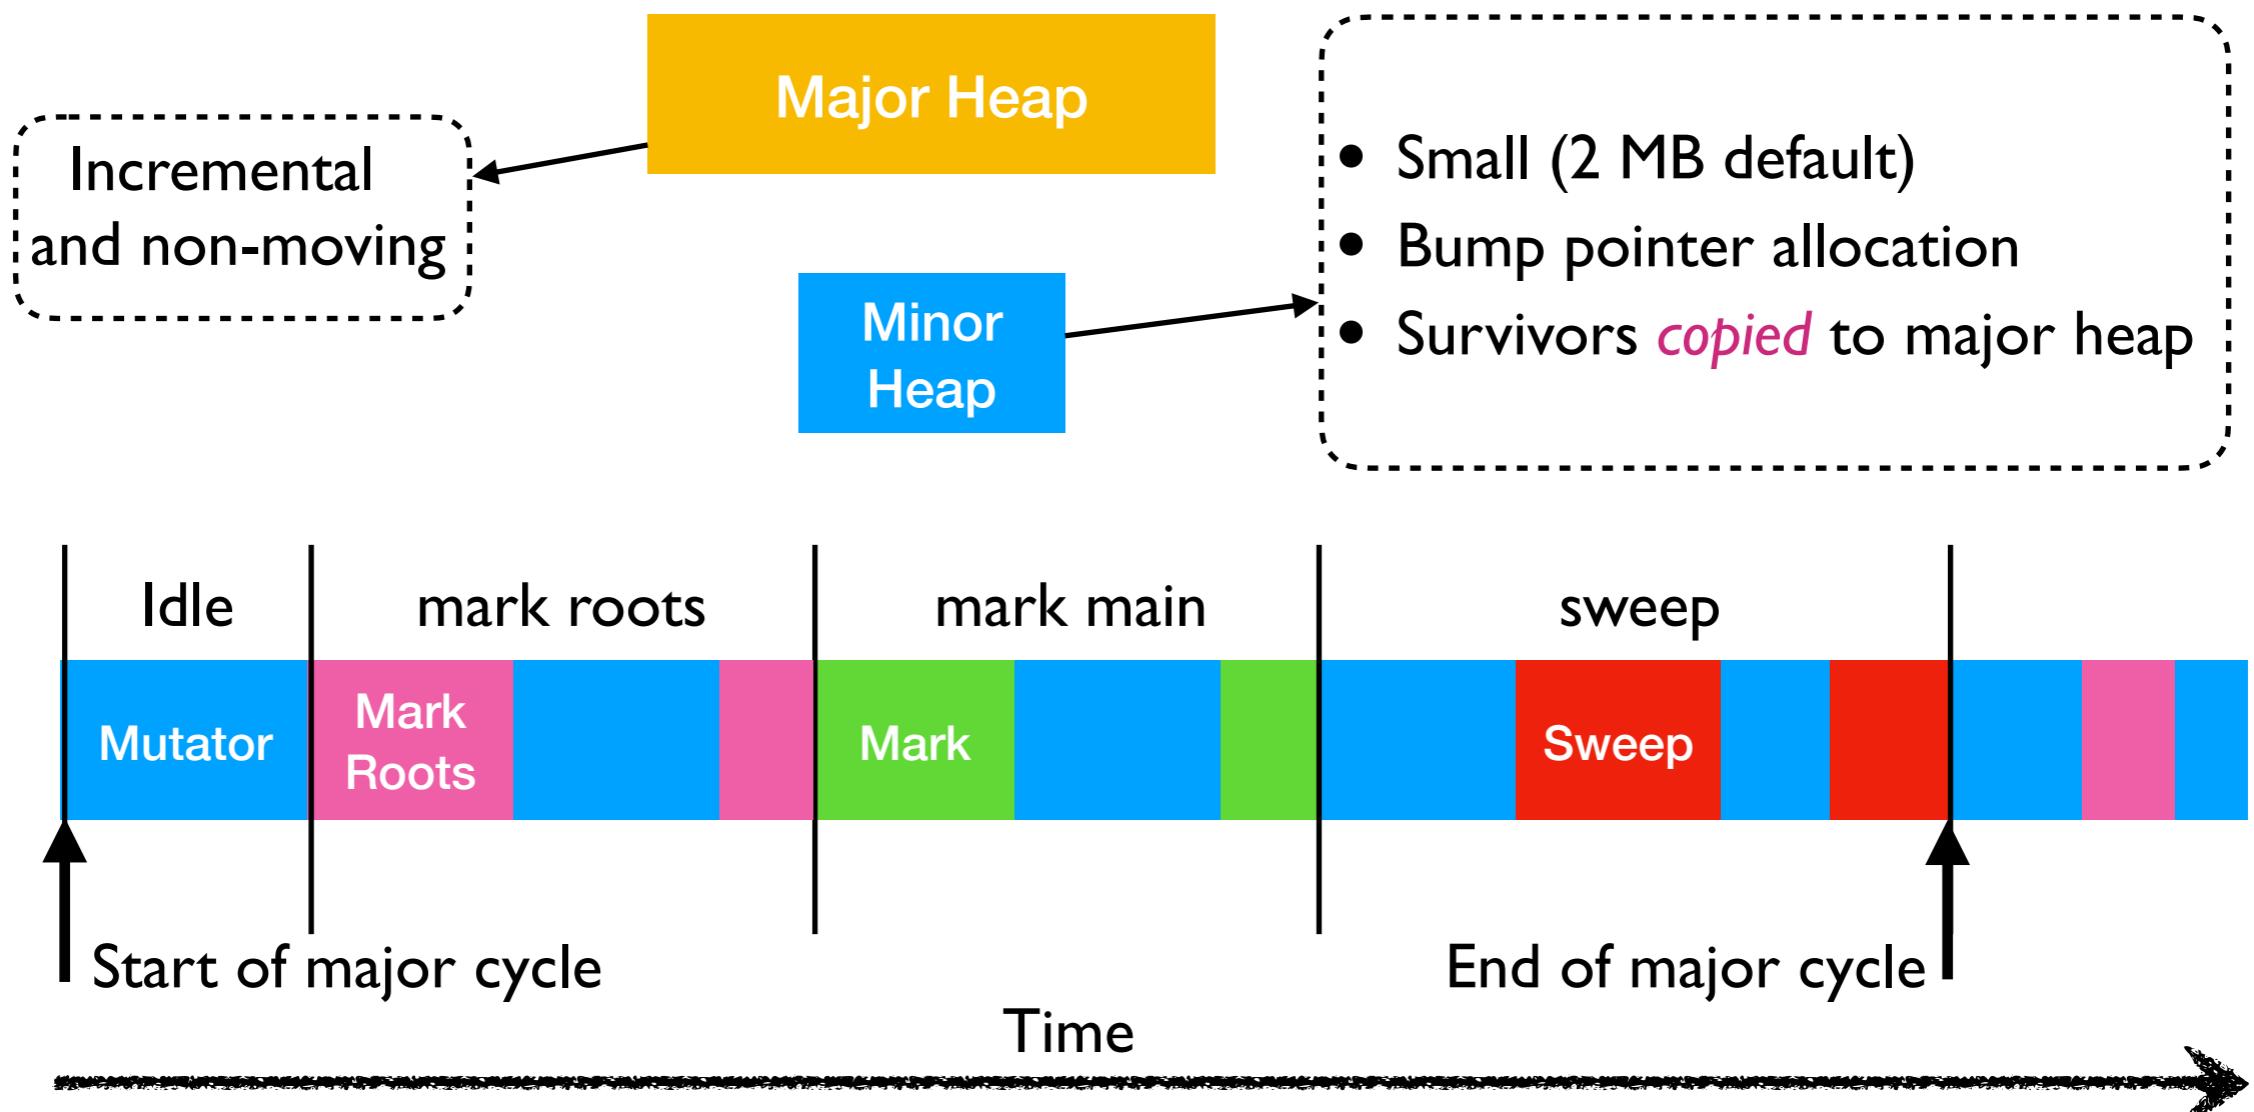
Stock OCaml GC

- A generational, non-moving, incremental, mark-and-sweep GC



Stock OCaml GC

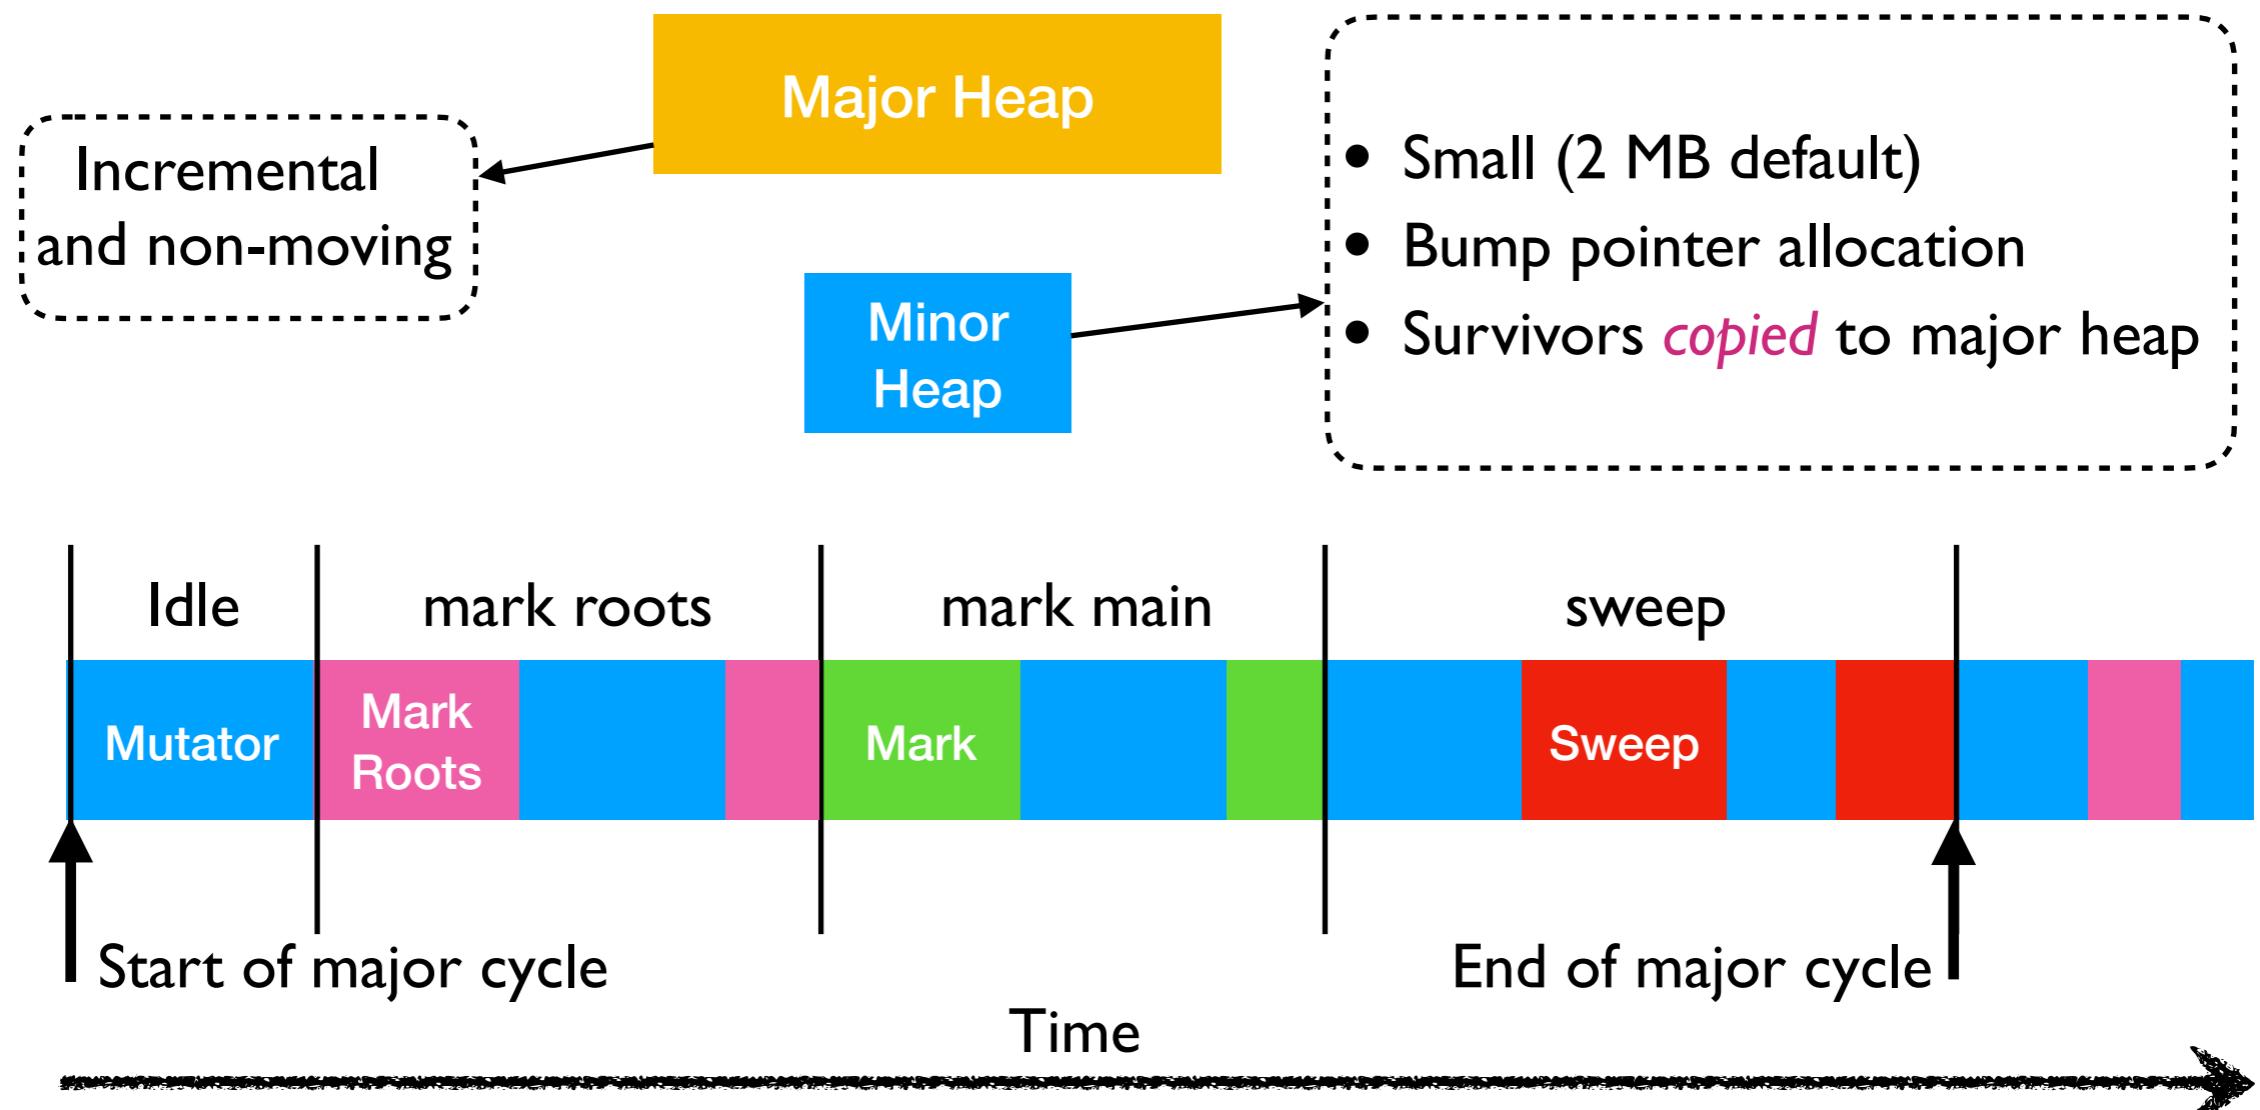
- A generational, non-moving, incremental, mark-and-sweep GC



- Fast allocations

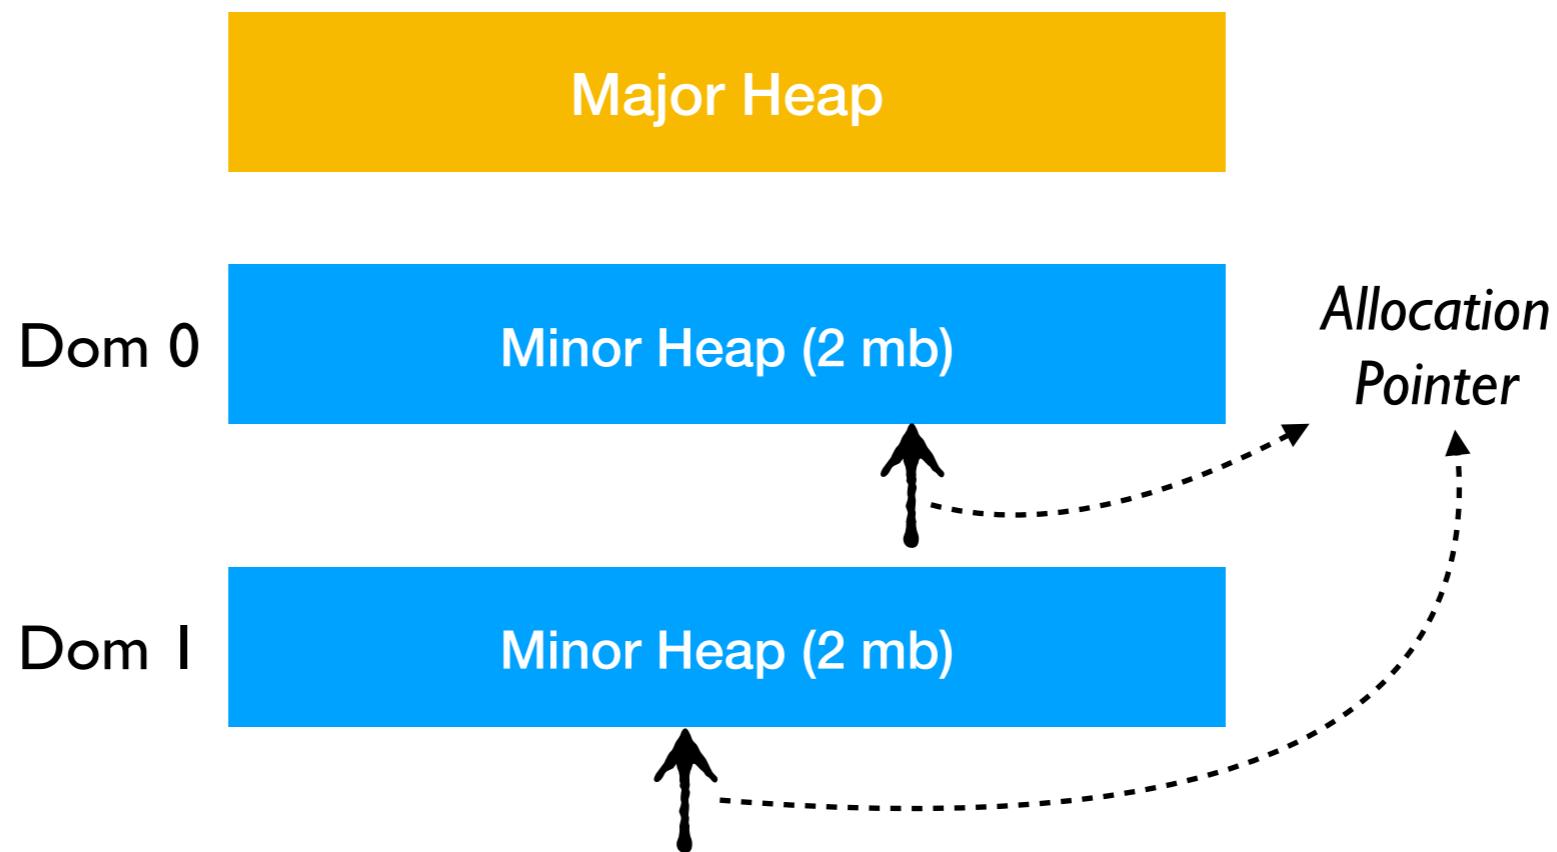
Stock OCaml GC

- A generational, non-moving, incremental, mark-and-sweep GC



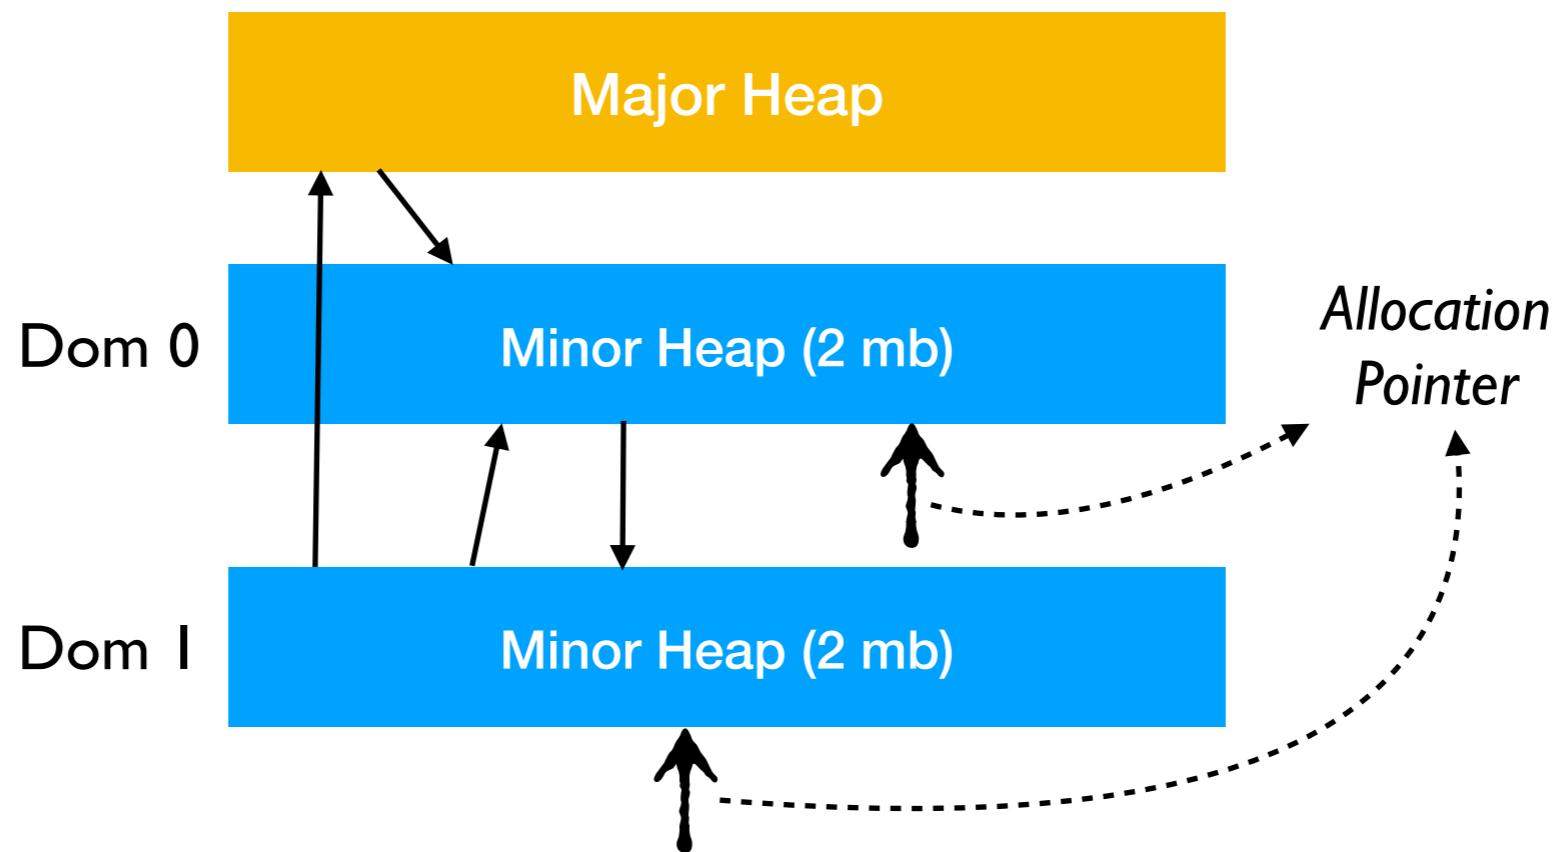
- Fast allocations
- Max GC latency < 10 ms, 99th percentile latency < 1 ms

Multicore OCaml: Minor GC



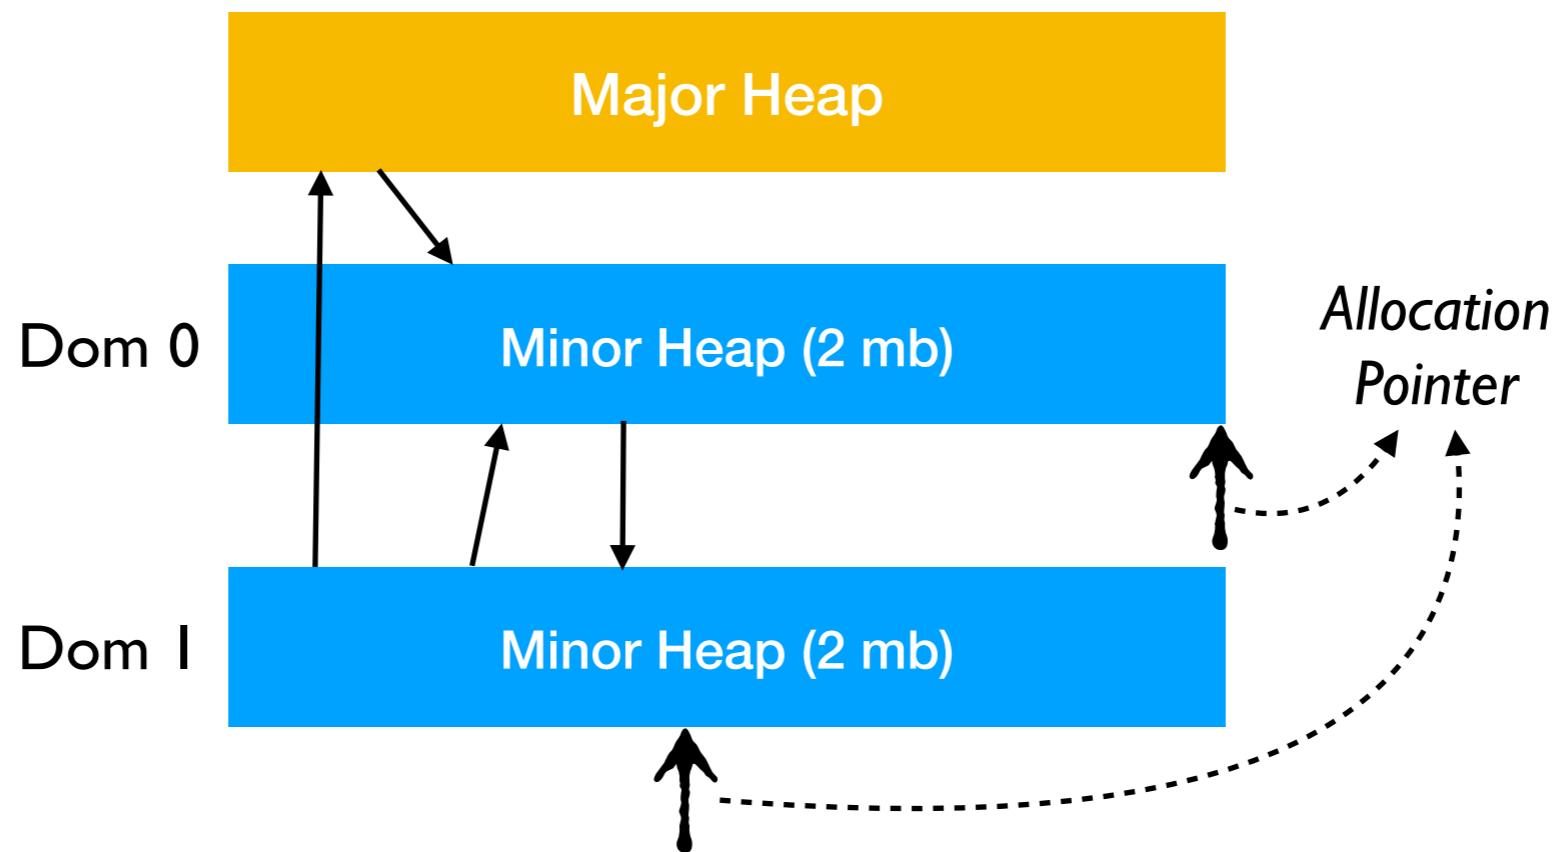
- Private minor heap arenas per domain
 - ♦ Fast allocations without synchronisation

Multicore OCaml: Minor GC



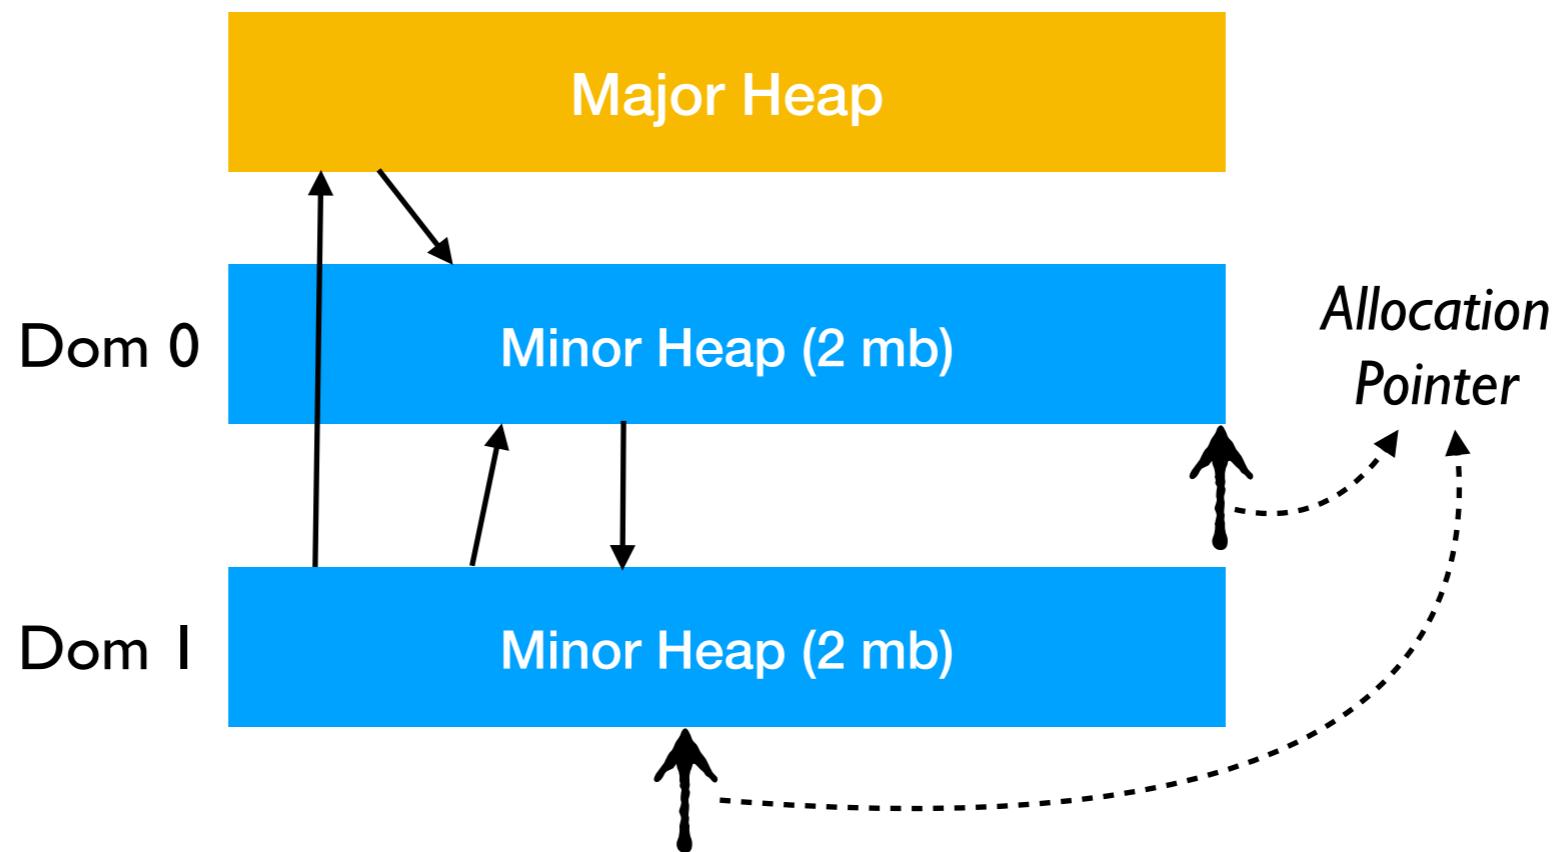
- Private minor heap arenas per domain
 - ♦ Fast allocations without synchronisation
- Pointers permitted between minor arenas and major heap

Multicore OCaml: Minor GC



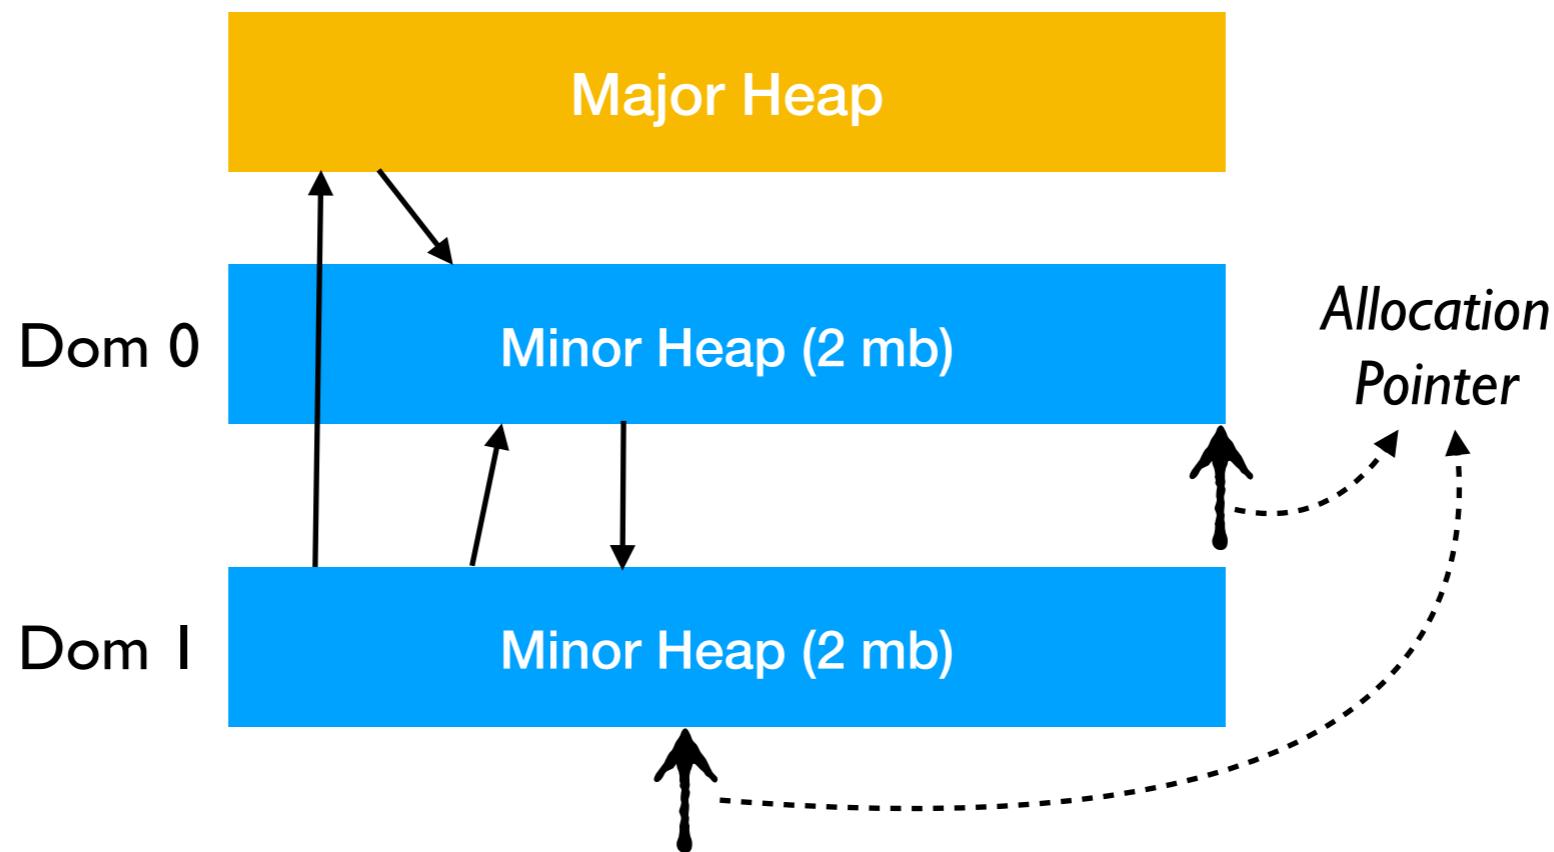
- *Stop-the-world parallel minor collection* when a domain runs out of minor heap
 - ◆ 2 global barriers / minor gc

Multicore OCaml: Minor GC



- *Stop-the-world parallel minor collection* when a domain runs out of minor heap
 - ◆ 2 global barriers / minor gc
- Bringing domains to a stop surprisingly fast
 - ◆ Poll points inserted into non-allocating loops

Multicore OCaml: Minor GC



- *Stop-the-world parallel minor collection* when a domain runs out of minor heap
 - ◆ 2 global barriers / minor gc
- Bringing domains to a stop surprisingly fast
 - ◆ Poll points inserted into non-allocating loops
- On 24 cores, ~ 10 ms pauses for completing stop-the-world minor GC

Multicore OCaml: Allocator

- Multicore-aware allocator for major heap
 - ◆ Based on Streamflow [Schneider et al. 2006]

Multicore OCaml: Allocator

- Multicore-aware allocator for major heap
 - ◆ Based on Streamflow [Schneider et al. 2006]
- *Thread-local, size-segmented free lists* for small objects
 - ◆ Most allocations in OCaml are small (99% of objects < 5 words in size)
 - ◆ Malloc for large allocations

Multicore OCaml: Allocator

- Multicore-aware allocator for major heap
 - ◆ Based on Streamflow [Schneider et al. 2006]
- *Thread-local, size-segmented free lists* for small objects
 - ◆ Most allocations in OCaml are small (99% of objects < 5 words in size)
 - ◆ Malloc for large allocations
- Most allocations do not need synchronisation

Multicore OCaml: Allocator

- Multicore-aware allocator for major heap
 - ◆ Based on Streamflow [Schneider et al. 2006]
- *Thread-local, size-segmented free lists* for small objects
 - ◆ Most allocations in OCaml are small (99% of objects < 5 words in size)
 - ◆ Malloc for large allocations
- Most allocations do not need synchronisation
- Sequential performance on par with OCaml's recent *best-fit* allocator

Multicore OCaml: Major GC

- *Mostly-concurrent mark-and-sweep* for major collection
 - ♦ Based on VCGC [Huelsbergen and Winterbottom 1998]

Multicore OCaml: Major GC

- *Mostly-concurrent mark-and-sweep* for major collection

- *Based on VCGC [Huelsbergen and Winterbottom 1998]*

- GC colours

Unmarked

Marked

Garbage

Free

Multicore OCaml: Major GC

- *Mostly-concurrent mark-and-sweep* for major collection
 - ◆ Based on VCGC [Huelsbergen and Winterbottom 1998]
- GC colours Unmarked Marked Garbage Free
- Domains alternate between running the mutator and the GC

Multicore OCaml: Major GC

- *Mostly-concurrent mark-and-sweep* for major collection

- Based on VCGC [Huelsbergen and Winterbottom 1998]

- GC colours

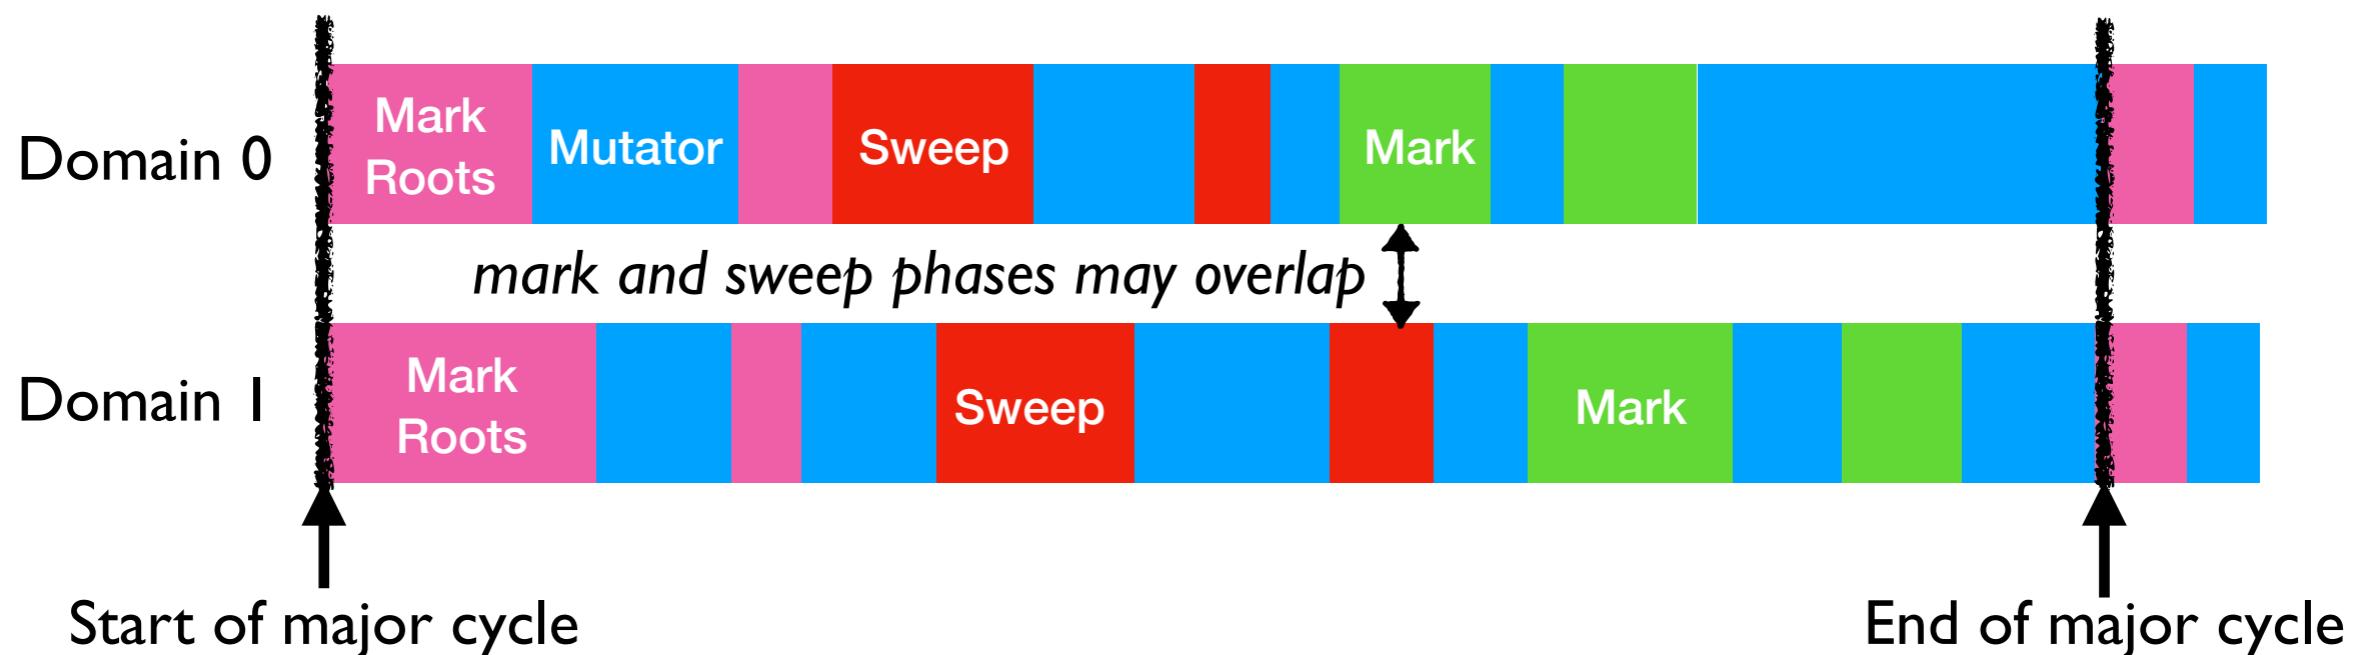
Unmarked

Marked

Garbage

Free

- Domains alternate between running the mutator and the GC



Multicore OCaml: Major GC

- *Mostly-concurrent mark-and-sweep* for major collection

- Based on VCGC [Huelsbergen and Winterbottom 1998]

- GC colours

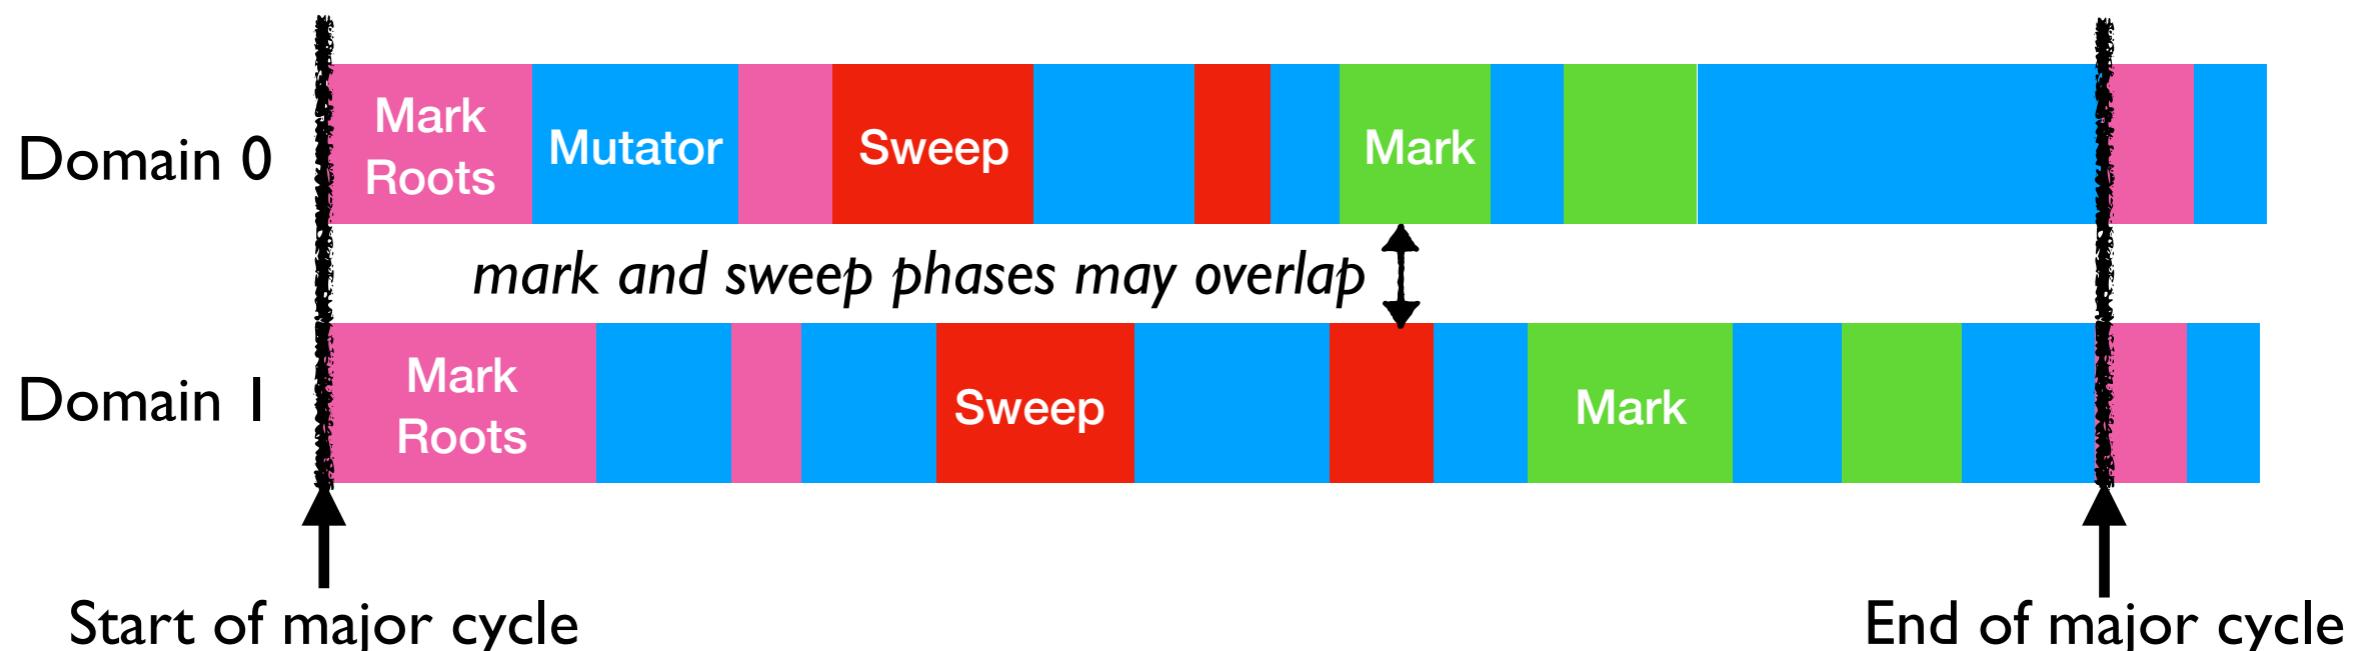
Unmarked

Marked

Garbage

Free

- Domains alternate between running the mutator and the GC



- Marking:

Unmarked

 →

Marked

 Sweeping:

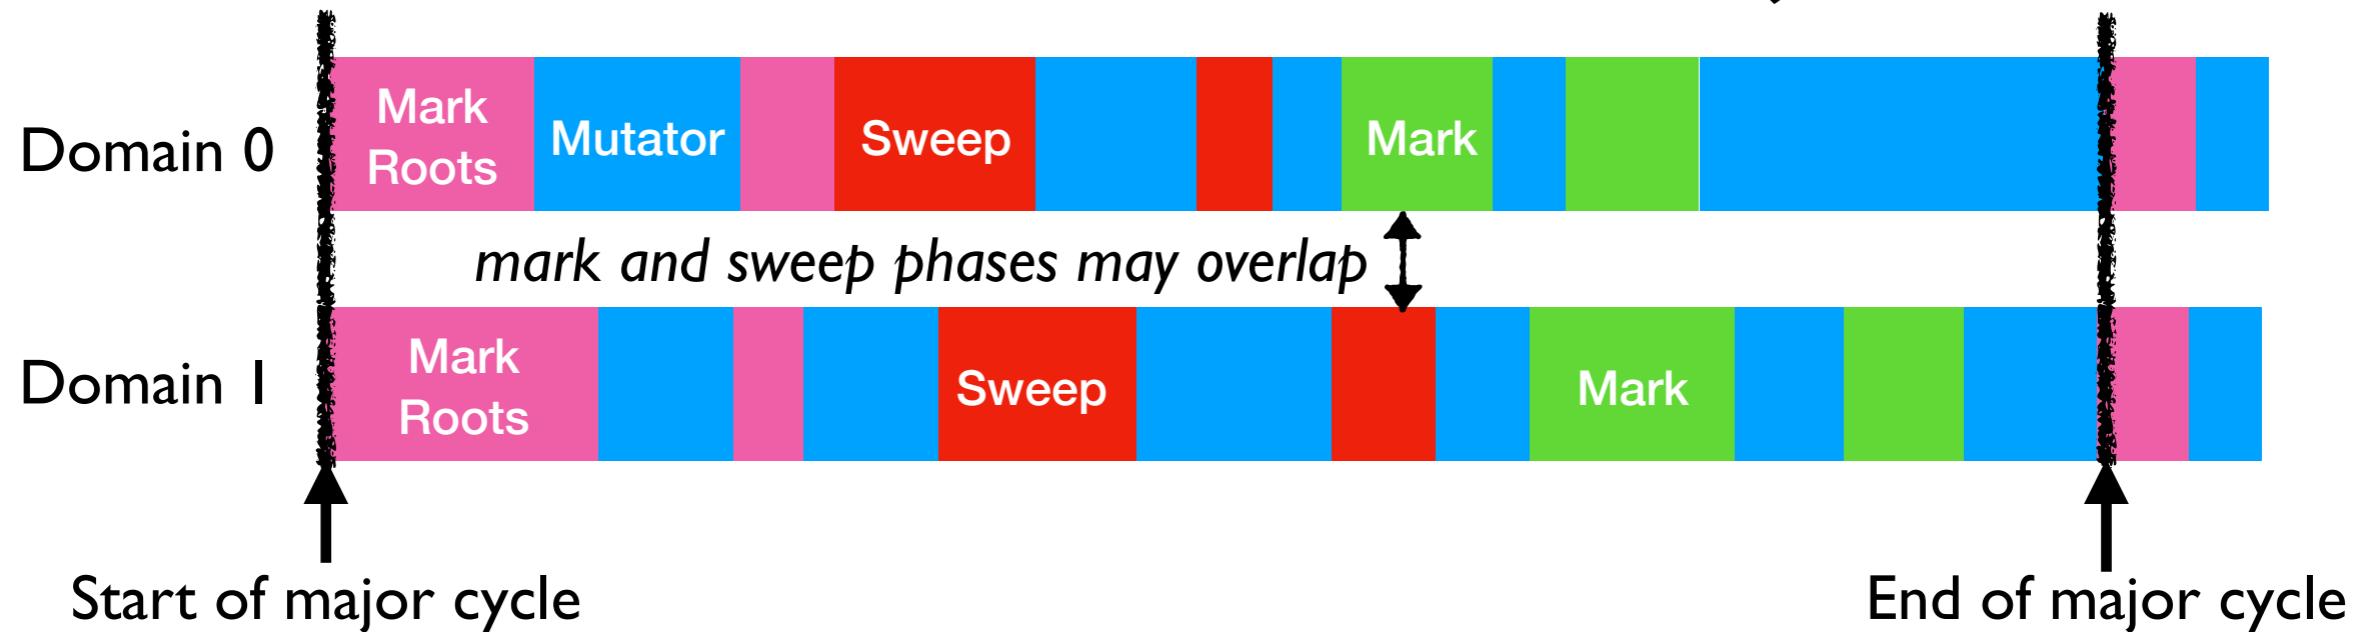
Garbage

 →

Free

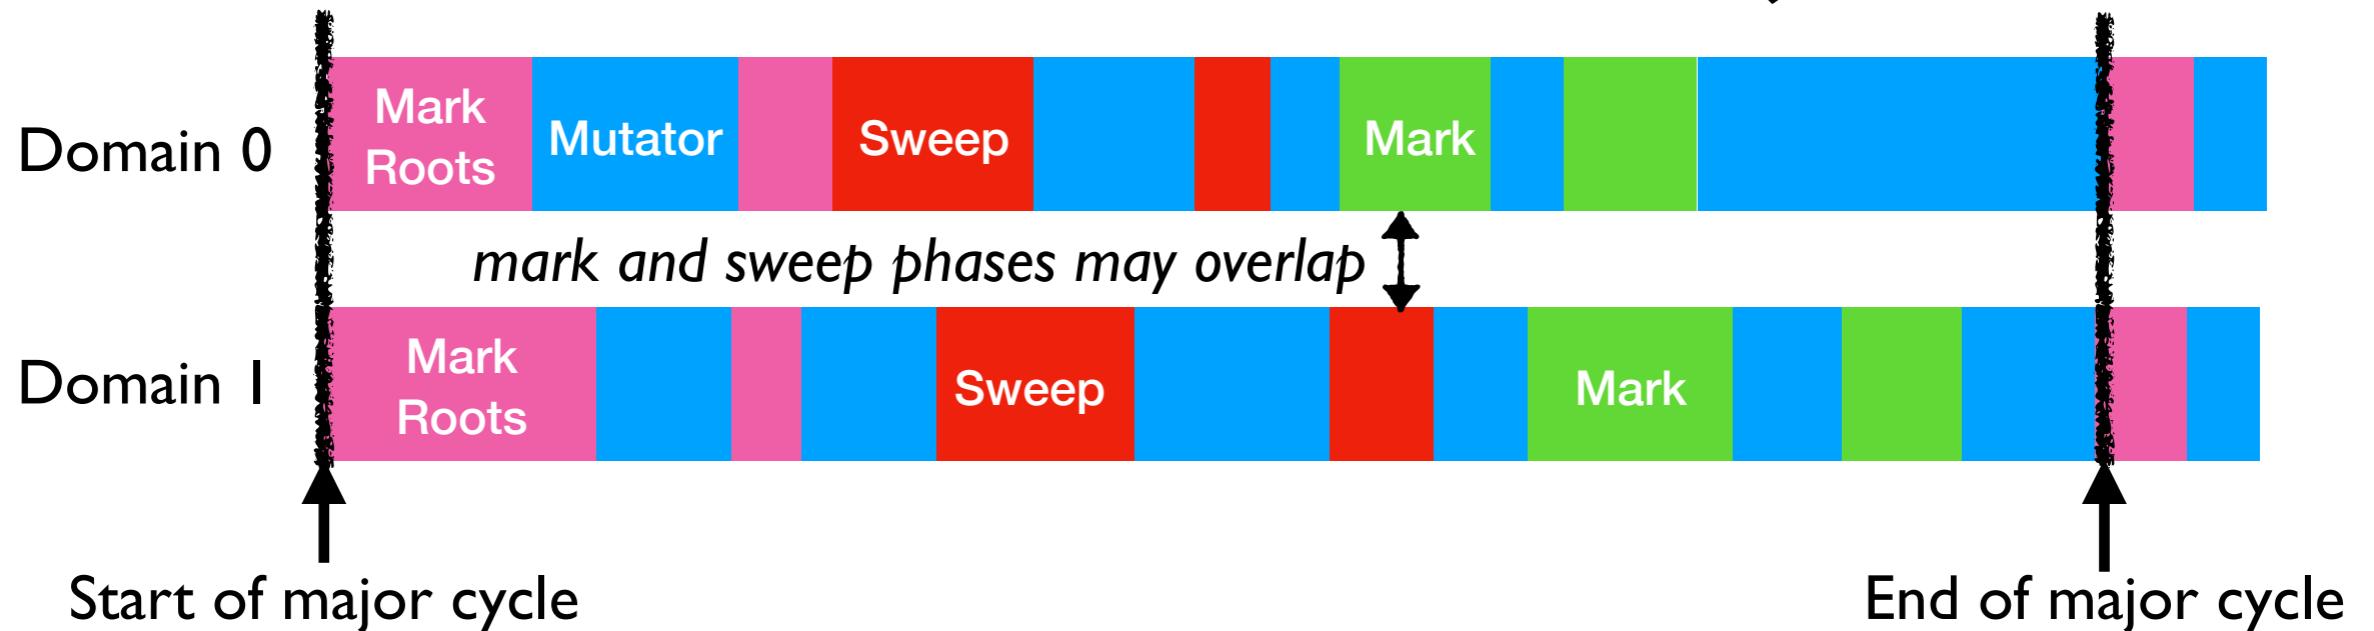
- Marking is *racy* but uses plain *non-atomic stores and idempotent*

Multicore OCaml: Major GC

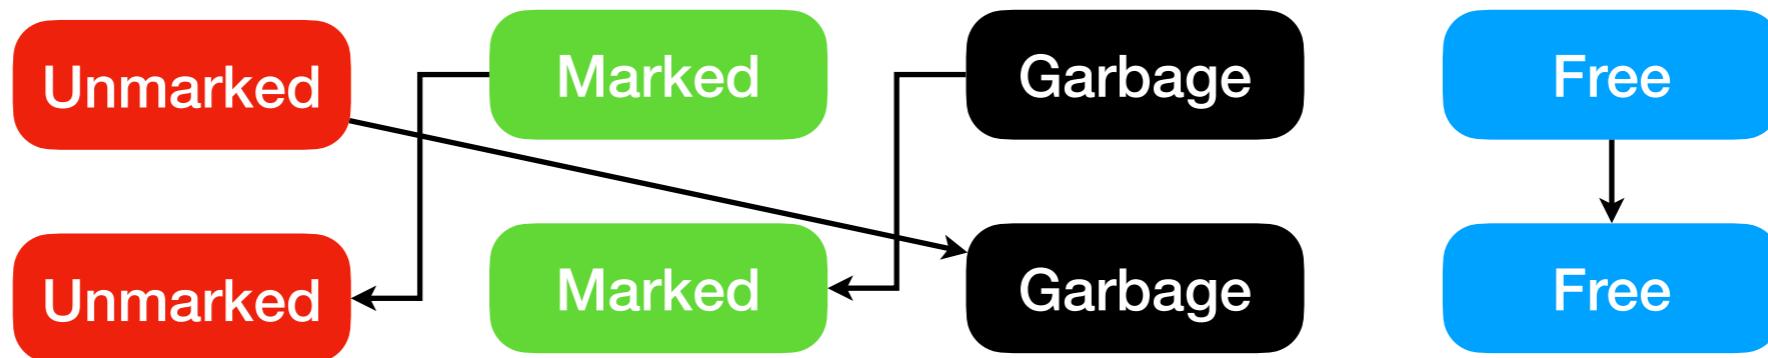


- Marking & sweeping done \Rightarrow stop-the-world

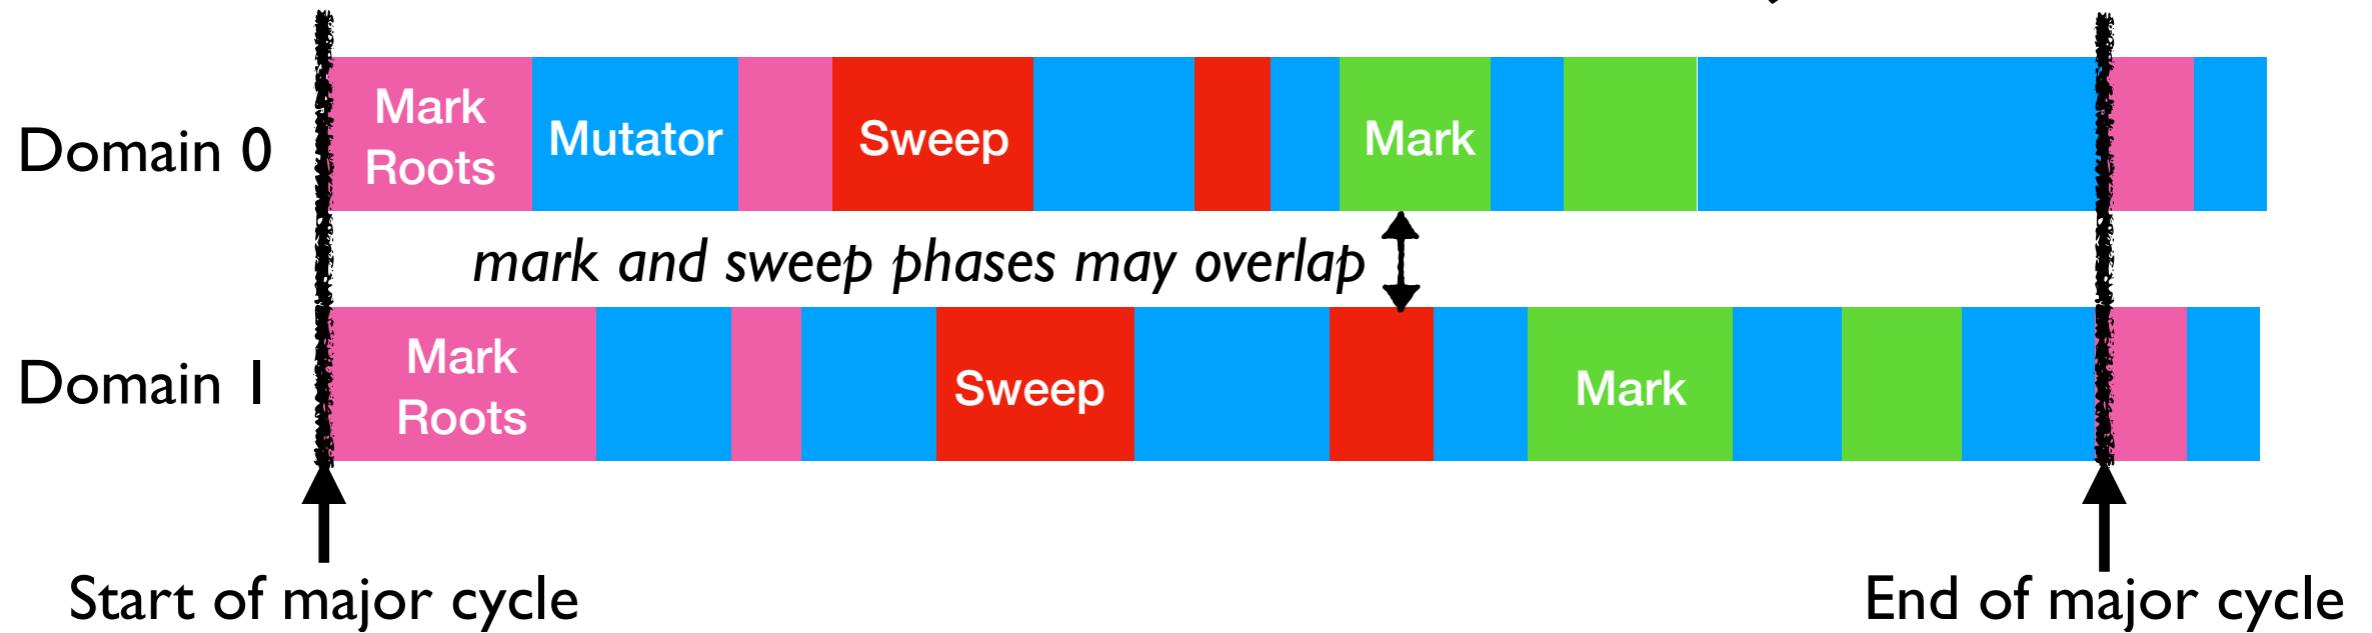
Multicore OCaml: Major GC



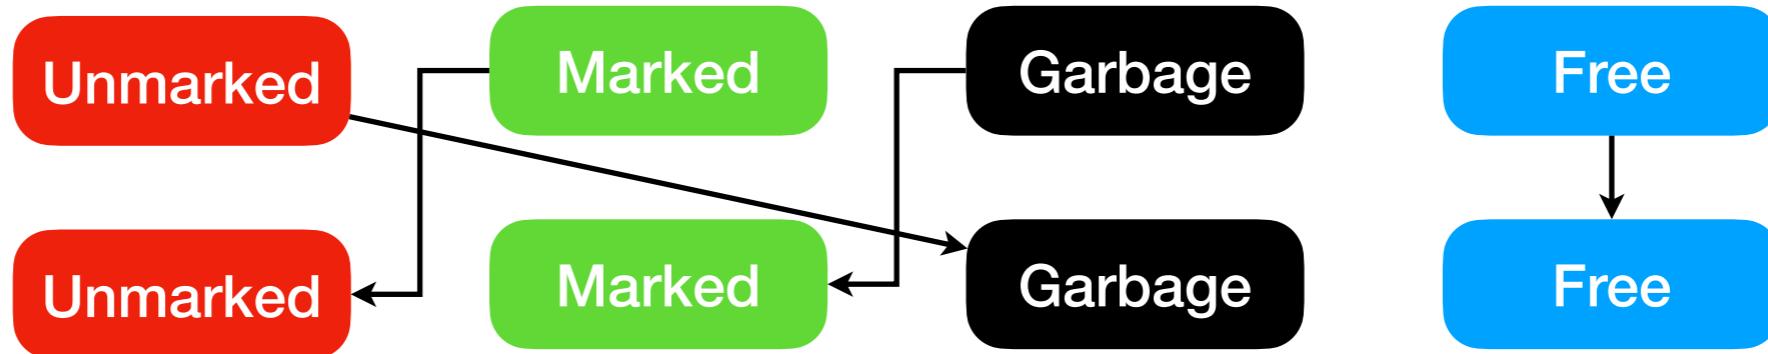
- Marking & sweeping done \Rightarrow stop-the-world



Multicore OCaml: Major GC

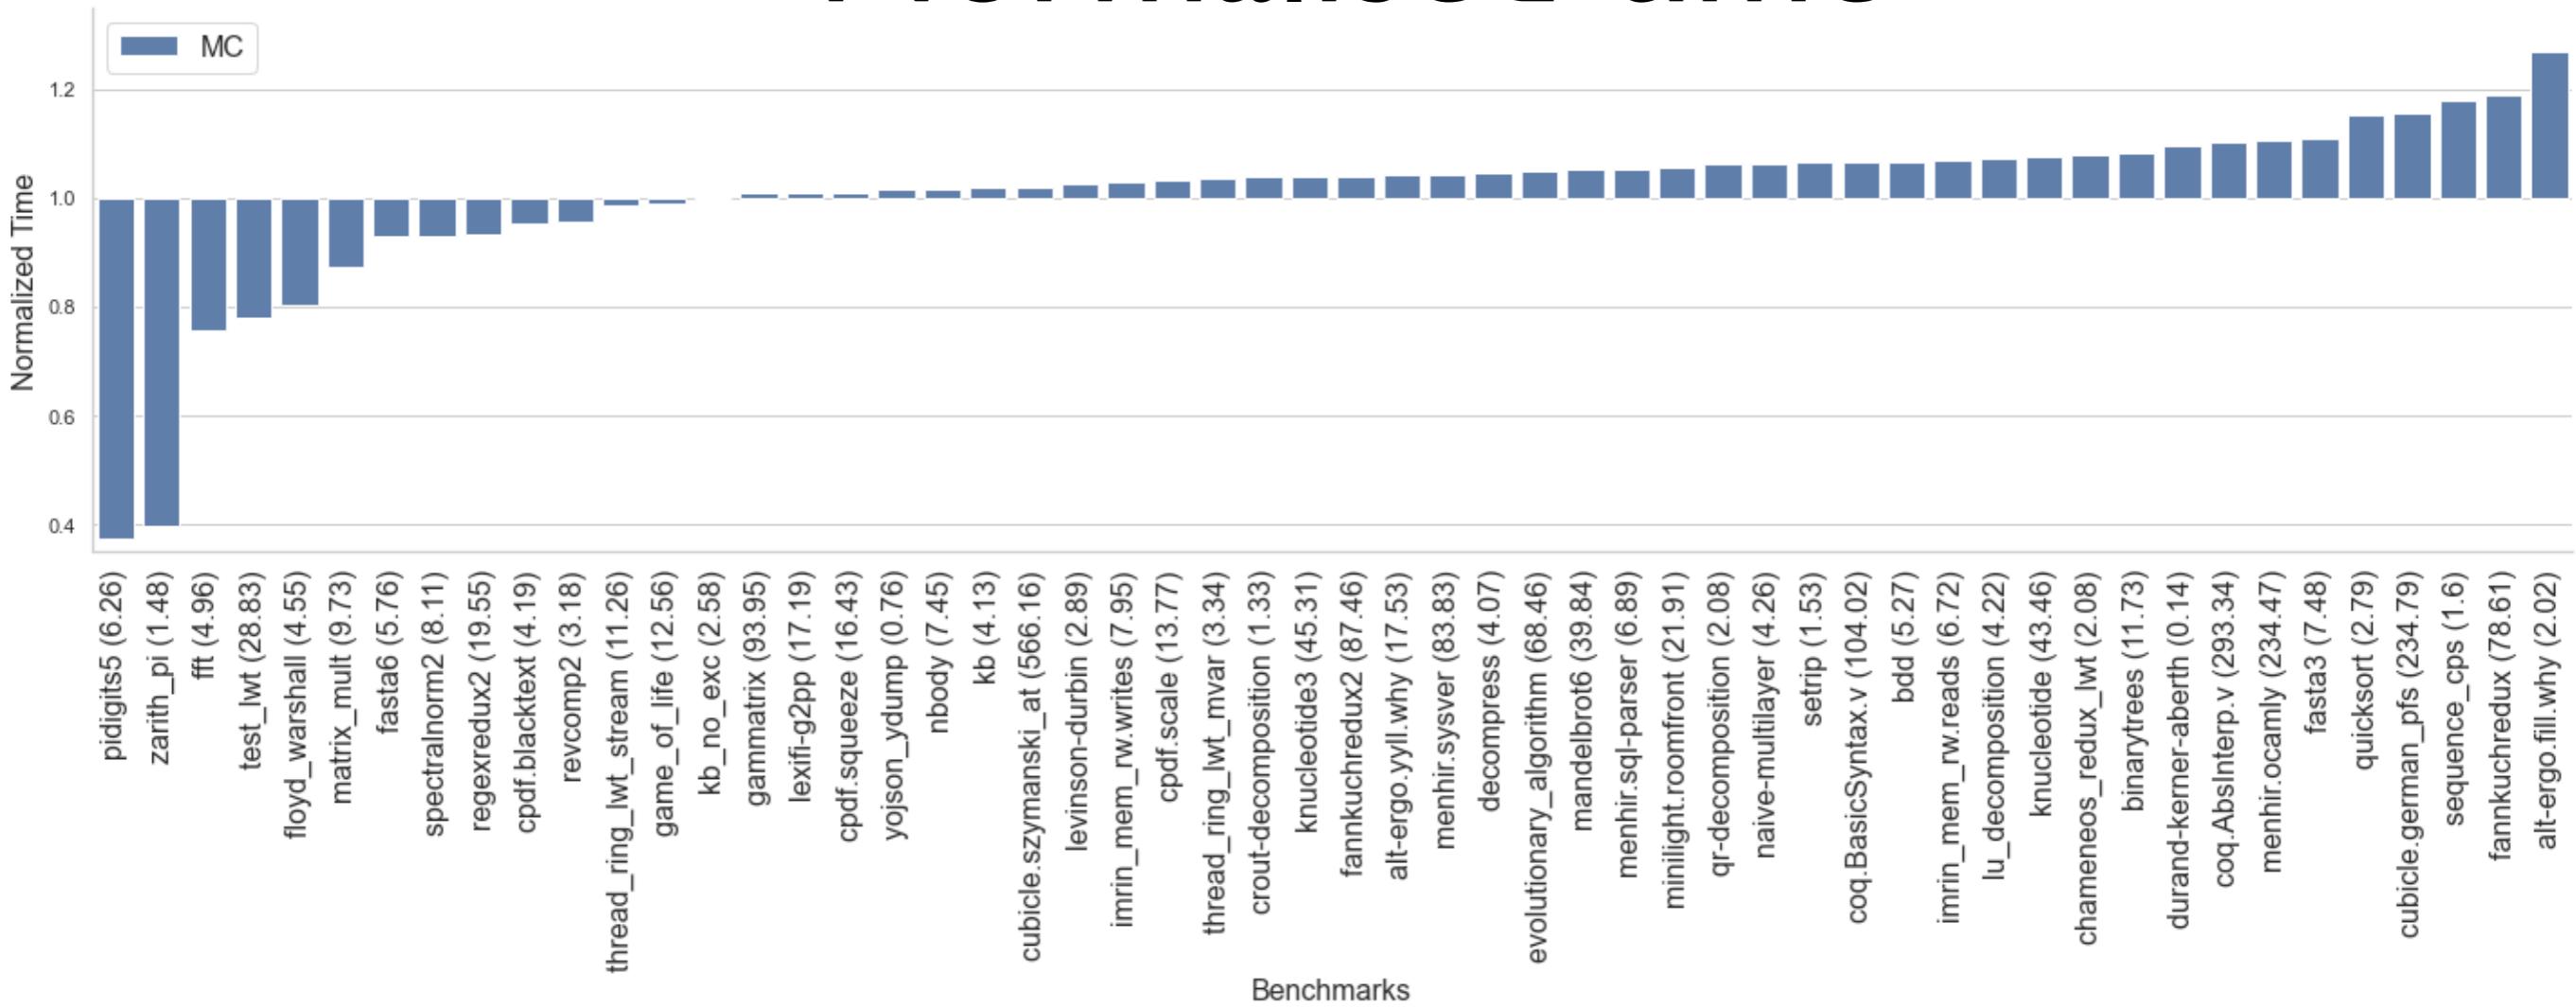


- Marking & sweeping done \Rightarrow stop-the-world

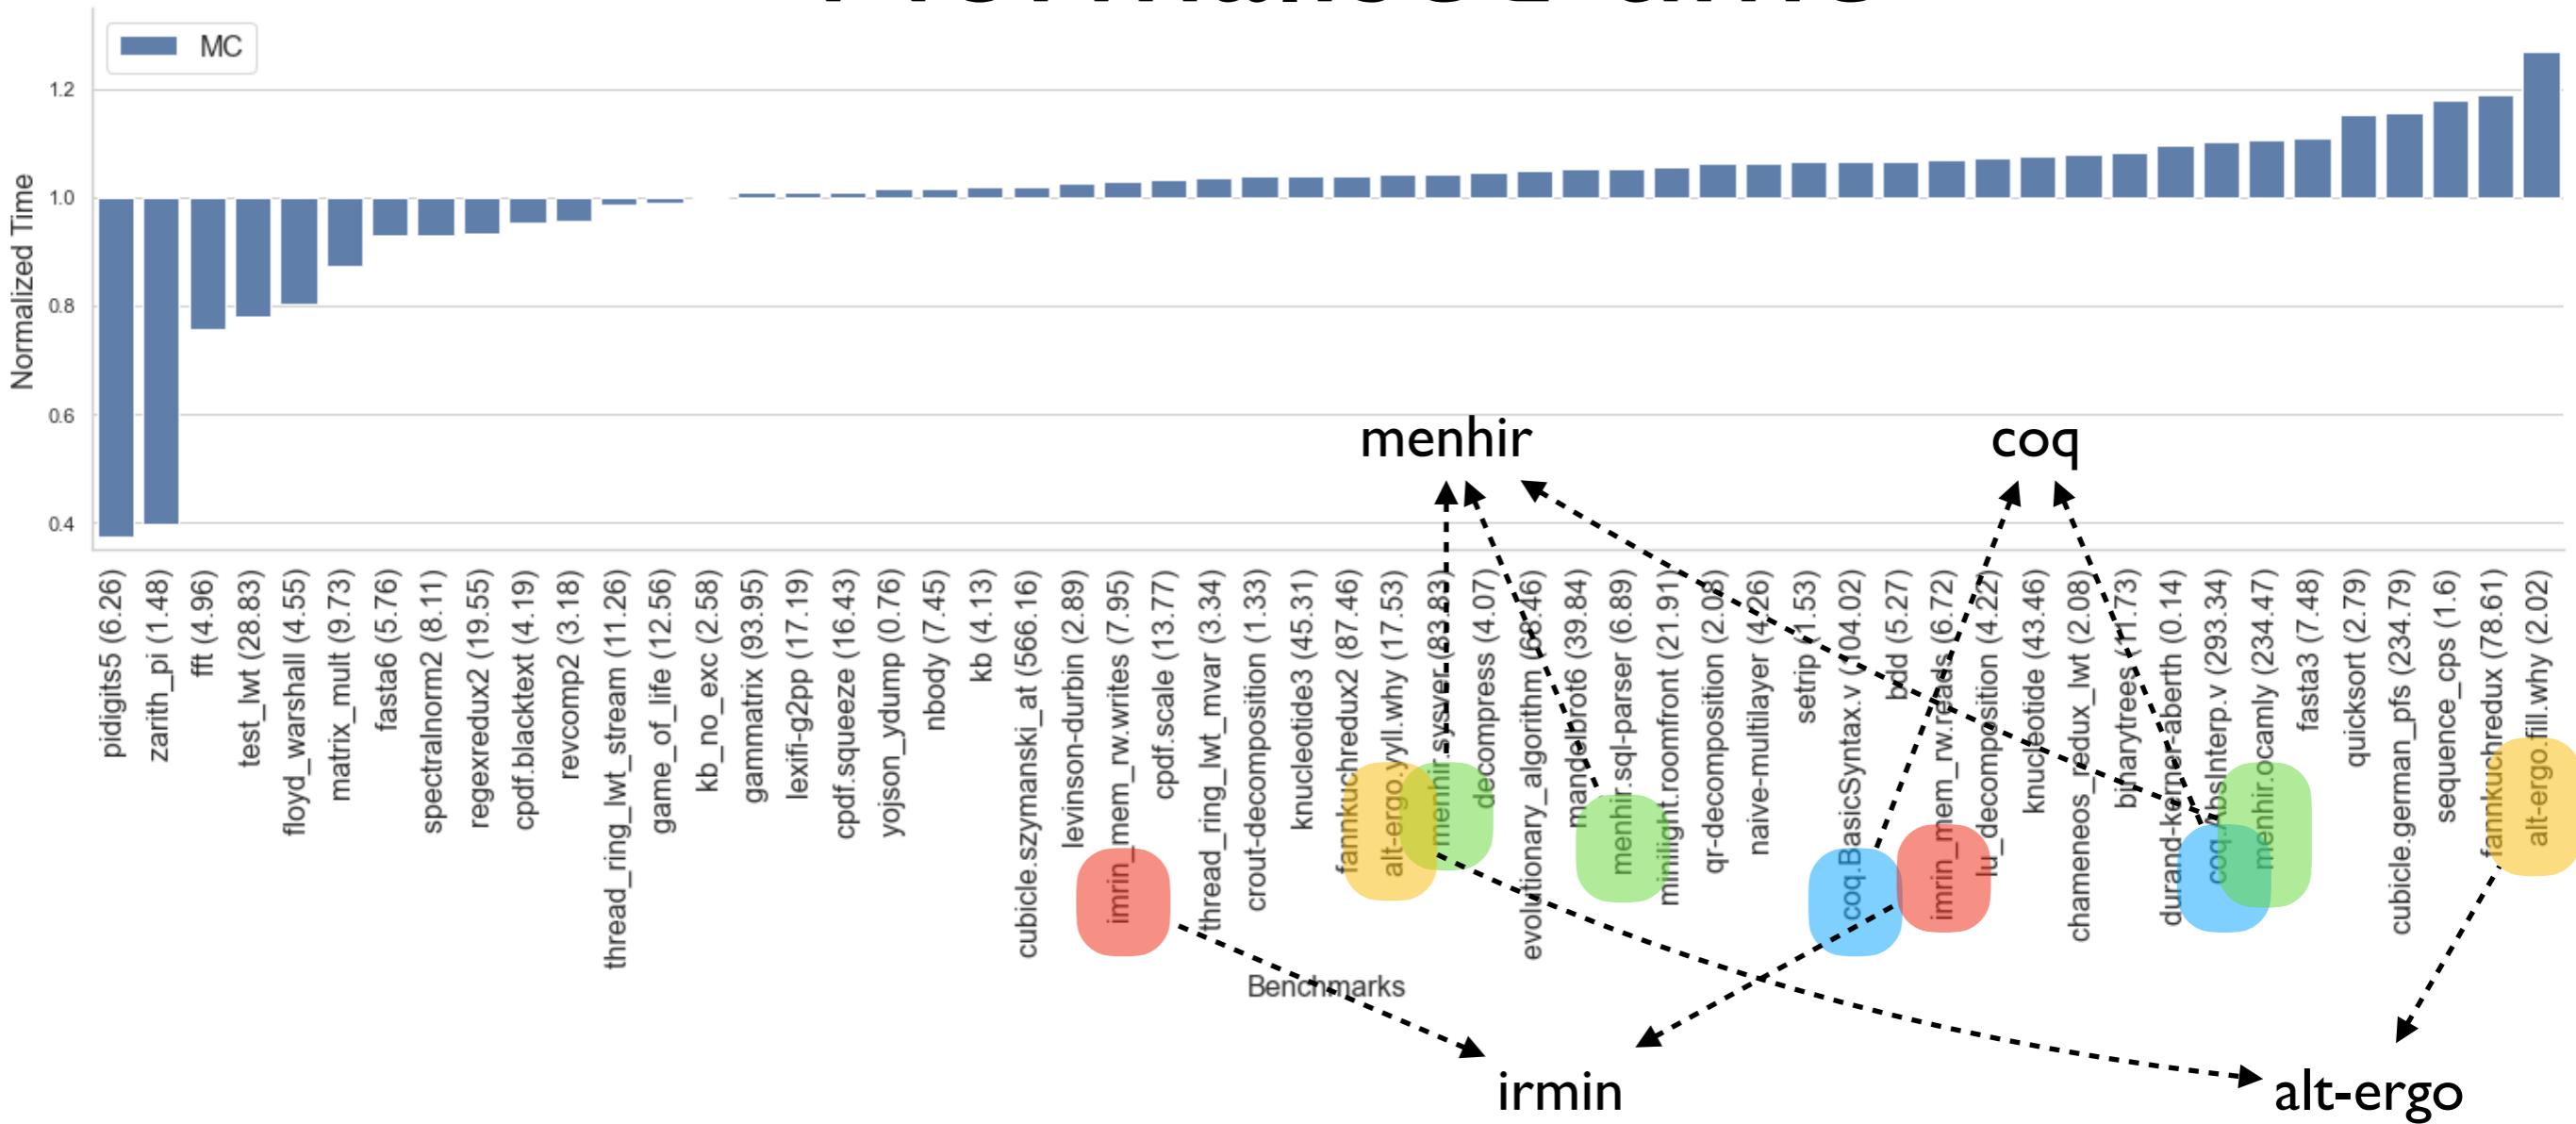


- Stop-the-world pauses are ~ 5 ms on 24 cores

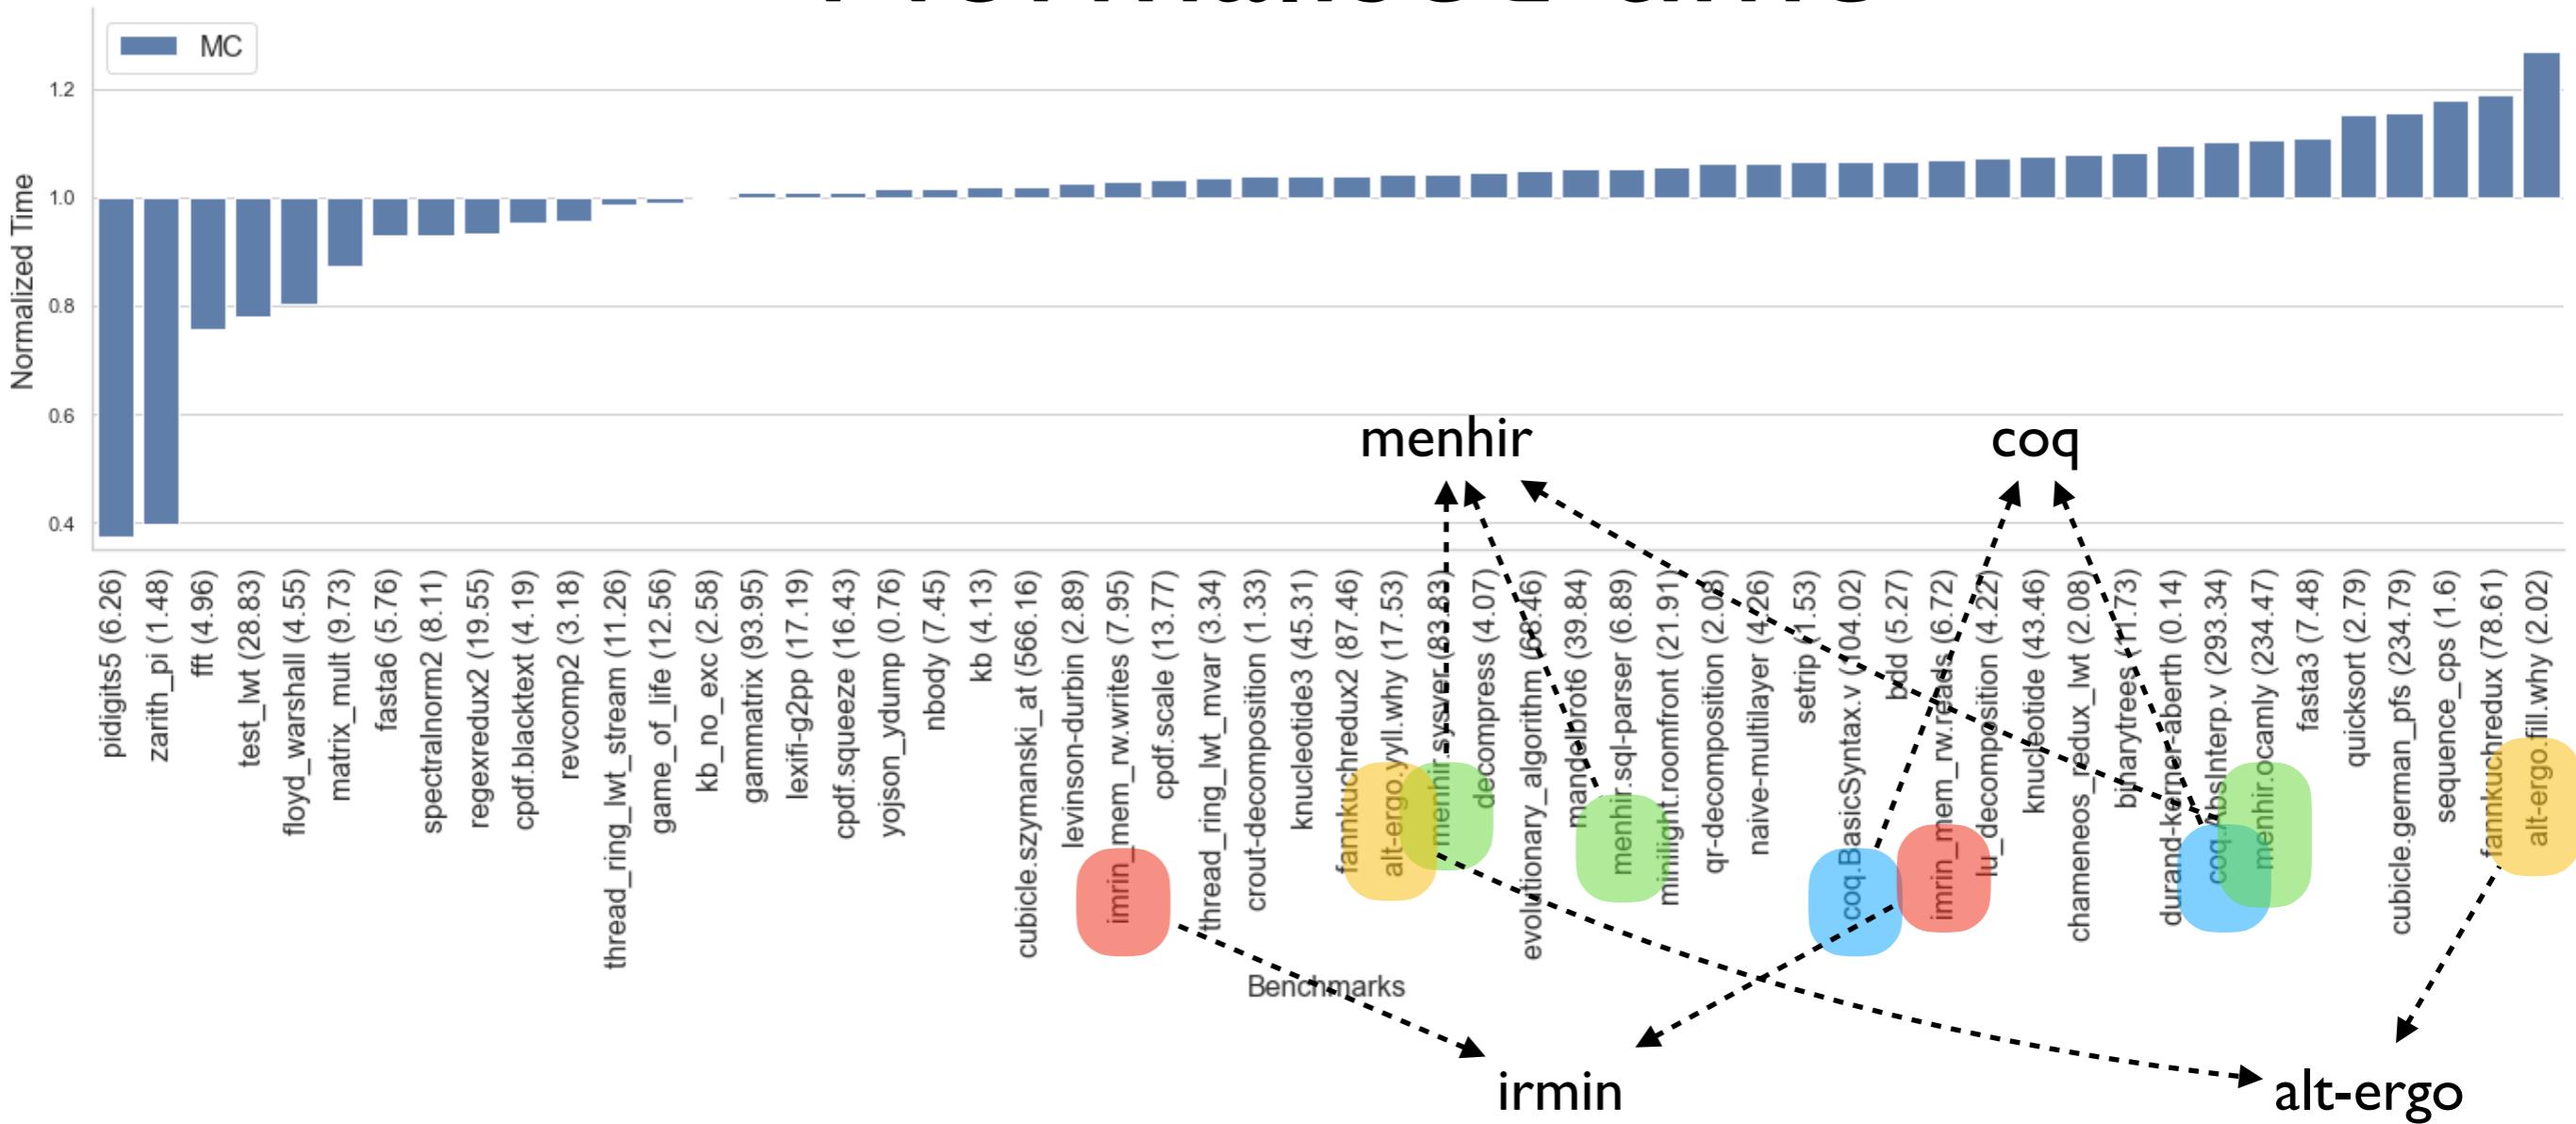
Sequential Performance: Normalised time



Sequential Performance: Normalised time

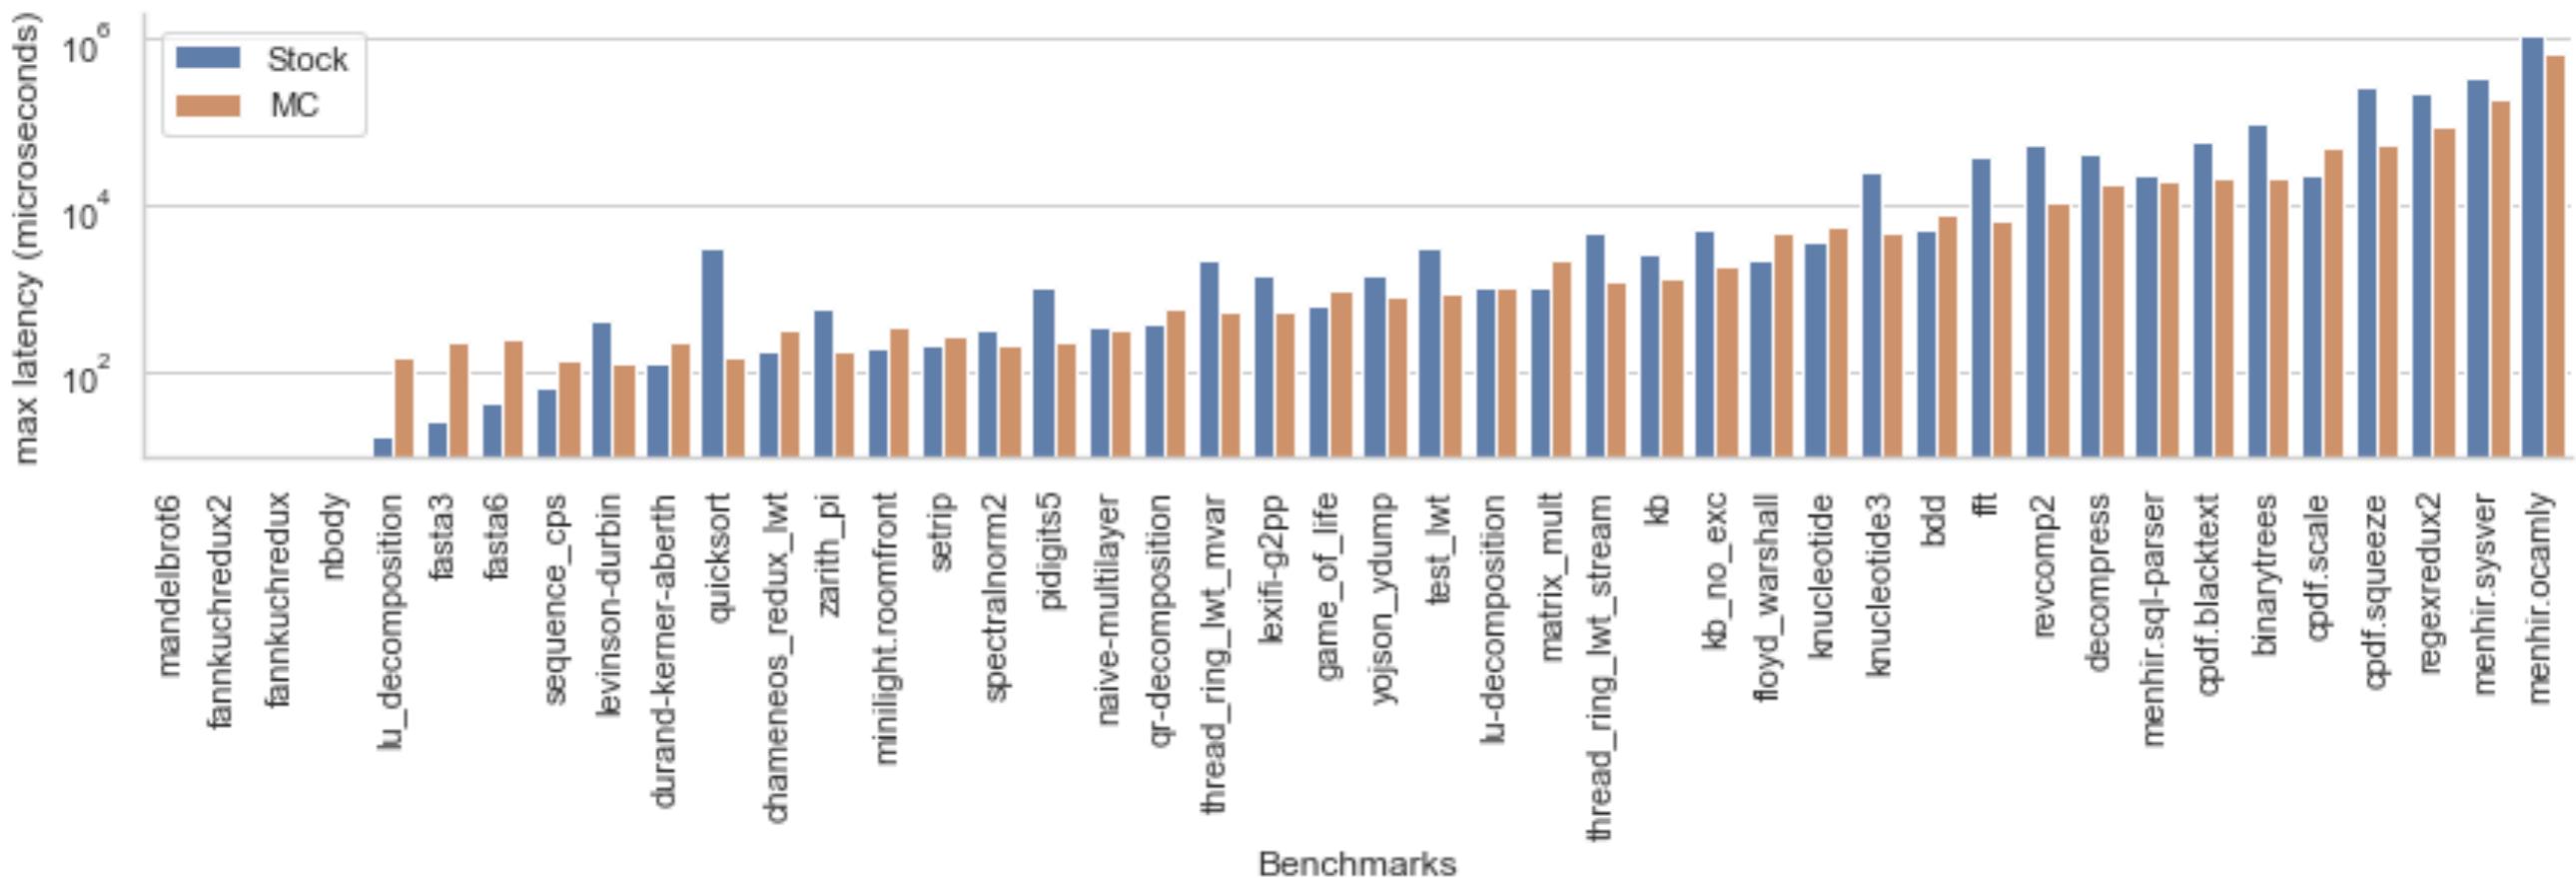


Sequential Performance: Normalised time

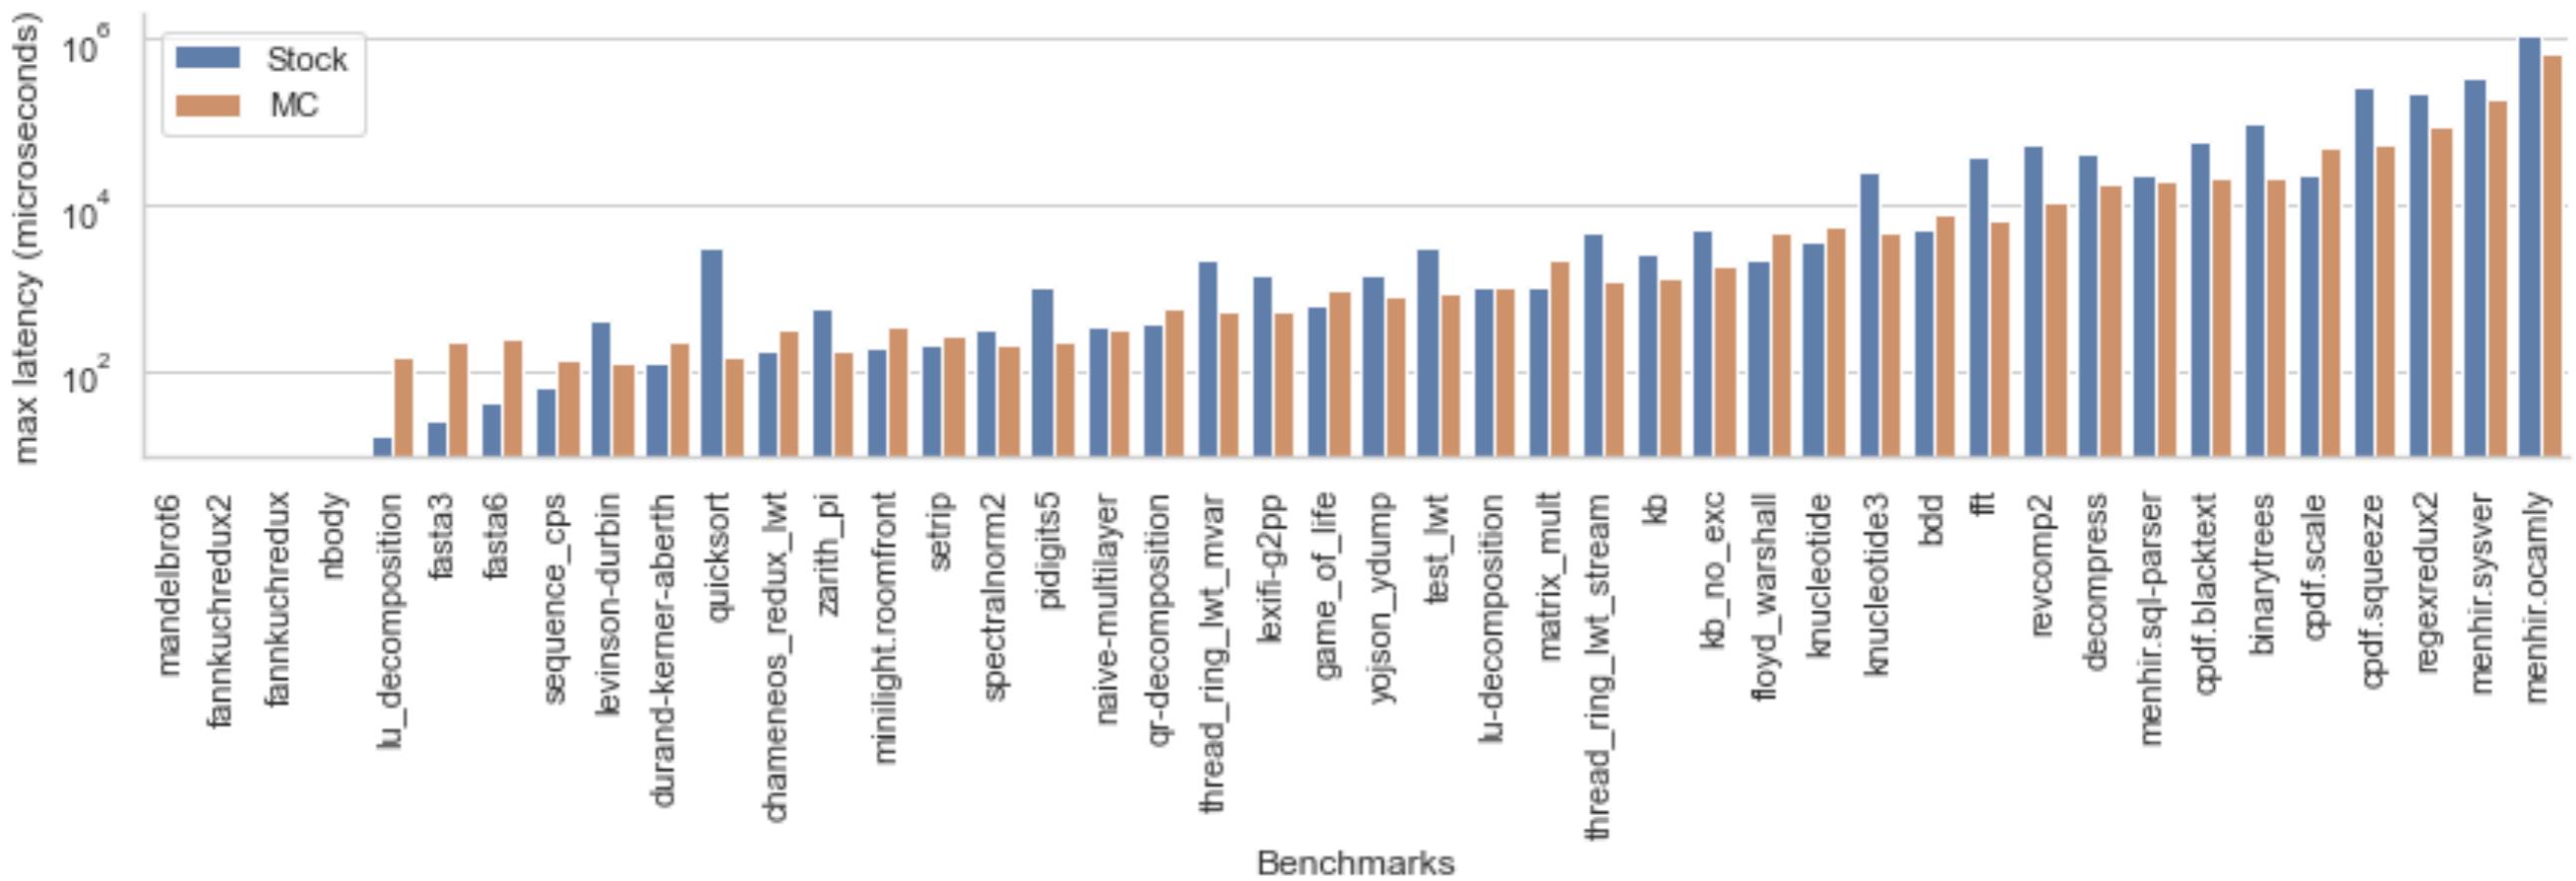


- **~1% faster than stock** (geomean of normalised running times)
 - ◆ Difference under measurement noise mostly
 - ◆ Outliers due to difference in allocators

Sequential Performance: Max pause time



Sequential Performance: Max pause time



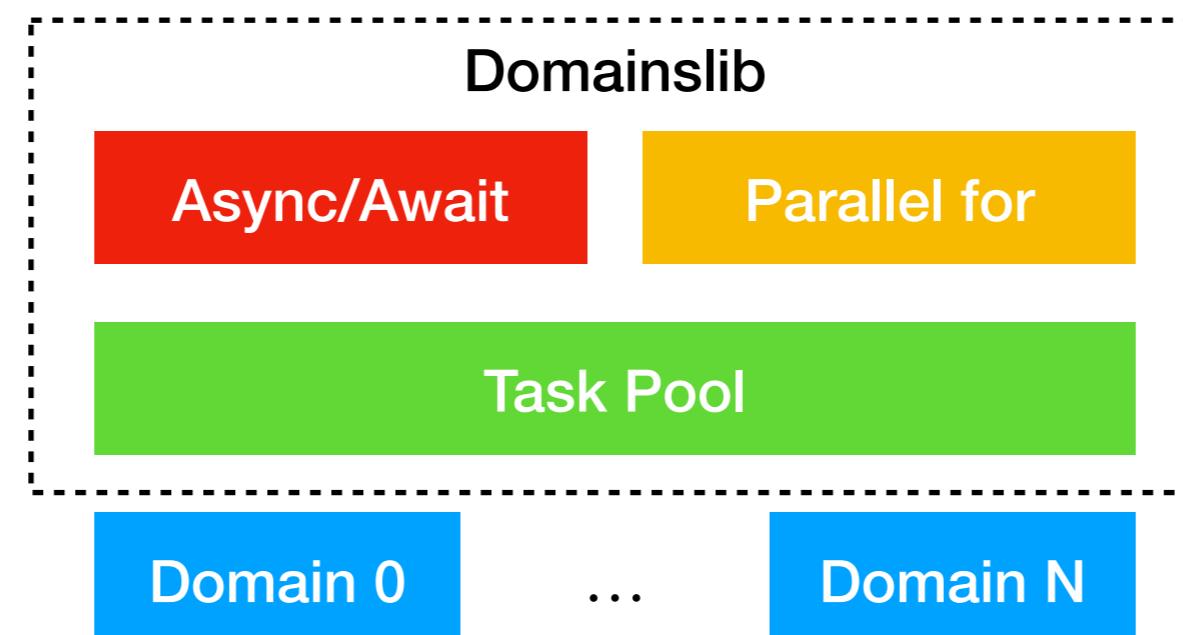
- *Pausetimes are lower under Multicore OCaml than stock OCaml*

Domainslib for parallel programming

- Domain API exposed by the compiler is too low-level
 - ◆ Mutex, condition variables, ...

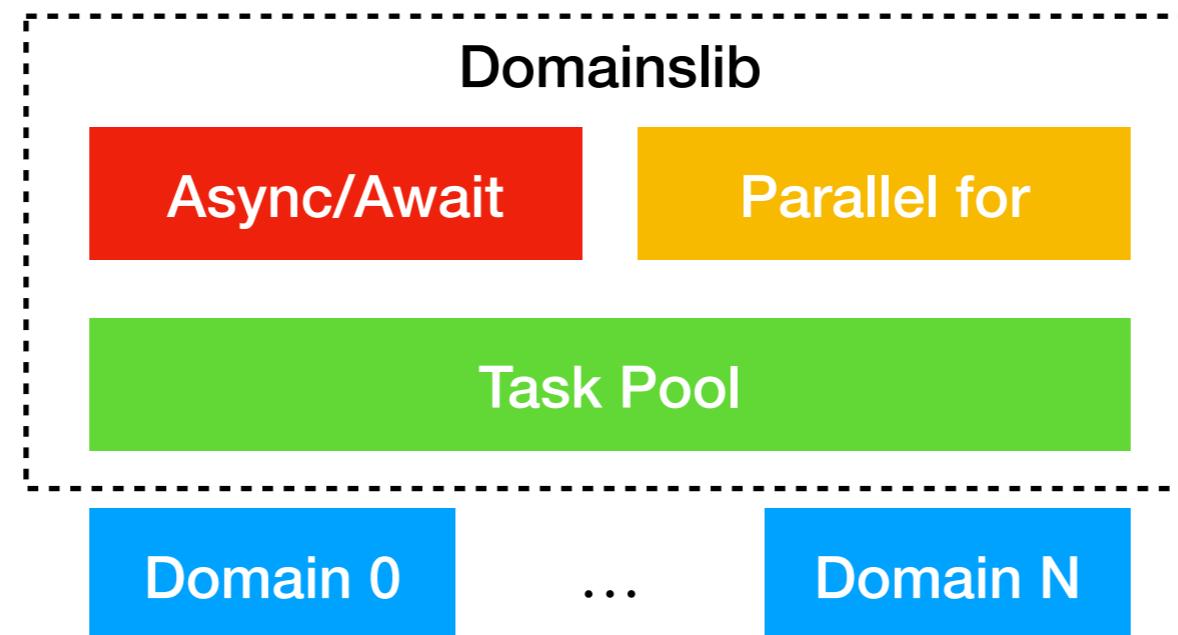
Domainslib for parallel programming

- Domain API exposed by the compiler is too low-level
 - ◆ Mutex, condition variables, ...
- Domainslib - <https://github.com/ocaml-multicore/domainslib>



Domainslib for parallel programming

- Domain API exposed by the compiler is too low-level
 - ◆ Mutex, condition variables, ...
- Domainslib - <https://github.com/ocaml-multicore/domainslib>



Let's look at examples!

Recursive Fibonacci - Sequential

```
let rec fib n =  
  if n < 2 then 1  
  else fib (n-1) + fib (n-2)
```

Recursive Fibonacci - Parallel

```
module T = Domainslib.Task
```

```
let fib n =
  let pool = T.setup_pool
    ~num_additional_domains:(num_domains - 1) in
  let res = fib_par pool n in
  T.teardown_pool pool;
  res
```

Recursive Fibonacci - Parallel

```
module T = Domainslib.Task
```

```
let rec fib_par pool n =
  if n <= 40 then fib_seq n
  else
    let a = T.async pool (fun _ -> fib_par pool (n-1)) in
    let b = T.async pool (fun _ -> fib_par pool (n-2)) in
    T.await pool a + T.await pool b
```

```
let fib n =
  let pool = T.setup_pool
    ~num_additional_domains:(num_domains - 1) in
  let res = fib_par pool n in
  T.teardown_pool pool;
  res
```

Recursive Fibonacci - Parallel

```
module T = Domainslib.Task

let rec fib_seq n =
  if n < 2 then 1
  else fib_seq (n-1) + fib_seq (n-2)

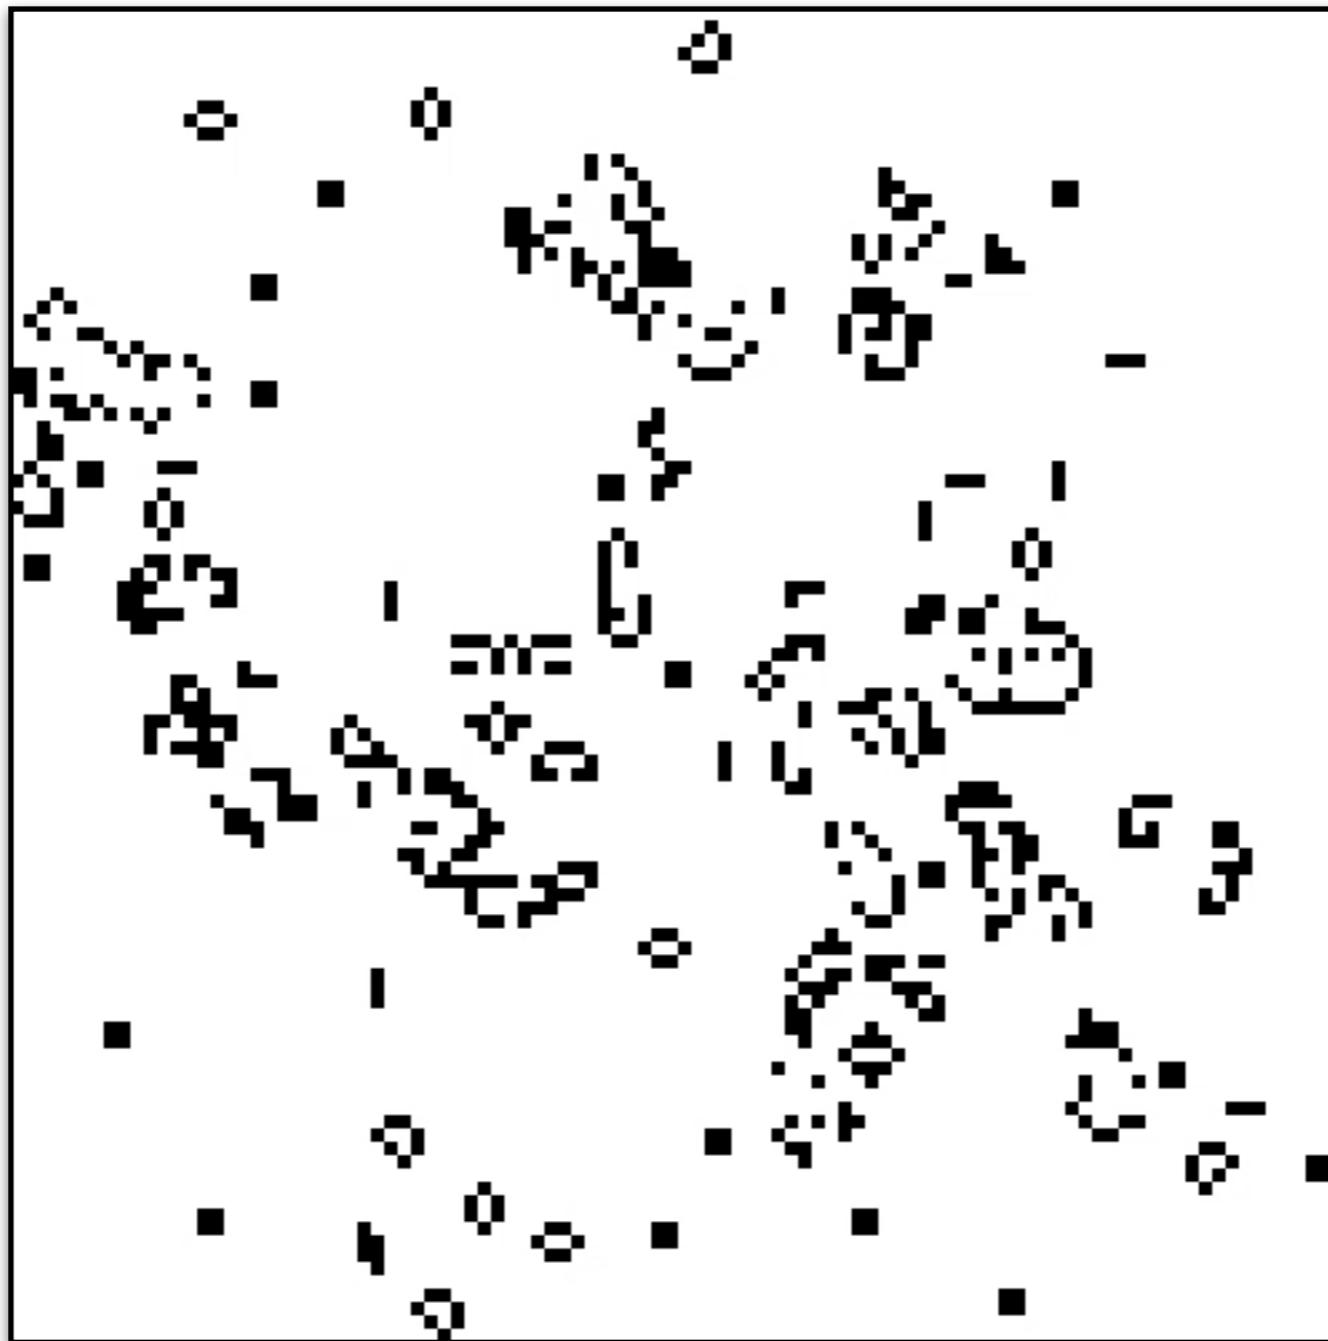
let rec fib_par pool n =
  if n <= 40 then fib_seq n
  else
    let a = T.async pool (fun _ -> fib_par pool (n-1)) in
    let b = T.async pool (fun _ -> fib_par pool (n-2)) in
    T.await pool a + T.await pool b

let fib n =
  let pool = T.setup_pool
    ~num_additional_domains:(num_domains - 1) in
  let res = fib_par pool n in
  T.teardown_pool pool;
  res
```

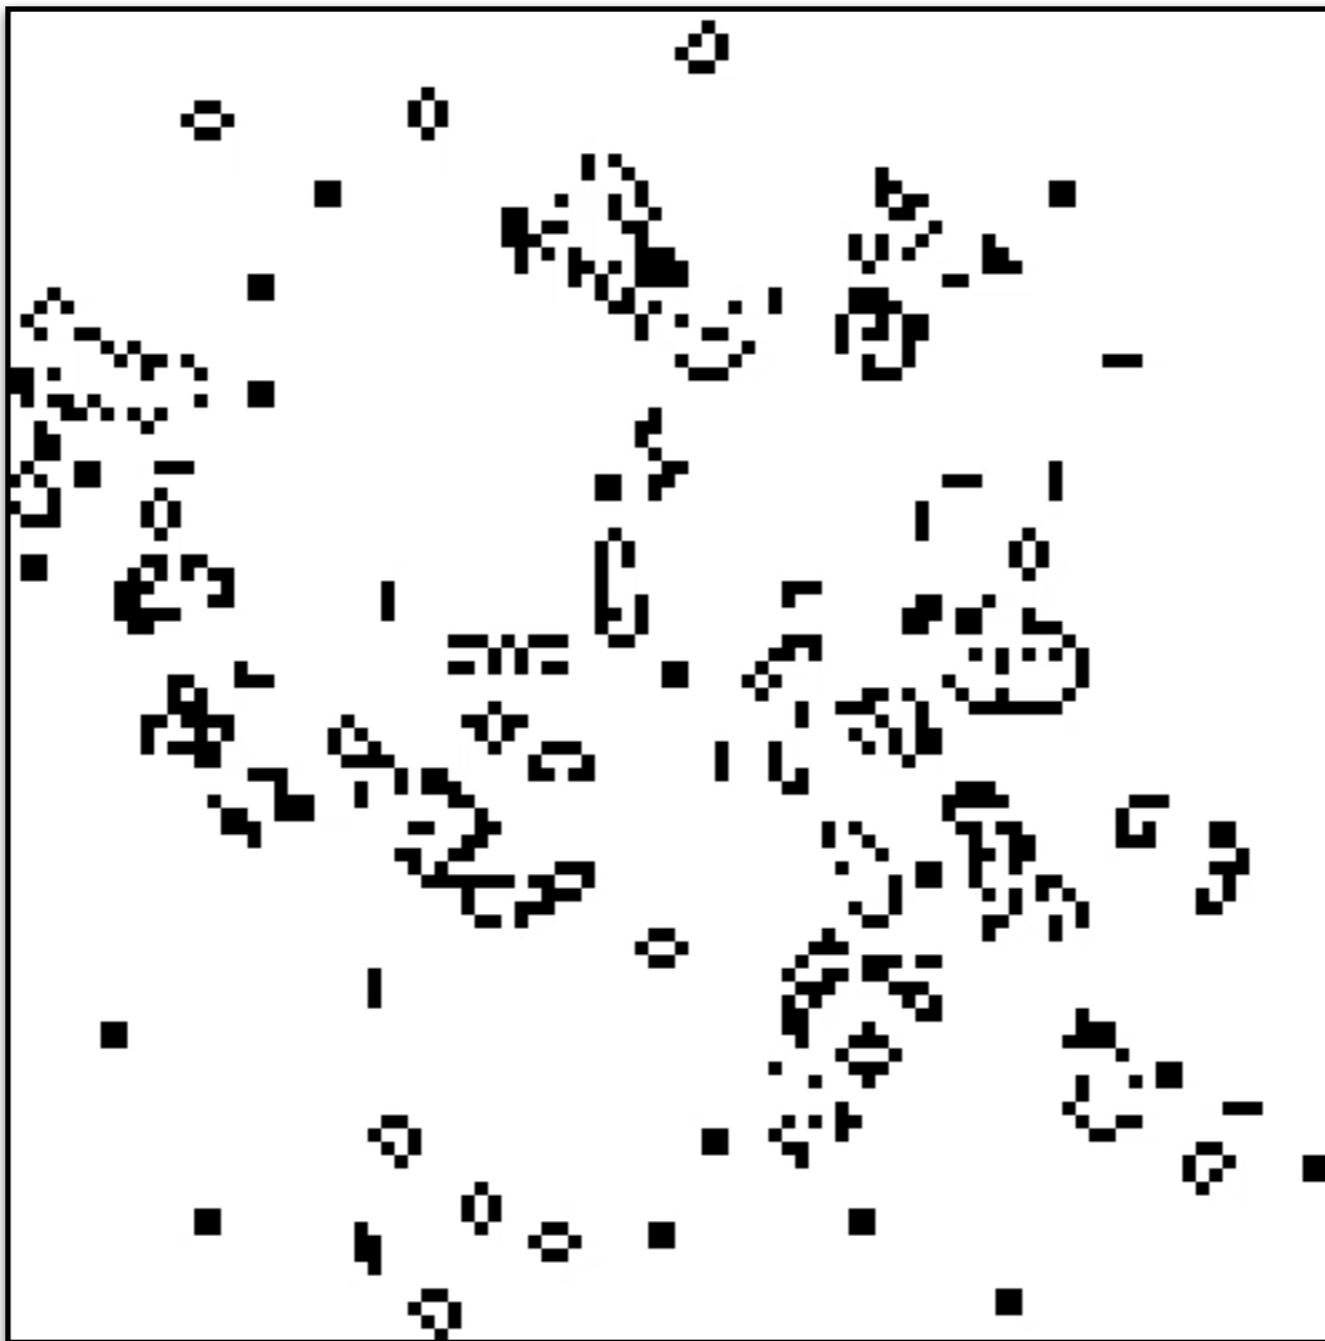
Performance: `fib(48)`

Cores	Time (Seconds)	Vs Serial	Vs Self
1	37.787	0.98	1
2	19.034	1.94	1.99
4	9.723	3.8	3.89
8	5.023	7.36	7.52
16	2.914	12.68	12.97
24	2.201	16.79	17.17

Conway's Game of Life



Conway's Game of Life



Conway's Game of Life

```
let next () =
  ...
  for x = 0 to board_size - 1 do
    for y = 0 to board_size - 1 do
      next_board.(x).(y) <- next_cell cur_board x y
    done
  done;
  ...

```

Conway's Game of Life

```
let next () =
  ...
  for x = 0 to board_size - 1 do
    for y = 0 to board_size - 1 do
      next_board.(x).(y) <- next_cell cur_board x y
    done
  done;
  ...
  ...

let next () =
  ...
  T.parallel_for pool ~start:0 ~finish:(board_size - 1)
    ~body:(fun x ->
      for y = 0 to board_size - 1 do
        next_board.(x).(y) <- next_cell cur_board x y
      done);
  ...
  ...
```

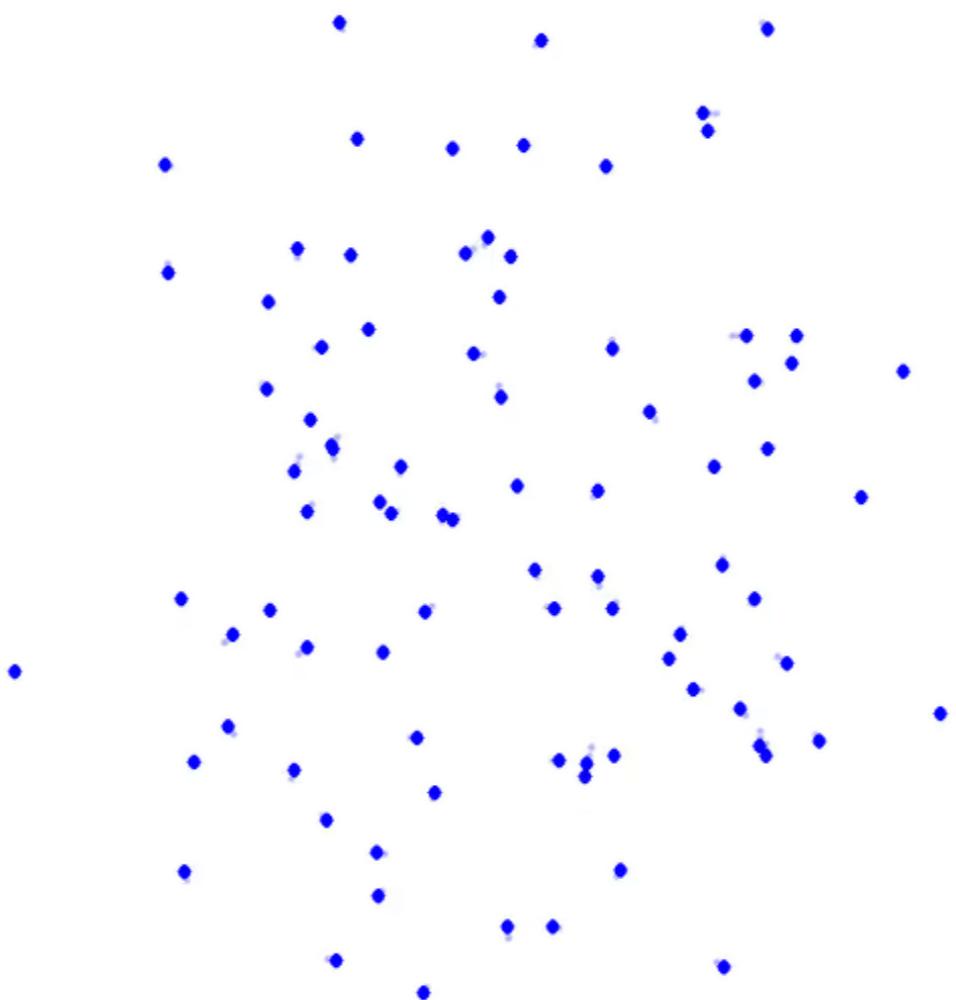
Performance: Game of Life

Board size = 1024, Iterations = 512

Cores	Time (Seconds)	Vs Serial	Vs Self
1	24.326	1	1
2	12.290	1.980	1.98
4	6.260	3.890	3.89
8	3.238	7.51	7.51
16	1.726	14.09	14.09
24	1.212	20.07	20.07

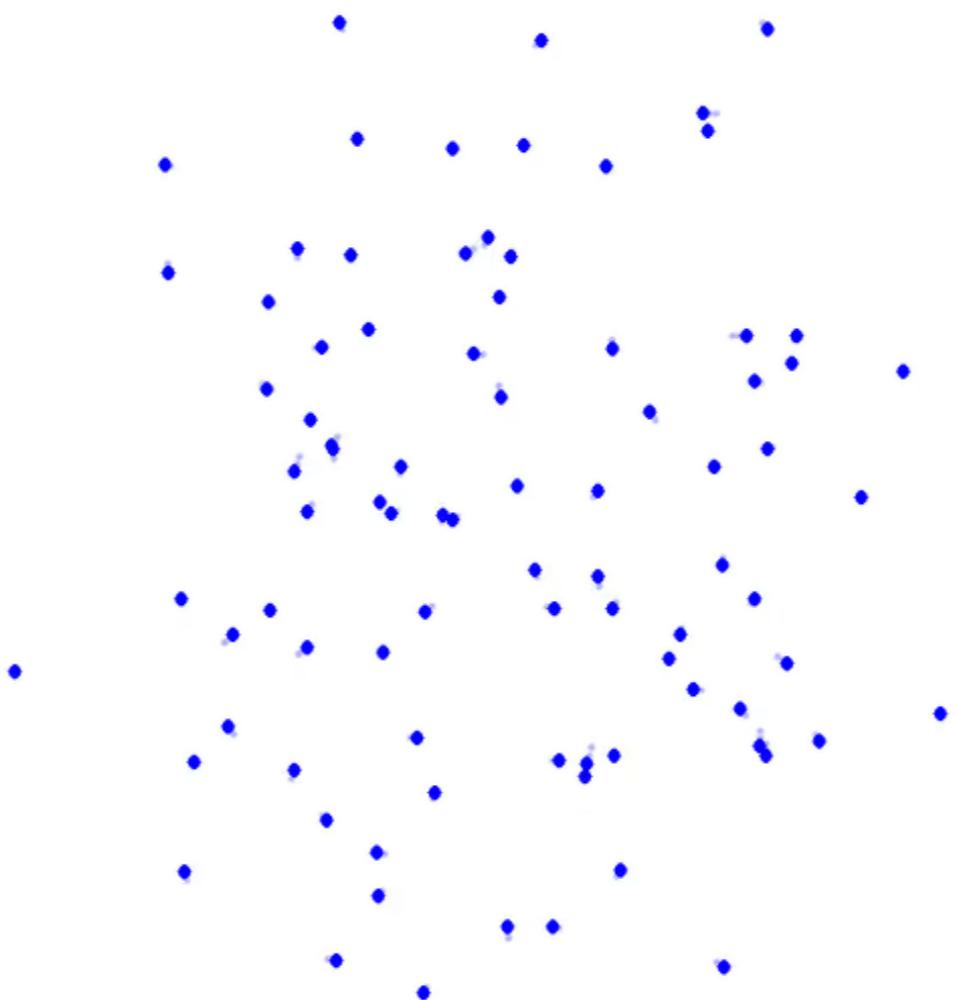
Performance Hacking: N-body

- Simulates the orbits of large number of astronomical objects
- Taken from computer language benchmarks game



Performance Hacking: N-body

- Simulates the orbits of large number of astronomical objects
- Taken from computer language benchmarks game



Sequential N-body

```
18 let advance bodies dt =
19  for i = 0 to Array.length bodies - 1 do
20    let b = bodies.(i) in
21    for j = 0 to Array.length bodies - 1 do
22      let b' = bodies.(j) in
23      if (i!=j) then begin
24        let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
25        let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
26        let mag = dt /. (dist2 *. sqrt(dist2)) in
27        b.vx <- b.vx -. dx *. b'.mass *. mag;
28        b.vy <- b.vy -. dy *. b'.mass *. mag;
29        b.vz <- b.vz -. dz *. b'.mass *. mag;
30      end
31    done
32  done;
```

Sequential N-body

```
18 let advance bodies dt =
19  for i = 0 to Array.length bodies - 1 do
20    let b = bodies.(i) in
21    for j = 0 to Array.length bodies - 1 do
22      let b' = bodies.(j) in
23      if (i!=j) then begin
24        let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
25        let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
26        let mag = dt /. (dist2 *. sqrt(dist2)) in
27        b.vx <- b.vx -. dx *. b'.mass *. mag;
28        b.vy <- b.vy -. dy *. b'.mass *. mag;
29        b.vz <- b.vz -. dz *. b'.mass *. mag;
30      end
31    done
32  done;
```

```
$ perf record ./nbody.exe
$ perf report
```

Sequential N-body

```
18 let advance bodies dt =
19  for i = 0 to Array.length bodies - 1 do
20    let b = bodies.(i) in
21    for j = 0 to Array.length bodies - 1 do
22      let b' = bodies.(j) in
23      if (i!=j) then begin
24        let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
25        let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
26        let mag = dt /. (dist2 *. sqrt(dist2)) in
27        b.vx <- b.vx -. dx *. b'.mass *. mag;
28        b.vy <- b.vy -. dy *. b'.mass *. mag;
29        b.vz <- b.vz -. dz *. b'.mass *. mag;
30      end
31    done
32  done;
```

```
$ perf record ./nbody.exe
$ perf report
```

Samples: 17K of event 'cycles:uppp', Event count (approx.): 12363186877			
Overhead	Command	Shared Object	Symbol
99.87%	nbody.exe	nbody.exe	[.] camlDune__exe__Nbody__advance_230
0.12%	nbody.exe	nbody.exe	[.] camlDune__exe__Nbody__energy_246
0.00%	nbody.exe	nbody.exe	[.] camlStdlib__domain__get_355
0.00%	nbody.exe	[kernel]	[k] 0xfffffffffb44009e7

Parallel N-Body

```
21 let advance pool bodies dt =
22   T.parallel_for pool
23     ~start:0
24     ~finish:(num_bodies - 1)
25     ~body:(fun i ->
26       let b = bodies.(i) in
27       for j = 0 to Array.length bodies - 1 do
28         let b' = bodies.(j) in
29         if (i!=j) then begin
30           let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
31           let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
32           let mag = dt /. (dist2 *. sqrt(dist2)) in
33           b.vx <- b.vx -. dx *. b'.mass *. mag;
34           b.vy <- b.vy -. dy *. b'.mass *. mag;
35           b.vz <- b.vz -. dz *. b'.mass *. mag;
36         end
37       done);
```

Parallel N-Body

```
21 let advance pool bodies dt =
22   T.parallel_for pool
23     ~start:0
24     ~finish:(num_bodies - 1)
25     ~body:(fun i ->
26       let b = bodies.(i) in
27       for j = 0 to Array.length bodies - 1 do
28         let b' = bodies.(j) in
29         if (i!=j) then begin
30           let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
31           let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
32           let mag = dt /. (dist2 *. sqrt(dist2)) in
33           b.vx <- b.vx -. dx *. b'.mass *. mag;
34           b.vy <- b.vy -. dy *. b'.mass *. mag;
35           b.vz <- b.vz -. dz *. b'.mass *. mag;
36         end
37       done);
```

- $\sim 5\times$ speedup on 8 cores compared to sequential version.

Parallel N-Body

```
21 let advance pool bodies dt =
22   T.parallel_for pool
23     ~start:0
24     ~finish:(num_bodies - 1)
25     ~body:(fun i ->
26       let b = bodies.(i) in
27       for j = 0 to Array.length bodies - 1 do
28         let b' = bodies.(j) in
29         if (i!=j) then begin
30           let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
31           let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
32           let mag = dt /. (dist2 *. sqrt(dist2)) in
33           b.vx <- b.vx -. dx *. b'.mass *. mag;
34           b.vy <- b.vy -. dy *. b'.mass *. mag;
35           b.vz <- b.vz -. dz *. b'.mass *. mag;
36         end
37       done);
```

- $\sim 5\times$ speedup on 8 cores compared to sequential version.
- *Can we do better?*

Parallel N-Body

```
21 let advance pool bodies dt =
22   T.parallel_for pool
23     ~start:0
24     ~finish:(num_bodies - 1)
25     ~body:(fun i ->
26       let b = bodies.(i) in
27       for j = 0 to Array.length bodies - 1 do
28         let b' = bodies.(j) in
29         if (i!=j) then begin
30           let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
31           let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
32           let mag = dt /. (dist2 *. sqrt(dist2)) in
33           b.vx <- b.vx -. dx *. b'.mass *. mag;
34           b.vy <- b.vy -. dy *. b'.mass *. mag;
35           b.vz <- b.vz -. dz *. b'.mass *. mag;
36         end
37       done);
```

- **~5X** speedup on 8 cores compared to sequential version.
- *Can we do better?*
- All the domains writing to the same shared array *bodies* in the inner loop
 - ◆ Leads to poor cache behaviour

Parallel N-Body (cache friendly)

```
21 let advance pool bodies dt =
22   T.parallel_for pool
23     ~start:0
24     ~finish:(num_bodies - 1)
25     ~body:(fun i ->
26       let b = bodies.(i) in
27       let vx, vy, vz = ref b.vx, ref b.vy, ref b.vz in
28       for j = 0 to Array.length bodies - 1 do
29         let b' = bodies.(j) in
30         if (i!=j) then begin
31           let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
32           let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
33           let mag = dt /. (dist2 *. sqrt(dist2)) in
34           vx := !vx -. dx *. b'.mass *. mag;
35           vy := !vy -. dy *. b'.mass *. mag;
36           vz := !vz -. dz *. b'.mass *. mag;
37         end
38       done;
39       b.vx <- !vx;
40       b.vy <- !vy;
41       b.vz <- !vz);
```

Parallel N-Body (cache friendly)

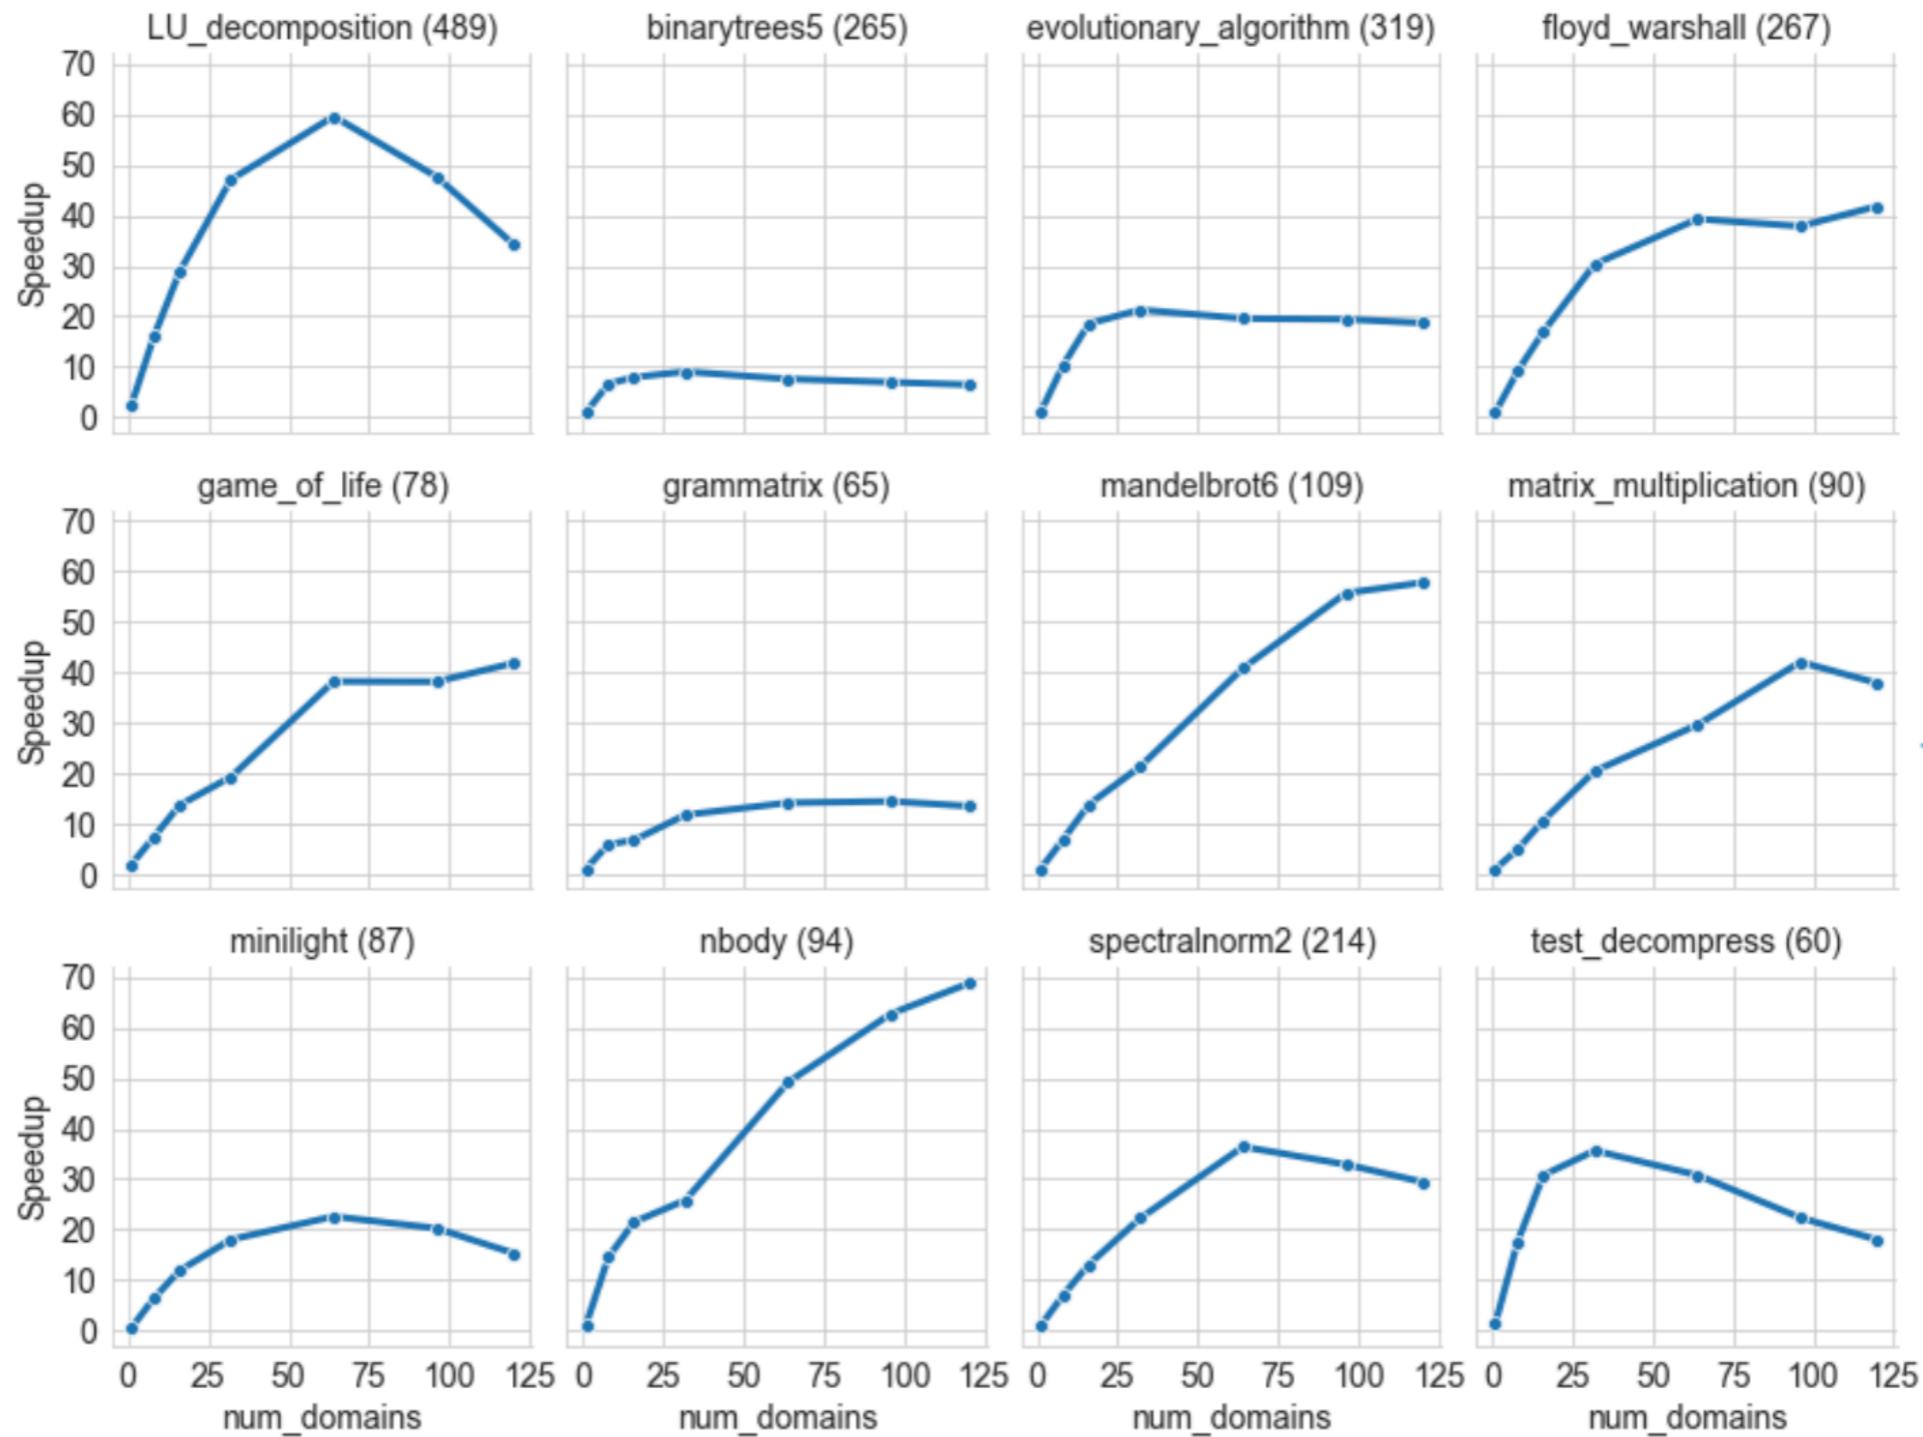
```
21 let advance pool bodies dt =
22   T.parallel_for pool
23     ~start:0
24     ~finish:(num_bodies - 1)
25     ~body:(fun i ->
26       let b = bodies.(i) in
27       let vx, vy, vz = ref b.vx, ref b.vy, ref b.vz in
28       for j = 0 to Array.length bodies - 1 do
29         let b' = bodies.(j) in
30         if (i!=j) then begin
31           let dx = b.x -. b'.x and dy = b.y -. b'.y and dz = b.z -. b'.z in
32           let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
33           let mag = dt /. (dist2 *. sqrt(dist2)) in
34           vx := !vx -. dx *. b'.mass *. mag;
35           vy := !vy -. dy *. b'.mass *. mag;
36           vz := !vz -. dz *. b'.mass *. mag;
37         end
38       done;
39       b.vx <- !vx;
40       b.vy <- !vy;
41       b.vz <- !vz);
```

- 22% faster than the unoptimised version

Parallel N-Body (cache friendly)

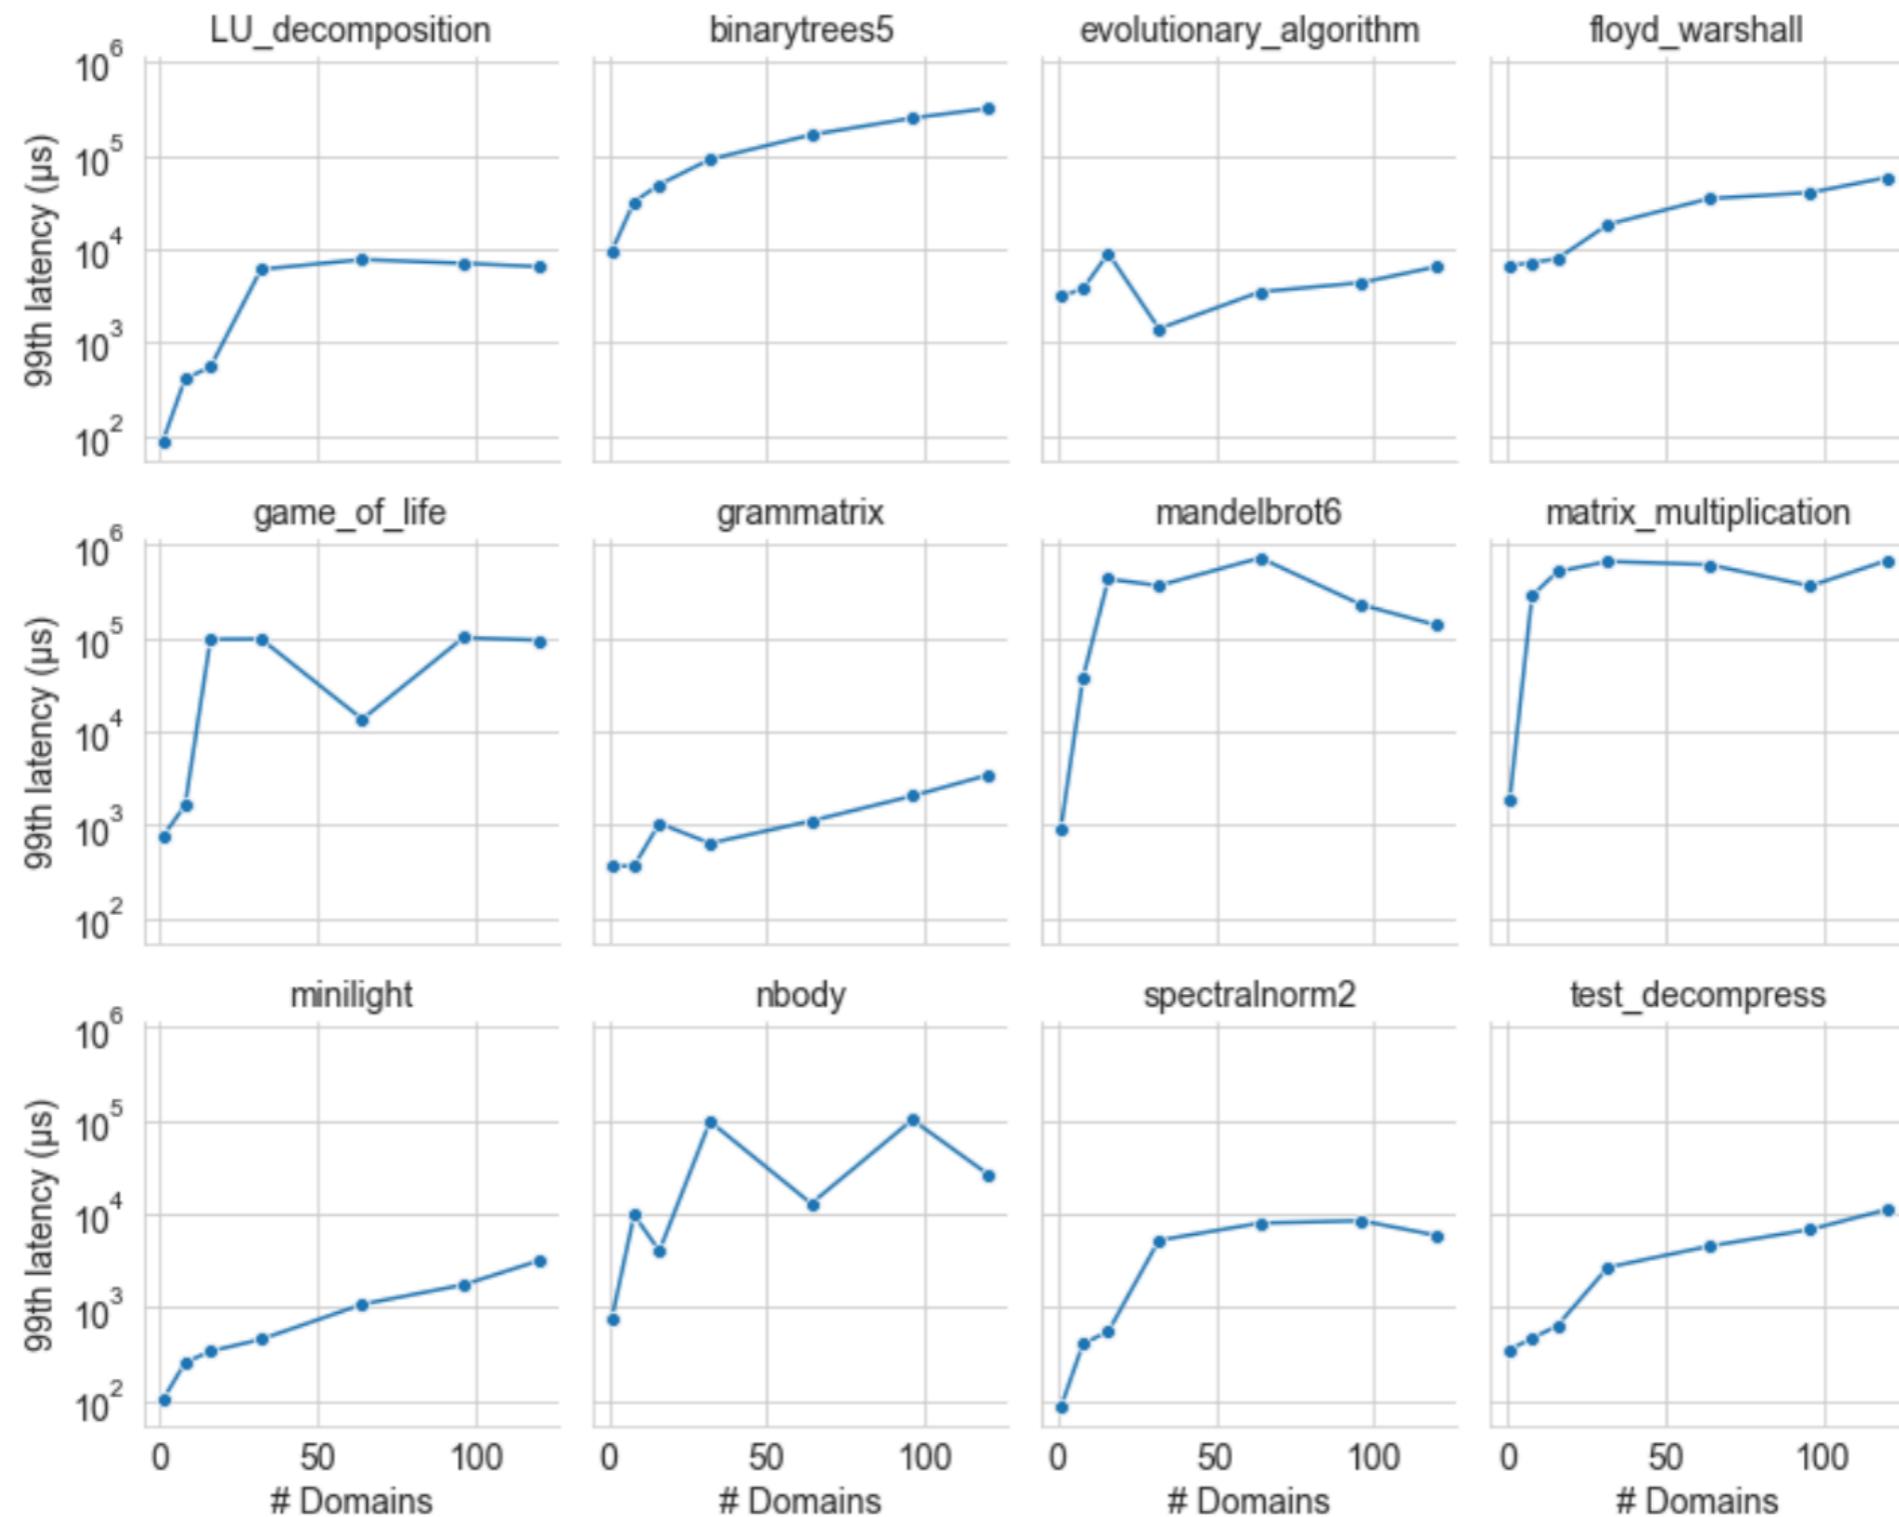
```
kc@kaveri:~/temp/nbody$ perf stat -e cycles,cycle_activity.cycles_no_execute,cycle_activity.stalls_mem_any,  
mem_load_uops_l3_miss_retired.remote_hitm ./_build/default/nbody_multicore.exe 8  
10745700105645904.000000000  
10745765317367518.000000000  
  
Performance counter stats for './_build/default/nbody_multicore.exe 8':  
  
16,919,777,849      cycles  
4,099,415,910      cycle_activity.cycles_no_execute  
3,864,941,761      cycle_activity.stalls_mem_any  
1,374,487          mem_load_uops_l3_miss_retired.remote_hitm  
  
0.861178677 seconds time elapsed  
  
kc@kaveri:~/temp/nbody$ perf stat -e cycles,cycle_activity.cycles_no_execute,cycle_activity.stalls_mem_any,  
mem_load_uops_l3_miss_retired.remote_hitm ./_build/default/nbody_multicore_optimised.exe 8  
10745700105645904.000000000  
10745765317367518.000000000  
  
Performance counter stats for './_build/default/nbody_multicore_optimised.exe 8':  
  
13,993,448,825      cycles  
2,926,815,349      cycle_activity.cycles_no_execute  
2,844,213,127      cycle_activity.stalls_mem_any  
409,388            mem_load_uops_l3_miss_retired.remote_hitm  
  
0.701971803 seconds time elapsed
```

Parallel Scalability



On 4 x 32-core AMD EPYC 7551

Parallel Latency



Summary

- Multicore OCaml brings shared-memory parallelism to OCaml
 - ◆ A new concurrent GC that optimises for *latency* and retains *throughput*

Summary

- Multicore OCaml brings shared-memory parallelism to OCaml
 - ◆ A new concurrent GC that optimises for *latency* and retains *throughput*
- *Domainslib* for high-level parallel programming
 - ◆ Async/await & parallel for loops
 - ◆ Scales well to large number (128) of cores

Summary

- Multicore OCaml brings shared-memory parallelism to OCaml
 - ◆ A new concurrent GC that optimises for *latency* and retains *throughput*
- *Domainslib* for high-level parallel programming
 - ◆ Async/await & parallel for loops
 - ◆ Scales well to large number (128) of cores
- OCaml 5.0 will be x86-64 only
 - ◆ Arm64 and Power after the 5.0 release
 - ◆ Linux, MacOS and Windows (mingw-w64 only)

Summary

- Multicore OCaml brings shared-memory parallelism to OCaml
 - ◆ A new concurrent GC that optimises for *latency* and retains *throughput*
- *Domainslib* for high-level parallel programming
 - ◆ Async/await & parallel for loops
 - ◆ Scales well to large number (128) of cores
- OCaml 5.0 will be x86-64 only
 - ◆ Arm64 and Power after the 5.0 release
 - ◆ Linux, MacOS and Windows (mingw-w64 only)
- Sivaramakrishnan et al, “*Retrofitting Parallelism onto OCaml*”, ICFP 2020