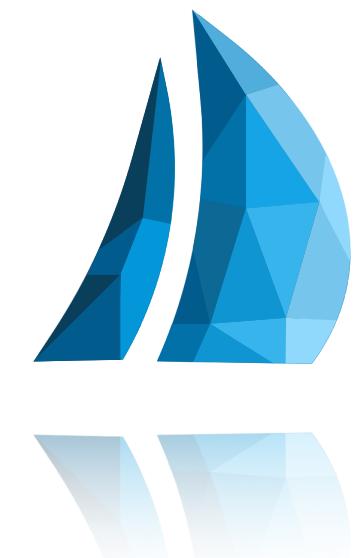


# Runtime Systems Research

KC Sivaramakrishnan ([kcsrk.info](http://kcsrk.info))

FSTTCS Mentoring workshop  
16th December 2025



Tarides

# Who am I – KC Sivaramakrishnan

- CS Prof at IIT Madras
  - Programming languages, formal verification and systems
- A core maintainer of the *OCaml* programming language
- CTO at Tarides
  - Building functional systems using *OCaml*
  - Maintainers of the OCaml compiler and platform tools

# Programming language research

- What is PL research?
- A bridge between **humans** and **machines**

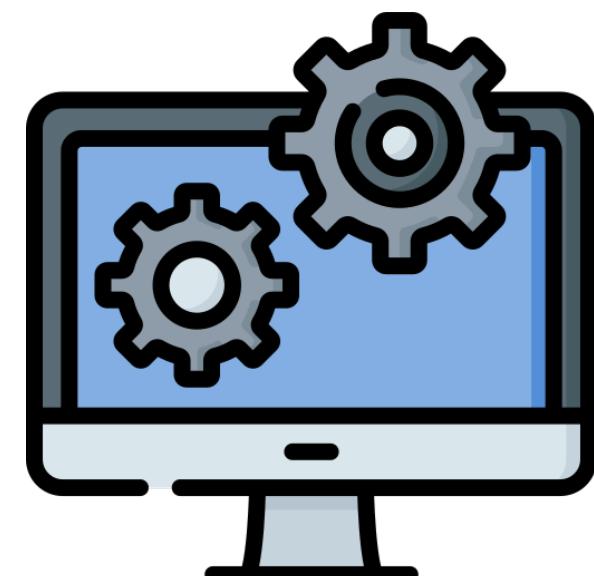
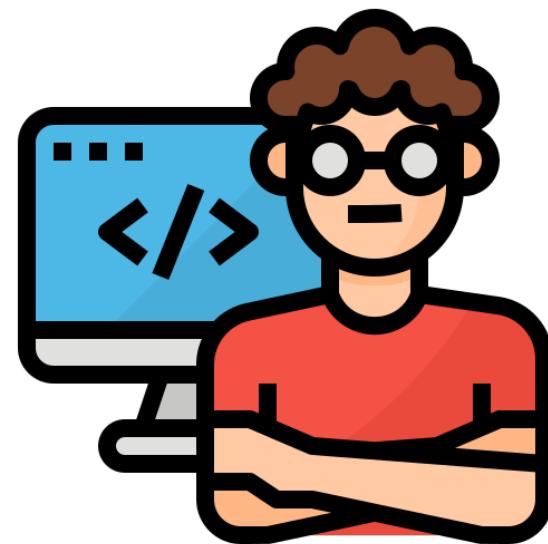
PL research is about helping people build *reliable software* by *designing abstractions and systems* that make *intent precise, mistakes harder to make and behaviour efficient but predictable.*

- Tools
  - Language design (Expression), Formal semantics (Meaning), Type systems & Verification (Guarantees), Compilers & Runtime systems (Execution)

# Runtime Systems Research

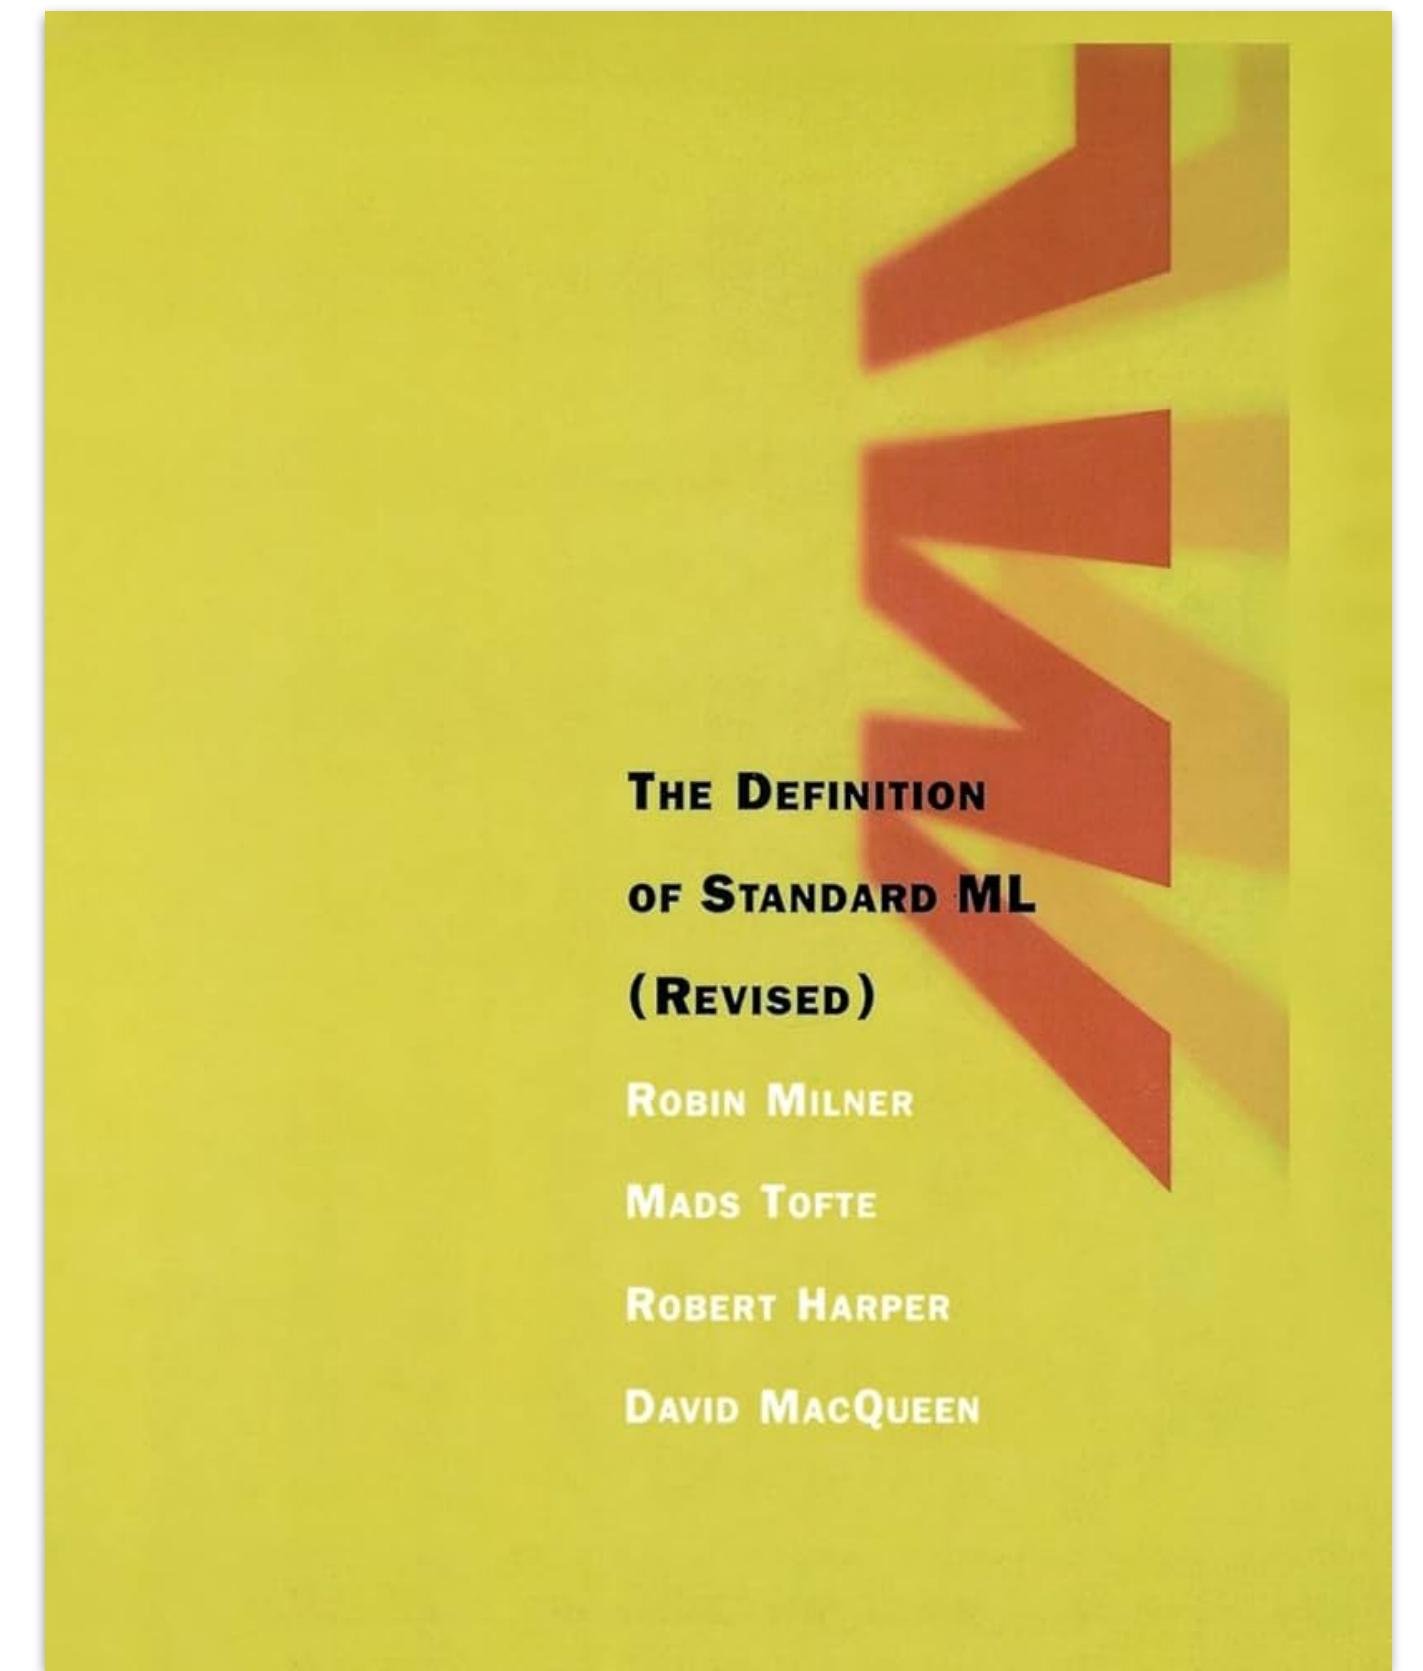
*How do we make the **promises** of a language hold in **reality**?*

- Themes
  - Memory management, JIT, Concurrency, Performance, etc.
  - Overlaps with Networks, Storage, OS research
- Runtime systems sit at the boundary between **beautiful ideas** and **messy reality**
  - .... and so does a researcher's career.
- Goal of the talk
  - Insight into runtime systems research through my own journey



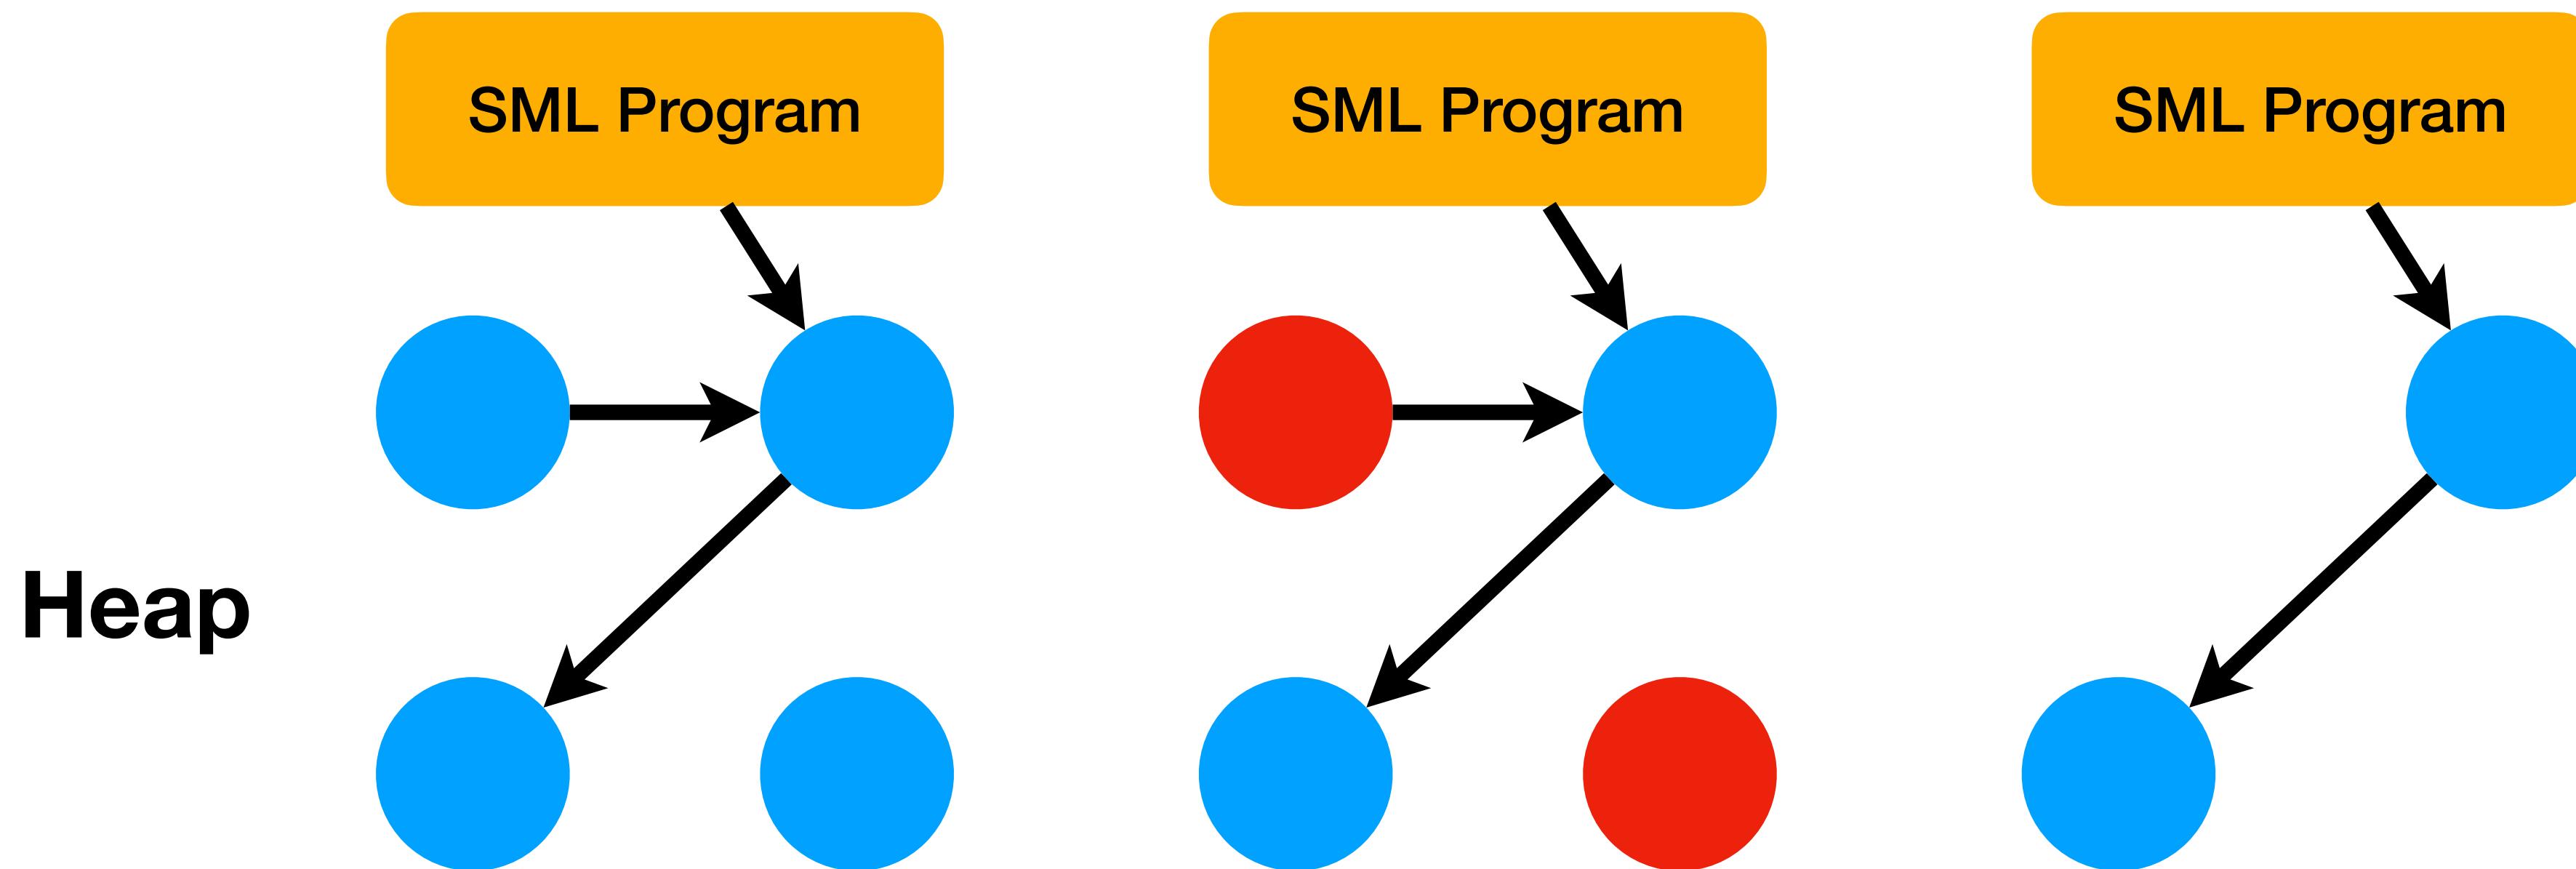
# MultiMLton

- PhD starting project @ Purdue University
- MultiMLton – a multicore-aware extension of MLton Standard ML compiler
  - **Standard ML:** a rigorously specified FP language in ML family
  - **MLton:** a whole-program optimising compiler for the Standard ML language
- Extend this to take advantage of multi-core
  - Language design ← Lukasz Ziarek
  - Parallelism support in the runtime ← me



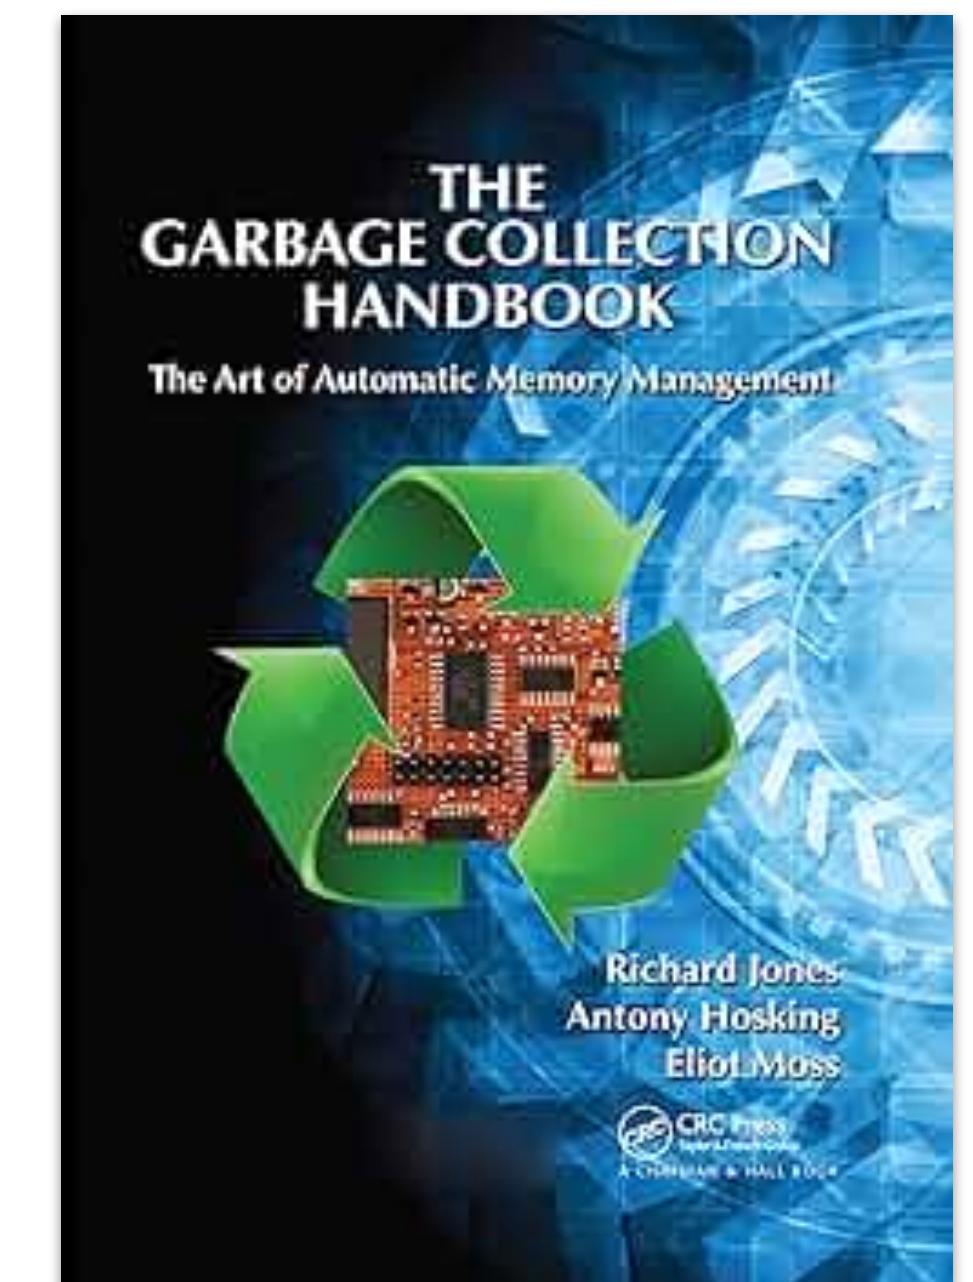
# Parallelism support in the runtime

- Automatic memory management with a garbage collector (GC)

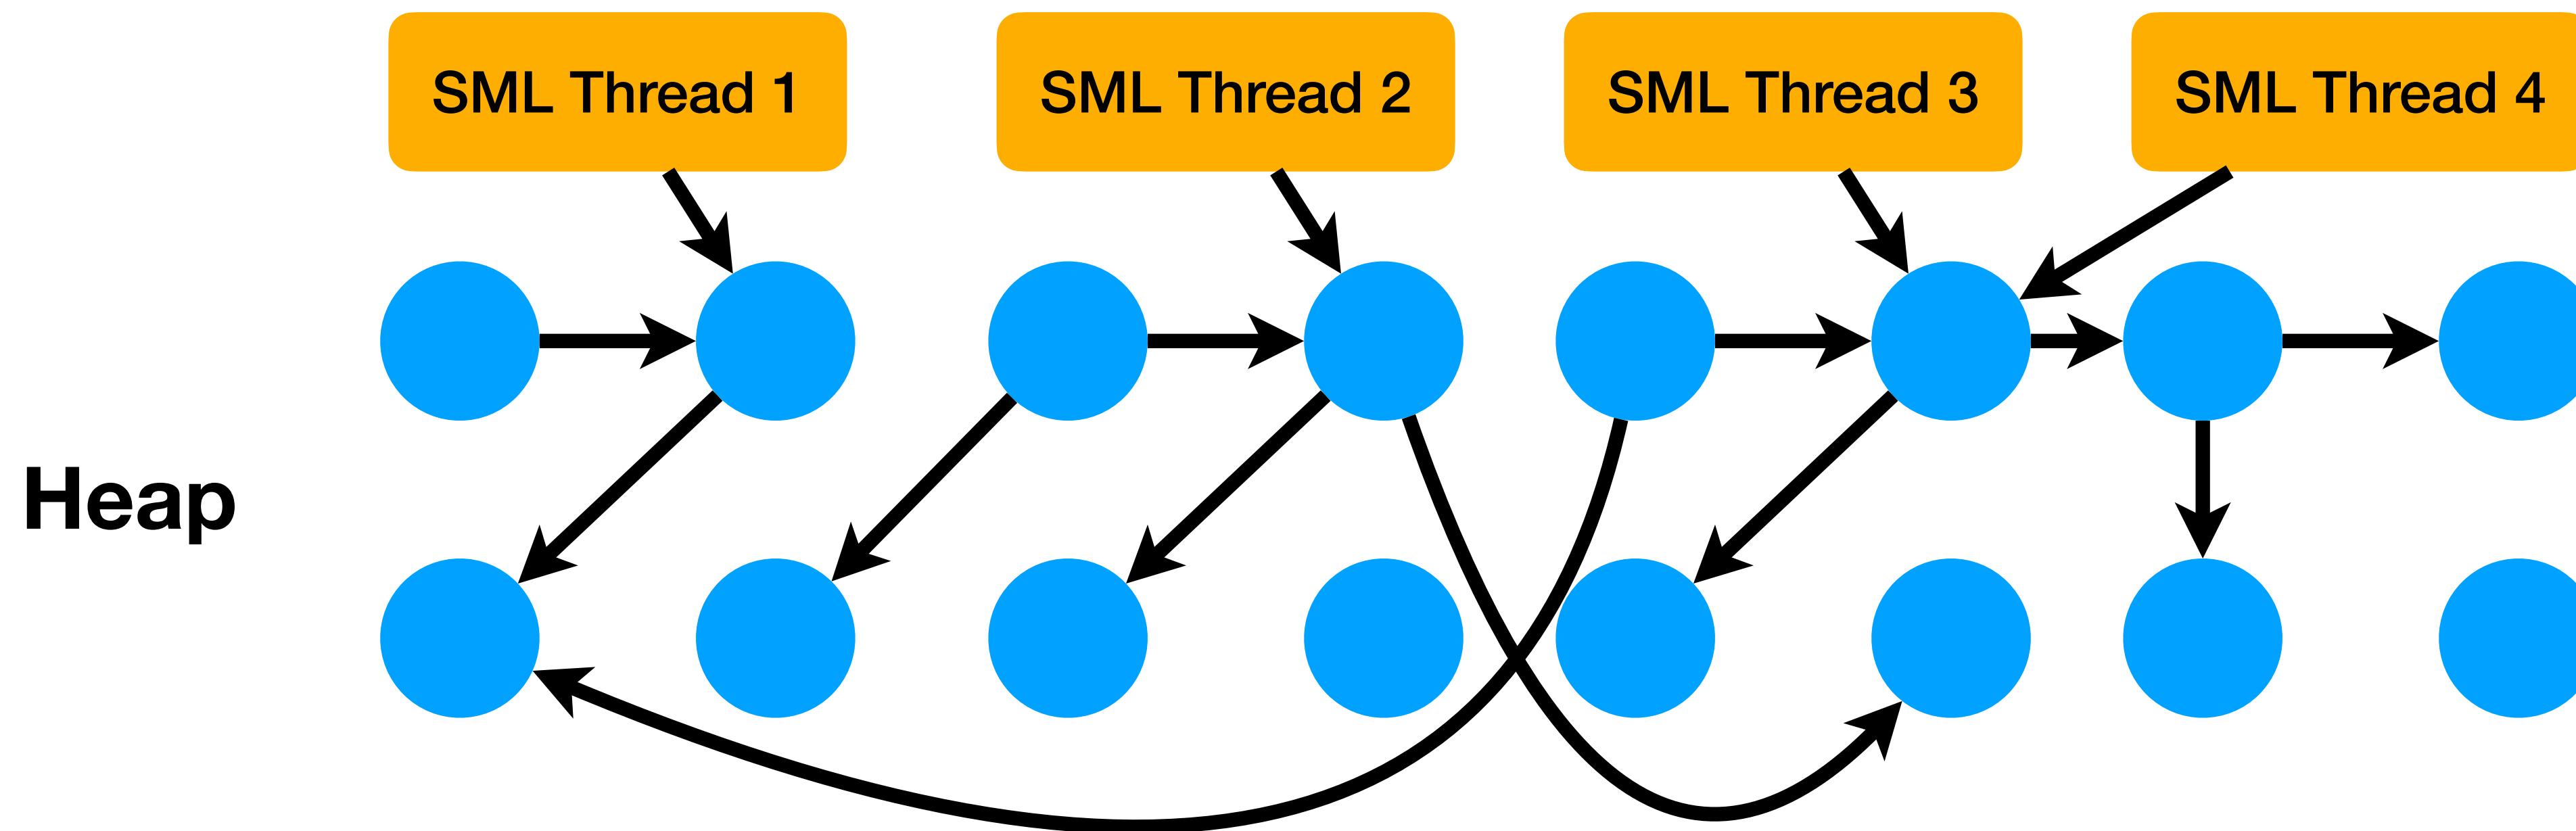


- No “perfect” GC
  - Trade-offs – throughput, latency, memory usage, complexity,...

... also Java, .NET,  
Python, Go,  
JavaScript, OCaml,  
Haskell, JavaScript,  
etc.



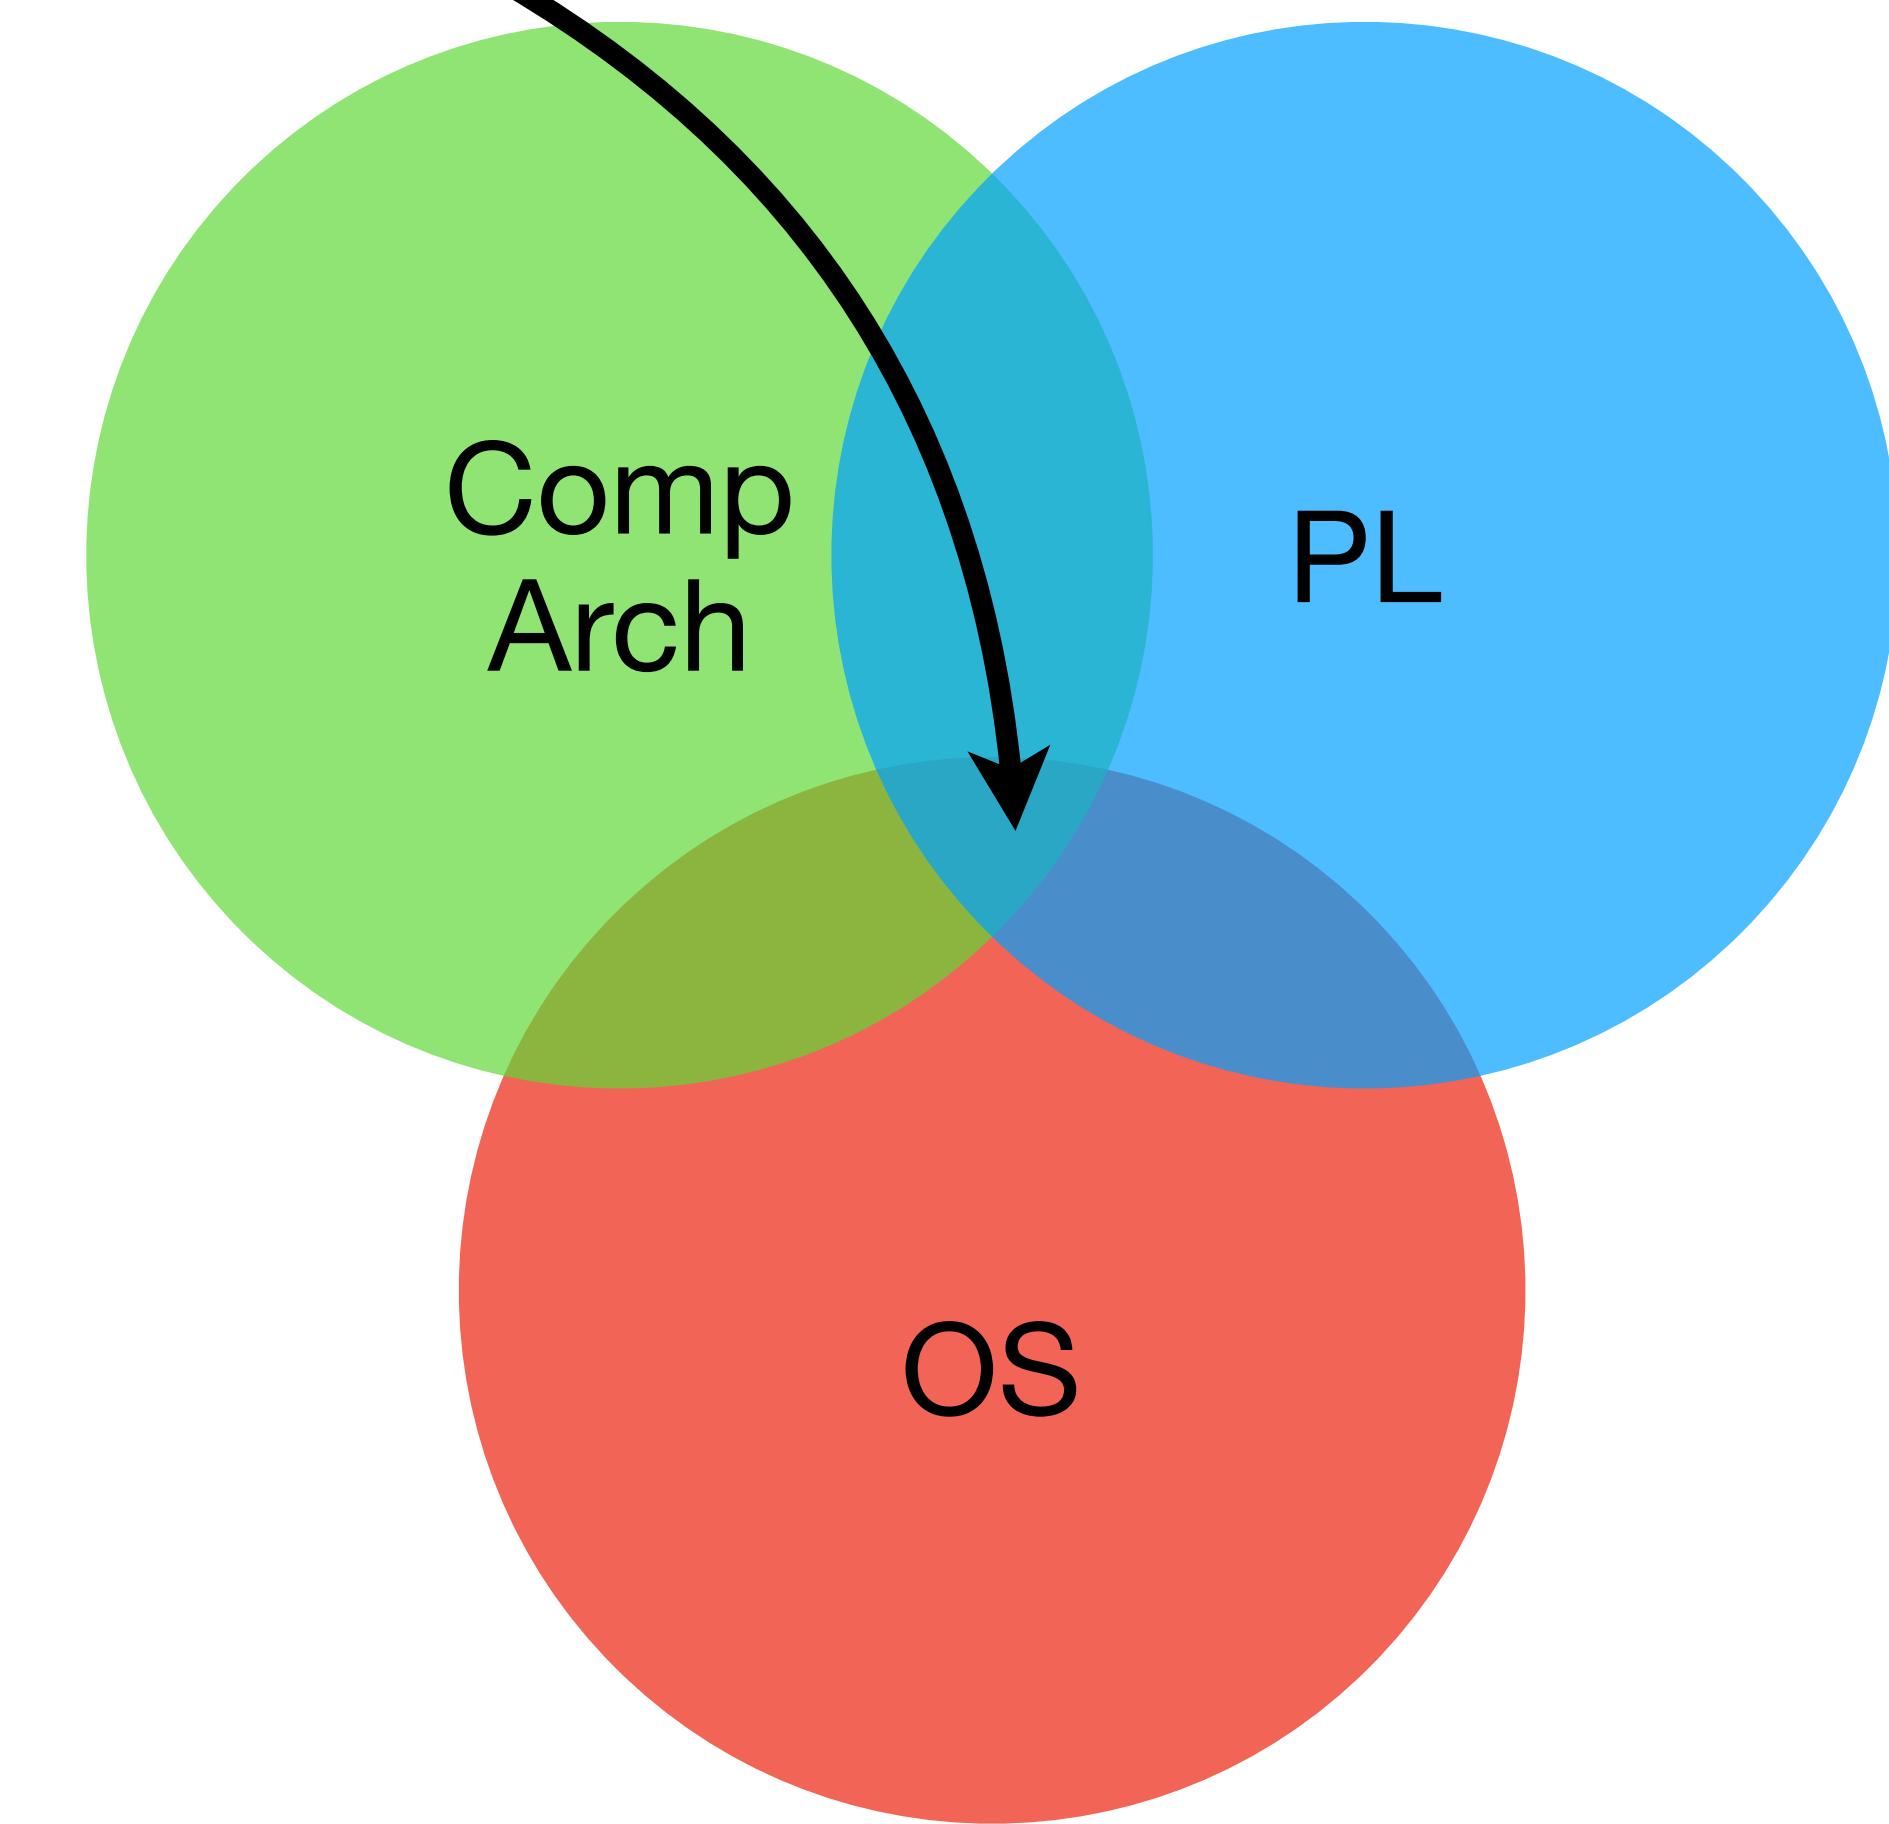
# Parallel Garbage Collector



- Should we ***stop all*** the SML threads?
- Should we collect all the garbage at once or ***incrementally***?
- Should the GC run ***concurrently*** with the SML threads?
- How do the GC ***interact*** with the OS and (micro-)architectural quirks?

# GC research area —

- Quite challenging/rewarding
- Many **open problems**, especially as the compute stack gets heterogeneous and distributed
- **Impactful** – any improvement can benefit all users of the language
- **But...**
  - Requires massive, long-term engineering effort
    - Bugs are rare, timing-dependent, and hard to reproduce
    - Research ideas are tightly entangled with infrastructure
  - Many industrial-strength GCs; “Solved” problem?
  - Easy to get stuck building; hard to know when to publish



# PhD Research

## Partial Memoization of Concurrency and Communication

Lukasz Ziarek

KC Sivaramakrishnan

Suresh Jagannathan

Department of Computer Science

Purdue University

{lziarek, chandras, suresh}@cs.purdue.edu

**ICFP 2009**

tract

ization is a well-known optimization technique used to eliminate redundant calls for pure functions. If a call to a function  $f$  with argument  $v$  yields result  $r$ , a subsequent call to  $f$  with  $v$  can be

### 1. Introduction

Eliminating redundant computation is an important optimization supported by many language implementations. One instance of this optimization class is memoization.

Synchronous communication is an important feature of concurrent systems, building *composable* abstractions

## Composable Asynchronous Events

Lukasz Ziarek, KC Sivaramakrishnan, Suresh Jagannathan

Purdue University

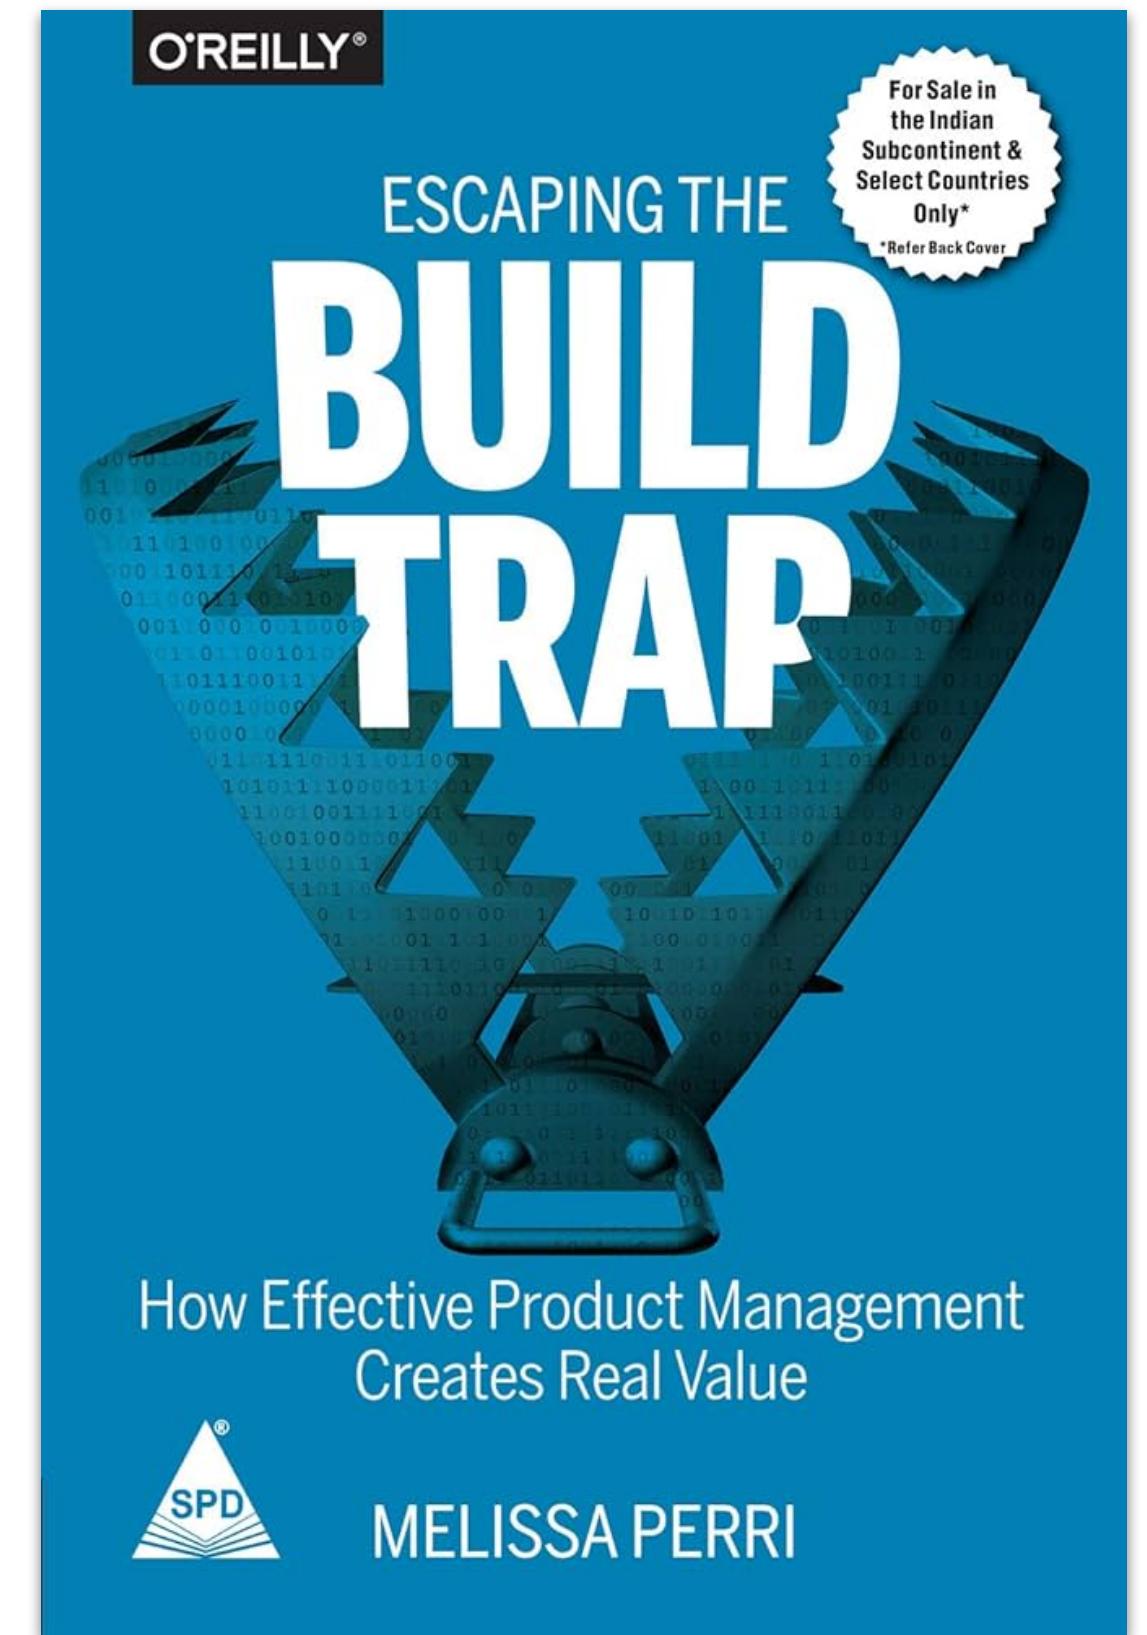
{lziarek, chandras, suresh}@cs.purdue.edu

**PLDI 2011**

- Started my PhD in 2008
- I was supporting language research projects in my group
- ***What was my research goal?***
  - Build Trap!

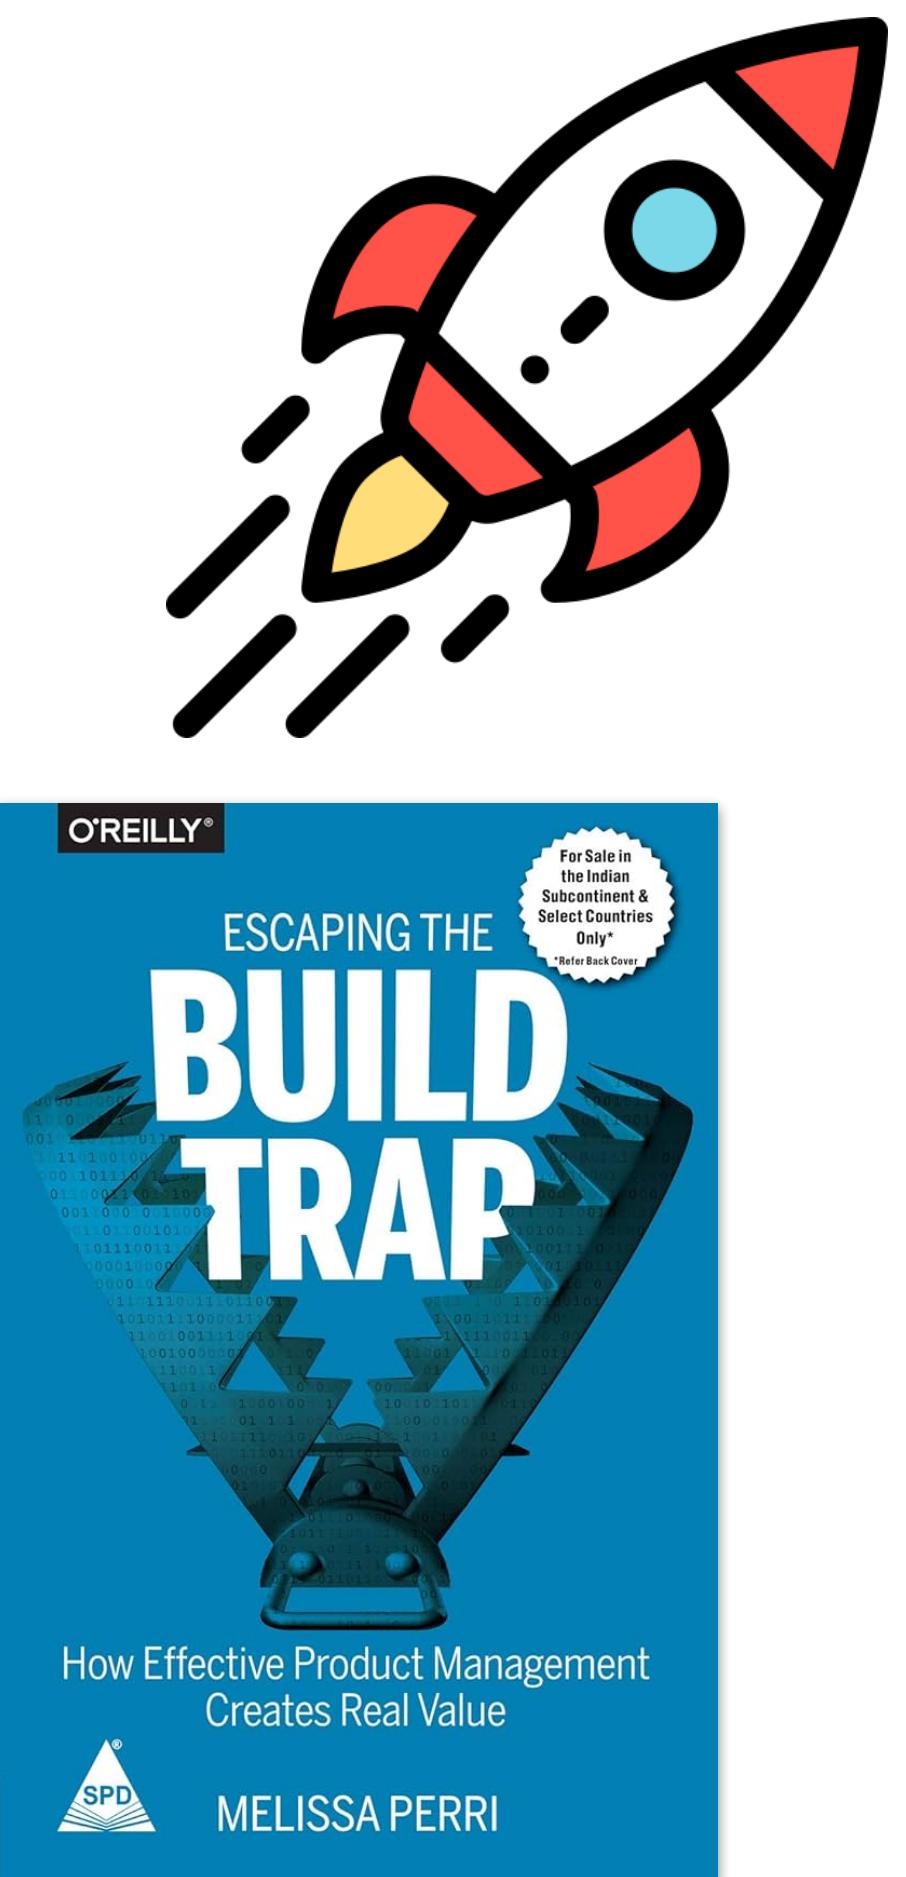
# Build Trap!

- ***Keep building, value will emerge on its own!***
- Companies falling into measuring success
  - by **outputs** — features shipped, code written, systems built
  - Instead of **outcomes** — customer value, impact, learning
- For runtime systems research
  - **Outputs** — code written, systems built, bugs fixed
  - **Outcomes** — improving the state of the art ⇒ ***publishing papers***
- Risks in systems research
  - High-upfront cost, complexity of supporting infra, performance rabbit holes, insufficient checkpoints



# Escaping the build trap

- **Start with the research question/hypothesis that you want to test**
  - Starting to build without one is a definite path to the build trap
- **Treat the system as an instrument, not the product**
  - “Avoid Success at All Costs”, SPJ about GHC
  - Research PL — Koka, Effekt, Links, Flix, Hazel...
- **Force early articulation of the paper**
  - If you can’t write the introduction, you don’t yet know why you’re building
  - SPJ, “Writing a research paper is the way to do research”
- **Stop when the marginal build effort doesn’t sharpen the claim**
  - Engineering progress feels tangible, but insight is fuzzier
  - Reviewers reward the latter!



# Did escape the build trap eventually 😅

## Eliminating Read Barriers through Procrastination and Cleanliness

KC Sivaramakrishnan

Lukasz Ziarek

Suresh Jagannathan

Purdue University

{chandras, lziarek, suresh}@cs.purdue.edu

**ISMM 2012**

*Using pervasive concurrency in the language  
to trade off GC overheads*

- Not Core A\* (😅) PL systems conferences – POPL, PLDI, ICFP, OOPSLA, ASPLOS...
  - ... and that was ok, I had fun doing this research

## *MultiMLton: A multicore-aware runtime for standard ML*

K.C. SIVARAMAKRISHNAN

*Purdue University, West Lafayette, IN, USA*  
(e-mail: chandras@purdue.edu)

LUKASZ ZIAREK

*SUNY Buffalo, NY, USA*  
(e-mail: lziarek@buffalo.edu)

SURESH JAGANNATHAN

*Purdue University, West Lafayette, IN, USA*  
(e-mail suresh@cs.purdue.edu)

**JFP 2014**

*Complete language and runtime system;  
evaluation on 768-core behemoth*

# Finding my research focus



- PL research takes a long time to mature
  - GC – 60s (Lisp), 00s (Java)
  - Strong static typing – 70s (ML), 00s (Java), 10s (TypeScript)
  - Concurrency – Research 80s, 10s (Go, Rust)
  - Static memory safety – Research 90s, 10s (Rust)

# Finding my research focus



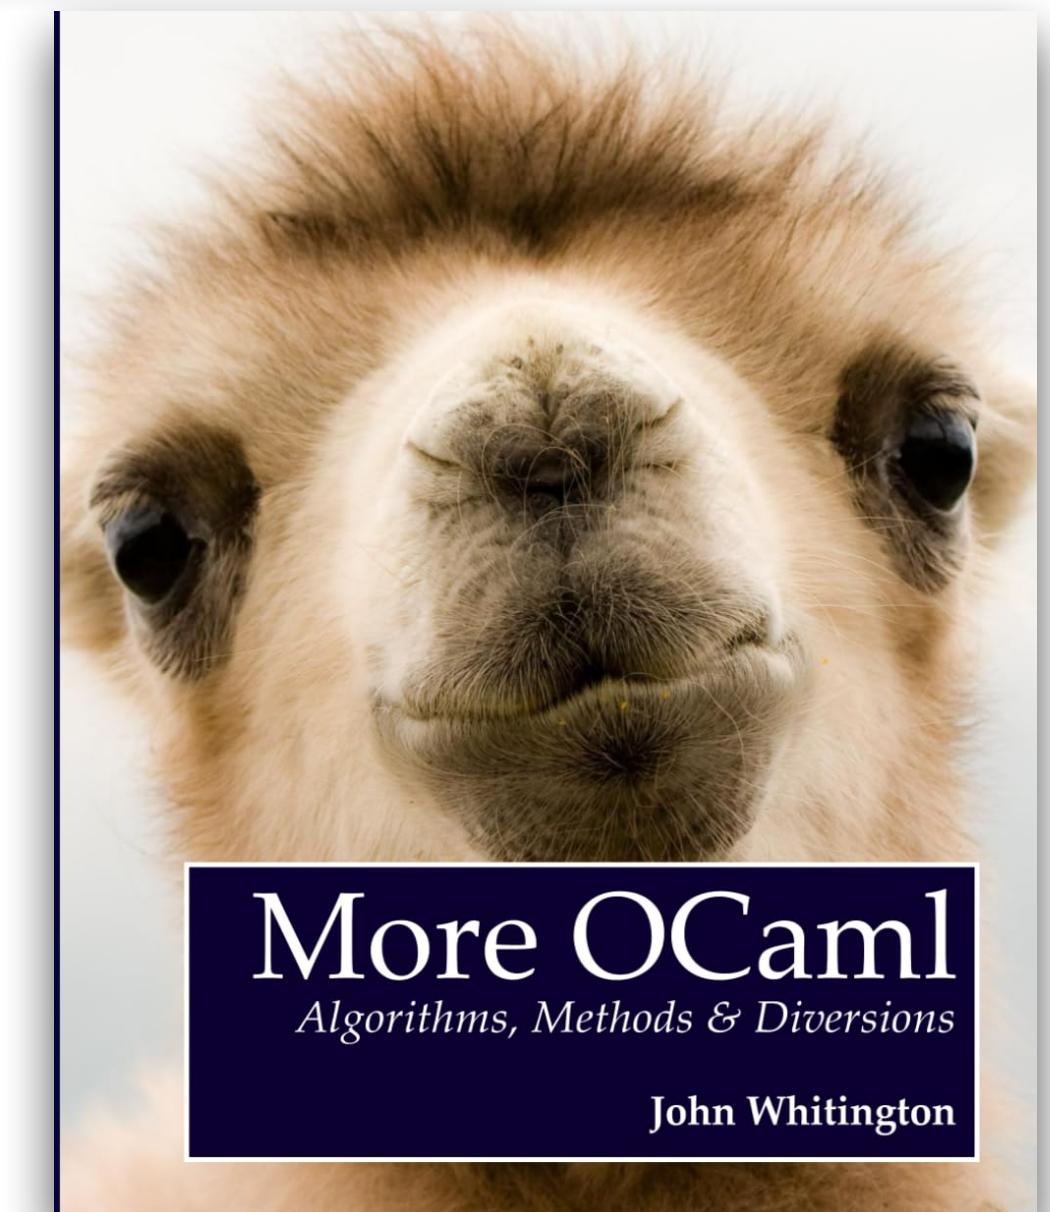
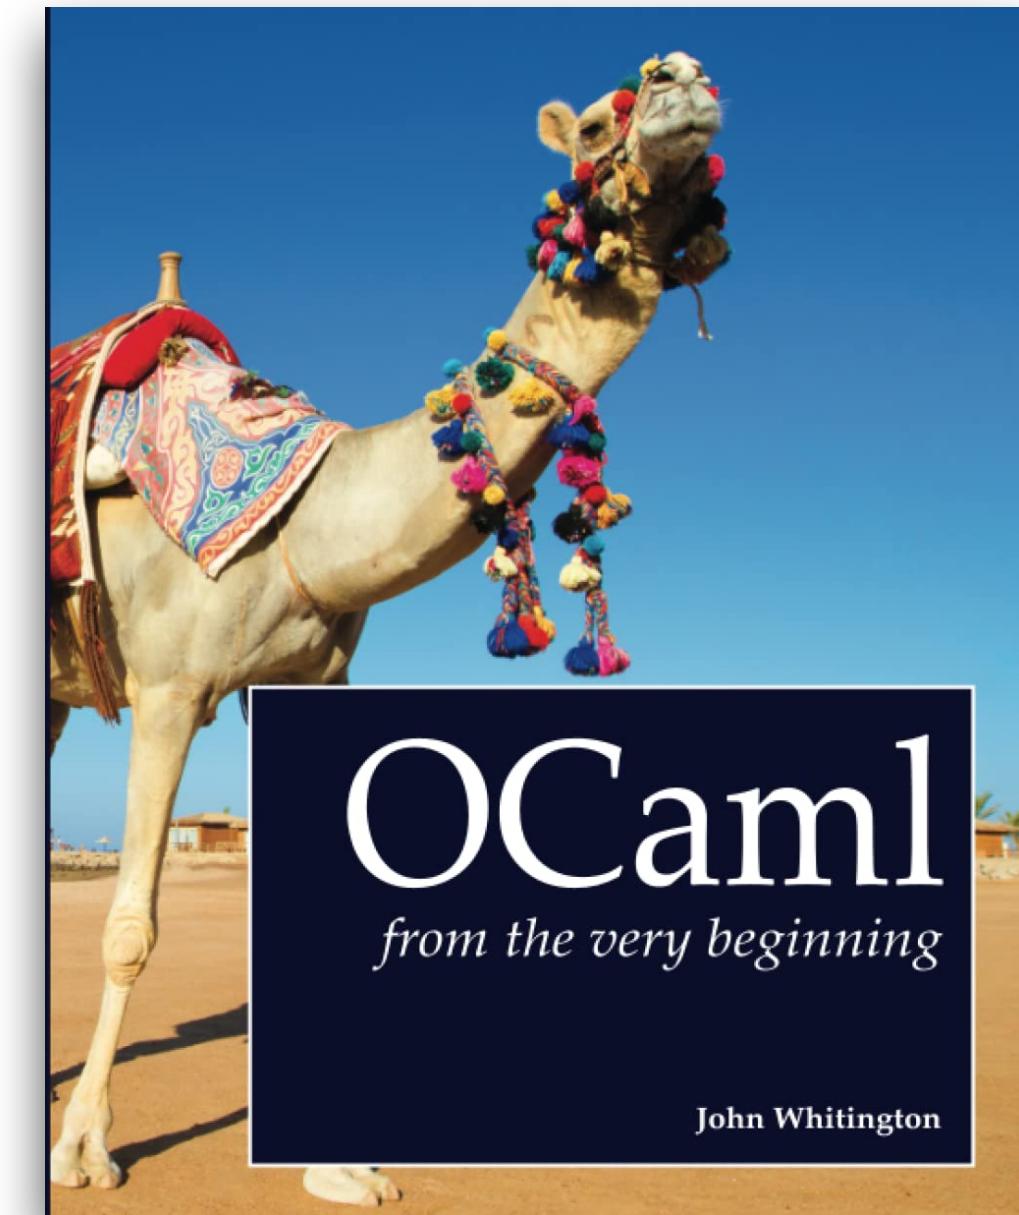
- Loved being in the area of *translating research to practice*
  - Runtime systems are naturally amenable to this
- Have an *impact* outside of research papers
  - Benefit “real” users, not “imagined” ones
  - *MultiMLton hasn’t been developed since 2014*
- Like the *academic freedom* to move about in the spectrum

# Multicore OCaml

- Post-doc @ U Cambridge
- **Multicore OCaml** — native support for concurrency and parallelism to the OCaml programming language



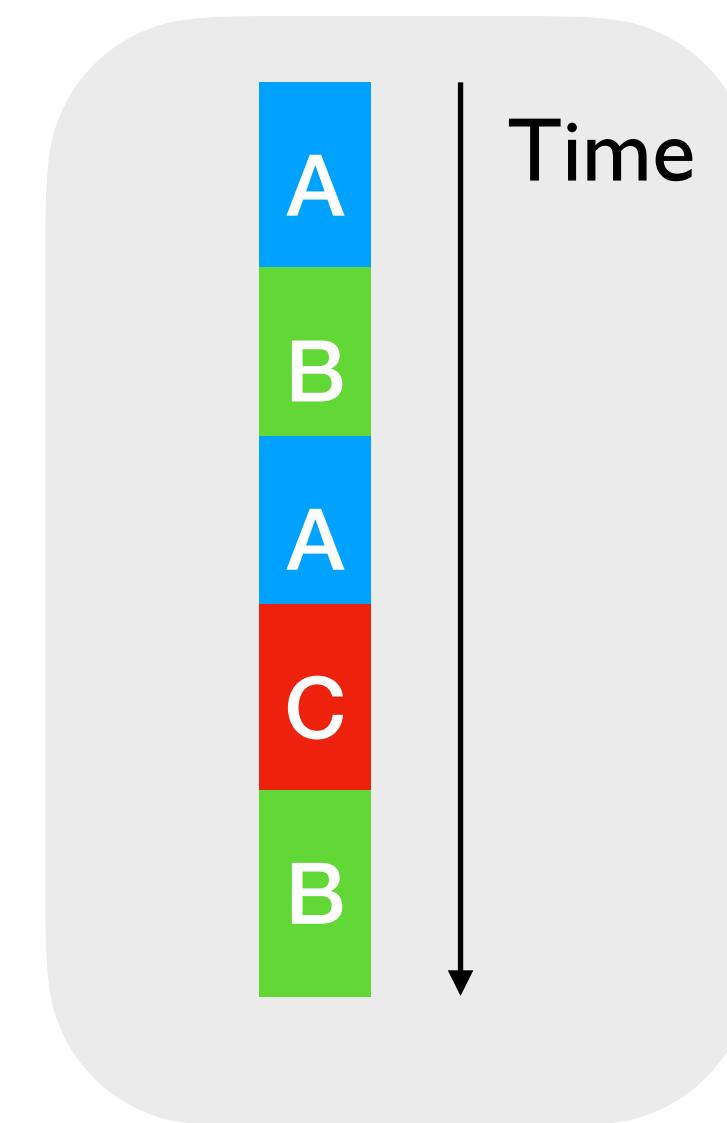
- A functional-first programming language in the ML family
- Projects — Rocq, Frama-C, Why3, F\*
- Industrial Users — Jane Street, Meta, ARM, SemGrep, Microsoft
- ***Still sequential in 2014***
- *Promise of translating learning from MultiMLton to a widely-used language*



# Multicore OCaml

- Native support for concurrency and parallelism to OCaml

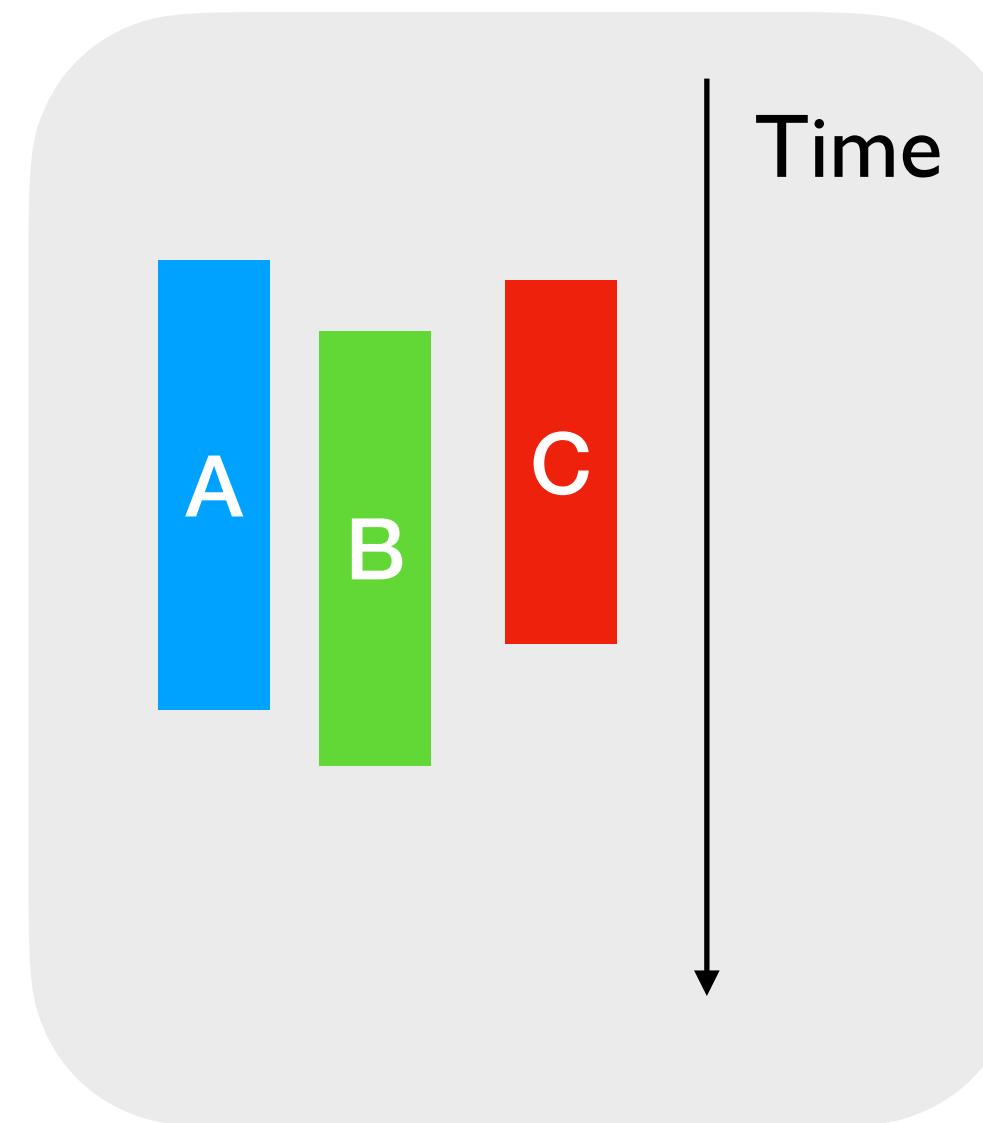
*Concurrency*



*Interleaved  
execution*

*Effect Handlers*

*Parallelism*



*Simultaneous  
execution*

*Domains*

# Challenges

- A new multicore garbage collector and multicore runtime system
  - Replacing a car engine with a new one!
- Make the language itself thread-safe
  - OCaml is a safe language! (Unlike C/C++, Go)
- Maintain feature and performance backwards compatibility!
  - Most OCaml programs will continue to remain single-threaded



**XKCD published in 2014**  
Today, *bird recognition* is a commodity ML task.

# Research Focus

- The goal was **upstreaming** multicore features to OCaml
  - Publishing papers is a means to **build credibility** for upstreaming
  - Conscious tradeoff to have an impact beyond papers
- Building in the open
  - Liberally licensed open-source software
  - Quality >>> research-prototypes, < production (...initially)

Lindsey Kuper

gasche 11 years ago

I think the only reasonable definition of "research" is "what researchers do" (researchers being defined, of course, as "those that research"). All the things listed in your "Instead," sentence \*are\* research in my book, and even those posts "that weren't research" could solidly be argued to be part of some people's research (including at least you, apparently). For example, your interesting discussions about Hacker School count as "diffusion", helping other people to appreciate your discipline, and we have a checkbox about that in my lab's yearly Activity Report.

# Starting out

## Multicore OCaml

Stephen Dolan

Leo White

Anil Madhavapeddy

*Currently, threading is supported in OCaml only by means of a global lock, allowing at most one thread to run OCaml code at any time. We present ongoing work to design and implement an OCaml runtime capable of shared-memory parallelism.*

### 1 Introduction

Adding shared-memory parallelism to an existing lan-

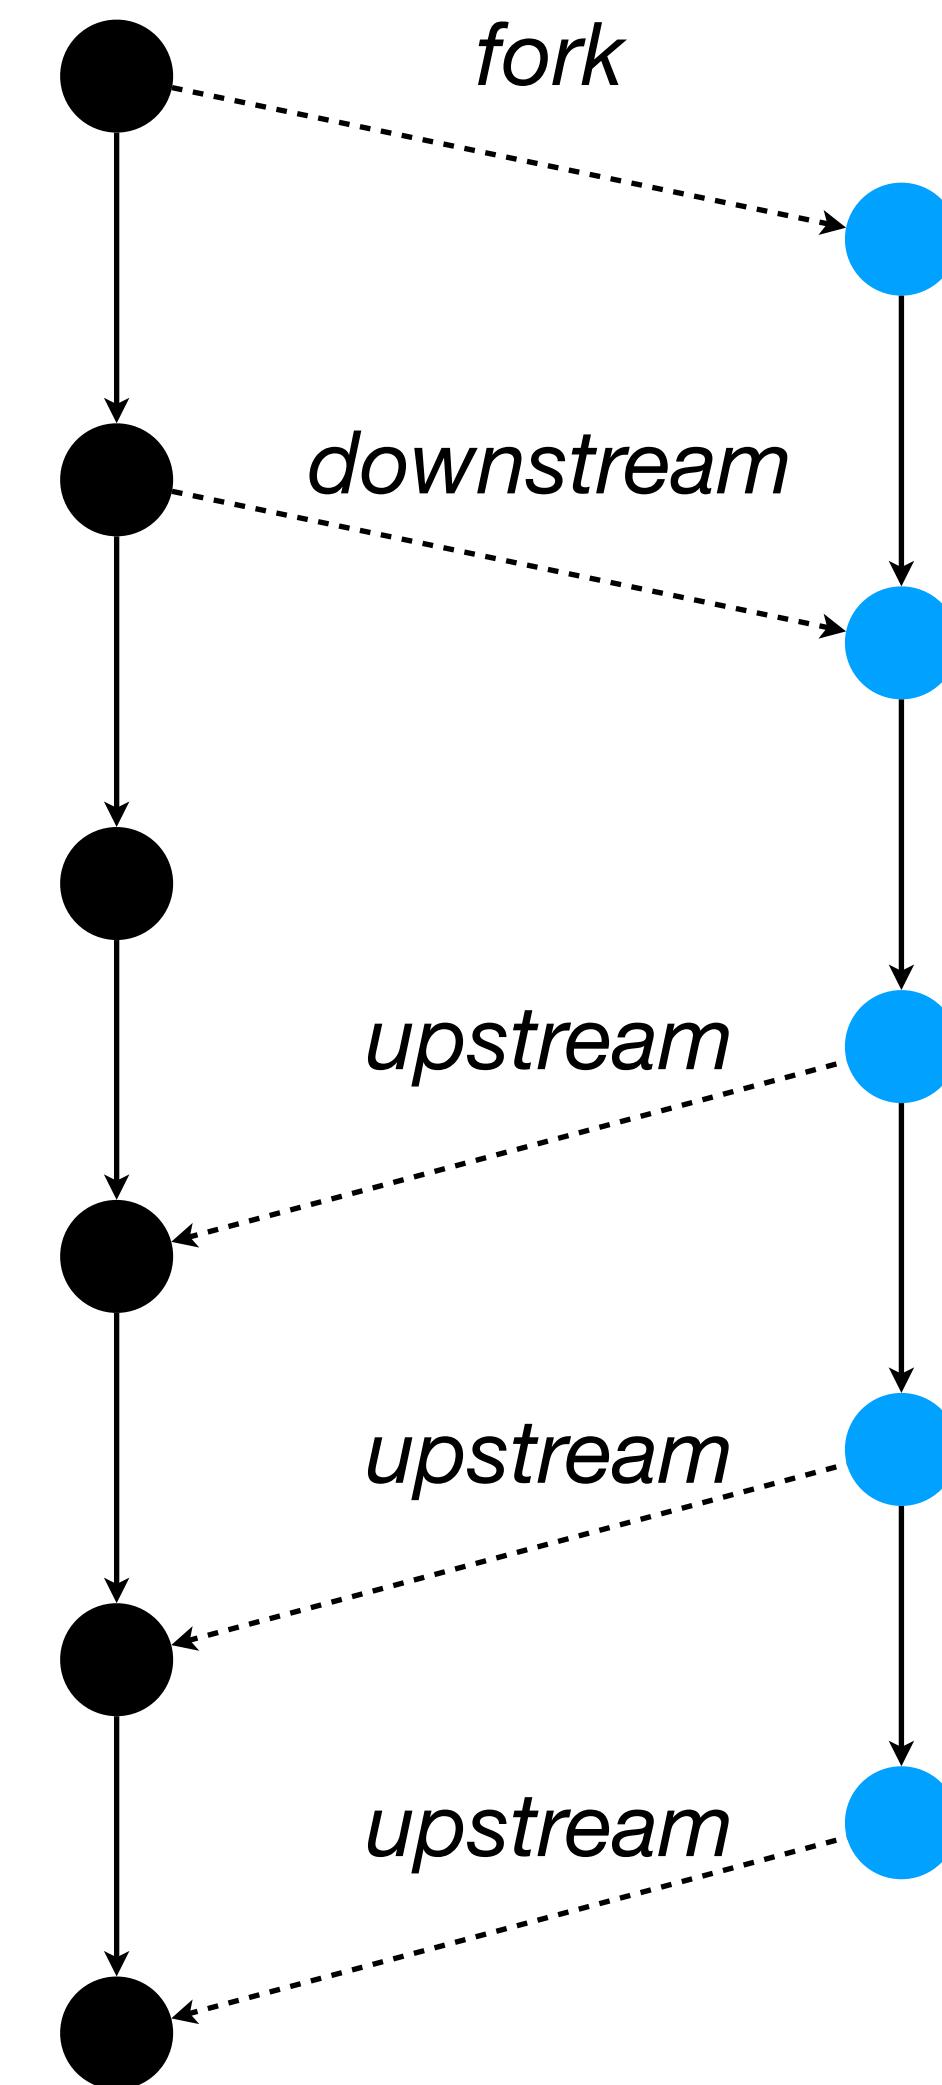
all objects reachable from it to be promoted to the shared heap en masse. Unfortunately this eagerly promotes many objects that were never really shared: just because an object is pointed to by a shared object does not mean another thread is actually going to attempt to access it.

Our design is similar but lazier, along the lines of the multicore Haskell work [2], where objects are promoted to the shared heap whenever another thread

**OCaml Workshop 2014**

## Upstream OCaml

## Multicore OCaml



# Building confidence through papers

Multicore GC and  
runtime system

## Retrofitting Parallelism onto OCaml

ICFP 2020 🏆

## Bounding Data Races in Space and Time

(Extended version, with appendices)

PLDI 2018

Relaxed Memory  
Model

Concurrency  
story

KC  
STE  
LEC  
SAD  
TO/  
AND  
SUB  
ATU  
ANI  
OCa  
OCa

Uni  
**Abstract**  
We propose  
programs t  
of data rac  
antees that  
memory parallel pro  
f

KC Sivaramakrishnan  
IIT Madras  
Chennai, India  
kcsrk@cse.iitm.ac.in

Tom Kelly  
OCaml Labs  
Cambridge, UK  
tom.kelly@cantab.net

### Abstract

Effect handlers have been gathering momentum as a mechanism for modular programming with user-defined effects.

Stephen Dolan  
OCaml Labs  
Cambridge, UK  
stephen.dolan@cl.cam.ac.uk

Sadiq Jaffer  
Opsian and OCaml Labs  
Cambridge, UK  
sadiq@toao.com

### 1 Introduction

Effect handlers [45] provide a modular foundation for user-defined effects. The key idea is to separate the definition of

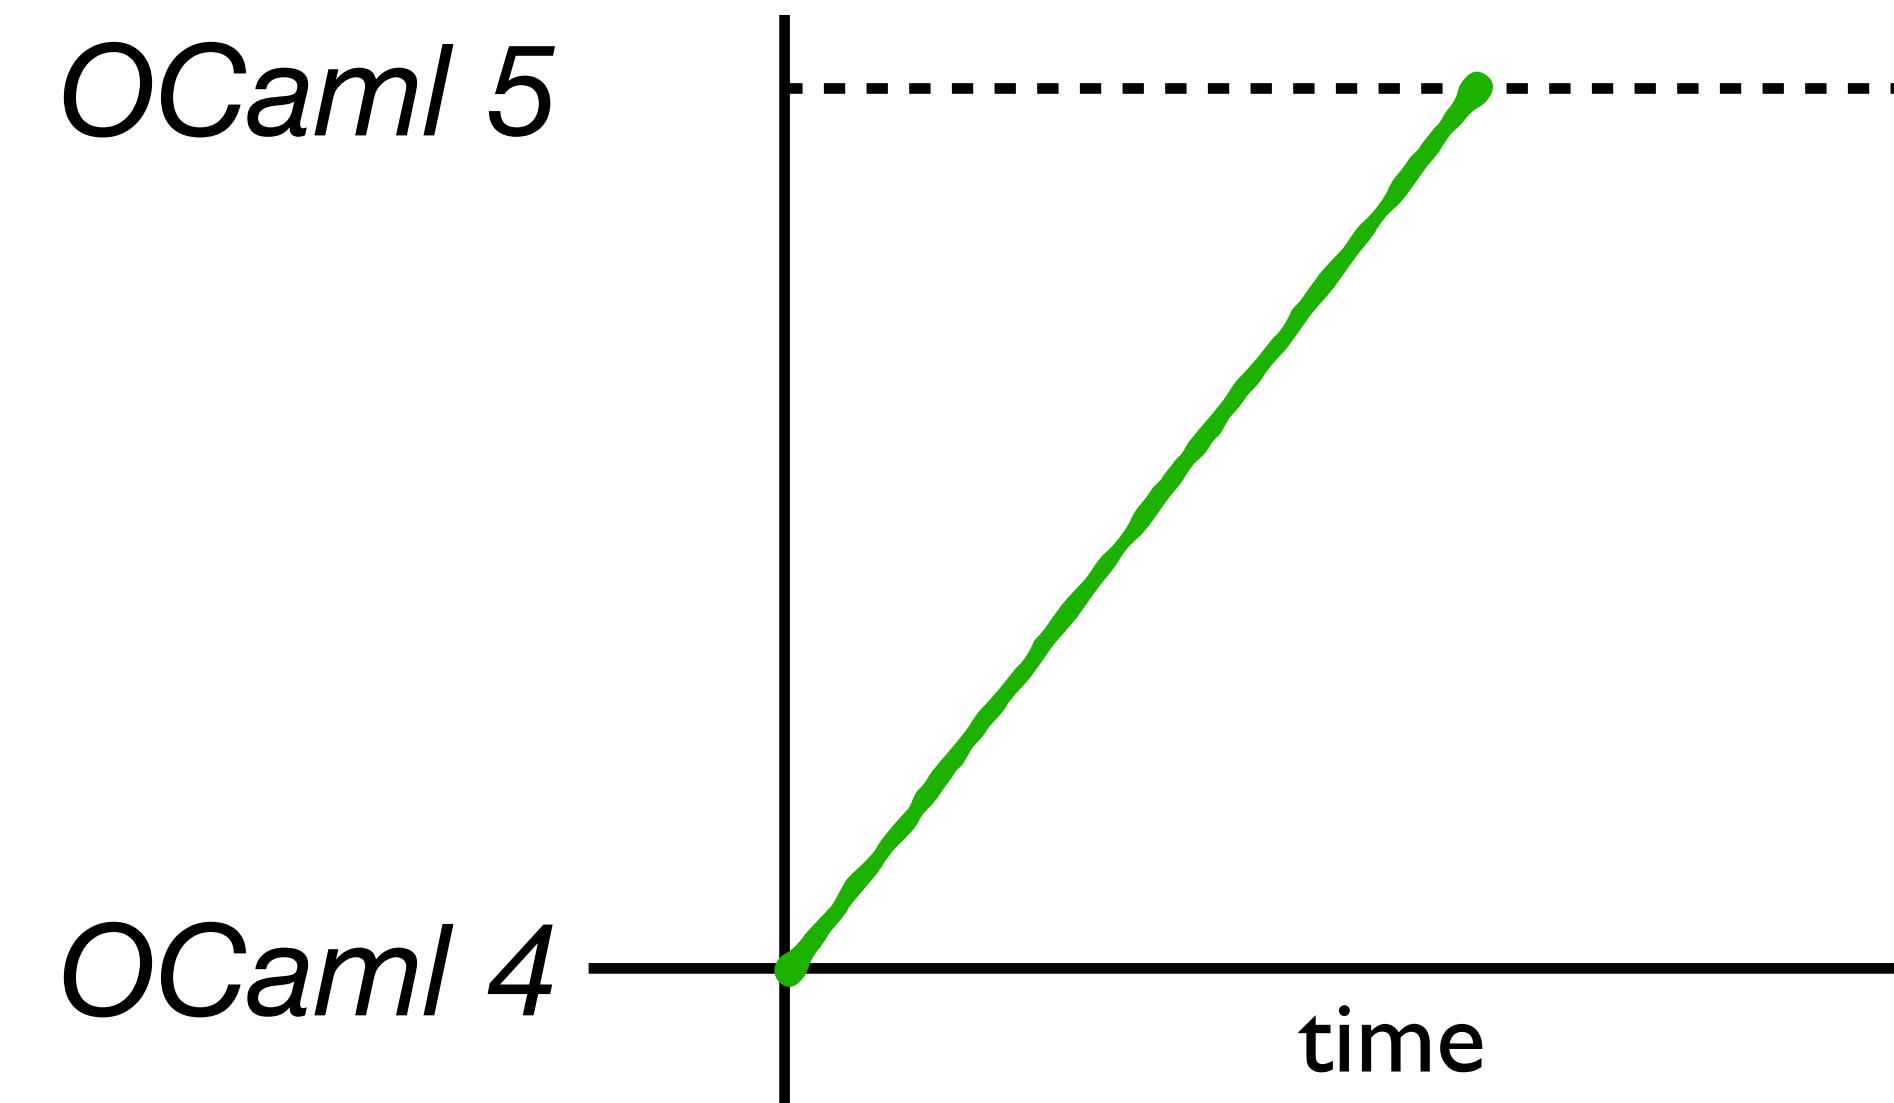
Leo White  
Jane Street  
London, UK  
leo@lpw25.net

Anil Madhavapeddy  
University of Cambridge and OCaml Labs  
Cambridge, UK  
avsm2@cl.cam.ac.uk

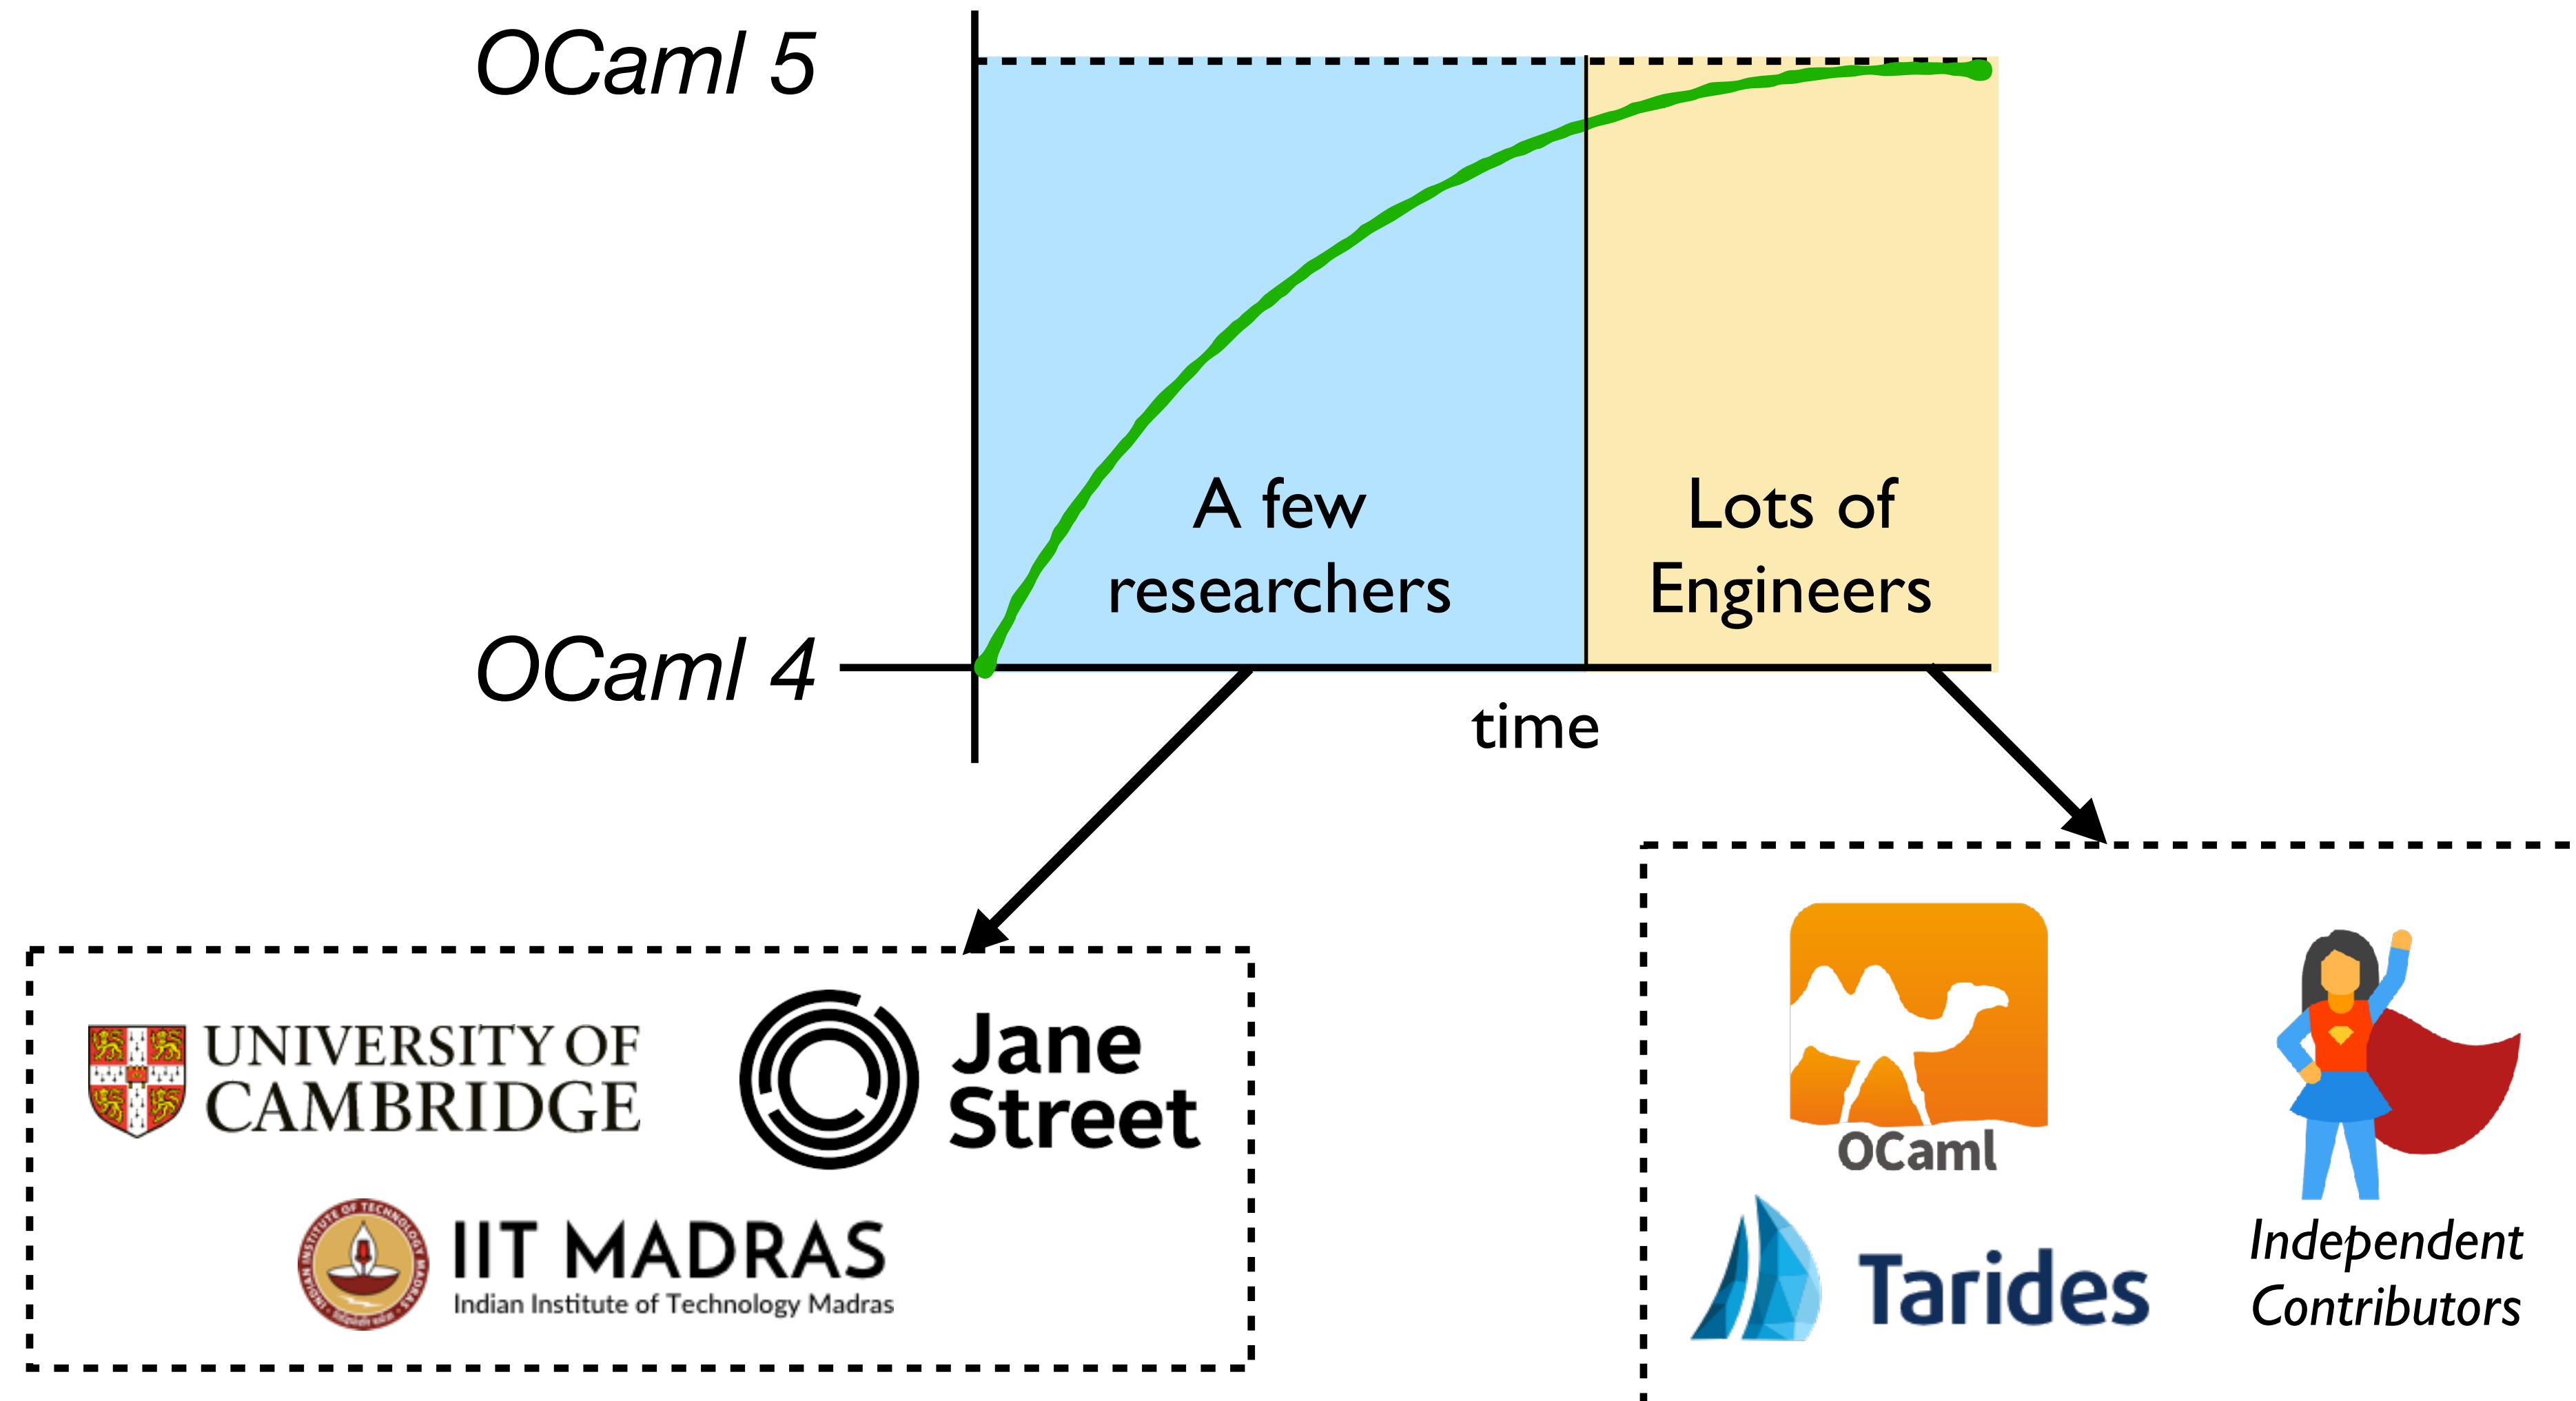
PLDI 2021

**Peer-reviewed ideas build confidence**

# Growing the language



# Growing the language



# Upstream and Release

## Multicore OCaml #10831

**Merged** xavierleroy merged 4,103 commits into `ocaml:trunk` from `ocaml-multicore:multicore-pr` on Jan 10, 2022

Conversation 393 Commits 250 Checks 0 Files changed 300+ +22,955 -14,062

**kayceesrk** commented on Dec 21, 2021 · edited Member ...

This PR adds support for shared-memory parallelism through domains and direct-style concurrency through effect handlers (without syntactic support). It intends to have backwards compatibility in terms of language features, C API, and also the performance of single-threaded code.

**For users**

If you want to learn more about Multicore OCaml, please have a look at the [multicore](#)

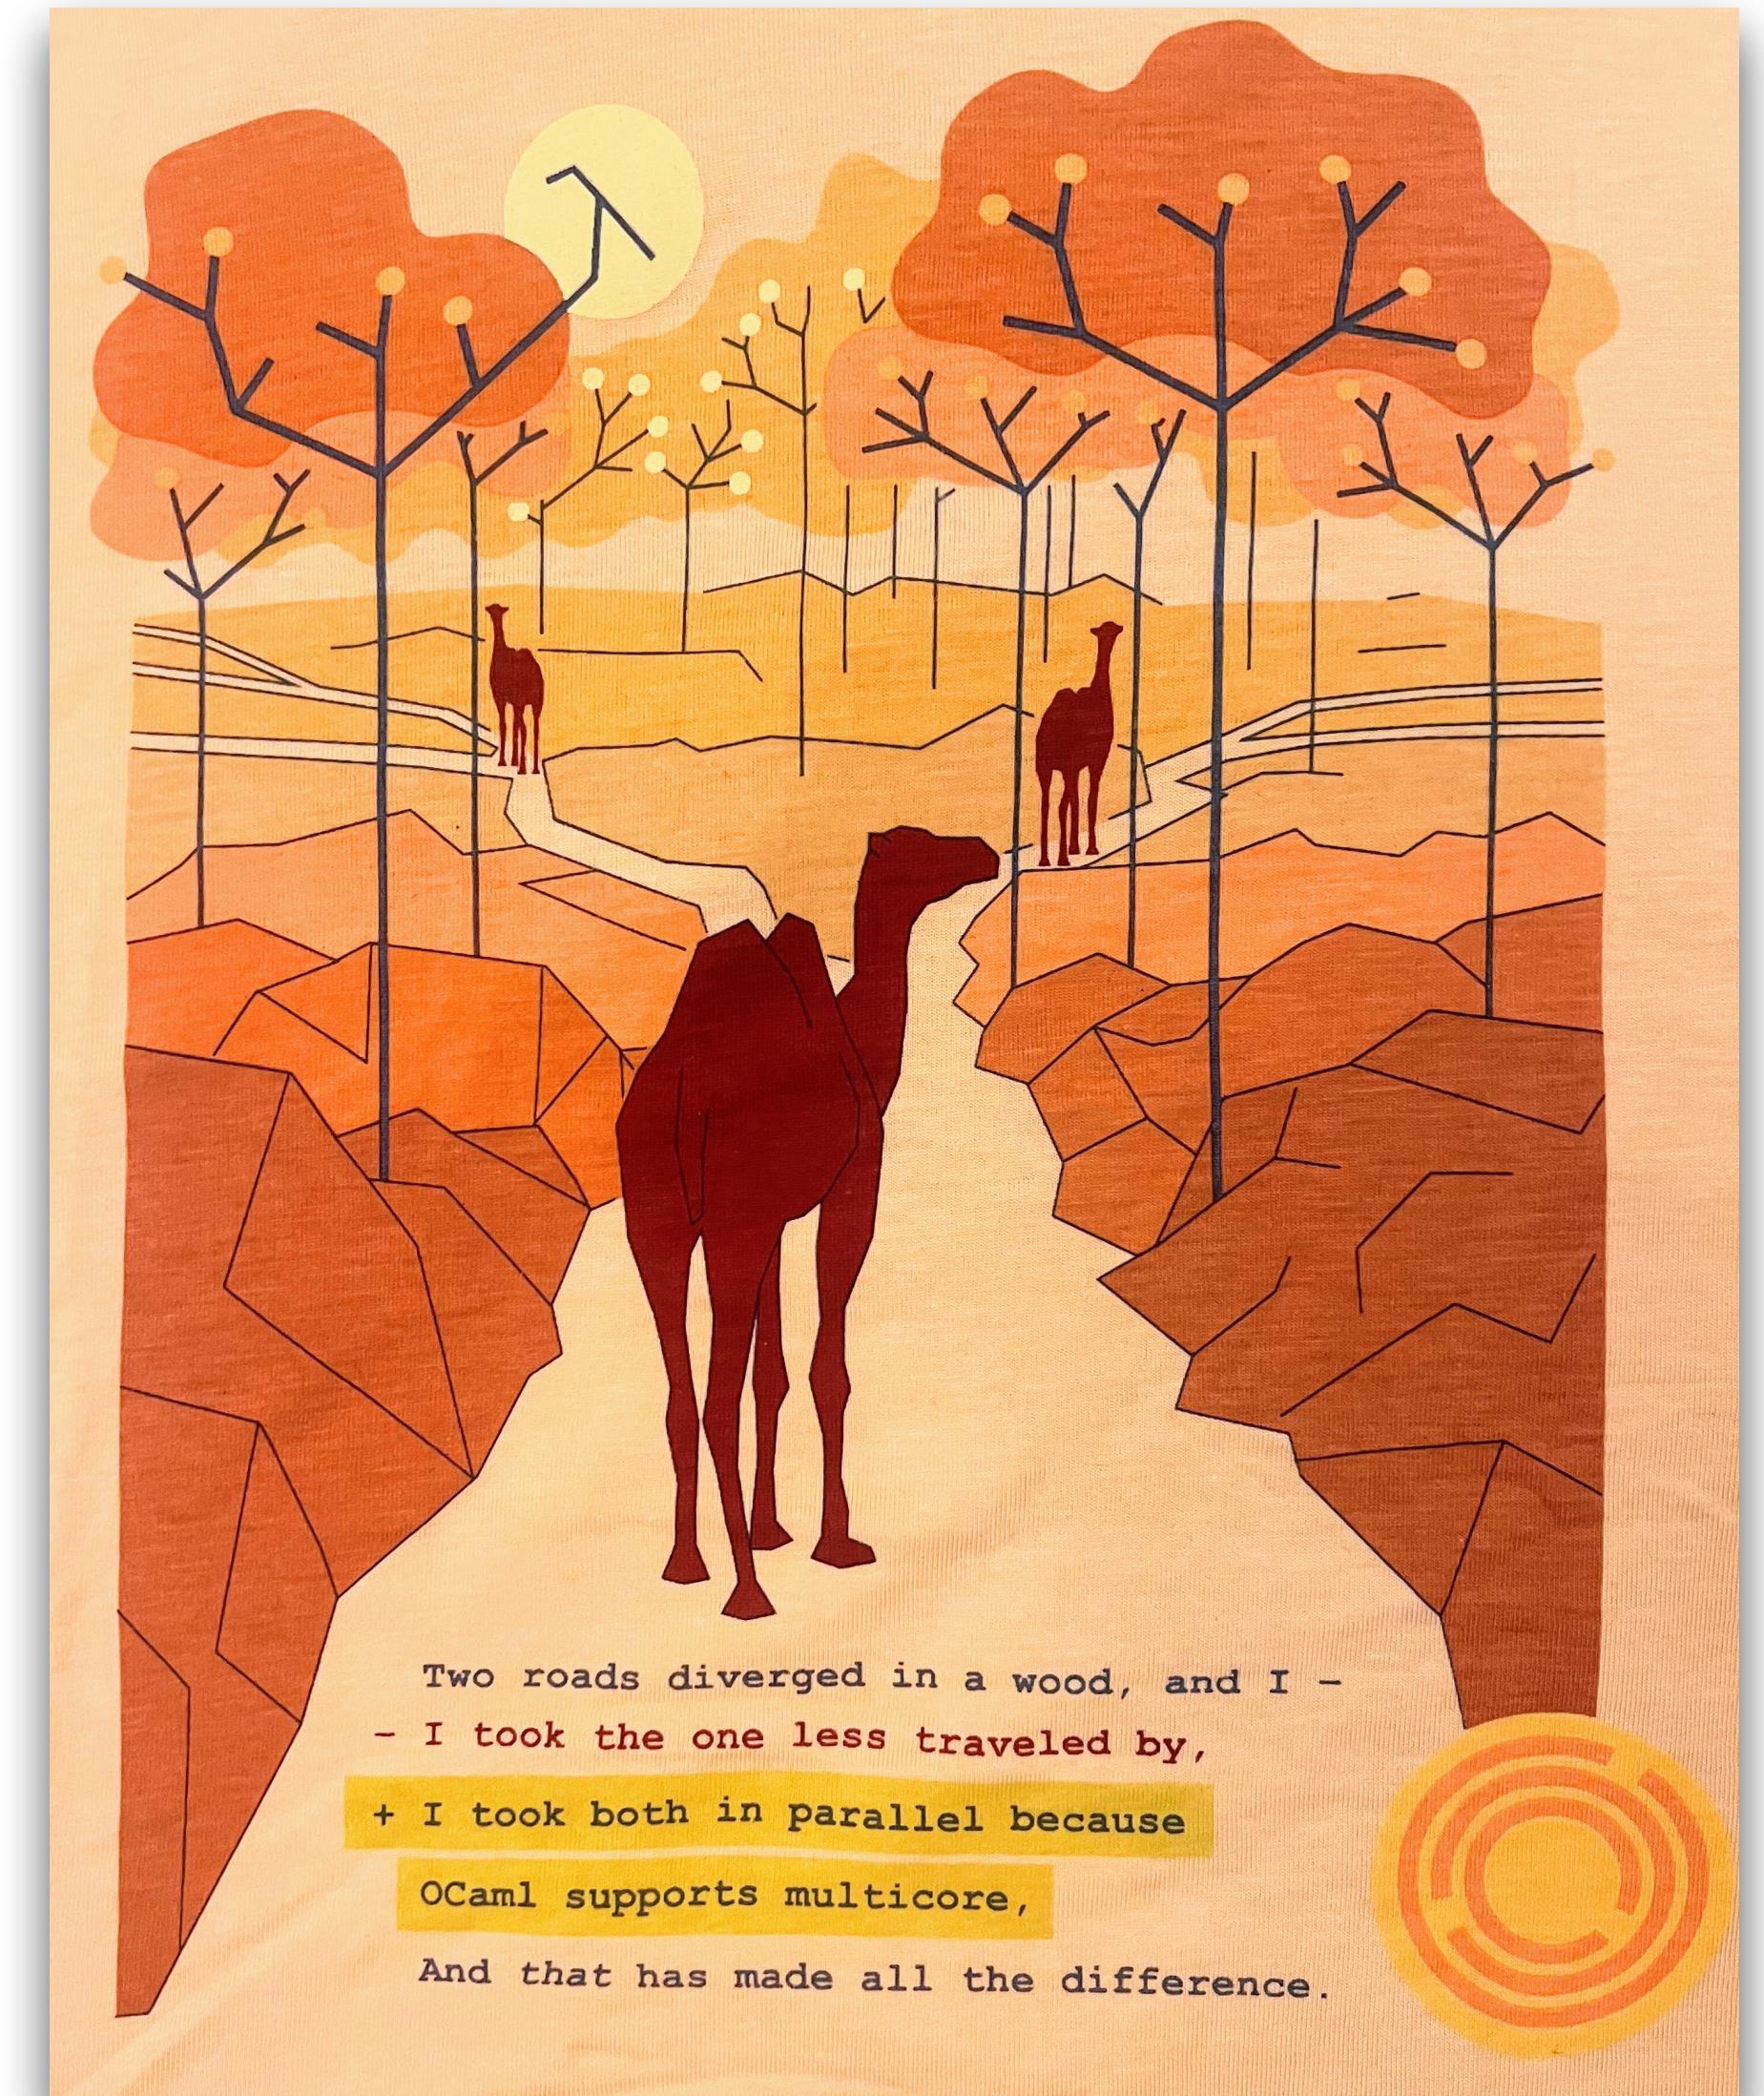
**Reviewers**

- abbysmal
- gasche
- sadiqj
- avsm
- xavierleroy
- damiendoliguez
- dra27

**Started – Mar 2014, Merged – Jan 2022**

# Upstream and Release

- **Released** – Dec 16 2022, as OCaml 5.0
- **Long tail** of adding missing features, bug fixes and performance improvements
  - 5.1 – Sep 2023
  - 5.2 – May 2024
  - 5.3 – Jan 2025
  - 5.4 – Sep 2025



# Adoption

- Several **severe performance regressions** were observed by industrial users
- Despite our efforts around
  - **Rigorous, continuous** benchmarking on **real-world programs**
  - sandmark.tarides.com – Benchmark suite, Infra and runners
- Missing gap
  - Open-source workloads do not fully characterise production workloads
- Jane Street, SemGrep are running OCaml 5 in production! 

[ICFP/SPLASH 2025 \(series\)](#) / [REBASE \(series\)](#) / [REBASE](#) /

## The Saga of Multicore OCaml

### Track

REBASE

### When

Sat 18 Oct 2025 16:00 - 17:00 at [Peony SW - REBASE](#) Chair(s): [Filip Křikava](#), [Ben L. Titzer](#)

### Abstract

In December 2022, after nearly a decade of development, OCaml 5 was released with a multi-core capable garbage collector. This was an exciting milestone, finally making it possible to write shared-memory parallel programs in OCaml. The new runtime was designed to be easy to adopt: it didn't disturb OCaml's FFI, and

## Advancing Performance via a Systematic Application of Research and Industrial Best Practice

[WENYU ZHAO](#), Australian National University, Australia

[STEPHEN M. BLACKBURN](#), Google and Australian National University, Australia

[KATHRYN S. MCKINLEY](#), Google, United States

[MAN CAO](#), Google, United States

[SARA S. HAMOUDA](#)\*, Canva, Australia

An elusive facet of high-impact research is translation to production. Production deployments are *intrinsically complex and specialized*, whereas research exploration requires stripping away incidental complexity and extraneous requirements to create *clarity and generality*. Conventional wisdom suggests that promising research rarely holds up once simplifying assumptions and missing features are addressed. This paper describes a productization methodology that led to a striking result: outperforming the mature and highly optimized state of the art by more than 10%.

Concretely, this experience paper captures lessons from translating a high-performance research garbage collector published at PLDI'22, called LXR, to a hyperscale revenue-critical application. Key to our success was

# What's next for OCaml?

- **OxCaml** – Bridging the performance and safety gap between OCaml and Rust
  - *Data-race-free parallelism* through *modes*
  - Better control over object layout, allocations and GC
- Draws lessons from Multicore OCaml execution
  - Several award-winning papers at POPL, ICFP, OOPSLA
- But different in other ways...
  - In production at Jane Street
  - Valuable user-feedback-oriented design

The screenshot shows the homepage of the OxCaml website. At the top, there is a navigation bar with the Jane Street logo and links for "About", "Documentation", and "Get OxCaml". Below the navigation bar is a large blue header featuring a white cartoon ox standing on the left and the text "OCaml, Oxidized!" in white on the right. The ox has "0xCAML" written on its side. The main content area below the header contains a section titled "OxCaml" with a brief description: "OxCaml is a fast-moving set of extensions to the OCaml programming language. It is both Jane Street's production compiler, as well as a laboratory for experiments focused towards making OCaml better for performance-oriented programming. Our hope is that these extensions can over time be contributed to upstream OCaml." At the bottom right of the page, there is a link to the website: <https://oxcaml.org>.

# FP Launchpad @ IIT Madras

- A Centre for Functional Systems Research and Education
- Mission — Robust, high-performance systems using O(x)Caml
  - **PL & compilers:** language design, semantics, optimisation
  - **Pragmatic verification:** Type systems, deductive verification, testing, model checking
  - **Runtime systems:** GC, performance tooling, parallelism
  - **Hardware-software co-design:** Correct, secure, fast systems end-to-end
- Opportunities
  - Post-bacc fellowships, MS/PhD positions, Research Staff

If interested, talk to me