

A Mechanically Verified Garbage Collector for OCaml

KC Sivaramakrishnan
kcsrk.info

IARCS Verification Seminar Series
18th November 2025

IIT
MADRAS

SADAM



Tarides

CARDID



Language



- Functional-first but multi-paradigm (imperative, OO)
- Static-type system with Hindley-Milner type inference
- Advanced features – powerful module system, GADTs, Polymorphic variants
- Multicore support and effect handlers

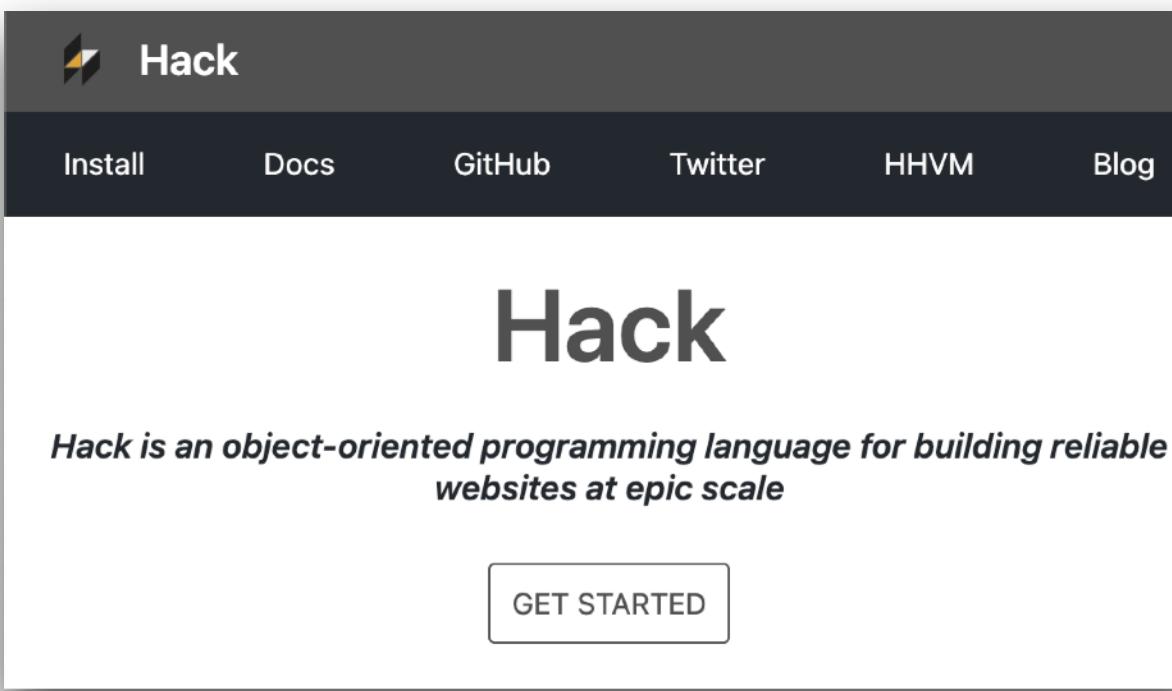


Platform

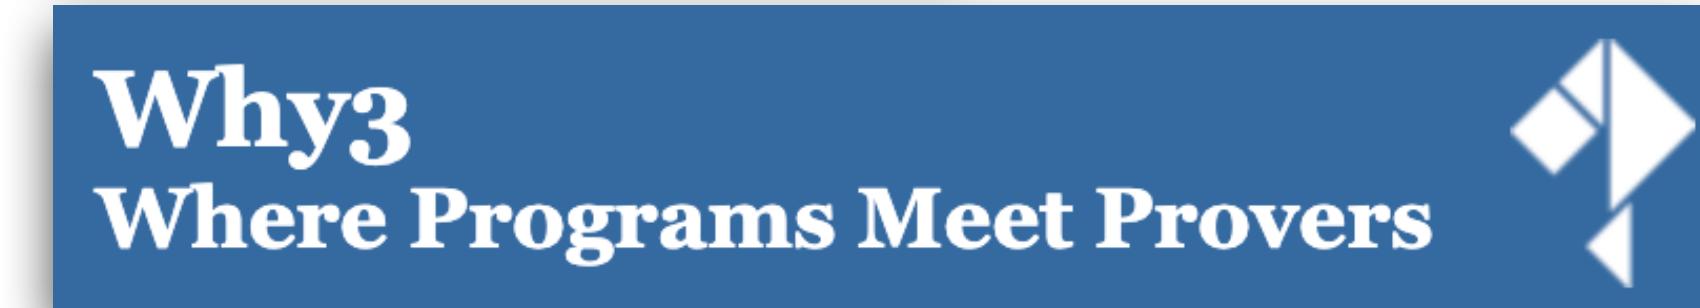
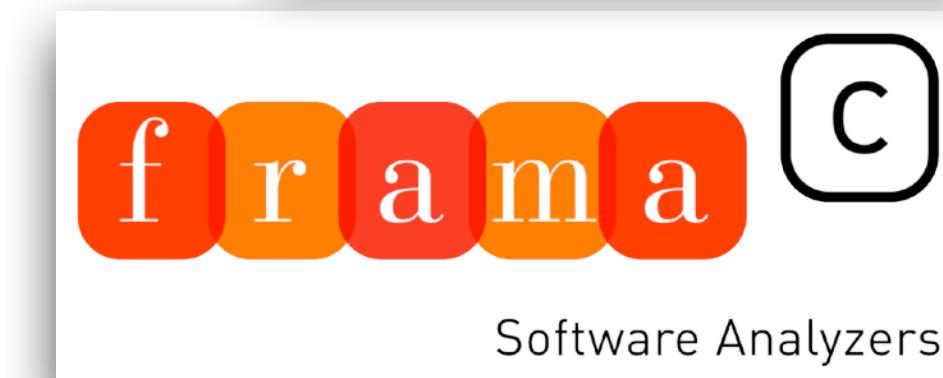
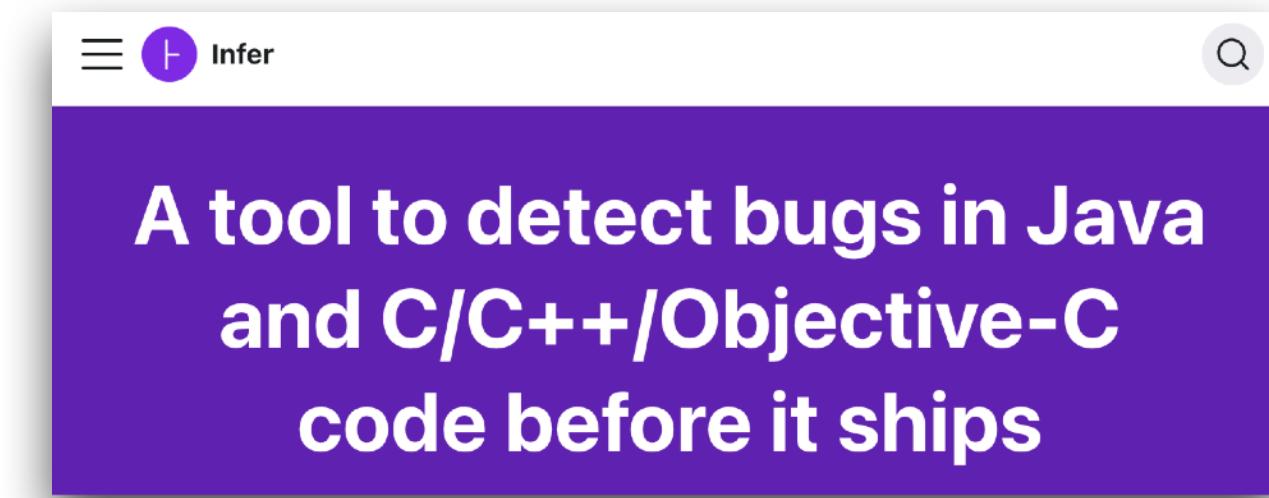
- Fast, native code— x86, ARM, RISC-V, etc.
- JavaScript and WebAssembly (using WasmGC) compilation

High dynamic range

From scripts to scalable systems, research prototypes to production infrastructure



Compilers



Verification & Static Analysis

High dynamic range

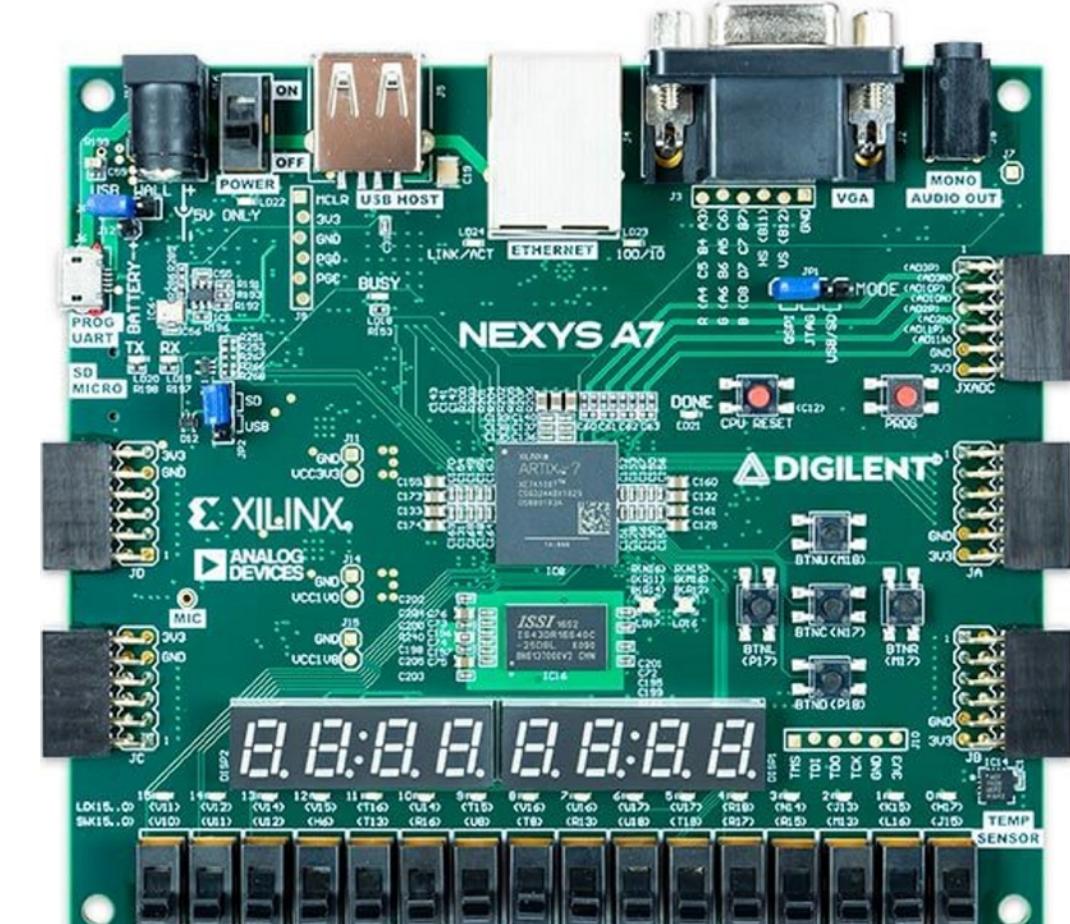
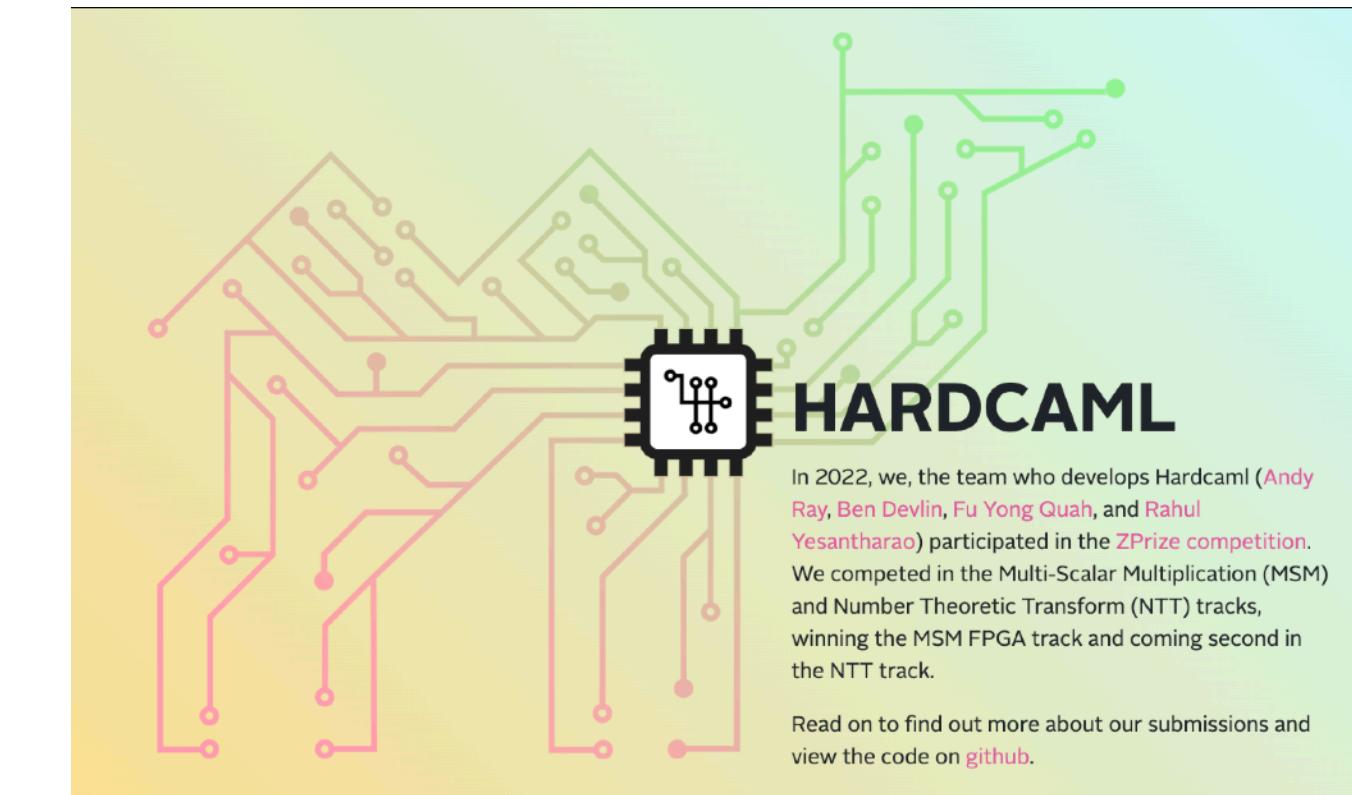
From scripts to scalable systems, research prototypes to production infrastructure



Jane
Street

Bloomberg

60+M lines of
OCaml code!

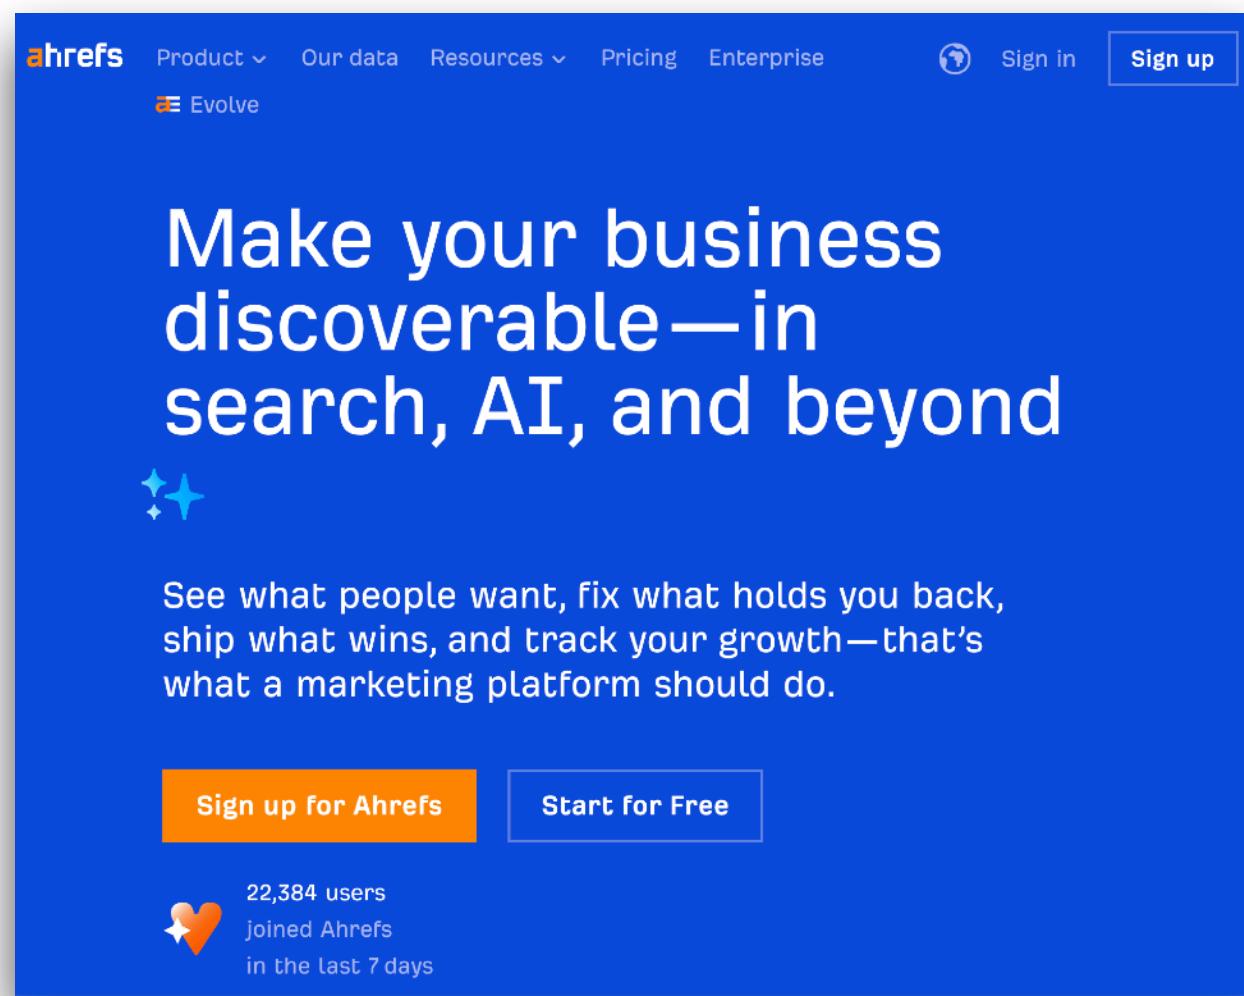


Finance

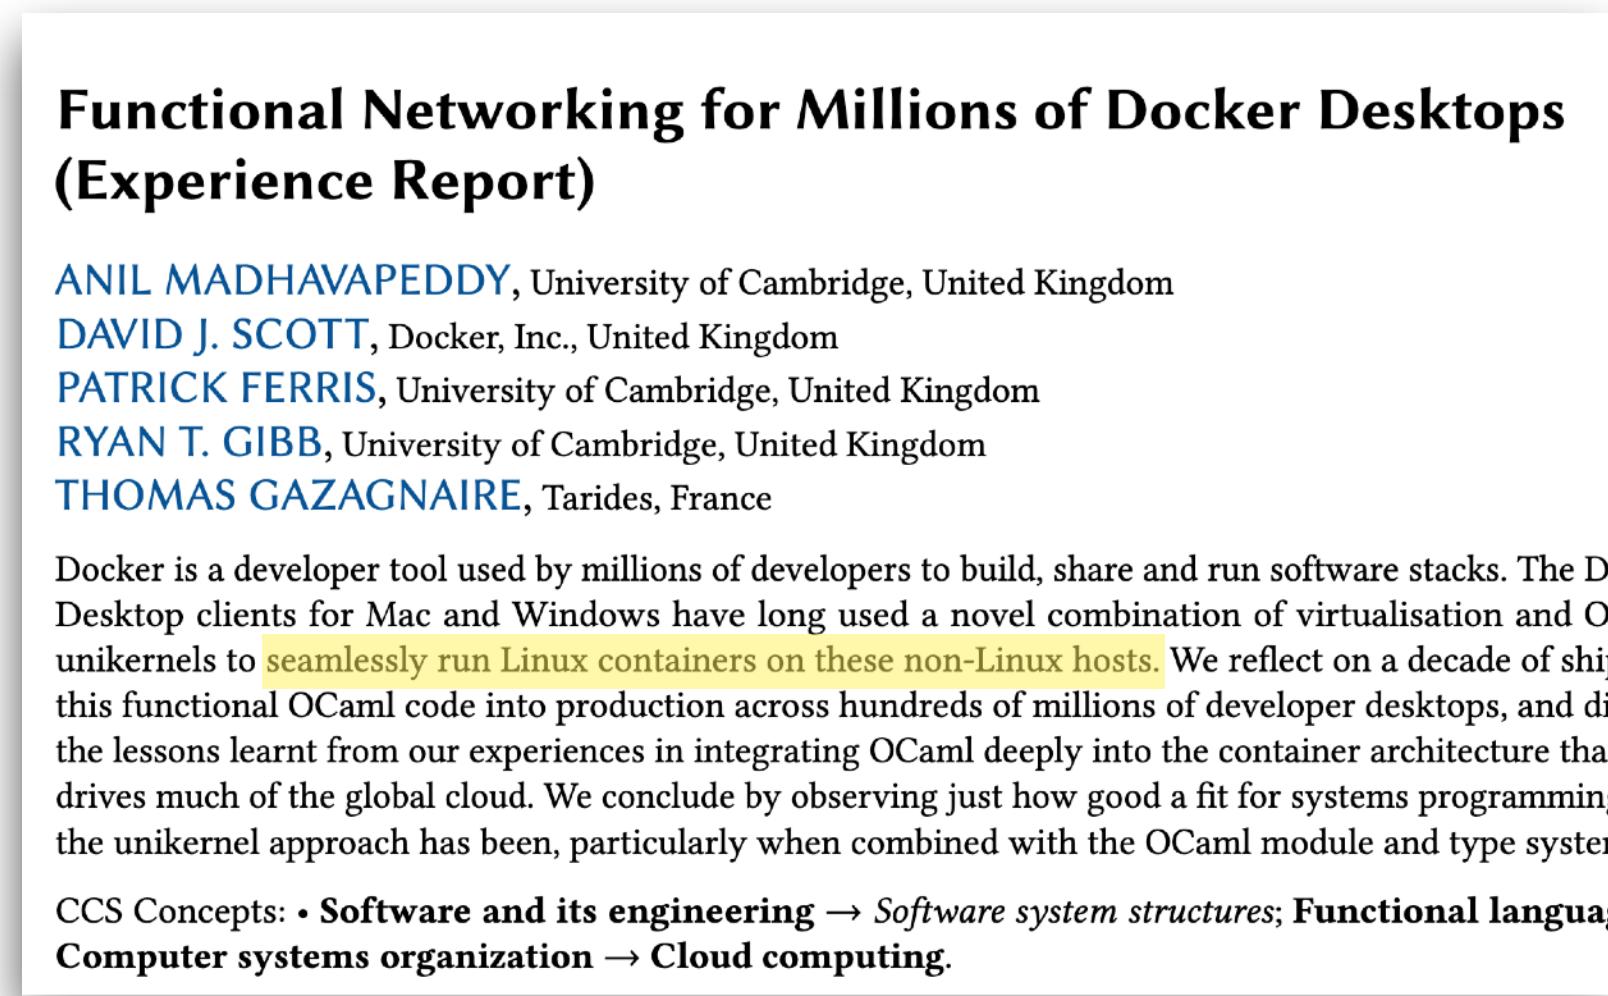
Hardware design

High dynamic range

From scripts to scalable systems, research prototypes to production infrastructure



The Ahrefs homepage features a blue header with navigation links for Product, Our data, Resources, Pricing, Enterprise, Sign in, and Sign up. Below the header, a large white text area says "Make your business discoverable—in search, AI, and beyond" with a small star icon. A quote follows: "See what people want, fix what holds you back, ship what wins, and track your growth—that's what a marketing platform should do." At the bottom, there are two buttons: "Sign up for Ahrefs" and "Start for Free". A small stats box shows "22,384 users joined Ahrefs in the last 7 days".

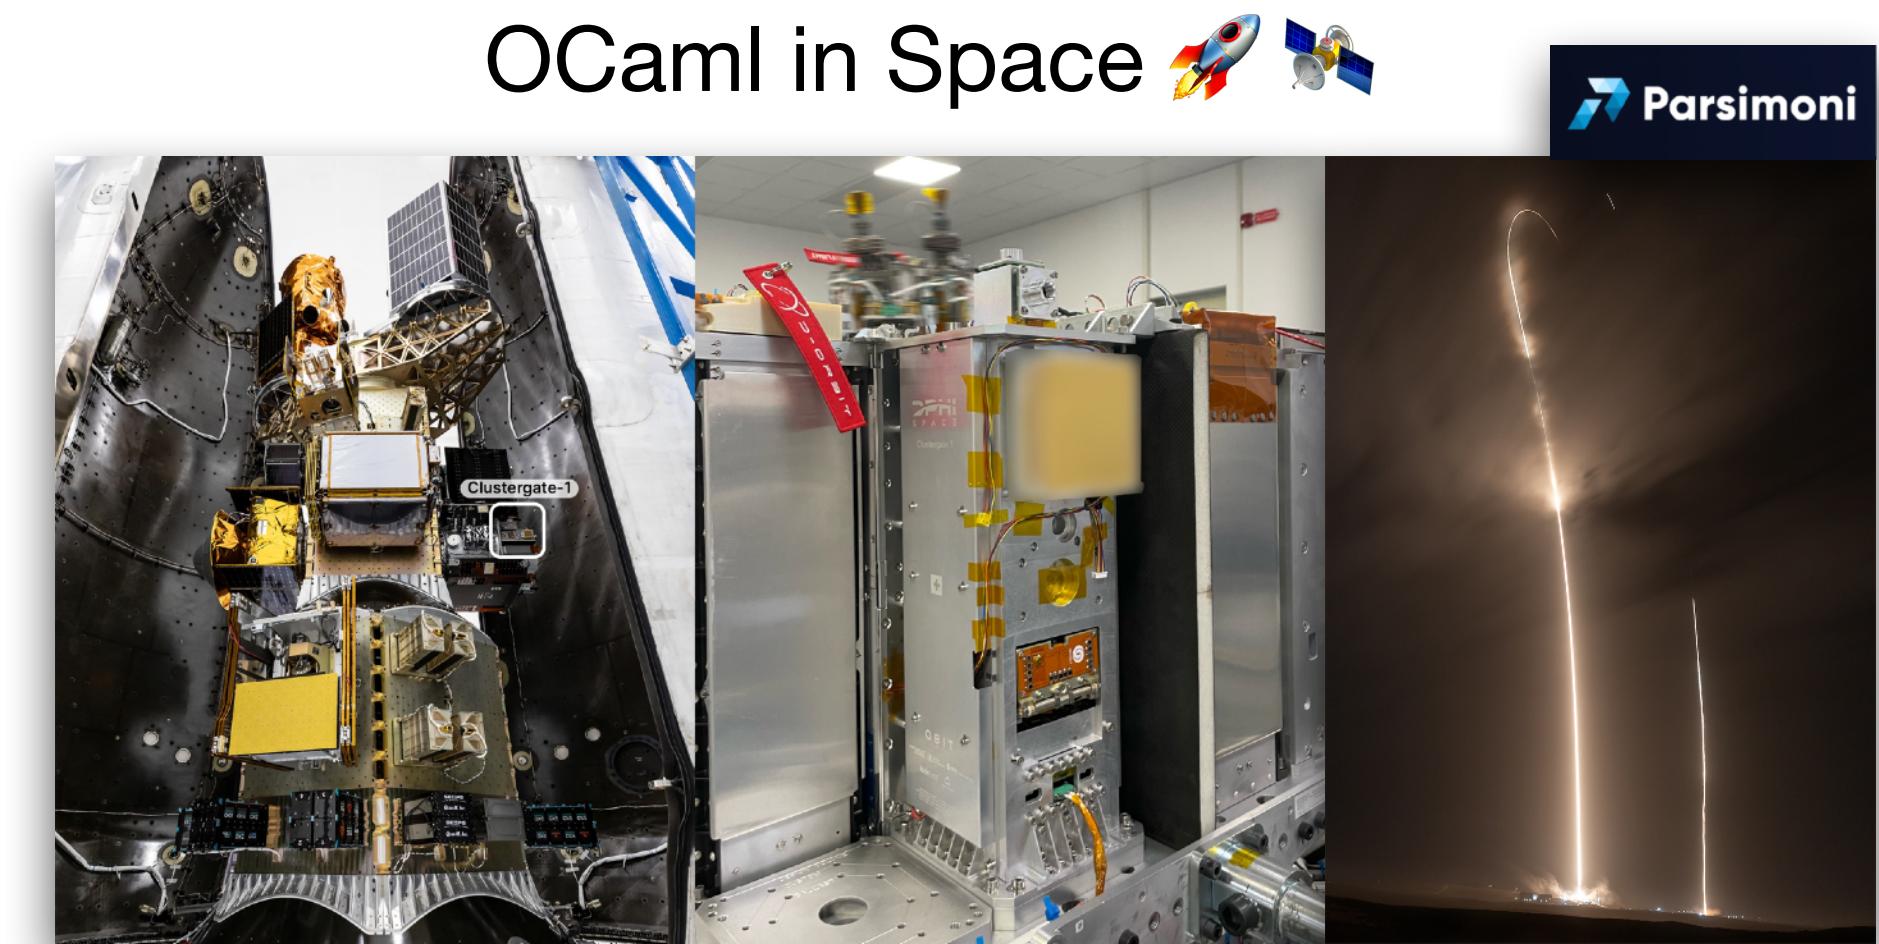


Functional Networking for Millions of Docker Desktops (Experience Report)

ANIL MADHAVAPEDDY, University of Cambridge, United Kingdom
DAVID J. SCOTT, Docker, Inc., United Kingdom
PATRICK FERRIS, University of Cambridge, United Kingdom
RYAN T. GIBB, University of Cambridge, United Kingdom
THOMAS GAZAGNAIRE, Tarides, France

Docker is a developer tool used by millions of developers to build, share and run software stacks. The Docker desktop clients for Mac and Windows have long used a novel combination of virtualisation and OCaml unikernels to seamlessly run Linux containers on these non-Linux hosts. We reflect on a decade of shipping this functional OCaml code into production across hundreds of millions of developer desktops, and discuss the lessons learnt from our experiences in integrating OCaml deeply into the container architecture that drives much of the global cloud. We conclude by observing just how good a fit for systems programming the unikernel approach has been, particularly when combined with the OCaml module and type system.

CCS Concepts: • Software and its engineering → Software system structures; Functional languages; Computer systems organization → Cloud computing.



Web Frontend

Networking

Virtualisation

Trusted Computing Base

- Unsurprisingly, the OCaml compiler is written in OCaml
 - But includes *a runtime system written in C*

```
[kc@KCx-MacBook-Pro-2 ocaml % cloc .  
 6167 text files.  
 4025 unique files.  
 4354 files ignored.  
  
github.com/AlDanial/cloc v 2.02 T=2.41 s (1669.2 files/s, 314367.8 lines/s)  
-----  
Language      files    blank   comment     code  
-----  
OCaml        2974    62693   115033    393861  
C            308     8629    10315     48137  
Bourne Shell 72      7502    8900      45563  
m4           12      1489    111       12448  
Assembly     18      561     2194      5869  
C/C++ Header 91      2034    3580      5600  
make          23      922     610       3633  
Markdown     29      940     34        2828  
D             410     0        0        2168  
AsciiDoc     12      518     0        1682  
CSV           1       0        0        1512
```

~400k lines of OCaml

~60k lines of **C & Assembly**

All of it in the runtime system,
much of it in the GC

OCaml GC

- OCaml is a garbage-collected (GCed) language.
 - Low-latency, high-throughput with good space-time tradeoff
- Memory management subsystem = Allocator + GC
- GC in OCaml 5
 - is a **complex piece of software**
 - A generational, concurrent and parallel!
 - Support for weak references, finalisers (2 kinds), ephemeros, etc.
- **Tons of bugs** during development (see Multicore OCaml project)
 - C is excellent for writing unsafe, hard-to-reason-about code! :-(

Can we build a correct-by-construction GC for OCaml?

Verified GC desiderata

- **Correct-by-construction**
 - End-to-end proof that the GC preserves safety and liveness
- **Pluggable**
 - Should be able to *slot in* for the existing GC
- **Extensible**
 - Accommodate improvements without rearchitecting from scratch
- **Efficient**
 - Be competitive with existing GC



Everything?

Our work

- A verified stop-the-world ***mark-and-sweep*** GC for OCaml
- Written in F* and its subset Low*, proof-oriented programming languages
- Extracted to memory-safe C and integrated with OCaml 4.14.1 (non-multicore)
- Competitive performance with vanilla OCaml

[Home](#) > [Journal of Automated Reasoning](#) > Article

A Mechanically Verified Garbage Collector for OCaml

[Open access](#) | Published: 14 May 2025
Volume 69, article number 11, (2025) [Cite this article](#)

[Download PDF](#)  You have full access to this [open access](#) article

[Sheera Shamsu](#) , [Dipesh Kafle](#), [Dhruv Maroo](#), [Kartik Nagar](#), [Karthikeyan Bhargavan](#) & [KC Sivaramakrishnan](#)

 997 Accesses [Explore all metrics](#) →

Abstract

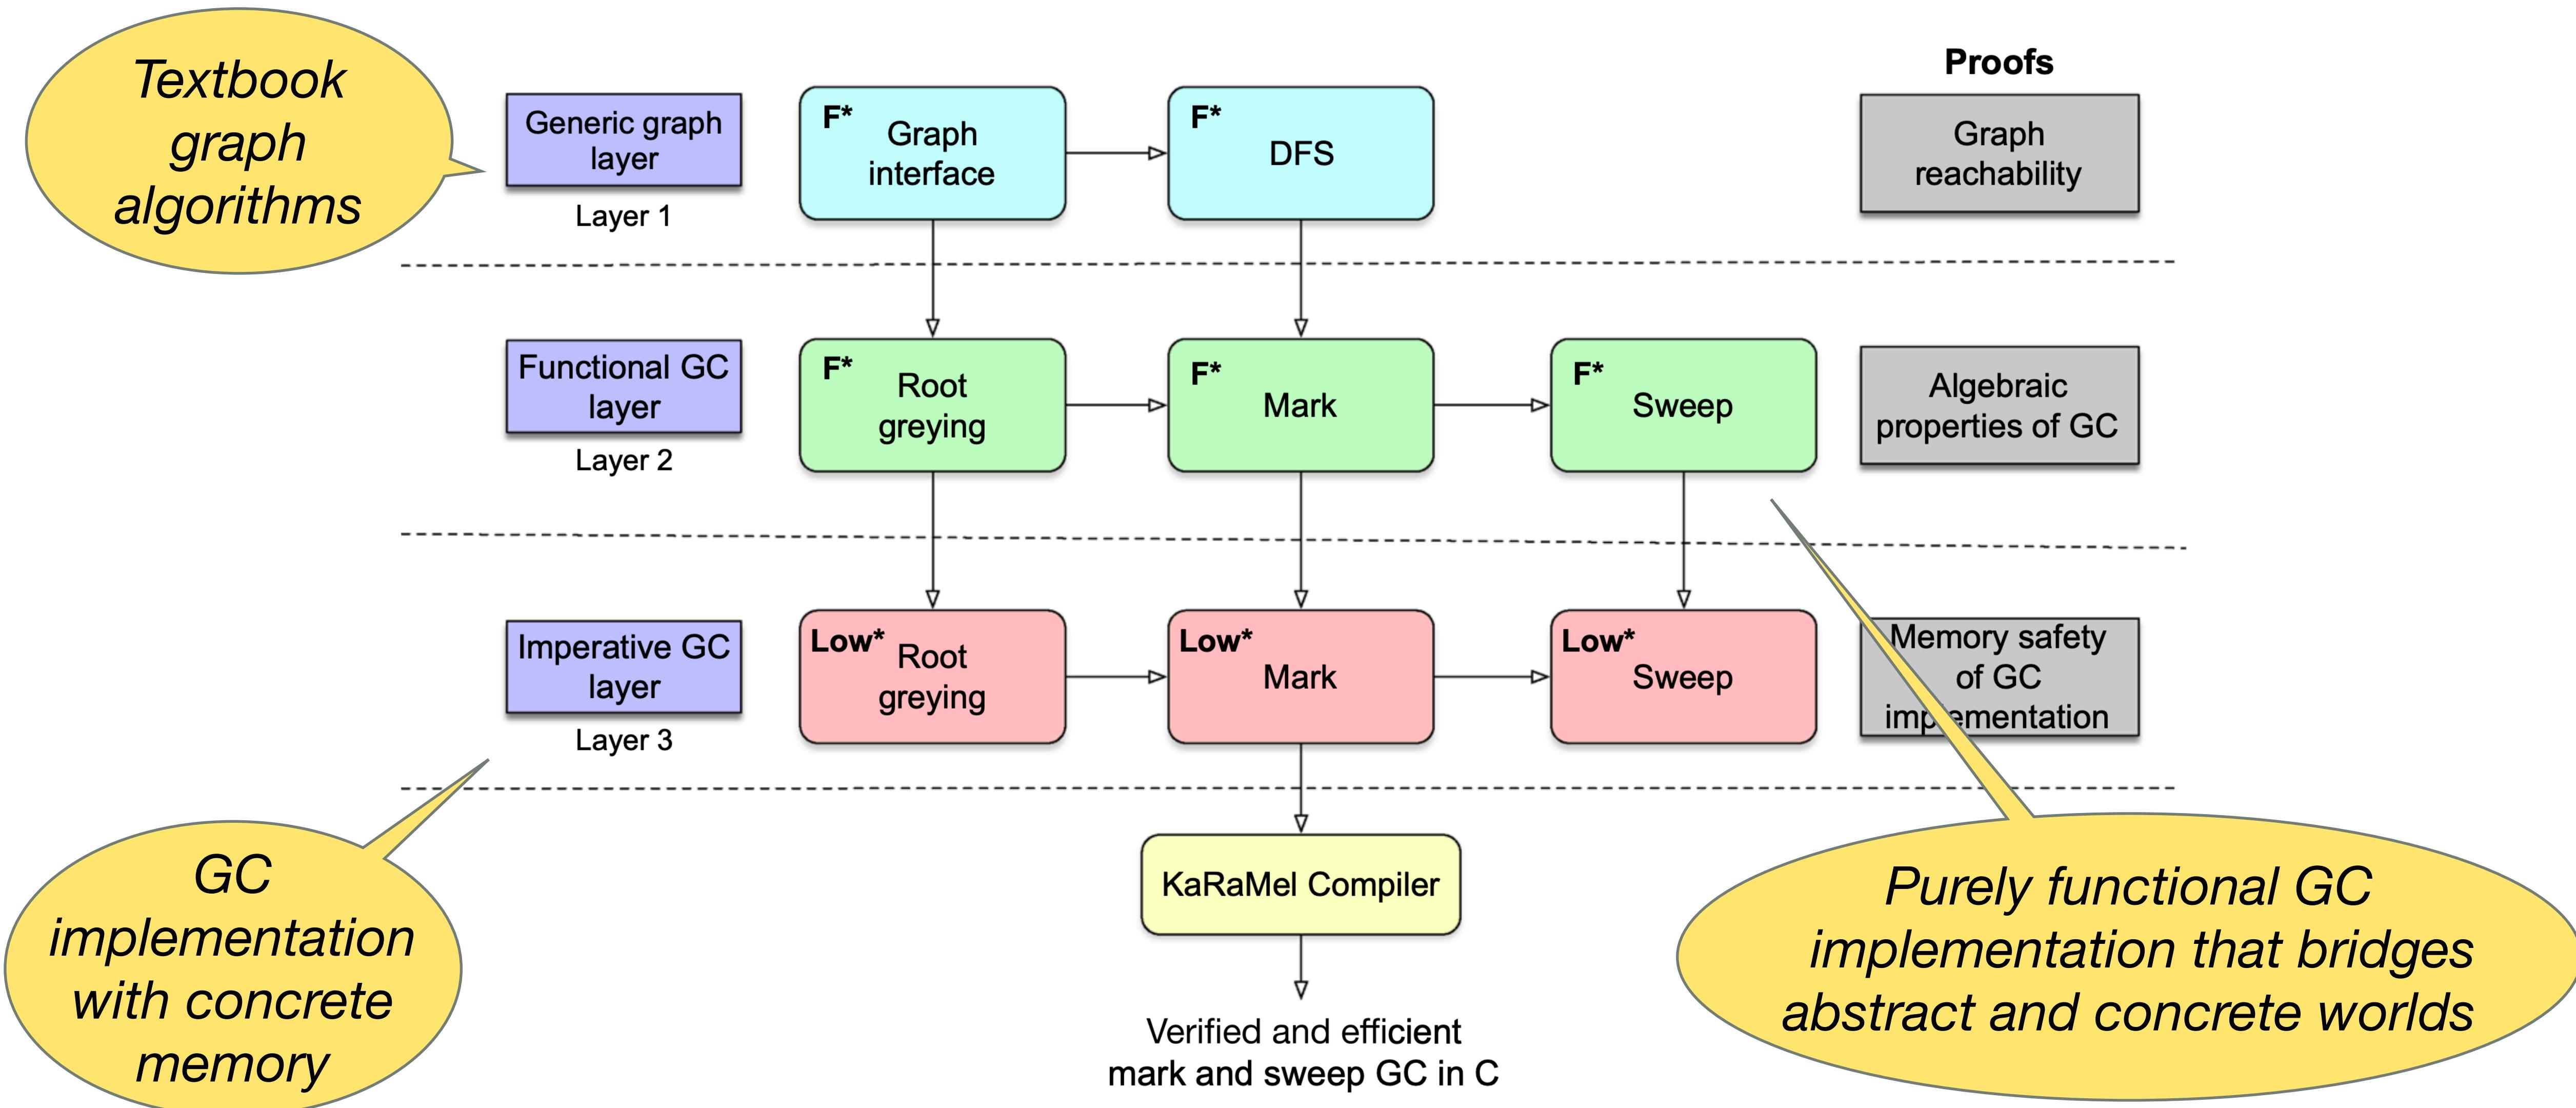
The OCaml programming language finds application across diverse domains, including systems programming, web development, scientific computing, formal verification, and symbolic mathematics. OCaml is a memory-safe programming language that uses a

Journal of Automated Reasoning, 2025

GC Correctness

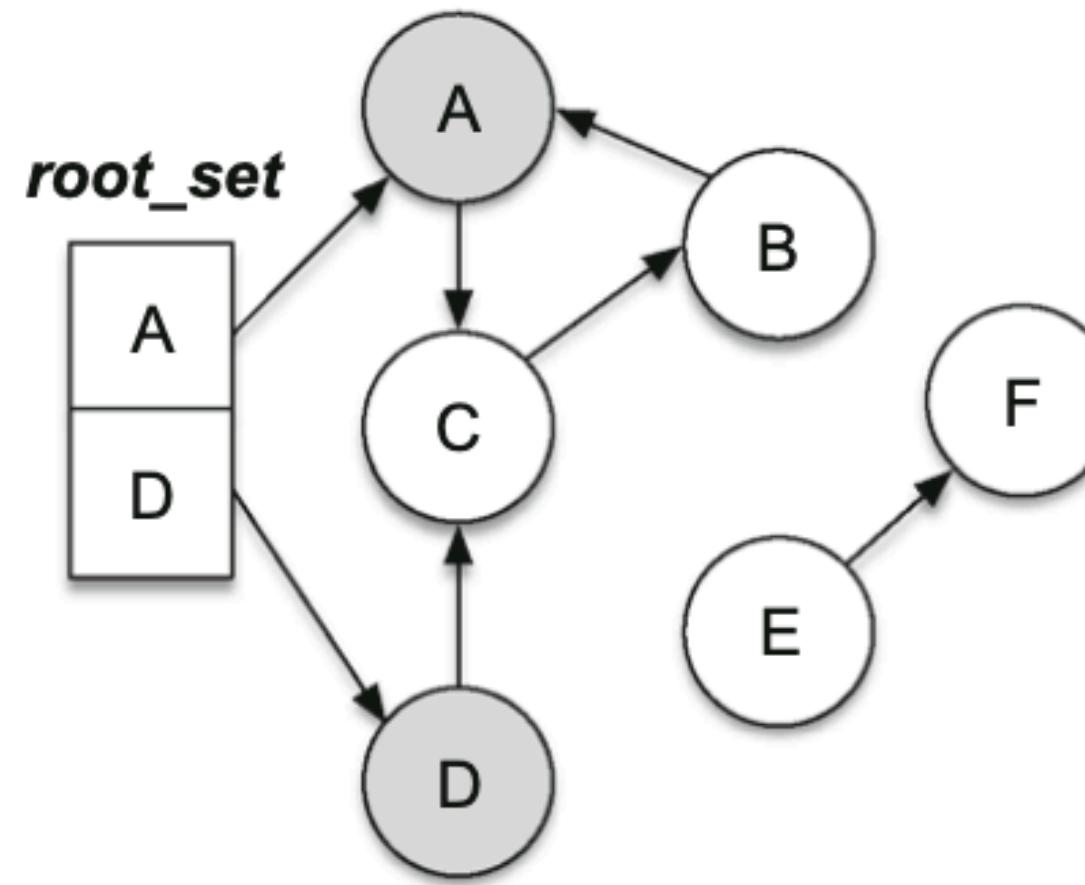
- Correctness properties
 - Safety – GC preserves all ***unreachable*** objects
 - Liveness – GC frees all ***unreachable*** objects
- How do we approach this?
 - Reachability is a property of the ***graph*** formed by objects in the heap.
 - Reason about correctness by connecting ***reachability*** to the ***GC operations***

GC verification – Layered approach



Tricolour Mark-and-sweep GC

White = Unmarked, Grey = Marking, Black = Marked

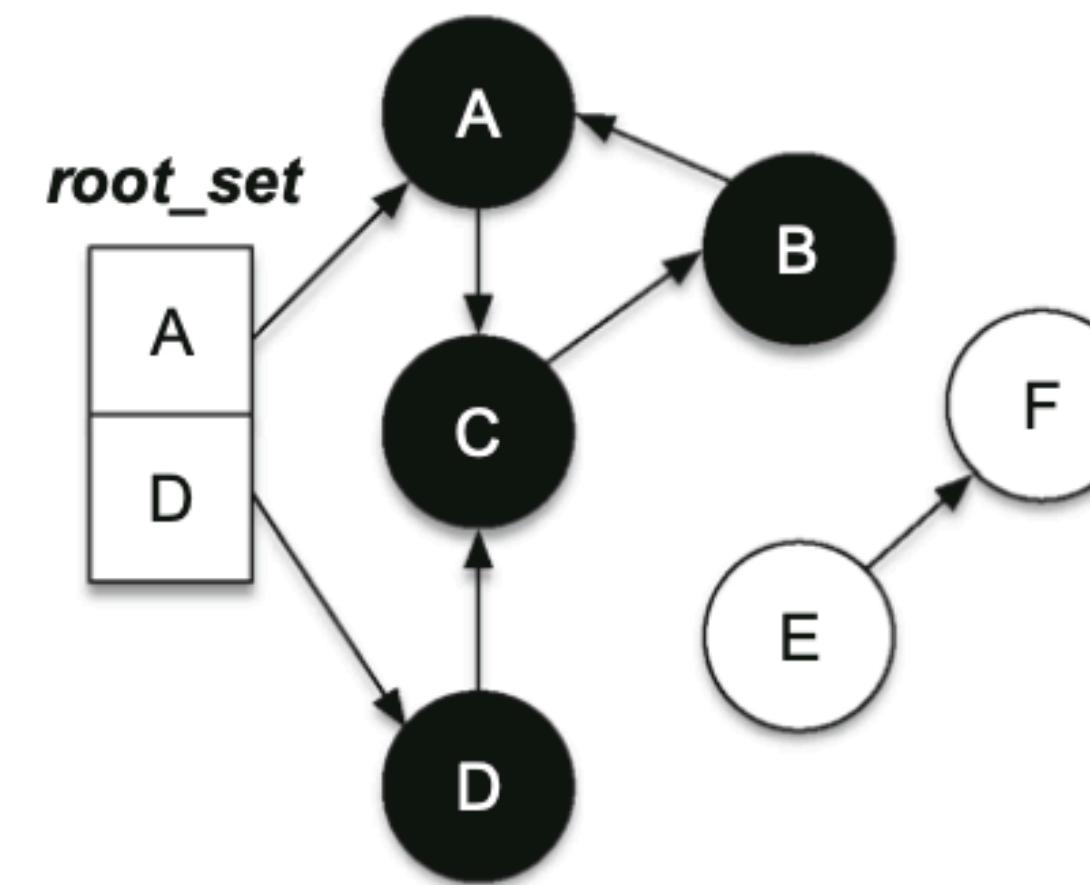


Whites = {B, C, E, F}

Blacks = {}

Greys = {A, D}

(a) Start of mark

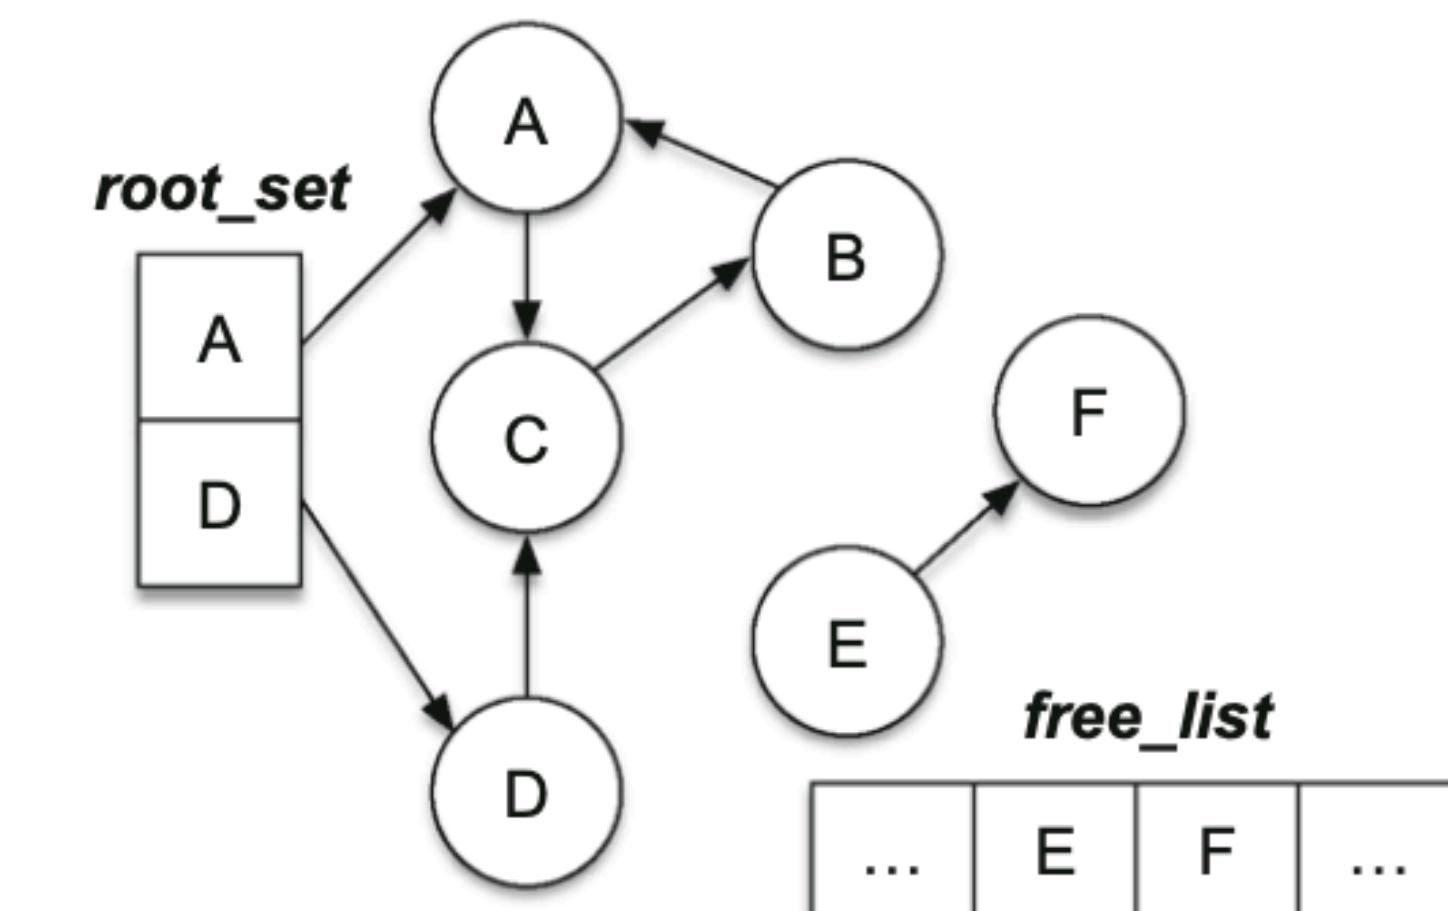


Whites = {E, F}

Blacks = {A, B, C, D}

Greys = {}

(b) After mark



Whites = {A, B, C, D}

Blacks = {}

Greys = {}

(c) After sweep

Layer 1 – Graph and reachability

```
noeq type graph (#a:eqtype) = {  
    vertices : v: vertex_set #a;  
    (* [vertices] are a sequence of type a with no duplicates *)  
    edges : e: edge_set #a {edge_ends_are_vertices vertices e}  
    (* [edges] are a sequence of type (a,a) with no duplicates *)  
}
```

```
type reach: (g:graph) → (x:vertex) → (y:vertex) → Type =  
(* reachability is reflexive *)  
| ReachRefl : (g:graph) → (x:vertex) → (reach g x x)  
(* reachability is transitive *)  
| ReachTrans : (g:graph) → (x:vertex) → (z:vertex) →  
    (reach g x z) →  
    (* [edge g z y] is a type refinement which mandates  
     that [(z,y)] is an edge in [g] *)  
    (y:vertex {edge g z y}) → (reach g x y)
```

Layer 1 – Functional Depth-first search

```
(* dfs calls dfs_body until stack empty.  
   Inputs are graph, stack and visited set. *)  
let rec dfs (g:graph) (st:seq U64.t) (vl:seq U64.t)  
  : Pure (seq U64.t)  
  (requires ...)  
  (ensures (λ res → ...))  
  (decreases (length g.vertices - length vl; length st)) =  
  if length st = 0 then vl  
  else  
    let st1 ,vl1 = dfs_body g st vl in  
    dfs g st1 vl1  
  
let dfs_body g st vl  
  : Pure ...  
  (requires ...)  
  (ensures (λ res → ...)) =  
    let x   = stack_top st in  
    let xs  = stack_rest st in  
    let s   = successors g x in  
    let vl1 = set_insert x vl in  
    let st1 = push_unvisited s xs vl1 in  
    (st1, vl1)
```

Layer 1 – Reachability \equiv DFS

```
(* r_list is the root set, stack is filled with r_list initially *)
val dfs_reachability_lemma (g:graph) (st:seq obj_addr)
                            (vl:seq obj_addr) (r_list:seq obj_addr)

: Lemma

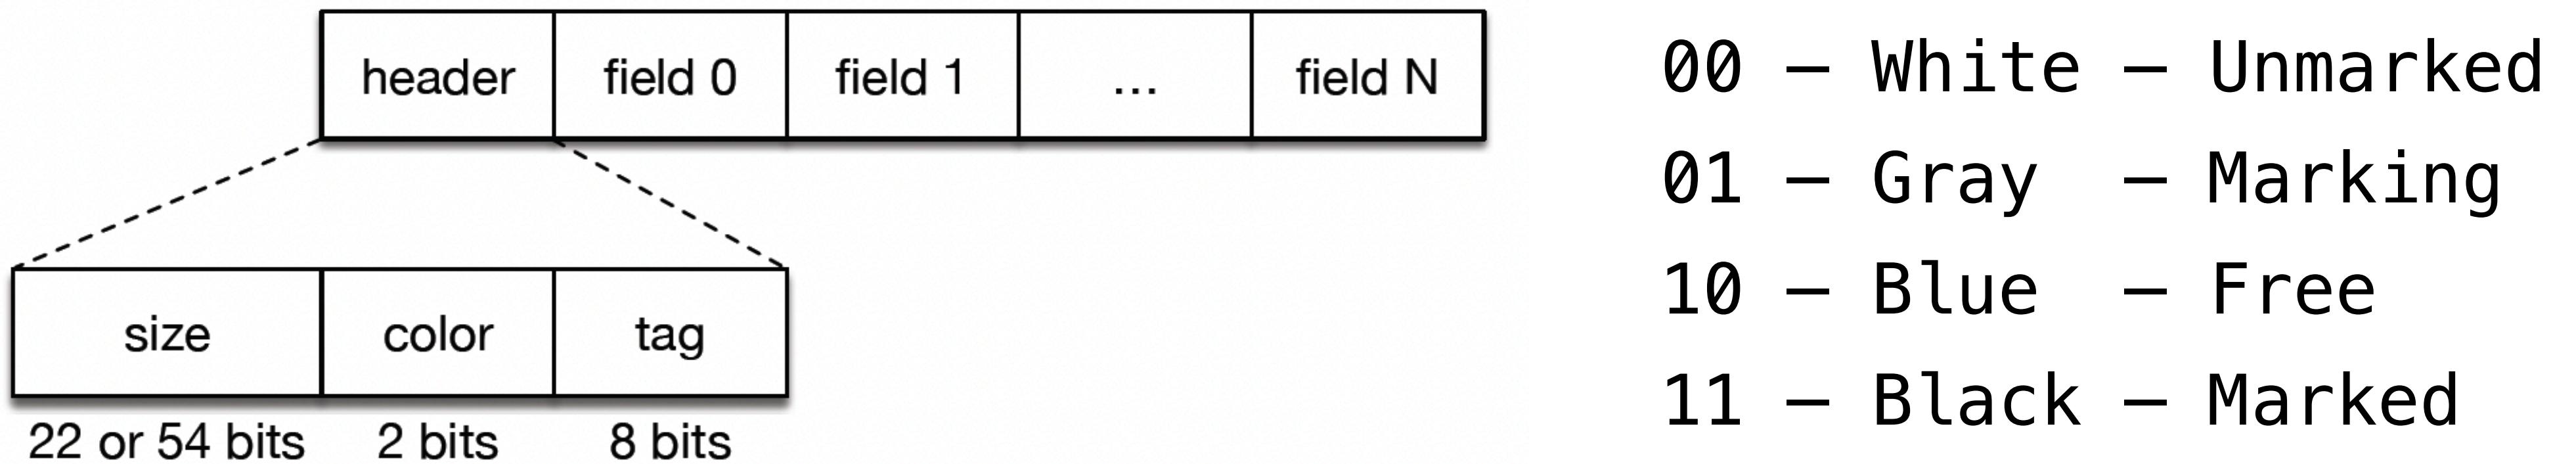
(requires
  (* Pre-conditions required to prove forward direction *)
  (*F1*) mutually_exclusive_sets st vl ∧
  (*F2*) (forall y.y ∈ st ==> (exists x.x ∈ r_list ∧ reach g x y)) ∧
  (*F3*) (forall y.y ∈ vl ==> (exists x.x ∈ r_list ∧ reach g x y)) ∧

  (* Pre-conditions required to show the backward direction *)
  (*B1*) (forall x y.x ∈ r_list ∧ reach g x y ==>
           (exists z.z ∈ st ∧ reach g z y) ∨ y ∈ vl)
  (*B2*) (forall x y.x ∈ vl ∧ reach g x y ==>
           (exists z.z ∈ st ∧ reach g z y) ∨ y ∈ vl)

(ensures (forall y.y ∈ (dfs g st vl) <=> (exists x.x ∈ r_list ∧ reach g x y)))
```

Layer 2 – Heap

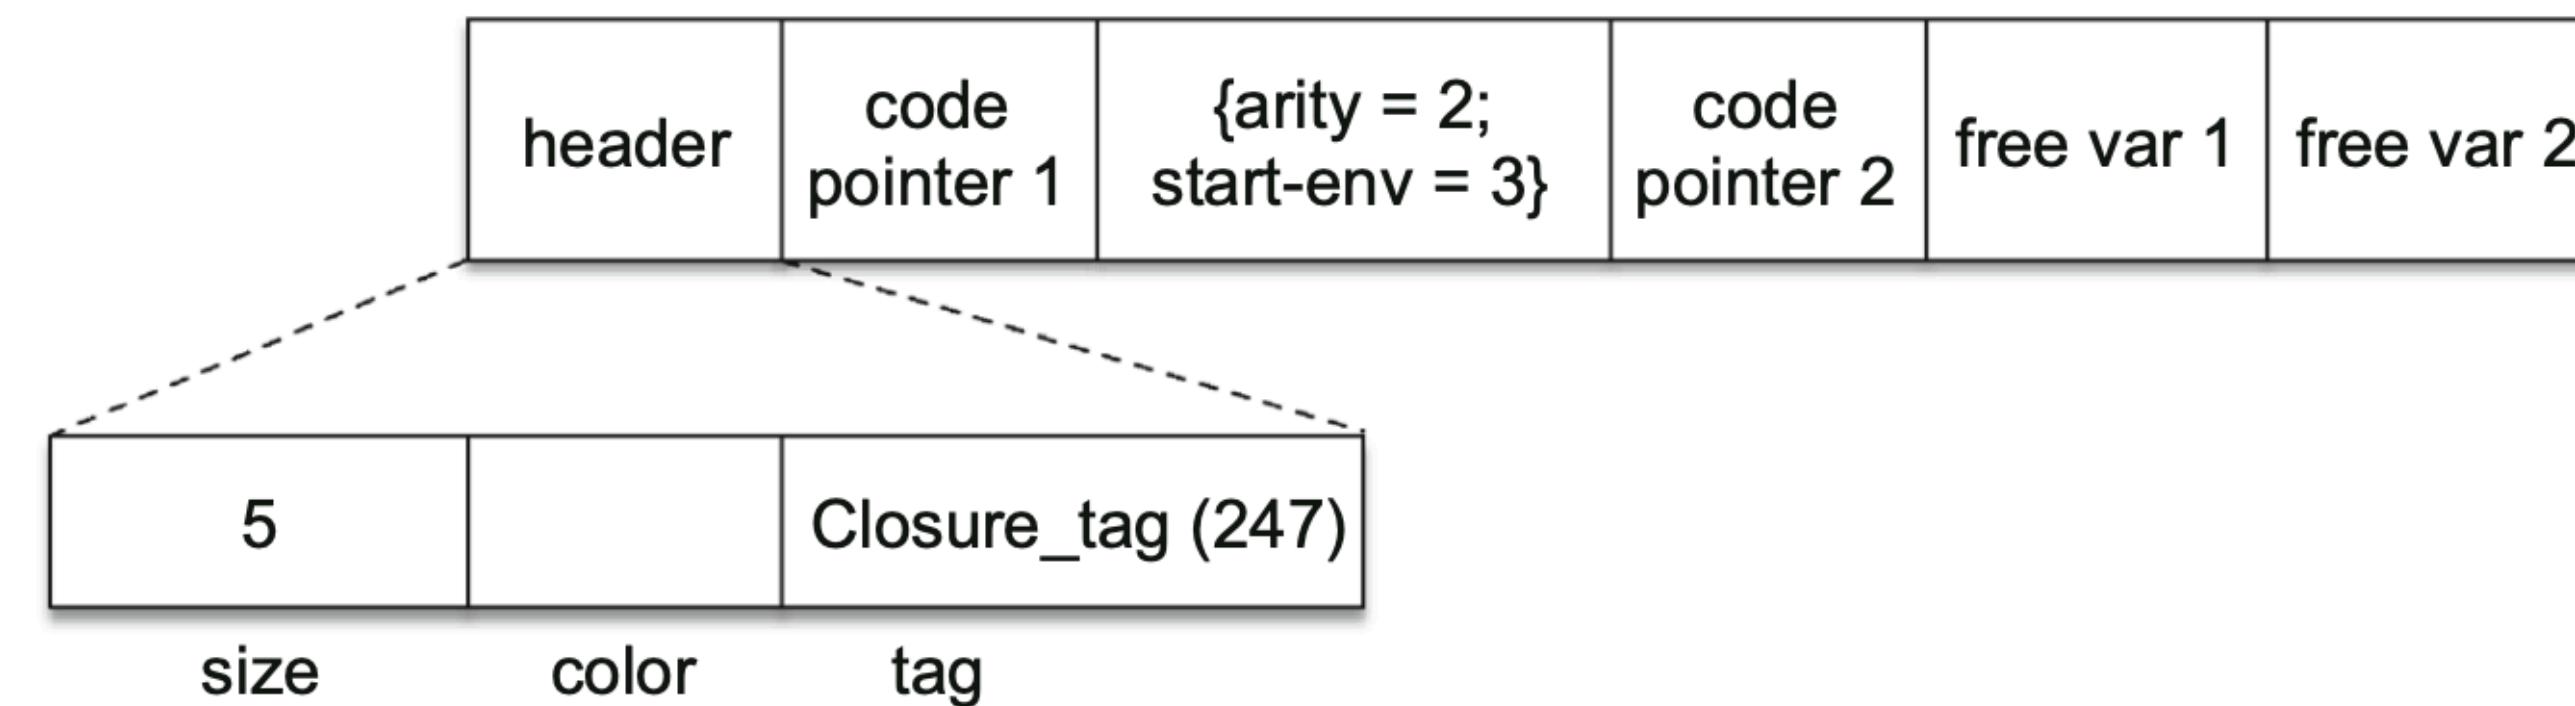
- A single contiguous buffer packed with objects
- Objects follow OCaml object layout



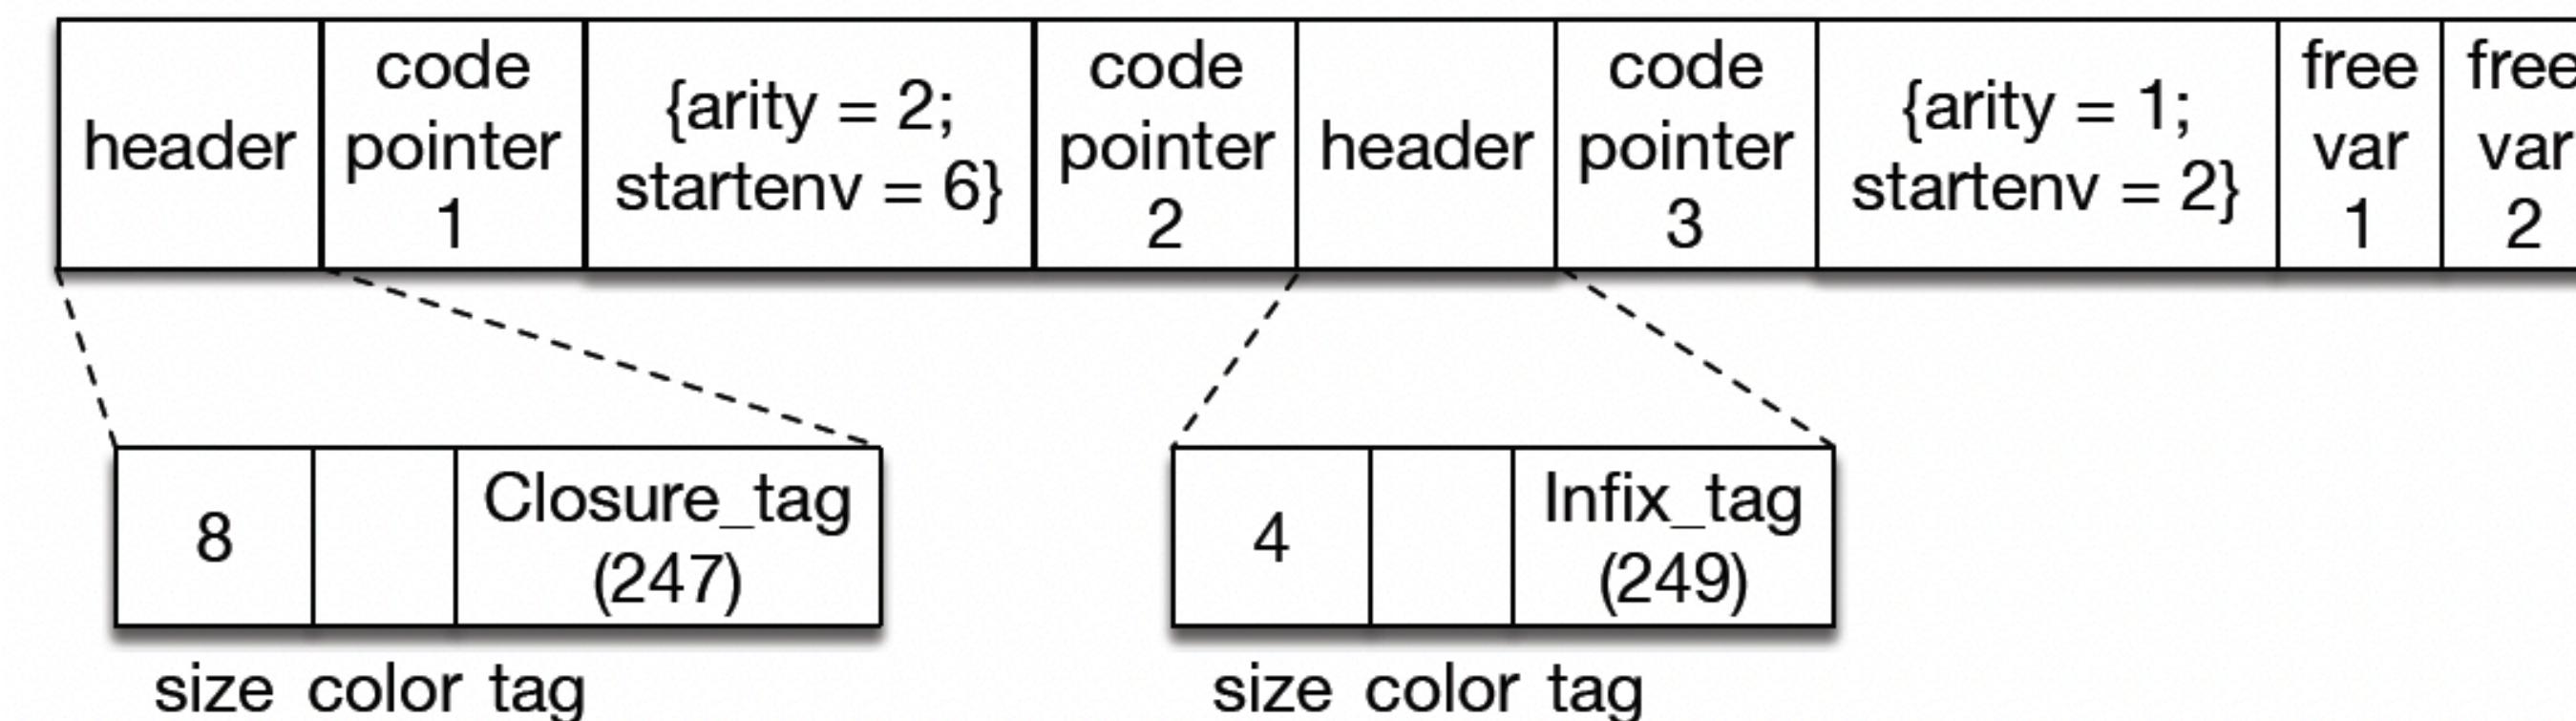
- Uniform value representation – Tagged ints
 - Makes it easy for GC to scan the object

Layer 2 – Heap

- Special Objects
 - NO SCAN objects – Strings, Float Arrays, Abstract, etc; contain no pointers.
 - Closure Objects



- Mutually-recursive closure objects



Layer 2 – Well-formed heap $\omega(h)$

- A well-formed heap satisfies a few properties.
 - Objects are non-overlapping
 - Objects are within the heap
 - Objects fill the heap
 - Pointers in the objects point to other allocated objects (non-blue)
 - Objects satisfy their layout requirements (think Closure and Infix objects).

Layer 2 – We’re still functional!

```
//Machine integers
module U64 = FStar.UInt64
module U8 = FStar.UInt8
let mword = 8UL
val heap_size : n:int{ n `mod` U64.v mword == 0 ∧ n >= 16 ∧
                        n < 1099511627776}
(* heap is a sequence of 8 bit unsigned machine integers *)
type heap = h:seq U8.t{length h == heap_size}
(* A valid heap address *)
type hp_addr = addr:U64.t {U64.v addr < heap_size ∧
                           is_multiple_of_mword addr}
```

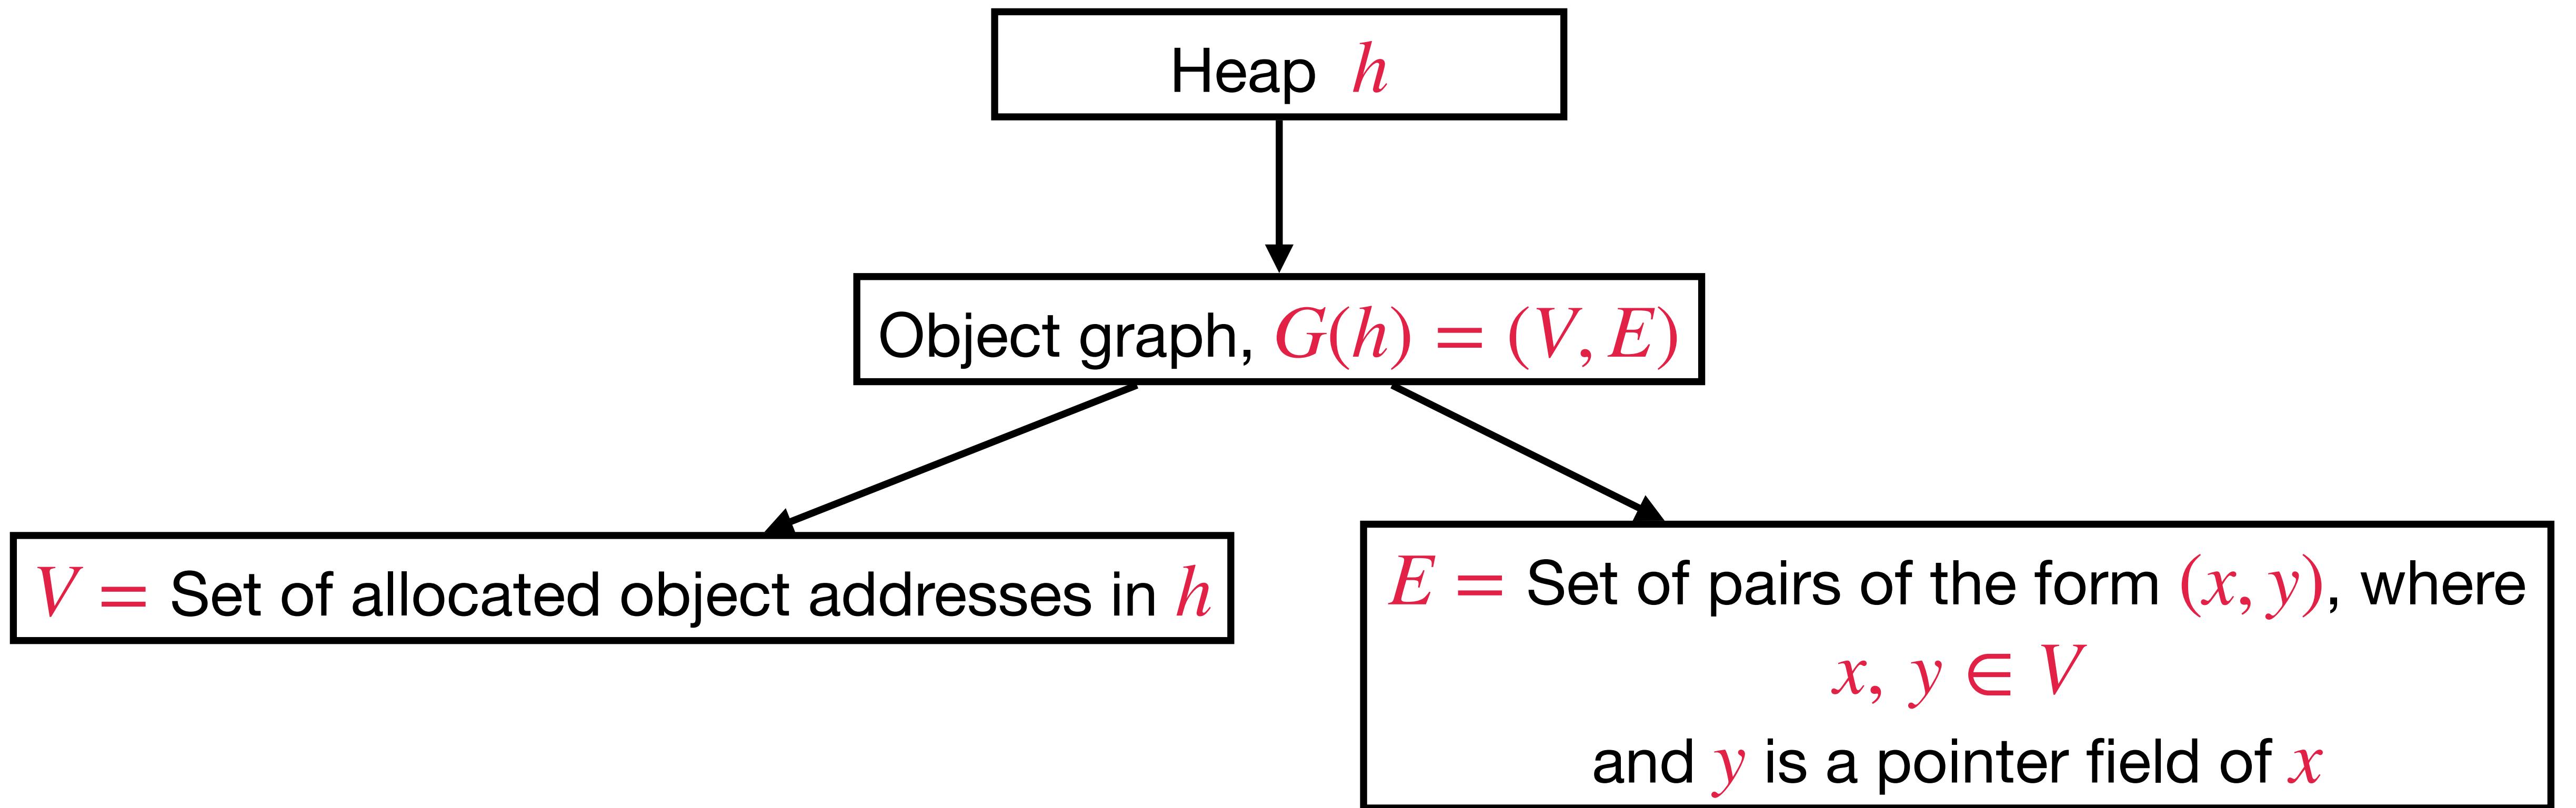
Layer 2 – DFS and Mark

```
(* dfs calls dfs_body until stack empty.  
Inputs are graph, stack and visited set. *)  
let rec dfs (g:graph) (st:seq U64.t) (vl:seq U64.t)  
: Pure (seq U64.t)  
(requires ...)  
(ensures (λ res → ...))  
(decreases (length g.vertices - length vl; length st)) =  
if length st = 0 then vl  
else  
  let st1,vl1 = dfs_body g st vl in  
    dfs g st1 vl1  
  
let dfs_body g st vl  
: Pure ...  
(requires ...)  
(ensures (λ res → ...)) =  
  let x = stack_top st in  
  let xs = stack_rest st in  
  let s = successors g x in  
  let vl1 = set_insert x vl in  
  let st1 = push_unvisited s xs vl1 in  
  (st1, vl1)  
  
(* mark calls mark_body until stack empty.  
Inputs are heap and stack *)  
let rec mark (h:heap) (st:seq obj_addr)  
: Pure (heap)  
(requires ...)  
(ensures (λ res → ...))  
(decreases (length allocs h - length blacks h;  
           length st)) =  
if length st = 0 then h  
else  
  let st1, h1 = mark_body h st in  
  mark h1 st1  
  
let mark_body (h:heap) (st:seq obj_addr)  
: Pure ...  
(requires ...)  
(ensures (λ res → ...)) =  
  let x = stack_top st in  
  let xs = stack_rest st in  
  let h1 = colorHeader h x black in  
  let st1 = darken h1 xs x 1UL in  
  (st1, h1)
```

Layer 2 – Mark specification

```
val mark (h:heap) (st:seq obj_addr)
  : Pure (heap)
  (requires (* Only core conditions shown *)
    (*1*) well_formed_heap (h) ∧
    (*2*) (forall x. x ∈ st ⇔ hd_address x ∈ greys(h)))
  (ensures (* Only core conditions shown *)
    (*1*) (λ h1 → well_formed_heap (h1) ∧
    (*2*) (forall x i. (hd_address x) ∈ h_objs(h) ⇒
                  field x h i = field x h1 i) ∧
    (*3*) (forall x. x ∈ h_objs(h1) ⇒ (color (hd_address x h)1 ≠ grey)))
```

Relating Layer 2 and Layer 1 – Graph(Heap)



Layer 2 – DFS ≡ Mark

```
val dfs_mark_equivalence_lemma (h:heap) (st:seq obj_addr)
                                (vl:seq obj_addr) (h_list:seq obj_addr)
  : Lemma
    (* Only important properties shown *)
    (requires (*1*) mutually_exclusive_sets st vl ∧
     (*2*) well_formed_heap(h) ∧
     (* stack invariant *)
     (*3*) (∀ x. x ∈ st ⇔ (hd_address x) ∈ greys(h))
     (* visited-list invariant *)
     (*4*) (∀ x. x ∈ vl ⇔ (hd_address x) ∈ blacks(h))

    (ensures (∀ x. x ∈ (dfs (graph_from_heap h) st vl) ⇔
              (hd_address x) ∈ blacks(mark h st))))
```

- Proof strategy
 - At every recursive call to dfs and mark, the stack **st** remains the same
 - Both terminate when the stack **st** is empty

Layer 2 – Reachability \equiv Mark

```
val mark_reachability_lemma (h:heap) (st:seq obj_addr)
                             (r_list:seq obj_addr)
  : Lemma
  (requires
   (*1*) well_formed_heap (h)
   (*2*) ( $\forall$  x. x  $\in$  st  $\iff$  hd_address x  $\in$  greys(h))  $\wedge$ 
   (*3*) well_formed_heap (mark h st))

  (ensures
   (*1*) (graph_from_heap (mark h st) = graph_from_heap h)  $\wedge$ 
   (*2*) ( $\forall$  x y.y  $\in$  r_list  $\wedge$ 
          reach (graph_from_heap h) y x  $\iff$  x  $\in$  blacks(mark h st)))
```

- Already proved Reachability \equiv DFS and DFS \equiv Mark
 - Hence, Reachability \equiv Mark

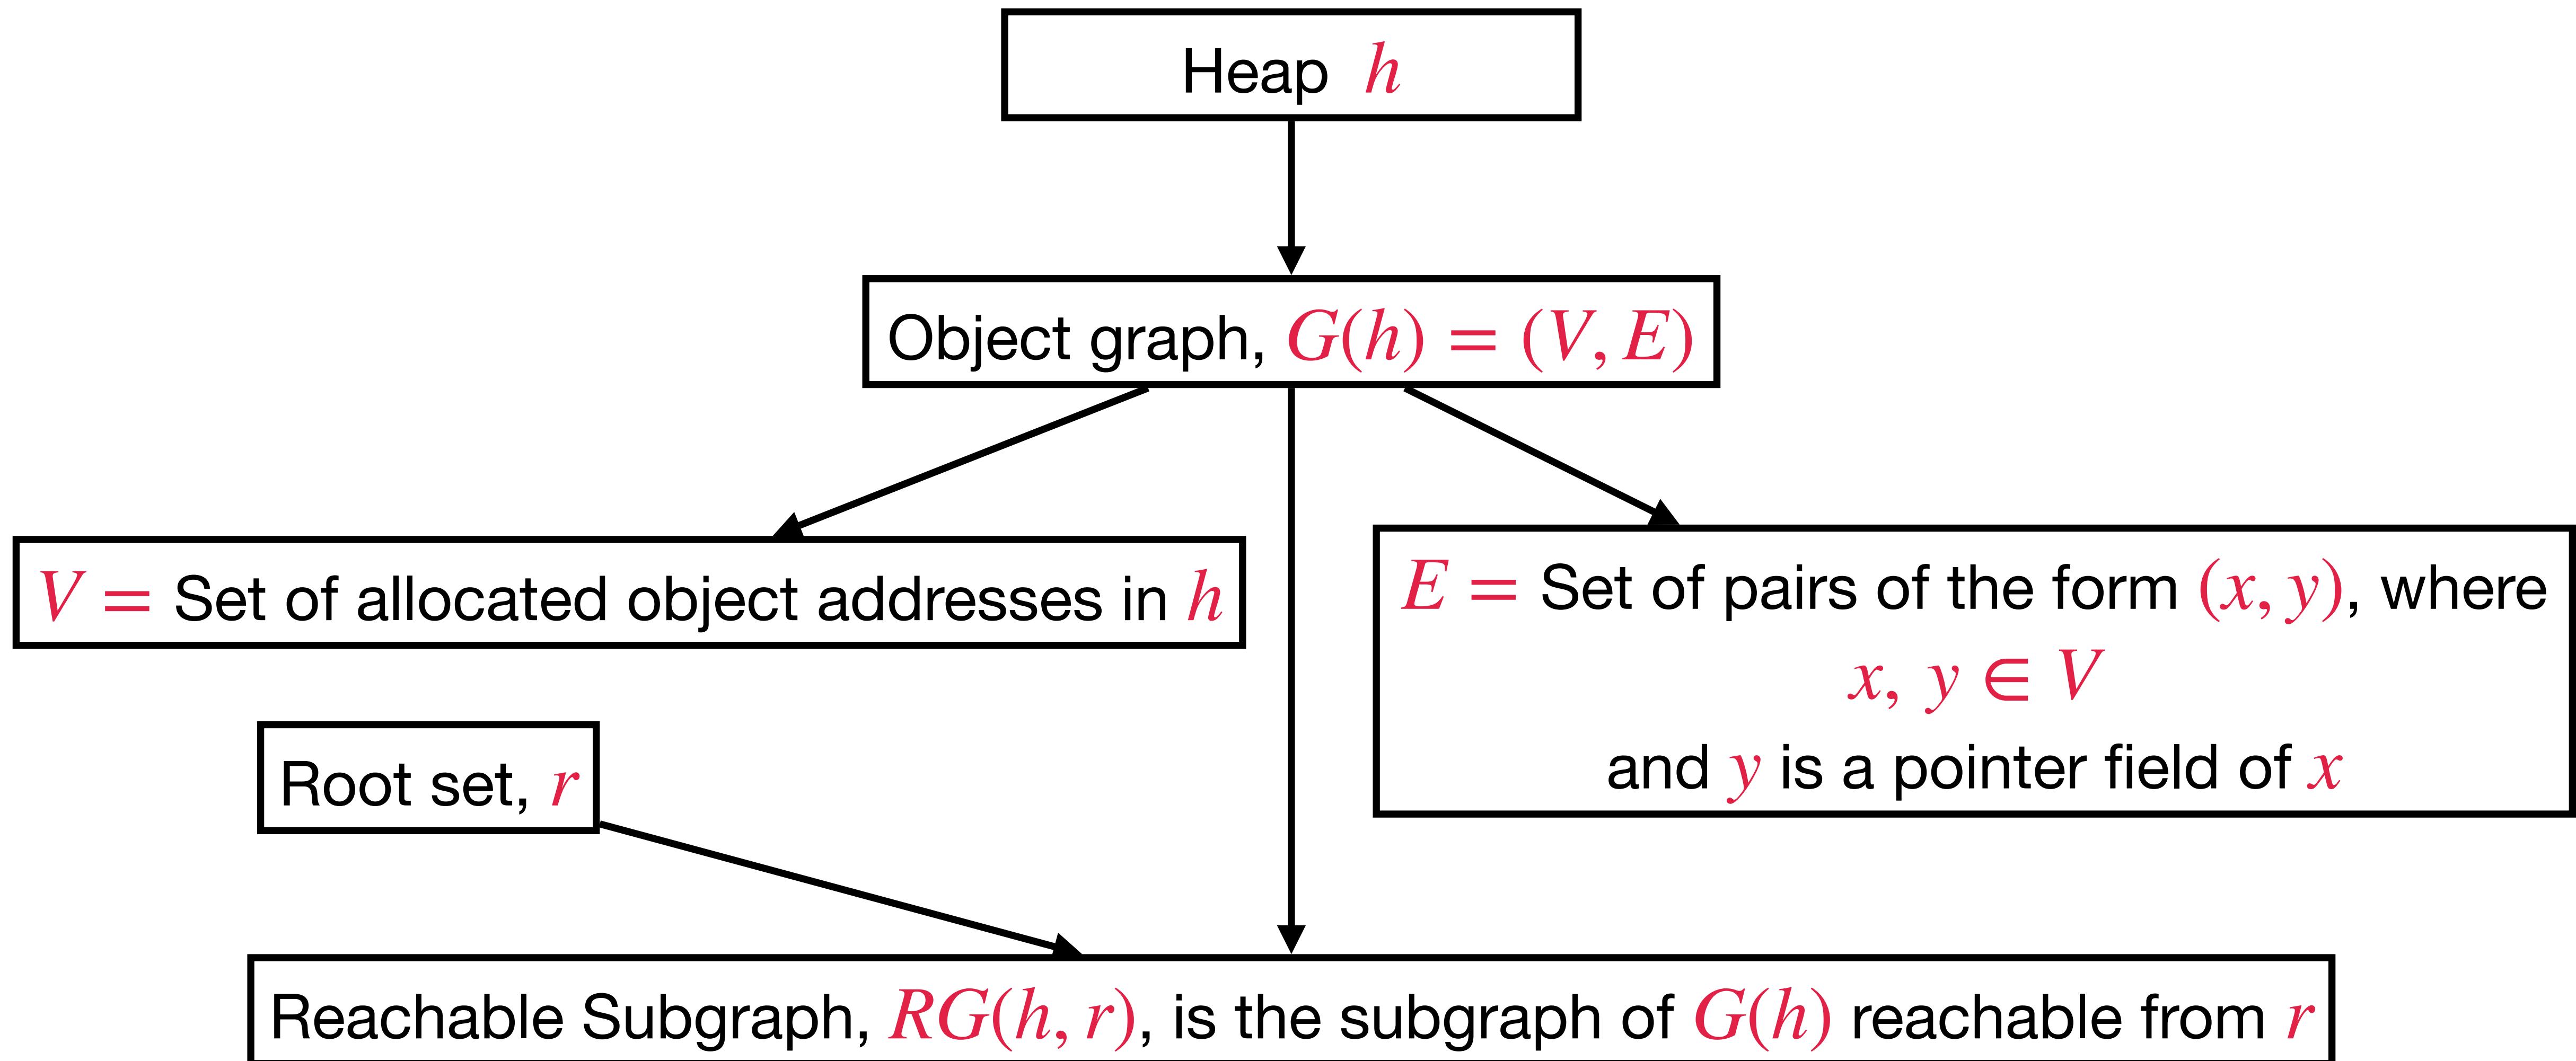
Layer 2 – Sweep specification

```
val sweep_subgraph_lemma (h:heap) (r_list:seq obj_addr)
                          (curr_ptr:obj_addr) (fp:obj_addr)
  : Lemma
  (requires
    (*1*) well_formed_heap (h) ∧
    (*2*) (forall x. x ∈ h_objs(h) ==> (color (hd_address x h) ≠ grey)) ∧
    (*3*) well_formed_heap (sweep h curr_ptr fp))

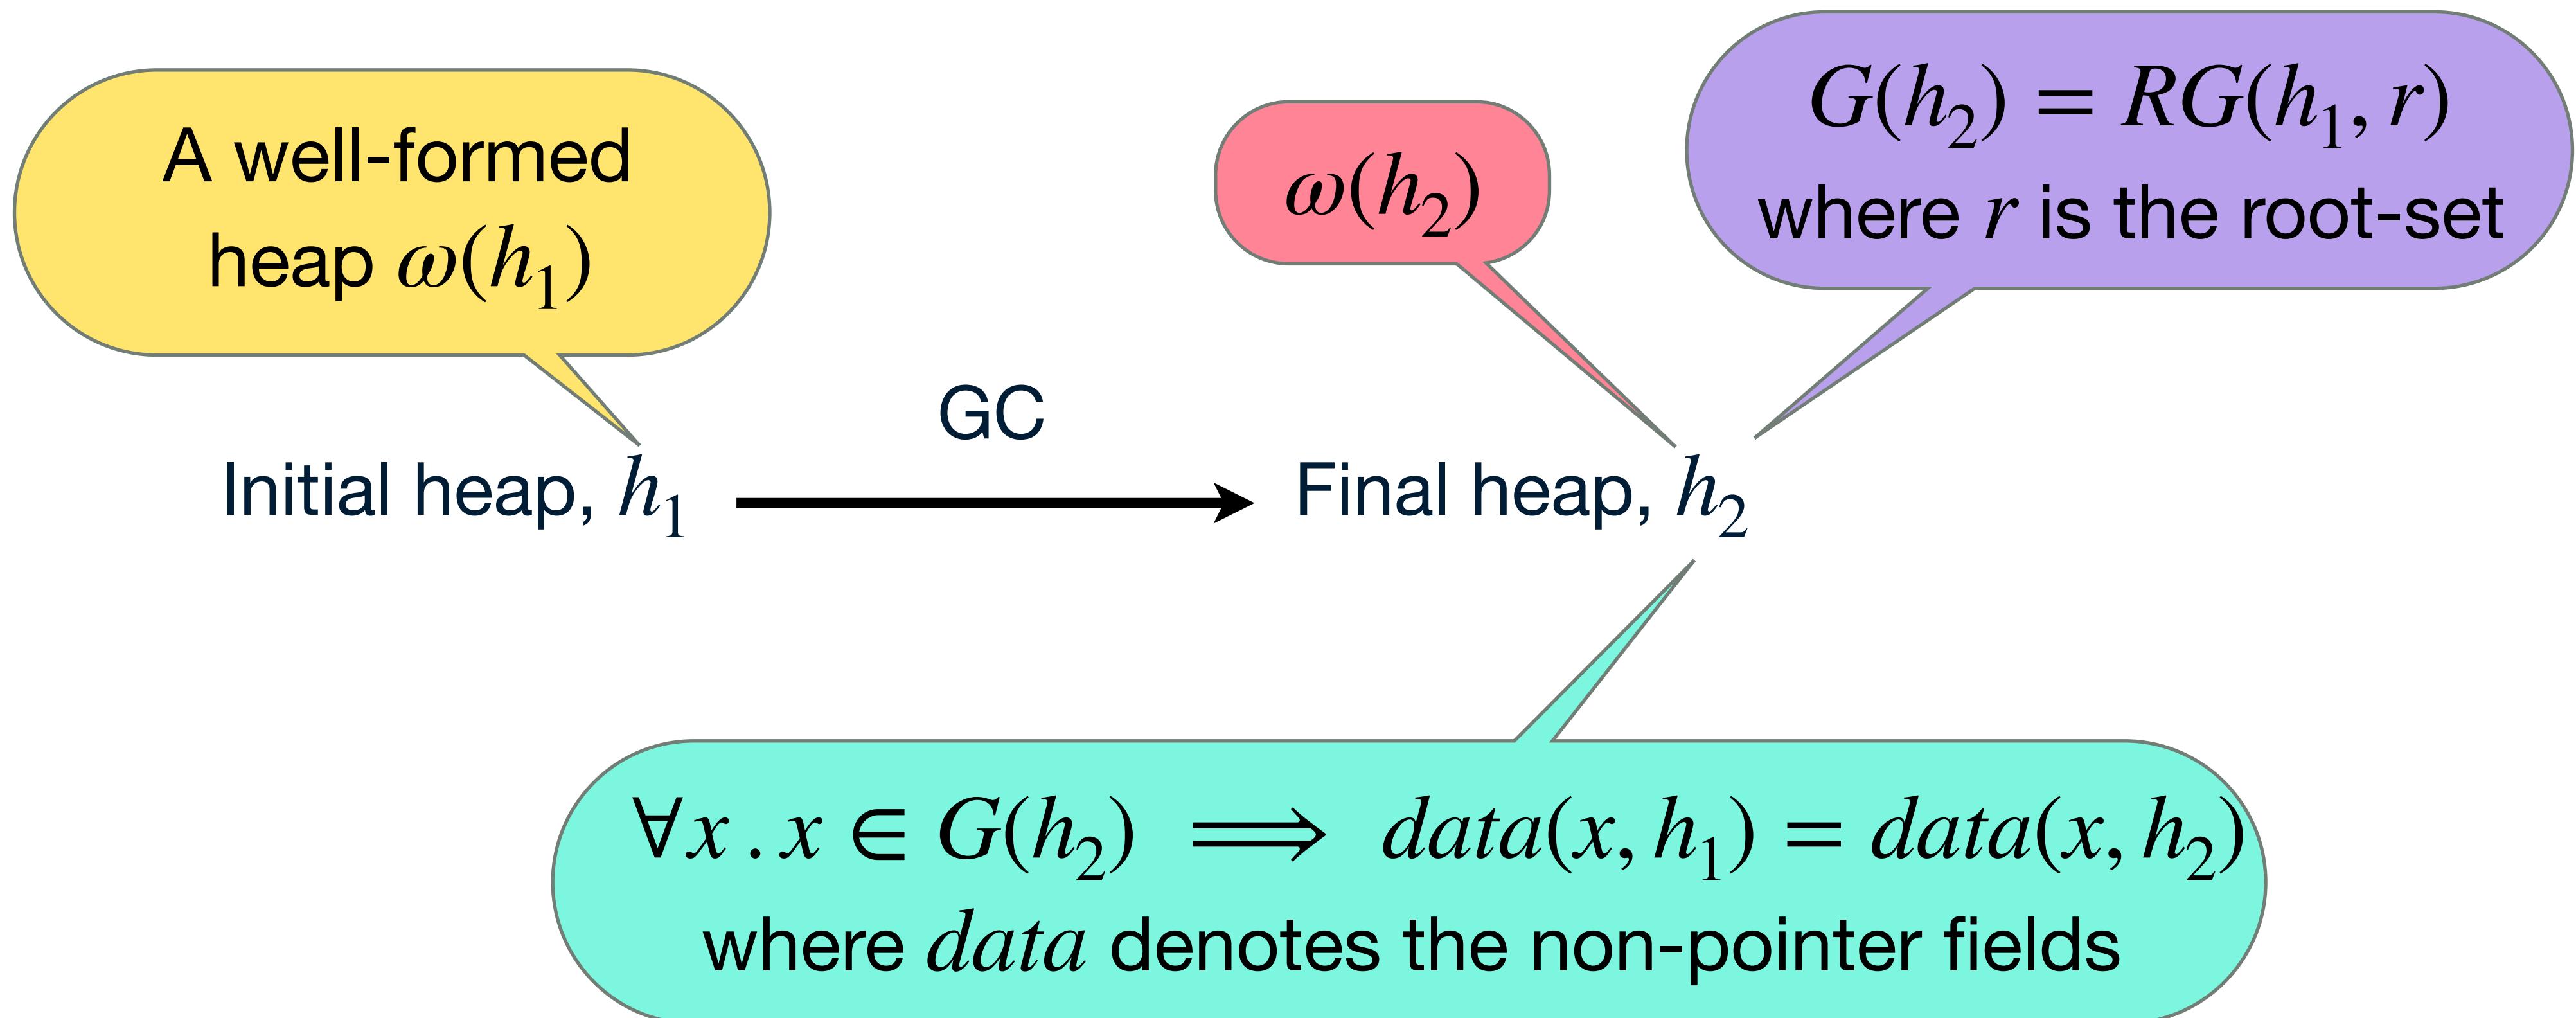
  (ensures
    (*1*) (forall x. x ∈ graph_from_heap (sweep h curr_ptr fp).vertices ==>
            x ∈ graph_from_heap (h).vertices ∧
            (exists y. y ∈ r_list ∧ reach (graph_from_heap h) y x)) ∧

    (*2*) (forall x y.
            (x,y) ∈ graph_from_heap (sweep h curr_ptr fp).edges ==>
            x ∈ graph_from_heap (sweep h curr_ptr fp).vertices ∧
            y ∈ graph_from_heap (sweep h curr_ptr fp).vertices ∧
            (x,y) ∈ graph_from_heap (h).edges))
```

Relating Layer 2 and Layer 1 – Graph(Heap)



GC Correctness



GC Correctness

```
val end_to_end_correctness_theorem
  (* Initial heap *)
  (h_init:heap{well_formed_heap h_init})
  (* mark stack - contains grey objects *)
  (st: seq Usize.t {pre_conditions_on_stack h_init st })
  (* root set *)
  (roots : seq Usize.t{pre_conditions_on_root_set h_init roots})
  (* free list pointer *)
  (fp:hp_addr{pre_conditions_on_free_list h_init fp})
: Lemma
( requires
  (* Pre-conditions elided for brevity. Important ones are:
     + The mark stack [st] contains all the [roots].
     + All the grey objects in the heap are in the mark stack [st].
   *)
  ( ensures
    (* heap after mark *)
    let h_mark = mark h_init st in
    (* heap after sweep *)
    let h_sweep = fst (sweep h_mark mword fp) in
    (* graph at init *)
    let g_init = graph_from_heap h_init in
    (* graph after sweep *)
    let g_sweep = graph_from_heap h_sweep in
      (* GC preserves well-formedness of the heap *)
      (* 1 *) well_formed_heap h_sweep ∧
      (* GC preserves reachable object set *)
      (* 2 *) ((∀ x. x ∈ g_sweep.vertices ⇔
                (∃ o. mem o roots ∧ reach g_init o x))) ∧
      (* GC preserves pointers between objects *)
      (* 3 *) ((∀ x. mem x (g_sweep.vertices) ⇒
                (successors g_init x) ==
                (successors g_sweep x))) ∧
      (* The resultant heap objects are either white or blue only *)
      (* 4 *) ((∀ x. mem x (h_objs h_sweep) ⇒
                color x h_sweep == white ∨
                color x h_sweep == blue) ∧
      (* No object field (either pointer or immediate) is modified *)
      (* 5 *) field_reads_equal h_init h_sweep )
```

Layer 3 – Imperative GC

- Layer 3 is in Low^* , a low-level explicit memory subset of F^*
 - Explicitly reason about lifetime and aliasing
 - Tedious (Low^* does not use separation logic) but mechanical
 - Low^* programs are memory-safe by construction
 - Can be extracted to C using KaRaMel
- Our Idea is to have 1:1 correspondence between the functions in layer 2
 - Tail-recursive functions \rightarrow loops
 - Invariants on tail-recursive functions \rightarrow loop invariants
 - Mark Stack is a list \rightarrow Mark stack is a fixed-sized buffer
 - Assumed to be the size of the heap (*limitation*)

Layer 3 – Extracted Sweep

```
void sweep (uint8_t *g, uint64_t *sw, uint64_t *fp,
            uint64_t limit, uint64_t mword) {
    while (*sw < limit) {
        uint64_t curr_obj_ptr = *sw;
        uint64_t curr_header = hd_address(curr_obj_ptr);
        uint64_t wz = wosize_of_block(curr_header, g);
        uint64_t next_header = curr_header + (wz + 1ULL) * mword;
        uint64_t next_obj_ptr = next_header + mword;
        sweep_body (g, sw, fp);
        sw[OU] = next_obj_ptr;
    }
}

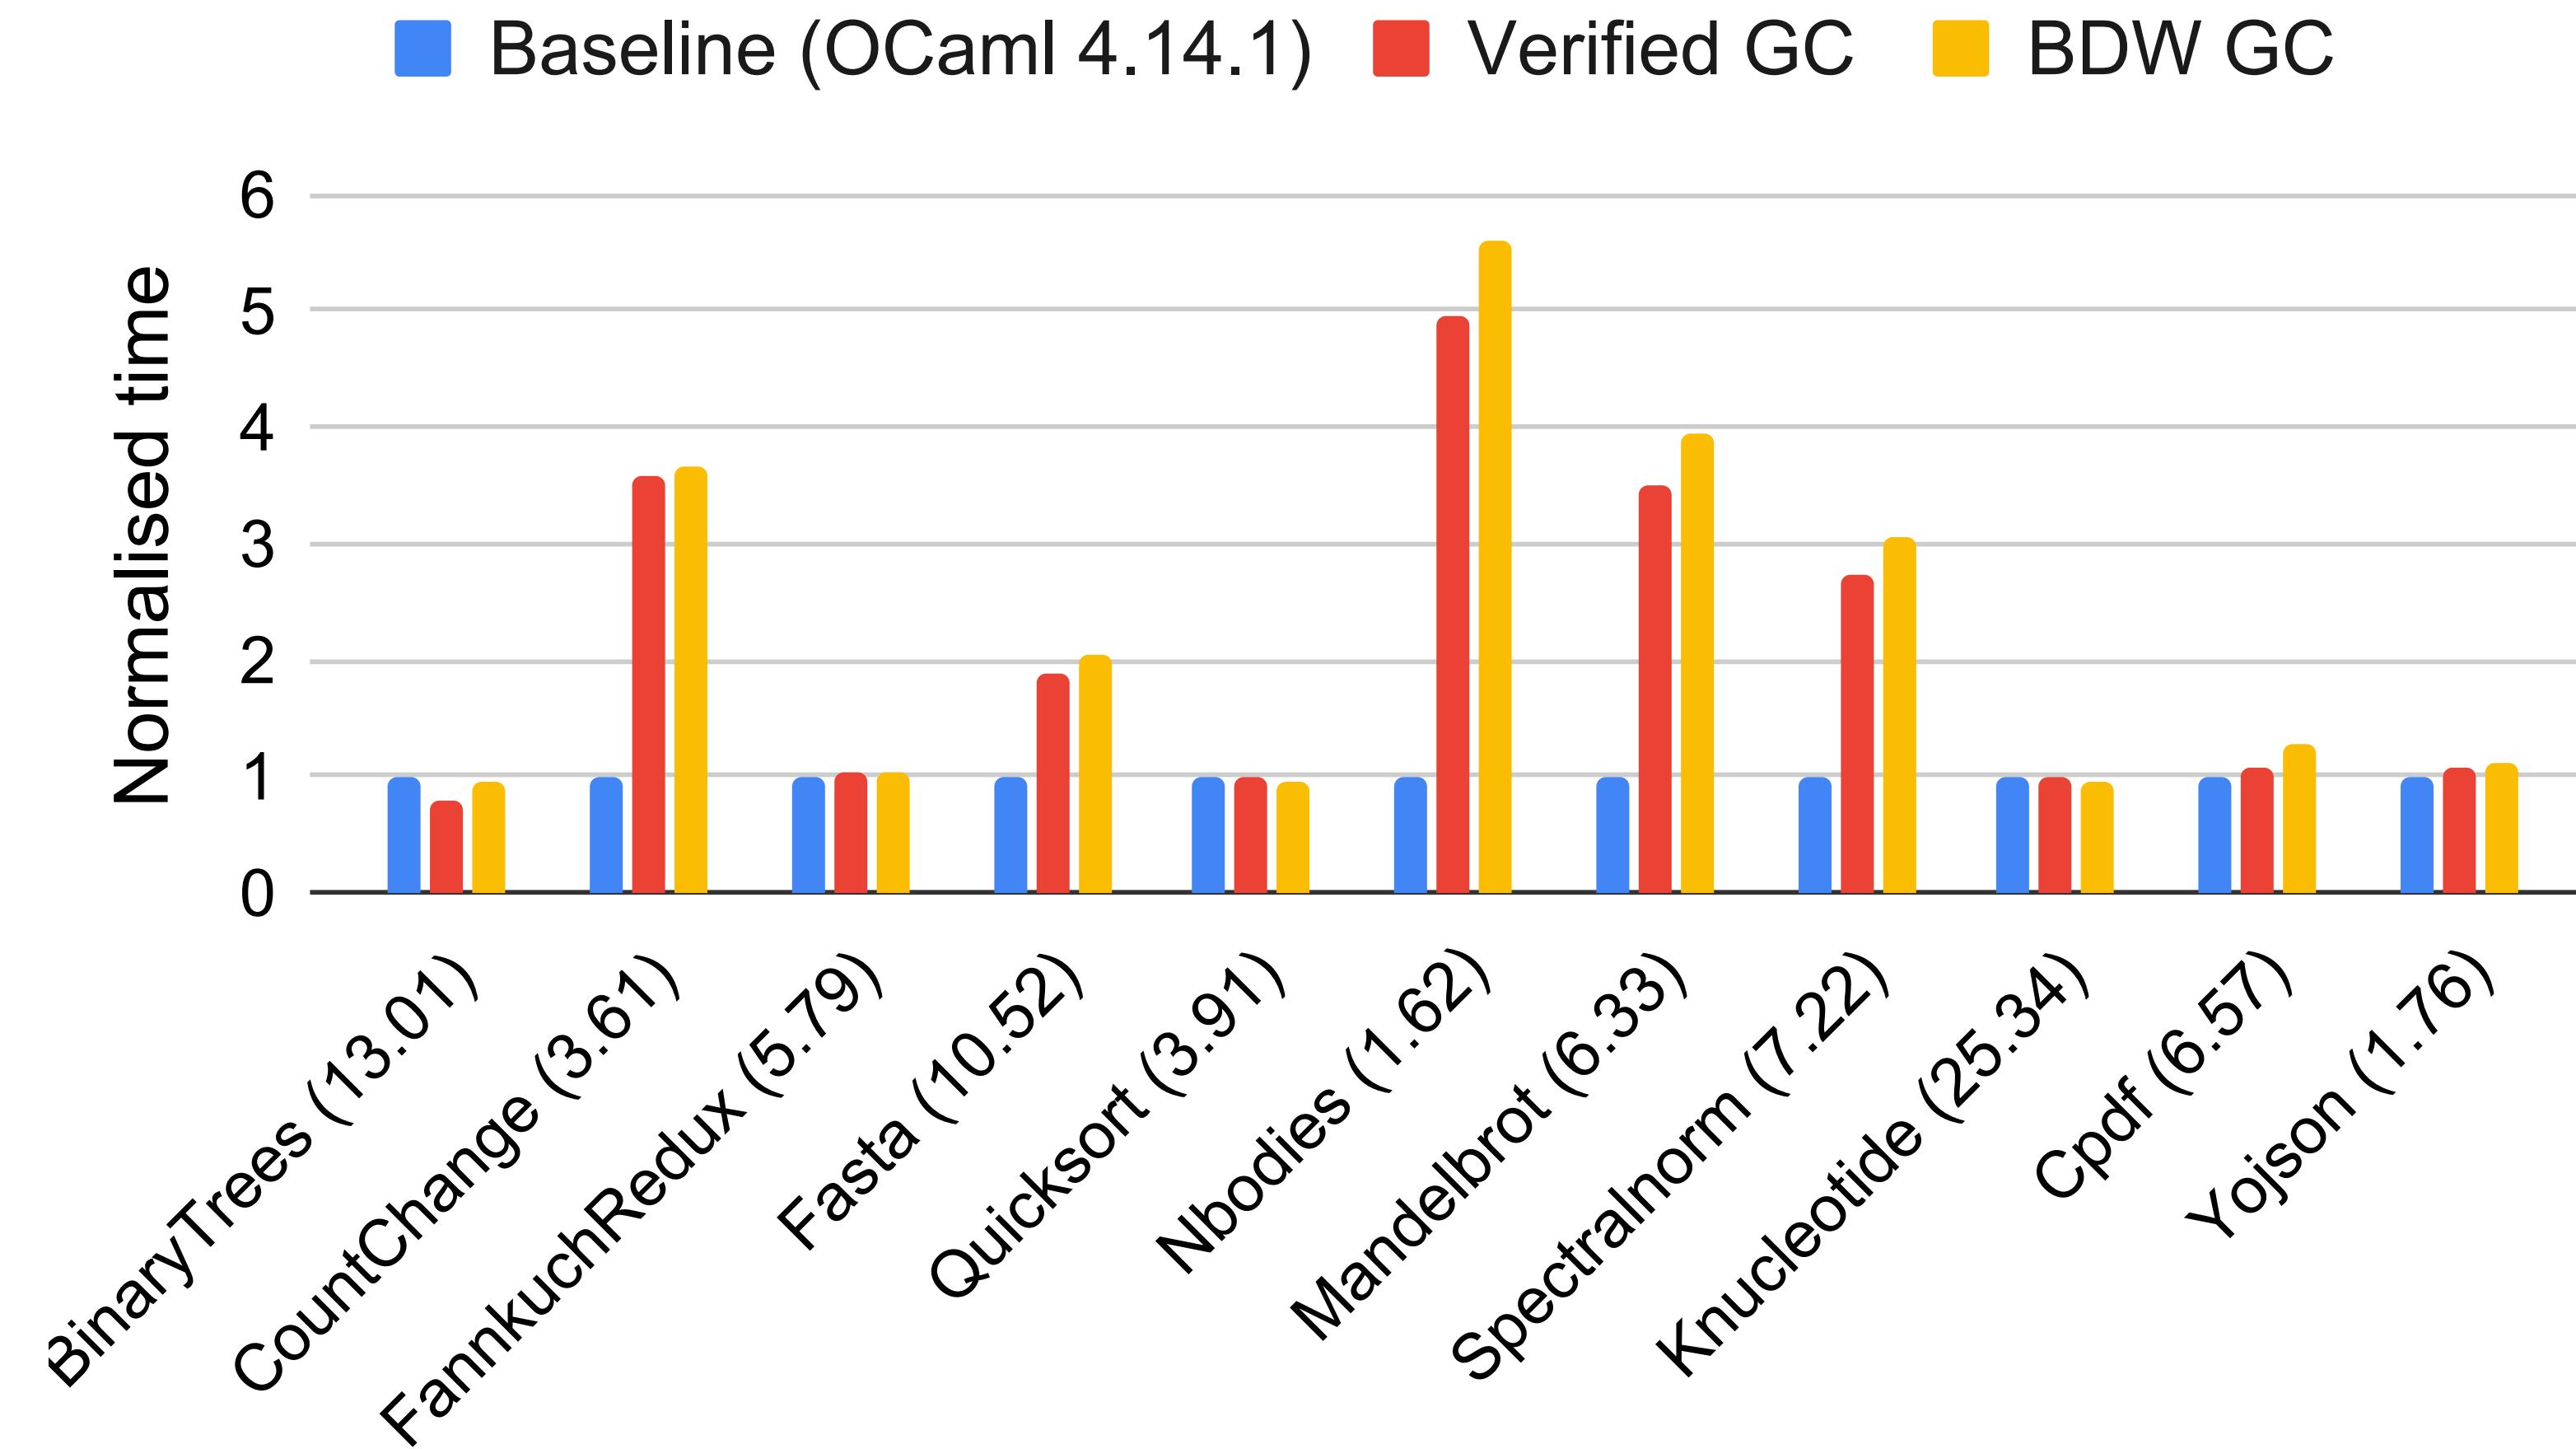
void sweep_body (uint8_t *g, uint64_t *sw, uint64_t *fp) {
    uint64_t curr_obj_ptr = *sw;
    uint64_t curr_header = hd_address(curr_obj_ptr);
    uint64_t c = color_of_block(curr_header, g);
    uint64_t wz = wosize_of_block(curr_header, g);

    if (c == white || c == blue) {
        colorHeader(g, curr_header, blue);
        uint64_t fp_val = *fp;
        uint32_t x1 = fp_val;
        store64_le(g + x1, curr_obj_ptr);
        fp[OU] = curr_obj_ptr;
    } else {
        colorHeader(g, curr_header, white);
    }
}
```

Verification effort

Modules	#Lines	#Defns	#Lemmas	Time
Graph	4653	72	81	2m3s
DFS	657	1	9	2m5s
Functional GC	18401	65	218	120m
Imperative GC	2734	19	19	27m43s

Performance



Extensions

- Extended to support coalescing during sweep
 - No two adjacent free/blue objects after sweep
 - Removes fragmentation
- In the paper,
 - Specification sketches for incremental and copying collectors
 - Not implemented fully

Limitations and Future Work

- Mark stack is assumed to be the size of the heap
 - Mark stack overflow handling is tricky!
- Allocator is still unverified
- No support for weak references, ephemerons, finalisers
 - Even specifying their correctness is challenging
- OCaml 5 is multicore but the verified GC is sequential only
 - Need a framework with concurrent separation logic (?)
- Although specifications are extensible, proof burden still very large
 - Transplant ideas to reusable GC framework such as MMTk (?)

Summary

- A verified stop-the-world ***mark-and-sweep*** GC for OCaml
- Written in F* and its subset Low*, proof-oriented programming languages
- Extracted to memory-safe C and integrated with OCaml 4.14.1 (non-multicore)
- Competitive performance with vanilla OCaml
- All the artifacts have been open-sourced

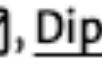
https://github.com/prismlab/verified_ocaml_gc

[Home](#) > [Journal of Automated Reasoning](#) > Article

A Mechanically Verified Garbage Collector for OCaml

[Open access](#) | Published: 14 May 2025
Volume 69, article number 11, (2025) [Cite this article](#)

[Download PDF](#)  You have full access to this [open access](#) article

[Sheera Shamsu](#) , [Dipesh Kafle](#), [Dhruv Maroo](#), [Kartik Nagar](#), [Karthikeyan Bhargavan](#) & [KC Sivaramakrishnan](#)

 997 Accesses [Explore all metrics](#) →

Abstract

The OCaml programming language finds application across diverse domains, including systems programming, web development, scientific computing, formal verification, and symbolic mathematics. OCaml is a memory-safe programming language that uses a

Journal of Automated Reasoning, 2025